

Coding the Impossible

Z80 Demoscene Techniques for Modern Makers

Alice Vinogradova

Alice Vinogradova — v25 (2026-03-27)

Contents

Chapter 1: Thinking in Cycles	2
T-States: The Currency of the Z80	2
Machine Cycles and Memory Access	3
The Frame: Your Canvas	4
Pentagon vs. Wait Machines	5
The Original Sinclair Machines	5
The Pentagon: Clean Timing	5
Thinking in Budgets	6
Practical: Setting Up Your Development Environment	7
What You Need	7
Project Structure	7
Build Configuration	8
Practical: The Timing Harness	8
Reading the Stripe	9
Variations	10
What Fits in a Frame?	10
Historical Note: Dark's Advice	11
The Computation Scheme	11
Summary	12
Try It Yourself	13
 Chapter 2: The Screen as a Puzzle	 14
The Memory Map: 6,912 Bytes of Screen	14
The Interleave: Where the Rows Live	16
Why?	17
The Bit Layout: Decoding (x, y) into an Address	17
The address calculation in Z80	18
DOWN_HL: Moving One Pixel Row Down	19
Introspec's Optimisation	20
Attribute Memory: 768 Bytes That Changed Everything	22
The Attribute Clash	23
The Border: More Than Decoration	24
Border Effects	24
Practical: The Checkerboard Fill	25
What to try	26
Navigating the Screen: A Practical Summary	26
Moving right one byte (8 pixels)	27
Moving down one pixel row	27
Moving down one character row (8 pixels)	27
Moving up one pixel row	27

Computing the attribute address from a pixel address	28
Putting It Together: What the Screen Layout Means for Code	30
Summary	30
Try It Yourself	31
Chapter 3: The Demoscener's Toolbox	32
Unrolled Loops and Self-Modifying Code	32
The cost of looping	32
Unrolling: trade RAM for speed	32
Self-modifying code: the Z80's secret weapon	33
Self-modifying variables: the \$+1 pattern	34
The Stack as a Data Pipe	35
Why PUSH is the fastest write on the Z80	35
The technique	35
POP as a fast read	36
Where PUSH tricks are used	37
LDI Chains	37
LDI vs LDIR	37
When LDI chains shine	38
Bit Tricks: SBC A,A and Friends	38
SBC A,A as a conditional mask	38
ADD A,A vs SLA A	39
Code Generation	39
Code generation: writing the program that draws	39
Offline: generating assembly from a higher-level language	39
Runtime: the program writes machine code during execution	40
The cost of generation	40
RET-Chaining	41
Turning the stack into a dispatch table	41
Three strategies for the render list	42
The trade-offs	42
Sidebar: "Code is Dead" (Introspec, 2015)	43
Putting It All Together	43
Try It Yourself	44
Chapter 4: The Maths You Actually Need	46
Multiplication on Z80	46
Method 1: Shift-and-Add from LSB	46
Method 2: Square Table Lookup	47
Signed Multiply	49
Two's Complement in Practice	50
Sign Extension: The rla / sbc a,a Idiom	50
mul_signed — Signed 8×8 Multiply	50
mul_signed_c — Thin Wrapper for Backface Culling	51
Cost Comparison	52
Going Further: The Fix-Uncigned Technique	52
Division on Z80	53
Method 1: Shift-and-Subtract (Restoring Division)	53
Method 2: Logarithmic Division	54
Sine and Cosine	55
The Parabolic Approximation	55
Bresenham's Line Drawing	56

The Classic Algorithm and Xopha Modification	56
Dark's Matrix Method: 8x8 Pixel Grids	56
Trap-Based Termination	57
Fixed-Point Arithmetic	57
Format 8.8	58
Fixed-Point Multiplication	58
Why Fixed-Point Matters	59
Theory and Practice	59
What Dark Got Right	60
Random Numbers: When Tables Won't Do	60
The R Register Trick	60
Four Generators from the Community	61
Elite's Tribonacci Approach	63
Elite's Galaxy Generator: A Deeper Look	64
Shaped Randomness	64
Seeds and Reproducibility	65
Comparison Table	65
Chapter 5: 3D on 3.5 MHz	68
The Problem: Twelve Multiplications Per Vertex	68
The Midpoint Method	68
The Cube as Foundation	69
Deriving Vertices by Averaging	69
Building Complex Objects	69
The Virtual Processor	70
Architecture	70
Execution	70
Writing Programs	71
Why a Virtual Processor?	72
Rotation	72
Z-Axis Rotation	72
Cost Per Basis Vertex	74
Projection	74
Parallel Projection	74
Perspective Projection	74
Solid Polygons	75
Backface Culling	75
Z-Sorting	77
Convex Polygon Filling	77
Putting It All Together	78
The Shape of Objects	79
Practical: A Spinning Solid Object	80
Step 1: Define the Object	80
Step 2: Write the Midpoint Program	80
Step 3: Define Faces	80
Step 4: The Frame Loop	80
Historical Context: From Magazine to Demo	82
Summary	82
Chapter 6: The Sphere — Texture Mapping on 3.5 MHz	84
The Problem: A Round Object on a Square Screen	84
The Insight: Code That Writes Code	85

Inside the Disassembly	85
Counting T-States	86
The Code Generation Pass	87
What Dark Knew: Spectrum Expert and the Building Blocks	88
The Hype Debate: Inner Loops vs. Mathematics	88
Practical: A Simplified 56x56 Spinning Sphere	89
Step 1: Precompute the Sphere Geometry	89
Step 2: Build Skip Tables	89
Step 3: Generate the Rendering Code	90
Step 4: Execute and Display	91
Step 5: Source Image Layout	91
The Larger Pattern	91
Summary	92
Chapter 7: Rotozoomer and Chunky Pixels	93
What a Rotozoomer Actually Does	93
Fixed-Point Stepping on the Z80	94
Chunky Pixels: Trading Resolution for Speed	94
Why 2x2 Is the Sweet Spot	95
The \$03 Encoding Trick	95
The Inner Loop from Illusion	96
Self-Modifying Code at the Byte Level	96
Per-Frame Code Generation	97
Buffer to Screen Transfer	98
Deep Dive: 4x4 Chunky Pixels (sq, Hype 2022)	98
Performance Comparison	99
Historical Roots: Born Dead #05 and the Scene Lineage	99
Practical: Building a Simple Rotozoomer	100
Texture Design and Boundary Handling	101
Why Page-Aligned, Why 256 Columns	101
Choosing Texture Size	101
What About Screen Boundaries?	102
The Design Space	102
Three Approaches to Texture Rotation	103
Variant 1: Monochrome Bitmap (Full Pixel Resolution)	103
Variant 2: Chunky Rotozoomer (2x2 or 4x4 Blocks)	104
Variant 3: Attribute Rotozoomer (8x8 Block “Pixels”)	104
Comparison	105
The Rotozoomer in Context	105
Summary	106
Chapter 8: Multicolor — Breaking the Attribute Grid	107
The ULA’s Perspective	108
The LDPUSH Insight	108
LD DE,nn / PUSH DE	109
How much fits in a scanline?	109
The stack pointer as a cursor	109
The GLUF Engine: Multicolor in a Real Game	110
Double buffering	111
The two-frame architecture	111
What the player sees	111
Ringo: A Different Kind of Multicolor	112

The 11110000b pattern	112
Two-screen switching	112
Where the T-states go	113
Horizontal scrolling	113
Traditional Multicolor: The Interrupt-Driven Approach	114
Sidebar: Black Crow #05 — Early Multicolor	115
The Spectrum's Palette and Multicolor	115
Practical: A Multicolor Game Screen	116
Step 1: The display code buffer	116
Step 2: Attribute changes within the display code	116
Step 3: Tile rendering via operand patching	117
Step 4: Sprite overlay	117
Step 5: The main loop	118
What DenisGrachev's Work Means	119
Summary	119
Try It Yourself	119
Chapter 9: Attribute Tunnels and Chaos Zoomers	121
The Attribute Grid as a Framebuffer	121
Plasma: The Colour Engine	122
Four-Fold Symmetry: Divide and Conquer	123
The Chaos Zoomer	124
Code Generation: Processing Writes Z80	125
The Zapilator Question	125
A Taste of the Scripting Engine	126
The Making Of: Timeline and Inspiration	127
Practical: Building a Simplified Attribute Tunnel	127
Step 1: Fill Pixel Memory with a Pattern	127
Step 2: Sine Table	128
Step 3: Plasma for One Quarter	128
Step 4: Four-Way Copy	129
Step 5: Main Loop	129
Key Insight	130
Sources	130
Chapter 10: The Dotfield Scroller and 4-Phase Colour	131
Part 1: The Dotfield Scroller	131
What the Viewer Sees	131
The Font as Texture	131
Stack-Based Address Tables	132
The Inner Loop	133
The Budget Arithmetic	134
How the Bounce Is Encoded	135
Beyond Simple Sine: Lissajous, Helix, and Multi-Wave Patterns	135
Part 2: 4-Phase Colour Animation	136
The Colour Problem	136
The Trick	136
The Maths of Perception	136
Why Inversion Is Essential	137
Practical Cost	138
Text Overlay	138
Demoscene Lineage	139

The Shared Principle: Temporal Cheating	140
Practical 1: A Bouncing Dot-Matrix Text Scroller	140
Practical 2: A 4-Phase Colour Cycling Animation	141
Summary	141
Try It Yourself	142
Chapter 11: Sound Architecture - AY, TurboSound, and Triple AY	143
11.1 The AY-3-8910 Register Map	143
Complete Register Table	144
Tone Channels (R0-R5): How Pitch Works	144
The Noise Generator (R6)	146
R7: The Mixer - The Most Important Register	146
Volume Registers (R8-R10)	147
Envelope Generator (R11-R13)	147
11.2 Chiptune Techniques on 3 Channels	148
Arpeggios: Fake Chords	148
Buzz-Bass: The Envelope Trick	150
The Period Alignment Problem	150
Natural Tuning: Table #5	151
Drum Synthesis	154
Ornaments: Per-Frame Modulation	155
11.3 TurboSound: 2 x AY	155
Chip Selection	156
6 Channels, True Stereo	156
Engine Modification	156
11.4 Triple AY on ZX Spectrum Next	157
Enhanced Features	157
9 Channels: Orchestral Thinking	157
11.5 Music Engine Architecture	158
The Interrupt-Driven Player	158
The Player Loop	159
Frame Budget	159
Formats and Trackers	160
11.6 Sound Effects System	160
Channel Stealing	160
Procedural SFX Tables	162
11.7 Putting It Together: The Working Example	163
Sidebar: Beeper - A Brief History of Impossibility	163
Sidebar: Agon Light 2 - VDP Sound System	164
11.8 Practical Exercises	165
Summary	165
Chapter 12: Digital Drums and Music Sync	166
12.1 Digital Drums on the AY	166
The Problem: The AY Cannot Play Samples	166
The Cost: CPU Annihilation	167
n1k-o's Insight: The Hybrid Drum	167
The Frame Budget: Two Frames Per Hit	169
12.2 Asynchronous Frame Generation	170
The Naive Approach Fails	170
Introspec's Solution: Stockpile Frames	170
The Buffer Dynamics	171

12.3 The Scripting Engine	172
Why You Need a Script	172
Outer Script: The Sequence of Effects	172
Inner Script: Variations Within an Effect	172
kWORK: The Key Command	172
12.4 GABBA's Innovation: The Video Editor as Timeline Tool	174
The Problem with Code-Based Sync	174
diver4d's Answer: Luma Fusion	174
What This Changes	176
12.5 Z80 Threading: A Different Path	176
The Problem, Restated	176
IM2-Based Context Switching	176
How It Works in Practice	177
The Threading Model	177
Practical Considerations	178
12.6 Practical: A Minimal Scripted Demo Engine	178
What We Build	178
The Memory Map	178
The Timeline Script	179
The Main Engine Loop	179
The Display ISR	181
The Effect Generator Dispatch	182
Observations	182
12.7 Practical Exercises	183
Summary	183
Chapter 13: The Craft of Size-Coding	185
13.1 What is Size-Coding?	185
The Z80 Size-Coder's Toolkit	186
13.2 Anatomy of a 256-Byte Intro: NHBF	188
The Music	188
The Visual	188
The Puzzle: Finding Overlaps	188
Art-Top's Discovery	188
The Byte Budget	189
Key Techniques at 256 Bytes	190
13.3 Famous 256-Byte Intros: What Made Them Clever	190
13.4 The LPRINT Trick	191
How It Works	191
The Transposition Effect	192
From Pirate Loaders to Demo Art	192
Why It Matters for Size-Coding	192
13.5 512-Byte Intros: Room to Breathe	192
What Each Size Tier Enables	192
Common 512-Byte Patterns	193
Self-Modifying Tricks	193
The ORG Trick	194
13.6 4K Intros: The Mini-Demo	194
Compression Becomes Viable	194
Music Fits	195
Multi-Effect Structure	195

GOA4K, inal, and Megademica	196
The 4K Tradeoffs	196
Competition Categories	197
13.7 Practical: Writing a 256-Byte Intro Step by Step	197
Step 1: The Unoptimised Version	197
Step 2: Replace CALL with RST	197
Step 3: Overlap Data with Code	198
Step 4: Exploit Register State	198
Step 5: Smaller Encodings Everywhere	198
Step 6: Counting Bytes Precisely	198
The Final Push	199
13.8 Size-Coding Music: Bytebeat on AY	199
The Minimal AY Formula Engine	200
Techniques for Better-Sounding Formulas	200
Bytebeat vs. Sequenced Music	201
13.9 Procedural Graphics: The Rendered GFX Compo	201
Why the Spectrum Is Interesting for This	201
Common Approaches	202
The Byte Budget for Art	202
13.10 Size-Coding as Art	202
Summary	203
Try It Yourself	204

Chapter 14: Compression — More Data in Less Space 205

The Memory Problem	205
Compression as bandwidth amplifier	206
The Benchmark	206
The corpus	206
The results	206
The tradeoff triangle	208
How LZ Compression Works	209
ZX0 — The Size Coder's Choice	210
Why ZX0 exists	210
The decompressor	210
When to use ZX0	211
RLE and Delta Coding	211
RLE: Run-Length Encoding	211
Delta coding: store what changed	213
Pre-Compression Data Preparation	214
Entropy: the theoretical floor	214
The second derivative: sinusoidal and quadratic data	214
Transposition: column-major for tabular data	215
Interleaving planes: masks and pixels	215
Pattern detection: when not to compress	216
Practical transforms for common demo data	216
The Practical Pipeline	216
From asset to binary	217
Makefile integration	217
Example: loading screen with ZX0	217
Choosing the right compressor	218
The MegaLZ Revival	218

What the Numbers Mean in Practice	218
Summary: The Compressor Cheat Sheet	219
Try It Yourself	219
Chapter 15: Anatomy of Two Machines	221
15.1 ZX Spectrum 128K: Memory Map	221
The \$7FFD Port: Bank Switching	222
A Practical Memory Map for a 128K Game	222
15.2 Contended Memory: The Practical Truth	223
What Gets Contended	223
How Much Slower?	225
The Practical Response	225
15.3 ULA Timing	226
Frame Structure	226
Tact-Maps: Frame Regions	226
Scanline Timing	227
Total vs Practical Budget	227
15.4 Floating Bus, ULA Snow, and the \$7FFD Bug	228
Floating Bus	228
The \$7FFD Read Bug	228
ULA Snow	228
15.5 Clone Differences	229
Pentagon 128	229
Scorpion ZS-256	229
ZX Spectrum Next	230
15.6 Agon Light 2: A Different Beast	231
Dual-Processor Architecture	231
eZ80 Memory Model: 24-Bit Flat	232
ADL Mode vs Z80 Mode	232
MOS API: The Operating System	233
VDP Commands: Talking to the Screen	234
15.7 Comparing the Platforms	234
15.8 Practical: Memory Inspector Utility	235
Spectrum Version	235
Agon Version	238
Summary	239
Chapter 16: Fast Sprites	241
Method 1: XOR Sprites	241
The simplest approach	241
When XOR works	242
When XOR fails	243
Method 2: OR+AND Masked Sprites	243
The industry standard	243
Data format	243
The inner loop	244
Cycle counting	245
Byte alignment and the shift problem	245
Method 3: Pre-Shifted Sprites	246
The memory-for-speed trade-off	246
Memory calculation	246
Practical compromise	247

Method 4: Stack Sprites (The PUSH Method)	247
The fastest output on the Z80	247
The inner loop	248
Making it work: pre-calculated SP chain	248
The costs	249
When to use stack sprites	249
Method 5: Compiled Sprites	250
The sprite is the code	250
How it works	250
16x16 compiled sprite	251
The trade-offs	252
Compiled sprites with masking	252
Method 6: Dirty Rectangles	253
The background problem	253
The save/restore cycle	253
Save/restore routine	254
Full frame budget	255
Draw order and overlap	255
Optimising the Inner Loops	255
Eliminating pointer management	255
Agon Light 2: Hardware VDP Sprites	258
Moving a sprite	258
Scanline limits	259
The trade-off	259
Practical: 8 Animated Sprites at 25 fps	259
Spectrum implementation	259
Agon implementation	261
Summary	261
Try It Yourself	262
Chapter 17: Scrolling	264
The Budget	264
Vertical Pixel Scrolling	265
The Interleave Problem	265
The Algorithm: Scroll Up by One Pixel	265
Cost Analysis	266
Partial Scrolling: The Practical Approach	266
Scrolling by 8 Pixels (One Character Row)	267
Horizontal Pixel Scrolling	267
Why Horizontal Scrolling Is Expensive	267
The Full Budget Calculation	268
Can We Do Better?	269
Attribute (Character) Scrolling	269
Scrolling Attributes with LDIR	269
The Combined Method: Character Scroll + Pixel Offset	270
How It Works	271
The Simple Combined Method	271
Implementation: The Edge-Column Pixel Shift	272
The Rendering Pipeline	272
Scrolling the Pixel Data by One Character Column	273
The Shadow Screen Trick	274

Shadow Screen Scrolling Strategy	275
Comparison: Scrolling Methods on ZX Spectrum	276
Scrolling Right (and the Direction Problem)	276
Agon Light 2: Hardware Scrolling	278
Hardware Scroll Offsets	278
Tilemap Scrolling	278
Ring-Buffer Column Loading	279
Comparison: Spectrum vs. Agon Scrolling	279
Practical: Horizontal Side-Scrolling Level	280
Spectrum Version: Combined Character + Pixel Scroll	280
Agon Version: Hardware Tilemap Scrolling	282
Vertical + Horizontal: Combined Scrolling	282
Optimisation Tips	282
1. Use a Screen Address Lookup Table	282
2. Scroll Only What Is Visible	283
3. Use PUSH for the Character Scroll	283
4. Split the Character Scroll Across Frames	283
5. Palette and Attribute Tricks	283
Summary	283

Chapter 18: Game Loop and Entity System	285
18.1 The Main Loop	285
The Frame Budget, Revisited	286
18.2 The Game State Machine	287
State Definitions	287
The Jump Table	287
The Dispatcher	289
Why Not a Chain of Comparisons?	291
State Transitions	291
18.3 Input: Reading the Player	292
ZX Spectrum Keyboard	292
Kempston Joystick	293
Edge Detection: Press vs Hold	294
Agon Light 2: PS/2 Keyboard via MOS	294
18.4 The Entity Structure	295
Structure Layout	295
Why 16 Bytes?	296
Both Axes in 8.8 Fixed-Point	297
The 8.8 Fixed-Point System	297
18.5 The Entity Array	298
Entity Slot Allocation	298
Iterating Entities	298
Update Dispatch by Type	299
The Player Update Handler	300
18.6 The Object Pool	300
Spawning a Bullet	301
Deactivating an Entity	302
Bullet Update Handler	302
Pool Sizing	303
18.7 Explosion and Effect Entities	303
18.8 Putting It All Together: The Game Skeleton	304

18.9 Agon Light 2: The Same Architecture, More Room	316
18.10 Design Decisions and Trade-Offs	316
Fixed vs Variable Frame Rate	316
Entity Size: 16 Bytes vs Larger	317
When to Use HL Instead of IX	317
Summary	318
Try It Yourself	318
Chapter 19: Collisions, Physics, and Enemy AI	320
Part 1: Collision Detection	320
AABB: The Only Shape You Need	320
Ordering the Tests for Fastest Rejection	323
Tile Collisions: The Tilemap as a Collision Surface	324
Sliding Collision Response	327
Part 2: Physics	327
Gravity: Falling Convincingly	327
Why Fixed-Point Matters Here	328
Jump: The Anti-Gravity Impulse	328
Variable-Height Jumps	329
Friction: Slowing Down on the Ground	330
Applying Velocity to Position	331
The Physics Loop	331
Part 3: Enemy AI	332
The Finite State Machine	332
The JP Table	332
Patrol: The Dumb Walk	333
Chase: The Relentless Follower	335
Attack: Fire and Cooldown	336
Retreat: The Reverse Chase	337
Death: Animate and Remove	338
Optimisation: Update AI Every 2nd or 3rd Frame	339
Part 4: Practical – Four Enemy Types	340
Wiring It Into the Game Loop	342
Agon Light 2 Notes	343
Tuning Guide	343
Summary	344
Try It Yourself	345
Chapter 20: Demo Workflow — From Idea to Compo	346
20.1 What “Design” Means in a Demo	346
20.2 Lo-Fi Motion: A Complete Case Study	347
The Concept: “Belarusian Pixel”	347
The Scene Table	347
Fourteen Effects	348
The Effects Themselves	349
The Toolchain	349
The Build Pipeline	349
The Timeline: Two Weeks of Evenings	350
20.3 Making-of Culture	351
Eager: The Technical NFO	351
GABBA: A Different Workflow	351
NHBF: The Puzzle	352

20.4 The Toolchain in Detail	352
Assembler: sjasmplus	352
Emulators	353
Graphics and Code Generation	353
Build Automation and CI	353
Synchronisation and Compositing	353
20.5 Compo Culture	355
The Major Parties	356
How to Enter Your First Compo	356
20.6 The Community	357
Hype (hype.retroscene.org)	357
ZXArt (zxart.ee)	357
Pouet (pouet.net)	358
20.7 Project Management for Demo Makers	358
The Minimum Viable Demo	358
Working with Collaborators	358
Debugging and Testing	359
20.8 Transcending the Platform: Introspec's "MORE"	359
20.9 Your First Demo: A Practical Roadmap	360
Week 1: Foundation	361
Week 2: Effects	361
Week 3: Polish	361
Week 4: Release	362
Summary	362
Chapter 21: Full Game - ZX Spectrum 128K	364
21.1 Project Architecture	364
Directory Layout	365
The Build System	365
21.2 Memory Map: 128K Bank Assignments	366
Ironclaw Bank Allocation	367
The Stack	370
21.3 The State Machine	370
21.4 The Gameplay Frame	371
Reading Input	372
The Scroll Engine	373
21.5 Sprite Integration	374
Background Restore (Dirty Rectangles)	375
21.6 Collisions, Physics, and AI in Context	376
Physics-Collision Loop	376
Tile Collision	377
Enemy AI	378
21.7 Sound Integration	379
Music	379
Sound Effects	380
21.8 Loading: Tape and DivMMC	381
The .tap File and BASIC Loader	381
esxDOS Loading (DivMMC)	382
21.9 Loading Screen, Menu, and High Scores	385
Loading Screen	385
Title Screen and Menu	385

High Scores	387
21.10 Level Loading and Decompression	388
21.11 Profiling with DeZog	390
What is DeZog?	390
The Profiling Workflow	390
DeZog Configuration for Ironclaw	392
21.12 The Data Pipeline in Detail	393
Tilesets (PNG to Spectrum pixel format)	393
Sprite Sheets (PNG to pre-shifted sprite data)	393
Level Maps (Tiled JSON to binary tilemap)	394
Music (Vortex Tracker II to PT3)	394
Putting It Together	394
21.13 Release Format: Building the .tap	394
Testing the Release	396
21.14 Final Polish	396
Summary	397
Chapter 22: Porting — Agon Light 2	399
Same ISA, Different World	399
The Architecture at a Glance	400
ADL Mode vs Z80-Compatible Mode	401
Why You Might Use Z80 Mode	402
The Mode Switching Mechanism	402
The Practical Rule	403
The MBASE Trap	403
What Transfers Directly	403
Game Logic and Entity System	403
AABB Collision Detection	404
Fixed-Point Arithmetic	405
State Machine	405
What Needs Rewriting	406
Rendering: From Framebuffer to VDP Commands	406
What Needs Rethinking	409
Memory Architecture	409
Loading	409
Sound	410
Input	411
eZ80 at 18 MHz: What Still Matters, What Does Not	412
What Still Matters: Inner Loop Efficiency	412
What Becomes Irrelevant: Memory Conservation Tricks	413
What Becomes Irrelevant: Self-Modifying Code	413
What Becomes Irrelevant: Stack Tricks for Rendering	414
The Comparison Table	414
The Porting Process: Step by Step	415
Step 1: Set Up the Agon Project	415
Step 2: Replace the Rendering Layer	415
Step 3: Translate the Game Logic	416
Step 4: Rewrite Sound	416
Step 5: Rewrite Input	416
Step 6: Rewrite Loading	416
Step 7: Test and Tune	416

What Each Platform Forces You To Do Better	416
A Note on eZ80 Instructions	416
Summary	417
Chapter 23: AI-Assisted Z80 Development	419
23.1 The Historical Parallel: HiSoft C on ZX Spectrum	419
23.2 The Claude Code Feedback Loop	420
The Loop	420
A Concrete Example	421
What Makes the Loop Fast	422
What Makes the Loop Slow	423
23.3 DeZog Integration: The Other Half of the Loop	423
The AI + DeZog Workflow	423
Memory Inspection for Data-Heavy Code	424
What DeZog Cannot Do (Yet)	424
23.4 When AI Helps, When It Does Not	424
AI Helps: High Confidence	424
AI Helps: Medium Confidence	425
AI Does Not Help: Low Confidence	425
23.5 Case Study: Building MinZ	426
What MinZ Is	426
Where AI Helped Build MinZ	426
Where AI Did Not Help Build MinZ	428
The MinZ Verdict	428
23.5b Sidebar: The Other AI — Brute-Force Superoptimisation	429
23.6 Honest Take: “Z80 They Still Don’t Know”	430
What the AI Gets Wrong, Specifically	430
Where Introspec Is Right	431
Where Introspec Is Not Quite Right	431
23.7 The “Antique Toy” Demo: AI in Practice	432
23.8 The Feedback Loop in Practice	432
23.9 Building Your Own AI-Assisted Workflow	433
Prompt Engineering for Z80	433
23.10 The Broader Picture	434
Summary	434
Try It Yourself	435
Glossary	436
A. Timing & Performance	436
B. Hardware — Sinclair & Clones	438
C. Hardware — Soviet/Post-Soviet Ecosystem	441
D. Hardware — Agon Light 2	442
E. Techniques	443
F. Assembly Notation & Directives	447
G. Demoscene & Culture	449
Key People	451
Key Demos & Productions	452
Key Publications	452
H. Algorithms & Compression	453
Appendix A: Z80 Instruction Quick Reference	456
8-Bit Load Instructions	456

16-Bit Load Instructions	457
8-Bit Arithmetic and Logic	457
16-Bit Arithmetic	459
Rotate and Shift	459
Bit Manipulation	461
Jump, Call, Return	461
I/O Instructions	462
Block Instructions	463
Exchange and Misc	464
The Demoscene “Fast” Instructions	466
Fastest Register-to-Register Move	466
Fastest Way to Zero a Register	466
Fastest Memory Read	466
Fastest Memory Write	466
Fastest 2-Byte Write	466
Fastest 2-Byte Read	466
Fastest Block Copy	467
Fastest I/O	467
Fastest Pointer Swap	467
Fastest Conditional Loop	467
Fastest Indirect Jump	467
Undocumented Instructions	467
IXH, IXL, IYH, IYL (Half-Index Registers)	467
SLL r (Shift Left Logical)	468
OUT (C),0	468
CB-Prefix Undocumented Bit Operations on (IX+d)	468
Flag Effects Cheat Sheet	469
Instructions That Set All Arithmetic Flags (S, Z, H, P/V, N, C)	469
Instructions That Set Z and S (but NOT Carry)	469
Instructions That Set ONLY Carry-Related Flags	469
Instructions That Set NO Flags	469
Practical Tricks	469
Register Architecture	470
Main Register Set	470
Special Registers	470
Shadow Registers	470
Register Pairing for Instructions	470
Common Instruction Sequences	471
Pixel Address Calculation (Screen Address from Y,X)	471
DOWN_HL: Move One Pixel Row Down	472
8x8 Unsigned Multiply (Shift-and-Add)	472
AY Register Write	472
16-Bit Compare (HL vs DE)	473
Stack-Based Screen Fill	473
Fast Pixel-Row Iteration (Split Counters)	473
Quick Cost Comparisons	474
Instruction Encoding Size Reference	475
Appendix B: Sine Table Generation and Trigonometric Tables	476
The Standard Format	476
Comparison of Approaches	477

Approach 1: Full 256-Byte Table	477
Approach 2: Quarter-Wave Table	478
Approach 3: Parabolic Approximation (Dark's Method)	479
Approach 4: Second-Order Delta Encoding (The Deep Trick)	479
The Key Insight	479
Full Table via 2-Bit Second-Order Deltas	480
Quarter-Wave via 2-Bit Second-Order Deltas	480
Approach 5: Bhaskara I Approximation (7th Century)	481
The Formula	482
Accuracy	482
Z80 Implementation	482
Bhaskara I + Correction Bitmap (Exact)	483
When to Use Bhaskara I	483
Practical Recommendations	483
The Decision Tree	484
What Does Not Work	484
Raider's Commandments	484
Reference: The Full 256-Byte Table	485
Appendix C: Compression Quick Reference	487
Compressor Comparison Table	487
Decision Tree: Which Compressor?	491
Compressibility of Common ZX Spectrum Data Types	491
Pre-compression tricks	494
Minimal RLE Decompressor	494
ZX0 Standard Decompressor (Z80)	495
Integration Patterns	497
Pattern 1: Decompress to screen at startup	497
Pattern 2: Decompress to buffer between effects	497
Pattern 3: Stream decompression during playback	498
Pattern 4: Bank-switched compressed data (128K)	498
Build Pipeline: From Asset to Binary	498
Quick Formulas	499
See Also	500
Appendix D: Development Environment Setup	501
1. The Assembler: sjasmplus	501
Installing from Source	501
Version Pinning	502
Key Flags	502
File Extension	502
Hex Notation	502
2. The Editor: VS Code	503
Essential Extensions	503
Build Task	503
Recommended Settings	504
3. Emulators	504
ZEsarUX – Feature-Rich Debugging	504
Fuse – Lightweight and Accurate	505
Unreal Speccy – Windows, Pentagon-Focused	505
For Agon Light 2	505
Which Emulator Should I Use?	506

4. The Debugger: DeZog	506
Installation	506
Connecting to ZEsarUX	506
Using the Internal Simulator	507
Key DeZog Features	507
5. Building the Book's Examples	508
Clone the Repository	508
Prerequisites	508
Build Commands	508
Running an Example	509
6. Project Structure for Your Own Code	509
Minimal Makefile	510
Include Convention	510
7. Alternative Tools	511
Other Assemblers	511
Other VS Code Extensions	511
CSpect - Next-Focused Emulator	511
SpectrumAnalyzer	512
8. Troubleshooting	512
"sjasmpus: command not found"	512
Compilation errors in the book's examples	512
DeZog cannot connect to ZEsarUX	512
Emulator shows garbled screen	512
Build output is empty or zero bytes	512
Tool Reference	512
See Also	513

Appendix E: eZ80 Quick Reference 514

1. Architecture Overview	514
2. Mode System	515
The Two Modes	515
Mode Suffixes	515
Mode-Switching Calls and Jumps	515
The Practical Rule	516
The MBASE Trap	516
3. New Instructions	516
Arithmetic and Test	516
Address Computation	516
I/O and System	517
4. MLT — The Game Changer	517
What MLT Enables	518
MLT Limitations	518
5. Key Differences Summary	518
6. Agon Light 2 Specifics	519
Hardware	519
Frame Budget	519
MOS API	520
VDP Command Protocol	520
7. Porting Checklist	521
See Also	521

Appendix F: Z80 Variants — Extended Instruction Sets 523

1. The Z80 Family Tree	523
2. Z80N — The Demoscener’s Wish List	524
Screen Navigation (the DOWN_HL problem)	524
Sprite Rendering (the masked blit problem)	525
Arithmetic (the multiply problem)	526
Bit Manipulation	526
Barrel Shifts (the multi-bit shift problem)	526
Convenience Instructions	527
The Big Picture	528
3. eZ80 — The Enterprise Extension	528
4. R800 — The MSX turboR Speedster	529
5. Soviet Clones — Behind the Iron Curtain	529
6. Comparison Table	530
Effective Multiply Performance	531
7. What This Means for the Book	531
See Also	532

Appendix G: AY-3-8910 / TurboSound / Triple AY Register Reference 533

I/O Ports on ZX Spectrum 128K	533
The Write Sequence	533
Reading a Register	534
Bulk Register Write	534
Complete Register Map	534
Overview	534
R0-R1: Channel A Tone Period (12-bit)	536
R2-R3: Channel B Tone Period (12-bit)	536
R4-R5: Channel C Tone Period (12-bit)	536
R6: Noise Period (5-bit)	536
R7: Mixer Control (The Most Important Register)	537
R8-R10: Volume Registers (5-bit)	538
R11-R12: Envelope Period (16-bit)	538
R13: Envelope Shape (4-bit, Write-Only)	539
Tone Period to Frequency Conversion	541
Formula	541
AY Clock by Platform	541
Complete Note Frequency Table	542
Compact Note Table for Z80 Code	545
Table #5: Natural Tuning (Just Intonation)	545
TurboSound: 2 x AY	547
Chip Selection	547
TurboSound Engine Architecture	548
Stereo Configuration	548
ZX Spectrum Next: Triple AY (3 x AY)	549
Chip Selection on Next	549
Per-Channel Stereo Panning	549
Common Patterns in Z80 Code	550
Silence the AY	550
Play a Single Note on Channel A	550
Trigger a Buzz-Bass Note	550
Trigger Envelope (Restart)	551
Digital Drum Hit (4-bit DAC)	551

Tracker Format Notes	552
ProTracker 3 (.pt3)	552
Other Tracker Formats	553
Vortex Tracker II	553
AY-3-8910 vs YM2149: Differences That Matter	554
Quick Reference Card	554
Register Summary (Tear-Out)	554
Appendix H: Storage APIs — TR-DOS and esxDOS	556
1. TR-DOS (Beta Disk 128)	556
Hardware	556
Disk Format	556
WD1793 Port Map	557
WD1793 Commands	557
ROM API: Loading a File	558
Direct Sector Access	558
Disk Detection	559
2. esxDOS (DivMMC / DivIDE)	560
Hardware	560
API Pattern	560
Function Reference	560
File Open Modes	561
Seek Whence Values	561
Code Example: Load a File	561
Code Example: Streaming Data from File	562
Detecting esxDOS	563
3. +3DOS (Amstrad +3)	564
4. Practical Patterns	564
Loading Screen from Disk (TR-DOS)	564
Loading Screen from SD (esxDOS)	565
Dual-Mode Loader	565
Streaming Compressed Data	566
5. File Format Reference	566
6. See Also	567
Appendix I: Bytebeat and AY-Beat - Generative Sound on Z80	568
1. Classic Bytebeat: The PCM Tradition	568
Famous Formulas	568
Why It Works	569
On the Spectrum: The Beeper Dead End	569
2. AY-Beat: Bytebeat Reimagined for a Tone Generator	570
Basic AY-Beat Architecture	570
What Changes from PCM Bytebeat	571
3. Drone: Envelope + Tone (E+T Mode)	571
The Drone Recipe	572
Byte Cost	573
4. Noise Percussion	573
Basic Kick Pattern	573
Percussion Character by Noise Period	574
Rhythmic Variety from Bit Masks	574
Using the Envelope for Drum Decay	575
5. Multi-Channel Harmony	575

Three Voices from One Formula	575
Why Rotation Creates Harmony	576
Volume Formulas for Multi-Channel	576
6. Formula Cookbook	577
Tone Period Formulas	577
Volume Formulas	578
How to Read the Table	579
7. Putting It Together: A Complete AY-Beat Engine	579
What This Produces	581
Customisation Points	581
8. Advanced: Combining Techniques	582
Architecture for a 256-Byte Intro	582
Register Budget	582
Size Breakdown	582
9. AY-as-DAC: Classic Bytebeat Through the Volume Register	583
Three Output Paths Compared	584
10. Music Theory for Algorithms	584
Scale Tables: Constraining Output to Pleasant Notes	584
Octave Derivation: Free Pitch Range	585
Arpeggio: Chord Tones in Sequence	586
Step Ornaments: Trills, Mordents, and Slides	586
Chord Progressions: Harmonic Movement	587
Total Data Budget for Rich Music	587
11. L-System Grammars: Fractal Melodies	588
The Concept	588
Why L-Systems Work for Music	588
Useful L-System Rules	589
Z80 Implementation	589
Melody as Motion, Not Absolute Notes	591
Tribonacci: Three Symbols for Richer Patterns	592
PRNG Melodies with Curated Seeds	592
Combining L-Systems with Other Techniques	593
Other Grammars for Music	593
See Also	593

Appendix J: Modern Tools for Retro Demo Production 595

J.1 Two Worlds, One Philosophy	595
J.2 Sync Tools	596
The Sync Problem	596
GNU Rocket	596
Vortex Tracker II	597
Blender VSE (Video Sequence Editor)	598
Blender Graph Editor	598
Motion Canvas	599
Comparison Table	599
J.3 Pre-visualisation and Storyboarding	600
Blender as Storyboard Tool	600
Video Editors for Rough-cut Assembly	600
The GABBA Workflow	601
Kolnogradov on Sync Planning	601
J.4 Data Generation Engines	601

Unity as a Data Generator	601
Unreal Engine as a Data Generator	602
Blender as a Data Generator	602
The Export Pipeline	602
When to Use Which	603
J.5 The PC Demoscene Toolchain: A Brief History	603
Farbrausch (1999–2012)	603
TiXL (2024–present)	604
Bonzomatic	605
Crinkler	605
Squishy	605
Shader Minifier	606
z80-optimizer	606
The Common Philosophy	606
J.6 Music and Sound Tools	606
Sointu	607
4klang	607
WaveSabre	607
Oidos	607
Furnace	607
Comparison Table	608
J.7 Practical Recipes	608
Recipe 1: GNU Rocket → Z80 Sync Table	608
Recipe 2: Blender VSE → Frame Number Table	609
Recipe 3: Unity VR → Trajectory Data	610
Recipe 4: Furnace → AY Music	611
Recipe 5: packbench → Pre-compression Analysis	613
Further Reading	613

Appendix K: GPU Superoptimization — When Brute Force Beats Human

Intuition	614
K.1 The Idea	614
K.2 Constant Multiplication	614
16-bit Multiply ($u8 \times K \rightarrow HL$)	615
K.3 Constant Division	615
K.4 Branchless Idioms	616
K.5 16-bit Idioms	617
K.6 The Packed Arithmetic Cassette	617
K.7 How It Works	618
K.8 Source & Data	618

What's New	619
v20 (2026-02-28)	619
Community feedback (Aki, Rombor via SinDiKat Slack)	619
Translation Memory tool + Russian edition	619
v19 (2026-02-28)	620
v18 (2026-02-27)	620
v17 (2026-02-27)	620
v16 (2026-02-26)	621
v12 (2026-02-24)	621
v11 (2026-02-24)	622
v10 (2026-02-24)	622

v0.8 (2026-02-24)	622
v0.7 (2026-02-24)	623
v0.6 (2026-02-23)	623

!!! UNDER CONSTRUCTION !!!

This book is a work in progress.
Content is incomplete, may contain errors,
and is subject to change without notice.

How this book is written:

1. With assistance from LLM (Claude Opus 4.6) based on open publications and data
2. By Alice directly — where personal expertise is “good enough” or no available sources or in new areas
3. Corrections and contributions on topic are welcome — PRs open

Factual errors are possible. If you spot one, please open an issue or PR.

Chapter 1: Thinking in Cycles

“Coder effects are always about evolving a computation scheme.” –
Introspec (spke), Life on Mars

You have 71,680 clock ticks. That is your canvas, your budget, your entire world. Every instruction you write costs some of those ticks. Every frame, the counter resets and you get another 71,680 – no more, no less. Miss the deadline and the screen tears, the music stutters, the illusion breaks.

This chapter is about learning to see your code the way a Z80 demoscener does: not as text, not as algorithms, but as a *budget*.

T-States: The Currency of the Z80

The Z80 CPU does not execute instructions in uniform chunks. Every instruction takes a specific number of **T-states** (time states) – the fundamental clock cycles of the processor. At 3.5 MHz, one T-state lasts approximately 286 nanoseconds. That number is not important. What is important is that instructions have wildly different costs, and you need to know those costs by heart.

Here is a handful of instructions you will use constantly:

Instruction	What it does	T-states
NOP	Nothing	4
LD A,B	Copy B into A	4
LD A, (HL)	Load byte from memory address in HL	7
LD (HL),A	Store A to memory address in HL	7
LD A,n	Load an immediate byte into A	7
INC HL	Increment HL	6
ADD A,B	Add B to A	4
PUSH HL	Push HL onto the stack	11
DJNZ label	Decrement B, jump if not zero	13 (taken) / 8 (falls through)
LDIR	Block copy, per byte	21 (repeating) / 16 (last byte)
OUT (n),A	Write A to I/O port	11

Look at the range. A register-to-register LD A,B costs 4 T-states – the minimum for any instruction. A memory read LD A, (HL) costs 7, because the CPU needs extra machine cycles to put the address on the bus and wait for the RAM to respond. LDIR, the block copy instruction that every Spectrum coder reaches for instinctively, costs 21 T-states per byte it copies (except the last byte, which costs 16). That is over five times the cost of a NOP.




Figure 1: T-state costs for common Z80 instructions

Why does this matter? Because when you are filling a screen, or updating sprite data, or computing the next frame of a plasma effect, every instruction is eating into your budget. The difference between a 4-T-state instruction and a 7-T-state instruction, multiplied across ten thousand iterations in an inner loop, is the difference between an effect that runs at 50 frames per second and one that does not.

Machine Cycles and Memory Access

Each T-state is one tick of the CPU clock, but the Z80 does not talk to memory on every tick. Instructions are broken into **machine cycles** (M-cycles), each of which takes 3-6 T-states. The first machine cycle of every instruction is the **opcode fetch** (M1), which always takes 4 T-states: the CPU puts the program counter on the address bus, reads the opcode byte, and simultaneously refreshes DRAM. Further machine cycles read additional bytes (operands, memory data) or write results.

This is why LD A,B takes exactly 4 T-states – it is a single-byte instruction that completes entirely within the opcode fetch. But LD A, (HL) takes 7 T-states: 4 for the opcode fetch, then 3 more for the memory read cycle where the CPU puts HL on the address bus and reads the byte at that address.

You do not need to memorize the internal machine cycle breakdown of every instruction. But understanding the pattern – opcode fetch + operand reads + memory accesses = total cost – helps you develop intuition for *why* instructions cost what they cost. A PUSH HL at 11 T-states makes sense when you realise the CPU

must do an opcode fetch (5T for this one, since it also decrements SP), then two separate memory write cycles (3T each) to store the high and low bytes of HL onto the stack.

The Frame: Your Canvas

The ZX Spectrum generates a PAL video signal at approximately 50 frames per second. Every frame, the ULA chip reads video memory and paints the screen, line by line. At the end of each frame, the ULA fires a maskable interrupt. The CPU executes a HALT instruction to wait for that interrupt, does its work, and then HALTs again to wait for the next frame. This is the heartbeat of every Spectrum program.

The number of T-states between one interrupt and the next – the **frame budget** – depends on the machine:

Machine	T-states per frame	Scanlines	Hz
ZX Spectrum 48K	69,888	312	50.08
ZX Spectrum 128K	70,908	311	50.02
Pentagon 128	71,680	320	48.83

Those are *total* T-states between interrupts. The practical budget is less – subtract the interrupt handler cost (a PT3 music player typically consumes 3,000–5,000 T-states per frame), HALT overhead, and on non-Pentagon machines, contention penalties. On a Pentagon with a music player, expect roughly 66,000–68,000 T-states for your main loop. Chapter 15 has the detailed tact-maps.




Figure 2: Frame budget breakdown across ZX Spectrum models

If your main loop – input handling, game logic, sound update, screen rendering – takes more T-states than one frame, you drop frames. Things slow down. The border stripe trick we will build later in this chapter will make that painfully visible.

To put these numbers in perspective: a single LDIR copying 6,912 bytes (one full screen of pixel data) costs approximately $6,912 \times 21 = 145,152$ T-states. That is

more than two entire frames on a 48K Spectrum. You cannot even copy the screen once per frame with the simplest possible method. This is the kind of constraint that forces creativity.

Pentagon vs. Wait Machines

You will notice that the frame budgets above differ between machines. The difference is not just in the numbers – it reflects a fundamental architectural split that shaped the ZX Spectrum demoscene.

The Original Sinclair Machines

On the original 48K and 128K Spectrums, the screen memory lives at addresses \$4000-\$5AFF (pixel data) and \$5800-\$5B00 (colour attributes). This memory region – the entire \$4000-\$7FFF range, in fact – is **contended memory**. The ULA (Uncommitted Logic Array), which generates the video signal, needs to read this memory to paint the screen. The CPU and the ULA share the same memory bus, and when both want to read at the same time, the ULA wins. The CPU is forced to wait.

During the 192 active display lines, every CPU access to the \$4000-\$7FFF range is potentially delayed. The delay follows a repeating 8-T-state pattern: 6, 5, 4, 3, 2, 1, 0, 0 extra wait states, cycling across each scanline. An instruction that should take 7 T-states might take 13 if it lands on the worst phase of the contention cycle.

This makes cycle-counting on original Spectrums a nightmare. Your carefully calculated inner loop runs at a different speed depending on where in the frame it executes, and whether the code or data it touches happens to be in the contended range. Introspec documented this in his “GO WEST” articles on Hype (2015): during screen rendering, each byte access to contended memory costs an average of 2.625 extra T-states. For stack operations writing to screen memory, expect roughly 1.3 extra T-states per byte.

The Pentagon: Clean Timing

The Pentagon 128, the most popular Soviet ZX Spectrum clone, took a different approach. Its designers gave the ULA its own memory access window that does not conflict with the CPU. **There is no contended memory on the Pentagon.** Every instruction takes exactly the number of T-states listed in the datasheet, regardless of where the code lives or what memory it accesses.

This is why the Pentagon has a different frame length – 71,680 T-states, 320 scanlines. The ULA timing is slightly different because there is no need to interleave CPU and ULA access. But the payoff is enormous: you can count cycles with absolute confidence. When your inner loop says it costs 36 T-states per iteration, it costs 36 T-states per iteration, every single time, everywhere in the frame.

This clean timing is why the Pentagon became the standard platform for the ZX Spectrum demoscene, particularly in the Former Soviet Union where these clones were ubiquitous. When you watch demos from groups like X-Trade, 4D+TBK (Triebkraft), or Life on Mars, they are overwhelmingly targeting Pentagon timing.

When Introspec wrote his legendary technical teardown of Illusion by X-Trade, the cycle counts he quoted assumed Pentagon.

For learning, the Pentagon model is ideal: you can focus on understanding what instructions cost without worrying about contention effects. All the T-state tables in this book assume Pentagon timing unless stated otherwise. When we need to discuss the differences (and we will, in Chapter 15), we will be explicit.

The practical rule: place your time-critical code in uncontended memory (\$8000-\$FFFF on a 48K), and your cycle counts will be correct on both Pentagons and original Spectrums.

Thinking in Budgets

Now that you know the frame size, you can start doing the arithmetic that defines Z80 demoscene thinking.

Say you want to fill the entire screen with a calculated colour every frame – a simple plasma effect, updating only the 768 bytes of attribute memory at \$5800. At 50 fps, you need to compute and write 768 colour values every 71,680 T-states.

If your inner loop per attribute byte looks like this:

```
ld    a,c          ; 4T    column index
add   a,b          ; 4T    add row index (diagonal pattern)
add   a,d          ; 4T    add frame counter (animation)
and   7            ; 7T    clamp to 0-7
ld    (hl),a       ; 7T    write attribute
inc   hl           ; 6T    next attribute address
; --- 32T per byte
```

That is 32 T-states per byte. For 768 bytes: $32 \times 768 = 24,576$ T-states. Add loop overhead (maintaining counters for rows and columns, the DJNZ for the inner loop), and you might land around 28,000-30,000 T-states. That leaves over 40,000 T-states for everything else – music playback, input handling, whatever you need.

But what if you wanted to update every *pixel* byte, all 6,144 of them? At 32 T-states per byte, that is 196,608 T-states – nearly three frames. Suddenly you are looking at a 17 fps update rate instead of 50 fps. You either need a faster inner loop, a smaller update region, or a different approach entirely.

This is how Z80 programmers think. Every design decision starts with arithmetic: how many bytes, how many T-states per byte, how many T-states in the frame budget, does it fit? When it does not fit, you do not reach for a faster machine – you reach for a cleverer algorithm.

Agon Light 2 Sidebar

The Agon Light 2 runs a Zilog eZ80 at 18.432 MHz. The eZ80 executes the same Z80 instruction set (it is a direct architectural descendant), but most instructions run in fewer clock cycles – many single-byte instructions complete in just 1 cycle instead of 4. At 18.432 MHz with a 50 Hz frame rate, you get approximately **368,640 T-states per frame**.

That is just over 5 times the Pentagon's budget. The same Z80 assembly language, the same registers, the same instruction mnemonics – but with five times the room to breathe. An inner loop that consumes 70% of a Pentagon frame might use only 14% of an Agon frame.

This does not make the Agon “easy.” It has its own constraints: no ULA-style video memory (the display is managed by an ESP32 coprocessor running the VDP), flat 24-bit addressing instead of banked memory, and a completely different I/O model. But if you have ever wished you could just have a *little more room* in your frame budget to try something ambitious, the Agon is where the same Z80 thinking scales up.

Throughout this book, we will note where the Agon's larger budget changes the calculus. For now, just remember the number: **~368,000 T-states**. Same ISA, five times the canvas.

Practical: Setting Up Your Development Environment

Before we write our first timing harness, you need a working toolchain. The setup described here follows sq's guide from Hype (2019), which has become the community standard.

What You Need

1. **VS Code** – your editor and integrated environment.
2. **Z80 Macro Assembler extension** by mborik (mborik.z80-macroasm) – syntax highlighting, code completion, symbol resolution for Z80 assembly. Install from the VS Code marketplace.
3. **Z80 Assembly Meter** by Nestor Sancho – displays the byte count and cycle count of the currently selected instruction(s) in the status bar. This is invaluable. Select a block of code and see its total T-state cost instantly.
4. **sjasmpplus** – the assembler itself. Cross-platform, open source, supports macros, Lua scripting, multiple output formats. Download from <https://github.com/z00m128/sjasmpplus> and place the binary somewhere in your PATH.
5. **Unreal Speccy** (Windows) or **Fuse** (cross-platform) – the emulator. Unreal Speccy is preferred for demo development because it emulates Pentagon timing accurately and has a built-in debugger.

Project Structure

Create a directory for your Chapter 1 experiments. A minimal project looks like this:

```
main.a80      -- your source file
build.bat     -- (Windows) sjasmpplus main.a80
Makefile      -- (macOS/Linux) make target
```

Build Configuration

In VS Code, set up a build task (`.vscode/tasks.json`) so you can compile with `Ctrl+Shift+B`:

```
"version": "2.0.0",
"tasks": [
  {
    "label": "Assemble Z80",
    "type": "shell",
    "command": "sjasmpplus",
    "args": [
      "--fullpath",
      "--nologo",
      "--msg=war",
      "${file}"
    ],
    "group": {
      "kind": "build",
      "isDefault": true
    },
    "problemMatcher": {
      "owner": "z80",
      "fileLocation": "absolute",
      "pattern": {
        "regexp": "^(.*)\\((\\d+)\\):\\s+(error|warning):\\s+(.*)$",
        "file": 1,
        "line": 2,
        "severity": 3,
        "message": 4
      }
    }
  }
]
```

Press `Ctrl+Shift+B`. If `sjasmpplus` is in your `PATH` and there are no errors, you will get a `.sna` or `.tap` file (depending on your source directives) that you can open directly in your emulator.

For integration with Unreal Speccy, Alex_Rider's 2024 extension adds an `F5-to-launch` binding – the emulator opens your compiled snapshot automatically. If you are on macOS or Linux and using Fuse, a simple Makefile rule does the same:

```
fuse --machine pentagon main.sna
```

Practical: The Timing Harness

This is the most important debugging tool you will build in this entire book. It is dead simple, it requires no special hardware, and you will use it constantly.

The idea: change the border colour to red immediately before the code you want

to measure, and change it back to black immediately after. The Spectrum's border is drawn by the ULA in real time, synchronised with the electron beam. A wider red stripe means more T-states spent in your code.

Here is the complete harness:

```

    ORG $8000

start:
    ; Wait for the frame interrupt
    halt

    ; --- Border RED: code under test begins ---
    ld  a, 2          ; 7T  red = colour 2
    out ($FE), a      ; 11T write to border port

    ; ===== CODE UNDER TEST =====
    ; Replace this block with whatever you want to measure.
    ; Example: 256 iterations of a NOP loop.

    ld  b, 0          ; 7T  B=0 wraps to 256 iterations
.loop:
    nop               ; 4T
    nop               ; 4T
    nop               ; 4T
    nop               ; 4T -- 16T per iteration body
    djnz .loop        ; 13T taken, 8T on final iteration
    ; Total: 256 * (16+13) - 5 = 7,419 T-states

    ; ===== END CODE UNDER TEST =====

    ; --- Border BLACK: idle ---
    xor a             ; 4T  A=0 (black), shorter than LD A,0
    out ($FE), a      ; 11T

    ; Loop forever
    jr  start

```

Load this into your emulator. You will see a red stripe across the border. The height of that stripe is directly proportional to the number of T-states your test code consumed.

Timing harness output — red border stripes show T-states consumed by the code under test, black gaps show idle time

Figure 3: Timing harness output — red border stripes show T-states consumed by the code under test, black gaps show idle time

Reading the Stripe

Each scanline takes 224 T-states (on Pentagon). So if your red stripe is N scanlines tall, your code took approximately $N \times 224$ T-states. The example above uses about 7,419 T-states, which is roughly 33 scanlines – you should see a red band about

one-sixth of the way down the border.

Now try replacing the NOP loop with something heavier. Replace the four NOPs with:

```
.loop:
    ld    a,(hl)        ; 7T
    add   a,(hl)        ; 7T
    ld    (de),a        ; 7T
    inc   hl            ; 6T  -- 27T per iteration body
    djnz  .loop         ; 13T taken
    ; Total: 256 * (27+13) - 5 = 10,235 T-states
```

The red stripe grows noticeably. That visual difference – you can see it without a debugger, without a profiler, without any tooling at all – is 2,816 T-states. About 12 scanlines.

This is how Spectrum demo coders have profiled their effects since the 1980s. The border is your oscilloscope.

Variations

You can use different colours to mark different phases of your code:

```
ld    a, 2            ; red
out   ($FE), a
call  render_sprites
ld    a, 1            ; blue
out   ($FE), a
call  update_music
ld    a, 4            ; green
out   ($FE), a
call  game_logic
xor   a               ; black
out   ($FE), a
```

Now the border shows a red band (sprite rendering), then blue (music), then green (game logic), then black (idle time). You can see at a glance which subsystem is eating your frame budget.

A note on `xor a` versus `ld a, 0`: both set A to zero. `XOR A` takes 4 T-states and 1 byte. `LD A, 0` takes 7 T-states and 2 bytes. In a timing harness the difference is negligible, but it is worth noticing – this kind of micro-awareness is what Z80 coding is made of.

What Fits in a Frame?

Let us use our budget arithmetic to answer some practical questions.

How many sprites can you draw per frame? A 16x16 masked sprite using the OR+AND method takes roughly 144 T-states per row (mask+draw for two bytes, row advance, loop control) for 16 rows — about 2,400 T-states per sprite (see Chapter 16 for the instruction-level breakdown). On a Pentagon, that is 71,680 /

2,400 = ~29 sprites, if sprite rendering were the *only* thing you did. In practice, with music, game logic, and everything else, 8-12 full sprites per frame is typical.

How many bytes can LDIR copy per frame? At 21 T-states per byte: $71,680 / 21 = 3,413$ bytes. Not even half the screen.

How many multiply operations? A fast square-table 8x8 multiply takes about 54 T-states. $71,680 / 54 = 1,327$ multiplications per frame. A single-point 3D rotation needs 9 multiplies. So you could rotate about 147 points per frame *if you did nothing else*. Practical limit with a full demo engine: 30-50 points.

Every design question reduces to this arithmetic. Can I do it? How many can I do? What do I have to give up to make room?

Historical Note: Dark's Advice

In 1997, a programmer named Dark from the group X-Trade published a series of articles in *Spectrum Expert* #01, a Russian electronic magazine for ZX Spectrum developers. These articles covered multiplication, division, sine/cosine generation, and line-drawing algorithms in Z80 assembly – the fundamental building blocks that power every demo effect.

Dark opened with this advice:

“Read a maths textbook – derivatives, integrals. Knowing them, you can create a table of practically any function in assembly.”

This was not empty theory. Dark was not just a writer – he was a coder. X-Trade's demo *Illusion*, released at ENLIGHT'96, featured a textured spinning sphere, a rotozoomer, a 3D engine, and a bouncing dot scroller, all running on a 3.5 MHz Z80. The algorithms Dark described in his magazine articles were the same algorithms powering his demo's effects.

Twenty years later, Introspec (spke) published a detailed technical teardown of *Illusion* on Hype, analysing the inner loops instruction by instruction, counting every T-state. The magazine articles from 1997 and the reverse engineering from 2017 tell the same story from both sides: the author explaining his building blocks, and a peer measuring the finished machine. We will follow this thread throughout the book.

Dark's point stands: the maths is not optional. You do not need a degree in mathematics, but you need to understand how to turn a mathematical function into a table, how to approximate expensive operations with cheap ones, and how to think about error versus speed. Chapter 4 will walk through Dark's algorithms in detail. For now, remember his advice. It is the starting point of everything.

The Computation Scheme

Introspec, writing about what makes a good demo effect, distilled the philosophy into a single sentence:

“Coder effects are always about evolving a computation scheme.”

This is the deepest insight in this chapter. A demo effect is not a picture; it is a *process*. Each frame, the computation scheme produces the next state from the previous one. The art is in choosing a scheme that produces visually compelling evolution while fitting within the frame budget.

A plasma is a computation scheme: sum sine waves at each grid position, offset by time. A tunnel is a computation scheme: look up angle and distance from pre-computed tables, offset by time. A spinning 3D object is a computation scheme: multiply vertex coordinates by a rotation matrix that changes each frame. The particular scheme determines the visual result, the T-state cost, and the memory requirements – all at once, all interlocked.

When you sit down to write an effect, you are not asking “how do I draw this picture.” You are asking “what computation, evolved frame by frame, produces this visual?” That shift in thinking – from image to process, from output to scheme – is the Z80 programmer’s worldview.

And the first constraint on any scheme is the budget. 71,680 T-states. Can you evolve your computation within that budget? If not, can you find a cheaper scheme that produces a similar visual? Can you precompute part of the scheme into tables? Can you spread the computation across multiple frames? Can you exploit symmetry to compute half the screen and mirror the other half?

These are the questions that drive every chapter in this book. They start here, with counting T-states.

Summary

- Every Z80 instruction has a specific T-state cost. Learn the common ones by heart: NOP = 4, LD A,B = 4, LD A,(HL) = 7, PUSH HL = 11, LDIR = 21/16, OUT (n),A = 11.
 - The **frame budget** is your hard constraint: 69,888 T-states (48K), 70,908 (128K), or 71,680 (Pentagon). At 50 fps, everything must fit.
 - **Pentagon has no contended memory**, making cycle counts reliable and predictable. This is why it became the demoscene standard.
 - The **Agon Light 2** (eZ80 @ 18.432 MHz) gives ~368,000 T-states per frame – same instruction set, five times the room.
 - The **border colour timing harness** is your oscilloscope: red before, black after, read the stripe width.
 - Z80 programming is **budget arithmetic**: bytes x T-states per byte vs. frame budget. Every design decision starts here.
 - Effects are **computation schemes evolved over time**. The art is finding a scheme that fits the budget and looks good.
-

Try It Yourself

1. Build the timing harness from this chapter. Replace the NOP loop with a LDIR that copies 256 bytes and compare the stripe width to the NOP version. Calculate the expected T-state difference and verify it visually.
2. Write a loop that fills all 768 bytes of attribute memory (\$5800-\$5AFF) with a single colour value. Measure it with the harness. Now try filling it using LDIR instead of a byte-by-byte loop. Which is faster? By how many scanlines?
3. Open the Z80 Assembly Meter in VS Code. Select different code blocks and watch the T-state counter in the status bar. Get used to checking costs as you write.
4. Set up the multi-colour border profiler (red / blue / green / black) with three dummy loops of different lengths. Adjust the loop counts until you can visually distinguish all three bands. This is your calibration exercise for reading border timing.

Next: Chapter 2 - The Screen as a Puzzle. We will dive into the Spectrum's notoriously scrambled video memory layout and learn why INC H moves you one pixel down.

Chapter 2: The Screen as a Puzzle

“Why do the rows go in that order?” – Every ZX Spectrum programmer, at some point

Open any emulator, type `PEEK 16384` and you are reading the first byte of the Spectrum’s screen. But which byte is it? Not the top-left of the screen in any simple sense. The pixel at coordinate (0,0) is there, yes – but the pixel at (0,1), the very next row down, lives 256 bytes away. The pixel at (0,8), the top row of the second character cell, lives only 32 bytes from the start. And the pixel at (0,64) – the first row of the screen’s middle third – lives exactly 2,048 bytes from the start, at \$4800.

This is the Spectrum’s most famous puzzle. The screen memory layout is not sequential, not intuitive, and not an accident. It is a consequence of hardware design choices made in 1982, and it shapes every piece of code that touches the display. Understanding this layout – and learning the tricks that make it fast to navigate – is fundamental to everything that follows in this book.

The Memory Map: 6,912 Bytes of Screen

The Spectrum’s display occupies a fixed region of memory:

\$5800 - \$5AFF Attributes 768 bytes (32 x 24 colour cells)


The pixel area holds the bitmap: 256 pixels across, packed 8 per byte, giving 32 bytes per row. With 192 rows, that is $32 \times 192 = 6,144$ bytes. Each byte represents 8 horizontal pixels, with bit 7 as the leftmost pixel and bit 0 as the rightmost.

The attribute area holds the colour information: one byte per 8x8 character cell. There are 32 columns and 24 rows, giving $32 \times 24 = 768$ bytes.


Together: $6,144 + 768 = 6,912$ bytes. That is the entire display.

The pixel data and attribute data serve different purposes but are tightly coupled. Each pixel byte controls 8 dots on screen; the attribute byte for the corresponding 8x8 cell controls what colour those dots appear in. Change the pixel and you change the shape. Change the attribute and you change the colour. But you can only change colour for an entire 8x8 block – not per pixel. This is the “attribute clash” that defines the Spectrum’s visual character, and we will return to it shortly.

First, the puzzle: why are the pixel rows scrambled?



Address Bit Layout: 010TTSSS LLLCCCC



Example: Third 1, scanline 3, char row 5, column 10

High: 010 | 01 | 011 = \$4B (TT=01, SSS=011)
 Low: 101 | 01010 = \$AA (LLL=101, CCCCC=01010)
 Address: \$4BAA → pixel row 107 (third 1, row 5, line 3), column 10

Figure 4: ZX Spectrum screen memory layout with thirds, character cells, and attribute area

The Interleave: Where the Rows Live

If the Spectrum stored its pixel rows sequentially, row 0 would be at \$4000, row 1 at \$4020, row 2 at \$4040, and so on. Each row is 32 bytes, so row N would simply be at \$4000 + N * 32. Simple, fast, sensible.

That is not what happens.

The screen is divided into three **thirds**, each 64 pixel rows tall. Within each third, the rows are interleaved by character cell row. Here is where the first 16 rows actually live:

Row 1:	\$4100	Third 0, char row 0, scan line 1
Row 2:	\$4200	Third 0, char row 0, scan line 2
Row 3:	\$4300	Third 0, char row 0, scan line 3
Row 4:	\$4400	Third 0, char row 0, scan line 4
Row 5:	\$4500	Third 0, char row 0, scan line 5
Row 6:	\$4600	Third 0, char row 0, scan line 6
Row 7:	\$4700	Third 0, char row 0, scan line 7
Row 8:	\$4020	Third 0, char row 1, scan line 0
Row 9:	\$4120	Third 0, char row 1, scan line 1
Row 10:	\$4220	Third 0, char row 1, scan line 2
Row 11:	\$4320	Third 0, char row 1, scan line 3
Row 12:	\$4420	Third 0, char row 1, scan line 4
Row 13:	\$4520	Third 0, char row 1, scan line 5
Row 14:	\$4620	Third 0, char row 1, scan line 6
Row 15:	\$4720	Third 0, char row 1, scan line 7

Look at the pattern. The first 8 rows are the 8 scanlines of character row 0 – but they are 256 bytes apart, not 32. Within those 8 rows, the high byte of the address increments by 1 each time: \$40, \$41, \$42, ... \$47. Then row 8 jumps to \$4020 – back to a high byte of \$40, but with the low byte advanced by 32.

Here is the complete picture for the top third of the screen:

Char row 1:	scan lines at \$4020, \$4120, \$4220, \$4320, \$4420, \$4520, \$4620,
↪	\$4720
Char row 2:	scan lines at \$4040, \$4140, \$4240, \$4340, \$4440, \$4540, \$4640,
↪	\$4740
Char row 3:	scan lines at \$4060, \$4160, \$4260, \$4360, \$4460, \$4560, \$4660,
↪	\$4760
Char row 4:	scan lines at \$4080, \$4180, \$4280, \$4380, \$4480, \$4580, \$4680,
↪	\$4780
Char row 5:	scan lines at \$40A0, \$41A0, \$42A0, \$43A0, \$44A0, \$45A0, \$46A0,
↪	\$47A0
Char row 6:	scan lines at \$40C0, \$41C0, \$42C0, \$43C0, \$44C0, \$45C0, \$46C0,
↪	\$47C0
Char row 7:	scan lines at \$40E0, \$41E0, \$42E0, \$43E0, \$44E0, \$45E0, \$46E0,
↪	\$47E0

The middle third starts at \$4800 and follows the same pattern. The bottom third starts at \$5000.

Why?

The reason is the ULA – the Uncommitted Logic Array that generates the video signal. The ULA reads one byte of pixel data and one byte of attribute data for every 8-pixel character cell it draws. It needs both bytes at specific moments as it rasters across the screen.

The interleaved layout meant that the ULA's address counter logic could be built with fewer gates. As the ULA scans left to right across a character row, it increments the low 5 bits of the address (the column). When it reaches the right edge, it increments the high byte to move to the next scanline within the same character row. When it finishes all 8 scanlines, it wraps the high byte and advances the low-byte row bits.

This is elegant from a hardware perspective. The ULA's address generation is a simple combination of counters – no multiplication, no complex address arithmetic. The PCB routing was simpler, the gate count was lower, and the chip was cheaper to manufacture.

The programmer pays the price.

The Bit Layout: Decoding (x, y) into an Address

ZX Spectrum screen memory layout — interleaved thirds with colour-coded address bit mapping

Figure 5: ZX Spectrum screen memory layout — interleaved thirds with colour-coded address bit mapping

To understand the interleave precisely, look at how the Y coordinate maps into the 16-bit screen address. Consider a pixel at column x (0-255) and row y (0-191). The byte containing that pixel is at:

Low byte: L L L C C C C C

Where: - TT = which third of the screen (0, 1, or 2). Bits 7-6 of y . - SSS = scanline within the character cell (0-7). Bits 2-0 of y . - LLL = character row within the third (0-7). Bits 5-3 of y . - CCCCC = column in bytes (0-31). This is $x / 8$, or equivalently bits 7-3 of x .

The crucial thing: the bits of y are not in order. Bits 7-6 go to one place, bits 5-3 go to another, and bits 2-0 go to yet another. The y coordinate is sliced up and distributed across the address.

Let us visualise this with a concrete example. Pixel (80, 100):

```
y = 100:    binary = 01100100
            TT  = 01      (third 1, the middle third)
            LLL = 100     (char row 4 within the third)
            SSS = 100     (scan line 4 within the char cell)
```

```
High byte:  0  1  0  0  1  1  0  0  = $4C
Low byte:   1  0  0  0  1  0  1  0  = $8A
```

Address: \$4C8A

The bit within that byte is determined by the low 3 bits of x. Bit 7 is the leftmost pixel, so pixel position (x AND 7) maps to bit 7 - (x AND 7).

The address calculation in Z80

Converting (x, y) to a screen address is something you need to do fast and often. Here is a standard routine:

```
; Input:  B = y (0-191), C = x (0-255)
; Output: HL = screen address, A = bit mask
;
pixel_addr:
    ld    a, b          ; 4T    A = y
    and   $07           ; 7T    A = SSS (scan line within char)
    or    $40           ; 7T    A = 010 00 SSS (add screen base)
    ld    h, a          ; 4T    H = high byte (partial)

    ld    a, b          ; 4T    A = y again
    rla           ; 4T    \  shift bits 5-3 of y
    rla           ; 4T    /  left to bits 7-5
    and   $E0           ; 7T    mask to get LLL 00000
    ld    l, a          ; 4T    L = LLL 00000 (partial)

    ld    a, b          ; 4T    A = y again
    and   $C0           ; 7T    A = TT 000000
    rra           ; 4T    \
    rra           ; 4T    | shift bits 7-6 of y
    rra           ; 4T    /  to bits 4-3
    or    h            ; 4T    combine with SSS
    ld    h, a          ; 4T    H = 010 TT SSS (complete)

    ld    a, c          ; 4T    A = x
    rra           ; 4T    \
    rra           ; 4T    | x / 8
    rra           ; 4T    /
    and   $1F           ; 7T    mask to CCCCC
    or    l            ; 4T    combine with LLL 00000
    ld    l, a          ; 4T    L = LLL CCCCC (complete)
; --- Total: ~87 T-states
```

87 T-states is not cheap. In a tight inner loop processing thousands of pixels, you would not call this routine per pixel. Instead, you calculate the starting address once and then navigate the screen using fast pointer manipulation – which brings us to the most important routine in Spectrum graphics programming.

Pixel plotting demo — individual pixels placed on screen using the address calculation routine

Figure 6: Pixel plotting demo — individual pixels placed on screen using the address calculation routine

DOWN_HL: Moving One Pixel Row Down

You have a pointer in HL to some byte on the screen. You want to move it one pixel row down - to the byte at the same column, one scanline lower. How hard can this be?

On a linear framebuffer, you add 32 (the number of bytes per row). One ADD HL, DE with DE = 32: 11 T-states, done.

On the Spectrum, it is a puzzle within the puzzle. Moving one pixel row down means:

1. **Within a character cell** (scanlines 0-6 to 1-7): increment H. The scanline bits are in the low 3 bits of H, so INC H moves you one scanline down.
2. **Crossing a character cell boundary** (scanline 7 to scanline 0 of the next row): reset the scanline bits of H back to 0, and add 32 to L to move to the next character row.
3. **Crossing a third boundary** (bottom of char row 7 in one third to top of char row 0 in the next): reset L back, and add 8 to H to move to the next third. Equivalently, add \$0800 to the address.

The classic routine handles all three cases:

```
; DOWN_HL: move HL one pixel row down on the Spectrum screen
; Input:  HL = current screen address
; Output: HL = screen address one row below
;
down_hl:
    inc h            ; 4T    try moving one scan line down
    ld  a, h        ; 4T
    and 7            ; 7T    did we cross a character boundary?
    ret nz           ; 11/5T no: done

    ; Crossed a character cell boundary.
    ; Reset scan line to 0, advance character row.
    ld  a, l        ; 4T
    add a, 32        ; 7T    next character row (L += 32)
    ld  l, a        ; 4T
    ret c           ; 11/5T if carry, we crossed into next third

    ; No carry from L, but we need to undo the H increment
    ; that moved us into the wrong third.
    ld  a, h        ; 4T
    sub 8            ; 7T    back up one third in H
    ld  h, a        ; 4T
    ret             ; 10T
```

This routine takes different amounts of time depending on which case it hits:

Case	Frequency	T-states
Within a character cell	7 out of 8 rows	$4 + 4 + 7 + 11 = \mathbf{26}$
Character boundary, same third	7 out of 64 rows	$4 + 4 + 7 + 5 + 4 + 7 + 4 + 5 + 4 + 7 + 4 + 10 = \mathbf{65}$
Third boundary	2 out of 192 rows	$4 + 4 + 7 + 5 + 4 + 7 + 4 + 11 = \mathbf{46}$

The common case – staying within a character cell – is fast: 26 T-states (a conditional RET that fires costs 11T, not 5T). The uncommon case (crossing a character row boundary within the same third) is 65 T-states. Averaged over all 192 rows, the cost works out to about **30.5 T-states per call**.

That average hides a problem. If you are iterating down the full screen and calling DOWN_HL on every row, those occasional 65-T-state calls spike your per-line timing unpredictably. For a demo effect that needs consistent timing per scanline, this jitter is unacceptable.

Introspec's Optimisation

In December 2020, Introspec (spke) published a detailed analysis on Hype titled “Once more about DOWN_HL” (Eshchy raz pro DOWN_HL). The article examined the problem of iterating down the full screen efficiently – not just the cost of one call, but the total cost of moving HL through all 192 rows.

The naive approach – calling the classic DOWN_HL routine 191 times – costs **5,922 T-states** for a full screen traversal. Introspec's goal was to find the fastest way to iterate through all 192 rows, visiting every screen address in top-to-bottom order.

His key insight was to use **split counters**. Instead of testing the address bits after every increment to detect boundary crossings, he structured the loop to match the screen's three-level hierarchy directly:

```

For each third (3 iterations):
  For each character row within the third (8 iterations):
    For each scan line within the character cell (8 iterations):
      process this row
      INC H                      ; next scan line
    undo 8 INC H's, ADD 32 to L  ; next character row
  undo 8 ADD 32's, advance to next third

```

The innermost operation is just INC H – 4 T-states. No testing, no branching. The character-row and third transitions happen at fixed, predictable points in the loop, so there is no conditional logic in the inner loop at all.

The result: **2,343 T-states** for a full screen traversal. That is a 60% improvement over the classic approach, and the per-line cost is absolutely predictable – no jitter.

There was also an elegant variation attributed to RST7, using a dual-counter approach where the outer loop maintains a pair of counters that naturally track the character-row and third boundaries. The inner loop body reduces to a single

INC H, and the boundary handling is folded into the counter manipulation at the outer loop level.

The practical lesson: when you need to iterate through the Spectrum's screen in order, do not call a general-purpose DOWN_HL routine 191 times. Restructure your loop to match the screen's natural hierarchy, and the branching disappears.

Here is a simplified version of the split-counter approach:

```
; Iterate all 192 screen rows using split counters
; HL = $4000 at entry (top-left of screen)
;
iterate_screen:
    ld    hl, $4000        ; 10T  start of screen
    ld    c, 3             ; 7T   3 thirds

.third_loop:
    ld    b, 8             ; 7T   8 character rows per third

.row_loop:
    push hl                ; 11T  save start of this char row

    ; --- Process 8 scan lines within this character cell ---
    REPT 7
        ; ... your per-row code here, using HL ...
        inc h              ; 4T   next scan line
    ENDR
    ; ... process the 8th (last) scan line ...

    pop  hl                ; 10T  restore char row start
    ld   a, l              ; 4T
    add  a, 32             ; 7T   next character row
    ld   l, a              ; 4T

    djnz .row_loop         ; 13T/8T

    ; Advance to next third
    ld   a, h              ; 4T
    add  a, 8              ; 7T   next third ($0800 higher)
    ld   h, a              ; 4T

    dec  c                 ; 4T
    jr   nz, .third_loop   ; 12T/7T
```

The REPT 7 directive (supported by sjasmpplus) repeats the block 7 times at assembly time – a partial unroll. Inside that block, moving down one scanline is a single INC H. No testing, no branching. The character-row advance and third-advance happen at the fixed outer loop boundaries.

Attribute Memory: 768 Bytes That Changed Everything

Below the pixel data, at \$5800-\$5AFF, sits the attribute memory. It is 768 bytes – one for each 8x8 character cell on the screen, arranged sequentially from left to right, top to bottom. Unlike the pixel area, the attribute layout is entirely linear: cell (col, row) is at $\$5800 + \text{row} * 32 + \text{col}$.

Each attribute byte has this layout:

```
+-----+-----+-----+-----+
|  F  |  B  | PAPER |  INK  |
+-----+-----+-----+-----+
```

F = Flash (0 = off, 1 = flashing at ~1.6 Hz)

B = Bright (0 = normal, 1 = bright)

PAPER = Background colour (0-7)

INK = Foreground colour (0-7)


The 3-bit colour codes map to:

1 = Blue	5 = Cyan
2 = Red	6 = Yellow
3 = Magenta	7 = White

With the BRIGHT bit, each colour has a normal and bright variant. Black stays black whether bright or not, so the total palette is 15 distinct colours:

Bright: Black Blue Red Magenta Green Cyan Yellow White
(brighter versions of each)

ZX Spectrum Attribute Byte Layout



Example:

\$41 = 01000001

= BRIGHT Blue ink on Black paper



Figure 7: Attribute byte bit layout showing flash, bright, paper, and ink fields

An attribute byte of \$47 = 01000111: flash off (bit 7 = 0), bright **on** (bit 6 = 1), paper

= 000 (black), ink = 111 (white). Bright white text on a black background. The non-bright version is \$07 = 00000111 – the Spectrum’s default after BORDER 0: PAPER 0: INK 7.

This kind of bit-level detail matters when you are constructing attribute values at speed. A common pattern:

```
; Build an attribute byte: bright white ink on blue paper
; Bright = 1, Paper = 001 (blue), Ink = 111 (white)
; = 01 001 111 = $4F
    ld    a, $4F
```

The Attribute Clash

Here is the defining constraint of the ZX Spectrum: within each 8x8 pixel cell, you can only have **two colours** – ink and paper. Every set pixel (1) displays in the ink colour. Every clear pixel (0) displays in the paper colour. You cannot have three colours, or gradients, or per-pixel colouring, within a single cell.

This means that if a red sprite overlaps with a green background, the 8x8 cell containing the overlap must choose: all set pixels in this cell are either red or green. You cannot have some red and some green set pixels in the same cell. The visual result is a jarring block of colour that “clashes” with its surroundings – the infamous attribute clash.

```
+-----+-----+
| Red   | Red on |
| sprite| green   |
| pixels| back-  |
|       | ground |
+-----+-----+
```

With attribute clash (Spectrum reality):

```
+-----+-----+
| Red   | Either |
| sprite| ALL red |
| pixels| or ALL |
|       | green  |
+-----+-----+
```

The overlapping cell cannot have both colours.

Many early Spectrum games simply avoided the problem: monochrome graphics, or characters carefully designed to align with the 8x8 grid. Games like Knight Lore and Head Over Heels used a single ink/paper pair for the entire play area, eliminating clash entirely at the cost of colour.

But the demoscene saw this differently. Attribute clash is not just a limitation – it is a **creative constraint**. The 8x8 grid forces a particular aesthetic: bold blocks of colour, sharp geometric patterns, deliberate use of contrast. Demo effects that work entirely in attribute space – tunnels, plasmas, scrollers – can update 768 bytes per frame instead of 6,144, freeing enormous amounts of cycle budget for computation.

When your entire display is attribute-driven, clash becomes irrelevant because you are not mixing sprites with backgrounds – the attributes *are* the graphics.

Introspec’s demo Eager (2015) built its visual language entirely around this insight. The tunnel effect, the chaos zoomer, and the colour cycling animation all operate on attributes, not pixels. The result is an effect that runs at full frame rate with room to spare for digital drums and a sophisticated scripting engine. Clash is not a problem because the constraint was embraced from the start.

The Border: More Than Decoration

The 256x192 pixel display area sits in the centre of the screen, surrounded by a wide border. The border colour is set by writing to port \$FE:

```
ld    a, 1          ; 7T    blue = colour 1
out   ($FE), a      ; 11T   set border colour
```

Only bits 0-2 of the byte written to \$FE affect the border colour. There are 8 colours (0-7), with no bright variants – the border palette is the non-bright set. Bits 3 and 4 of port \$FE control the MIC and EAR outputs (tape interface and beeper sound), so you should mask or set those bits appropriately if you are not intending to make noise.

The border colour change takes effect immediately – on the very next scanline being drawn. This is what makes the border so useful as a debugging tool. As we saw in Chapter 1, changing the border colour before and after a section of code creates a visible stripe whose height reveals the code’s T-state cost. The border is your oscilloscope.

Border Effects

Because border colour changes are visible on the next scanline, precisely timed OUT instructions can create multicolour stripes, raster bars, and even crude graphics in the border area.

The basic principle: the ULA draws one scanline every 224 T-states (on Pentagon). If you execute an OUT (\$FE), A instruction at the right moment, you change the border colour at a specific horizontal position on the current scanline. By executing a rapid sequence of OUT instructions with different colour values, you can paint horizontal stripes of colour in the border.

```
; Simple border stripes
; Assumes we are synced to the start of a border scanline
```

```
ld    a, 2          ; 7T    red
out   ($FE), a      ; 11T
; ... delay to fill this scanline ...
ld    a, 5          ; 7T    cyan
out   ($FE), a      ; 11T
; ... delay to fill next scanline ...
ld    a, 6          ; 7T    yellow
out   ($FE), a      ; 11T
```


More advanced border effects can create gradient bars, scrolling text, or even low-resolution images. The challenge is extreme: you have 224 T-states per scanline, and each colour change costs at minimum 18 T-states (7 for LD A,n + 11 for OUT). That gives you roughly 12 colour changes per scanline, which means at most 12 horizontal colour bands per line.

Demo coders have pushed this to remarkable extremes. By pre-loading multiple registers with colour values and using faster sequences like OUT (C), A followed by register swaps, they squeeze more colour changes per line. The border becomes a display unto itself – a canvas outside the canvas.

For our purposes, the border's most important role is the one from Chapter 1: a free, always-available timing visualiser. When you are optimising the screen-fill routine later in this chapter, the border is how you will see your progress.

Practical: The Checkerboard Fill

The example at `chapters/ch02-screen-as-puzzle/examples/fill_screen.a80` fills the pixel area with a checkerboard pattern and the attributes with bright white on blue. Let us walk through it section by section.

```
ORG $8000

SCREEN EQU $4000      ; pixel area start
ATTRS  EQU $5800      ; attribute area start
SCRLEN EQU 6144        ; pixel bytes (256*192/8)
ATTLEN EQU 768         ; attribute bytes (32*24)
```

The code is placed at \$8000 – safely in uncontended memory on all Spectrum models. The constants name the key addresses and sizes.

```
start:
    ; --- Fill pixels with checkerboard pattern ---
    ld  hl, SCREEN
    ld  de, SCREEN + 1
    ld  bc, SCRLEN - 1
    ld  (hl), $55      ; checkerboard: 01010101
    ldir
```

This uses the classic LDIR self-copy trick. It writes \$55 (binary 01010101) to the first byte at \$4000, then copies from each byte to the next for 6,143 bytes. The result: every byte of the pixel area is \$55, which produces alternating set/clear pixels – a checkerboard. Because the pattern is the same in every byte, the interleaved row order does not matter – every row gets the same pattern regardless.

Cost: LDIR copies 6,143 bytes. The last iteration costs 16T, all others 21T: $(6,143 - 1) \times 21 + 16 = 128,998$ T-states. Nearly two full frames on a Pentagon. This is fine for a one-time setup, but you would never do this in a per-frame rendering loop.

```
    ; --- Fill attributes: white ink on blue paper ---
    ; Attribute byte: flash=0, bright=1, paper=001 (blue), ink=111 (white)
    ; = 01 001 111 = $4F
    ld  hl, ATTRS
```

```
ld    de, ATTRS + 1
ld    bc, ATTLEN - 1
ld    (hl), $4F
ldir
```

Same technique for the attributes. The value \$4F decodes as: flash off (0), bright on (1), paper blue (001), ink white (111). Every 8x8 cell gets bright white ink on blue paper. The checkerboard pixels are set/clear, so you see alternating white and blue dots – a classic ZX Spectrum visual pattern.

Cost: LDIR copies 767 bytes – $(767 - 1) \times 21 + 16 = 16,102$ T-states.

```
; --- Border: blue ---
ld    a, 1
out   ($FE), a

; Infinite loop
.wait:
    halt
    jr   .wait
```

Sets the border to blue (colour 1) to match the paper colour, creating a visually clean frame. Then loops forever, halting between frames. The HALT waits for the next maskable interrupt, which fires once per frame – this is the idle heartbeat of every Spectrum program.

Screen fill with alternating bytes — checkerboard pattern in bright white on blue

Figure 8: Screen fill with alternating bytes — checkerboard pattern in bright white on blue

What to try

Load `fill_screen.a80` in your assembler and emulator. Then experiment:

- Change \$55 to \$AA for the inverse checkerboard, or to \$FF for solid fill, or \$81 for vertical bars.
- Change \$4F to \$07 to see the same pattern without BRIGHT, or to \$38 for white paper with black ink (the inverse of default).
- Try \$C7 – that sets the flash bit. Watch the characters alternate between ink and paper colours at about 1.6 Hz.
- Replace the LDIR pixel fill with a DOWN_HL loop that writes different patterns to different rows. Now you will see the interleave in action: if you write \$FF to rows 0-7 (the first character cell's scanlines), the filled area will appear as 8 horizontal stripes separated by gaps – because those rows are 256 bytes apart, not 32.

Navigating the Screen: A Practical Summary

Here are the essential pointer operations for the Spectrum screen, collected in one place. These are the building blocks of every graphics routine.

Moving right one byte (8 pixels)

```
inc l          ; 4T
```

This works within a character row because the column is in the low 5 bits of L. If you need to cross byte boundaries at the right edge (column 31 to column 0 of the next row), you need the full DOWN_HL plus reset of L – but typically you do not, because your loops are 32 bytes wide.

Moving down one pixel row

```
inc h          ; 4T    (within a character cell)
```

This works for 7 out of 8 rows. On the 8th row, you need the full boundary-crossing logic from the DOWN_HL routine above.

Moving down one character row (8 pixels)

```
ld  a, l        ; 4T
add a, 32        ; 7T
ld  l, a         ; 4T    total: 15T (if no third crossing)
```

This advances by one character row within a third. If L overflows (carry set), you have crossed into the next third and need to add 8 to H.

Moving up one pixel row

```
dec h          ; 4T    (within a character cell)
```

The inverse of INC H. Same boundary issues at character cell and third boundaries. Here is the full UP_HL routine, the mirror of DOWN_HL:

```
; UP_HL: move HL one pixel row up on the Spectrum screen
; Input:  HL = current screen address
; Output: HL = screen address one row above
;
; Classic version:
up_hl:
    dec h          ; 4T    try moving one scan line up
    ld  a, h        ; 4T
    and 7           ; 7T    did we cross a character boundary?
    cp  7           ; 7T
    ret nz          ; 11/5T no: done

    ; Crossed a character cell boundary upward.
    ld  a, l        ; 4T
    sub 32          ; 7T    previous character row (L -= 32)
    ld  l, a         ; 4T
    ret c          ; 11/5T if carry, crossed into prev third

    ld  a, h        ; 4T
    add a, 8         ; 7T    compensate H
    ld  h, a         ; 4T
    ret            ; 10T
```

There is a subtle optimisation here, contributed by Art-top (Artem Topchiy): replacing `and 7 / cp 7` with `cpl / and 7`. After `DEC H`, if the low 3 bits of `H` wrapped from 000 to 111, we crossed a character boundary. The classic test checks `AND 7` then compares with 7. The optimised version complements first: if the bits are 111, `CPL` makes them 000, and `AND 7` gives zero. This saves 1 byte and 3 T-states in the boundary-crossing path:

```
; UP_HL optimised (Art-top)
; Saves 1 byte, 3 T-states on boundary crossing
;
up_hl_opt:
    dec h            ; 4T
    ld a, h          ; 4T
    cpl              ; 4T    complement: 111 -> 000
    and 7             ; 7T    zero if we crossed boundary
    ret nz            ; 11/5T

    ld a, l           ; 4T
    sub 32            ; 7T
    ld l, a           ; 4T
    ret c             ; 11/5T

    ld a, h           ; 4T
    add a, 8          ; 7T
    ld h, a           ; 4T
    ret               ; 10T
```

The same `CPL / AND 7` trick works in `DOWN_HL` too, though the boundary condition there tests for 000 (which `CPL` turns into 111, also non-zero after `AND`), so it does not help going downward. It is specifically the *upward* direction where the classic code needs the extra `CP 7` that the optimisation eliminates.

Computing the attribute address from a pixel address

If `HL` points to a byte in the pixel area, the corresponding attribute address can be calculated. Recall the pixel address structure: `H = 010TTSSS`, `L = LLLCCCC`. The attribute address for the same character cell is `$5800 + TT * 256 + LLL * 32 + CCCCC`. Since `L` already encodes `LLL * 32 + CCCCC` (which ranges 0-255), the attribute address is simply `($58 + TT) : L`. All we need to do is extract the two `TT` bits from `H`, combine them with `$58`, and leave `L` unchanged:

```
; Convert pixel address in HL to attribute address in HL
; Input:  HL = pixel address ($4000-$57FF)
; Output: HL = corresponding attribute address ($5800-$5AFF)
;
    ld a, h          ; 4T
    rrca              ; 4T
    rrca              ; 4T
    rrca              ; 4T
    and 3             ; 7T
    or $58            ; 7T
    ld h, a           ; 4T
    ; L unchanged    --- Total: 34T
```

This works because L already contains LLL CCCCC – the character row within the third (0-7) combined with the column (0-31) – and that is exactly the low byte of the attribute address. The high byte just needs the third number added to \$58. Elegant.

Special case: when H has scanline bits = 111. If you are iterating through a character cell top-to-bottom and have just processed the last scanline (scanline 7), the low 3 bits of H are 111. In this case there is a faster 4-instruction conversion, contributed by Art-top:

```
; Pixel-to-attribute when H low bits are %111
; (e.g., after processing the last scanline of a character cell)
; Input:  HL where H = 010TT111
; Output: HL = attribute address
;
    srl  h          ; 8T   010TT111 -> 0010TT11
    rrc  h          ; 8T   0010TT11 -> 10010TT1
    srl  h          ; 8T   10010TT1 -> 010010TT
    set  4, h       ; 8T   010010TT -> 010110TT = $58+TT
; L unchanged.    --- Total: 32T, 4 instructions
```

This is 2 T-states faster than the general method and avoids the AND / OR sequence. The trade-off is that it only works when the scanline bits are 111 – but that is exactly the situation after a top-to-bottom character cell render loop, which is one of the most common use cases.

Agon Light 2 Sidebar

The Agon Light 2's display is managed by a VDP (Video Display Processor) – an ESP32 microcontroller running the FabGL library. The eZ80 CPU communicates with the VDP over a serial link, sending commands to set graphics modes, draw pixels, define sprites, and manage palettes.

There is no interleaved memory layout. There is no attribute clash. The VDP supports multiple bitmap modes at various resolutions (from 640x480 down to 320x240 and below), with 64 colours or full RGBA palettes depending on the mode. Hardware sprites (up to 256) and tile maps are supported natively.

What changes for the programmer:

- **No address puzzle.** Pixel coordinates map linearly to buffer positions. You do not need DOWN_HL or split-counter screen traversal.
- **No attribute clash.** Each pixel can be any colour. The 8x8 grid constraint does not exist.
- **No direct memory access to the framebuffer.** The CPU cannot write directly to video memory the way a Spectrum CPU writes to \$4000. Instead, you send VDP commands over the serial link. Drawing a pixel means sending a command sequence, not storing a byte. This introduces latency – the serial link runs at 1,152,000 baud – but it also means the CPU is free during rendering.
- **No cycle-level border tricks.** The VDP handles display timing independently. You cannot create raster effects by timing OUT instructions, because the display pipeline is decoupled from the CPU clock.

For a Spectrum programmer, the Agon feels freeing and frustrating in equal measure. The constraints that forced creative solutions on the Spectrum simply do not exist – but neither do the direct-hardware tricks that those constraints enabled. You trade the puzzle for an API.

Putting It Together: What the Screen Layout Means for Code

Every technique in the rest of this book is shaped by the screen layout described in this chapter. Here is why each piece matters:

Sprite drawing requires computing a screen address for the sprite's position, then iterating down through the sprite's rows. Each row means `INC H` (7 out of 8 times) or the full character-boundary crossing. A 16-pixel-tall sprite spans exactly 2 character cells – you will cross one boundary. A 24-pixel sprite spans 3 cells, crossing 2 boundaries. The boundary-crossing cost is a fixed tax on every sprite.

Screen clearing (Chapter 3) uses the `PUSH` trick – setting `SP` to `$5800` and pushing data downward through the pixel area. The interleave does not matter for clearing because every byte gets the same value. But for *patterned* clears (striped backgrounds, gradient fills), the interleave means you must think carefully about which rows get which data.

Scrolling (Chapter 17) is where the layout hurts most. Scrolling the screen up by one pixel means moving each row's 32 bytes to the address of the row above it. On a linear framebuffer, this is one big block copy. On the Spectrum, the source and destination addresses for each row are related by the `DOWN_HL` logic – not by a fixed offset. A scroll routine must navigate the interleave for every row it copies.

Attribute effects (Chapters 8–9) are where the layout helps. Because the attribute area is linear and small (768 bytes), updating colours is fast. A full-screen attribute update with `LDIR` costs about 16,000 T-states – less than a quarter of a frame. This is why attribute-based effects (tunnels, plasmas, colour cycling) are a staple of Spectrum demoscene work.

Summary

- The Spectrum's 6,912-byte display consists of **6,144 bytes of pixel data** at `$4000–$57FF` and **768 bytes of attributes** at `$5800–$5AFF`.
- Pixel rows are **interleaved** by character cell: the address encodes `y` as `010 TT SSS` (high byte) and `LLL CCCCC` (low byte), where the bits of `y` are shuffled across the address.
- Moving **one pixel row down** within a character cell is just `INC H` (4 T-states). Crossing character and third boundaries requires additional logic.
- The classic `DOWN_HL` routine handles all cases but costs up to 65 T-states at boundaries. For full-screen iteration, **split-counter loops** (Introspec's approach) reduce total cost by 60% and eliminate timing jitter.
- Each attribute byte encodes **Flash, Bright, Paper, and Ink** in the format `FBPPPIII`. Only **two colours per 8x8 cell** – this is the attribute clash.

- Attribute clash is not just a limitation but a **creative constraint** that defined the Spectrum's visual aesthetic and led to efficient attribute-only demo effects.
 - The **border** colour is set by OUT (\$FE), A (bits 0-2) and changes are visible on the next scanline, making it a **timing debug tool** and a canvas for demoscene raster effects.
 - The **Agon Light 2** has no interleaved layout, no attribute clash, and no direct framebuffer access – it replaces the puzzle with a VDP command API.
-

Try It Yourself

1. **Map the addresses.** Pick 10 random (x, y) coordinates and calculate the screen address by hand using the 010TTSSS LLLCCCC bit layout. Then write a small Z80 routine that plots a single pixel at each coordinate and verify your calculations match.
 2. **Visualise the interleave.** Modify fill_screen.a80 to write different values to the first 8 rows. Write \$FF (solid) to row 0 and \$00 (empty) to rows 1-7. Because rows 0-7 are at \$4000, \$4100, ..., \$4700, you will need to change H to reach each row. The result should be a single bright line at the very top, with a gap of 7 empty lines before the next solid line at row 8.
 3. **Time DOWN_HL.** Use the border-colour timing harness from Chapter 1. Call the classic DOWN_HL routine 191 times (for a full screen traversal) and measure the stripe. Then implement the split-counter version and compare. The split-counter version should produce a visibly shorter stripe.
 4. **Attribute painter.** Write a routine that fills the attribute area with a gradient: column 0 gets colour 0, column 1 gets colour 1, and so on (cycling through 0-7). Each row should have the same pattern. Then modify it so each row shifts the pattern by one position – a diagonal rainbow. This is the seed of an attribute-based demo effect.
 5. **Border stripes.** After a HALT, execute a tight loop that changes the border colour on every scanline for 64 lines. Use the 8 border colours in sequence (0, 1, 2, 3, 4, 5, 6, 7, repeat). You should see horizontal rainbow stripes in the top border. Adjust the timing delay between OUT instructions until the stripes are clean and stable.
-

Sources: Introspec “Eshchy raz pro DOWN_HL” (Hype, 2020); Introspec “GO WEST Part 1” (Hype, 2015) for contended memory effects at screen addresses; Introspec “Making of Eager” (Hype, 2015) for attribute-based effect design; the Spectrum's ULA documentation for memory layout rationale; Art-top (personal communication, 2026) for the optimised UP_HL and fast pixel-to-attribute conversion.

Next: Chapter 3 – The Demoscener's Toolbox. Unrolled loops, self-modifying code, the stack as a data pipe, and the techniques that let you do the impossible within the budget.

Chapter 3: The Demoscener's Toolbox

Every craft has its bag of tricks — patterns that practitioners reach for so instinctively they stop thinking of them as tricks at all. A Z80 demoscener reaches for the techniques in this chapter.

These patterns — unrolled loops, self-modifying code, the stack as a data pipe, LDI chains, code generation, and RET-chaining — appear in almost every effect we will build in Part II. They are what separates a demo that fits in one frame from one that takes three. Learn them here, and you will recognise them everywhere.

Unrolled Loops and Self-Modifying Code

The cost of looping

Consider the simplest possible inner loop: clearing 256 bytes of memory.

```
; Looped version: clear 256 bytes at (HL)
    ld  b, 0                ; 7 T   (B=0 means 256 iterations)
    xor a                  ; 4 T
.loop:
    ld  (hl), a             ; 7 T
    inc hl                 ; 6 T
    djnz .loop              ; 13 T  (8 on last iteration)
```

Each iteration costs $7 + 6 + 13 = 26$ T-states to store a single byte. Only 7 of those T-states do the work — the rest is overhead. That is 73% waste. For 256 bytes: $256 \times 26 - 5 = 6,651$ T-states. On a machine where you have 71,680 T-states per frame, those wasted T-states hurt.

Unrolling: trade RAM for speed

The solution is brutal and effective: write out the loop body N times and delete the loop.

```
; Unrolled version: clear 256 bytes at (HL)
    xor a                  ; 4 T
    ld  (hl), a            ; 7 T
    inc hl                 ; 6 T
    ld  (hl), a            ; 7 T
```



```

inc hl          ; 6 T
ld  (hl), a     ; 7 T
inc hl          ; 6 T
; ... repeated 256 times total

```

Each byte now costs $7 + 6 = 13$ T-states. No DJNZ. No loop counter. Total: $256 \times 13 = 3,328$ T-states — half the looped version.

The cost is code size: 256 repetitions occupy 512 bytes vs. 7 for the loop. You are trading RAM for speed.

When to unroll: Inner loops that execute thousands of times per frame — screen clearing, sprite drawing, data copying.

When NOT to unroll: Outer loops that run once or twice per frame. Saving 5 T-states across 24 iterations buys you 120 T-states — less than three NOPs. Not worth the bloat.

The practical middle ground is *partial unrolling*: unroll 8 or 16 iterations inside the loop, keep DJNZ for the outer count. The `push_fill.a80` example in this chapter's `examples/` directory does exactly this: 16 PUSHes per iteration, 192 iterations.

Self-modifying code: the Z80's secret weapon

The Z80 has no instruction cache, no prefetch buffer, no pipeline. When the CPU fetches an instruction byte from RAM, it reads whatever is there *right now*. If you changed that byte one cycle ago, the CPU sees the new value. This is a guaranteed property of the architecture.

Self-modifying code (SMC) means writing to instruction bytes at runtime. The classic pattern is patching an immediate operand:

```

; Self-modifying code: fill with a runtime-determined value
ld  a, (fill_value)      ; load the fill byte from somewhere
ld  (patch + 1), a        ; overwrite the operand of the LD below
patch:
ld  (hl), $00             ; this $00 gets replaced at runtime
inc hl
; ...

```

The `ld (patch + 1), a` writes into the immediate operand of the next `ld (hl), $00`, changing it to `ld (hl), $AA` or whatever you loaded. The CPU executes whatever bytes it finds. Some common SMC patterns:

Patching opcodes. You can even replace the instruction itself. Need a loop that sometimes increments HL and sometimes decrements it? Before the loop, write the opcode for INC HL (\$23) or DEC HL (\$2B) into the instruction byte. Inside the inner loop, there is no branch at all — the right instruction is already in place. Compare this to a branch-per-iteration approach that would cost 12 T-states (JR NZ) on every single pixel.

Saving and restoring the stack pointer. This pattern appears constantly when using PUSH tricks (below):

```

ld  (restore_sp + 1), sp    ; save SP into the operand below
; ... do stack tricks ...

```

```
restore_sp:
    ld    sp, $0000                ; self-modified: the $0000 was overwritten
```

The `ld (nn), sp` saves the current SP directly into the operand of the later `ld sp, nn`. No temporary variable. This is idiomatic Z80 demoscene code.

Self-modifying variables: the \$+1 pattern

The most pervasive SMC pattern on the ZX Spectrum is not patching opcodes or saving SP — it is embedding a *variable* directly inside an instruction's immediate operand. The idea is simple: instead of storing a counter in a named memory location and loading it with `LD A, (nn)` at 13 T-states, you let the instruction's own operand byte *be* the variable.

```
.smc_counter:
    ld    a, 0                    ; 7T — this 0 is the "variable"
    inc   a                      ; 4T
    ld    (.smc_counter + 1), a   ; 13T — write back to the operand byte
```

The `ld a, 0` fetches its operand as part of the normal instruction decode — 7 T-states total, and the value is already in A. Compare that to loading from a separate memory address: `ld a, (counter)` costs 13 T-states, plus you still need a separate `ld (counter), a` at 13 T-states to write it back. The SMC version reads the variable for free (it is part of the instruction fetch) and only pays the 13 T-states once for the write-back.

In `sjasmpplus`, you can place a label at `$+1` to give the embedded variable a readable name:

```
    ld    a, 0                    ; 7T
.scroll_pos EQU $ - 1            ; .scroll_pos names the operand byte above
    add   a, 4                    ; 7T — advance by 4 pixels
    ld    (.scroll_pos), a        ; 13T — store back into the operand
```

This pattern appears everywhere in ZX Spectrum code: scroll positions, animation frame counters, effect phase accumulators, direction flags. Any single-byte value that persists between calls is a candidate. You will see it constantly in Parts II and V — practically every effect routine in this book uses at least one self-modifying variable.

The convention is to prefix these labels with `.smc_` or place them immediately after the instruction they modify. Either way, the intent should be clear to anyone reading the source. As we noted in Chapter 2, local labels (`.label`) prevent naming collisions when multiple routines each have their own embedded variables.

A word of caution. SMC is safe on the Z80, the eZ80, and every Spectrum clone. It is *not* safe on modern cached CPUs (x86, ARM) without explicit cache flush instructions. If you port to a different architecture, this is the first thing that breaks.

The Stack as a Data Pipe

Why PUSH is the fastest write on the Z80

The PUSH instruction writes 2 bytes to memory and decrements SP, all in 11 T-states. Let us compare the alternatives for writing data to a screen address:

Method	Bytes written	T-states	T-states per byte
ld (hl), a + inc hl	1	13	13.0
ld (hl), a + inc l	1	11	11.0
ldi	1	16	16.0
ldir (per byte)	1	21	21.0
push hl	2	11	5.5

PUSH writes two bytes in 11 T-states — 5.5 T-states per byte. Nearly 4x faster than LDIR. The catch: PUSH writes to where SP points, and SP is normally your stack. To use PUSH as a data pipe, you must hijack the stack pointer.

The technique

The pattern is always the same:

1. Disable interrupts (DI). If an interrupt fires while SP points to the screen, the CPU will push the return address into your pixel data. Chaos follows.
2. Save SP. Use self-modifying code to stash it.
3. Set SP to the *end* of your target area. The stack grows downward — PUSH decrements SP before writing. So if you want to fill from \$4000 to \$57FF, you set SP to \$5800.
4. Load your data into register pairs and PUSH repeatedly.
5. Restore SP and re-enable interrupts (EI).

graph TD

```

A["DI – disable interrupts"] --> B["Save SP via self-modifying code"]
B --> C["Set SP to screen bottom ($5800)"]
C --> D["Load register pairs with fill data"]
D --> E["PUSH loop: 16 PUSHes per iteration<br>11T × 16 = 176T → 32 bytes"]
E --> F{All 192<br>iterations done?}
F -- No --> E
F -- Yes --> G["Restore SP from self-modified LD SP,nn"]
G --> H["EI – re-enable interrupts"]

```

```

style E fill:#f9f,stroke:#333
style A fill:#fdd,stroke:#333
style H fill:#fdd,stroke:#333

```

Why PUSH wins: LD (HL),A + INC HL writes 1 byte in 13T (13.0 T/byte). PUSH HL writes 2 bytes in 11T (**5.5 T/byte**) — nearly 2.4× faster per byte. The cost: interrupts must be disabled while SP is hijacked.

Here is the core of the push_fill.a80 example from this chapter's examples/ directory:

stack_fill:

```
di ; critical: no interrupts while SP is moved
```

```

ld    (restore_sp + 1), sp    ; self-modifying: save SP

ld    sp, SCREEN_END        ; SP points to end of screen ($5800)
ld    hl, $AAAA             ; pattern to fill

ld    b, 192                 ; 192 iterations x 16 PUSHes x 2 bytes = 6144
.loop:
push hl                      ; 11 T  \
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  | 16 PUSHes = 32 bytes
push hl                      ; 11 T  | = 176 T-states
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  |
push hl                      ; 11 T  /
djnz .loop                   ; 13 T (8 on last)

restore_sp:
ld    sp, $0000             ; self-modified: restores original SP
ei
ret

```

The 16-PUSH inner body writes 32 bytes in 176 T-states. Total for the full 6,144-byte pixel area: roughly 36,000 T-states. Compare LDIR: $6,144 \times 21 - 5 = 129,019$ T-states. The PUSH method is about 3.6x faster — the difference between fitting in one frame and bleeding into the next.

PUSH-based screen fill — the entire pixel area filled in a single frame using the stack trick

Figure 9: PUSH-based screen fill — the entire pixel area filled in a single frame using the stack trick

POP as a fast read

PUSH is the fastest write, but POP is the fastest *read*. POP loads 2 bytes from (SP) into a register pair in 10 T-states — that is 5.0 T-states per byte. Compare the alternatives:

Method	Bytes read	T-states	T-states per byte
ld a, (hl) + inc hl	1	13	13.0
ld a, (hl) + inc l	1	11	11.0
ldi (as a read+write)	1	16	16.0

Method	Bytes read	T-states	T-states per byte
pop hl	2	10	5.0

The pattern: pre-build a table of 16-bit values in memory, point SP at the start of the table, and POP into register pairs. Each POP advances SP by 2, walking through the table automatically. This is the read-side complement of the PUSH write trick.

Combine POP and PUSH and you get a fast memory-to-memory pipe: POP a value from a source table (10T), process the register pair if needed, then PUSH it to the destination (11T). Total: 21 T-states for 2 bytes — the same throughput as LDIR, but with the register pair available for processing between the read and write. You can mask bits, add offsets, swap bytes, or apply any register-to-register transformation at no extra memory-access cost. This POP-process-PUSH pipeline is the backbone of many compiled sprite routines.

Where PUSH tricks are used

- **Screen clearing.** The most common use. Every demo needs to clear the screen between effects.
- **Compiled sprites.** The sprite is compiled into a sequence of PUSH instructions with pre-loaded register pairs. The fastest possible sprite output on the Z80.
- **Fast data output.** Any time you need to blast a block of data to a contiguous address range: attribute fills, buffer copies, display list construction.

The price you pay: interrupts are off. If your music player runs from an IM2 interrupt, it will miss a beat during a long PUSH sequence. Demo coders plan around this — schedule PUSH fills during border time, or split them across multiple frames.

LDI Chains

LDI vs LDIR

LDI copies one byte from (HL) to (DE), increments both, and decrements BC. LDIR does the same but repeats until BC = 0. The difference is timing:

Instruction	T-states	Notes
LDI	16	Copies 1 byte, always 16 T
LDIR (per byte)	21	Copies 1 byte, loops back. Last byte: 16 T

LDIR costs 5 extra T-states per byte for its internal loop-back check. Those 5 T-states add up fast.

For 256 bytes: - LDIR: $255 \times 21 + 16 = 5,371$ T-states - 256 x LDI: $256 \times 16 = 4,096$ T-states - Savings: 1,275 T-states (24%)

A chain of individual LDI instructions is just 256 repetitions of the two-byte opcode \$ED \$A0. That is 512 bytes of code to save 24% — the same RAM-for-speed trade-off as loop unrolling.

When LDI chains shine

The sweet spot is copying blocks of known size. A chain of 32 LDIs saves 160 T-states over LDIR for a sprite row. Across 24 rows, that is 3,840 T-states per frame.

But the real power emerges when you combine LDI chains with *entry point arithmetic*. If you have a chain of 256 LDIs and want to copy only 100 bytes, jump into the chain at position 156. No loop counter, no setup. This technique is used in Introspec's chaos zoomer in Eager (2015):

```
; Chaos zoomer inner loop (simplified from Eager)
; Each line copies a different number of bytes from a source buffer.
; Entry point into the LDI chain is calculated per line.
    ld    hl, source_data
    ld    de, dest_screen
    ; ... calculate entry point based on zoom factor ...
    jp    (ix)                ; jump into the LDI chain at the right point

ldi_chain:
    ldi                    ; byte 255
    ldi                    ; byte 254
    ldi                    ; byte 253
    ; ... 256 LDIs total ...
    ldi                    ; byte 0
    ; falls through to next line setup
```

This variable-length copy with zero per-byte loop overhead is a technique you simply cannot achieve with LDIR. It is one reason LDI is everyone's best friend in demoscene code.

LDI chain vs LDIR — red stripes show LDI chain timing, blue stripes show LDIR; thinner red proves LDI is faster

Figure 10: LDI chain vs LDIR — red stripes show LDI chain timing, blue stripes show LDIR; thinner red proves LDI is faster

Bit Tricks: SBC A,A and Friends

SBC A,A as a conditional mask

After any instruction that produces a carry flag, SBC A,A converts that flag into a full byte: \$FF if carry was set, 00 if not. The cost: $4T$ — states. Compare this to the branching alternative — `JRC, .set ^ LDA, 0 ^ JR.done ^ .set : LDA, FF /.done:` — which costs 17-22 T-states depending on which path is taken, plus the pipeline disruption of a conditional branch.

The canonical use case is *bit-to-byte expansion*. Given a byte where each bit represents a pixel (the Spectrum's pixel format), you can expand each bit into a

full attribute byte:

```
rlc  (hl)          ; rotate top bit into carry    - 15T
sbc  a, a          ; A = $FF if set, $00 if not    - 4T
and  $47           ; A = bright white ($47) or $00 - 7T
```

Three instructions, 26 T-states, no branches. To select between two *arbitrary* values rather than zero and a mask, use the pattern SBC A,A : AND mask : XOR base. The AND selects which bits change between the two values, and the XOR flips them to the desired base. This pattern replaces every “if bit set then value A else value B” test in your inner loops.

ADD A,A vs SLA A

Both instructions shift A left by one bit. But ADD A,A is 4 T-states and 1 byte, while SLA A is 8 T-states and 2 bytes. There is no situation where SLA A is preferable — ADD A,A is strictly faster and smaller. Similarly, ADD HL,HL shifts HL left in 11 T-states (1 byte), replacing the two-instruction sequence SLA L : RL H at 16 T-states (4 bytes). For a 16-bit left shift inside an inner loop running 192 times per frame, that substitution alone saves 960 T-states — over four scanlines of border time.

These are not tricks. They are vocabulary. Just as a fluent speaker does not pause to conjugate common verbs, a Z80 programmer reaches for ADD A,A and SBC A,A without conscious thought. If you find yourself writing SLA A or a conditional branch to select between two values, stop and reach for the shorter form. The T-states add up.

Code Generation

Code generation: writing the program that draws

Everything above is a fixed optimisation — the code runs the same way every frame. Code generation goes further: your program writes the program that draws the screen. There are two variants: offline (before assembly) and runtime (during execution).

Offline: generating assembly from a higher-level language

Introspec used Processing (a Java-based creative coding environment) to generate Z80 assembly for the chaos zoomer in Eager (2015). A chaos zoomer changes magnification every frame — different source pixels map to different screen locations. Rather than computing these mappings at runtime, the Processing script pre-calculated every mapping and output .a80 source files containing optimised LDI chains and LD instructions.

The workflow: a Processing script calculates, for each frame, which source byte maps to which screen byte. It outputs Z80 assembly source — sequences of ld hl, source_addr and ldi instructions — which the assembler (sjasmplus) builds alongside the hand-written engine code. At runtime, the engine simply calls into the pre-generated code for the current frame.

This is not “cheating.” It is the fundamental insight that division of labour between compile time and runtime can eliminate branches, lookups, and arithmetic from the inner loop entirely. The Processing script does the hard maths once, slowly, on a modern machine. The Z80 does the easy part — copying bytes — as fast as physically possible.

Runtime: the program writes machine code during execution

Sometimes parameters change every frame, so offline generation is not enough. The sphere-mapping routine in X-Trade’s Illusion (ENLiGHT’96) generates machine code into a RAM buffer at runtime. The sphere geometry changes as it rotates — different pixels need different skip distances. Before each frame, the engine emits opcode bytes into a buffer, then executes them:

```
; Runtime code generation (conceptual, simplified from Illusion)
; Generate an unrolled rendering loop for this frame's sphere slice

    ld    hl, code_buffer
    ld    de, sphere_table        ; per-frame skip distances

    ld    b, SPHERE_WIDTH
.gen_loop:
    ld    a, (de)                ; load skip distance for this pixel
    inc   de

    ; Emit: ld a, (hl) -- opcode $7E
    ld    (hl), $7E
    inc   hl

    ; Emit: add a, N -- opcodes $C6, N
    ld    (hl), $C6
    inc   hl
    ld    (hl), a                ; the skip distance, as immediate operand
    inc   hl

    djnz  .gen_loop

    ; Emit: ret -- opcode $C9
    ld    (hl), $C9

    ; Now execute the generated code
    call  code_buffer
```

The generated code is a straight-line sequence with no branches, no lookups, no loop overhead — but it is *different code every frame*. Instead of “if pixel_skip == 3 then...” at 12 T-states per branch, you emit the exact instructions needed and execute them branch-free.

The cost of generation

Runtime code generation is not free. Look at the generator loop above: each emitted instruction requires loading an opcode byte, storing it, advancing the

pointer, and possibly loading an operand — roughly 30-50 T-states per emitted byte, depending on complexity. Call it ~40 T-states on average. For a generated routine of 100 instruction bytes, that is about 4,000 T-states of generation overhead.

The break-even point: generation pays off when the generated code runs more than once per frame, or when it replaces branching logic that costs more than the generation itself. In the Illusion sphere mapper, each generated rendering pass executes once per frame — but it replaces per-pixel conditional branches that would cost far more. AloneCoder documented a similar trade-off in his rotation engine: generating a sequence of INC H/INC L instructions for coordinate stepping costs roughly 5,000 T-states to emit, but eliminates coordinate arithmetic that would cost approximately 146,000 T-states if computed inline. The generation overhead is under 4% of the cost it replaces.

The rule of thumb: if you find yourself writing a loop that contains branches selecting between different instruction sequences based on per-pixel or per-line data, that loop is a candidate for code generation. Emit the right instructions once, execute them branch-free, and regenerate only when the parameters change.

When to generate code: If the same operations happen every frame with only data changes, self-modifying code (patching operands) suffices. If the *structure* changes — different numbers of iterations, different instruction sequences — generate the code. If you can precompute the variations on a modern machine, prefer offline generation: it is debuggable, verifiable, and imposes zero runtime cost. Runtime generation pays off when the generated code executes far more often than it costs to generate.

RET-Chaining

Turning the stack into a dispatch table

In 2025, DenisGrachev published a technique on Hype developed for his game Dice Legends. The problem: rendering a tile-based playfield requires drawing dozens of tiles per frame. The naive approach uses CALL:

```
; Naive approach: call each tile renderer
    call draw_tile_0
    call draw_tile_1
    call draw_tile_2
    ; ...
```

Each CALL costs 17 T-states. For a 30 x 18 playfield (540 tiles), that is 9,180 T-states on dispatch alone.

DenisGrachev's insight: set SP to a *render list* — a table of addresses — and end each tile-drawing procedure with RET. RET pops 2 bytes from (SP) into PC. If SP points to your render list, RET does not return to the caller — it jumps to the next routine in the list.

```
; RET-chaining: zero call overhead
    di
    ld  (restore_sp + 1), sp    ; save SP
    ld  sp, render_list        ; SP points to our dispatch table
```

```

; "Call" the first tile routine by falling into it or using RET:
ret                                ; pops first address from render_list

; Each tile routine ends with:
draw_tile_N:
; ... draw the tile ...
ret                                ; pops NEXT address from render_list

; The render list is a sequence of addresses:
render_list:
    dw  draw_tile_42                ; first tile to draw
    dw  draw_tile_7                 ; second tile
    dw  draw_tile_42                ; third tile (same tile type, different
↪ position)
    ; ... one entry per tile on screen ...
    dw  render_done                 ; sentinel: address of cleanup code

render_done:
restore_sp:
    ld  sp, $0000                    ; self-modified: restore SP
    ei

```

Each dispatch now costs 10 T-states (RET) instead of 17 (CALL). For 540 tiles: 3,780 T-states saved. But the real gain is free dispatch — each entry can point to a different procedure (wide tile, blank tile, animated tile). No jump table, no indirect call. The render list *is* the program.

Three strategies for the render list

DenisGrachev explored three approaches to constructing the render list:

1. **Map as render list.** The tilemap itself is the render list: each cell contains the address of the drawing routine for that tile type. Simple but inflexible — changing a tile means rewriting 2 bytes in the map.
2. **Address-based segments.** The screen is divided into segments. Each segment's render list is a block of addresses copied from a master table. Changing tiles means copying a new address block.
3. **Byte-based with 256-byte lookup tables.** Each tile type is a single byte (the tile index). A 256-byte lookup table maps tile indices to routine addresses. The render list is built by iterating through the tilemap bytes and looking up each address. This is the approach DenisGrachev chose for Dice Legends.

Using the byte-based approach, he expanded the playfield from 26 x 15 tiles (the limit of his previous engine) to 30 x 18 tiles while maintaining the target frame rate. The savings from eliminating CALL overhead, combined with the zero-cost dispatch, freed enough T-states to render 40% more tiles.

The trade-offs

Like all stack tricks, interrupts must be disabled while SP is hijacked. Each tile routine must be self-contained — ending with RET and not using CALL, since the

real stack is unavailable. In practice, tile routines are short enough that this is not a limitation.

Sidebar: “Code is Dead” (Introspec, 2015)

In January 2015, Introspec published a short, provocative essay on Hype titled “Code is Dead” (Kod myortv). The argument draws a parallel to Roland Barthes’ “Death of the Author”: just as Barthes argued that a text’s meaning belongs to the reader, not the writer, Introspec argues that demo code only truly lives when someone reads it — in a debugger, in a disassembly listing, in source code shared on a forum.

The uncomfortable truth: modern demos are consumed as visual media. People watch them on YouTube. They vote on Pouet based on video captures. Nobody sees the inner loops. A brilliant optimisation that saves 3 T-states per pixel is invisible to 99% of the audience. “Writing code purely for its own sake,” Introspec wrote, “has lost relevance.”

And yet.

You are reading this book. We are opening the debugger. We are counting T-states. We are looking inside. The techniques in this chapter are not museum exhibits. They are living tools, and the fact that most people will never see them does not diminish their craft.

Introspec’s essay is a challenge, not a surrender. He went on to publish some of the most detailed technical analyses the ZX scene has ever produced — including the Illusion teardown and the compression benchmarks referenced throughout this book. Code may be dead to the YouTube viewer. But to the reader with a disassembler and a curious mind, it is very much alive.

Putting It All Together

The techniques in this chapter are not independent. In practice, they compose:

- **Screen clearing** combines *unrolled loops* with *PUSH tricks*: a partially unrolled loop of 16 PUSHes per iteration, with SP hijacked via *self-modifying code*.
- **Compiled sprites** combine *code generation* (each sprite compiles to executable code), *POP reads* and *PUSH output* (the fastest way to move pixel data through registers), *bit tricks* (SBC A,A for mask expansion), and *self-modification* (patching screen addresses per frame).
- **Tile engines** combine *RET-chaining* for dispatch with *LDI chains* inside each tile routine for fast data copy.
- **Chaos zoomers** combine *offline code generation* (Processing scripts emitting assembly) with *LDI chains* (the generated code is mostly LDI sequences) and *self-modification* (patching source addresses per frame).
- **Attribute effects** combine *POP reads* from pre-computed tables with *bit tricks* (SBC A,A to expand bitmasks into colour values) and *PUSH writes* for fast

attribute output.

The common thread: every technique eliminates something from the inner loop. Unrolling eliminates the loop counter. Self-modification eliminates branches. PUSH eliminates per-byte write overhead. POP eliminates per-byte read overhead. LDI chains eliminate the LDIR repeat penalty. Bit tricks eliminate conditional branches. Code generation eliminates the entire distinction between code and data. RET-chaining eliminates CALL overhead.

The Z80 runs at 3.5 MHz. You have 71,680 T-states per frame. Every T-state you save in the inner loop is a T-state you can spend on more pixels, more colours, more motion. The toolbox in this chapter is how you get there.

In the chapters that follow, you will see each of these techniques at work in real demos — the textured sphere of Illusion, the attribute tunnel of Eager, the multicolor engine of Old Tower. The goal of this chapter was to give you the vocabulary. Now let us see what the masters built with it.

Try It Yourself

1. **Measure the difference.** Take the timing harness from Chapter 1 and measure three versions of a 256-byte fill: (a) the `ld (hl), a : inc hl : djnz` loop, (b) a fully unrolled `ld (hl), a : inc hl x 256`, and (c) the PUSH-based fill from `examples/push_fill.a80`. Compare the border stripe widths. The PUSH version's stripe should be visibly shorter.
2. **Build a self-modifying clear.** Write a screen-clear routine that takes the fill pattern as a parameter and patches it into a PUSH-based fill loop using self-modifying code. Call it twice with different patterns and watch the screen alternate.
3. **Time an LDI chain.** Write a 32-byte copy using LDIR and another using 32 x LDI. Measure both with the border-colour technique. The LDI chain should save 160 T-states — visible if you run the copy in a tight loop.
4. **Experiment with entry points.** Build a 128-entry LDI chain and a small routine that calculates an entry point based on a value in register A (0-128). Jump into the chain at different points. This is a simplified version of the variable-length copy used in real chaos zoomers.
5. **Variable-length copier with calculated entry.** Build a 256-entry LDI chain and a front-end that accepts a byte count in register B (1-256). Calculate the entry point: each LDI is 2 bytes, so the offset is $(256 - B) \times 2$ from the start of the chain. Add this to the chain's base address, then JP (HL) into it. Wrap the whole thing in the border-colour harness and compare the stripe width against LDIR for the same byte count. For small counts (under 16), the difference is slim. For counts above 64, the LDI chain pulls visibly ahead.
6. **Bit-to-attribute unpacker.** Write a routine that reads a byte from (HL), rotates each bit out with RLC (HL), and uses `SBC A,A : AND $47` to expand each bit into an attribute byte (bright white or black). Store the 8 resulting attribute bytes to a destination buffer using `(DE) / INC DE`. This is the seed of a compiled

sprite's attribute writer — in later chapters you will see this pattern generate entire sprite routines.

Sources: DenisGrachev “Tiles and RET” (Hype, 2025); Introspec “Making of Eager” (Hype, 2015); Introspec “Technical Analysis of Illusion” (Hype, 2017); Introspec “Code is Dead” (Hype, 2015)

Chapter 4: The Maths You Actually Need

“Read a maths textbook – derivatives, integrals. You will need them.” –
Dark, Spectrum Expert #01 (1997)

In 1997, a teenager in St. Petersburg sat down to write a magazine article about multiplication. Not the kind you learn in school – the kind that makes a wireframe cube spin on a ZX Spectrum at 50 frames per second. His name was Dark, he coded for the group X-Trade, and his demo *Illusion* had already won first place at ENLiGHT’96. Now he was writing *Spectrum Expert*, an electronic magazine distributed on floppy disk, and he was going to explain exactly how his algorithms worked.

What follows is drawn directly from Dark’s “Programming Algorithms” article in Spectrum Expert #01. These are the routines that powered *Illusion* – the same multiply that rotated vertices, the same sine table that drove the rotozoomer, the same line drawer that rendered wireframes at full frame rate. When Introspec reverse-engineered *Illusion* twenty years later on the Hype blog, he found these exact algorithms at work inside the binary.

Multiplication on Z80

The Z80 has no multiply instruction. Every time you need A times B – for rotation matrices, perspective projection, texture mapping – you must synthesize it from shifts and adds. Dark presents two methods, and he is characteristically honest about the trade-off between them.

Method 1: Shift-and-Add from LSB

The classic approach. Scan through the bits of the multiplier from LSB to MSB. For each set bit, add the multiplicand into an accumulator. After each bit, shift the accumulator right. After eight iterations, the accumulator holds the full product.

Here is Dark’s 8x8 unsigned multiply. Input: B times C. Result in A (high byte) and C (low byte):

```
; MULU112 -- 8x8 unsigned multiply
; Input:  B = multiplicand, C = multiplier
; Output: A:C = B * C (16-bit result, A=high, C=low)
; Cost:   196-204 T-states (Pentagon)
```

```

;
; From Dark / X-Trade, Spectrum Expert #01 (1997)

mulu112:
    ld    a, 0            ; clear accumulator (high byte of result)
    ld    d, 8            ; 8 bits to process

.loop:
    rr    c                ; shift LSB of multiplier into carry
    jr    nc, .noadd       ; if bit was 0, skip addition
    add   a, b              ; add multiplicand to accumulator
.noadd:
    rra                    ; shift accumulator right (carry into bit 7,
                           ; bit 0 into carry -- this carry feeds
                           ; back into C via the next RR C)

    dec   d
    jr    nz, .loop
    ret

```

Study this carefully. The RRA instruction shifts A right, but also pushes A's lowest bit into the carry flag. On the next iteration, RR C rotates that carry into the top of C. So the low bits of the product gradually assemble in C, while the high bits accumulate in A. After eight iterations, the full 16-bit result sits in A:C.

The cost is 196 to 204 T-states depending on how many multiplier bits are set – each set bit costs one extra ADD A,B (4 T-states). The example at [chapters/ch04-maths/examples/multiply8.a80](#) shows a variant returning the result in HL.

For 16x16 producing a 32-bit result, Dark's MULU224 runs in 730 to 826 T-states. In practice, demoscene 3D engines avoid full 16x16 multiplies by keeping coordinates in 8.8 fixed-point and using 8x8 multiplies where possible.

Method 2: Square Table Lookup

Dark's second method trades memory for speed, exploiting an algebraic identity that every demoscener eventually discovers:

Pre-compute a table of $n^2/4$ values, and multiplication becomes two lookups and a subtraction – approximately 61 T-states, more than three times faster than shift-and-add.

You need a 512-byte table of $(n^2/4)$ for $n = 0$ to 511, page-aligned for single-register indexing. The table must be 512 bytes because $(A+B)$ can reach 510.

```

; MULU_FAST -- Square table multiply
; Input:  B, C = unsigned 8-bit factors
; Output: HL = B * C (16-bit result)
; Cost:   ~61 T-states (Pentagon)
; Requires: sq_table = 512-byte table of n^2/4, page-aligned
;
; A*B = ((A+B)^2 - (A-B)^2) / 4

```

```

mulu_fast:

```

Shift-and-Add Multiply: 13 x 11 = 143

Multiplier = 13 = 00001101 Multiplicand = 11 = 00001011

Step	Bit	Multiplier bits	Action	Acc	Acc (binary)
0	b0 = 1	0000110 1	ADD 11	11	00001011
1	b1 = 0	000011 0 1	skip	11	00001011
2	b2 = 1	0000 1 101	ADD 44	55	00110111
3	b3 = 1	0000 1 101	ADD 88	143	10001111
4	b4 = 0	000 0 1101	skip	143	10001111
5	b5 = 0	00 0 01101	skip	143	10001111
6	b6 = 0	0 0 001101	skip	143	10001111
7	b7 = 0	0 0001101	skip	143	10001111

Result: 13 x 11 = 143 (\$8F = 10001111)

Bit set: ADD multiplicand

Bit clear: skip (no addition)

Figure 11: Shift-and-add 8x8 multiply walkthrough


```

    ld    h, sq_table >> 8    ; high byte of table address
    ld    a, b
    add   a, c                  ; A = B + C (may overflow into carry)
    ld    l, a
    ld    e, (hl)              ; look up (B+C)^2/4 low byte
    inc   h
    ld    d, (hl)              ; look up (B+C)^2/4 high byte

    ld    a, b
    sub   c                    ; A = B - C (may go negative)
    jr    nc, .pos
    neg
    ; take absolute value
.pos:
    ld    l, a
    dec   h
    ld    a, e
    sub   (hl)                 ; subtract (B-C)^2/4 low byte
    ld    e, a
    inc   h
    ld    a, d
    sbc   a, (hl)              ; subtract (B-C)^2/4 high byte
    ld    d, a

    ex    de, hl               ; HL = result
    ret

```

The trade-off? Dark is characteristically honest: **“Choose: speed or accuracy.”** The table stores integer values of $n^2/4$, so there is a rounding error of up to 0.25 per lookup. For large values this is negligible. For the small coordinate deltas in 3D rotation, the error produces visible vertex jitter. With shift-and-add, the rotation is perfectly smooth.

For texture mapping, plasma, scrollers – use the fast multiply. For wireframe 3D where the eye tracks individual vertices – stick with shift-and-add. Dark knew this because he had tried both in *Illusion*.

Generating the square table is a one-time startup cost. Dark suggests using the derivative method: since $d(x^2)/dx = 2x$, you can build the table incrementally by adding a linearly increasing delta at each step. In practice, most coders compute the table in a BASIC loader or initialisation routine and move on.

Signed Multiply

A chapter that teaches unsigned multiplication and then uses it to rotate 3D coordinates has a gap: rotation matrices operate on signed values. X can be -40 or +40, sine values range from -128 to +127. Every call `mul_signed` in Chapter 5 depends on the routine you are about to see. As Ped7g put it during his review: “a chapter that teaches 3D rotation without showing signed multiply is like a cookbook that lists ingredients but forgets the oven.”

Two's Complement in Practice

The Z80 represents signed integers in two's complement. The rules are simple:

- Bit 7 is the sign bit: 0 = positive, 1 = negative
- Positive values are the same as unsigned: \$00 = 0, \$01 = 1, ..., \$7F = 127
- Negative values count down from \$FF: \$FF = -1, \$FE = -2, ..., \$80 = -128
- NEG computes the absolute value of a negative number (negate A: $A = 0 - A$). Cost: 8T

The critical insight for arithmetic: **ADD and SUB do not care about signedness.** Adding \$FF (-1) to \$03 (+3) gives \$02 (+2) — correct in both signed and unsigned interpretation. The hardware addition is identical. Only multiplication requires explicit sign handling, because the shift-and-add loop treats the multiplier bits as unsigned positional values.

Sign Extension: The rla / sbc a,a Idiom

When you multiply an 8-bit signed value by another 8-bit signed value, you need to know the signs. The cheapest way to extract the sign bit on the Z80:

```
; Cost: 8T, 2 bytes. Branchless.
    rla                ; 4T rotate sign bit into carry
    sbc a, a           ; 4T A = 0 if carry clear, $FF if set
```

After `sbc a,a`, A is \$00 for positive inputs or \$FF for negative inputs. This is the standard sign-extension byte used across the Z80 demoscene.

mul_signed — Signed 8×8 Multiply

The algorithm: XOR the two inputs to determine the result sign, take absolute values, multiply unsigned, negate the result if the sign was negative. This is the routine that Chapter 5 calls six times per vertex rotation and twice per backface cull.

```
; mul_signed — 8x8 signed multiply
; Input:  B = signed multiplicand, C = signed multiplier
; Output: HL = signed 16-bit result
; Cost:   ~240-260 T-states (Pentagon)
;
; Algorithm: determine sign, abs both, unsigned multiply, negate if needed.
```

```
mul_signed:
    ld  a, b
    xor c                ; 4T bit 7 = result sign (1 = negative)
    push af              ; 11T save sign flag

    ; Absolute value of B
    ld  a, b
    or  a
    jp  p, .b_pos        ; 10T skip if positive
    neg                ; 8T A = |B|
.b_pos:
    ld  b, a
```

```

; Absolute value of C
ld  a, c
or  a
jp  p, .c_pos
neg
.c_pos:
ld  c, a

; Unsigned 8x8 multiply: B * C -> A:C (high:low)
ld  a, 0
ld  d, 8
.mul_loop:
rr  c
jr  nc, .noadd
add a, b
.noadd:
rra
dec d
jr  nz, .mul_loop

; A:C = unsigned product. Move to HL.
ld  h, a
ld  l, c

; Negate result if sign was negative
pop af          ; 10T recover sign
or  a
jp  p, .done    ; 10T skip if positive
; Negate HL: HL = 0 - HL
xor  a
sub  l
ld  l, a
sbc  a, a
sub  h
ld  h, a
.done:
ret

```

The core is the same shift-and-add loop from `mulu112`, wrapped with sign detection and conditional negation. The overhead is ~40-60 T-states beyond the unsigned multiply, depending on how many operands need negation.

`mul_signed_c` — Thin Wrapper for Backface Culling

Chapter 5's backface culling passes the first operand in A rather than B. A thin wrapper avoids restructuring the caller:

```

; mul_signed_c — signed multiply with A,C inputs
; Input:  A = signed multiplicand, C = signed multiplier
; Output: HL = signed 16-bit result
; Cost:   ~250-270 T-states (Pentagon)

```

```
mul_signed_c:
    ld    b, a          ; 4T
    jr    mul_signed    ; 12T  fall through to mul_signed
```

Cost Comparison

Routine	Input	Result	T-states	Notes
mulu112 (unsigned)	B, C	A:C (16-bit)	196-204	Chapter 4 shift-and- add
mulu_fast (square table)	B, C	HL (16-bit)	~61	Needs 512-byte table; rounding error
mul_signed	B, C (signed)	HL (signed 16-bit)	~240-260	Sign handling adds ~40-60T
mul_signed_c	A, C (signed)	HL (signed 16-bit)	~250-270	Wrapper for backface culling

The signed multiply is roughly 25% more expensive than unsigned. For a wireframe cube with 8 vertices and 6 multiplies per axis rotation (12 total per full 3D rotation), the per-vertex cost is ~3,120 T-states — still comfortably within the frame budget.

The rotation matrices in Chapter 5 call `mul_signed` six times per vertex for Z-axis and perspective, and `mul_signed_c` twice per face for backface culling. Now you know exactly what those calls do.

Credit: The signed arithmetic gap was identified by Ped7g (Peter Helcmanovsky) during his review of the book.

Going Further: The Fix-Unsigned Technique

The $2 \times \text{abs}$ approach above is clear and correct, but there is a more elegant method. The key mathematical identity: for an N-bit two's complement number, $a_{\text{signed}} = a_{\text{unsigned}} - 2^N \times \text{sign_bit}$. So for 8×8 multiply:

$$a_s \times b_s = a_u \times b_u - 256 \times \text{sign}(a) \times b_u - 256 \times \text{sign}(b) \times a_u$$

(The $+65536 \times \text{sign}(a) \times \text{sign}(b)$ term overflows the 16-bit result and vanishes.)

In practice: **do the unsigned multiply, then fix the high byte.** If A was negative, subtract B from the high byte. If B was negative, subtract A from the high byte. No absolute values, no conditional negate of the result.

On pure Z80, the savings over $2 \times \text{abs}$ are modest — the shift-and-add loop dominates the cost either way. But on Z80N, where `MUL DE` does an 8×8 unsigned multiply in a single instruction, the fix-unsigned approach shrinks signed multiply to ~70 T-states and 16 bytes:

```
; Input:  DE = x (16-bit signed), L = y (8-bit signed)
; Output: AL = signed 16-bit result
; Cost:   ~70T, 16 bytes. Requires Z80N MUL instruction.
```

```
    xor  a
    bit  7, l          ; is y negative?
    jr   z, .y_pos
    sub  e              ; compensate: high byte -= low(x)
.y_pos:
    ld   h, d           ; H = x_high, L = y
    ld   d, l           ; D = y,      E = x_low
    mul  de             ; DE = x_low * y (unsigned)
    add  a, d           ; accumulate high byte
    ex   de, hl         ; L = low result
    mul  de             ; DE = x_high * y (unsigned)
    add  a, e           ; final high byte → A:L
    ret
```

The same principle works for pure Z80 — save the original operands before the shift-and-add loop, then apply the two conditional subtracts to the high byte. The code is slightly shorter than $2 \times \text{abs} + \text{neg}$, and avoids PUSH/POP for the sign flag.

Credit: Fix-unsigned technique and Z80N implementation by Ped7g.
Pure Z80 variants by base and busy (SinDiKat).

Division on Z80

Division on the Z80 is even more painful than multiplication. No divide instruction, and the algorithm is inherently serial – each quotient bit depends on the previous subtraction. Dark again presents two methods: accurate and fast.

Method 1: Shift-and-Subtract (Restoring Division)

Binary long division. Start with a zeroed accumulator. The dividend shifts in from the right, one bit per iteration. Try subtracting the divisor; if it succeeds, set a quotient bit. If it fails, restore the accumulator – hence “restoring division.”

```
; DIVU111 -- 8-bit unsigned divide
; Input:  B = dividend, C = divisor
; Output: B = quotient, A = remainder
; Cost:   236-244 T-states (Pentagon)
;
; From Dark / X-Trade, Spectrum Expert #01 (1997)

divu111:
    xor  a              ; clear accumulator (remainder workspace)
    ld   d, 8           ; 8 bits to process

.loop:
    sla  b              ; shift dividend left -- MSB into carry
    rla                    ; shift carry into accumulator
```

```

cp    c                ; try to subtract divisor
jr    c, .too_small    ; if accumulator < divisor, skip
sub   c                ; subtract divisor from accumulator
inc   b                ; set bit 0 of quotient (B was just shifted,
                        ;   so bit 0 is free)
.too_small:
dec   d
jr    nz, .loop
ret                    ; B = quotient, A = remainder

```

The INC B to set the quotient bit is a neat trick: B was just shifted left by SLA B, so bit 0 is guaranteed zero. INC B sets it without affecting other bits – cheaper than OR or SET.

The 16-bit version (DIVU222) costs 938 to 1034 T-states. A thousand T-states for a single divide. With a frame budget of ~70,000 T-states, you can afford perhaps 70 divides per frame – doing nothing else. This is why demoscene 3D engines go to extreme lengths to avoid division.

Method 2: Logarithmic Division

Dark's faster alternative uses logarithm tables:

$$A / B = \text{AntiLog}(\text{Log}(A) - \text{Log}(B))$$

With two 256-byte lookup tables – Log and AntiLog – division becomes two lookups, a subtraction, and a third lookup. Cost drops to roughly 50-70 T-states. For perspective division (dividing by Z to project 3D points onto screen), this is a game-changer.

Generating the log table is where things get interesting. Dark proposes building it using derivatives – the same incremental technique as the square table. The derivative of $\log_2(x)$ is $1/(x * \ln(2))$, so you accumulate fractional increments step by step, starting from $\log_2(1) = 0$ and working upward. The constant $1/\ln(2) = 1.4427$ needs to be scaled to fit the table's 8-bit range.

And here is where Dark's honesty shines through. After deriving the generation formula, he attempts to compute a correction coefficient for the table scaling and arrives at 0.4606. He then writes – in a published magazine article – *"Something is not right here, so it is recommended to write a similar one yourself."*

A seventeen-year-old in 1997, publishing in a disk magazine read by his peers across the Russian Spectrum scene, openly saying: I got this working, but my derivation has a hole in it, figure out the clean version yourself. That honesty is rare in technical writing at any level, and it is one of the things that makes Spectrum Expert such a remarkable document.

In practice, the log tables work. Rounding errors from compressing a continuous function into 256 bytes are acceptable for perspective projection. Dark's 3D engine in *Illusion* uses exactly this technique.

Sine and Cosine

Rotation, scrolling, plasma – every effect that curves needs trigonometry. On the Z80, you pre-compute a lookup table. Dark’s approach is beautifully pragmatic: a parabola is close enough to a sine wave for demo work.

The Parabolic Approximation

Half a period of cosine, from 0 to π , curves from +1 down to -1. A parabola $y = 1 - 2*(x/\pi)^2$ follows almost the same path. Maximum error is about 5.6% – terrible for engineering, invisible in a demo at 256x192 resolution.

Dark generates a 256-byte signed cosine table (-128 to +127), indexed by angle: 0 = 0 degrees, 64 = 90 degrees, 128 = 180 degrees, 256 wraps to 0. The power-of-two period means the angle index wraps naturally with 8-bit overflow, and cosine becomes sine by adding 64.

```
; Generate 256-byte signed cosine table (-128..+127)
; using parabolic approximation
;
; The table covers one full period: cos(n * 2*pi/256)
; scaled to signed 8-bit range.
;
; Approach: for the first half (0..127), compute
;   y = 127 - (x^2 * 255 / 128^2)
; approximated via incrementing differences.
; Mirror for second half.
```

```
gen_cos_table:
    ld    hl, cos_table
    ld    b, 0                ; x = 0
    ld    de, 0                ; running delta (fixed-point)

    ; First quarter: cos descends from +127 to 0
    ; Second quarter: continues to -128
    ; ...build via incremental squared differences

    ; In practice, the generation loop runs ~30 bytes
    ; and produces the table in a few hundred cycles.
```

The key insight: you do not need to compute x^2 for each entry. Since $(x+1)^2 - x^2 = 2x + 1$, you build the parabola incrementally – start at the peak, subtract a linearly increasing delta. No multiplication, no division, no floating point.

The resulting table is a piecewise parabolic approximation. Plot it against true sine and you will struggle to see the difference. For wireframe 3D or a bouncing scroller, it is more than good enough.

Sidebar: Raider’s 9 Commandments of Sine Tables

In the Hype comments on Introspec’s analysis of *Illusion*, veteran coder Raider dropped a list of rules for sine table design that became informally known as the “9 commandments.” The key principles:

- Use a power-of-two table size (256 entries is canonical).

- Align the table to a page boundary so `H` holds the base and `L` is the raw angle – indexing is free.
- Store signed values for direct use in coordinate arithmetic.
- Let the angle wrap naturally via 8-bit overflow – no bounds checking.
- Cosine is just sine offset by a quarter period: load angle, add 64, look up.
- If you need higher precision, use a 16-bit table (512 bytes) but you rarely do.
- Generate the table at startup rather than storing it in the binary – saves space, costs nothing.
- For 3D rotation, pre-multiply by your scaling factor and store the scaled values.
- Never compute trigonometry at runtime. If you think you need to, you are wrong.

These commandments reflect decades of collective experience. Follow them and your sine tables will be fast, small, and correct.

Bresenham's Line Drawing

Every edge of a wireframe object is a line from (x_1, y_1) to (x_2, y_2) , and you need to draw it fast. Dark's treatment in *Spectrum Expert #01* is the longest section of his article, working through three progressively faster approaches.

The Classic Algorithm and Xopha Modification

Bresenham's algorithm steps along the major axis one pixel at a time, maintaining an error accumulator for minor-axis steps. On the Spectrum, "set a pixel" is expensive – the interleaved screen memory means computing a byte address and bit position costs real T-states. The ROM routine takes over 1000 T-states per pixel. Even a hand-optimised Bresenham loop costs ~80 T-states per pixel.

Dark mentions Xopha's improvement: maintain a screen pointer (HL) and advance it incrementally rather than recomputing from scratch. Moving right means rotating a bit mask; moving down means the multi-instruction `DOWN_HL` adjustment. Better, but the core problem remains.

Dark's Matrix Method: 8x8 Pixel Grids

Then Dark makes his key observation: **"87.5% of checks are wasted."**

In a Bresenham loop, at every pixel you ask: should I step sideways? For a nearly horizontal line, the answer is almost always no. On average, seven out of eight checks produce no side-step. You are burning T-states on a conditional branch that almost never fires.

Dark's solution: pre-compute the pixel pattern for each line slope within an 8x8 pixel grid, and unroll the drawing loop to output entire grid cells at once. A line segment within an 8x8 area is fully determined by its slope. For each of the eight octants, enumerate all possible 8-pixel patterns as straight sequences of SET bit, (HL) instructions with address increments between them.


```
; Example: one unrolled 8-pixel segment of a nearly-horizontal line
; (octant 0: moving right, gently sloping down)
;
; The line enters at the left edge of an 8x8 character cell
; and exits at the right edge, dropping one pixel row partway through.
```

```
set 7, (hl)      ; pixel 0 (leftmost bit in byte)
set 6, (hl)      ; pixel 1
set 5, (hl)      ; pixel 2
set 4, (hl)      ; pixel 3
set 3, (hl)      ; pixel 4
; --- step down one pixel row ---
inc h            ; next screen row (within character cell)
set 2, (hl)      ; pixel 5
set 1, (hl)      ; pixel 6
set 0, (hl)      ; pixel 7 (rightmost bit in byte)
```

No conditional branches. No error accumulator. SET bit, (HL) takes 15 T-states; eight of them plus a couple of INC H operations gives ~130 T-states per 8-pixel segment, or about 16 T-states per pixel. With lookup and cell-advance overhead, Dark achieves approximately **48 T-states per pixel** – nearly half the classical Bresenham cost.

The price is memory: a separate unrolled routine for each slope per octant, about **3KB total**. On a 128K Spectrum, a modest investment for a massive speed gain.

Trap-Based Termination

Instead of checking a loop counter at every pixel, Dark plants a sentinel where the line ends. When the drawing code hits the sentinel, it exits – eliminating the DEC counter / JR NZ overhead entirely.

The complete system – octant selection, segment lookup, unrolled drawing, trap termination – is one of the most impressive pieces of code in Spectrum Expert #01. When Introspec disassembled *Illusion* in 2017, he found this matrix method at work, drawing the wireframes at full frame rate.

Fixed-Point Arithmetic

Every algorithm in this chapter assumes something we have not yet made explicit: fixed-point numbers.

The Z80 has no floating-point unit. Every register holds an integer. But demo effects need fractional values – rotation angles, sub-pixel velocities, scale factors. The solution is fixed-point: pick a convention for where the “decimal point” lives within an integer, then do all arithmetic in integers while tracking the scaling mentally.

Format 8.8

The most common format on the Z80 is **8.8**: high byte = integer part, low byte = fractional part. One 16-bit register pair holds one fixed-point number:

L = fractional part (0..255, representing 0/256 to 255/256)

HL = \$0180 represents 1.5 (H=1, L=128, and $128/256 = 0.5$). HL = \$FF80 signed is -0.5 (H=\$FF = -1 in two's complement, L=\$80 adds 0.5).

The beauty: **addition and subtraction are free** – just normal 16-bit operations:

```
; Fixed-point 8.8 addition: result = a + b
; HL = first operand, DE = second operand
  add hl, de          ; that's it. 11 T-states.
```

```
; Fixed-point 8.8 subtraction: result = a - b
  or  a              ; clear carry
  sbc hl, de         ; 15 T-states.
```

The processor does not care that you are treating these as fixed-point. Binary addition is the same whether the bits represent integers or 8.8 values.

Fixed-Point Multiplication

Multiplying two 8.8 numbers produces a 16.16 result – 32 bits. You want 8.8 back, so you take bits 8..23 of the product (effectively shifting right by 8). In practice, with small integer parts (coordinates, rotation factors between -1 and +1), you can decompose the multiply into partial products:

```
; Fixed-point 8.8 multiply (simplified)
; Input:  BC = first operand (B.C in 8.8)
;         DE = second operand (D.E in 8.8)
; Output: HL = result (H.L in 8.8)
;
; Full product = BC * DE (32 bits), we want bits 8..23
;
; Decomposition:
;   BC * DE = (B*256+C) * (D*256+E)
;            = B*D*65536 + (B*E + C*D)*256 + C*E
;
; In 8.8 result (bits 8..23):
;   H.L = B*D*256 + B*E + C*D + (C*E)/256
;
; For small B,D (say -1..+1), B*D*256 is the dominant term.
; C*E/256 is a rounding correction.
; Total cost: ~200 T-states using the shift-and-add multiplier.
```

```
fixmul88:
  ; Multiply B*E -> add to result high
  ld  a, b
  call mul8          ; A = B*E (assuming 8x8->8 truncated)
  ld  h, a

  ; Multiply C*D -> add to result
```

```

ld    a, c
ld    b, d
call mul8          ; A = C*D
add   a, h
ld    h, a

; For higher precision, also compute B*D and C*E
; and combine. In practice, the two middle terms
; are often sufficient for demo work.

ld    l, 0          ; fractional part (approximate)
ret

```

For sine-table-driven rotation where sine values are 8-bit signed (-128 to +127, representing -1.0 to +0.996), multiplying an 8-bit coordinate by a sine value via `mulu112` gives a 16-bit result already in 8.8 format – high byte is the rotated integer coordinate, low byte is the fraction.

Why Fixed-Point Matters

Format 8.8 is the sweet spot for the Z80: fits in a register pair, add/subtract are free, multiply costs ~200 T-states, and precision is sufficient for screen-resolution effects. Other formats exist – 4.12 for more fractional precision, 12.4 for more integer range – but 8.8 covers the vast majority of use cases. The game development chapters later in this book use 8.8 exclusively.

Theory and Practice

These algorithms are not isolated techniques. They form a system. Multiply feeds the rotation matrix. Rotation outputs coordinates needing perspective division. Division uses log tables. Projected vertices connect with lines drawn by the matrix method. All of it runs on fixed-point arithmetic, with sine values from the parabolic table.

Dark designed these as components of a single engine – the engine that powered *Illusion*. A wireframe cube spinning at full frame rate exercises every routine in this chapter:

1. **Read the rotation angle** from the sine table (parabolic approximation, ~20 T-states per lookup)
2. **Multiply** vertex coordinates by rotation factors (shift-and-add for accuracy, or square-table for speed – ~200 or ~60 T-states per multiply, 12 multiplies per vertex)
3. **Divide** by Z for perspective projection (log tables, ~60 T-states per division)
4. **Draw lines** between projected vertices (matrix Bresenham, ~48 T-states per pixel)

For a simple cube (8 vertices, 12 edges), the total per-frame cost is roughly:

- Rotation: 8 vertices x 12 multiplies x 200 T-states = 19,200 T-states
- Projection: 8 vertices x 1 divide x 60 T-states = 480 T-states
- Line drawing: 12 edges x ~40 pixels x 48 T-states = 23,040 T-states

- **Total: ~42,720 T-states** – comfortably within the ~70,000 T-state frame budget

Switch to the fast square-table multiply and rotation drops to 5,760 T-states. The vertices jitter slightly, but you now have headroom for more complex objects. Speed or accuracy – in a demo, you make that choice for every effect, every frame.

What Dark Got Right

Looking back at Spectrum Expert #01 from nearly thirty years' distance, what strikes you is not just the quality of the algorithms but the quality of the thinking. Dark presents each one, explains the trade-offs honestly, admits when his derivation has gaps, and trusts the reader to fill those gaps.

He was writing for Spectrum coders in Russia in the late 1990s – a community building some of the most impressive 8-bit demos in the world, on hardware the rest of the world had abandoned. These are the building blocks they used. When you write your first 3D engine for the Spectrum, these routines will make it possible.

In the next chapter, Dark and STS extend this mathematical foundation into a complete 3D system: the midpoint method for vertex interpolation, backface culling, and solid polygon rendering. The maths here is the foundation. Chapter 5 is the architecture built on top.

Random Numbers: When Tables Won't Do

Everything so far in this chapter is deterministic. Given the same inputs, the same multiply, the same sine lookup, the same line draw – you get the same output. That is exactly what you want for a spinning wireframe cube or a smooth plasma.

But sometimes you need chaos. Stars twinkling in a star field. Particles scattering from an explosion. Noise textures for terrain generation. A shuffled order for loading screens. In size-coding competitions (256 bytes or less), a good random number generator can produce surprisingly complex visual effects from almost no code.

The Z80 has no hardware random number generator. You must synthesize randomness from arithmetic, and the quality of that arithmetic matters more than you might think.

The R Register Trick

The Z80 has a built-in source of entropy that many coders reach for first: the R register. It increments automatically with each instruction fetch (every M1 cycle), cycling through 0-127. You can read it in 9 T-states:

```
ld  a, r          ; 9 T -- read refresh counter
```

This is *not* a PRNG. The R register is fully deterministic – it advances by one per instruction, and its value at any point depends entirely on the code path taken since

reset. In a demo with a fixed main loop, R produces the same sequence every time. But it is useful as a seed source: read R once at startup (when timing depends on how long the user waited before pressing a key) and feed that unpredictable value into a proper PRNG.

Some coders mix R into their generator at every call, adding genuine instruction-timing entropy. The Ion generator below uses exactly this trick.

Four Generators from the Community

In 2024, Gogin (of the Russian ZX scene) assembled a collection of Z80 PRNG routines and shared them for evaluation. Gogin tested them systematically, filling large bitmaps to reveal statistical patterns. The results are instructive – not all “random” routines are equally random.

Here are four generators from that collection, ordered from best to worst quality.

Patrik Rak (Raxoft)’s CMWC Generator (Best Quality)

This is a **Complement Multiply-With-Carry** generator by Patrik Rak (Raxoft), using the multiplier 253 and an 8-byte circular buffer. The mathematics behind CMWC are well-studied: George Marsaglia proved that certain multiplier/buffer combinations produce sequences with enormous periods. With multiplier 253 and buffer size 8, the theoretical period is $(253^8 - 1) / 254$ – approximately 2^{66} values before repeating.

```
; Patrik Rak's CMWC PRNG
; Quality: Excellent -- passes visual bitmap tests
; Size:    ~30 bytes code + 8 bytes table
; Output:  A = pseudo-random byte
; Period:  ~2^66

patrik_rak_cmcw_rnd:
    ld    hl, .table
.smc_idx:
    ld    bc, 0                ; 10 T -- i (self-modifying)
    add   hl, bc                ; 11 T
    ld    a, c                  ; 4 T
    inc   a                    ; 4 T
    and   7                     ; 7 T -- wrap index to 0-7
    ld    (.smc_idx+1), a       ; 13 T -- store new index
    ld    c, (hl)               ; 7 T -- y = q[i]
    ex    de, hl                ; 4 T
    ld    h, c                  ; 4 T -- t = 256 * y
    ld    l, b                  ; 4 T
    sbc   hl, bc                ; 15 T -- t = 255 * y
    sbc   hl, bc                ; 15 T -- t = 254 * y
    sbc   hl, bc                ; 15 T -- t = 253 * y
.smc_car:
    ld    c, 0                  ; 7 T -- carry (self-modifying)
    add   hl, bc                ; 11 T -- t = 253 * y + c
    ld    a, h                  ; 4 T
    ld    (.smc_car+1), a       ; 13 T -- c = t / 256
```

```

ld    a, l                ; 4 T -- x = t % 256
cpl                   ; 4 T -- x = ~x (complement)
ld    (de), a             ; 7 T -- q[i] = x
ret                                ; 10 T

```

```
.table:
```

```
DB    82, 97, 120, 111, 102, 116, 20, 12
```

The algorithm multiplies the current buffer entry by 253, adds a carry value, stores the new carry, and complements the result. The 8-byte circular buffer means the generator's state space is vast – 8 bytes of buffer plus 1 byte of carry plus the index, giving far more internal state than any single-register generator can achieve.

Gogin's verdict: **best quality** in the collection. When filling a 256x192 bitmap, no visible patterns emerge even at large scales.

Ion Random (Second Best)

Originally from Ion Shell for the TI-83 calculator, adapted for Z80. This generator mixes the R register with a feedback loop, achieving surprisingly good randomness from just ~15 bytes:

```

; Ion Random
; Quality: Good -- minor patterns visible only at extreme scale
; Size:    ~15 bytes
; Output:  A = pseudo-random byte
; Origin:  Ion Shell (TI-83), adapted for Z80

```

```
ion_rnd:
```

```

.smc_seed:
    ld    hl, 0            ; 10 T -- seed (self-modifying)
    ld    a, r             ; 9 T -- read refresh counter
    ld    d, a             ; 4 T
    ld    e, (hl)          ; 7 T
    add   hl, de           ; 11 T
    add   a, l             ; 4 T
    xor   h                ; 4 T
    ld    (.smc_seed+1), hl ; 16 T -- update seed
    ret                    ; 10 T

```

The R register injection means this generator produces different sequences depending on the calling context – how many instructions execute between calls affects R, which feeds back into the state. For a demo main loop with fixed timing, R advances predictably, but the nonlinear mixing (ADD + XOR) still produces good output. In a game where player input varies the call pattern, the R contribution adds genuine unpredictability.

Gogin's verdict: **second best**. Very compact, good quality for its size.

XORshift 16-bit (Mediocre)

A 16-bit XORshift generator – the Z80 adaptation of Marsaglia's well-known family:

```

; 16-bit XORshift PRNG
; Quality: Mediocre -- visible diagonal patterns in bitmap tests
; Size:    ~25 bytes
; Output:  A = pseudo-random byte (H or L)
; Period:  65535

xorshift_rnd:
.smc_state:
    ld    hl, 1                ; 10 T -- state (self-modifying, must not be 0)
    ld    a, h                 ; 4 T
    rra                        ; 4 T
    ld    a, l                 ; 4 T
    rra                        ; 4 T
    xor    h                   ; 4 T
    ld    h, a                 ; 4 T
    ld    a, l                 ; 4 T
    rra                        ; 4 T
    ld    a, h                 ; 4 T
    rra                        ; 4 T
    xor    l                   ; 4 T
    ld    l, a                 ; 4 T
    xor    h                   ; 4 T
    ld    h, a                 ; 4 T
    ld    (.smc_state+1), hl ; 16 T -- update state
    ret                        ; 10 T

```

XORshift generators are fast and simple, but with only 16 bits of state the period is at most 65,535. More problematically, the bit-rotation pattern creates visible diagonal streaks when the output is mapped to pixels. For a quick star field or particle effect this may be acceptable. For anything that fills large screen areas with “noise,” the patterns become obvious.

Patrik Rak’s CMWC Variant (Mediocre)

A second CMWC variant by Patrik Rak (Raxoft), similar in principle to his version above but with a different buffer arrangement. Gogin found it produced **visible patterns at scale** – likely due to the way the carry propagation interacts with the buffer indexing. We include it in the compilable example (examples/prng.a80) for completeness, but for production use, his 8-byte-buffer version above is strictly superior.

Elite’s Tribonacci Approach

Worth a brief mention: the legendary *Elite* (1984) used a Tribonacci-like sequence for its procedurally generated galaxy. Three registers feed back into each other in a cycle, producing deterministic but well-distributed sequences. The key insight was reproducibility – given the same seed, the same galaxy generates every time, which meant the entire universe could “fit” in a few bytes of generator state. David Braben and Ian Bell used this to generate 8 galaxies of 256 star systems each from a handful of seed bytes. The technique is closer to a hash function than a PRNG, but the principle – small state, large apparent complexity – is the same one that drives demoscene size-coding.

Elite's Galaxy Generator: A Deeper Look

The Tribonacci approach deserves more detail because it illustrates a key principle: **a PRNG is not just a random number source - it is a compression algorithm.**

David Braben and Ian Bell needed 8 galaxies of 256 star systems, each with a name, position, economy, government type, and tech level. Storing all of that explicitly would consume kilobytes. Instead, they stored only a 6-byte seed per galaxy and a deterministic generator that expanded each seed into the full star system data. The generator was a three-register feedback loop - each step rotates and XORs three 16-bit values:

```
; Elite's galaxy generator (conceptual, 6502 origin):
; seed = [s0, s1, s2] (three 16-bit words)
; twist: s0' = s1, s1' = s2, s2' = s0 + s1 + s2 (mod 65536)
; repeat twist for each byte of star system data
```

On the Z80, the same principle works with three register pairs. The “twist” operation produces deterministic but well-distributed values. The crucial property: given the same seed, the same galaxy generates every time. Navigation between stars is just re-seeding and re-generating.

This idea - **small state, large apparent complexity** - drives demoscene size-coding too. A 256-byte intro that fills the screen with intricate patterns is doing exactly what Elite did: expanding a tiny seed into a large, complex output through a deterministic process.

Shaped Randomness

Sometimes you want numbers that are random but follow a specific distribution. A flat uniform PRNG gives every value equal probability, but real-world phenomena are rarely uniform: enemy spawn rates, particle speeds, terrain heights - all tend to cluster around preferred values.

Common tricks on the Z80:

- **Triangular distribution** - add two uniform random bytes and shift right. The sum clusters around the centre (128), producing “natural-looking” variation. Cost: two PRNG calls + ADD + SRL = ~20 extra T-states.

```
; Triangular random: result clusters around 128
call patrik_rak_cmw_rnd ; A = uniform random
ld b, a
call patrik_rak_cmw_rnd ; A = another uniform random
add a, b ; sum (wraps at 256)
rra ; divide by 2 → triangular distribution
```

- **Rejection sampling** - generate a random number, reject values outside your desired range. For power-of-two ranges this is free (just AND with a mask). For arbitrary ranges, loop until the value fits.
- **Weighted tables** - store a 256-byte lookup table where each output value appears in proportion to its desired probability. Index with a uniform random byte. The table costs 256 bytes but the lookup is instant (7 T-states). Perfect when the distribution is complex and fixed.

- **PRNG as hash function** – feed structured data (coordinates, frame numbers) through the PRNG to get deterministic noise. This is how size-coded plasma and noise textures work: `random(x XOR y XOR frame)` gives a different-looking value per pixel per frame, but it is entirely reproducible.

Seeds and Reproducibility

In a demo, reproducibility is usually desirable: the effect should look the same every time it runs, because the coder choreographed the visuals to match the music. Seed the PRNG once with a fixed value and the sequence is deterministic.

In a game, unpredictability matters. Common seeding strategies:

- **FRAMES system variable (\$5C78)** – the Spectrum ROM maintains a 3-byte frame counter at address \$5C78 that increments every 1/50th of a second from power-on. Reading it gives a time-dependent seed that varies with how long the machine has been running. Art-top (Artem Topchiy) recommends using it to initialise Patrik Rak's CMWC table:

```
; Seed Patrik Rak CMWC from FRAMES system variable
ld  hl, $5C78                ; FRAMES (3 bytes, increments at 50 Hz)
ld  a, (hl)                  ; low byte -- most variable
ld  de, patrik_rak_cmwcrnd.table
ld  b, 8
.seed_loop:
xor  (hl)                    ; mix with FRAMES
ld  (de), a                  ; write to table
inc  de
rlca                          ; rotate for variety
add  a, b                    ; add loop counter
djnz .seed_loop
```

- **Read R at a user-input moment** – the exact instruction count between reset and the player pressing a key varies each run. LD A,R at that moment captures timing entropy.
- **Frame counter accumulation** – XOR the R register into an accumulator every frame during the title screen; use the accumulated value as seed when the game starts.
- **Combine multiple sources** – XOR together R, the low byte of FRAMES, and a byte from the floating bus (on 48K Spectrums, reading certain ports returns whatever the ULA is currently fetching from RAM – a source of positional entropy).

For demos, simply initialise the generator's state to a known value and leave it. The compilable example (`examples/prng.a80`) shows all four generators with fixed seeds.

Comparison Table

Algorithm	Size (bytes)	Speed (T-states)	Quality	Period	Notes
Patrik Rak CMWC	~30 + 8 table	~170	Excellent	~2 ⁶⁶	Best overall; 8-byte buffer
Ion Random	~15	~75	Good	Depends on R	Compact; mixes R register
XORshift 16	~25	~90	Mediocre	65,535	Visible diagonal patterns
Patrik Rak CMWC (alt)	~35 + 10 table	~180	Mediocre	~2 ⁶⁶	Patterns visible at scale
LD A,R alone	2	9	Poor	128	NOT a PRNG; use as seed only

For most demoscene work, **Patrik Rak's CMWC** is the clear winner: excellent quality, reasonable size, and a period so long it will never repeat during a demo. If code size is critical (size-coding, 256-byte intros), **Ion Random** packs remarkable quality into 15 bytes. XORshift is a fallback when you need something quick and do not care about visual quality.

Credits: PRNG collection, quality assessment, and bitmap testing by **Gogin**. Patrik Rak's CMWC generator is based on George Marsaglia's Complementary Multiply-With-Carry theory. Ion Random originates from **Ion Shell** for the TI-83 calculator.

PRNG output — random attribute colours fill the screen, revealing the generator's statistical quality

Figure 12: PRNG output — random attribute colours fill the screen, revealing the generator's statistical quality

All cycle counts in this chapter are for Pentagon timing (no wait states). On a standard 48K Spectrum or Scorpion with contended memory, expect higher counts for code executing in the lower 32K of RAM (\$4000-\$7FFF). Chapter 15.2 explains contention patterns in detail and provides strategies for keeping critical loops out of contended address space. See also Appendix A for the complete timing reference.

Sources: Dark / X-Trade, "Programming Algorithms" (Spectrum Ex-

pert #01, 1997); Gogin, PRNG collection and quality assessment; Patrik Rak (Raxoft), CMWC generator; Ped7g (Peter Helcmanovsky), signed arithmetic gap identification and review

Chapter 5: 3D on 3.5 MHz

“Calculate only what you must. Derive the rest.” — Dark & STS, Spectrum Expert #02 (1998)

The previous chapter gave you the building blocks: multiplication, division, sine tables, line drawing. Now we put them together. The goal is a spinning three-dimensional solid object on a ZX Spectrum — filled polygons, backface culling, correct depth ordering — at a usable frame rate.

This is where you hit the wall.

The Problem: Twelve Multiplications Per Vertex

Rotating a point in three-dimensional space around all three axes requires a sequence of trigonometric multiplications. If you rotate sequentially — first around Z, then Y, then X — each axis involves four multiplications and two additions to transform two coordinates. Three axes, four multiplications each: twelve multiplications per vertex.

Take the shift-and-add multiply from Chapter 4, which costs roughly 200 T-states. Twelve of those give you 2,400 T-states to rotate a single vertex. A simple cube has 8 vertices: 19,200 T-states just for rotation. Chapter 4 showed that this fits in the frame budget — barely.

Now try something more interesting. A sphere approximated by 20 vertices and 36 faces:

That is 67% of the Pentagon’s 71,680 T-state frame budget, consumed before you have drawn a single pixel. You still need perspective projection, backface culling, polygon sorting, and the actual fill. There is no room. The object cannot be more complex than a cube unless you find a fundamentally cheaper way to compute vertex positions.

Dark and STS found one.

The Midpoint Method

The insight is geometric. Not every vertex in an object carries independent information. Many vertices sit at structurally predictable positions — midpoints of edges,

centers of faces, reflections of other vertices. If you can express those relationships explicitly, you can replace expensive multiplications with cheap averages.

The Cube as Foundation

Consider a cube centered at the origin. It has 8 vertices, but they are not 8 independent points. They are 4 pairs of diametrically opposite vertices. If you know one vertex of a pair, the other is its negation through the center:

$$\begin{aligned} v1 &= (x, y, -z) & \rightarrow & & v6 &= (-x, -y, z) \\ v2 &= (x, -y, z) & \rightarrow & & v5 &= (-x, y, -z) \\ v3 &= (x, -y, -z) & \rightarrow & & v4 &= (-x, y, z) \end{aligned}$$

Rotate 4 vertices using the full 12-multiplication procedure. Negate them to get the other 4. Negation on the Z80 is NEG — 8 T-states for one coordinate, 24 T-states for all three. Compare that to 2,400 T-states for a full rotation. You have cut the vertex computation nearly in half.

But the midpoint method goes much further than mirroring.

Deriving Vertices by Averaging

The key operation is the average: given two already-computed points, their midpoint is simply the mean of their coordinates.

On the Z80, this is an addition and a shift:

```
; Average two signed 8-bit coordinates
; A = first coordinate, B = second coordinate
; Result in A = (A + B) / 2
```

```
add a, b          ; 4 T-states
sra a             ; 8 T-states
                  ; ----
                  ; 12 T-states total
```

SRA (Shift Right Arithmetic) preserves the sign bit, so this works correctly for negative coordinates. For all three coordinates (x, y, z), averaging costs 36 T-states per derived vertex. Compare that to 2,400 T-states for a full rotation.

The ratio: averaging is **66 times cheaper** than rotation.

This means you can build complex objects from a small set of “basis” vertices that you rotate fully, then derive all remaining vertices through chains of averages. The more vertices you can derive, the more time you save.

Building Complex Objects

Suppose you want a 20-vertex object. With the midpoint method:

1. Rotate 4 basis vertices fully: $4 \times 2,400 = 9,600$ T-states
2. Mirror 4 vertices by negation: $4 \times 24 = 96$ T-states
3. Derive 12 vertices by averaging: $12 \times 36 = 432$ T-states
4. **Total: 10,128 T-states**

Without the midpoint method, the same 20 vertices would cost 48,000 T-states. You have saved 37,872 T-states — more than half the frame budget freed up for projection, culling, and rendering.

The constraint is topological: you can only derive a vertex by averaging if it genuinely lies at the midpoint of two other vertices (or close enough that the error is invisible at 256x192 resolution). This shapes how you design your 3D models. You do not model freely and then optimize — you design the model *around* the midpoint structure from the start.

Dark and STS give examples of the derivation chains:

```
v9  = (v3 + v7) / 2
v10 = (v2 + v6) / 2
v11 = (v8 + v9) / 2      ; derived from two already-derived vertices
```

Notice that v11 is derived from v8 and v9, which are themselves derived. Chains can go several levels deep. Each level adds only 36 T-states per vertex, so the cost stays negligible regardless of depth.

The Virtual Processor

Here is where Dark does something that feels anachronistic for 1998. Rather than hardcoding the derivation chains for each specific object, he designs a tiny interpreter — a virtual processor — that executes “programs” describing how to compute vertices.

Architecture

The virtual processor has:

- **One register** (a working register holding one 3D point — three bytes: x, y, z)
- **64 cells of RAM** (each cell holds one 3D point — 192 bytes total)
- **4 instructions**

Opcode	Bits	Name	Operation
00	00nnnnnn	Load	register <- cell[n]
01	01nnnnnn	Store	cell[n] <- register
10	10nnnnnn	Average	register <- (register + cell[n]) / 2
11	11-----	End	halt execution

Each instruction is encoded in a single byte: 2 bits for the opcode, 6 bits for the cell number (0-63). The entire instruction set fits in 256 possible values.

Execution

The interpreter loop is compact:

```
; Virtual processor main loop
; IX points to the program (sequence of 1-byte instructions)
```

```

; Point RAM at a fixed address, 3 bytes per cell

vp_loop:
    ld    a, (ix+0)          ; fetch instruction
    inc   ix
    ld    b, a               ; save full instruction
    and   %11000000         ; extract opcode (top 2 bits)

    cp    %11000000         ; END?
    ret   z                  ; yes – halt

    ld    a, b
    and   %00111111         ; extract cell number (bottom 6 bits)
    ; ... compute cell address from cell number ...
    ; ... dispatch based on opcode ...

    jr    vp_loop

```

The **Load** instruction copies a cell's x, y, z values into the working register. **Store** copies the working register back into a cell. **Average** adds the cell's coordinates to the working register and shifts each result right by one — the midpoint operation. **End** terminates the program.

Writing Programs

A vertex derivation chain becomes a simple sequence of bytes. Dark uses a compact notation in the article:

```

; Example: derive v8 = (v4 + v5) / 2, then store it
; Cell 4 = v4, Cell 5 = v5, Cell 8 = destination

```

```

DB  4          ; LOAD cell[4]      (opcode 00, cell 4)
DB  128+5      ; AVG  cell[5]      (opcode 10, cell 5 = %10000101)
DB  64+8       ; STORE cell[8]     (opcode 01, cell 8 = %01001000)

```

The notation 128+5 encodes %10000101 — opcode 10 (Average) with cell number 5. 64+8 encodes %01001000 — opcode 01 (Store) with cell number 8. Raw numbers, packed into data bytes, forming a tiny domain-specific program.

A complete object description might look like:

```

; Midpoint program for a 12-vertex object
; Cells 0-3: basis vertices (rotated by main code)
; Cells 4-7: mirrored vertices (negated by main code)
; Cells 8-11: derived via midpoint averaging

```

```

midpoint_program:
    DB  0          ; LOAD v0
    DB  128+1      ; AVG  v1          -> register = (v0+v1)/2
    DB  64+8       ; STORE v8

    DB  2          ; LOAD v2
    DB  128+3      ; AVG  v3          -> register = (v2+v3)/2
    DB  64+9       ; STORE v9

```

```

DB  4          ; LOAD v4
DB 128+5       ; AVG  v5      -> register = (v4+v5)/2
DB 64+10       ; STORE v10

DB  6          ; LOAD v6
DB 128+7       ; AVG  v7      -> register = (v6+v7)/2
DB 64+11       ; STORE v11

DB 192         ; END          (%11000000)

```

Thirteen bytes describe the computation of four derived vertices. The virtual processor executes them in roughly $13 \times 30 = 390$ T-states (each instruction takes approximately 25–35 T-states depending on type). Four fully rotated vertices would have cost 9,600 T-states. The savings are enormous.

Why a Virtual Processor?

You might ask: why not just write the averaging code directly in Z80 assembly? Inline the additions and shifts, skip the interpreter overhead. It would be slightly faster per vertex.

The answer is flexibility. The virtual processor separates the *description* of an object's topology from the *execution* of the vertex computation. Change the object? Write a new program — a new sequence of data bytes. The interpreter code stays the same. You can store programs for multiple objects and switch between them at zero code cost. You can even generate programs algorithmically.

This is, in essence, a domain-specific bytecode interpreter — a pattern that modern programmers would recognize from game engines, shader compilers, and scripting languages. Dark designed it in 1998, on a ZX Spectrum, to save T-states on vertex computation. The architecture is clean.

Rotation

With the midpoint method handling most vertices, you still need to rotate the basis vertices properly. Dark and STS use sequential rotation around the three axes, applied in order: Z, then Y, then X. Each rotation uses the sine and cosine tables from Chapter 4.

Z-Axis Rotation

Rotation around Z affects only X and Y:

$$Y' = -X * \sin(Az) + Y * \cos(Az)$$

In Z80 assembly, using the 8x8 signed multiply and 256-entry sine/cosine tables:

```

; Rotate point around Z axis
; Input:  (px), (py) = coordinates; (angle_z) = rotation angle
; Output: (px), (py) updated
; Uses:   cos_table, sin_table (page-aligned, signed 8-bit)

```



```

rotate_z:
    ld    a, (angle_z)
    ld    l, a
    ld    h, cos_table >> 8
    ld    d, (hl)          ; D = cos(Az)
    ld    h, sin_table >> 8
    ld    e, (hl)          ; E = sin(Az)

    ; X' = X*cos(Az) + Y*sin(Az)
    ld    a, (px)
    ld    b, a
    ld    c, d              ; B=X, C=cos
    call  mul_signed        ; HL = X * cos(Az)
    push  hl

    ld    a, (py)
    ld    b, a
    ld    c, e              ; B=Y, C=sin
    call  mul_signed        ; HL = Y * sin(Az)
    pop   de
    add   hl, de            ; HL = X*cos + Y*sin
    ld    a, h              ; take high byte as new X'
    ld    (px), a

    ; Y' = -X*sin(Az) + Y*cos(Az)
    ld    a, (px_original) ; need the original X, not the updated one
    neg
    ld    b, a
    ld    c, e              ; B=-X, C=sin
    call  mul_signed        ; HL = -X * sin(Az)
    push  hl

    ld    a, (py)
    ld    b, a
    ld    c, d              ; B=Y, C=cos
    call  mul_signed        ; HL = Y * cos(Az)
    pop   de
    add   hl, de            ; HL = -X*sin + Y*cos
    ld    a, h
    ld    (py), a

    ret

```

The same pattern repeats for Y-axis rotation (affecting X and Z) and X-axis rotation (affecting Y and Z). Dark wraps all three into a single ROTATE procedure that takes three angle parameters and transforms a point in place.

Note the detail about preserving the original X value. The second formula uses the pre-rotation X, not the X' you just computed. A common bug is to use the already-updated coordinate, which produces a skewed rotation. Dark addresses this explicitly in the article.

Cost Per Basis Vertex

Each axis rotation requires 4 multiplications and 2 additions. At 200 T-states per multiply and 11 T-states per 16-bit add:

Three axes: $3 \times 822 = 2,466$ T-states per vertex

With 4 basis vertices: roughly 9,864 T-states for rotation. Add the midpoint program execution and you have the full vertex computation for an arbitrarily complex object at a fraction of the naive cost.

Projection

Once all vertices are rotated in 3D space, you need to project them onto the 2D screen.

Parallel Projection

The simplest approach: ignore the Z coordinate entirely. Just use X and Y as screen coordinates (with appropriate offset to center the object on screen).

```
; Parallel projection: screen coords = rotated X, Y + offset
ld  a, (px)
add a, 128           ; center horizontally (128 = half of 256)
ld  (screen_x), a

ld  a, (py)
add a, 96            ; center vertically (96 = half of 192)
ld  (screen_y), a
```

Cost: essentially zero. The result looks flat — objects do not appear to recede into the distance. Parallel projection is useful for wireframe previews and effects where the rotation itself provides the illusion of depth, but it lacks the visceral punch of perspective.

Perspective Projection

Perspective makes near objects larger and far objects smaller, producing the depth cue that makes 3D convincing:

$$Y_{\text{screen}} = (Y * \text{Scale}) / (Z + Z_{\text{distance}}) + Y_{\text{offset}}$$

Scale controls the field of view. Zdistance is the distance from the camera to the projection plane — it prevents division by zero when Z approaches the camera and controls how aggressively objects scale with depth. Xoffset and Yoffset center the projection on screen.

The expensive operation here is the division. One divide per coordinate, two coordinates per vertex. Using the logarithmic division from Chapter 4 (~60 T-states per divide), the cost is modest:

```
; Perspective projection for one vertex
; Input:  (px), (py), (pz) = rotated 3D coordinates
```

; Output: (screen_x), (screen_y)

perspective:

```
; Compute denominator: Z + Zdistance
ld  a, (pz)
add a, ZDISTANCE      ; Z + viewing distance
ld  c, a               ; C = denominator

; Xscreen = (X * Scale) / (Z + Zdist) + Xoffset
ld  a, (px)
ld  b, SCALE
call mul_signed        ; HL = X * Scale
ld  a, h               ; take high byte as numerator
call log_divide        ; A = A / C (using log tables)
add a, XOFFSET
ld  (screen_x), a

; Yscreen = (Y * Scale) / (Z + Zdist) + Yoffset
ld  a, (py)
ld  b, SCALE
call mul_signed        ; HL = Y * Scale
ld  a, h               ; A = A / C
call log_divide        ; A = A / C
add a, YOFFSET
ld  (screen_y), a

ret
```

Each vertex costs two multiplications (400 T-states) and two log divides (120 T-states), plus overhead — roughly 600 T-states per vertex. For 20 vertices: 12,000 T-states. Combined with the midpoint rotation, we are at roughly 22,000 T-states for all vertex computation and projection. Less than a third of the frame budget.

Solid Polygons

A wireframe object is a collection of edges. A solid object is a collection of filled polygons. Going from wire to solid requires three additional capabilities: determining which faces are visible, sorting them by depth, and filling them with pixels.

Backface Culling

A closed 3D object has faces that point toward the viewer and faces that point away. The back-facing polygons are hidden and need not be drawn. Skipping them saves both rendering time and produces a correct solid appearance without requiring a full depth buffer.

The test is geometric. For each face, compute the Z-component of the surface normal using the cross product of two edge vectors:

Edge vectors:

$V = v1 - v0 = (Vx, Vy)$ (in screen coordinates)
 $W = v2 - v0 = (Wx, Wy)$

Z-component of normal = $Vx * Wy - Vy * Wx$

If the result is positive, the face is oriented toward the viewer and should be drawn. If negative, the face points away — cull it. If zero, the face is edge-on and invisible.

; Backface culling test for one face
 ; Input: three projected vertices (x0,y0), (x1,y1), (x2,y2)
 ; Output: S flag set (M condition) if face is back-facing (should be culled)

```
backface_test:
    ; V = v1 - v0
    ld    a, (x1)
    sub   (ix+x0)
    ld    d, a                ; D = Vx = x1 - x0

    ld    a, (y1)
    sub   (ix+y0)
    ld    e, a                ; E = Vy = y1 - y0

    ; W = v2 - v0
    ld    a, (x2)
    sub   (ix+x0)
    ld    b, a                ; B = Wx = x2 - x0

    ld    a, (y2)
    sub   (ix+y0)
    ld    c, a                ; C = Wy = y2 - y0

    ; Normal Z = Vx * Wy - Vy * Wx
    ld    a, d
    call  mul_signed_c        ; HL = Vx * Wy (D * C)
    push  hl

    ld    a, e
    ld    c, b
    call  mul_signed_c        ; HL = Vy * Wx (E * B)

    pop   de
    ex    de, hl
    or    a
    sbc   hl, de              ; HL = Vx*Wy - Vy*Wx
    ret                                ; S flag: negative = back-facing = cull
```

Two multiplications and a subtraction per face. At 400 T-states for the multiplies plus overhead, the test costs roughly 500 T-states per face. For a 12-face object, that is 6,000 T-states — and for each culled face, you save the entire cost of filling it.

On a typical rotating solid, roughly half the faces are back-facing at any time. Culling them halves your fill workload.

Z-Sorting

For a convex object (a cube, a tetrahedron), backface culling alone produces correct results: every visible face is fully visible, with no overlapping. For concave or multi-object scenes, you need to draw faces in back-to-front order so that nearer faces overwrite farther ones — the painter's algorithm.

Dark and STS compute a depth value for each visible face (typically the average Z of its vertices) and sort the face list accordingly. A simple insertion sort is adequate for the small face counts involved — sorting 6–12 faces takes negligible time compared to filling them.

```
; Simplified depth sort: compute average Z for each visible face,
; sort face indices by descending Z (farthest first)
```

```
sort_faces:
    ; For each visible face:
    ;   average_z = (z[v0] + z[v1] + z[v2] + z[v3]) / 4
    ;   store (average_z, face_index) in sort buffer
    ; Then insertion-sort the buffer by average_z
    ; ...
```

Convex Polygon Filling

Once you know which faces to draw and in what order, you need to fill them. A convex polygon (all interior angles less than 180 degrees) can be filled with a simple scanline approach:

1. Find the topmost and bottommost vertices.
2. Walk down the left edge and the right edge simultaneously, one scanline at a time.
3. For each scanline, draw a horizontal line from the left edge to the right edge.

The edge-walking uses Bresenham-style incremental stepping — no division needed per scanline, just additions and conditional increments. The horizontal fill itself is a tight loop of byte writes:

```
; Fill one scan line from x_left to x_right at screen row Y
; Screen address already computed in HL
```

```
fill_scanline:
    ld    a, (x_right)
    sub   (ix+x_left)
    ret   c                ; nothing to fill if right < left
    ret   z
    ld    b, a              ; B = pixel count

    ; For byte-aligned fills: write whole bytes
    ld    a, $FF            ; solid fill
.fill_loop:
    ld    (hl), a
    inc   l                  ; next byte (within same screen line)
    djnz .fill_loop
    ret
```

This is simplified — real polygon fillers must handle partial bytes at the left and right edges, where the polygon boundary falls within a byte rather than on a byte boundary. Those edge cases add complexity but not much cost, since they occur only twice per scanline.

Putting It All Together

The complete frame loop for a spinning 3D solid object follows this sequence:




Figure 13: 3D rendering pipeline: model, rotation, projection, screen

2. For each basis vertex:
 - Rotate through Z, Y, X axes [~2,400 T per vertex]
3. Negate basis vertices to get mirrors [~24 T per vertex]
4. Run midpoint program to derive rest [~36 T per derived vertex]
5. Project all vertices (perspective) [~600 T per vertex]
6. For each face:
 - Backface test [~500 T per face]
 - If visible: compute average Z
7. Sort visible faces by Z [~200 T for small lists]
8. For each visible face (back to front):
 - Fill polygon [varies with area]
9. Wait for next frame (HALT)

For a 20-vertex, 18-face object with 4 basis vertices, the per-frame budget breaks down as:

Stage	Vertices/Faces	Cost per	Total
Rotation (basis)	4	2,466	9,864
Negation (mirrors)	4	24	96
Midpoint derivation	12	36	432
Projection	20	600	12,000

Stage	Vertices/Faces	Cost per	Total
Backface test	18	500	9,000
Z-sort	~9 visible	-	~200
Polygon fill	~9 visible	~1,500 avg	~13,500
Total			~45,092

Forty-five thousand T-states out of 71,680 available. Tight, but workable — you have 26,000 T-states remaining for screen clearing, angle updates, and the line drawing or edge-highlighting that makes the object look crisp. And this is for a 20-vertex object, far more complex than anything you could afford with naive rotation.

The Shape of Objects

The midpoint method influences how you think about 3D models. You do not design a mesh and then optimize it — you start from the topology that the method requires.

A good midpoint object begins with a small basis. Four fully rotated points define a tetrahedron-like skeleton. Negation doubles that to eight. Midpoint averaging fills in the rest. The art is choosing basis vertices that produce useful derived points.

Consider building a 14-vertex object from scratch:

```
Mirrors:  v4, v5, v6, v7          (4 negated)
Derived:
v8  = (v0 + v1) / 2              edge midpoint
v9  = (v2 + v3) / 2              edge midpoint
v10 = (v4 + v5) / 2              edge midpoint on mirrored side
v11 = (v6 + v7) / 2              edge midpoint on mirrored side
v12 = (v0 + v2) / 2              cross-edge midpoint
v13 = (v8 + v10) / 2             second-level derivation
```

The virtual processor program for this is 19 bytes:

```
object_14v_program:
DB  0, 128+1, 64+8      ; v8  = avg(v0, v1)
DB  2, 128+3, 64+9      ; v9  = avg(v2, v3)
DB  4, 128+5, 64+10     ; v10 = avg(v4, v5)
DB  6, 128+7, 64+11     ; v11 = avg(v6, v7)
DB  0, 128+2, 64+12     ; v12 = avg(v0, v2)
DB  8, 128+10, 64+13    ; v13 = avg(v8, v10)
DB 192                  ; END
```

Nineteen bytes of data replace what would otherwise be 14,400 T-states of rotation multiplications for the 6 derived vertices.

You can push this further. A second level of averaging (deriving from already-derived points) costs nothing extra per instruction — the virtual processor does not care whether a cell holds a rotated or a derived point. Dark and STS describe chains three or four levels deep, building objects with 30 or more vertices from just 3–4 basis points.

The constraint is precision. Each averaging step introduces a rounding error of up to 0.5 units (from the integer shift). After three levels of derivation, the cumulative error can reach 1.5 units — noticeable on an object that spans 60 pixels, invisible on one that spans 120. Design your objects large enough that the rounding is below the screen resolution.

Practical: A Spinning Solid Object

Here is the outline for building a complete spinning 3D solid using everything in this chapter.

Step 1: Define the Object

Start with basis vertices. A truncated octahedron works well with the midpoint method:

```
; 4 basis vertices in signed 8-bit coordinates
basis_vertices:
    DB    30,  0,  30      ; v0 (x, y, z)
    DB    0,  30,  30      ; v1
    DB    30,  30,  0      ; v2
    DB    0,   0,   0      ; v3 (at origin for center reference)
```

Step 2: Write the Midpoint Program

```
midpoint_prog:
    ; Mirrors: cells 4-7 are pre-negated by the main loop
    ; Derive additional vertices:
    DB    0, 128+1, 64+8      ; v8 = avg(v0, v1)
    DB    2, 128+3, 64+9      ; v9 = avg(v2, v3)
    DB    0, 128+2, 64+10     ; v10 = avg(v0, v2)
    DB    1, 128+3, 64+11     ; v11 = avg(v1, v3)
    DB    192                  ; END
```

Step 3: Define Faces

```
; Face table: each face is a list of vertex indices + attribute byte
; Vertex order must be consistent (clockwise when front-facing)
face_table:
    DB    4, 0, 1, 8, 10      ; face 0: quad (4 vertices)
    DB    4, 2, 3, 9, 11      ; face 1: quad
    ; ... remaining faces ...
    DB    0                    ; end marker
```

Step 4: The Frame Loop

```
main_loop:
    halt                      ; wait for vsync (IM1)

    ; Clear the screen area (or use double buffering)
```



```

call clear_viewport

; Update angles
ld hl, angle_z
inc (hl)
ld hl, angle_x
ld a, (hl)
add a, 2
ld (hl), a

; Rotate basis vertices
ld b, 4 ; 4 basis vertices
ld ix, basis_vertices
ld iy, point_ram ; cell 0 onwards
.rotate_basis:
push bc
call rotate_xyz ; rotate point at (IX) by current angles
; store result at (IY)

ld bc, 3
add ix, bc
add iy, bc
pop bc
djnz .rotate_basis

; Negate for mirrors (cells 4-7 = negation of cells 0-3)
call negate_basis

; Run midpoint program
ld ix, midpoint_prog
call virtual_processor

; Project all vertices
call project_all

; Backface cull and sort
call cull_and_sort

; Draw visible faces
call draw_faces

jr main_loop

```

This is the skeleton. Each `call` hides a procedure built from the techniques in this chapter and Chapter 4. The frame loop itself is clean — update, compute, draw, repeat.

Rotating wireframe cube rendered on the ZX Spectrum using the midpoint method and perspective projection

Figure 14: Rotating wireframe cube rendered on the ZX Spectrum using the midpoint method and perspective projection

Historical Context: From Magazine to Demo

Dark and STS published the midpoint method in Spectrum Expert #02 in 1998. They were young coders in St. Petersburg, writing for a disk magazine distributed within the Russian ZX Spectrum community. The articles are written in the direct, practical style of people teaching their peers — here is the problem, here is the trick, here is the code.

But Spectrum Expert was not an academic exercise. Dark was from X-Trade, the same group that produced *Illusion* — the demo that won first place at ENLiGHT'96. The algorithms in the magazine are not theoretical proposals; they are the building blocks of real, competition-winning demo code. The sine tables from Chapter 4 drove the rotozoomer. The line drawer rendered the wireframes. And the midpoint method powered the 3D objects.

The virtual processor approach is particularly striking in retrospect. In 1998, the dominant paradigms in professional game development were shifting toward hardware acceleration and away from software rendering. On the Spectrum, hardware acceleration did not exist. Everything was software, and software had to be *designed* — not just written, but architected. Dark's bytecode interpreter for vertex computation is a piece of software architecture that would not feel out of place in a modern game engine's animation system or shader compiler. It separates data from execution, enables rapid iteration on object design, and keeps the hot loop tight.

The connection to *Illusion* runs deeper than shared authorship. When Introspec disassembled *Illusion* twenty years later, he found the exact same mathematical infrastructure: the shift-and-add multiplier, the parabolic sine table, the log-table divider. The midpoint method and the virtual processor are extensions of that infrastructure — the same thinking applied to a different problem. Dark was not just publishing algorithms; he was documenting the engineering philosophy behind his own award-winning demo.

In the next chapter, we will see one of *Illusion*'s most spectacular effects up close: the texture-mapped sphere. It uses the same sine tables and the same fixed-point arithmetic from Chapter 4, combined with the self-modifying code techniques from Chapter 3 and a completely different approach to the rendering problem. The midpoint method and the sphere are siblings — born from the same coder, the same tools, the same relentless drive to make the impossible fit inside 71,680 T-states.

Summary

- Rotating 3D objects naively requires 12 multiplications per vertex — too expensive for complex objects on a 3.5 MHz Z80.
- The **midpoint method** rotates only a few basis vertices fully, then derives the rest through averaging. Averaging costs ~36 T-states per vertex versus ~2,400 for full rotation — 66 times cheaper.
- A **virtual processor** with 4 instructions (Load, Store, Average, End) executes compact “programs” that describe vertex derivation chains. Object topology is data, not code.

- **Rotation** uses sequential Z/Y/X axis transforms with sine/cosine lookup tables from Chapter 4.
- **Perspective projection** uses the logarithmic division from Chapter 4 for the Z-divide.
- **Backface culling** via cross-product normal test eliminates invisible faces at ~500 T-states each.
- **Z-sorting** with the painter's algorithm handles overlapping faces for non-convex objects.
- A 20-vertex solid object can be rendered in ~45,000 T-states per frame — tight but feasible within the 71,680 T-state budget.
- These techniques were published in Spectrum Expert #02 (1998) by the same team that built *Illusion*. The magazine articles document the engineering behind the demo.

All cycle counts in this chapter are for Pentagon timing (no wait states). On a standard 48K Spectrum with contended memory, expect higher counts for code executing in the lower 32K of RAM (\$4000-\$7FFF). Chapter 15.2 explains contention patterns in detail and provides strategies for keeping critical loops out of contended address space. See also Appendix A for the complete timing reference.

Sources: Dark & STS, "Programming: 3D Graphics" (Spectrum Expert #01, 1997); Dark & STS, "Programming: 3D Graphics — Midpoint Method" (Spectrum Expert #02, 1998). The virtual processor design and midpoint derivation examples are drawn directly from the SE#02 article.

Chapter 6: The Sphere — Texture Mapping on 3.5 MHz

*“Coder effects are always about evolving a computation scheme.” —
Introspec, 2017*

It is 1996, and a demo called *Illusion* takes first place at ENLiGHT’96 in St. Petersburg. The audience watches a monochrome image wrap itself around a spinning sphere, rotating smoothly in real time, on a ZX Spectrum running at 3.5 MHz with no hardware acceleration whatsoever. No blitter. No GPU. No co-processor. Just a Z80, 48 kilobytes of contiguous RAM, and whatever a twenty-year-old coder named Dark could squeeze out of them.

Twenty years later, in March 2017, Introspec sits down with a copy of the binary and a disassembler. He picks apart the rendering loop instruction by instruction, counts T-states, maps memory addresses to data structures, and publishes his findings on Hype. What follows is one of the most detailed public teardowns of a demoscene effect ever written for the ZX Spectrum — and, in the comment thread that erupts beneath it, a debate about what really matters in real-time rendering on constrained hardware.

This chapter follows Introspec’s analysis. We will look over his shoulder as he traces the code, understand *why* the sphere works the way it does, and then build a simplified version ourselves.

The Problem: A Round Object on a Square Screen

A sphere on screen is not a sphere. It is a circle filled with a distorted image. The distortion follows the rules of spherical projection: pixels near the equator are spaced evenly, pixels near the poles are compressed horizontally, and the entire mapping curves to create the illusion of a three-dimensional surface.

The source image in *Illusion* is stored as a monochrome bitmap — one byte per pixel, where each byte is either 0 or 1. This is extravagant by Spectrum standards, where screen memory packs eight pixels per byte, but it buys something essential: the rendering code can treat pixels as arithmetic values rather than bit positions.

The task, then, is this: read pixels from the source image, sample them according to a spherical projection, pack eight of them into a single screen byte, and write

that byte to video memory. Do this for every visible byte of the sphere. Do it fast enough to animate. Do it on a 3.5 MHz Z80.

The Insight: Code That Writes Code

The first question any Z80 programmer asks is: what does the inner loop look like? On a machine where a single NOP takes 4 T-states and you have roughly 70,000 T-states per frame, the inner loop *is* the program. Everything else — setup, table generation, frame management — is overhead that happens once or infrequently. The inner loop runs thousands of times per frame.

Dark's solution is to not have a fixed inner loop at all.

Instead, the rendering code is *generated at runtime*. For each horizontal line of the sphere, the program constructs a sequence of Z80 instructions tailored to that line's geometry. The generated code reads source pixels in order, accumulates them into screen bytes via shifts and additions, and advances through the source data by distances that vary with the sphere's curvature. Different lines of the sphere produce different code.

This is a technique that appears throughout the demoscene: self-generating code, sometimes called “compiled sprites” when applied to sprite rendering, or “unrolled loop generation” in the general case. What makes the sphere version distinctive is the variability. A compiled sprite is fixed — once generated, it draws the same shape every time. The sphere code changes with the rotation angle, because different source pixels become visible as the sphere turns.

Inside the Disassembly

Introspec traced the rendering engine to a block of generated code and a set of lookup tables starting at address \$6944. The tables encode the sphere's geometry as a series of *skip distances*: for each position along a scanline of the sphere, how many source pixels should be skipped before the next one is sampled.

At the equator, the skip distances are roughly uniform — the source image maps onto the sphere with minimal distortion. Near the poles, the horizontal compression of the projection means larger skips between sampled pixels. At the very top and bottom, only a few pixels are visible per line, and the skips can be substantial.

The generated inner loop has a repeating structure. For each screen byte (eight packed pixels), it executes a sequence like this:

```
; --- Accumulating 8 source pixels into one screen byte ---
; HL points into the source image (one byte per pixel)
; A is the accumulator, building the screen byte bit by bit

    add  a,a          ; shift accumulator left (make room for next pixel)
    add  a,(hl)       ; add source pixel (0 or 1) into lowest bit
    inc  l            ; advance to next source pixel
    ; ... possibly more INC L instructions here,
    ; depending on how many pixels to skip

    add  a,a          ; shift again
```

```

add a,(hl)      ; sample next pixel
inc l
inc l           ; skip one extra pixel (sphere curvature)

add a,a
add a,(hl)
inc l

; ... six more times, for 8 pixels total ...

```

The key detail: between each `add a,(hl)`, the number of `inc l` instructions varies. One pixel position might need a single `inc l` (sample adjacent pixels). Another might need three or four (skip over compressed regions of the projection). The lookup tables at \$6944 encode exactly how many `inc l` instructions to insert at each position.

Let us look more carefully at what happens with a single pixel:

```

add a,a          ; 4 T-states (shift A left by 1)
add a,(hl)       ; 7 T-states (add source pixel into bit 0)
inc l           ; 4 T-states (advance source pointer)

```

That is the minimum cost: 15 T-states to shift the accumulator and sample one pixel, plus 4 T-states for each source byte skipped. After eight such sequences, the accumulator holds a complete screen byte.

Notice that the source pointer is advanced using `inc l` rather than `inc hl`. This is deliberate. `INC HL` takes 6 T-states; `INC L` takes 4. By constraining the source data to sit within a single 256-byte page (so that only the low byte of the address changes), Dark saves 2 T-states per advance. When you are doing this thousands of times per frame, those 2 T-states add up.

There is a subtlety here that is easy to miss. The source image is stored as one byte per pixel, and `INC L` wraps around within a 256-byte page. This means each scanline of source data must fit within 256 bytes, and the source buffer must be page-aligned. The constraint shapes the entire memory layout of the demo.

Counting T-States

Introspec calculated the cost per output byte as:

101 + 32x T-states

where x is the average number of extra `INC L` instructions per pixel beyond the mandatory one. Let us verify this.

The fixed cost per pixel is:

Instruction	T-states
<code>add a,a</code>	4
<code>add a,(hl)</code>	7
<code>inc l</code>	4
Subtotal	15

For 8 pixels, the fixed cost is $8 \times 15 = 120$ T-states. But there is additional overhead per byte: the code must write the completed byte to screen memory and set up for the next. Let us assume the output sequence looks something like:

```
ld    (de),a      ; 7 T-states  (write screen byte)
inc   e           ; 4 T-states  (advance screen pointer)
```

There may also be setup for the accumulator (an `xor a` or similar) at the start of each byte. Taking Introspec's measured figure of 101 T-states as the fixed base cost per byte, the overhead beyond the raw pixel sampling accounts for roughly $101 - 120 = \dots$ which means the base figure already includes the output instructions and some of the pixel work is interleaved differently than the naive count suggests.

The cleaner way to read the formula: 101 T-states of fixed overhead (output, pointer management, any per-byte setup), plus 32 T-states per extra skip. The "32" comes from 8 pixels times 4 T-states per extra `INC L`, which gives us x as the average number of extra skips per pixel position in that byte. When the sphere is near the equator, x is small — the projection is close to uniform. Near the poles, x is large, and the rendering slows down. But the poles also have fewer bytes to draw (the sphere is narrower there), so the total workload roughly balances.

Is this fast enough? The Spectrum's frame is approximately 70,000 T-states (more on Pentagon: 71,680). A 56-pixel-diameter sphere occupies roughly 7 bytes across at its widest. Over the full height, perhaps 200–250 bytes need to be rendered. At 101 T-states per byte (equatorial, x near zero), that is roughly 25,000 T-states — comfortably within a single frame budget, with room left for screen clearing, table lookups, and all the other housekeeping. Even near the poles where x might average 2–3, the cost per byte rises to 165–197 T-states, but fewer bytes need drawing. The arithmetic works out. It fits.

The Code Generation Pass

Before the inner loop runs, a *code generation pass* constructs it. This pass reads the lookup tables at \$6944, which encode the sphere geometry for the current rotation angle, and emits Z80 instructions into a buffer:

1. For each scanline of the sphere, read the skip distances from the table.
2. Emit `add a,a` followed by `add a,(hl)` for each pixel.
3. Emit the appropriate number of `inc l` instructions based on the skip distance.
4. After every 8 pixels, emit the output instruction to write the accumulated byte to screen memory.
5. At the end of each scanline, emit a return or jump to the next line's handler.

The generated code block is then called directly. The CPU executes the instructions as if they were a normal subroutine, but they were written moments ago by the code generator. This is self-modifying code in the most literal sense — the program generates the program that draws the screen.

The code generation pass itself is not free, but it runs once per frame (or once per rotation step), while the generated inner loop runs hundreds of times. The amortized cost is negligible.

Texture-mapped sphere prototype — skip tables encode spherical projection, runtime code generation emits pixel sequences

Figure 15: Texture-mapped sphere prototype — skip tables encode spherical projection, runtime code generation emits pixel sequences

What Dark Knew: Spectrum Expert and the Building Blocks

There is a detail in this story that transforms it from a technical curiosity into a narrative arc. Dark — the coder behind Illusion’s sphere effect — is the same Dark who wrote the *Programming Algorithms* articles in Spectrum Expert #01, published in 1997.

Those articles cover multiplication (shift-and-add vs. square-table lookup), division (restoring and logarithmic), sine table generation via parabolic approximation, and Bresenham line drawing with optimized 8x8 matrix blocks. They are tutorial material, written for the ZX Spectrum programming community, explaining fundamental techniques that any demo coder would need.

And they are, quite precisely, the building blocks used in Illusion.

The sphere requires: trigonometric lookup tables for computing the projection (sine/cosine, the parabolic approximation from Dark’s article). Fixed-point multiplication for scaling. Careful memory layout for speed (the same cycle-counting discipline that Dark teaches throughout the articles). The skip-table approach to encoding sphere geometry is a direct application of the kind of precomputation-driven thinking that Dark advocates.

Dark wrote the demo first — Illusion won at ENLIGHT’96. Then, in 1997–98, he published the textbook that explained every technique he had used. Twenty years later, Introspec reverse-engineered the demo and found exactly the algorithms Dark had documented. We have both sides of the story: the practitioner explaining his methods after the fact, and the analyst confirming that those methods are precisely what the finished product contains.

The Hype Debate: Inner Loops vs. Mathematics

Introspec’s 2017 article on Hype sparked a long comment thread. Among the most substantive exchanges was a debate between kotsoft and Introspec about where the real work of an effect like this lies.

kotsoft argued that the mathematical approach to the projection — how you compute which source pixel maps to which screen position — is the critical design decision. Get the projection wrong, or use a naive algorithm, and no amount of inner-loop optimization will save you. The mathematical model determines whether the effect is even *feasible* on the hardware.

Introspec countered that the inner loop is where the T-states are actually spent. You can have a beautiful mathematical model, but if the rendering code costs 200 T-states per byte instead of 100, you have halved your frame rate. The mathematical approach defines *what* to compute; the inner loop determines *whether you can compute it in time*.

Both are right, and the tension between them illuminates something fundamental about demoscene coding. A demo effect is not pure mathematics and not pure engineering. It is the intersection: an elegant computation scheme (the sphere projection encoded as skip tables) married to an efficient execution strategy (generated unrolled loops with `INC L` advances). Neither alone is sufficient.

Introspec's summary captures it: "coder effects are always about evolving a computation scheme." The word *evolving* is key. You do not start with a textbook algorithm and optimize it until it fits. You evolve the algorithm and the implementation together, each constraining and enabling the other, until you find a form that works within the hardware's budget.

Practical: A Simplified 56x56 Spinning Sphere

Let us sketch how you would build a simplified version of this effect. We will target a 56x56 pixel sphere — 7 bytes wide at the equator, 56 scanlines tall. The goal is not to reproduce Illusion's full rendering engine, but to understand the core technique well enough to implement it.

Step 1: Precompute the Sphere Geometry

For each scanline y (from -28 to +27, centered on the sphere), compute the visible arc:

This gives the half-width of the sphere at that scanline. For each pixel position x within that arc, compute the corresponding longitude and latitude on the sphere surface:

$\text{longitude} = \arcsin(x / \text{radius_at_y}) + \text{rotation_angle}$

These give the (u, v) coordinates into the source texture for each screen pixel.

Step 2: Build Skip Tables

Rather than storing full (u, v) pairs for every pixel (prohibitively expensive in memory), compute the *difference in source position* between adjacent screen pixels. For each scanline, you need a list of skip values: how many source pixels to advance between consecutive screen samples.

Near the equator, consecutive screen pixels map to nearly adjacent source pixels — skips of 1. Near the poles, the projection compresses, and you skip over more source pixels — skips of 2, 3, or more.

Store these as a table. For our 56x56 sphere, you need at most 56 entries per line (the widest line), times 56 lines, times one byte per entry. That is at most 3,136 bytes for a single rotation angle — but in practice, you can exploit vertical symmetry (top half mirrors bottom half) and store only half the table.

For animation, you need skip tables for multiple rotation angles. With 32 rotation steps, you could fit the tables into $32 \times 1,568 =$ about 49KB. That overflows available memory, so in practice you would use fewer rotation steps, coarser angular resolution, or regenerate tables on the fly from a compact representation.

Step 3: Generate the Rendering Code

For each frame, read the skip table for the current rotation angle and generate Z80 code:

```
; Code generator pseudocode (in Z80 assembly, this would be
; a loop that writes opcodes into a buffer)
```

```
generate_sphere_code:
```

```
    ld    iy,skip_table      ; pointer to skip distances
    ld    ix,code_buffer     ; pointer to output code buffer
```

```
.line_loop:
```

```
    ; For each scan line...
```

```
    ld    b,bytes_this_line  ; number of output bytes (e.g. 7 at equator)
```

```
.byte_loop:
```

```
    ; For each output byte, emit code for 8 pixels:
```

```
    ld    c,8                ; 8 pixels per byte
```

```
.pixel_loop:
```

```
    ; Emit: ADD A,A
```

```
    ld    (ix+0),$87         ; opcode for ADD A,A
```

```
    inc    ix
```

```
    ; Emit: ADD A,(HL)
```

```
    ld    (ix+0),$86         ; opcode for ADD A,(HL)
```

```
    inc    ix
```

```
    ; Emit INC L instructions based on skip distance
```

```
    ld    a,(iy+0)           ; read skip distance
```

```
    inc    iy
```

```
.emit_inc_l:
```

```
    or     a
```

```
    jr     z,.pixel_done
```

```
    ld    (ix+0),$2C         ; opcode for INC L
```

```
    inc    ix
```

```
    dec    a
```

```
    jr     nz,.emit_inc_l
```

```
.pixel_done:
```

```
    dec    c
```

```
    jr     nz,.pixel_loop
```

```
    ; Emit: LD (DE),A (write byte to screen)
```

```
    ld    (ix+0),$12         ; opcode for LD (DE),A
```

```
    inc    ix
```

```
    ; Emit: INC E
```

```
    ld    (ix+0),$1C         ; opcode for INC E
```

```
    inc    ix
```

```

dec  b
jr   nz,.byte_loop

; Emit line transition code here (advance DE to next screen line)
; ...

jr   .line_loop

```

This is simplified — the actual Illusion code is more tightly integrated, and Dark likely used a more compact and efficient code generator. But the principle is the same: read skip distances, emit opcodes.

Step 4: Execute and Display

Once the code buffer is filled, call it as a subroutine:

```

ld  hl,source_image      ; source texture (page-aligned, 1 byte/pixel)
ld  de,screen_address    ; start of sphere area in video memory
call code_buffer         ; execute the generated rendering code

```

The generated code runs through the entire sphere, reading source pixels, packing them into screen bytes, and writing them to video memory. When it returns, the sphere is drawn.

Sphere outline rendered on the ZX Spectrum — monochrome texture mapped onto a rotating sphere using skip-table code generation

Figure 16: Sphere outline rendered on the ZX Spectrum — monochrome texture mapped onto a rotating sphere using skip-table code generation

For animation, increment the rotation angle, load the corresponding skip table (or regenerate it), regenerate the code, and render again.

Step 5: Source Image Layout

The source texture must be organized for fast sequential access. Since the rendering code uses `INC L` to advance through it, the texture must be page-aligned (start at an address where the low byte is \$00), and each row must fit within 256 bytes. A 256-pixel-wide texture stored as one byte per pixel fits this constraint perfectly: each row occupies one page.

For the monochrome case, each pixel is \$00 or \$01. This means `ADD A,(HL)` either adds 0 (pixel off) or 1 (pixel on) to the accumulator's lowest bit, right after `ADD A,A` has shifted everything up. The result is a bit-packed screen byte where each bit corresponds to a sampled source pixel.

The Larger Pattern

The sphere in Illusion is a specific instance of a general demoscene pattern that appears throughout this book. The pattern has three parts:

Precomputation. Expensive mathematical work — projection, trigonometry, coordinate transforms — is done once (or once per frame) and stored as compact tables. The tables encode *what* to render without encoding *how*.

Code generation. The rendering code itself is generated from the tables. This eliminates branches, loop counters, and conditional logic from the inner loop. Every instruction in the generated code does useful work. There is no overhead for “figuring out what to do next” — that decision was made during generation.

Sequential memory access. The inner loop reads data sequentially, advancing a pointer with single-byte increments. This is the fastest possible access pattern on the Z80, where register-indirect loads (`LD A, (HL)`) are cheap and indexed addressing (`LD A, (IX+d)`) is expensive.

The rotozoomer in the next chapter uses the same pattern. So does the dotfield scroller in Chapter 10. So do the attribute tunnels in Chapter 9. The specifics differ — different tables, different generated code, different data formats — but the architecture is the same. Introspec recognized this when he wrote that coder effects are about “evolving a computation scheme.” The sphere, the rotozoomer, the tunnel: they are all evolved from the same fundamental approach. The evolution is in the details — which computation, which table layout, which inner loop — but the skeleton is shared.

Dark understood this in 1996. He encoded it in his Spectrum Expert articles in 1997. Introspec confirmed it by disassembly in 2017. The pattern is as valid now as it was then, on any platform where T-states are scarce and every instruction must earn its keep.

Summary

- The sphere effect in Illusion maps a monochrome source image onto a rotating sphere using dynamically generated Z80 code.
- Lookup tables encode the sphere’s geometry as pixel skip distances. The rendering code is generated from these tables at runtime.
- The inner loop uses `ADD A,A` and `ADD A, (HL)` to accumulate pixels into screen bytes, with variable numbers of `INC L` instructions to advance through the source data.
- Performance: $101 + 32x$ T-states per output byte, where x depends on position.
- The approach exemplifies a general demoscene pattern: precompute geometry, generate code, access memory sequentially.
- Dark applied these algorithms in Illusion (1996), then documented them in Spectrum Expert (1997–98). Introspec reverse-engineered the result twenty years later, confirming the techniques.

Sources: Introspec, “Technical Analysis of Illusion by X-Trade” (Hype, 2017); Dark, “Programming Algorithms” (Spectrum Expert #01, 1997). The Hype comment thread includes contributions from kotsoft, Raider, and others.

Chapter 7: Rotozoomer and Chunky Pixels

“The trick is that you don’t rotate the screen. You rotate your walk through the texture.” – paraphrasing the core insight behind every rotozoomer ever written

There is a moment in *Illusion* where the screen fills with a pattern – a texture, monochrome, repeating – and then it begins to turn. The rotation is smooth and continuous, the zoom breathes in and out, and the whole thing runs at a pace that makes you forget you are watching a Z80 push pixels at 3.5 MHz. It is not the most technically demanding effect in the demo. The sphere (Chapter 6) is harder mathematically. The dotfield scroller (Chapter 10) is tighter in its cycle budget. But the rotozoomer is the one that looks effortless, and on the Spectrum, making something look effortless is the hardest trick of all.

This chapter traces two threads. The first is Introspec’s 2017 analysis of the rotozoomer from *Illusion* by X-Trade. The second is sq’s 2022 article on Hype about chunky pixel optimisation, which pushes the approach to 4x4 pixels and catalogues a family of rendering strategies with precise cycle counts. Together, they map the design space: how chunky pixels work, how rotozoomers use them, and the performance trade-offs that determine whether your effect runs at 4 frames per screen or 12.

What a Rotozoomer Actually Does

A rotozoomer displays a 2D texture rotated by some angle and scaled by some factor. The naive approach: for every screen pixel, compute its corresponding texture coordinate via a trigonometric rotation:

$$tx = -sx * \sin(\theta) * scale + sy * \cos(\theta) * scale + offset_x$$

At 256x192, that is 49,152 pixels each needing two multiplications. Even with a 54-T-state square-table multiply (Chapter 4), you exceed five million T-states – roughly 70 frames’ worth of CPU time. The effect is mathematically trivial and computationally impossible.

The key insight is that the transformation is *linear*. Moving one pixel right on screen always adds the same (dx, dy) to the texture coordinates. Moving one

pixel down always adds the same (dx', dy'). The per-pixel cost collapses from two multiplications to two additions:

Step down: $dx' = \sin(\theta) * scale, \quad dy' = \cos(\theta) * scale$

Start each row at the correct texture coordinate and step by (dx, dy) for every pixel. The inner loop becomes: read the texel, advance by (dx, dy), repeat. Two additions per pixel, no multiplications. The per-frame setup is four multiplications to compute the step vectors from the current angle and scale. Everything else follows from linearity.

This is the fundamental optimisation behind every rotozoomer on every platform. On the Amiga, on the PC, on the Spectrum.

Fixed-Point Stepping on the Z80

On a 16-bit or 32-bit platform, dx and dy would be fixed-point values: the integer part selects the texel, and the fractional part accumulates sub-pixel precision. On the Z80, we lack the registers and the bandwidth for true fixed-point inner loops. The classic Spectrum solution is to collapse the step to integer increments – always exactly +1, -1, or 0 per axis – and control the *ratio* of steps between axes to approximate the angle.

Consider a rotation of 30 degrees. The exact step vector would be (cos 30, -sin 30) = (0.866, -0.5). On a machine with fixed-point arithmetic, you would add 0.866 to the column coordinate and subtract 0.5 from the row coordinate per pixel. On the Z80, the inner loop instead alternates between two integer steps: some pixels step (+1 column, 0 rows) and others step (+1 column, -1 row). If you distribute these in a roughly 2:1 ratio – two column-only steps for every diagonal step – the average direction approximates the 0.866:0.5 ratio of a 30-degree walk. This is Bresenham's line algorithm applied to texture traversal.

The zoom factor determines how many texels you skip per screen pixel. At scale 1.0, every texel maps to one screen pixel. At scale 2.0, you skip every other texel, effectively zooming in. On the Spectrum, this is controlled by doubling the walk instructions: instead of one INC L per pixel, you execute two, stepping by 2 texels and producing a 2x zoom. Intermediate zoom levels again use Bresenham-like distribution: some pixels step by 1, others by 2, with the ratio controlled by an error accumulator.

The per-frame cost of computing these parameters is negligible: four lookups into a sine table, a few multiplications (or table lookups, see Chapter 4), and a Bresenham setup pass. All the heavy work is in the inner loop, which has been reduced to nothing but register increments and memory reads.

Chunky Pixels: Trading Resolution for Speed

Even at two additions per pixel, writing 6,144 bytes to the Spectrum's interleaved video memory per frame is impractical – not if you also want to update the angle and leave time for music. Chunky pixels solve this by reducing the effective resolution. Instead of one texel per screen pixel, you map one texel to a 2x2, 4x4, or 8x8 block.

Illusion uses 2x2 chunky pixels: effective resolution 128x96, a 4x reduction in work. The effect looks blocky up close, but at the speed the texture sweeps across the screen, motion hides the coarseness. The eye forgives low resolution when everything is moving.

Why 2x2 Is the Sweet Spot

The choice of chunk size involves a three-way tradeoff: visual quality, rendering speed, and memory. At 2x2, you get 128x96 effective pixels – enough to read text and recognise patterns in the texture. At 4x4, the 64x48 grid is noticeably coarser; fine details in the texture become unreadable, but the effect still “reads” as a coherent rotating surface. At 8x8, you are down to 32x24 blocks, which is the attribute grid resolution – any texture detail is lost, and the effect looks like coloured rectangles. The last case can be useful for colour-only effects (attribute tunnels, Chapter 9), but for a pixel-rendered rotozoomer, 2x2 or 4x4 is the practical range.

The memory cost matters too. Each chunky pixel stores one byte, so a 2x2 rotozoomer at 128x96 needs 12,288 texels per frame. With a 256-byte texture row (the natural width for 8-bit wrapping), the texture itself occupies 256 bytes per row times however many rows you need. A 4x4 version only processes 3,072 texels, which means the inner loop runs one-quarter as many iterations – but the visual cost is significant.

In practice, Spectrum demos land on 2x2 for featured rotozoomer effects and reserve 4x4 for situations where the rotozoomer shares the screen with other effects (bumpmapping overlays, split-screen compositions).

The \$03 Encoding Trick

The encoding is designed for the inner loop. Each chunky pixel is stored as \$03 (on) or \$00 (off). This value is not arbitrary – it encodes exactly the two low bits set: %00000011. Watch what happens as four pixels accumulate in the A register:

```
After 2x shift: A = %00001100          ($0C)
After pixel 2:  A = %00001100 + %00000011 ($0F)
After 2x shift: A = %00111100          ($3C)
After pixel 3:  A = %00111100 + %00000011 ($3F)
After 2x shift: A = %11111100          ($FC)
After pixel 4:  A = %11111100 + %00000011 ($FF)
```

If all four pixels are “on”, the result is \$FF – all bits set. If all four are “off” (\$00), the shifts and additions produce \$00. Mixed patterns produce the correct 2-bit-per-pixel stripe: for example, on-off-on-off gives %11001100 = \$CC. Each pair of bits in the output byte corresponds to one chunky pixel. Since each chunky pixel is 2 screen pixels wide (2x2), the 8-bit output byte covers exactly 8 screen pixels: four chunky columns times two pixels each.

The critical property: because we only ever add \$03 or \$00, there is no carry between pixel fields. The two-bit groups never overflow into each other. This is what makes the encoding branchless – no masking needed, no OR operations, just ADD A,A and ADD A,(HL).

The Inner Loop from Illusion

Introspec's disassembly reveals the core rendering sequence. HL walks through the texture; H tracks one axis and L the other:

```
; Inner loop: combine 4 chunky pixels into one output byte
ld  a,(hl)      ; 7T -- read first chunky pixel ($03 or $00)
inc  l          ; 4T -- step right in texture
dec  h          ; 4T -- step up in texture
add  a,a        ; 4T -- shift left
add  a,a        ; 4T -- shift left (now shifted by 2)
add  a,(hl)     ; 7T -- add second chunky pixel
```

The sequence repeats for the third and fourth pixels. The inc l and dec h together trace a diagonal path through the texture – and diagonal means rotated. The specific combination of increment and decrement instructions determines the rotation angle.

Step	Instructions	T-states
Read pixel 1	ld a,(hl)	7
Walk	inc l : dec h	8
Shift + Read pixel 2	add a,a : add a,a : add a,(hl)	15
Walk	inc l : dec h	8
Shift + Read pixel 3	add a,a : add a,a : add a,(hl)	15
Walk	inc l : dec h	8
Shift + Read pixel 4	add a,a : add a,a : add a,(hl)	15
Walk	inc l : dec h	8
Output + advance	ld (de),a : inc e	~11
Total per byte		~95

Introspec measured approximately 95 T-states per 4 chunks.

The critical observation: the walk direction is hardcoded into the instruction stream. A different rotation angle requires different instructions. Eight primary directions are possible using combinations of inc l, dec l, inc h, dec h, and nop. This means the rendering code changes every frame.

Self-Modifying Code at the Byte Level

“Per-frame code generation” sounds exotic, but the mechanism is mundane. Each walk instruction is a single byte in memory. INC L is opcode \$2C. DEC L is \$2D. INC H is \$24. DEC H is \$25. NOP is \$00. To change the walk direction from “right and up” (INC L + DEC H) to “pure right” (INC L + NOP), you write \$00 to the byte where \$25 currently sits. That is the entire code generation step: LD A,\$00 : LD (walk_target),A. A few stores into the instruction stream, and the inner loop now walks in a different direction.

The targets are known at assembly time. Each SMC site is labelled (e.g., .smc_walk_h_0:) and the patching code uses those labels as literal addresses. There is no dynamic memory allocation, no instruction parsing, no runtime disassembly. You are writing known opcodes to known addresses. The Z80 has no instruction

cache to invalidate, no pipeline to flush. The write takes effect immediately on the next fetch from that address.

In a fully unrolled inner loop (which Illusion uses for its 16-byte rows), there would be 64 walk-instruction sites to patch: 4 walk pairs per output byte times 16 bytes per row. Patching 64 bytes costs about $64 \times 13 = 832$ T-states (each LD (nn),A is 13 T-states), which is negligible compared to the 100,000+ T-states the rendering pass takes. The code generator is cheap. The generated code is what matters.

Per-Frame Code Generation

The rendering code is generated fresh each frame, with walk-direction instructions patched for the current angle:

Angle range	H step	L step	Direction
~0 degrees	nop	inc l	Pure right
~45 degrees	dec h	inc l	Right and up
~90 degrees	dec h	nop	Pure up
~135 degrees	dec h	dec l	Left and up
~180 degrees	nop	dec l	Pure left
~225 degrees	inc h	dec l	Left and down
~270 degrees	inc h	nop	Pure down
~315 degrees	inc h	inc l	Right and down

For intermediate angles, the generator distributes steps unevenly using Bresenham-like error accumulation. A 30-degree rotation alternates between inc l : nop and inc l : dec h at roughly a 2:1 ratio, approximating the 1.73:1 tangent of 30 degrees. The resulting code is an unrolled loop where each iteration has its own specific walk pair, tuned to the current angle.

The rendering cost for 128x96 at 2x2 chunky. The 128x96 area is 96 pixel rows, but each 2x2 texel covers two pixel rows, giving 48 texel rows. Each texel row produces 16 output bytes (128 pixels / 8 bits per byte, with 4 chunky pixels packed per byte):

1,520 x 48 texel rows = 72,960 T-states total

Roughly 1 frame on a Pentagon (71,680 T-states per frame). But this is the bare inner loop only. A complete accounting adds:

Row setup (per row):	~ 800 T	(48 rows x ~17 T each)
Buffer-to-screen copy:	~ 20,000 T	(stack trick, 1,536 bytes)
Sine table lookups:	~ 200 T	
Frame overhead:	~ 500 T	(HALT, border, angle update)

Inner loop:	72,960 T	
Total per frame:	~ 95,460 T	(= 1.33 Pentagon frames)

On a standard 48K/128K Spectrum at 69,888 T-states per frame, the rendering takes roughly 1.4 frames. Introspec's estimate of 4-6 frames per screen accounts for the more complex code path in Illusion (which handles the full 256x192 screen,

not just a 128x96 strip) and the cost of the music engine running in the interrupt. On a Pentagon with its slightly longer frame (71,680 T-states) and no contention, the inner loop runs about 3% faster.

Memory contention on the 48K/128K Spectrum adds another hidden cost. During the top 192 scanlines, the ULA steals cycles from the CPU when accessing the lower 16KB of RAM (\$4000-\$7FFF). The inner loop reads from the texture (which should be above \$8000, out of contended memory) and writes to a buffer (also above \$8000), so it avoids contention entirely. The buffer-to-screen transfer, however, writes directly to video RAM and will be slowed by contention if it overlaps with the display period — expect roughly **20% more T-states** for the transfer on a 48K/128K machine compared to Pentagon. This is why demos synchronise the screen transfer to the border period or to the bottom of the display. See Chapter 15.2 for the full contention model and scheduling strategies.

Buffer to Screen Transfer

The rotozoomer renders into an off-screen buffer, then transfers to video memory. The interleaved screen layout makes direct rendering painful, and buffering avoids tearing.

The transfer uses the stack:

```
pop  hl                ; 10T -- read 2 bytes from buffer
ld   (screen_addr),hl  ; 16T -- write 2 bytes to screen
```

Screen addresses are embedded as literal operands, pre-calculated for the Spectrum's interleaving – another instance of code generation. At 26 T-states per two bytes, a full 1,536-byte transfer costs under 20,000 T-states. The rendering pass is the bottleneck, not the transfer.

Deep Dive: 4x4 Chunky Pixels (sq, Hype 2022)

sq's article pushes chunky pixels to 4x4 – effective resolution 64x48. The visual result is coarser, but the performance gain opens up effects like bumpmapping and interlaced rendering. The article is a study in optimisation methodology: start straightforward, iteratively improve, measure at each step.

Approach 1: Basic LD/INC (101 T-states per pair). Load chunky value, write to buffer, advance pointers. The bottleneck is pointer management: INC HL at 6 T-states adds up over thousands of iterations.

Approach 2: LDI variant (104 T-states – slower!). LDI copies a byte and auto-increments both pointers in one instruction. But it also decrements BC, consuming a register pair. The save/restore overhead makes it *slower* than the naive approach. A cautionary tale: on the Z80, the “clever” instruction is not always the fast one.

Approach 3: LDD dual-byte (80 T-states per pair). By arranging source and destination in reverse order, LDD's auto-decrement works in your favour. A combined two-byte sequence exploits this for a 21% improvement over baseline.

Approach 4: Self-modifying code (76-78 T-states per pair). Pre-generate 256 rendering procedures, one per possible byte value, each with the pixel value baked in as an immediate operand:

```
; One of 256 pre-generated procedures
proc_A5:
    ld    (hl), $A5      ; 10T -- value baked into instruction
    inc   l              ; 4T
    ld    (hl), $A5      ; 10T -- 4x4 block spans 2 bytes horizontally
    ; ... handle vertical repetition ...
    ret                ; 10T
```

The 256 procedures occupy approximately 3KB. Per-pixel rendering drops to 76-78 T-states – 23% faster than baseline, 27% faster than LDI.

Performance Comparison

Approach	Cycles/pair	Relative	Memory
Basic LD/INC	101	1.00x	Minimal
LDI variant	104	0.97x	Minimal
LDD dual-byte	80	1.26x	Minimal
Self-modifying (256 procs)	76-78	1.30x	~3KB

The self-modifying approach wins, but the margin over LDD is narrow. In a 128K demo, 3KB is easily available. In a 48K production, the LDD approach might be the better engineering decision.

Historical Roots: Born Dead #05 and the Scene Lineage

sq notes these techniques build on work published in Born Dead #05, a Russian demoscene newspaper from approximately 2001. Born Dead was one of several Russian-language disk magazines that served as technical journals for the ZX Spectrum demoscene. Unlike Western PC demoscene publications that could assume 486-class hardware, the Spectrum magazines operated under the constraints of a community that was still actively developing new techniques for a machine from 1982. The foundational article described basic chunky rendering – the idea that you could treat the Spectrum’s bit-mapped display as a lower-resolution chunky-pixel buffer and gain speed at the expense of resolution.

sq’s contribution, twenty-one years later, was the systematic optimisation and the pre-generated procedure variant. But between Born Dead #05 and sq’s 2022 article, the chunky rotozoomer appeared in numerous Spectrum demos. X-Trade’s Illusion (ENLiGHT’96) was among the earliest full implementations. Other notable examples include Exploder^XTM’s GOA4K and Refresh, 4D’s productions, and later work from the Russian and Polish scenes. The technique spread partly through disassembly – Introspec’s 2017 analysis of Illusion is itself an example of the scene’s tradition of learning by reverse engineering – and partly through the informal knowledge network of disk magazines, BBS postings, and direct communication between coders.

This is how scene knowledge evolves: a technique surfaces in an obscure disk magazine, circulates within the community, and twenty-one years later someone

revisits it with fresh measurements and new tricks. The chain from Born Dead to sq to this chapter is unbroken.

Practical: Building a Simple Rotozoomer

Here is the structure for a working rotozoomer with 2x2 chunky pixels and a checkerboard texture.

Texture. A 256-byte page-aligned table where each byte is \$03 or \$00, generating 8-pixel-wide stripes. The H register provides the second dimension; XORing H into the lookup creates a full checkerboard:

```

ALIGN 256
texture:
    LUA ALLPASS
    for i = 0, 255 do
        if math.floor(i / 8) % 2 == 0 then
            sj.add_byte(0x03)
        else
            sj.add_byte(0x00)
        end
    end
end
ENDLUA

```

Sine table and per-frame setup. A 256-entry page-aligned sine table drives the rotation. Each frame reads `sin(frame_counter)` and `cos(frame_counter)` (cosine via a 64-index offset) to compute the step vectors, then patches the inner loop's walk instructions with the correct opcodes.

The rendering loop. The outer loop sets the starting texture coordinate for each row (stepping perpendicular to the walk direction). The inner loop walks through the texture:

```

.byte_loop:
    ld    a,(hl)           ; read texel 1
    inc   l                ; walk (patched per-frame)
    add   a,a : add a,a    ; shift
    add   a,(hl)           ; read texel 2
    inc   l                ; walk
    add   a,a : add a,a    ; shift
    add   a,(hl)           ; read texel 3
    inc   l                ; walk
    add   a,a : add a,a    ; shift
    add   a,(hl)           ; read texel 4
    inc   l                ; walk
    ld    (de),a           ; write output byte
    inc   de
    djnz  .byte_loop

```

The `inc l` instructions are the targets of the code generator. Before each frame, they are patched to the appropriate combination of `inc l/dec l/inc h/dec h/nop` based on the current angle. For non-cardinal angles, a Bresenham error accumulator

distributes the minor-axis steps across the row, so each walk instruction in the unrolled loop may be different from its neighbours.

Rotozoomer output — the texture rotates and scales in real-time, rendered with 2x2 chunky pixels

Figure 17: Rotozoomer output — the texture rotates and scales in real-time, rendered with 2x2 chunky pixels

Main loop. HALT for vsync, compute step vectors, generate walk code, render to buffer, stack-copy buffer to screen, increment frame counter, repeat.

Texture Design and Boundary Handling

The texture is the most constrained data structure in the rotozoomer. Every design decision in the inner loop – the page alignment, the wrapping behaviour, the power-of-two sizing – traces back to how the texture is laid out in memory.

Why Page-Aligned, Why 256 Columns

The texture is page-aligned so that H selects the row and L selects the column. This is not merely convenient; it makes the inner loop possible. INC L and DEC L wrap at the 256-byte page boundary automatically – when L overflows from \$FF to \$00, H is unchanged. The texture wraps horizontally for free, with zero branch overhead. If the texture were not page-aligned, L increments would carry into H, corrupting the row address. You would need explicit masking (AND \$3F after every step), which would add 4-8 T-states per pixel and destroy the tight inner loop.

The vertical axis (H) also wraps, but over the full range of rows allocated to the texture. If you allocate 64 rows (pages), H ranges from the texture base page to base+63. INC H and DEC H will happily walk past the end of the texture into whatever memory follows. Illusion handles this by masking H to the texture height at the start of each row (not per pixel – per-pixel masking would be too expensive). This works because within a single 16-byte row, the H coordinate changes by at most 16 steps, and if the texture is tall enough relative to the row width, an overflow within a row cannot reach memory that produces visual garbage. A 64-row texture with 16 H-steps per row has a comfortable margin.

Choosing Texture Size

The texture must be a power-of-two in width (always 256, since L is 8 bits) and ideally a power-of-two in height for easy masking. Common choices:

- **256x256** (64KB): fills all of upper RAM on a 128K Spectrum. Maximum resolution, but leaves no room for code or buffers.
- **256x64** (16KB): the practical choice. Fits in one 16KB bank on 128K hardware. The 6-bit height mask (AND \$3F) is fast and tiles seamlessly.
- **256x32** (8KB): fits on a 48K Spectrum with room for everything else. The texture repeats more visibly, but for a checkerboard or stripe pattern, repetition is the design.
- **256x16** (4KB): minimal. Works for very simple patterns like single-axis stripes.

For non-repeating textures (images, logos), the height should be at least as large as the effective screen height divided by the scale factor. A 2x2 rotozoomer with 96 effective rows needs at least 96 texture rows to avoid visible tiling when the zoom is at 1:1. At higher zoom levels, fewer rows are needed because the camera is “closer” to the texture surface.

What About Screen Boundaries?

The Spectrum’s 256x192 screen is 32 bytes wide by 192 lines. If your rotozoomer fills a 128x96 strip in the centre, you never approach the edge of video memory. But a full-screen rotozoomer at 256x192 (or even 128x192 with 2x2 chunky) must handle the case where the output address reaches the attribute area at \$5800. The simplest approach: render into a buffer and only copy the portion that fits. A more aggressive approach: clip the row count to the visible area during code generation, which avoids wasted computation but adds complexity to the row loop.

In practice, most Spectrum rotozoomers render a strip smaller than the full screen. The visual framing – a border, a title bar, a music credit – hides the cropping and buys back T-states for other effects.

The Design Space

The chunky pixel size is the most consequential design decision in a rotozoomer:

Parameter	2x2 (Illusion)	4x4 (sq)	8x8 (attributes)
Resolution	128x96	64x48	32x24
Texels/frame	12,288	3,072	768
Inner loop cost	~73,000 T	~29,000 T	~7,300 T
Frames/screen	~1.3	~0.5	~0.1
Visual quality	Good motion	Chunky but fast	Very blocky
Use case	Featured effects	Bumpmapping, overlays	Attribute-only FX

The 4x4 version fits within a single frame with room for a music engine and other effects. The 2x2 version takes roughly 1.3-1.5 frames (including overhead) but looks substantially better. The 8x8 case is the attribute tunnel from Chapter 9.

Once you have a fast chunky renderer, the rotozoomer is just one application. The same engine drives **bumpmapping** (read height differences instead of raw texels, derive shading), **interlaced effects** (render odd/even rows on alternating frames, doubling effective frame rate at the cost of flicker), and **texture distortion** (vary the walk direction per row for wavy or ripple effects). A 4x4 rotozoomer can share a frame with a scrolltext, a music engine, and a screen transfer. sq’s work was motivated by exactly this versatility.

Three Approaches to Texture Rotation

Everything above treats the rotozoomer as one technique with a tuneable chunk size. But “rotozoomer” on the Spectrum is really a family of three distinct approaches, each with different inner loops, different visual character, and different performance profiles. They share the same mathematical foundation – the linear step vectors, the Bresenham-style angle distribution – but diverge completely at the rendering level.

Variant 1: Monochrome Bitmap (Full Pixel Resolution)

The purest form: every screen pixel maps to one texel. The texture is monochrome – one bit per pixel – so reading a texel means testing a single bit, and writing to the screen means setting or clearing a single bit. No chunky encoding, no block grouping. The result is a rotated texture at the full 256x192 resolution of the Spectrum display.

The inner loop skeleton looks something like this:

```
; For each screen pixel:
; DE = texture pointer, HL = screen pointer
    ld  a,(de)          ; 7T  -- read texture byte
    and n               ; 7T  -- test texture bit at current coords
    jr  z,.pixel_off    ; 12/7T
    set m,(hl)          ; 15T  -- set screen bit
    jr  .pixel_done     ; 12T
.pixel_off:
    res m,(hl)          ; 15T  -- clear screen bit
.pixel_done:
    ; advance texture coords (inc e / dec d / etc.)
    ; advance screen bit position
    ; ... next pixel
```

Note that SET and RES only work with (HL), (IX+d), or (IY+d) – not (DE) or (BC). This forces HL to serve as the screen pointer, while DE handles the texture coordinates.

The per-pixel cost is brutal: 35-45 T-states minimum, with branching on every pixel. Across 49,152 pixels, that is 1.5 to 2 million T-states for the rendering pass alone – roughly 21-28 frames on a standard Spectrum. A full-screen monochrome rotozoomer at 50fps is not happening.

But nobody said you need to fill the whole screen. The technique shines when applied to a smaller region – a 128x64 strip, a circular viewport, a masked area – or when you accept a lower frame rate in exchange for the visual impact of full-resolution rotation. It also works beautifully for distortion effects where the “rotation” is not uniform: varying the step vectors per scanline produces wave distortions, barrel effects, and the “sonic ripple” look seen in parts of *Illusion by Dark/X-Trade*. The coordinate mapping is no longer a simple rotation but a per-line warp through the texture. The maths is the same – fixed-point stepping along a direction – but the direction itself changes every row.

The visual payoff is striking. Where a 2x2 chunky rotozoomer looks like a rotating mosaic, the monochrome bitmap version looks like a rotating *image*. On a machine

where every effect fights the same 69,888 T-state budget, dedicating multiple frames to full-resolution rendering is a deliberate aesthetic choice.

Variant 2: Chunky Rotozoomer (2x2 or 4x4 Blocks)

This is the technique covered in the bulk of this chapter. Each screen block (2x2 or 4x4 pixels) maps to one texel. The \$03/\$00 encoding, the add a,a : add a,(hl) accumulation, the walk-instruction patching – all of it targets this approach.

At 2x2 (128x96 effective resolution), the inner loop runs at approximately 95 T-states per output byte, producing the smooth, recognisable rotozoomer seen in Illusion. At 4x4 (64x48), sq's pre-generated procedure variant eliminates the loop overhead entirely, bringing the cost down to 76-78 T-states per output pair and leaving room for multi-effect compositions within a single frame.

The chunky rotozoomer occupies the middle ground: fast enough for real-time, detailed enough to carry a featured effect. It is the workhorse of the Spectrum rotozoomer repertoire.

Variant 3: Attribute Rotozoomer (8x8 Block “Pixels”)

The Spectrum's attribute area at \$5800-\$5AFF stores colour information for each 8x8 pixel character cell: 32 columns by 24 rows, 768 bytes total. Each byte encodes INK, PAPER, BRIGHT, and FLASH for a single 8x8 block. The attribute rotozoomer ignores the bitmap entirely and treats these 768 attribute cells as the display surface. Each cell becomes one “pixel” in a 32x24 image.

The inner loop is structurally identical to the chunky version – step through texture coordinates, read a value, write it to the output – but the output is the attribute area, and the “texel” value is a colour attribute byte rather than a bit pattern. The effective resolution is just 32x24, which means the entire rendering pass is 768 iterations of the stepping loop.

The maths:

```
~10 T-states per cell (read texel + write attribute + step)
768 x 10 = ~7,680 T-states total
```

That is roughly 11% of a single frame. You could run the attribute rotozoomer nine times over and still have room for a music engine. The cost is so low that the effect is essentially free.

But the visual payoff is different from the bitmap variants. You are not rotating pixels – you are rotating coloured blocks. At 32x24, fine detail is invisible. What you get instead is a sweeping field of colour, a vivid mosaic that turns and breathes. The attribute rotozoomer in Illusion uses exactly this: a boldly coloured texture (not a monochrome bitmap) mapped through the attribute grid, producing the characteristic “stained glass” look of rotating colour fields that Illusion is known for. The PAPER and INK fields in each attribute byte give you two colours per cell, so a carefully designed texture can pack more visual information than the raw resolution suggests.

The attribute rotozoomer is perfect for backgrounds, transitions, or as a base layer with pixel-level effects composited on top. Because it only writes to the attribute area, the bitmap can be used simultaneously for a different effect – a scroller, a

logo, a particle field – running at its own pace. This layered approach is a hallmark of multi-effect demo screens on the Spectrum.

Comparison

Vari- ant	Effective resolution	Bytes written/frame	~T-states (render)	Colour	Typical use
Monochrome bitmap	256x192 (or subregion)	6,144 (full screen)	1,500,000- 2,000,000	1-bit	Hero effect, distortion, warp
Chunky 2x2	128x96	1,536	~73,000	1-bit	Featured roto- zoomer
Chunky 4x4	64x48	384	~29,000	1-bit	Multi- effect, overlay
At- tribute	32x24	768	~7,700	INK+PA- PER (2 colours/cell)	Back- ground, colour cell transition

The progression from top to bottom is a smooth trade: resolution for speed, detail for headroom. The monochrome bitmap gives you everything the Spectrum's display can show, at a cost that demands dedication. The attribute version gives you almost nothing in resolution, but it runs so fast that the rotozoomer becomes just another instrument in a multi-effect composition rather than the main event.

All four rows in this table share the same core algorithm. The step vectors are computed the same way. The Bresenham distribution works the same way. The difference is only where you write and how many iterations you run. Once you have built one rotozoomer, you have built all of them.

The Rotozoomer in Context

The rotozoomer is not a rotation algorithm. It is a *memory traversal pattern*. You walk through a buffer in a straight line, and the walk direction determines what you see. Rotation is one choice of direction. Zoom is a choice of step size. The Z80 does not know trigonometry. It knows INC L and DEC H. Everything else is the programmer's interpretation.

In Illusion, the rotozoomer sits alongside the sphere and the dotfield scroller. All three share the same architecture: precomputed parameters, generated inner loops, sequential memory access. The sphere uses skip tables and variable INC L counts. The rotozoomer uses direction-patched walk instructions. The dotfield uses stack-based address tables. Three effects, one engine philosophy.

Dark built all of them. Introspec traced all of them. The pattern that connects them is the lesson of Part II: compute what you need before the inner loop starts, generate code that does nothing but read-shift-write, and keep the memory access sequential.

Summary

- A rotozoomer displays a rotated and zoomed texture by walking through it at an angle. Linearity reduces per-pixel cost from two multiplications to two additions.
 - Chunky pixels (2x2, 4x4) reduce effective resolution and rendering cost proportionally. Illusion uses 2x2 at 128x96; sq's system uses 4x4 at 64x48.
 - Illusion's inner loop: `ld a,(hl) : add a,a : add a,a : add a,(hl)` with walk instructions between reads. Cost: ~95 T-states per byte for 4 chunky pixels.
 - Walk direction changes per frame, requiring code generation – the rendering loop is patched before each frame.
 - sq's 4x4 optimisation journey: basic LD/INC (101 T-states) to LDI (104 T-states, slower) to LDD (80 T-states) to self-modifying code with 256 pre-generated procedures (76-78 T-states, ~3KB). Based on earlier work in Born Dead #05 (~2001).
 - Buffer-to-screen transfer via `pop hl : ld (nn),hl` at ~26 T-states per two bytes.
 - The rotozoomer shares its architecture with the sphere (Chapter 6) and dotfield (Chapter 10): precomputed parameters, generated inner loops, sequential memory access.
-

Sources: Introspec, “Technical Analysis of Illusion by X-Trade” (Hype, 2017); sq, “Chunky Effects on ZX Spectrum” (Hype, 2022); Born Dead #05 (~2001, original chunky pixel techniques).

Chapter 8: Multicolor — Breaking the Attribute Grid

“The multicolor will be conquered.” — DenisGrachev, Hype, 2019

Every ZX Spectrum coder knows the rule. Two colours per 8x8 cell. Ink and paper. That is what the ULA gives you, and that is all it gives you. If your character has a red hat and a blue coat and the two colours fall inside the same attribute cell, one of them loses. The result — garish fringing, miscoloured sprites, characters that change colour as they walk past scenery — is attribute clash, and it is the defining visual constraint of the platform.

Attribute clash is so fundamental to the Spectrum’s identity that many coders simply accept it. They design around it. They pick their palettes to minimise it. They restrict sprite sizes or avoid certain colour combinations. For thirty years, the 8x8 grid has been a fact of life.

But the ULA does not know this.

The ULA reads attribute bytes as it draws the screen, one scanline at a time. It does not read all 768 attribute bytes at once. It reads each row of 32 attributes exactly when it needs them, eight scanlines later it reads the same row again for the next pixel line within that character row, and so on. The attribute for any given cell is read eight times per frame — once for each pixel row in that cell.

The trick is obvious once you see it: if you change the attribute byte between reads, the ULA will apply a different colour to different pixel rows within the same cell. Instead of two colours for all eight rows, you get two colours per *group of rows*. The 8x8 grid does not break because the hardware was redesigned. It breaks because you rewrote the data faster than the hardware could consume it.

This is multicolor. It has been known since at least the early 2000s, when the Russian ZX magazine Black Crow published an algorithm and example code in its fifth issue. But for years, multicolor remained a curiosity — impressive in demos, impractical in games, because the CPU spent so many T-states changing attributes that nothing was left for game logic.

Normal 8x8 attributes vs multicolour 8x2 beam-racing — the right panel shows how rewriting attributes between scanlines multiplies colour resolution

Figure 18: Normal 8x8 attributes vs multicolour 8x2 beam-racing — the right panel shows how rewriting attributes between scanlines multiplies colour resolution

Then DenisGrachev figured out how to make games with it.

The ULA's Perspective

To understand multicolor, you need to see the screen from the ULA's point of view.

The ULA draws 192 visible scanlines per frame, top to bottom. Each scanline takes 224 T-states (on Pentagon). For each scanline, the ULA reads 32 pixel bytes and 32 attribute bytes from memory. The pixel bytes determine which dots are ink and which are paper. The attribute bytes determine which colours “ink” and “paper” actually are.

Within a character row (8 pixel lines), the ULA reads the same 32 attribute bytes for every scanline. It does not cache them — it reads them fresh each time. This means you have a window of opportunity between one scanline's attribute read and the next to change the attribute data.

The “traditional” multicolor approach exploits this directly. After a HALT (which synchronises the CPU with the frame interrupt), you count T-states to know exactly when the ULA will read each attribute row. Then, in the gap between reads, you overwrite the attribute bytes with new values. When the ULA reads them for the next scanline, it sees the new colours.

The limitation is brutal: counting T-states precisely, changing 32 bytes between scanlines, then waiting for the next opportunity. The CPU spends almost all its time on this bookkeeping. In a typical traditional multicolor engine, you can change attributes every 2 or 4 scanlines, giving 8x2 or 8x4 colour resolution. But the cycle budget consumed by the attribute-changing code itself leaves almost nothing for game logic, sprite rendering, or sound.

This is why multicolor stayed in demos. Demos can afford to spend 100% of the CPU on visual effects. Games cannot.

The LDPUSH Insight

In January 2019, DenisGrachev published an article on Hype titled “Multicolor Will Be Conquered” (Mul'tikolor budet pobezhdon). The title was a statement of intent. He had been developing Old Tower, a game for the ZX Spectrum with 8x2 multicolor — attributes changing every two pixel rows — and he wanted to explain how he had solved the cycle budget problem that made multicolor impractical in games.

The key insight is one of those ideas that seems inevitable in hindsight: the code that outputs pixel data is the display buffer.

Traditional multicolor separates code and data. You have a buffer of attribute bytes somewhere in memory, and rendering code that copies them to the screen at the right moment. DenisGrachev's technique fuses the two. The “buffer” is a sequence of Z80 instructions — specifically, LD DE,nn followed by PUSH DE — and executing those instructions writes the display data directly to screen memory via the stack pointer.

Here is how it works.

LD DE,nn / PUSH DE

The instruction LD DE,nn loads a 16-bit immediate value into register pair DE. It takes 10 T-states and is 3 bytes long: the opcode byte \$11, followed by two data bytes (the value to load, low byte first). The instruction PUSH DE decrements SP by 1, writes the high byte of DE, decrements SP by 1 again, then writes the low byte. The result: SP ends up 2 lower, with the high byte at the higher address and the low byte at the lower address. It takes 11 T-states.

Together, LD DE,nn : PUSH DE costs 21 T-states, is 4 bytes long, and writes 2 bytes of data to the screen. The “data” is the immediate operand of the LD instruction. To change what gets drawn, you do not rewrite a display buffer — you patch the operand bytes inside the LD instructions themselves.

```
; One LDPUSH pair: writes 2 bytes to screen memory
ld    de,$AA55      ; 10 T  load pixel data
push de             ; 11 T  write to (SP), SP = SP - 2
                    ; ---
                    ; 21 T total, 2 bytes output
```

A scanline of pixel data is 32 bytes wide. But you cannot fill all 32 bytes with PUSH alone, because PUSH writes downward (SP decrements) and you need the data to appear left-to-right on screen. The Spectrum’s screen memory layout handles this: within a single scanline, consecutive bytes are at ascending addresses. PUSH writes to descending addresses. So the data comes out backwards — the last byte pushed appears at the lowest address, which is the leftmost byte on screen.

This means you build your LDPUSH sequence in reverse display order. The first LD DE,nn : PUSH DE in the code writes the rightmost two bytes of the scanline. The last writes the leftmost two bytes. When executed, the pushes fill the scanline right-to-left.

How much fits in a scanline?

DenisGrachev does not fill all 32 bytes. The GLUF engine (which powers Old Tower and other games) uses a 24-character-wide game area, padded with a border on each side. That is 24 bytes per scanline in the game area.

At 4 bytes of code per 2 bytes of output, you need 48 bytes of code to fill 24 bytes of screen. But there is an additional byte for the initial LD SP,nn setup and for the attribute change between scanline groups. DenisGrachev reports 51 bytes per scanline of generated code.

The beauty of this approach: there is no separate rendering pass. The instructions that populate the screen ARE the executable code. When you need to update a tile or sprite, you patch the immediate operands of the LD instructions. When the display code runs, it outputs the patched data. Code is data. Data is code.

The stack pointer as a cursor

During execution, SP is pointed at the right edge of the current scanline in screen memory. Each PUSH writes 2 bytes and decrements SP by 2, moving the “write

cursor” leftward across the scanline. At the end of one scanline’s output, SP points to the left edge. The code then adjusts SP to the right edge of the next scanline and repeats.

The fundamental constraint: interrupts must be disabled while SP is hijacked. If an interrupt fired, the CPU would push the return address into screen memory, corrupting the display. This means the multicolor rendering code runs with DI and re-enables interrupts with EI only after SP is restored to the real stack. The entire rendering pass — all 192 visible scanlines — happens in one uninterruptible block.

Multicolor and contention (48K/128K Spectrum). The LDPUSH technique writes directly to screen memory via SP, so every PUSH hits contended addresses (\$4000–\$5AFF). On a standard 48K or 128K Spectrum, the ULA inserts wait states during the active display period, adding ~1 T per PUSH. More critically, cycle-exact beam racing — aligning attribute changes to specific scanlines — becomes unreliable because contention delays shift the timing by variable amounts depending on the ULA’s pixel clock phase. DenisGrachev’s GLUF engine was developed and tested on Pentagon, where contention does not exist. **On non-Pentagon machines**, the Ringo dual-screen approach (switching pages via port \$7FFD between scanline groups) is more robust: it does not require cycle-exact timing because the page switch takes effect at the next scanline boundary regardless of when the OUT instruction executes. Treat LDPUSH multicolor as a Pentagon-targeted technique unless you can verify timing on your target hardware. See Chapter 15.2 for the contention model.

The GLUF Engine: Multicolor in a Real Game

Old Tower was a proof of concept. GLUF (DenisGrachev’s game framework) was the production engine. Here are the numbers:

Parameter	Value
Multicolor resolution	8x2 (attributes change every 2 pixel rows)
Game area	24x16 characters (192x128 pixels)
Buffering	Double-buffered (two sets of display code)
Sprite size	16x16 pixels
Tile size	16x16 pixels
Sound	25 Hz (every other frame)
Cycles per frame	~70,000 (nearly the entire Pentagon budget)

The 8x2 resolution means the engine changes attributes four times per character row instead of once. Each character row is 8 scanlines tall; changing attributes every 2 scanlines gives four distinct colour bands within a single character cell. A character that would normally be limited to ink-on-paper suddenly has up to eight colours (two per band, four bands). In practice, the effect is striking — sprites and tiles display far more colour detail than the Spectrum’s architecture was designed to allow.

Double buffering

GLUF maintains two complete sets of LDPUSH display code. While one set executes (drawing the current frame), the other is being patched with new tile and sprite data for the next frame. This eliminates the flicker that would result from modifying display code while it is executing.

The cost is memory. Each set of display code covers the 24x16 character game area at 51 bytes per scanline times 128 scanlines: roughly 6,500 bytes per buffer. Two buffers consume about 13,000 bytes. On a 128K Spectrum with banked memory, this is manageable but significant — it means careful memory planning for the rest of the game's assets.

The two-frame architecture

Here is where the engineering gets brutal. GLUF does not render a complete frame every 1/50th of a second. It uses a two-frame architecture:

Frame 1: Change attributes for the multicolor effect, then render as many tiles as possible into the display code buffer. The attribute changes are the time-critical part — they must happen at precisely the right moment relative to the raster beam.

Frame 2: Finish rendering any remaining tiles, then overlay sprites onto the display code buffer.

The split is necessary because the workload simply does not fit in one frame. Rendering tiles into the LDPUSH buffer means patching operand bytes inside the display code — for each tile pixel, you calculate which LD instruction it affects and write the new value into the right byte offset. This is not a simple block copy. The interleaved structure of the display code (opcode-data-data-opcode-data-data...) means tile rendering involves scattered writes, not sequential ones.

The total rendering budget is approximately 70,000 T-states per frame — nearly the entire Pentagon budget of 71,680. What remains is barely enough for game logic, input handling, and the periodic sound update.

Sound runs at 25 Hz instead of 50 Hz. Every other frame, the engine skips the sound update entirely to reclaim those T-states for rendering. The player does not notice the halved update rate for simple sound effects. For music, the 25 Hz rate means each note lasts twice as many frames, which requires the music engine to be written specifically for this constraint.

What the player sees

The player does not see any of this. They see a side-scrolling game with more colour than a Spectrum should have. Sprites move over scenery without the usual attribute clash. Tiles display shading, texture, and multi-coloured detail that would be impossible with standard 8x8 attributes. The game feels like it belongs on a more capable platform.

This is the point DenisGrachev's work makes forcefully: multicolor is not a demo trick. It is a game engine technique. The engineering is extreme — double buffers, two-frame rendering, 25 Hz sound — but the result is a playable game with visuals that genuinely break the Spectrum's perceived limits.

Multicolour effect — 8x2 attribute resolution gives each character cell up to eight colours instead of two

Figure 19: Multicolour effect — 8x2 attribute resolution gives each character cell up to eight colours instead of two

Ringo: A Different Kind of Multicolor

In December 2022, DenisGrachev published a second article on Hype: “Ringo Render 64x48.” Where GLUF extended the attribute grid to 8x2, Ringo threw it away entirely and built something closer to a chunky-pixel framebuffer with per-pixel colour.

The approach is conceptually simple and technically devious.

The 11110000b pattern

Fill every pixel byte in screen memory with the value \$F0 — binary 11110000. The left four pixels of each byte are set (ink colour), and the right four are clear (paper colour). Now, if you change the attribute for that cell, the left half displays the ink colour and the right half displays the paper colour. A single 8-pixel-wide cell displays two distinct colours, side by side.

With standard 8x8 attributes, this gives you a 64x24 grid of independently coloured “pixels,” each 4 real pixels wide and 8 real pixels tall. Not bad, but not revolutionary.

Two-screen switching

The trick that makes Ringo work: the ZX Spectrum 128K has two screen buffers. Screen 0 lives at \$4000, Screen 1 at \$C000 (in bank 7). A single OUT to port \$7FFD switches which screen the ULA displays.

Ringo prepares both screens with the 11110000b pattern but with *different* attributes on each screen. It then switches between the two screens every 4 scanlines. The effect: within each 8-pixel-tall character row, the top 4 scanlines show Screen 0’s attributes, and the bottom 4 show Screen 1’s. Each half can specify independent ink and paper colours.

Combined with the 11110000b pixel pattern, this gives:

- 2 colour columns per character cell (left 4 pixels = ink, right 4 = paper)
- 2 colour rows per character cell (top 4 scanlines from Screen 0, bottom 4 from Screen 1)
- Total: 4 independently coloured sub-cells per 8x8 character cell

Over the full screen: 64 columns x 48 rows = **3,072 independently coloured pixels**, each 4x4 real pixels in size. The effective resolution is 64x48 with full per-pixel colour from the Spectrum’s 15-colour palette.

Dual-screen multicolour on the 128K Spectrum — Ringo’s two-screen switching produces a 64x48 grid with independent per-pixel colour

Figure 20: Dual-screen multicolour on the 128K Spectrum — Ringo’s two-screen switching produces a 64x48 grid with independent per-pixel colour

This is a fundamentally different approach from GLUF's 8x2 multicolor. GLUF changes attributes in sync with the beam, requiring precise timing and consuming massive cycle budgets. Ringo uses dual-screen hardware switching, which requires only a single OUT instruction every 4 scanlines. The CPU overhead for the screen switching itself is minimal.

Where the T-states go

The cheap screen switching means more T-states are available for game logic. But rendering into two screens simultaneously is not free. Every tile and sprite update must be written to both Screen 0 and Screen 1, because the player sees a composite of both.

Ringo's sprites are 12x10 "pixels" in the 64x48 grid, which means 12 attribute bytes wide and 10 attribute rows tall (split across the two screens). Each sprite occupies 120 bytes of data. Sprite rendering uses fixed-cycle macros — sequences of instructions with known, constant execution time, critical for maintaining synchronisation with the screen switching.

Tile rendering is more involved. DenisGrachev pre-generates tile rendering code in memory pages, using pop af : or (hl) patterns:

```
; Tile rendering fragment (conceptual)
    pop  af          ; 10 T  load tile data from stack-based source
    or   (hl)        ;  7 T  combine with existing screen data
    ld   (hl),a       ;  7 T  write back
    inc  l            ;  4 T  next attribute column
                    ; ---
                    ; 28 T per attribute byte
```

The pop af is a stack trick: the tile data is arranged as a stack-format table in memory. SP points to the tile data, and POP reads two bytes at a time. The or (hl) combines the tile colour with whatever is already on screen, allowing transparent tiles and layered backgrounds.

Horizontal scrolling

Ringo implements horizontal scrolling with half-character displacement. Since each "pixel" in the 64x48 grid is 4 real pixels wide, scrolling by one "pixel" means shifting the 11110000b pattern by 4 bits. But the pattern is fixed — you cannot easily shift it without corrupting the colour trick.

Instead, DenisGrachev scrolls by moving the attribute data. A one-pixel scroll shifts all attributes one column to the left or right and draws the new column at the edge. Because the attributes are the only thing that changes (the pixel pattern stays fixed at \$F0), the scroll is just a block copy of attribute bytes. For the 64x48 grid, this is 48 bytes per column shift (one byte per row), far cheaper than pixel-level scrolling.

For sub-"pixel" scrolling — smooth movement within a 4-pixel-wide column — DenisGrachev alternates between \$F0 and \$E0 (or similar shifted patterns) in the pixel data. This requires more bookkeeping but achieves half-character displacement, giving the illusion of 128-column horizontal resolution.

Traditional Multicolor: The Interrupt-Driven Approach

Before we move to the practical, it is worth understanding the “classic” approach that GLUF and Ringo displaced. Traditional multicolor is conceptually the simplest: change attributes at precisely the right moment, and the ULA will display different colours on different scanlines.

The technique works like this:

1. Execute HALT to synchronise with the frame interrupt. After HALT, the CPU is at a known T-state position relative to the start of the display.
2. Count T-states from the HALT. The ULA reads each row of 32 attributes at a known point in each scanline. By padding with NOPs or other known-length instructions, you can position your code at exactly the right moment.
3. At the precise T-state when the ULA has finished reading attributes for the current scanline (but before it reads them for the next), overwrite the 32 attribute bytes with new values.
4. Wait for the ULA to read the new values, then overwrite again for the next change.

The timing is brutal. Each scanline takes 224 T-states on Pentagon. The ULA reads 32 attribute bytes early in each scanline, and the CPU must change all 32 bytes in the gap before the next read. With LD (HL),A : INC L at 11 T-states per byte, writing 32 bytes takes 352 T-states — more than one entire scanline. You cannot change every scanline. At best, you can change every other scanline (8x2 resolution) if you use the fastest possible output method (PUSH-based), and even then the timing margins are razor-thin.

sequenceDiagram

```

    participant CPU
    participant ULA as ULA (electron beam)
    participant ATTR as Attribute RAM ($5800)
  
```

```

    Note over ULA: Scanline N begins (224T)
    ULA->>ATTR: Read 32 attribute bytes for row
    Note over ULA: Display scanline N using attrs
  
```

```

    rect rgb(200, 230, 255)
      Note over CPU: ← Window: CPU rewrites attributes
      CPU->>ATTR: Write 32 new attribute bytes
      Note over CPU: (must complete before next read!)
    end
  
```

```

    Note over ULA: Scanline N+2 begins
    ULA->>ATTR: Read 32 attribute bytes (sees NEW values)
    Note over ULA: Display scanline N+2 with new colours
  
```

```

    rect rgb(200, 230, 255)
      Note over CPU: ← Window: CPU rewrites again
      CPU->>ATTR: Write 32 more attribute bytes
    end
  
```

Note over ULA: Scanline N+4 begins
ULA->>ATTR: Read attrs (sees NEWEST values)
Note over ULA: Display with third colour band

The race: At 224 T-states per scanline, writing 32 bytes via LD (HL),A : INC L costs 352T — more than one scanline. This is why 8×2 resolution (change every 2 scanlines) is the practical limit with CPU-based attribute rewriting, and why LDPUSH merges pixel and attribute output to avoid a separate attribute pass.

The practical result: traditional multicolor consumes 80–90% of the CPU on attribute management. In a demo, where the multicolor *is* the effect, this is acceptable. In a game, it is lethal. No T-states remain for game logic, collision detection, or sound.

DenisGrachev's LDPUSH technique solves this by merging the attribute output with the pixel output. The same code that writes pixel data also writes attributes, and both are embedded in executable instructions. There is no separate "attribute management" phase eating the budget. The rendering pass handles everything.

Sidebar: Black Crow #05 — Early Multicolor

The technique of changing attributes between scanlines was documented as early as 2001 in Black Crow #05, a Russian ZX Spectrum scene magazine distributed on TRD disk images. The article presented the basic algorithm — synchronise with the raster, count T-states, change attributes — along with working example code.

Black Crow is significant as a historical marker. By 2001, the ZX Spectrum demoscene had evolved well past the platform's commercial lifespan, and coders were systematically cataloguing tricks that pushed the hardware beyond its design specifications. Multicolor was one of many techniques shared through the scene magazine ecosystem: Spectrum Expert, ZX Format, Born Dead, Black Crow, and later the online platform Hype.

DenisGrachev's contribution, nearly two decades later, was not inventing multicolor but solving its practical problems for game development. The magazine articles documented what was possible. DenisGrachev showed what was *usable*.

The Spectrum's Palette and Multicolor

A brief note on colour mechanics, because multicolor's visual impact depends entirely on the palette.

The ZX Spectrum has 15 colours: 8 base colours (black, blue, red, magenta, green, cyan, yellow, white) and 7 BRIGHT variants (bright black is the same as black, so only 7 additional). Each attribute byte specifies an ink colour (3 bits), a paper colour (3 bits), a BRIGHT flag (1 bit, applies to both ink and paper simultaneously), and a FLASH flag (1 bit).

With 8x2 multicolor, each character cell gets four attribute rows instead of one. Each row specifies independent ink and paper colours. That is up to eight colours per cell (two per row, four rows) — though in practice, the BRIGHT constraint (it applies to both ink and paper) limits the effective combinations.

With Ringo's 64x48 approach, each sub-cell is fully independent. The 15-colour palette is available at every one of the 3,072 positions. The result is closer to what 8-bit home computers with more capable hardware — the MSX2, the Amstrad CPC — could achieve natively. On the Spectrum, it is achieved entirely in software, by exploiting the timing relationship between the CPU and the ULA.

Practical: A Multicolor Game Screen

Let us build a simplified multicolor renderer. The goal: a 24-character-wide game area with 8x2 colour resolution and one moving sprite on top. This will not match GLUF's full feature set, but it will demonstrate the core LDPUSH technique and the two-frame rendering approach.

Step 1: The display code buffer

First, we need a block of memory filled with LDPUSH instruction pairs. For each scanline in the 24-character game area, we need 12 pairs of LD DE,nn : PUSH DE (each pair outputs 2 bytes, 12 pairs output 24 bytes = the full game area width).

```
; Structure of one scanline's display code (conceptual)
; SP is pre-set to the right edge of this scanline in screen memory

    ld    de,$0000      ; 10 T  rightmost 2 bytes (will be patched)
    push de              ; 11 T
    ld    de,$0000      ; 10 T  next 2 bytes leftward
    push de              ; 11 T
    ld    de,$0000      ; 10 T
    push de              ; 11 T
    ; ... 12 pairs total ...
    ld    de,$0000      ; 10 T  leftmost 2 bytes
    push de              ; 11 T
    ; --- 252 T per scanline (12 x 21) ---

    ; Then: adjust SP for the next scanline
    ; Then: change attributes (every 2nd scanline)
```

The total display code for 128 scanlines (16 character rows x 8 scanlines each) at approximately 51 bytes per scanline is about 6,500 bytes.

Step 2: Attribute changes within the display code

Every two scanlines, the display code must include attribute writes. Between the LDPUSH sequences for scanlines N and N+2, insert code that overwrites the 32 attribute bytes for the current character row:

```

; After outputting scanline N...
; Attribute change for the next 2-scanline band

    ld    sp,attr_row_end      ; point SP at end of attribute row
    ld    de,attr_data_0      ; 10 T  rightmost 2 attribute bytes
    push de                    ; 11 T
    ld    de,attr_data_1      ; 10 T
    push de                    ; 11 T
    ; ... 16 pairs for 32 attribute bytes ...

    ld    sp,next_scanline_end ; point SP at next scanline's right edge
    ; Continue with pixel LDPUSH pairs for scanlines N+2, N+3

```

The attribute changes are embedded directly in the display code stream. They execute at exactly the right moment because they are *positioned* at exactly the right point in the instruction sequence. No T-state counting required. No NOP padding. The structure of the code guarantees the timing.

Step 3: Tile rendering via operand patching

To draw a tile into the game area, you must patch the operand bytes of the LD instructions in the display code buffer. A 16x16 pixel tile covers 2 bytes wide and 16 scanlines tall. In the display code, those 2 bytes are the operand of a specific LD DE instruction. To update the tile:

```

; Patch one scanline of a 16x16 tile into the display buffer
; IX points to the LD DE instruction for this position in the buffer
; HL points to the tile's pixel data for this scanline

    ld    a,(hl)               ; 7 T  read tile byte 0
    ld    (ix+1),a             ; 19 T  patch into LD DE operand (low byte)
    inc    hl                  ; 6 T
    ld    a,(hl)               ; 7 T  read tile byte 1
    ld    (ix+2),a             ; 19 T  patch into LD DE operand (high byte)
    inc    hl                  ; 6 T
    ; advance IX to the next scanline's LD DE instruction
    ; (stride depends on display code structure)

```

At 19 T-states per IX-indexed write, this is not cheap. For a full 16x16 tile (16 scanlines x 2 bytes x 2 patches per byte): roughly 1,200 T-states per tile. In a 24x16 character game area with 2-character-wide tiles, there are up to 192 tiles. Even with partial updates (only redrawing tiles that changed), tile rendering dominates the budget.

This is why GLUF splits rendering across two frames. Frame 1 handles the time-critical attribute changes and renders as many tiles as possible. Frame 2 finishes tiles and composites sprites on top.

Step 4: Sprite overlay

Sprites are rendered on top of tiles using the same operand-patching technique, but with an additional step: save the original operand bytes before overwriting them, so the sprite can be erased on the next frame by restoring the saved data.

```

; Sprite rendering: save background, patch sprite data
ld  a,(ix+1)           ; read current (background) byte
ld  (save_buffer),a    ; save for later restoration
ld  a,(sprite_data)    ; load sprite pixel
ld  (ix+1),a           ; patch into display code

```

The save/restore mechanism is the multicolor equivalent of dirty-rectangle sprite rendering. You save what was there, draw the sprite, display it, then restore the saved bytes to erase the sprite before drawing it in its new position.

Step 5: The main loop

```

main_loop:
    halt                ; synchronise with frame

    ; --- Frame 1: attributes + tiles ---
    di
    ld  (restore_sp+1),sp ; save real SP
    call execute_display  ; run the LDPUSH display code
restore_sp:
    ld  sp,$0000         ; restore SP
    ei

    call update_tiles     ; patch changed tiles into buffer B
    call read_input       ; handle player input
    call update_game_logic ; move entities, check collisions

    halt                ; synchronise with next frame

    ; --- Frame 2: remaining tiles + sprites ---
    di
    ld  (restore_sp2+1),sp
    call execute_display  ; display the current frame
restore_sp2:
    ld  sp,$0000
    ei

    call finish_tiles     ; patch any remaining tiles
    call erase_old_sprite  ; restore saved bytes
    call render_sprite    ; patch sprite into new position
    call update_sound     ; sound at 25 Hz (every other frame pair)

    jr  main_loop

```

This skeleton captures the essential rhythm: two frames per logical game frame, display code executing with interrupts disabled, tile and sprite rendering happening between display passes. The interleaving of rendering and game logic within the two-frame structure is the engineering heart of a multicolor game engine.

What DenisGrachev's Work Means

DenisGrachev's achievement was not inventing multicolor — the technique was known. It was solving the engineering problem of fitting multicolor rendering, tile engines, sprite overlay, double buffering, sound, and game logic into the same frame budget. The two-frame architecture, the merged code-as-data display buffer, and the 25 Hz sound compromise are game engine decisions, not demo tricks. Ringo pushed further by trading colour resolution (64x48 vs GLUF's 192x128 pixel grid) for a cheaper rendering path via dual-screen switching.

Summary

- **Attribute clash** (two colours per 8x8 cell) is the Spectrum's defining visual constraint. The ULA reads attributes per scanline, not per frame — if you change them between reads, you get more colours per cell.
 - The **LDPUSH technique** fuses display data with executable code: `LD DE,nn : PUSH DE` sequences write pixel data to screen memory when executed, and the immediate operands serve as the display buffer. Patching the operands changes what gets drawn.
 - **GLUF** achieves 8x2 multicolor in a 24x16 character game area with double-buffered display code, 16x16 sprites and tiles, and a two-frame architecture that splits rendering across 2/50ths of a second.
 - **Ringo** uses the 11110000b pixel pattern with dual-screen switching every 4 scanlines to achieve a 64x48 grid with per-pixel colour — a fundamentally different trade-off that favours colour independence over spatial resolution.
 - **Traditional multicolor** (interrupt-driven attribute changes) is conceptually simpler but consumes 80–90% of the CPU, making it impractical for games.
 - The **Black Crow #05** magazine documented multicolor as early as 2001. DenisGrachev's contribution was making it practical for game development.
 - DenisGrachev's work demonstrates that demoscene techniques are engineering tools, not just demo tricks. The distinction between "possible in a demo" and "usable in a game" is the engineering challenge.
-

Try It Yourself

1. **Build the display buffer.** Write a program that generates a block of `LD DE,nn : PUSH DE` pairs in memory, points SP at screen memory, and executes the block. You should see a pattern on screen corresponding to the immediate values you chose. Change the values and re-execute to see the screen update.
2. **Add attribute changes.** Extend your display buffer to include attribute writes every 2 scanlines. Fill alternate bands with different colours. You should see horizontal colour stripes within a single character row — proof that the multicolor effect is working.
3. **Patch a tile.** Write a routine that takes a 16x16 pixel tile and patches it into the display buffer at a specified position by modifying the LD operand bytes. Draw several tiles to fill the game area.

4. **Move a sprite.** Implement save/restore background patching: before drawing the sprite, save the operand bytes it will overwrite. After displaying the frame, restore the saved bytes. Move the sprite one character per frame and verify it moves cleanly without leaving trails.
5. **Measure the budget.** Use the border-colour timing harness from Chapter 1 to measure how much of the frame your display code consumes. On Pentagon, the red stripe should nearly fill the border — GLUF uses $\sim 70,000$ of the available 71,680 T-states. See how much room remains for game logic.

Sources: DenisGrachev, “Multicolor Will Be Conquered” (Hype, 2019); DenisGrachev, “Ringo Render 64x48” (Hype, 2022); Black Crow #05 (ZXArt, 2001). Game engines: Old Tower, GLUF, Ringo by DenisGrachev.

Chapter 9: Attribute Tunnels and Chaos Zoomers

“This is a VERY BUGGY demo. It is THE hardest thing I’ve ever done in a demo – no joke.” – Introspec, file_id.diz for the party version of Eager (to live), 3BM Open Air 2015

In the summer of 2015, Introspec sat down to build something he had never attempted before. The demo that emerged – Eager (to live), released under the Life on Mars label at 3BM Open Air – won first place in the ZX Spectrum demo compo. It ran for two minutes, looping, and every frame of it rendered at 50Hz with true digital drums blended into the AY chip’s output. The visual centrepiece was a tunnel that appeared to bore into the screen, colours rippling outward in organic waves. When people saw it, many assumed heavy pixel manipulation was involved. In fact, the tunnel never touched a single pixel. The entire effect lived in attribute memory.

This chapter is a making-of. We will work through the two core visual effects from Eager – the attribute tunnel and the chaos zoomer – tracing the creative reasoning alongside the code. Along the way, we will encounter a scripting engine that hints at the synchronisation architecture covered in Chapter 12, and a philosophical argument about precalculation that divided the ZX scene for years. But we begin where Introspec began: staring at 768 bytes and realising they were enough.

The Attribute Grid as a Framebuffer

Every ZX Spectrum coder knows the attribute area at \$5800–\$5AFF. Each of the 768 bytes controls the ink and paper colour for an 8x8 block of pixels, giving a 32x24 grid. In game development, attributes are the source of colour clash headaches. In Chapter 8, we saw how multicolor engines rewrite attributes in sync with the raster beam to fight the 8x8 grid. The attribute tunnel does the opposite: it embraces the grid.

The insight is disarmingly simple. If you fill pixel memory with a fixed pattern – say, alternating ink/paper stripes or a checkerboard – then the attribute byte alone determines what the viewer sees in each 8x8 cell. Change the ink and paper colours, and the visual content of that cell changes completely. You now have 32x24 “pixels” of colour, each one an attribute byte. Writing a full frame means writing

768 bytes. No screen address interleaving. No bit manipulation. No pixel plotting. Just a linear block copy into attribute RAM.

At 32x24, the resolution is terrible by any normal standard. But Introspec was not building a normal effect. He was building a tunnel.

Think about what a tunnel looks like from the viewer's perspective. The "mouth" of the tunnel – the centre of the screen – is where the eye is drawn. The walls recede toward the edges. Near the centre, individual details are small and blurred by depth. Near the edges, the walls are close and you can see texture. This maps beautifully onto a variable-resolution display: coarse resolution in the centre (where the tunnel is far away and detail does not matter) and finer resolution at the edges (where it does).

Introspec pushed this further with pseudo-chunky rendering. In the centre of the screen, several attribute cells share the same colour, creating larger "pixels." Toward the edges, each 8x8 cell gets its own value. The eye accepts the blocky centre because that is where the tunnel mouth is – depth naturally destroys detail there. Peripheral vision picks up the finer resolution at the edges, creating an impression of higher fidelity than the data actually contains.

This is the first lesson of Eager: the attribute grid is not a limitation to work around. It is a framebuffer to work with.

Plasma: The Colour Engine

The tunnel's colours do not come from a stored texture mapped onto a tube. They come from a plasma calculation – the classic sum-of-sines approach that has been a demoscene staple since the Amiga days, here adapted for the Spectrum's attribute palette.

The basic idea: for each position on the 32x24 grid, sum several sine waves with different frequencies and phases. The result, after clamping to the available colour range, determines the attribute byte. Vary the phases over time and the plasma animates, creating that organic, rippling flow.

On the Z80, this means table lookups. A 256-byte sine table, page-aligned so you can index it with a single register, provides the base function. For each cell, you look up $\sin(x * \text{freq1} + \text{phase1}) + \sin(y * \text{freq2} + \text{phase2}) + \dots$ where the multiplications by frequency are really just additions to the index (multiply by 2 = look up every other entry, multiply by 3 = add the index to itself twice). The accumulated value indexes into a colour map that produces an attribute byte.

The tunnel shape is implicit, not explicit. There is no distance-from-centre calculation, no angle table, no polar coordinate transform. Instead, the frequency and phase parameters of the plasma are arranged so that the resulting colour pattern naturally forms concentric rings when viewed on screen. The rings emerge from the interference of the sine waves, just as Moire patterns emerge from overlapping grids. Adjust the parameters and the rings contract toward the centre, creating the illusion of depth – of looking down a tunnel.

graph TD

```

A["For each attribute cell (row, col)"] --> B["Look up sin(col × freq1 +
↪ phase1)<br>from 256-byte sine table"]
B --> C["Look up sin(row × freq2 + phase2)"]
C --> D["Look up sin(col+row + phase3)"]
D --> E["Sum the three sine values"]
E --> F["Index into colour map<br>(sum → attribute byte)"]
F --> G["Write attribute to buffer"]
G --> H{More cells<br>in quadrant?}
H -- Yes --> A
H -- No --> I["4-fold symmetry copy:<br>top-left →
↪ top-right,<br>bottom-left, bottom-right"]
I --> J["Increment phase1, phase2, phase3<br>→ next frame"]
J --> A

style F fill:#f9f,stroke:#333
style I fill:#bfb,stroke:#333

```

Key insight: There is no distance-from-centre calculation, no angle table, no polar coordinate transform. The tunnel shape emerges from sine wave interference — concentric rings appear naturally from overlapping frequencies. Only one quarter (16×12) is computed; the rest is mirrored.

This is cheaper than a true geometric tunnel (which would require per-pixel distance and angle lookups) and produces a visually rich result. The trade-off is less geometric precision, but at 32x24 resolution, geometric precision was never on the table anyway.

Plasma algorithm — true-colour sine sum reference (left) vs ZX Spectrum 15-colour attribute output (right), showing how four-fold symmetry halves computation

Figure 21: Plasma algorithm — true-colour sine sum reference (left) vs ZX Spectrum 15-colour attribute output (right), showing how four-fold symmetry halves computation

Four-Fold Symmetry: Divide and Conquer

Even at 32x24, calculating plasma for all 768 cells every frame is expensive on a 3.5MHz Z80. Introspec cut the workload by a factor of four with a classic optimisation: exploit the tunnel's natural symmetry.

A tunnel viewed head-on is symmetric about both the horizontal and vertical axes. If you calculate one quarter of the screen – the top-left 16x12 block – you can copy it to the other three quarters by mirroring. Top-left to top-right is a horizontal flip. Top-left to bottom-left is a vertical flip. Top-left to bottom-right is both.

The copy routine is tight. In Introspec's implementation, HL points to the source byte in the top-left quarter, and the three destination addresses (top-right, bottom-left, bottom-right) are maintained in a combination of absolute addresses and register pair BC:

```

ld a,(hl)      ; read source byte from top-left quarter
ld (nn),a      ; write to upper-right quarter (mirrored)

```

```

ld (mm),a      ; write to lower-left quarter (mirrored)
ld (bc),a      ; write to lower-right quarter (mirrored)
ldi           ; copy source to its own destination AND advance HL, DE

```

The `ld (nn),a` and `ld (mm),a` instructions use absolute addressing – the target addresses are embedded directly in the code, patched via self-modification or code generation for each cell position. The `ldi` instruction at the end does double duty: it copies the byte from (HL) to (DE) for the top-left quarter’s own position in the attribute buffer, and it auto-increments both HL and DE while decrementing BC. This means the loop counter, the source pointer advance, and one of the four writes are all folded into a single two-byte instruction.

The total cost: 56 T-states per source byte (7 + 13 + 13 + 7 + 16), which works out to under 15 T-states per *output* byte since each source byte is written to four destinations. For 192 source bytes (one quarter of the 768-byte attribute area), that is roughly **10,752 T-states** to fill the entire screen — about 15% of the ~70,000 T-state frame budget. This leaves the majority of the frame for the plasma calculation, the music engine, and the digital drum playback that made Eager distinctive.

The addresses (nn) and (mm) are literal two-byte values baked into `LD (addr),A` instructions, patched via self-modification or code generation for each cell position. This is standard demoscene practice: the Z80’s lack of an instruction cache means self-modifying code executes reliably.

The Chaos Zoomer

The second major visual effect in Eager is the chaos zoomer. Where the tunnel is smooth and organic, the zoomer is jagged and fractal-like – a field of attribute data zooming toward or away from the viewer, with new detail emerging at the edges as the zoom progresses.

The “chaos” comes from the visual result, not the algorithm. The effect zooms into a region of attribute data, magnifying the centre while the edges scroll inward. Because the source data has patterns at multiple scales, the zooming reveals self-similar structure that reads as fractal to the eye.

The implementation relies on unrolled `ld hl,nn : ldi` sequences. Each `ld hl,nn` loads a new source address – the position in the source buffer to sample for this particular output cell. The following `ldi` copies from (HL) to (DE), advancing DE to the next output position. The source addresses are arranged so that cells near the centre of the screen sample from nearby positions in the source data (magnification), while cells near the edges sample from widely spaced positions (compression). Vary the mapping over time and the zoom animates.

```

; Unrolled chaos zoomer fragment
ld hl,src_addr_0    ; source for output cell 0
ldi                 ; copy to output, advance DE
ld hl,src_addr_1    ; source for output cell 1
ldi
ld hl,src_addr_2    ; source for output cell 2
ldi
; ... repeated for all 768 cells (or one quarter, with symmetry)

```

The key optimisation: since `ldi` auto-increments `DE`, you never need to calculate or load the destination address. The output is always written sequentially into attribute RAM. Only the source addresses vary, and they are embedded directly in the instruction stream as immediate operands. This makes the zoomer a long sequence of `ld hl,nn : ldi` pairs – conceptually simple, but each pair is just 5 bytes (3 for `ld hl,nn` + 2 for `ldi`) and 26 T-states. For a full quarter-screen of 192 cells, that is roughly 5,000 T-states of pure copying, plus the four-way symmetry copy on top.

The complication is that the source addresses change every frame as the zoom progresses. Updating 192 two-byte addresses embedded in the code would cost nearly as much as the copy itself. This is where code generation enters the picture.

Code Generation: Processing Writes Z80

Introspec did not write the zoomer’s unrolled code by hand. The address sequences are different for every zoom level, and calculating them at runtime would consume the frame budget. Instead, he wrote the code generator in Processing, the Java-based creative coding environment. A Processing sketch calculated, for each frame and each output cell, which source cell should be sampled, then output a complete `.a80` source file containing the unrolled `ld hl,nn : ldi` sequence with all addresses filled in. `sjsmplus` compiled this generated source alongside the hand-written engine code.

The pipeline: Processing calculates the zoom mapping, writes `.a80` source, the assembler compiles it, and at runtime the scripting engine selects which pre-generated frame to execute. The Z80 does not compute the mapping. It merely plays it back.

This trades memory for speed – the pre-generated code for all zoom frames occupies substantial RAM, hence the 128K requirement – but the runtime cost per frame is minimal.

The Zapilator Question

The ZX scene has a long and occasionally heated relationship with precalculation. The Russian demoscene coined the term *zapilator* – roughly, “precalculator” – for demos that rely heavily on pre-generated data rather than real-time computation. The word carries a faint whiff of disapproval. If the PC does all the interesting work, what is the Spectrum actually doing? Is it a demo or a slideshow?

Introspec’s answer is characteristically nuanced. The art, he argues, is not in the computation itself but in *designing what to precalculate*. Choosing the right zoom mapping, the right interpolation, the right way to decompose the problem so that the pre-generated code fits in memory and the playback runs at 50Hz – this is engineering. The Processing script does not write itself. The Z80 code structure that makes playback efficient does not emerge automatically. The creativity lives in the architecture, not in whether the inner loop contains an `add` or a `ldi`.

He has a point. The chaos zoomer's visual quality depends on the source data, the mapping function, the zoom curve, the colour palette, and the interplay with the music. All of these are artistic decisions. The fact that the address calculations happen at compile time rather than runtime is an implementation detail – one that enables a visual quality impossible with real-time computation at 3.5MHz. The machine's constraints – its memory, its instruction set, its timing – shaped every decision. That the shaping happened partly in Processing and partly in Z80 assembly does not diminish the result.

For this book's purposes, the takeaway is practical: code generation is a legitimate and powerful technique. If your effect requires calculations that exceed the Z80's frame budget, consider moving them to build time. Your assembler's macro language, a Lua script inside sjasmpplus, or an external program in Python or Processing can all serve as code generators. The Z80 gets to do what it does best: copy data at maximum speed.

A Taste of the Scripting Engine

Eager contains more than the tunnel and the zoomer. It runs for two minutes, with multiple visual variations, transitions, and the digital drum playback that gives the demo its rhythmic pulse. Coordinating all of this is a scripting engine – a topic we will explore in depth in Chapter 12 when we discuss music synchronisation. But a brief sketch here sets up that discussion.

Introspec's engine uses two levels of scripting. The **outer script** controls the sequence of effects: play the tunnel for N frames, transition to the zoomer, back to the tunnel with different parameters, and so on. The **inner script** controls variations within a single effect: changing the plasma frequencies, shifting the colour palette, adjusting the zoom speed.

A critical command in the scripting language is what Introspec calls **kWORK**: “generate N frames, then show them independently.” This is the key to Eager's asynchronous frame generation. The engine pre-renders several frames of the current effect into memory buffers. Then, while those frames are being displayed (one per screen refresh), the engine can do other work – like playing a digital drum sample through the AY chip.

This async architecture is what made Eager so difficult to build. When a drum hit occurs, the CPU is consumed by digital sample playback. Frame generation stalls. The visual effect survives on pre-rendered frames until the drum finishes and generation resumes. During rapid drum hits, the generator falls behind; between hits, it catches up. “My brain is not coping with asynchronous coding well,” Introspec wrote in the party version's file `_id.diz`. The honest exhaustion in that note captures the reality of interleaving a timing-critical audio routine with a frame generation pipeline on a machine with one thread and no operating system.

We will return to this architecture in Chapter 12, where we also examine n1k-o's digital drum technique and the double-buffered attribute frames that make the whole system work.

The Making Of: Timeline and Inspiration

Eager was developed between June and August 2015. Introspec has said the initial inspiration came from seeing the twister effect in **Bomb** by Atebit – a visual trick that exploited attribute manipulation to create the illusion of a three-dimensional rotating column. “What else could you do with attributes alone?” was the question that started the project.

The music came from n1k-o (of Skrju), whose track gave the demo its rhythmic structure. The hybrid drum technique – digital sample for the attack transient, AY envelope for the decay – was n1k-o’s innovation, and it drove the entire architectural decision to build an asynchronous frame generation engine. Without the drums, Eager could have been a simpler demo. With them, it became what Introspec called “the hardest thing I’ve done in a demo.”

The development was compressed into roughly ten weeks. The party version, submitted to 3BM Open Air 2015, still had bugs – the file_id.diz carried a note thanking diver of 4D+TBK “for the cool tip” and apologising for the instability. The final version fixed the timing issues across Spectrum models (128K, +2, +2A/B, +3, Pentagon – all at 3.5MHz only, no turbo). That cross-pollination – a scener from one group passing a technical insight to a coder from another – is how the ZX demoscene evolves.

Practical: Building a Simplified Attribute Tunnel

Let us build a stripped-down version of Eager’s attribute tunnel. We will implement:

1. A fixed pixel pattern in bitmap memory (so attributes have something to colour).
2. A plasma calculation over the 32x24 attribute grid.
3. Four-fold symmetry to reduce the calculation to one quarter.
4. Animation by incrementing the plasma phase each frame.

This will not match Eager’s visual sophistication – we are omitting the pseudo-chunky variable-size pixels, the code-generated zoomer, and the async architecture. But it will demonstrate the core principle: the attribute grid as your framebuffer.

Step 1: Fill Pixel Memory with a Pattern

We need a fixed pixel pattern so that ink and paper colours are both visible. A simple checkerboard works:

```
; Fill bitmap memory ($4000-$57FF) with checkerboard pattern
ld hl,$4000
ld de,$4001
ld bc,$17FF          ; 6143 bytes
ld a,$55             ; 01010101 binary -- alternating pixels
ld (hl),a
ldir
```

Every 8x8 cell will now display alternating ink and paper pixels. When we change the attribute, the checkerboard reveals both colours.

Step 2: Sine Table

Page-align a 256-byte sine table for fast indexing. This can be generated at assembly time using sjasmplus's Lua scripting, or pre-calculated and included as binary data:

```

ALIGN 256
sin_table:
    LUA ALLPASS
    for i = 0, 255 do
        -- Sine scaled to 0..63 (6-bit unsigned)
        sj.add_byte(math.floor(math.sin(i * math.pi / 128) * 31 + 32))
    end
ENDLUA

```

Step 3: Plasma for One Quarter

Calculate the plasma value for each cell in the top-left 16x12 quarter. The result is an index into a colour table that produces the attribute byte:

```

; Calculate plasma for top-left quarter (16 columns x 12 rows)
; Input: frame_phase is incremented each frame
; Output: attr_buffer filled with 192 attribute bytes

calc_plasma:
    ld iy,attr_buffer
    ld h,sin_table / 256      ; H = high byte of sine table page
    ld b,12                   ; 12 rows (half of 24)
.row_loop:
    ld c,16                   ; 16 columns (half of 32)
.col_loop:
    ; Plasma = sin(x*2 + phase1) + sin(y*3 + phase2) + sin(x+y + phase3)

    ; Term 1: sin(x*2 + phase1)
    ld a,c
    add a,a                    ; x * 2
    add a,(ix+0)               ; + phase1 (self-modifying or IX-indexed)
    ld l,a
    ld a,(hl)                  ; sin_table[x*2 + phase1]
    ld d,a                     ; accumulate in D

    ; Term 2: sin(y*3 + phase2)
    ld a,b
    add a,a
    add a,b                    ; y * 3
    add a,(ix+1)               ; + phase2
    ld l,a
    ld a,(hl)                  ; sin_table[y*3 + phase2]
    add a,d
    ld d,a

    ; Term 3: sin((x+y) + phase3)
    ld a,c
    add a,b                    ; x + y

```



```

add a,(ix+2)          ; + phase3
ld l,a
ld a,(hl)             ; sin_table[x+y + phase3]
add a,d               ; total plasma value

; Map to attribute byte
rrca
rrca
rrca                  ; shift to get ink bits in position
and %00000111        ; 8 ink colours
or  %00111000         ; white paper (bits 3-5 = 7)
ld (iy+0),a           ; store in buffer
inc iy

dec c
jr nz,.col_loop
djnz .row_loop
ret

```

This is intentionally simplified. A production version would use self-modifying code to inline the phase values (7 T-states per load versus 19 for IX-indexed) and a 256-byte colour lookup table instead of the `rrca` chain.

Step 4: Four-Way Copy

The simplified approach copies row by row, mirroring horizontally for the right half and indexing from the bottom for the lower half. But the production technique is Introspec's single-pass copy we saw earlier – it writes all four quadrants simultaneously from one source byte, using self-modifying `ld (nn),a` instructions with pre-patched addresses. The practical code for that pattern is in the chapter's examples directory as `tunnel_4way.a80`.

Step 5: Main Loop

```

main_loop:
    halt                ; wait for vsync

    call calc_plasma    ; calculate one quarter
    call copy_four_way  ; mirror to full screen

    ; Advance plasma phases
    ld hl,phase1
    inc (hl)
    inc (hl)            ; phase1 += 2
    ld hl,phase2
    inc (hl)            ; phase2 += 1
    ld hl,phase3
    dec (hl)            ; phase3 -= 1 (counter-rotating)

    jr main_loop

```

The different phase increments make the plasma terms rotate at different speeds.

Experiment with these values – even small changes produce dramatically different visual textures.

Attribute-based plasma effect — sine wave interference produces rippling colour patterns across the 32x24 attribute grid

Figure 22: Attribute-based plasma effect — sine wave interference produces rippling colour patterns across the 32x24 attribute grid

Key Insight

The attribute grid IS your framebuffer for this effect. You never touch pixel memory after the initial checkerboard fill. The entire animation consists of writing 768 bytes per frame to \$5800-\$5AFF. The Z80's screen address interleaving, which makes pixel manipulation so painful, is completely irrelevant. The attribute area is linear. The copy is fast. The visual result, at 50Hz, is smooth and surprisingly compelling.

This is the lesson Introspec drew from building Eager, and it applies far beyond tunnels. Any time the ZX Spectrum's colour system frustrates you, consider inverting the problem. Instead of fighting the attribute grid, use it. Those 768 bytes are the cheapest full-screen animation buffer on the machine.

Sources

- Introspec, "Making of Eager," Hype, 2015 (hype.retrospec.org/blog/demo/261.html)
- Introspec, `file_id.diz` from Eager (to live) party version, 3BM Open Air 2015
- Introspec, "Код мёртв" (Code is Dead), Hype, 2015
- Introspec, "За дизайн" (For Design), Hype, 2015

Next: Chapter 10 takes us to Illusion's dotfield scroller and Eager's four-phase colour animation – where two normal frames and two inverted frames create the illusion of a palette the Spectrum does not have.

Chapter 10: The Dotfield Scroller and 4-Phase Colour

“Two normal frames and two inverted frames. The eye sees the average.”
- Introspec, Making of Eager (2015)

The ZX Spectrum displays two colours per 8x8 cell. Text scrolls across a screen at whatever rate the CPU can manage. These are fixed constraints – the hardware does what it does, and no amount of cleverness will change the silicon.

But cleverness can change what the viewer *perceives*.

This chapter brings together two techniques from two different demos, separated by nearly twenty years but connected by a shared principle. The dotfield scroller from X-Trade’s *Illusion* (ENLiGHT’96) renders text as a bouncing cloud of individual dots, each placed at a cost of just 36 T-states. The 4-phase colour animation from Introspec’s *Eager* (3BM Open Air 2015) alternates four carefully constructed frames at 50Hz to trick the eye into seeing colours the hardware cannot produce. One exploits spatial resolution – placing dots wherever you want, unconstrained by character cells. The other exploits temporal resolution – cycling frames faster than the eye can follow. Together they demonstrate the two main axes of cheating on constrained hardware: space and time.

Part 1: The Dotfield Scroller

What the Viewer Sees

Picture a message – “ILLUSION BY X-TRADE” – rendered not in solid block characters but as a field of individual dots, each dot a single pixel. The text drifts horizontally across the screen in a smooth scroll. But the dots are not sitting on flat scanlines. They bounce. The entire dot field undulates in a sine wave, each column offset vertically from its neighbours, creating the impression of text rippling on the surface of water.

The Font as Texture

The font is stored as a bitmap texture in memory – one bit per dot. If the bit is 1, a dot appears on screen. If the bit is 0, nothing happens. The critical word is *transparent*. In a normal renderer, you write every pixel position. In the dotfield

scroller, transparent pixels are nearly free. You check the bit, and if it is zero, you skip. Only the set pixels require a write to video memory.

This means rendering cost is proportional to the number of visible dots, not the total area. A typical 8x8 character might have 20 set pixels out of 64. For a large scrolling message, this economy matters enormously. BC points to the font data; RLA shifts each bit into the carry flag to determine on or off.

Dotfield scroller prototype — each character rendered as individual pixels
bouncing on a sine wave, producing the classic demoscene text effect

Figure 23: Dotfield scroller prototype — each character rendered as individual pixels bouncing on a sine wave, producing the classic demoscene text effect

Stack-Based Address Tables

In a conventional scroller, each pixel's screen position is calculated from (x, y) coordinates using the Spectrum's interleaved address formula. That calculation involves shifts, masks, and lookups. Doing it for thousands of pixels per frame would consume the entire budget.

Dark's solution: pre-calculate every screen address and store them as a table that the stack pointer walks through. POP reads 2 bytes and auto-increments SP, all in 10 T-states. Point SP at the table instead of the real stack, and POP becomes the fastest possible address retrieval – no index registers, no pointer arithmetic, no overhead.

Compare POP to the alternatives. LD A, (HL) : INC HL fetches one byte in 11 T-states – you would need two such pairs (22 T) to fetch an address, plus LD L, A / LD H, A bookkeeping. An indexed load like LD L, (IX+0) : LD H, (IX+1) costs 38 T-states for the pair. POP fetches both bytes, increments the pointer, and loads a register pair – 10 T-states, no contest. The price is that you surrender the stack pointer to the renderer. Nothing else can use SP while the inner loop runs.

This means interrupts are fatal. If an interrupt fires while SP points into the address table, the Z80 pushes the return address onto the “stack” – which is actually your data table. Two bytes of carefully computed screen addresses get overwritten with a return address, and the interrupt service routine proceeds to execute whatever garbage sits at the corrupted location. The result is anything from a garbled frame to a hard crash. The solution is simple and non-negotiable: DI before hijacking SP, EI after restoring it. Every POP-trick routine in every Spectrum demo follows this pattern:

```
di
ld  (.smc_sp+1), sp ; save SP via self-modifying code
ld  sp, table_addr ; point SP at pre-computed data
; ... inner loop using POP ...
.smc_sp:
ld  sp, $0000      ; self-modified: restores original SP
ei
```

The save/restore uses self-modifying code because it is the fastest way to both save and restore SP in one step. EX (SP), HL requires a valid stack. LD (addr), SP exists (opcode ED 73, 20 T-states), but it saves SP to a fixed address – you would then need a separate LD SP, (addr) to restore it later (also 20 T-states), and the restore is

no faster than the self-modifying approach. The SMC technique writes SP's value directly into the operand field of a later LD SP,nnnn instruction: LD (.smc+1),SP costs 20 T-states for the save, and the restore (LD SP,nnnn with the patched operand) costs just 10 T-states. The combined save+restore is 30 T-states versus 40 T-states for the LD (addr),SP / LD SP,(addr) pair – a small saving that also avoids reserving a separate memory location.

One subtle consequence: the DI/EI window blocks the frame interrupt. If the inner loop runs long, HALT at the top of the main loop will still catch the next interrupt – but if the rendering overshoots an entire frame, you lose sync. This is why the budget arithmetic matters. You must know your worst-case timing before committing to the POP trick.

The bouncing motion is encoded entirely in the address table. Each entry is a screen address that already includes the vertical sine offset. The “bounce” does not happen at render time. It happened when the table was built. All three dimensions of the animation – scroll position, bounce wave, character shape – collapse into a single linear sequence of 16-bit addresses, consumed at full speed by POP.

The Inner Loop

Introspec's 2017 analysis of Illusion reveals the inner loop. One byte of font data contains 8 bits – 8 pixels. LD A,(BC) reads the byte once, then RLA shifts one bit at a time through 8 unrolled iterations:

```
; Dotfield scroller inner loop (unrolled for one font byte)
; BC = pointer to font/texture data, SP = pre-built address table

    ld  a,(bc)      ; 7 T  read font byte (once per 8 pixels)
    inc bc          ; 6 T  advance to next font byte

    ; Pixel 7 (MSB)
    pop hl          ; 10 T  get screen address from stack
    rla             ; 4 T  shift texture bit into carry
    jr  nc,.skip7   ; 12/7 T skip if transparent
    set 7,(hl)      ; 15 T  plot the dot
.skip7:
    ; Pixel 6
    pop hl          ; 10 T
    rla             ; 4 T
    jr  nc,.skip6   ; 12/7 T
    set 6,(hl)      ; 15 T
.skip6:
    ; ... pixels 5 through 0 follow the same pattern,
    ; with SET 5 through SET 0 ...
```

The per-pixel cost, excluding the amortised byte-fetch:

Path	Instructions	T-states
Opaque pixel	pop hl + rla + jr nc (not taken) + set ?, (hl)	36

Transparent pixel	<code>pop hl + rla + jr nc (taken)</code>	26
----------------------	---	-----------

Bouncing dotfield text scroller in action – text rendered as individual pixels undulating on a sine wave

Figure 24: Bouncing dotfield text scroller in action – text rendered as individual pixels undulating on a sine wave

The LD A, (BC) and INC BC cost 13 T-states amortised over 8 pixels – about 1.6 T per pixel. Introspec’s “36 T-states per pixel” is the worst-case cost within the unrolled byte, excluding that overhead.

The SET bit position changes for each pixel (7, 6, 5 ... 0), which is why the loop is unrolled 8 times rather than repeated. You cannot parameterise the bit position in SET without IX/IY indexing (far too slow) or self-modifying code (overhead). Unrolling is the clean solution.

The Budget Arithmetic

Let us work the numbers properly. The standard Spectrum 48K frame is 69,888 T-states (the Pentagon clone runs slightly longer at 71,680). Of that, the ULA steals T-states during the active display for memory contention, but the scroller writes to screen memory during the entire frame, not just during the border, so contention is a real factor. In practice, assume about 60,000 usable T-states on a 48K and 65,000 on a Pentagon. Subtract music playback (a typical AY player costs 3,000-5,000 T per frame), screen clearing, and table construction. That leaves roughly 40,000-50,000 T-states for the actual dot rendering.

Consider a display of 8 characters of 8x8 font = 512 font bits per frame (8 chars x 8 bytes x 8 bits). With a typical font fill rate of about 30%, roughly 154 bits are set (opaque) and 358 are clear (transparent). The inner loop cost:

- 154 opaque pixels at 36 T each = 5,544 T
- 358 transparent pixels at 26 T each = 9,308 T
- 64 byte-fetches (LD A, (BC) : INC BC) at 13 T each = 832 T
- Total: approximately 15,684 T-states

That is well within a single frame. You could render 20+ characters before hitting the budget ceiling. The bottleneck is not the inner loop – it is the table construction. Building 512 address entries with sine lookups and screen address calculation costs roughly 100-150 T-states per entry (depending on implementation), adding 50,000-75,000 T to the frame. Illusion solves this by pre-computing the entire table set into memory and cycling through offsets, or by building incrementally: when the scroll advances by one pixel, most table entries shift by one position and only the new column needs full recalculation.

The numbers work because two optimisations compound. Stack-based addressing eliminates all coordinate calculation from the inner loop. Texture-driven transparency eliminates all writes for empty pixels. The table build is expensive, but it runs outside the time-critical DI window and can be spread across the frame.

How the Bounce Is Encoded

The address table is where the art lives. To create the bouncing motion, a sine table offsets each column's vertical position:

The two frequency parameters control the visual character of the wave. `phase_freq` determines the spatial frequency – how many wave cycles fit across the visible dot columns. A value of 4 means each dot column advances 4 positions into the sine table, so $256/4 = 64$ columns span one full wave cycle. A value of 8 doubles the frequency, creating a tighter ripple. `speed_freq` controls how fast the wave propagates over time: higher values make the bounce scroll faster independently of the text scroll.

The sine table itself is a 256-byte array of signed offsets, page-aligned for fast lookup. Page alignment means the high byte of the table address is fixed; only the low byte changes, so the lookup reduces to:

```
ld    hl, sin_table    ; H = page, L = don't care
ld    l, a              ; A = (column * freq + phase) & $FF
ld    a, (hl)           ; 7 T – one memory read, no arithmetic
```

The values in the table are signed: positive offsets push the dot down, negative offsets push it up. The amplitude is baked into the table at generation time. A table with range -24 to +24 gives a bounce of 48 scanlines peak-to-peak. Generating the table is a one-time cost, typically done offline or during initialisation using a lookup or a simple approximation. On the Z80, computing true sine values at runtime is expensive, so demoscene coders either pre-compute tables externally or use quadrant symmetry: calculate one quarter-wave (64 entries), then mirror and negate to fill the remaining three quarters.

Given each dot's (x, y + y_offset), the Spectrum screen address is calculated and stored in the table. The table-building code runs once per frame, outside the inner loop. The inner loop sees only a stream of pre-computed addresses.

Beyond Simple Sine: Lissajous, Helix, and Multi-Wave Patterns

The beauty of the pre-computed table approach is that the inner loop does not care what shape the motion follows. It consumes addresses at a fixed cost regardless of the trajectory that generated them. This makes it trivial to experiment with different movement patterns – all the complexity lives in the table-building code.

A **Lissajous pattern** adds a horizontal sine offset as well as the vertical one. Instead of each column mapping to a fixed x byte on screen, the x position also oscillates:

```
y_offset = sin_table[(column * y_freq + phase_y) & 255]
```

When `x_freq` and `y_freq` are coprime (say 3 and 2), the dot field traces a Lissajous figure – the classic oscilloscope pattern. The text becomes a ribbon weaving through space. Different frequency ratios produce dramatically different shapes: 1:1 gives a circle or ellipse, 1:2 gives a figure-eight, 2:3 gives the trefoil pattern familiar from old analogue test equipment.

A **helix** or spiral effect uses a single phase that advances per column, but varies the amplitude:

```
y_offset = sin_table[(column * freq + phase) & 255] * amplitude / max_amp
```

This creates the illusion of dots receding into depth – the wave flattens at the “far” point of the spiral and expands at the “near” point.

Multi-wave superposition is the simplest technique with the most dramatic payoff. Add two sine terms with different frequencies:

The result is a complex, organic-looking undulation that never quite repeats. Advancing phase1 and phase2 at different speeds produces continuously evolving motion from just two table lookups per column. Three or more harmonics create waves that look almost fluid-dynamic. This is the cheapest possible way to generate complex motion – each additional harmonic costs one table lookup and one addition per column in the table builder, and the inner loop cost remains unchanged.

Part 2: 4-Phase Colour Animation

The Colour Problem

Each 8x8 cell has one ink colour (0-7) and one paper colour (0-7). Within a single frame, you get exactly two colours per cell. But the Spectrum runs at 50 frames per second, and the human eye does not see individual frames at that rate. It sees the average.

The Trick

Introspec’s 4-phase technique cycles through four frames:

1. **Normal A:** ink = C1, paper = C2. Pixel data = pattern A.
2. **Normal B:** ink = C3, paper = C4. Pixel data = pattern B.
3. **Inverted A:** ink = C2, paper = C1. Pixel data = pattern A (same pixels, swapped colours).
4. **Inverted B:** ink = C4, paper = C3. Pixel data = pattern B (same pixels, swapped colours).

At 50Hz, each frame displays for 20 milliseconds. The four-frame cycle completes in 80ms – 12.5 cycles per second, above the flicker-fusion threshold on CRT displays.

The Maths of Perception

Trace a single pixel that is “on” in pattern A and “off” in pattern B:

Frame	Pixel state	Displayed colour
Normal A	on (ink)	C1
Normal B	off (paper)	C4
Inverted A	on (ink)	C2
Inverted B	off (paper)	C3

The eye perceives the average: $(C1 + C2 + C3 + C4) / 4$.

Now check: a pixel “on” in both patterns sees C1, C3, C2, C4. A pixel “off” in both sees C2, C4, C1, C3. All cases produce the same average. The pixel pattern does not affect the perceived hue – only the choice of C1 through C4 does.

Then why have two patterns? Because the *intermediate* transitions matter. A pixel that alternates between bright red and bright green flickers noticeably at 12.5Hz. A pixel alternating between similar shades barely flickers at all. The dithering patterns – checkerboards, halftone grids, ordered matrices – control the *texture* of the flicker. Introspec chose patterns so that transitions between frames produced minimal visible oscillation. This is anti-clashing pixel selection: the careful arrangement of “on” and “off” bits to ensure that no pixel toggles between dramatically different colours in consecutive frames.

Why Inversion Is Essential

Without the inversion step, “on” pixels would always show ink and “off” pixels would always show paper. You would get exactly two visible colours per cell, flickering between two different pairs. The inversion ensures both ink and paper contribute to both pixel states across the cycle, mixing all four colours into the perceived output.

On the Spectrum, inversion is cheap. The attribute byte layout is FBPPPIII – Flash, Bright, 3 bits of paper colour, 3 bits of ink colour. Swapping ink and paper means rotating the lower 6 bits: paper moves to ink position, ink moves to paper position, while Flash and Bright stay put. In code:

```
; Swap ink and paper in attribute byte (A)
; Input:  A = F B P2 P1 P0 I2 I1 I0
; Output: A = F B I2 I1 I0 P2 P1 P0
    ld    b, a
    and   $C0          ; isolate Flash + Bright bits
    ld    c, a          ; save FB-----
    ld    a, b
    and   $38          ; isolate paper (--PPP---)
    rrca
    rrca
    rrca                ; paper now in ink position (-----PPP)
    ld    d, a          ; save ink-from-paper
    ld    a, b
    and   $07          ; isolate ink (-----III)
    rlca
    rlca
    rlca                ; ink now in paper position (--III---)
    or    d             ; combine: --IIIPPP
    or    c             ; combine: FBIIIPPP = swapped attribute
```

The alternative is to pre-compute both normal and inverted attribute buffers at initialisation and simply cycle buffer pointers at runtime. This trades 3,072 bytes of memory for zero per-frame computation – a worthwhile trade on 128K machines with memory to spare.

Practical Cost

Four pre-built attribute buffers, cycled once per frame. The per-frame cost is a block copy of 768 bytes into attribute RAM (\$5800-\$5AFF). Using LDIR, this costs 21 T-states per byte: $768 \times 21 = 16,128$ T-states. Using the stack trick (POP from the source buffer, switch SP, PUSH to attribute RAM, batching through register pairs and shadow registers), a realistic cost is around 11,000-13,000 T-states depending on batch size and loop overhead – a modest 1.2-1.5x speedup over LDIR. The gain is smaller than you might expect because each batch requires two SP switches (save source position, load destination, then swap back), and that overhead largely offsets the raw speed advantage of POP+PUSH over LDIR. For a *fill* (writing the same value to every byte), the PUSH trick is far more effective – load register pairs once, then PUSH repeatedly – but a copy from varying source data cannot avoid the read cost.

The cycle logic itself is trivial. A single variable holds the phase (0-3). Each frame, increment it and AND with 3 to wrap. Index into a 4-entry table of buffer base addresses:

```

ld    a, (phase)
inc   a
and   3
ld    (phase), a
add   a, a           ; phase * 2 (pointer table is 16-bit entries)
ld    hl, buf_ptrs
ld    e, a
ld    d, 0
add   hl, de
ld    a, (hl)
inc   hl
ld    h, (hl)
ld    l, a           ; HL = source buffer address
ld    de, $5800      ; DE = attribute RAM
ld    bc, 768
ldir                     ; copy attributes for this phase

```

Memory: $4 \times 768 = 3,072$ bytes for the buffers. On a 48K machine that is a significant chunk; on 128K you can place buffers in paged banks. The pixel patterns (A and B) are written once at initialisation and never touched again – only the attribute RAM changes each frame.

Text Overlay

In Eager, scrolling text overlays the colour animation. There are several approaches, each with different trade-offs.

The simplest is **cell exclusion**: reserve certain character cells for text, skip them during the colour cycle, and write fixed white-on-black attributes with actual font glyphs. This is easy to implement – just mask those cells out of the LDIR copy – but creates a hard visual boundary between the animated background and the static text region. The text looks pasted on.

A more sophisticated approach is **pattern integration**: the glyph shapes override specific bits in both pixel patterns A and B. Where the font has a set bit, both

patterns get that bit set (or cleared, depending on the desired text colour). This ensures the text pixel shows the same colour in all four phases – it does not flicker because it never transitions between different colour states. The surrounding pixels continue to cycle normally. The result is text that appears to float on the animated background, with colour bleeding up to the edges of each letterform. The cost is that you must regenerate (or patch) the pixel patterns whenever the text scrolls, which adds a few thousand T-states per frame depending on how many cells contain text.

A third option for 128K machines is **layer compositing**: maintain the 4-phase background in one set of memory pages and the text scroller in another, then combine them during the attribute copy. This keeps the two systems independent – the scroller does not need to know about the colour animation and vice versa – at the cost of a slightly more complex copy loop that masks text cells.

Demoscene Lineage

The dotfield scroller did not appear from nowhere. The technique sits in a lineage of ZX Spectrum effects that stretches from the mid-1980s to the present.

The earliest Spectrum scrollers were simple character-cell affairs: LDIR-based horizontal scrolls that shifted an entire line of character cells, one byte at a time. Pixel-smooth scrolling was harder – the Spectrum has no hardware scroll register, so every pixel shift requires rewriting the bitmap data. By the early 1990s, demo coders had developed several approaches: RL/RR-based pixel scrolling (shifting every byte in a screen line), look-up table scrollers (pre-shifted copies of each character), and the double-buffer technique (draw into a back buffer, copy to screen). All of these were limited by the fundamental cost of moving bytes in and out of video RAM.

The dotfield approach breaks from this tradition entirely. Instead of scrolling a contiguous block of pixels, it decomposes the text into individual dots and places each one independently. This was Dark's insight in the mid-1990s: if you give up the idea of a solid font and accept a pointillist rendering, you can use the POP trick to place each dot with minimal overhead. The visual result – text dissolving into a cloud of particles, bouncing on a sine wave – became one of the signature effects of the Russian demoscene.

X-Trade's *Illusion* (ENLiGHT'96) was the demo that made the technique famous in the Spectrum world. The dotfield scroller was its centrepiece effect, running smoothly alongside music and other visual elements. Dark published the algorithmic principles in *Spectrum Expert* issues #01 and #02 (1997-98), where he described the general approach to POP-based rendering and sine-table animation. Two decades later, Introspec's detailed reverse-engineering of the *Illusion* binary (published in *Hype* magazine, 2017) confirmed Dark's claims and provided the exact cycle counts that the community had long speculated about.

The 4-phase colour technique has a different pedigree. Colour-cycling on the Spectrum has been explored since the 1980s – simple two-frame alternation (flash-like effects) was common in games and demos. But the systematic four-phase approach, with its careful inversion step to ensure all four colours contribute equally, was refined by Introspec for *Eager* (3BM Open Air 2015). The party

version's `file_id.diz` explicitly mentions the technique, and Introspec's "Making of Eager" article in *Hype* (2015) describes the design process: choosing colours so that adjacent phases minimise visible flicker, and using dithering patterns that distribute the transitions evenly across the cell.

The broader principle – temporal multiplexing of colour – appears on other platforms too. The Atari 2600 famously alternates frames to create flickering pseudo-sprites. The Game Boy uses a similar trick for pseudo-transparency. On the Spectrum, the technique is particularly effective because the CRT phosphor persistence smooths the transitions more than an LCD would. This is worth noting for modern viewers: 4-phase colour looks substantially better on a real CRT or a good CRT emulator (with phosphor simulation) than on a raw pixel-perfect display.

The Shared Principle: Temporal Cheating

The dotfield scroller uses 50 frames per second for *spatial* flexibility. Each frame is a snapshot of dot positions at one instant; the viewer's brain interpolates between snapshots to perceive smooth motion. The CPU's job is to *place* dots as fast as possible, reading pre-computed addresses from the stack.

The 4-phase colour animation uses 50 frames per second for *colour* flexibility. Each frame displays one of four colour states; the viewer's retina averages them. No single frame contains the perceived result – it exists only in persistence of vision.

Both exploit the same physical reality: the CRT refreshes at 50Hz, and the human visual system cannot resolve individual frames at that rate. The Spectrum's *temporal* resolution is far richer than its spatial or colour resolution. Demoscene coders discovered that temporal resolution is the cheapest axis to exploit.

Both reduce their inner loops to the absolute minimum. The scroller to 36 T-states per dot. The colour animation to a single buffer copy per frame. Both move complexity out of the inner loop into pre-computation. And both produce results that look, to the casual viewer, like the hardware should not be capable of them.

This is what makes the demoscene a temporal art form. A screenshot of a dotfield scroller shows a scatter of pixels. A screenshot of a 4-phase colour animation shows two colours per cell, exactly as the hardware specifies. You have to see them *move* to see them work. The beauty is in the sequence, not the frame.

Practical 1: A Bouncing Dot-Matrix Text Scroller

Build a simplified dotfield scroller: a short text message rendered as a bouncing dot-matrix field using POP-based addressing.

Data structures. A page-aligned 8x8 bitmap font (the ROM font at \$3D00 works). A 256-byte sine table for the bounce offset. A RAM buffer for the address table (up to 4,096 x 2 bytes).

Table construction. Before each frame, iterate through the visible characters. For each bit in each font byte, calculate the screen address incorporating the sine-wave

bounce offset, and store it in the address table. This runs once per frame outside the inner loop.

Rendering. Disable interrupts. Save SP via self-modifying code. Point SP at the address table. Execute the unrolled inner loop: `ld a,(bc) : inc bc`, then 8 repetitions of `pop hl : rla : jr nc,skip : set N,(hl)` with N from 7 down to 0. Restore SP. Enable interrupts.

Main loop. `halt` (sync to 50Hz), clear the screen (PUSH-based clear from Chapter 3), build the address table, render the dotfield, advance scroll position and bounce phase.

Extensions. Partial screen clearing (track the bounding box). Double buffering via shadow screen on 128K. Multiple bounce harmonics. Variable dot density for a sparser, more ethereal look.

Practical 2: A 4-Phase Colour Cycling Animation

Build a 4-phase colour animation producing smooth gradients.

Pixel patterns. Fill bitmap memory with two complementary dither patterns. The simplest: even pixel lines get \$55 (01010101), odd lines get \$AA (10101010). For production quality, use an ordered 4x4 Bayer matrix.

Attribute buffers. Pre-calculate four 768-byte buffers. Buffers 0 and 1 hold normal attributes with two different colour schemes (varying ink/paper across the screen for a diagonal gradient). Buffers 2 and 3 are the inverted versions – ink and paper bits swapped. The swap is a bit rotation: three RRCAs to move ink bits to paper position, three RLCAs the other way, mask and combine.

Main loop. Each frame: `halt`, index into a 4-entry table of buffer pointers using a phase counter (AND 3), `LDIR` 768 bytes to \$5800, increment the phase counter. That is the entire runtime engine – about 16,000 T-states per frame.

Animation. For a moving gradient, regenerate one buffer per frame (the one about to become the oldest in the 4-frame cycle) with an advancing colour offset. This maintains a pipeline: display frame N while generating frame N+4. Alternatively, pre-compute all buffers across 128K banks for zero runtime cost.

Summary

- The **dotfield scroller** renders text as individual dots. The inner loop – `pop hl : rla : jr nc,skip : set ?, (hl)` – costs 36 T-states per opaque pixel, 26 per transparent pixel.
- **Stack-based addressing** encodes the bounce trajectory as pre-built screen addresses. POP retrieves them at 10 T-states each – the fastest random-access read on the Z80.
- **4-phase colour** cycles 4 attribute frames (2 normal + 2 inverted) at 50Hz. Persistence of vision averages the colours, creating the illusion of more than 2 colours per cell.

- The **inversion step** ensures all four colours contribute to every pixel position.
 - Both techniques exploit **temporal resolution** to create effects impossible in any single frame.
 - The scroller uses the stack for spatial flexibility; the colour animation uses frame alternation for colour flexibility – the two main axes of demoscene cheating.
-

Try It Yourself

1. Build the dotfield scroller. Start with a single static character plotted via the POP-based inner loop. Verify the expected timing with the border harness from Chapter 1. Then add the bounce table and watch it undulate.
 2. Experiment with bounce parameters. Change the sine amplitude, spatial frequency, and phase speed. Small changes produce dramatic visual differences.
 3. Build the 4-phase colour animation. Start with uniform colour (all cells the same in each phase). Verify you see a steady colour that is neither the ink nor the paper of any single frame. Then add the diagonal gradient.
 4. Try different dithering patterns. Checkerboard, 2x2 blocks, Bayer matrix, random noise. Which minimise visible flicker? Which produce the smoothest perceived gradients?
 5. Combine both techniques: 4-phase colour background with a monochrome dotfield scroller on top.
-

Sources: Introspec, “Technical Analysis of Illusion by X-Trade” (Hype, 2017); Introspec, “Making of Eager” (Hype, 2015); Dark, “Programming Algorithms” (Spectrum Expert #01, 1997). The inner loop disassembly and cycle counts follow Introspec’s 2017 analysis. The 4-phase colour technique is described in the Eager making-of article and the party version’s file_id.diz.

Chapter 11: Sound Architecture

- AY, TurboSound, and Triple AY

“The AY chip has three voices. That’s not a limitation – it’s a design constraint that shaped an entire musical genre.”

The AY-3-8912 is the voice of the ZX Spectrum 128K. General Instrument designed the AY-3-8910 family in 1978 as a Programmable Sound Generator (PSG) – a single chip that could produce music and sound effects without dedicating the CPU to sound generation. Investronica first put the AY into the Spanish Spectrum 128K, Sinclair adopted the design, and by the time Amstrad inherited it for the +2/+3, the chip was already a proven workhorse: Intellivision, MSX, Atari ST, and dozens of arcade machines all used it or its pin-compatible variants (YM2149 in the Atari ST). The Spectrum 128K specifically uses the AY-3-8912, which has fewer I/O ports than the 8910 but is otherwise identical in sound capabilities.

Fourteen registers. Three square-wave tone channels. One noise generator. One envelope generator. That is all you get. Everything you have ever heard in a Spectrum 128K chiptune – the pumping basslines, the rapid-fire arpeggiated chords, the snappy drums – comes from programming these fourteen registers fifty times per second.

This chapter takes you from bare register writes to a working music engine. By the end, you will understand exactly how a tracker player puts sound into the AY, and you will have the knowledge to write your own.

11.1 The AY-3-8910 Register Map

The AY has 14 registers, addressed R0 through R13. On the ZX Spectrum 128K, you access them through two I/O ports:

- **\$FFFD** – Register select (write the register number here first)
- **\$BFFD** – Data write (then write the value here)

Always select the register first, then write the data. This two-step process is fundamental:

```
; Write value E to AY register A
ay_write:
    ld    bc, $FFFD
```

```

out  (c), a      ; select register
ld   b, $BF      ; $BFFD high byte (C stays $FD)
out  (c), e      ; write data
ret

```

Note the trick: we only change the high byte of BC between the two OUT instructions. The low byte \$FD stays in C throughout. This saves loading a full 16-bit value twice.

Complete Register Table

Reg	Name	Bits	Description
R0	Tone A period, low	8	Lower 8 bits of channel A tone period
R1	Tone A period, high	4	Upper 4 bits of channel A tone period
R2	Tone B period, low	8	Lower 8 bits of channel B tone period
R3	Tone B period, high	4	Upper 4 bits of channel B tone period
R4	Tone C period, low	8	Lower 8 bits of channel C tone period
R5	Tone C period, high	4	Upper 4 bits of channel C tone period
R6	Noise period	5	Noise generator period (0-31)
R7	Mixer / I/O enable	8	Tone and noise enable per channel + I/O
R8	Volume A	5	Channel A volume (0-15) or envelope mode
R9	Volume B	5	Channel B volume (0-15) or envelope mode
R10	Volume C	5	Channel C volume (0-15) or envelope mode
R11	Envelope period, low	8	Lower 8 bits of envelope period
R12	Envelope period, high	8	Upper 8 bits of envelope period
R13	Envelope shape	4	Envelope waveform shape (0-15)

Tone Channels (R0-R5): How Pitch Works

Each of the three tone channels produces a square wave. The frequency is controlled by a 12-bit period value split across two registers:

On the Spectrum 128K, the AY clock is 1.7734 MHz (half the CPU clock of 3.5469 MHz). So for middle C (approximately 262 Hz):

AY-3-8910 Register Map




Figure 25: AY-3-8910 register map

The 12-bit period gives a range from 1 (110,837 Hz - inaudible ultrasonic) to 4095 (27 Hz - a deep bass rumble). Here is a practical note table for octave 4:

Note	Frequency (Hz)	Period (decimal)	R_LO	R_HI
C4	261.6	424	\$A8	\$01
C#4	277.2	400	\$90	\$01
D4	293.7	378	\$7A	\$01
D#4	311.1	357	\$65	\$01
E4	329.6	337	\$51	\$01
F4	349.2	318	\$3E	\$01
F#4	370.0	300	\$2C	\$01
G4	392.0	283	\$1B	\$01
G#4	415.3	267	\$0B	\$01
A4	440.0	252	\$FC	\$00
A#4	466.2	238	\$EE	\$00
B4	493.9	225	\$E1	\$00
C5	523.3	212	\$D4	\$00

To go up one octave, halve the period. To go down, double it. The full note table for all useful octaves lives in Appendix G.

The Noise Generator (R6)

Register R6 controls a single noise generator shared by all three channels. The 5-bit value (0-31) sets the noise “pitch” - lower values produce higher, hissier noise; higher values produce lower, rougher noise. The noise generator produces pseudo-random output using a 17-bit linear feedback shift register.

Useful ranges: - **0-5**: High hiss (hi-hat, cymbal) - **6-12**: Mid noise (snare body) - **13-20**: Low rumble (explosion) - **21-31**: Very low (wind, thunder)

R7: The Mixer - The Most Important Register

Register R7 is the heart of the AY. Six bits control which channels receive tone, noise, or both. Two more bits control the I/O port direction (irrelevant for sound, leave them as inputs = 1).

Bit 7: I/O port B direction (1 = input)
 Bit 6: I/O port A direction (1 = input)
 Bit 5: Noise C enable (0 = ON, 1 = off)
 Bit 4: Noise B enable (0 = ON, 1 = off)
 Bit 3: Noise A enable (0 = ON, 1 = off)
 Bit 2: Tone C enable (0 = ON, 1 = off)
 Bit 1: Tone B enable (0 = ON, 1 = off)
 Bit 0: Tone A enable (0 = ON, 1 = off)

The confusing part: 0 means ON. This trips up everyone the first time. The AY uses active-low logic for the mixer enables. A cleared bit enables the corresponding source.

Here is a table of useful mixer combinations:

Value	Binary (NcNbNa TcTbTa)	Effect
\$38	111 000	All three tones on, no noise
\$3E	111 110	Tone A only
\$3D	111 101	Tone B only
\$3B	111 011	Tone C only
\$36	110 110	Tone A + Noise A
\$07	000 111	All three noise, no tones
\$30	110 000	Tones A+B+C, Noise A
\$00	000 000	Everything on (cacophony)
\$3F	111 111	Everything off (silence)

A common pattern for music: tones A+B for melody and harmony, tone C + noise C for drums:

```
; Mixer: tone A on, tone B on, tone C on, noise C on
; Binary: 00 011 000 = noise C on (bit5=0) + all tones on (bits0-2=0)
;         noise A,B off (bits3,4=1), I/O bits=0
; = $18
ld  a, R_MIXER
ld  e, $18
call ay_write
```

Volume Registers (R8-R10)

Each channel has a 4-bit volume control (0-15, where 15 is loudest). But bit 4 is special: setting it switches that channel to **envelope mode**, where the volume is controlled automatically by the envelope generator instead of the fixed value.

Bit 4: 1 = use envelope generator, 0 = use bits 3-0
 Bits 3-0: Fixed volume level (0-15)

The volume curve is logarithmic on a real AY-3-8910 (each step is approximately 1.5 dB) but varies between chip revisions and clones. The YM2149 has a different volume curve, which is why the same tune sounds subtly different on an Atari ST versus a Spectrum.

Envelope Generator (R11-R13)

The AY has one envelope generator - shared across all channels that use it. It automatically modulates the volume according to a repeating waveform.

R11-R12: Envelope Period - A 16-bit value controlling the envelope speed:

At period = 1, the envelope cycles at about 6,927 Hz. At period = 65535, it cycles at roughly 0.11 Hz - about once every 9 seconds.

R13: Envelope Shape - Four bits select the waveform. Although 4 bits allow 16 values, only 10 shapes are unique:


AY-3-8910 Envelope Shapes (R13)

Figure 26: AY envelope shape waveforms

```

; Arpeggio: cycle C-E-G on channel A every frame
; Called once per frame from the interrupt handler
arpeggio:
    ld    a, (arp_pos)
    inc  a
    cp   3
    jr   nz, .no_wrap
    xor  a
.no_wrap:
    ld   (arp_pos), a

    ; Index into note table
    ld   hl, arp_notes
    add  a, a                ; x2 (each entry is a 16-bit period)
    ld   e, a
    ld   d, 0
    add  hl, de
    ld   e, (hl)
    inc  hl
    ld   d, (hl)

    ; Write period to channel A
    ld   a, R_TONE_A_LO
    call ay_write_de_lo
    ld   a, R_TONE_A_HI
    call ay_write_de_hi
    ret

arp_pos:    DB 0
arp_notes:  DW 424, 337, 283    ; C4, E4, G4

```

In a tracker, arpeggios are notated as effects applied to notes. Vortex Tracker II uses ornament tables that specify per-tick semitone offsets.

Buzz-Bass: The Envelope Trick

Here is the most distinctive chiptune bass sound: the “buzz.” It works by abusing the envelope generator. Set the envelope to a short repeating sawtooth (\$08 or \$0C), set the envelope period to match the desired bass note frequency, and restart the envelope every frame. The result is a buzzing, thick bass tone that sounds nothing like a simple square wave.

```
; Buzz-bass: play bass note using envelope generator
; DE = envelope period for desired note
buzz_bass:
    ; Set envelope period
    ld    a, R_ENV_LO
    call ay_write_de_lo
    ld    a, R_ENV_HI
    call ay_write_de_hi

    ; Set envelope shape to repeating sawtooth down
    ; Writing R13 restarts the envelope
    ld    a, R_ENV_SHAPE
    ld    e, $08          ; \\\\ repeating sawtooth down
    call ay_write

    ; Set channel volume to envelope mode (bit 4 = 1)
    ld    a, R_VOL_C
    ld    e, $10          ; bit 4 set = use envelope
    call ay_write
    ret
```

The envelope period for a given bass note is:

For a bass C2 (65.4 Hz): $\text{period} = 1,773,400 / (256 \times 65.4) = 106$.

The buzz-bass gives you a bass instrument that sounds fundamentally different from the tone channels, effectively adding a fourth voice to your arrangement.

The Period Alignment Problem

There is a catch with buzz-bass that equal-tempered note tables hide from you. Look at the formula again:

$$\text{envelope_period} = \text{tone_period} / 16$$

For the buzz to sound clean, the envelope period must be *exactly* $\text{tone_period} / 16$. But integer division truncates. If the tone period is not divisible by 16, the envelope period has a rounding error – and the envelope waveform drifts against the tone, producing audible beating.

Check our standard table. Octave 4:

Note	Period	Period mod 16	Envelope = Period / 16	Error?
C4	424	8	26 (should be 26.5)	Yes
D4	378	10	23 (should be 23.625)	Yes
E4	337	1	21 (should be 21.0625)	Yes
F4	318	14	19 (should be 19.875)	Yes
G4	283	11	17 (should be 17.6875)	Yes
A4	252	12	15 (should be 15.75)	Yes
B4	225	1	14 (should be 14.0625)	Yes

Not a single clean division! Every note in the equal-tempered scale produces a slightly detuned envelope. For short percussive buzz sounds the beating is masked, but for sustained bass notes it creates an unpleasant warble.




Figure 27: Tone + Envelope Phase Alignment: clean T+E with period divisible by 16 (top) vs beating T+E with rounding error (bottom)

Natural Tuning: Table #5

In June 2001, Ivan Roshin published “Частотная таблица с нулевой погрешностью” (A Frequency Table with Zero Error), arriving at the same conclusion that centuries of music theory had already established: replace equal temperament with *just intonation* – integer-ratio intervals that the AY hardware can divide cleanly.

The natural scale for C major / A minor uses these intervals:

This gives pure fifths (3:2 ratio) for C-G, E-B, A-E. Chromatic notes (sharps/flats) are calculated with the 16/15 ratio.




Figure 28: Just Intonation: interval structure of the natural scale with period divisibility table for buzz-bass

The resulting periods, computed for a *non-standard* AY clock of 1,520,640 Hz:

```
; Table #5: Natural tuning for AY clock = 1,520,640 Hz
; 96 notes (8 octaves), C major / A minor
; Ivan Roshin (concept, 2001), oisee/siril (VTi implementation, 2009)
natural_note_table:
    ; Octave 1: every period divisible by 16!
    DW 2880, 2700, 2560, 2400, 2304, 2160
    DW 2025, 1920, 1800, 1728, 1620, 1536
    ; Octave 2
    DW 1440, 1350, 1280, 1200, 1152, 1080
    DW 1013, 960, 900, 864, 810, 768
    ; Octave 3
    DW 720, 675, 640, 600, 576, 540
    DW 506, 480, 450, 432, 405, 384
    ; Octave 4
    DW 360, 338, 320, 300, 288, 270
    DW 253, 240, 225, 216, 203, 192
    ; Octave 5
    DW 180, 169, 160, 150, 144, 135
    DW 127, 120, 113, 108, 101, 96
    ; Octave 6
    DW 90, 84, 80, 75, 72, 68
    DW 63, 60, 56, 54, 51, 48
    ; Octave 7
    DW 45, 42, 40, 38, 36, 34
    DW 32, 30, 28, 27, 25, 24
```



```

; Octave 8
DW 23, 21, 20, 19, 18, 17
DW 16, 15, 14, 14, 13, 12

```

The key insight is that most main-scale periods are now divisible by 16. Here is octave 2 – the bass range that matters most for buzz:

Note	Period	mod 16	Envelope = Period/16	Clean?
C2	1440	0	90	Yes
C#2	1350	6	84.375 → 84	No
D2	1280	0	80	Yes
D#2	1200	0	75	Yes
E2	1152	0	72	Yes
F2	1080	8	67.5 → 68	~
F#2	1013	5	63.3 → 63	No
G2	960	0	60	Yes
G#2	900	4	56.25 → 56	No
A2	864	0	54	Yes
A#2	810	10	50.6 → 51	No
B2	768	0	48	Yes

Seven of twelve notes divide cleanly – all the natural notes of C major. Compare to the equal-tempered table where *none* do. On those seven notes the envelope and tone generators lock in phase, and the buzz-bass sounds pure.

The tradeoff: this table is only correct for C major / A minor. To play in other keys, you change the AY clock frequency:

Key	Chip Frequency (Hz)
C/Am	1,520,640
C#/A#m	1,611,062
D/Bm	1,706,861
D#/Cm	1,808,356
E/C#m	1,915,886
F/Dm	2,029,811
F#/D#m	2,150,510
G/Em	2,278,386
G#/Fm	2,413,866
A/F#m	2,557,401
A#/Gm	2,709,472
B/G#m	2,870,586

On real hardware the AY clock is fixed, so you cannot actually change keys at runtime. But in an emulator or tracker like Vortex Tracker II, the “chip frequency” is a setting. This is exactly what the Vortex Tracker Improved (VTi) modification by oisee did in 2009: it added Table #5 (the fifth table, index 4 counting from zero) with these natural periods, plus a per-module chip frequency setting that selects the key.

The autosiril MIDI-to-PT3 converter defaults to Table #5 precisely because of these clean envelope ratios – most converted tracks use buzz-bass extensively, and the natural tuning eliminates beating.

In practice: if you are writing a tracker module that relies heavily on buzz-bass, consider composing in C/Am with Table #5. The envelopes will lock perfectly to the tone. If you need a different key, either transpose the chip frequency (tracker-side) or accept the small rounding errors of equal temperament. For short percussive buzz sounds, the difference is inaudible; for sustained bass drones, it is very noticeable.

Drum Synthesis

With only one noise generator, drums need to share. The standard approach:

Snare drum: Noise channel + rapid envelope decay.

```
; Snare: noise burst with fast decay
drum_snare:
    ; Noise period: mid-range
    ld  a, R_NOISE
    ld  e, 8
    call ay_write

    ; Mixer: enable noise on channel C
    ld  a, R_MIXER
    ld  e, $18          ; tones A+B+C on, noise C on
    call ay_write

    ; Short envelope: fast decay
    ld  a, R_ENV_LO
    ld  e, 200
    call ay_write
    ld  a, R_ENV_HI
    ld  e, 0
    call ay_write

    ; Envelope shape: single decay
    ld  a, R_ENV_SHAPE
    ld  e, $00          ; \___ single decay to silence
    call ay_write

    ; Channel C to envelope mode
    ld  a, R_VOL_C
    ld  e, $10
    call ay_write
    ret
```

Kick drum: Rapid tone sweep downward. Set channel C to a low tone period, then increase the period over several frames:

```
; Kick: tone sweep down over 4 frames
; Call once per frame while kick_counter > 0
kick_update:
```

```

ld    a, (kick_counter)
or     a
ret    z

dec    a
ld     (kick_counter), a

; Sweep tone down (increase period)
ld     hl, (kick_period)
ld     de, 40          ; sweep speed
add    hl, de
ld     (kick_period), hl

; Write to channel C
ld     a, R_TONE_C_LO
ld     e, l
call   ay_write
ld     a, R_TONE_C_HI
ld     e, h
call   ay_write
ret

```

kick_counter: DB 0

kick_period: DW 0

Hi-hat: Very short noise burst, high noise frequency (low R6 value), immediate volume cut after 1-2 frames.

Ornaments: Per-Frame Modulation

An ornament is a table of per-frame offsets applied to the pitch or volume of a note. Ornaments give life to otherwise static square waves: vibrato, pitch slides, tremolo, attack/decay envelopes – all accomplished by table lookup.

```

; Example ornament: pitch vibrato
; Table of signed semitone offsets, applied once per frame
ornament_vibrato:
    DB 0, 0, 1, 1, 0, 0, -1, -1    ; 8-frame cycle
    DB $80                        ; end marker

```

The player engine applies the ornament offset to the base note's period value each frame, producing smooth modulation.

11.3 TurboSound: 2 x AY

Pentagon and Scorpion clones introduced TurboSound – two AY chips in one machine. The second chip is addressed by selecting it through a bit pattern written to port \$FFFD before register operations.

Chip Selection

On the NedoPC TurboSound card (the most common modern design), chip selection works through port \$FFFD:

```
; Select chip 0 (primary)
ld  bc, $FFFD
ld  a, $FF          ; bit pattern: select chip 0
out (c), a

; Select chip 1 (secondary)
ld  bc, $FFFD
ld  a, $FE          ; bit pattern: select chip 1
out (c), a
```

Once a chip is selected, all subsequent register reads and writes via *FFFD*/*BFFD* go to that chip. In practice, your music engine selects chip 0, updates all its registers, then selects chip 1 and updates its registers.

6 Channels, True Stereo

TurboSound doubles everything: 6 tone channels, 2 independent noise generators, 2 independent envelope generators. A typical stereo arrangement:

Chip	Channels	Stereo	Role
Chip 0	A0, B0, C0	Left or center	Lead melody, harmony, bass
Chip 1	A1, B1, C1	Right or center	Counter-melody, pads, drums

Or mix them for a wide stereo field: - Bass on both chips (center) - Lead on chip 0 (left) - Counter-melody on chip 1 (right) - Drums split: kick on chip 0, snare/hi-hat on chip 1

What TurboSound changes musically is significant: on a single AY, the composer is constantly making sacrifices. You cannot have a sustained bass note, a lead melody, and a drum hit at the same time without channel-stealing. With TurboSound, you have room. Dedicated bass channel, dedicated drums, and four remaining voices for melody and harmony. The compromise era ends.

Engine Modification

Adapting a single-AY engine to TurboSound is straightforward. Your interrupt handler becomes:

```
music_frame:
    ; Update chip 0
    ld  a, $FF
    ld  bc, $FFFD
    out (c), a          ; select chip 0
    call update_chip0    ; write 14 registers

    ; Update chip 1
    ld  a, $FE
```

```

ld    bc, $FFFD
out   (c), a          ; select chip 1
call  update_chip1    ; write 14 registers
ret

```

The register write loop writes all 14 registers from a buffer in RAM. For two chips, you maintain two 14-byte buffers and blast them out sequentially. Each register write requires selecting the register (LD A,reg + OUT) then writing the value (LD A,val + OUT), costing about 36 T-states per register. For 28 registers total: approximately 1,008 T-states, plus chip selection overhead.

11.4 Triple AY on ZX Spectrum Next

The ZX Spectrum Next takes it further: three AY-compatible sound chips, giving you 9 channels. But the Next's implementation goes beyond simple triplication.

Enhanced Features

The Next's AY chips include **per-channel stereo panning**. Each channel can be individually panned left, right, or center – something the original AY never supported. This is controlled through additional Next-specific registers.

The three chips are addressed via Next register \$06 (peripheral 2 setting) or through the standard \$FFFD port with chip select values.

9 Channels: Orchestral Thinking

Nine channels fundamentally change how you approach composition on 8-bit hardware. Instead of clever tricks to simulate complexity, you can think orchestrally:

Channel	Assignment	Panning
Chip 0, A	Bass line	Center
Chip 0, B	Rhythm guitar / pad	Left
Chip 0, C	Lead melody	Center
Chip 1, A	Harmony / counter-melody	Right
Chip 1, B	Arpeggio chord	Left
Chip 1, C	Percussion: kick + snare	Center
Chip 2, A	Hi-hat / cymbal	Right
Chip 2, B	Sound effects (reserved)	Center
Chip 2, C	Ambient / pad layers	Left

This is enough for genuinely rich arrangements. You have a dedicated SFX channel that never interrupts the music. You have independent noise generators for layered percussion. You can sustain chords without arpeggios. The AY's character remains – square waves are still square waves – but the compositional freedom approaches that of an Amiga MOD tracker with its four sample channels.

11.5 Music Engine Architecture

A music engine is the code that reads pattern data and writes AY registers at the correct tempo. On the Spectrum, it lives inside the interrupt handler.

The Interrupt-Driven Player

The ZX Spectrum's IM2 (Interrupt Mode 2) fires once per frame – every 1/50th of a second on PAL systems. The music engine hooks into this:

```
; Setup: install IM2 handler
setup_im2:
    di
    ld    a, $C0          ; interrupt vector table at $C000
    ld    i, a
    im    2

    ; Fill vector table at $C000 with $C1C1
    ; Handler at $C1C1
    ld    hl, $C000
    ld    de, $C001
    ld    bc, 256
    ld    (hl), $C1
    ldir

    ; Place JP at $C1C1
    ld    a, $C3          ; JP opcode
    ld    ($C1C1), a
    ld    hl, isr_handler
    ld    ($C1C2), hl

    ei
    ret

isr_handler:
    push af
    push bc
    push de
    push hl
    ; ... push all registers you use ...

    call music_play      ; <-- the player routine

    ; ... pop all registers ...
    pop    hl
    pop    de
    pop    bc
    pop    af
    ei
    reti
```

The Player Loop

The player routine, called 50 times per second, does the following:

1. Decrement the current note's duration counter
2. If zero, advance to the next note in the pattern
3. Apply ornaments and effects (arpeggio, vibrato, slide) to each channel
4. Calculate final period and volume for each channel
5. Write all 14 AY registers

A simplified skeleton:

```
music_play:
    ; Decrement speed counter
    ld  a, (speed_counter)
    dec a
    ld  (speed_counter), a
    ret nz                ; not time for a new row yet

    ; Reset speed counter
    ld  a, (song_speed)
    ld  (speed_counter), a

    ; Process each channel
    ld  ix, channel_a_data
    call process_channel
    ld  ix, channel_b_data
    call process_channel
    ld  ix, channel_c_data
    call process_channel

    ; Write all registers to AY
    call ay_flush_registers
    ret
```

Frame Budget

Here is the critical question: how many T-states can the music engine consume before it starves the main program?

The frame is 71,680 T-states on Pentagon (69,888 on 48K). The interrupt fires at the start of the frame. If the music player takes 5,000 T-states, the main program has 66,680 left for visuals.

Typical costs: - **Simple player** (no effects, no ornaments): ~1,500-2,500 T-states
 - **Pro Tracker 3 player**: ~3,000-5,000 T-states - **Full Vortex Tracker II player**
 with ornaments and effects: ~4,000-7,000 T-states - **TurboSound player** (2 chips):
 ~6,000-10,000 T-states

For a demo running a cycle-hungry effect, 7,000 T-states for music is significant - about 10% of the frame. Plan accordingly.

Formats and Trackers

The modern standard for AY music on the Spectrum is **Vortex Tracker II** (.pt3 format). It is a cross-platform tracker that runs on Windows and outputs files directly playable by proven Z80 player routines.

Format	Tracker	Features	Player Size
.pt3	Vortex Tracker II / Pro Tracker 3	Ornaments, samples, effects	~1.2-1.8 KB
.asc	ASC Sound Master	Simpler, smaller player	~0.8-1.0 KB
.sqt	SQ-Tracker	Compact, good compression	~0.6-0.8 KB
.stc	Sound Tracker	Basic, oldest	~0.5-0.7 KB

The Pipeline:

1. Compose in Vortex Tracker II on your PC
2. Export as .pt3 file
3. Include the .pt3 player source (Z80 assembly) in your project
4. Include the .pt3 data file as a binary blob
5. Call `music_init` at startup (passing the address of the song data)
6. Call `music_play` from your interrupt handler every frame

; In your main code:

```
ld    hl, song_data
call music_init
```

; In your interrupt handler:

```
call music_play
```

; At the end of the binary:

```
song_data:
    INCBIN "mysong.pt3"
```

This is the standard approach used by virtually every Spectrum 128K demo and game since the early 2000s.

11.6 Sound Effects System

Games need sound effects, and sound effects need channels. The standard approach is **priority-based channel stealing**: when a sound effect triggers, it temporarily takes over a channel from the music engine, then releases it when the effect finishes.

Channel Stealing

; Trigger a sound effect on the SFX channel

; HL = pointer to SFX data table

```
sfx_trigger:
```



```

ld    (sfx_pointer), hl
ld    a, 1
ld    (sfx_active), a
ld    a, 0
ld    (sfx_frame), a
ret

```

; Called every frame from the interrupt handler, AFTER music_play
sfx_update:

```

ld    a, (sfx_active)
or    a
ret    z                ; no active SFX

```

```

; Read current SFX frame data
ld    hl, (sfx_pointer)
ld    a, (sfx_frame)
ld    e, a
ld    d, 0

```

```

; Each SFX frame: [tone_lo, tone_hi, noise, volume, mixer_mask]
; 5 bytes per frame

```

```

push hl
ld    b, 5
call multiply_de_b    ; DE = frame * 5 (or use repeated add)
pop   hl
add   hl, de

```

```

ld    a, (hl)
cp    $FF                ; end marker?
jr    z, .sfx_done

```

```

; Override channel C with SFX data
ld    e, (hl) : inc hl
ld    a, R_TONE_C_LO
call ay_write

```

```

ld    e, (hl) : inc hl
ld    a, R_TONE_C_HI
call ay_write

```

```

ld    e, (hl) : inc hl
ld    a, R_NOISE
call ay_write

```

```

ld    e, (hl) : inc hl
ld    a, R_VOL_C
call ay_write

```

```

; Advance frame counter
ld    a, (sfx_frame)
inc   a

```

```

    ld    (sfx_frame), a
    ret

.sfx_done:
    xor   a
    ld    (sfx_active), a ; deactivate SFX
    ret

```

Procedural SFX Tables

Sound effects are defined as tables of per-frame register values. Here are four classic game sounds:

Explosion: Noise with decaying volume.

```

sfx_explosion:
    ; tone_lo, tone_hi, noise_period, volume, (unused)
    DB 0, 0, 15, 15, 0    ; frame 0: loud low noise
    DB 0, 0, 18, 13, 0    ; frame 1
    DB 0, 0, 20, 11, 0    ; frame 2
    DB 0, 0, 22, 9, 0     ; frame 3
    DB 0, 0, 25, 7, 0     ; frame 4
    DB 0, 0, 28, 5, 0     ; frame 5
    DB 0, 0, 30, 3, 0     ; frame 6
    DB 0, 0, 31, 1, 0     ; frame 7
    DB $FF                ; end

```

Laser: Fast tone sweep downward.

```

sfx_laser:
    DB 10, 0, 0, 14, 0    ; frame 0: high tone
    DB 30, 0, 0, 13, 0    ; frame 1: sweeping down
    DB 60, 0, 0, 12, 0    ; frame 2
    DB 100, 0, 0, 10, 0   ; frame 3
    DB 160, 0, 0, 8, 0    ; frame 4
    DB 240, 0, 0, 5, 0    ; frame 5
    DB 200, 1, 0, 3, 0    ; frame 6: into low range
    DB $FF                ; end

```

Jump: Short tone sweep upward.

```

sfx_jump:
    DB 200, 0, 0, 12, 0   ; frame 0: mid tone
    DB 150, 0, 0, 11, 0   ; frame 1: rising
    DB 100, 0, 0, 10, 0   ; frame 2
    DB 60, 0, 0, 8, 0     ; frame 3: high
    DB 40, 0, 0, 5, 0     ; frame 4
    DB $FF                ; end

```

Pickup: Rapid arpeggio upward.

```

sfx_pickup:
    DB $FC, 0, 0, 14, 0   ; frame 0: A4
    DB $D4, 0, 0, 13, 0   ; frame 1: C5
    DB $A0, 0, 0, 12, 0   ; frame 2: E5 (approx)

```

```

DB $6A,0, 0, 11, 0    ; frame 3: C6 (approx)
DB $6A,0, 0, 8,  0    ; frame 4: sustain
DB $6A,0, 0, 4,  0    ; frame 5: fade
DB $FF                ; end

```

11.7 Putting It Together: The Working Example

The file `chapters/ch11-sound/examples/ay_test.a80` contains a complete, assembling example that demonstrates the fundamentals: initializing the AY, setting the mixer, writing tone periods, and playing a melody. Study it alongside this chapter – every concept discussed here is exercised in that code.

The key patterns to notice in the example:

1. **Two-step port access:** Select register via `$FFFD`, then write data via `$BFFD`.
2. **Mixer configuration:** `$3E` enables only tone A (binary 111 110 – remember, 0 = ON).
3. **Note table:** Pre-calculated period values for each pitch.
4. **HALT-based timing:** Each HALT waits for one interrupt, giving 50 Hz timing resolution.

To extend this example into a real music player, you would replace the linear melody table with pattern-based data, add ornament processing, and move the playback into an IM2 interrupt handler so the main loop is free for visuals.

Sidebar: Beeper - A Brief History of Impossibility

Before the 128K and its AY chip, the original 48K Spectrum had exactly one bit of audio output. Pin 4 of port `$FE`. High or low. That is it.

One bit means one square wave at whatever frequency you toggle it. No volume control, no mixing, no hardware help. To play a note, you sit in a tight loop toggling the bit at the right frequency. To play *two* notes, you interleave two toggle loops. To play three, you interleave three. Each additional voice eats CPU time that could be doing something else – like, say, drawing graphics.

And yet.

Between 2010 and 2015, Shiru (Shiru Otaku) cataloged approximately 30 distinct beeper engines, each using a different technique to extract polyphony from a single bit. The approaches ranged from simple pin-compatible pulse interleaving (2-3 channels) to extraordinary feats of engineering:

- **ZX-16** by Jan Deak: 16-channel polyphony on a single bit. Sixteen. The CPU does nothing but toggle the speaker bit using a carefully timed schedule that approximates the sum of 16 independent waveforms through pulse-density modulation.

- **Octode XL:** 8-channel beeper engine that actually leaves enough CPU for visuals.
- **Rain** by Life on Mars (2016): A full demo running a 9-channel beeper engine *simultaneously with visual effects* on a 48K Spectrum. The entire production – music and graphics – runs without an AY chip, without bank switching, without anything beyond the base machine.

These are among the most remarkable engineering achievements in 8-bit computing. They prove that limits are more permeable than they appear. But they are also impractical for general use: most beeper engines consume 50-90% of CPU time, leaving almost nothing for gameplay or effects. The AY chip exists precisely to offload sound generation to dedicated hardware.

We cover beeper engines here as historical context and inspiration. For practical music in your demos and games, the AY is where you should focus your effort.

Sidebar: Agon Light 2 - VDP Sound System

The Agon Light 2 takes a completely different approach to sound. Its audio is generated by the ESP32 co-processor (the VDP), not by a dedicated sound chip. You send VDU commands over the serial link, and the ESP32 synthesizes audio in software.

Waveforms: The Agon's sound system offers multiple waveform types per channel – square, sine, triangle, sawtooth, and noise. This is already more flexible than the AY's square-only tone generators.

ADSR Envelopes: Each channel has a fully programmable Attack-Decay-Sustain-Release envelope. No sharing – every channel gets its own independent envelope, unlike the AY's single shared envelope generator.

Channel Count: The VDP audio system supports multiple simultaneous channels (the exact number depends on the firmware version, but typically 8 or more).

The Trade-off: The Agon's sound is controlled through VDU byte sequences sent over serial. This means: - Higher latency than direct register writes (serial transfer time) - Less precise timing (you cannot bit-bang exact T-state-level sync) - But much less CPU overhead (the eZ80 just sends commands; the ESP32 does all synthesis)

The paradigm is closer to MIDI than to register-level programming. You tell the VDP “play this note on channel 3 with a sine wave and this ADSR envelope,” and it handles the rest. The musical goals are the same – melody, bass, drums, effects – but the programming model is fundamentally different. No mixer register, no period calculations, no envelope shape tables. Just commands and parameters.

For cross-platform projects, consider abstracting your sound system behind a common API: `sound_play_note(channel, note, instrument)`. On

the Spectrum, the instrument lookup writes AY registers. On the Agon, it sends VDU commands. Same music data, different backends.

11.8 Practical Exercises

Exercise 1: Register Explorer. Write a program that lets you modify any AY register in real time using keyboard input. Display all 14 register values on screen. This is your single most useful debugging tool for sound work.

Exercise 2: Three-Channel Arrangement. Compose a simple 16-bar tune using all three channels: melody on A, bass (buzz-bass using envelope) on C, and arpeggiated chords on B. Use the HALT-based timing from the example as a starting point, then refactor it into an IM2-driven player.

Exercise 3: Drum Kit. Implement four drum sounds (kick, snare, hi-hat, crash) as procedural SFX tables. Write a simple drum pattern that plays them in sequence. Integrate with the tune from Exercise 2.

Exercise 4: Vortex Tracker Integration. Download Vortex Tracker II, compose a short tune, export as .pt3, and integrate the standard .pt3 player into a Spectrum program. Verify it plays correctly in an emulator.

Summary

The AY-3-8910 is simple on paper: 14 registers, 3 channels, basic waveforms. But the gap between “simple” and “limited” is filled by technique. Arpeggios fake chords. Envelope abuse creates bass. Noise shaping synthesizes drums. Ornaments breathe life into static tones. And when three channels are genuinely not enough, TurboSound doubles them and the Next triples them.

The architecture pattern is consistent across all configurations: an interrupt fires 50 times per second, a player routine reads pattern data and calculates register values, and those values are blasted to the AY in a tight loop. Whether you are writing your own player or integrating Vortex Tracker, the flow is the same. Understanding the registers means understanding the sound.

In the next chapter, we will put this to work: digital drum samples blended with AY playback, frame-accurate synchronization between music and visuals, and the scripting engines that tie a demo together.

Sources: Dark “GS Sound System” (Spectrum Expert #01, 1997); Dark “Music” (Spectrum Expert #02, 1998); Shiru “Beeper 20XX” (Hype 2016); Rain file_id.diz; Info Guide #14 ASC Sound Master docs (ZXArt 2024); Vortex Tracker II documentation

Chapter 12: Digital Drums and Music Sync

“My brain is not coping with asynchronous coding well.” – Introspec, file_id.diz for the party version of Eager (to live), 3BM Open Air 2015

A demo is not a slideshow of effects. A demo is a performance – one where every visual event lands on the beat, every transition breathes with the music, and the audience never suspects that behind the curtain, a 3.5MHz processor is juggling half a dozen competing demands with no operating system, no threads, and no safety net.

This chapter is about the architecture that makes that juggling act possible. We have spent the previous chapters building individual effects – tunnels, zoomers, scrollers, colour animations – and in Chapter 11 we learned how the AY chip produces music. Now we must wire everything together. The questions are no longer “how do I draw a tunnel?” or “how do I play a note?” but rather: How do I play a drum sample that consumes nearly all the CPU while keeping the visuals smooth? How do I synchronise effect changes to the beat of the music? How do I structure a two-minute demo so that it runs reliably from start to finish?

The answers come from three sources. Introspec’s Eager (2015) gives us digital drum synthesis and asynchronous frame generation. diver4d’s GABBA (2019) shows a radically different approach to music sync using a video editor as a timeline tool. And Robus’s threading system (2015) demonstrates that honest multithreading on the Z80 is possible, if rarely necessary.

Together, these three techniques represent the architectural thinking that separates a collection of effects from a finished demo.

12.1 Digital Drums on the AY

The Problem: The AY Cannot Play Samples

The AY-3-8910, as we covered in Chapter 11, is a PSG (Programmable Sound Generator). It generates square waves, noise, and envelope shapes. It has no sample playback capability, no DAC, no waveform RAM. Every sound it makes is built from those primitive sources in real time. If you want a realistic kick drum – the kind with a sharp transient punch followed by a resonant decay – the AY’s noise

generator and envelope can approximate it, but the result sounds unmistakably synthetic. It lacks the weight of a real percussive hit.

But there is a back door.

Registers R8, R9, and R10 control the volume of channels A, B, and C. Each is a 4-bit value (0-15). If you write to a volume register once per frame, you get a static volume level. But what if you write to it thousands of times per frame? What if you treat the volume register as a crude 4-bit DAC and feed it successive sample values from a digitised recording?

You get PCM playback. Crude, noisy, 4-bit, but recognisable. The AY becomes a sample player – not by design, but by brute force.

The Cost: CPU Annihilation

Here is the problem. To play a digitised drum sample at any reasonable quality, you need to update the volume register at audio rates. A sample rate of 8 kHz means one update every 125 microseconds. At 3.5 MHz, 125 microseconds is approximately 437 T-states. That is tight but feasible – you can do useful work in the gaps between sample writes.

But 8 kHz sounds terrible. For a punchy kick drum, you want at least the perception of higher fidelity. And here the economics collapse. At higher effective sample rates, you need an interrupt or a tight polling loop that fires every 125-250 T-states. At that frequency, there is almost no CPU time left for anything else. While the drum sample is playing, the processor is a dedicated audio playback engine. Video generation, scripting, input handling – everything stops.

A typical kick drum sample lasts 20-40 milliseconds for the critical attack portion. At 50 Hz, that is 1-2 frames. During those frames, the CPU is gone.

n1k-o's Insight: The Hybrid Drum

n1k-o, the musician behind Eager's soundtrack, found the solution. The key observation: a drum sound has two distinct phases. The **attack** – the initial transient, the sharp “click” or “thud” that gives a kick drum its punch – is short, complex, and impossible to synthesise convincingly on the AY. But the **decay** – the resonant tail that follows – is a smooth volume falloff, exactly the kind of thing the AY's envelope generator handles naturally.

The hybrid approach: play the attack as a digital sample (consuming CPU time for 1-2 frames), then hand off to the AY's envelope generator for the decay (consuming zero CPU time, since the hardware does the work automatically). Digital attack plus AY decay equals a drum sound that has the realistic punch of a sample and the smooth tail of hardware synthesis.

In practice, the implementation works like this:

```
; Play hybrid kick drum
; 1. Start digital sample playback for attack phase
; 2. When sample ends, configure AY envelope for decay

play_kick_drum:
    di                                ; disable interrupts -- timing critical
```

```

; --- Digital attack phase ---
; Play ~800 samples at ~8kHz = ~100ms = ~2 frames
ld hl, kick_sample      ; pointer to 4-bit sample data
ld b, 0                 ; 256 samples per loop pass
ld c, $FD               ; low byte of AY data port ($BFFD)

; Select volume register R8 (channel A)
ld a, 8
ld bc, $FFFD
out (c), a              ; select R8
ld c, $FD               ; prepare for $BFFD writes

.sample_loop:
ld a, (hl)              ; 7 T - load sample byte
inc hl                  ; 6 T - advance pointer
ld b, $BF               ; 7 T - high byte of $BFFD
out (c), a              ; 12 T - write volume = sample value
; ... timing padding to hit target sample rate ...
djnz .sample_loop      ; 13 T (approx 45 T per sample)

; --- AY decay phase ---
; Configure envelope for smooth volume decay
; The AY takes over -- zero CPU cost from here

ld a, R_ENV_LO
ld e, 200               ; envelope period: moderate decay speed
call ay_write
ld a, R_ENV_HI
ld e, 0
call ay_write
ld a, R_ENV_SHAPE
ld e, $00               ; \___ single decay to silence
call ay_write
ld a, R_VOL_A
ld e, $10               ; switch channel A to envelope mode
call ay_write

ei
ret

kick_sample:
; 4-bit PCM data: attack portion of a kick drum
; Each byte = one sample, value 0-15
DB 0, 2, 8, 15, 14, 12, 15, 13
DB 10, 14, 11, 8, 12, 9, 6, 10
; ... (typically 400-800 bytes for the full attack)

```

The sample data itself – those 400-800 bytes of 4-bit PCM – comes from a real drum recording, downsampled and quantised to 4 bits. The attack transient preserves the character of the original instrument: the beater hitting the drum head, the initial compression of air, the sharp onset that our ears use to identify the sound.

The AY's envelope then provides a clean, smooth decay that our ears accept as the natural resonance of the drum body.

The result is convincing. On a chip that has no sample playback capability at all, you hear something that sounds like a real kick drum. Not studio quality, not even Amiga quality, but worlds better than pure AY synthesis.

The Frame Budget: Two Frames Per Hit

The frame cost is concrete: two frames per drum hit. During these frames, approximately 140,000 T-states (two full frame periods on Pentagon) are consumed by the sample playback loop. The CPU does nothing else. The display continues showing whatever was in screen memory, but no new frames are generated. No music data is processed (the drum IS the music for those two frames). No scripting engine runs.

Two frames at 50 Hz is 40 milliseconds. For a musical track with kick drums on the beat at 130 BPM, that is roughly one drum hit every 23 frames. Two frames out of every 23 consumed by drum playback – about 9% of total CPU time, delivered in sharp bursts that completely monopolise the processor.

This is the architectural challenge that drives the rest of the chapter. How do you keep the visuals running smoothly when the audio steals the CPU for two frames at a time, unpredictably, dozens of times per minute?

Sidebar: MCC - More Than 16 Levels from One AY

The single-channel approach above gives you 16 volume levels (4-bit DAC). But the AY has three channels, and their analogue outputs are mixed together before reaching the speaker. What if you write different volume values to all three channels simultaneously? The combined signal has more amplitude steps than any single channel.

This is the MCC (Mixed Channel Covox) technique, documented by UnBEL!EVER (Born Dead #09, 1999). Two “edge” channels provide the main signal; the centre channel adds a correction value. With a pre-computed lookup table mapping 8-bit sample values to three register values, you can synthesise approximately 108 distinct amplitude levels – far better than 16.

The standard MCC playback loop runs at ~24 KHz (144 T per sample). A super-fast variant by MOn5+Er (Born Dead #0G, 2000) uses SP as the sample data pointer (POP HL reads two bytes in 10 T) and JP (IX) for the loop branch (8 T), achieving ~43.75 KHz at 80 T per sample – approaching telephone quality from a sound chip designed for square waves.

The catch: the MCC lookup table depends on the exact analogue output levels of each channel. The AY-3-8910 and YM2149F have different volume curves (the YM is closer to linear, the AY is more logarithmic), and even chips of the same model vary between production runs. UnBEL!EVER's 108-level table was calculated for the YM2149F. On a different chip, the combined levels shift, and the carefully calibrated amplitude steps become uneven. The centre channel contributes roughly 52% of

the mixed signal – if that ratio drifts, the correction values distort rather than improve the output.

For demo work targeting a specific machine (say, Pentagon with YM2149F), MCC works well. For cross-hardware compatibility, the single-channel 4-bit approach is safer. And for fun experiments – like feeding AY-beat formulas (Appendix I) through an averaged MCC table – the distortion is part of the charm.

Sources: UnBEL!EVER, “Воспроизведение оцифровок на AY (MCC),” Born Dead #09 (1999); MOn5+Er, Born Dead #0G (2000).

12.2 Asynchronous Frame Generation

The Naive Approach Fails

The simplest demo architecture is synchronous: generate one frame of the visual effect, wait for HALT (vsync), display it. Generate the next frame, HALT, display it. This is what we built in every practical exercise so far. It works perfectly when frame generation takes less than one frame period.

Now add digital drums. The music engine signals: “play kick drum on the next beat.” The sample playback routine seizes the CPU for two frames. During those two frames, no new video frames are generated. When the drum finishes, the display has been showing the same frame for three refreshes (the last generated frame was shown once normally, then twice during the drum hit). The visual effect stutters.

With one drum hit every 23 frames, the audience sees a brief freeze every half-second. It is noticeable. It is ugly. It is unacceptable for a competition demo.

Introspec’s Solution: Stockpile Frames

Introspec’s architecture in Eager decouples frame generation from frame display. The visual engine does not generate one frame and immediately show it. Instead, it generates frames into a buffer – as many as it can fit – and the display system shows them at a steady 50 Hz rate regardless of what the generator is doing.

The mechanism is double-buffered attribute frames. Two pages of attribute data exist in memory. While one page is being displayed (the ULA reads from it during the screen refresh), the generator writes the next frame into the other page. When a new frame is ready, the engine flips the pages: the newly generated frame becomes the display page, and the old display page becomes the new generation target.

```
Display:  [Frame 1] [Frame 2] [Frame 3] [Frame 4] [Frame 5]
Generator: —gen F2—|—gen F3—|—gen F4—|— DRUM —|—gen F5—
                ↑                ↑
                drum starts    drum ends
```

During the drum hit, the display shows Frame 4 (already generated). Frame 5 generation resumes immediately after the drum finishes.

But simple double buffering only gives you one frame of slack. If the drum consumes two frames, you need to have generated two frames ahead. This is where Introspec's asynchronous generation truly diverges from simple double buffering: the engine can **stockpile** multiple frames in advance.

On the 128K Spectrum, memory banking provides the space. Attribute frames are small – 768 bytes each. A single 16KB memory page can hold roughly 20 attribute frames. The generator runs as fast as it can, writing frame after frame into the buffer. The display system reads from the buffer at a steady 50 Hz. When the generator is faster than real time (which it usually is, since attribute plasma is cheap), the buffer fills up. When a drum hit pauses generation, the display system draws down the buffer. As long as the buffer does not run dry, the audience sees smooth 50 Hz animation.

The Buffer Dynamics

Think of it as a producer-consumer problem, but on a machine with no concurrency.

The **producer** is the plasma/tunnel/zoomer effect generator. It produces attribute frames at a variable rate – sometimes faster than 50 Hz (when the calculation is simple and no drums are playing), sometimes zero (during drum playback).

The **consumer** is the display system, reading one frame per screen refresh at exactly 50 Hz.

The **buffer** sits between them, absorbing the difference.

The dynamics are straightforward:

- **Between drum hits:** The generator runs faster than the display. The buffer fills up. If it reaches capacity, the generator idles (or the engine advances the scripting state).
- **During a drum hit:** The generator stops. The display drains the buffer at 50 Hz. A two-frame drum hit consumes two buffered frames.
- **After a drum hit:** The generator resumes, running as fast as possible to refill the buffer before the next hit.

The critical constraint: **the buffer must never run dry during a drum hit**. If two drum hits occur in rapid succession – say, a kick-snare pattern two frames apart – the buffer needs at least four frames of reserve. Introspec's scripting engine manages this by knowing the music timeline in advance. When a dense drum passage approaches, the engine generates extra frames to pad the buffer. When a quiet passage follows, the buffer naturally fills.

The catch: if the drum pattern is too dense – too many hits too close together – the generator cannot keep up. The buffer runs dry, and the display repeats a frame. This is a hard constraint of the architecture, and it influenced n1k-o's composition. The music was written with knowledge of the engine's capacity: drum hits are spaced far enough apart that the generator can always recover. The musician and the coder designed together, each understanding the other's constraints.

12.3 The Scripting Engine

Why You Need a Script

By this point, the list of things that need coordination is long:

- The visual effect generator (which effect is active, what parameters it uses)
- The music player (which pattern is playing, when drums trigger)
- The frame buffer (how full it is, when to generate more)
- Transitions between effects (fade out one, fade in the next)
- The overall timeline (the demo runs for two minutes – what happens when)

You could hardcode all of this in a monolithic main loop. Some demos do. But Introspec chose a different path: a two-level scripting system that separates *what happens* from *when it happens*.

Outer Script: The Sequence of Effects

The outer script is a linear sequence of commands that control the overall structure of the demo. Think of it as a setlist for a concert:

```
EFFECT  tunnel, params_set_1      ; start the tunnel effect
WAIT    200                      ; run for 200 frames (4 seconds)
EFFECT  zoomer, params_set_1      ; switch to chaos zoomer
WAIT    150                      ; 3 seconds
EFFECT  tunnel, params_set_2      ; tunnel again, different colours
WAIT    250                      ; 5 seconds
; ... and so on for the full demo
```

Each EFFECT command loads the generator function and its parameter block. Each WAIT tells the engine how many frames to run the current effect before advancing to the next command. Transitions between effects – crossfades, hard cuts, colour sweeps – are themselves scripted as effects.

Inner Script: Variations Within an Effect

Within a single effect, parameters change over time. The tunnel’s plasma frequencies shift, the colour palette rotates, the zoom speed accelerates. These variations are controlled by the inner script – a per-effect sequence of parameter changes keyed to frame numbers:

```
FRAME  0:  plasma_freq = 3, palette = warm
FRAME  50: plasma_freq = 5                ; frequency shift
FRAME 100: palette = cool                 ; colour change
FRAME 120: plasma_freq = 2, palette = hot ; both change
```

The inner script runs independently of the outer script. When the outer script says “run tunnel for 200 frames,” the inner script handles the visual evolution within those 200 frames.

kWORK: The Key Command

The most important command in the scripting system is what Introspec calls **kWORK**: “generate N frames, then show them independently of generation.” This

single command is the bridge between the scripting system and the asynchronous architecture.

When the engine encounters `kWORK 8`, it:

1. Generates 8 frames of the current effect into the frame buffer.
2. Hands those frames to the display system.
3. While the display system shows them (over $8/50 = 160\text{ms}$), the engine is free to do other work: process the next script command, prepare the next batch, or yield CPU time for drum playback.

This decoupling – generate now, display later – is the fundamental enabler for asynchronous operation. Without `kWORK`, the engine would be locked into a synchronous generate-display-generate-display cycle with no slack for drum interruptions.

In practice, the engine calls `kWORK` repeatedly, generating small batches of frames (4-8 at a time). Between batches, it checks whether a drum trigger is pending. If so, it lets the drum play, knowing that the display system has enough buffered frames to continue smoothly. After the drum finishes, it generates the next batch to replenish the buffer.

```
; Simplified engine loop (conceptual)
engine_loop:
    ; Check if drum is pending
    ld    a, (drum_pending)
    or    a
    jr    z, .no_drum
    call  play_drum          ; consumes 2 frames of CPU time
    xor   a
    ld    (drum_pending), a

.no_drum:
    ; Generate a batch of frames
    call  generate_batch     ; kWORK: produce N frames into buffer
    ; (generate_batch returns when batch is done)

    ; Check outer script for effect changes
    call  advance_script

    jr    engine_loop
```

The beauty of this architecture is its simplicity at the macro level. The engine is a loop: check for drums, generate frames, advance the script. All the complexity is inside `generate_batch` (which manages the buffer, handles the plasma calculation, and writes attribute data) and `play_drum` (which runs the digital sample routine from section 12.1). The scripting system provides the sequencing; the buffer provides the temporal decoupling; the drum routine provides the audio impact. Each component has a clear responsibility.

12.4 GABBA's Innovation: The Video Editor as Timeline Tool

In 2019, diver4d (of 4D+TBK) took first place at CAFE with GABBA, a gabber-themed demo with punishingly tight audio-visual synchronisation. The sync was so precise that every visual hit landed exactly on the musical beat, every transition matched a phrase boundary, and the whole production felt like a music video rather than a demo.

The technical surprise was in the workflow, not the code.

The Problem with Code-Based Sync

The traditional approach to music sync in ZX demos is to embed timing data in the code. You know that the kick drum hits at frame 47, so you write a script command that triggers the visual event at frame 47. Then you watch the demo, decide the timing is slightly off, change the number to 49, recompile, re-test, and repeat. For a two-minute demo at 50 fps, that is 6,000 frames of potential sync points. Getting them all right by trial and error takes weeks.

Introspec's Eager was built this way, and the development was gruelling. Every sync adjustment required recompilation – assembling the Z80 code, loading the binary into an emulator, watching the relevant section, noting what was off, editing the source, and repeating. The feedback loop was measured in minutes per iteration.

diver4d's Answer: Luma Fusion

diver4d bypassed the code-edit-compile-test cycle entirely. He used **Luma Fusion**, an iOS video editor, as his synchronisation tool.

The workflow:

1. **n1k-o composed the gabber track**, then exported it from Vortex Tracker into Excel. In the spreadsheet, he built a colour-coded visual map of the entire track: every row is one frame (= one pattern row), with columns for each musical layer — kick drums in blue, snare in red, melody in green, acid in purple, and so on. Extra columns held frame numbers and any sync data the coders needed. He also highlighted subtle effects that non-musicians might not hear. The result was a frame-accurate, human-readable map of the entire composition. The reason for this effort was practical: the coders heard gabber as a wall of sound and could not identify individual beats or transitions by ear. The spreadsheet made the musical structure visible. This workflow proved so effective that the team adopted it for all subsequent demos.
2. **diver4d recorded each visual effect** running at 50 fps in an emulator and exported the recordings as video clips.
3. **In Luma Fusion**, he arranged the video clips on a 50 fps timeline alongside the audio track. He could scrub through the demo frame by frame, seeing exactly how each visual aligned with each musical event. Moving a transition was as simple as dragging a clip on the timeline.
4. **Once the timing was right in the editor**, he extracted the frame numbers for each transition and effect change, and wrote those numbers into the Z80 script data.




Figure 29: n1k-o's demoplan for Unspoken – colour-coded sync map exported from Vortex Tracker into Excel. Each row is one frame, coloured columns mark musical layers: drums (navy), hi-hats (gold), melody (green), bass (purple). This is the same workflow pioneered for GABBA and used for all subsequent demos by the team.

The insight is straightforward: use the right tool for the job. A video editor is purpose-built for frame-level multimedia synchronisation. Z80 assembly is not. By doing the creative synchronisation work in the editor and the implementation work in assembly, *diver4d* separated the artistic decisions from the engineering constraints.

What This Changes

The immediate benefit is speed. Adjusting sync timing in a video editor takes seconds. Adjusting it in assembly takes minutes. Over hundreds of sync points, the cumulative time saving is enormous. But the deeper benefit is creative freedom. When iteration is cheap, you experiment more. You try the transition two frames early, see how it feels, try it two frames late. You notice that the visual works better hitting slightly *before* the beat (a technique borrowed from film editing, where cuts on the beat feel late because of human reaction time). You could never discover this insight through code-based iteration – the feedback loop is too slow.

The limitation is that this workflow works best for demos where the timing is fixed – where the demo always plays the same way. If you want interactive or generative elements that respond to runtime conditions, you need the code-based approach. But for the overwhelming majority of ZX demos, which are linear fixed-timeline productions, the video editor workflow is superior.

GABBA demonstrated that demoscene production tools do not have to be retro. The Z80 code is from 1985. The sync workflow can be from 2019. There is no contradiction.

12.5 Z80 Threading: A Different Path

Robus, writing in *Hype* in 2015, presented a technique that attacks the concurrency problem from a completely different angle: actual multithreading on the Z80.

The Problem, Restated

The fundamental tension in a demo engine is that multiple tasks need CPU time in the same frame: effect generation, music playback, drum samples, scripting, transitions. *Introspec*'s solution is cooperative: the engine manually interleaves these tasks using a scripting system and frame buffer. This works, but it requires careful manual scheduling and the entire asynchronous architecture we have been discussing.

What if the Z80 could run two tasks simultaneously?

IM2-Based Context Switching

It can, after a fashion. The Z80's IM2 interrupt provides a natural context switch point. Every frame, the interrupt fires. If the interrupt handler saves the current task's state and loads another task's state, you have preemptive multithreading.

Robus's `SwitchThread` procedure does exactly this:


```

; SwitchThread: save current thread, resume next thread
; Called from within the IM2 interrupt handler
SwitchThread:
    ; Save current thread's stack pointer
    ld    (thread_sp_save), sp

    ; Save current memory page configuration
    ld    a, (current_7ffd)
    ld    (thread_page_save), a

    ; Load next thread's state
    ld    a, (next_thread_page)
    ld    (current_7ffd), a
    ld    bc, $7FFD
    out   (c), a                ; switch memory page

    ld    sp, (next_thread_sp) ; switch stack pointer

    ; Execution continues in the next thread's context
    ; (it was previously suspended at this same point)
    ret

```

Each thread gets its own **128-byte stack** and a **dedicated memory page** (one of the 128K Spectrum's eight 16KB banks). The stack is small but sufficient – Z80 code rarely nests deeply. The dedicated memory page gives each thread its own workspace without interfering with the other.

How It Works in Practice

In Robus's WAYHACK demo, two threads run concurrently:

- **Thread 1:** Calculates the visual effect (a dungeon-crawling perspective renderer).
- **Thread 2:** Renders scrolling text along the bottom of the screen.

Neither thread knows about the other. Each runs in its own memory page with its own stack. Every frame, the IM2 interrupt fires and SwitchThread alternates between them. Thread 1 gets one frame of CPU time, then Thread 2 gets one frame, and so on.

The result: the text scroller runs at a steady 25 Hz (every other frame), and the visual effect runs at 25 Hz. Neither task needs to be aware of the other's existence. No cooperative scheduling, no yield points, no manual interleaving. The interrupt handles everything.

The Threading Model

The model is simple:

```

Frame 2: Interrupt → save Thread 1 → restore Thread 2 → Thread 2 runs
Frame 3: Interrupt → save Thread 2 → restore Thread 1 → Thread 1 runs
...

```

Each thread sees a consistent world: its registers, its stack, its memory page. The switch happens at a fixed point (the interrupt), so there are no race conditions on shared data. If the threads need to communicate (e.g., Thread 1 signals Thread 2 to change the text), they do so through a shared memory location that both threads can access – a simple flag or mailbox.

Practical Considerations

Robus’s own assessment is characteristically honest: **“Honest multithreading rarely requires more than two threads”** on the Z80. The overhead of context switching (saving and restoring SP plus a memory page switch) is modest – perhaps 100 T-states – but each additional thread halves the available CPU time per thread. With two threads, each gets 25 Hz. With three, each gets roughly 16.7 Hz. On a machine where visual smoothness demands close to 50 Hz, two threads is the practical limit.

The threading approach is orthogonal to Introspec’s asynchronous buffering approach. You could combine them: one thread generates effect frames into a buffer while the other handles music and drum playback. In practice, this combination is rare – the two techniques solve the same problem (interleaving CPU-hungry tasks) through different mechanisms, and most demo coders choose one or the other based on the specific demands of their production.

Threading works best when two tasks are truly independent and neither needs more than 25 Hz. The asynchronous buffer approach works best when one task (visuals) needs 50 Hz and the other (drums) needs unpredictable bursts. For Eager’s architecture, where visual smoothness was paramount and drum timing was dictated by the music, the buffer approach won. For WAYHACK’s architecture, where two steady-state tasks ran in parallel, threading won.

12.6 Practical: A Minimal Scripted Demo Engine

Let us build a minimal demo engine that ties together the concepts from this chapter. The goal is not Eager-level sophistication – it is a skeleton that demonstrates the architecture.

What We Build

- **Three simple effects:** plasma (attribute-based, from Chapter 9), colour bars (horizontal attribute stripes), and a text scroller.
- **AY music** playing via IM2 interrupt (using a .pt3 player, as described in Chapter 11).
- **A digital kick drum sample** that plays on the beat, stealing 2 frames of CPU.
- **A simple timeline script** that switches between effects at defined points.
- **Double-buffered attributes** to absorb the drum hit pauses.

The Memory Map

\$8000-\$9FFF	Music player + song data
\$A000-\$AFFF	Sine tables, colour maps, sample data

```

$B000-$BFFF   Frame ring buffer (attribute frames)
$C000-$DFFF   Shadow screen (second display page)
$E000-$FFFF   Stack + IM2 vector table + workspace

Bank 0-3:      Not used (available for larger effects)
Bank 5:        Normal screen ($4000-$5AFF display)
Bank 7:        Shadow screen ($C000-$DAFF display)

```

The Timeline Script

```

; Timeline script: sequence of (effect_id, duration_frames, param_ptr)
timeline:
    DB EFFECT_PLASMA,    0, 150    ; plasma for 150 frames (3 sec)
    DW plasma_params_1
    DB EFFECT_BARS,      0, 100    ; colour bars for 100 frames (2 sec)
    DW bars_params_1
    DB EFFECT_SCROLLER,  0, 200    ; text scroller for 200 frames (4 sec)
    DW scroller_params_1
    DB EFFECT_PLASMA,    0, 150    ; plasma again, different params
    DW plasma_params_2
    DB $FF                ; end marker: loop from start

EFFECT_PLASMA    EQU 0
EFFECT_BARS      EQU 1
EFFECT_SCROLLER  EQU 2

```

The Main Engine Loop

```

; Main engine loop
; Assumes IM2 is set up and music player runs in the ISR

engine_init:
    ; Set up display: fill pixel memory with checkerboard
    call fill_checkerboard

    ; Initialise ring buffer
    xor    a
    ld     (buf_write_idx), a
    ld     (buf_read_idx), a
    ld     (buf_count), a

    ; Load first effect from timeline
    ld     hl, timeline
    ld     (script_ptr), hl
    call load_next_effect

engine_main:
    ; === Step 1: Check for drum trigger ===
    ld     a, (drum_pending)
    or     a
    jr     z, .no_drum

```

```

; Play the drum -- this consumes ~2 frames
call play_kick_drum
xor a
ld (drum_pending), a
jr .after_drum

.no_drum:
; === Step 2: Generate a frame into the buffer ===
ld a, (buf_count)
cp BUF_CAPACITY ; buffer full?
jr nc, .buffer_full

; Generate one frame of the current effect
call generate_frame ; writes 768 bytes to ring buffer

; Advance buffer write pointer
ld a, (buf_write_idx)
inc a
cp BUF_CAPACITY
jr nz, .no_wrap_w
xor a
.no_wrap_w:
ld (buf_write_idx), a
ld a, (buf_count)
inc a
ld (buf_count), a

.buffer_full:
.after_drum:
; === Step 3: Advance timeline ===
ld hl, (frame_counter)
inc hl
ld (frame_counter), hl

; Check if current effect duration has elapsed
ld de, (effect_duration)
or a
sbc hl, de
jr c, .effect_continues

; Load next effect from timeline
call load_next_effect
ld hl, 0
ld (frame_counter), hl

.effect_continues:
; === Step 4: Wait if we are ahead of display ===
halt ; sync to frame boundary

jr engine_main

```

The Display ISR

```

; IM2 interrupt handler: runs every frame (50 Hz)
frame_isr:
    push af
    push bc
    push de
    push hl

    ; Play music (updates AY registers)
    call music_play

    ; Check if music engine signals a drum hit
    ld  a, (music_drum_flag)
    or  a
    jr  z, .no_drum_signal
    xor a
    ld  (music_drum_flag), a
    ld  a, 1
    ld  (drum_pending), a    ; signal main loop
.no_drum_signal:

    ; Display next frame from ring buffer
    ld  a, (buf_count)
    or  a
    jr  z, .no_frame        ; buffer empty, keep current frame

    ; Copy buffered attributes to display page
    call copy_buf_to_screen

    ; Advance read pointer
    ld  a, (buf_read_idx)
    inc a
    cp  BUF_CAPACITY
    jr  nz, .no_wrap_r
    xor a
.no_wrap_r:
    ld  (buf_read_idx), a
    ld  a, (buf_count)
    dec a
    ld  (buf_count), a

.no_frame:
    pop hl
    pop de
    pop bc
    pop af
    ei
    reti

BUF_CAPACITY EQU 8          ; 8 frames of buffer (8 x 768 = 6,144 bytes)

```

The Effect Generator Dispatch

```

; Generate one frame of the current effect
; Writes attribute data to the ring buffer
generate_frame:
    ld    a, (current_effect)
    or    a
    jr    z, .do_plasma
    cp    1
    jr    z, .doBars
    cp    2
    jr    z, .do_scroller
    ret

.do_plasma:
    call  calc_plasma          ; from Chapter 9 -- writes 768 bytes
    ret
.doBars:
    call  calc_colourBars      ; horizontal attribute stripes
    ret
.do_scroller:
    call  calc_text_scroll     ; text rendering into attributes
    ret

```

Observations

This skeleton is deliberately simple. A production engine would add:

- **Inner scripts** for parameter variation within each effect.
- **Transition effects** (crossfades between two attribute buffers).
- **Multiple drum sounds** (kick, snare, hi-hat), each with its own sample data.
- **Buffer level monitoring** so the generator can prioritise catching up after dense drum passages.
- **Memory banking** to store more frames and support larger effect data.

But even in this minimal form, the architecture demonstrates the key principles:

1. **Decoupled generation and display.** The generator and the display ISR communicate only through the ring buffer. Neither knows or cares about the other's timing.
2. **Drum hits are absorbed by the buffer.** When `play_kick_drum` consumes two frames, the display ISR continues showing buffered frames. The audience sees no stutter.
3. **The script drives the timeline.** Adding a new effect or changing the sequence means editing the timeline data table, not restructuring the engine code.
4. **The music player runs in the ISR.** It updates AY registers every frame regardless of what the main loop is doing. The only interaction is the `drum_pending` flag – a one-byte mailbox between the ISR and the main loop.

This is the architecture of a demo. Not the effects, not the music, not the art – the *plumbing* that makes all of those work together. It is the least visible part of a

demo and the hardest to get right. Introspec spent ten weeks on Eager, and the architecture consumed more of that time than any single effect.

12.7 Practical Exercises

Exercise 1: Basic Engine. Implement the skeleton above with a single effect (plasma from Chapter 9) and no drum samples. Verify that the ring buffer works correctly: the display shows smooth animation while the generator runs at its natural speed.

Exercise 2: Add the Drum. Record (or synthesise) a 4-bit kick drum sample (400-800 bytes). Add the `play_kick_drum` routine and trigger it every 25 frames. Verify that the display remains smooth during drum playback. What is the maximum drum rate before the buffer runs dry?

Exercise 3: Multi-Effect Timeline. Add a second effect (colour bars or text scroller). Write a timeline script that switches between effects every 3-4 seconds. Verify that transitions happen at the correct frame.

Exercise 4: Sync to Music. Load a short .pt3 tune and modify the player to set `music_drum_flag` when a particular pattern event occurs (e.g., a note on channel C below a certain pitch). Now the drums are driven by the music, not by a fixed frame counter. This is real music sync.

Exercise 5: Video Editor Workflow. Record your running demo in an emulator at 50 fps. Import the recording into a video editor (any editor that supports frame-level editing). Adjust the timeline script frame numbers based on what you see in the editor. Experience the difference in iteration speed compared to code-only sync.

Summary

This chapter was not about any single effect or technique. It was about architecture – the invisible structure that lets a demo exist as a coherent, synchronised, two-minute performance rather than a collection of disjointed screens.

The core problems are universal. Every demo engine must answer: How do I share the CPU between audio and video? How do I keep the display smooth when the audio steals processing time? How do I sequence effects and synchronise them to music? How do I manage the timeline of a multi-minute production?

The solutions we examined are complementary:

- **Digital drums** (n1k-o/Introspec) exploit the AY's volume registers as a crude DAC, blending digital samples with hardware synthesis to produce percussion that transcends the chip's designed capabilities.
- **Asynchronous frame generation** (Introspec) decouples video production from display through a ring buffer, absorbing the CPU bursts consumed by drum playback.

- **Scripted timelines** (Introspec) separate the *what* and *when* of a demo from the *how*, making it possible to design and adjust a two-minute production without restructuring the engine.
- **Video editor sync** (diver4d) moves the creative timing work to a tool purpose-built for it, dramatically accelerating the sync iteration cycle.
- **Z80 threading** (Robus) provides genuine concurrency for tasks that are independent and steady-state, at the cost of halving the frame rate for each task.

With this chapter, we close the circle on the demoscene section of the book. We have built effects (Parts I-II), made sound (Chapter 11), and now wired everything into a running engine. The reader who has followed from Chapter 1 has a complete picture: from T-state counting to a synchronised, scripted demo with digital drums.

In the next chapter, we shift gears entirely. Part IV takes us into size-coding – the art of fitting an entire production into 256 bytes. The architecture goes from “how do I manage a ring buffer?” to “how do I make every single byte do double duty?” The constraints tighten by three orders of magnitude, and the thinking changes to match.

Sources: Introspec, “Making of Eager,” Hype, 2015 (hype.retroscene.org/blog/demo/261.html); Introspec, `file_id.diz` from Eager (to live), 3BM Open Air 2015; diver4d, “Making of GABBA,” Hype, 2019 (hype.retroscene.org/blog/demo/948.html); Robus, “Threads on Z80,” Hype, 2015 (hype.retroscene.org/blog/dev/271.html); Eager source code excerpts courtesy of Introspec (Life on Mars)

Chapter 13: The Craft of Size-Coding

“It was like playing puzzle-like games – constant reshuffling of code to find shorter encodings.” – UriS, on writing NHBF (2025)

There is a category of demoscene competition where the constraint is not time but *space*. Your entire program – the code that draws the screen, produces the sound, handles the frame loop, holds whatever data it needs – must fit in 256 bytes. Or 512. Or 1K, or 4K, or 8K. Not a byte more. The file is measured, and if it is 257 bytes, it is disqualified.

These are **size-coding** competitions, and they produce some of the most remarkable work on the ZX Spectrum scene. A 256-byte intro that fills the screen with animated patterns and plays a recognisable melody is a form of compression so extreme it is hard to believe until you read the code. The gap between what the audience sees and the file size that produces it – that gap is the art.

This chapter is about the mindset, the techniques, and the specific tricks that make size-coding possible.

13.1 What is Size-Coding?

Demo competitions typically offer several size-limited categories:

Category	Size limit	What fits
256 bytes	256	One tight effect, maybe simple sound
512 bytes	512	An effect with basic music or two simple effects
1K intro	1,024	Multiple effects, proper music, transitions
4K intro	4,096	A short demo with several parts
8K intro	8,192	A polished mini-demo

The limits are absolute. The file is measured in bytes, and there is no negotiation.

What makes size-coding fascinating is that it inverts the normal optimisation hierarchy. In the cycle-counted world of demoscene effects, you optimise for *speed* – unrolling loops, duplicating data, generating code, all trading space for time. Size-coding inverts this. Speed does not matter. Readability does not matter. The only question is: can you make it one byte shorter?

UriS, who wrote the 256-byte intro NHBF for Chaos Constructions 2025, described the process as “playing puzzle-like games.” The description is exact. Size-coding is a puzzle where the pieces are Z80 instructions, the board is 256 bytes of RAM, and the best solutions involve moves that solve multiple problems simultaneously.

The mindset shift:

- **Every byte is precious.** A 3-byte instruction where a 2-byte one suffices is 0.4% of your entire program. At 256 bytes, one byte saved is like saving 250 bytes in a 64K program.
- **Code and data overlap.** The same bytes that execute as instructions can serve as data. The Z80 does not know the difference – only the program counter’s path through memory distinguishes code from data.
- **Instruction choice is driven by size, not speed.** RST \$10 costs 1 byte. CALL \$0010 does the same thing in 3 bytes. In a normal demo you would never notice. In 256 bytes, those 2 bytes are the difference between having sound or not.
- **Initial state is free data.** After boot, registers have known values. Memory at certain addresses contains known data. A size-coder exploits every bit of this free state.
- **Self-modifying code is not a trick – it is a necessity.** When you cannot afford a separate variable, you modify an instruction’s operand in place.

The Z80 Size-Coder’s Toolkit

Some tricks recur so often in size-coded intros that they form a shared vocabulary. Knowing these by heart – their byte costs, their side effects – is the prerequisite to serious size-coding.

Register initialization assumptions. When a Spectrum program launches from BASIC (via RANDOMIZE USR), the CPU state is not random. After CLEAR and before the USR call, A is typically 0, BC holds the USR address, DE and HL have known values from the BASIC interpreter, the stack pointer sits at the CLEAR address, and interrupts are enabled. Many of these are stable enough to rely on. If your program needs A = 0 at the start, do not write XOR A – it is already zero. If you need a 16-bit counter starting at 0, check whether DE or HL already holds 0 or a useful value. One byte saved here, two bytes saved there – these add up to the difference between 260 bytes and 256.

The system variables area (\$5C00-\$5CB5) is another source of free data. The BASIC interpreter maintains over 100 bytes of state at known addresses. If you need the value 2, you might find it at the address holding the current stream number. If you need \$FF, several system variable fields contain it. Reading from a fixed address costs 3 bytes (LD A, (nn)), but if it replaces a 2-byte load *plus* some computation, you come out ahead.

DJNZ as a short backwards jump. DJNZ label is 2 bytes, same as JR label – but it also decrements B. If B is nonzero and you need a backwards jump, DJNZ does both for free. Even when you do not care about B’s decrement, DJNZ is still 2 bytes, the same cost as JR. But if B happens to reach zero at the exact moment you want to fall through, you have merged a loop counter and a branch into a single instruction. Size-coders routinely structure loops so that B’s natural countdown aligns with the exit condition.

RST as a 1-byte CALL. The Z80 reserves eight restart addresses: \$00, \$08, \$10, \$18, \$20, \$28, \$30, \$38. RST *n* pushes the return address and jumps to the target – the same as CALL *n* – but in 1 byte instead of 3. On the Spectrum, the ROM places useful routines at several of these addresses:

- RST \$10 – print a character (ROM routine at \$0010)
- RST \$20 – collect next character from BASIC (less useful for demos)
- RST \$28 – enter the floating-point calculator (useful for math)
- RST \$38 – the maskable interrupt handler (IM 1 jumps here)

In a normal demo, these ROM routines are too slow to call in a tight loop. In a 256-byte intro, saving 2 bytes per call is worth the speed penalty. If your program calls RST \$10 six times to print characters, that is 12 bytes saved over six CALL \$0010 instructions. Twelve bytes is nearly 5% of 256.

Overlapping instructions. The Z80 decodes instructions byte by byte, with no alignment requirements. If you jump into the middle of a multi-byte instruction, the CPU decodes fresh from that point. This means you can hide one instruction inside another:

```
ld    a, $AF                ; opcode $3E, operand $AF
                                ; BUT: $AF is XOR A
```

If the CPU executes from the start, it sees LD A, \$AF (2 bytes). If another code path jumps to the second byte, it sees XOR A (1 byte). One byte serves two purposes. The technique is fragile – it demands perfect control of all execution paths – but in competition code, fragility is acceptable.

A common pattern: the byte \$18 is JR d (relative jump). If you need the value \$18 as data *and* need a branch at that location, the same byte does both. The operand that follows is both the jump offset and (from another perspective) the next piece of data.

Abusing flag state. Every arithmetic and logical instruction sets flags. Size-coders memorise which flags each instruction affects and exploit the results instead of computing them separately. After DEC B, the zero flag tells you whether B hit zero – no CP 0 needed. After ADD A, n, the carry flag tells you whether the result overflowed past 255. After AND mask, the zero flag tells you whether any masked bits were set.

The deepest flag trick is SBC A, A: if carry is set, A becomes \$FF; if carry is clear, A becomes \$00. One byte, no branch, a full bitmask from a flag. Compare this to the branching alternative:

```
    ; With branching: 6 bytes
    jr    nc, .zero          ; 2
    ld    a, $FF             ; 2
    jr    .done              ; 2
.zero:
    xor   a                  ; 1
.done:

    ; With SBC A,A: 1 byte
    sbc   a, a               ; 1 – carry -> $FF, no carry -> $00
```

Five bytes saved. In a 256-byte intro, that is two percent of the entire program.

13.2 Anatomy of a 256-Byte Intro: NHBF

NHBF (No Heart Beats Forever) was created by UriS for Chaos Constructions 2025, inspired by RED REDUX from Multimatograf 2025. It produces text with screen effects and music – looped square-wave power chords with random pentatonic melody notes – all in 256 bytes.

The Music

At 256 bytes, you cannot include a tracker player or note tables. NHBF drives the AY chip directly. The power chords are hardcoded as immediate values in the AY register-write instructions – the same bytes that form the LD A, n operand *are* the musical note. The melody channel uses a pseudo-random generator (typically LD A, R – read the refresh register – followed by AND to mask the range) to pick from a pentatonic scale. A pentatonic scale sounds pleasant regardless of which notes land next to each other, so the melody sounds intentional even though it is random. Two bytes for a “random” number; five notes that never clash.

The Visual

Printing text through the ROM – RST \$10 outputs a character for 1 byte per call – is the cheapest way to get pixels on screen. But even a 20-character string costs 40 bytes (character codes + RST calls). Size-coders look for ways to compress further: overlapping the string data with other code, or computing characters from a formula.

The Puzzle: Finding Overlaps

UriS describes the core process as constant reshuffling. You write a first version at 300 bytes, then stare at it. You notice the loop counter for the visual effect ends up holding the value you need as the AY register number. Remove the LD A, 7 that would set it – the loop already left A with 7. Two bytes saved. The screen-clearing routine uses LDIR, which decrements BC to zero. Arrange the code so the next section needs BC = 0, and save the LD BC, 0 – another 3 bytes.

Every instruction produces side effects – register values, flag states, memory contents – and the art is arranging instructions so that one routine’s side effects are another routine’s inputs.

Art-Top’s Discovery

During development, Art-Top noticed something remarkable: the register values left over from the screen-clearing routine happened to match the exact length needed for the text string. Not planned. UriS had written the screen clear, then the text output, and the two happened to share a register state that eliminated a separate length counter.

This kind of serendipitous overlap is the heart of 256-byte coding. You cannot plan for it. You can only create conditions where it can happen, by constantly rearranging code and watching for accidental alignments. When you find one, it

feels like discovering that two jigsaw pieces from different puzzles fit together perfectly.

The Byte Budget

When working at 256 bytes, a rough budget helps you plan before writing a single instruction. Here is a realistic breakdown for a typical ZX Spectrum 256-byte intro with both visuals and sound:

Component	Bytes	Notes
Pixel fill (dither/clear)	18-25	LD HL, LDIR or a compact fill loop
AY initialisation	16-22	Mixer, volume, initial tone — via port writes
Main loop frame sync	1	HALT
AY tone update per frame	10-14	Select register, write tone period
Visual effect core	30-50	The inner loop that computes and writes attributes
Outer loop / row control	8-12	Row counter, column counter, branches
Frame counter update (SMC)	6-8	Read, increment, write back into instruction
Loop back to main	2	JR main_loop
Total framework	~91-134	Before any effect-specific code

That leaves 122-165 bytes for the actual creative content – the visual formula, data tables, extra sound logic, text strings, or anything else that makes the intro *yours*. The framework is expensive. This is why size-coders fight so hard for every byte in the scaffolding: each byte saved in the framework is a byte gained for art.

Look at the companion example `intro256.a80`. Its pixel fill loop uses 18 bytes. The AY setup takes 20 bytes. The main loop framework (HALT, frame counter read, border update) is 8 bytes. The AY tone update is 13 bytes. The visual effect – a Moire interference pattern computed purely from register arithmetic – consumes 36 bytes. The frame counter writeback and loop jump take 8 bytes. Total: around 103 bytes of framework and 36 bytes of effect. That ratio – roughly 3:1 framework to effect – is typical. The better you compress the framework, the more room you have for creative expression.

Output of a 256-byte intro – animated Moire interference pattern with colour cycling, generated entirely from register arithmetic

Figure 30: Output of a 256-byte intro – animated Moire interference pattern with colour cycling, generated entirely from register arithmetic

Key Techniques at 256 Bytes

1. Use initial register and memory state. After a standard tape load, registers hold known values: A often holds the last byte loaded, BC the block length, HL points near the end of loaded data. The system variables area (\$5C00-\$5CB5) contains known values. Screen memory is clear after CLS. Every known value you exploit instead of loading explicitly saves 1-3 bytes.

2. Overlap code and data. The byte \$3E is the opcode for LD A, n and also the value 62 – an ASCII character, a screen coordinate, or an AY register value. If your program executes this byte as an instruction *and* reads it as data from a different code path, you have made one byte do two jobs. Common pattern: the immediate operand of LD A, n doubles as data that another routine reads with LD A, (addr) pointing at instruction_address + 1.

3. Choose instructions for size.

Large encoding	Small encoding	Savings
CALL \$0010 (3 bytes)	RST \$10 (1 byte)	2 bytes
JP label (3 bytes)	JR label (2 bytes)	1 byte
LD A, 0 (2 bytes)	XOR A (1 byte)	1 byte
CP 0 (2 bytes)	OR A (1 byte)	1 byte

The RST instructions are critical. RST n is a 1-byte CALL to one of eight addresses (\$00, \$08, \$10, \$18, \$20, \$28, \$30, \$38). On the Spectrum, RST \$10 calls the ROM character output, RST \$28 enters the calculator. In a normal demo these ROM routines are too slow. At 256 bytes, saving 2 bytes per CALL is everything.

Every JP in a 256-byte intro should be a JR – the whole program fits within the -128..+127 range.

4. Self-modifying code to reuse sequences. Need a subroutine to operate on two different addresses? Hardcode the first and patch the operand for the second call. Cheaper than parameter passing.

5. Mathematical relationships between constants. If your music needs tone period 200 and your effect needs loop count 200, use the same register. If one value is twice another, use ADD A, A (1 byte) instead of loading a second constant (2 bytes).

13.3 Famous 256-Byte Intros: What Made Them Clever

The ZX Spectrum 256-byte category has a rich history. Studying winning entries reveals what kinds of effects fit into 256 bytes and which creative strategies succeed.

Attribute-based effects dominate. The reason is arithmetic: the Spectrum's attribute area is 768 bytes (32 x 24), and you can fill it with a computed pattern using a tight nested loop of 15-20 bytes. Pixel-level effects require addressing 6,144 bytes of interleaved screen memory – much more code for address calculation alone. At 256 bytes, you simply cannot afford the overhead. So the vast majority of 256-byte intros work in attribute space: colour plasmas, interference patterns,

gradient animations, colour cycling. The pixel memory either stays blank, gets a one-time dither fill, or is left with whatever the ROM puts there.

Generative sound beats sequenced sound. A note table for a melody costs bytes – even a simple 8-note sequence is 8 bytes, plus the indexing logic. At 256 bytes, the winning strategy is to derive sound from the effect state. Use the frame counter as a tone period (pitch sweeps continuously). Use a byte from the visual computation as a noise parameter. Or use LD A, R – read the Z80’s refresh register, which increments on every instruction fetch – as a pseudo-random source, then mask it to a pentatonic range. The sound will not be a composition, but it will be *present*, and the audience will remember “that tiny intro that had music.”

The ROM is your library. Every byte of the Spectrum’s 16K ROM is available and does not count against your size limit. RST \$10 prints characters using the ROM’s full font rendering – 96 printable characters, 8x8 pixels each, with cursor management. That is thousands of bytes of rendering code available for 1 byte per call. RST \$28 accesses the floating-point calculator, which can compute sine, cosine, and square roots – operations that would cost dozens of bytes to implement. The cost is speed (the ROM routines are slow), but in a 256-byte intro running at 50fps with a simple effect, you often have cycles to spare.

The entries that win are the ones that look impossible at their size. Judges and audiences react to the gap between perceived complexity and file size. A 256-byte intro with a smooth colour plasma and a recognisable melody generates more applause than one with a slightly better visual but no sound. The trick is choosing an effect that *looks* complex but *encodes* cheaply. XOR-based interference patterns are perfect: visually intricate, mathematically trivial. Colour cycling through attributes is another: the eye perceives motion and depth, but the code is just incrementing values in a loop. Diagonal scrolling patterns, checkerboard animations, expanding rings – all can be produced with fewer than 20 bytes of inner-loop code if the formula is chosen carefully.

13.4 The LPRINT Trick

In 2015, diver4d published “Secrets of LPRINT” on Hype, documenting a technique older than the demoscene itself – one that first appeared in pirated cassette software loaders in the 1980s.

How It Works

The system variable at address 23681 (\$5C81) controls where BASIC’s output routines direct data. Normally it points to the printer buffer. Modify it to point to screen memory, and LPRINT writes directly to the screen:

That single POKE redirects the printer channel to \$4000 – the start of screen memory.

The Transposition Effect

The visual result is not just text on screen – it is *transposed* text. The Spectrum’s screen memory is interleaved (Chapter 2), but the printer driver writes sequentially. Data lands in screen memory according to the driver’s linear logic but *displays* according to the interleaved layout. The result cycles through 8 visual states as it progresses through the screen thirds – a cascade of data that builds in horizontal bands, shifting and recombining.

With different character data – graphical characters, UDGs, or carefully chosen ASCII sequences – the transposition produces striking visual patterns. The LPRINT statement handles all screen addressing, character rendering, and cursor advancement. Your program provides only the data.

From Pirate Loaders to Demo Art

diver4d traced the trick to pirated cassette loaders. Pirates adding custom loading screens needed visual effects in very few BASIC bytes – LPRINT was ideal. The technique fell out of use as the scene moved to machine code.

But in 2011, JtN and 4D released **BBB**, a demo that deliberately returned to LPRINT as an artistic statement. The old pirate-loader trick, framed with intention, became demo art. The constraint – BASIC, a printer redirect hack, no machine code – became the medium.

Why It Matters for Size-Coding

LPRINT achieves complex screen output for almost zero bytes of your own code. The ROM does the heavy lifting. Your contribution: a POKE to redirect output, data to print, and RST \$10 (or LPRINT) to trigger it. You leverage the Spectrum’s 16K ROM as a “free” screen-output engine – code that does not count against your size limit.

13.5 512-Byte Intros: Room to Breathe

Doubling from 256 to 512 bytes is not twice as much – it is qualitatively different. At 256, you fight for every instruction and sound is minimal. At 512, you can have a proper effect *and* proper sound, or two effects with a transition.

What Each Size Tier Enables

The jump between size categories is not linear. Each doubling opens qualitative new possibilities:

256 bytes is one effect and maybe primitive sound. You cannot afford a data table longer than about 16 bytes. Every variable lives in a register or in the instruction stream (self-modifying code). Text output is limited to a few characters. You have room for one nested loop with 2-3 arithmetic operations in the inner body. The visual will be attribute-based, generated purely from arithmetic. Sound, if present, is a tone sweep or random notes.

512 bytes lets you add a sine table (32-64 bytes), a real AY music engine (melody + bass on two channels), or a second visual effect with a transition. You can afford a proper frame-counted state machine that switches between two parts. Self-modifying code becomes structural rather than desperate. You might even have room for a short text string (10-20 characters) displayed with RST \$10.

1K (1,024 bytes) is a different world. You can have a tracker-style music player with a compressed pattern (one channel with a 32-step loop takes about 80-120 bytes including the player). Multiple effects with transitions become standard. Pixel-level effects – simple plasma in pixel space, scrolling text, raster bars – become feasible because you can afford the screen memory address calculation. You can include a 256-byte sine table, or generate one at startup and keep it in a buffer. At 1K, the constraint still shapes every decision, but the decisions are about *which features to include*, not about *which instructions you can afford*.

4K and 8K intros approach the territory of short demos. At 4K, compression becomes viable and you can fit multi-effect compositions with music – a qualitative leap covered in Section 13.6. An 8K intro is a polished mini-demo where the constraint is more about data compression than instruction-level size tricks. The techniques from this chapter still apply, but the focus shifts from “can I save one byte?” to “can I compress this data stream?”

The sweet spot for learning size-coding is 256 bytes. At that size, every technique in this chapter is mandatory. At 512, you have enough room to choose. At 1K, the size-coding mindset helps but does not dominate.

Common 512-Byte Patterns

Plasma via sine table sums. The sine table is the expensive part. A full 256-byte table consumes half your budget. Solutions: a 64-entry quarter-wave table mirrored at runtime (saves 192 bytes), or generate the table at startup using the parabolic approximation from Chapter 4 (~20 bytes of code instead of 256 bytes of data).

Tunnel via angle/distance lookup. At 512 bytes, you compute angle and distance on the fly using rough approximations. Lower visual quality than the Eager tunnel (Chapter 9), but recognisably a tunnel.

Fire via cellular automaton. Each cell averages its neighbours below, minus decay. A few instructions per pixel, convincing animation, and at 512 bytes you can add attributes for colour *and* a beeper sound.

Self-Modifying Tricks

Self-modification becomes structural at 512 bytes. Embed the frame counter *inside* an instruction:

```
frame_ld:
    ld    a, 0                ; this 0 is the frame counter
    inc  a
    ld    (frame_ld + 1), a    ; update the counter in place
```

No separate variable. The counter lives in the instruction stream.

Patch jump offsets to switch between effects:

```

effect_jump:
    jr    effect_1          ; this offset gets patched
    ; ...
effect_1:
    ; render effect 1, then:
    ld    a, effect_2 - effect_jump - 2
    ld    (effect_jump + 1), a    ; next frame jumps to effect 2

```

The ORG Trick

Choose your program's ORG address so that address bytes themselves are useful data. Place code at \$4000 and every JR/DJNZ targeting labels near the start generates small offset bytes – usable as loop counters, colour values, or AY register numbers. If your effect needs \$40 (the high byte of screen memory) as a constant, place code at an address where \$40 appears naturally in an address operand. The *encoding of the code itself* provides data you need elsewhere.

This is the deepest level of the size-coding puzzle.

13.6 4K Intros: The Mini-Demo

4096 bytes is where size-coding transitions from “one trick” to “mini-demo.” At 256 bytes, you have room for a single effect and maybe primitive sound. At 512 or 1K, you can have a proper effect with music. At 4K, you can have multiple effects, transitions between them, a full soundtrack, and a coherent narrative arc. The difference between 1K and 4K is qualitative, not just quantitative – it is the difference between “clever trick” and “tiny production.”

Compression Becomes Viable

The single biggest change at 4K is that data compression pays for itself. A good Z80 decompressor – ZX0, Exomizer, or similar – costs roughly 150-200 bytes of code. At 256 or 512 bytes, that overhead is catastrophic. At 4K, it is less than 5% of your budget, and the return is enormous: a 4K intro might contain 6-8K of uncompressed code and data, packed down to fit the limit. Your actual working space nearly doubles.

The workflow becomes a feedback loop: write code, assemble to a raw binary, compress with ZX0, check the output size, iterate. The number that matters is no longer the assembled size – it is the *compressed* size. This changes your optimisation strategy. You are no longer counting individual instruction bytes. You are thinking about what compresses well.

Code with repetitive patterns compresses better than code with high entropy. A table of sine values compresses well (smooth, predictable). A table of random bytes does not. Effect code that reuses similar instruction sequences across routines compresses better than code where every routine has a unique structure. This is a subtle shift: you optimise not just for *small code* but for *compressible code*.

Music Fits

At 256 bytes, sound is a luxury – a tone sweep or random pentatonic notes. At 4K, you can have a real soundtrack. A tiny AY player engine – something like Beepola’s output or a custom minimal tracker – takes 200-400 bytes. Add 500-1000 bytes of pattern data (compressed) and you have a full three-channel AY composition with melody, bass, and drums. These numbers compress well because music pattern data is highly repetitive.

The impact on the audience is disproportionate. Sound transforms a size-coding entry from a visual curiosity into an *experience*. At compo screenings, intros with music score dramatically higher than silent ones of equal visual quality. If you have 4K to work with and you are not including music, you are leaving points on the table.

Multi-Effect Structure

Unlike 256 bytes where you are locked into a single visual, 4K gives you room for 2-4 distinct effects with transitions. The structural framework is lightweight: a scene table mapping effect pointers to durations costs perhaps 30 bytes:

```
scene_table:
    DW effect_plasma      ; pointer to effect routine
    DB 150                ; duration in frames (3 seconds at 50fps)
    DW effect_tunnel
    DB 200
    DW effect_scroller
    DB 250
    DB 0                  ; end marker

scene_runner:
    ld  hl, scene_table
.next_scene:
    ld  e, (hl)
    inc hl
    ld  d, (hl)          ; DE = effect routine address
    inc hl
    ld  a, (hl)          ; A = duration
    or  a
    ret z                ; end of table
    inc hl
    push hl              ; save table pointer
    ld  b, a             ; B = frame counter
.frame_loop:
    push bc
    push de
    call .call_effect
    pop de
    pop bc
    halt                 ; wait for vsync
    djnz .frame_loop
    pop hl
    jr  .next_scene
```

```
.call_effect:
    push de
    ret                ; jump to DE via push+ret trick
```

Each individual effect might run 500-1000 bytes of code. At 4K compressed, you can afford three substantial effects, a scene table, a music player, and transition logic (fade to black between scenes is cheap – just zero the attribute area).

GOA4K, inal, and Megademica

GOA4K by Exploder^XTM is a landmark ZX Spectrum 128K 4K intro that demonstrates what is achievable when compression meets clever coding. It packs a chunk rotozoomer and other effects into 4096 bytes – visuals that would be respectable in a full-size demo, compressed down to a size you could fit in a single disk sector.

The story does not end there. **SerzhSoft** took GOA4K and remade it as **inal** – a 48K-only version in just 2980 bytes. The same visual impact, on a more constrained machine, in fewer bytes. This is the size-coding community at work: one coder sets a bar, another clears it from a harder starting position.

SerzhSoft went on to win the 4K intro compo at **Revision 2019** with **Megademica** – competing not in a ZX-specific category, but against all platforms at the world’s largest demoscene event. A ZX Spectrum 4K intro, judged alongside PC and Amiga entries, took first place. This is the trajectory that 4K size-coding enables: from local scene technique to global recognition.

Studying entries like these reveals a pattern: the best 4K intros choose effects that are visually impressive *and* compress well, then squeeze every byte through a tight pack-test-iterate cycle.

The 4K Tradeoffs

Working at 4K introduces tradeoffs that do not exist at smaller sizes:

Compression ratio drives effect choice. Not all effects compress equally. A plasma that relies on a smooth sine table compresses beautifully – the table data is predictable, and the rendering loop reuses similar instruction patterns. A pseudo-random dithering effect where every pixel is computed from a different formula produces high-entropy code that barely compresses at all. At 4K, you choose effects partly on their visual merit and partly on how well their implementation packs.

Boot time is visible. Decompression takes real time – typically 1-3 seconds on a 3.5MHz Z80 for a few kilobytes of data. The audience sees a pause before the intro starts. Most 4K intros mask this with a simple loading effect: fill the border with colour cycling, draw a quick pattern in attributes, or flash a single-frame title screen. The decompressor itself runs from a small uncompressed stub at the start of the file. Once decompression finishes, the stub jumps to the unpacked code and the real show begins.

You optimise for packed size, not runtime speed. In a 256-byte intro, the same code that runs is the code you measure. At 4K, you write code that decompresses into RAM and then executes from there. ROM constraints disappear – your unpacked code sits in free RAM. But the optimisation target shifts: you care about how many bytes the packed binary occupies, not the raw assembled size. An

effect that assembles to 900 bytes but compresses to 400 is better than one that assembles to 600 but compresses to 500.

Counting packed bytes. The build process gains a compression step. Assemble to binary, compress with ZX0 (or your compressor of choice), check the output file size. With sjasmplus:

```
zx0 build/intro4k.bin build/intro4k.zx0
ls -l build/intro4k.zx0    # this is the number that must be <= 4096
```

The decompressor stub prepended to the final file must also fit within the 4096-byte limit. Total file = decompressor stub + compressed payload. A typical ZX0 decompressor is about 70 bytes in its smallest form, leaving roughly 4026 bytes for compressed data.

Competition Categories

Demoscene parties offer various size-limited categories beyond the classic 256. Common competition tiers include 4K, 8K, and sometimes 16K, alongside the smaller 256 and 512. The specific categories vary by party – Chaos Constructions, DiHalt, and Forever have all hosted 4K compos for the Spectrum. Some parties combine platforms (a “4K intro” compo accepting entries for any 8-bit platform), while others are Spectrum-specific. Check the party rules before starting – the measurement method (raw file size vs. loaded memory image) and the exact byte limit matter.

At 8K and 16K, the approach is essentially the same as 4K but with more breathing room. An 8K intro is a polished mini-demo where the compression pipeline is standard and the creative challenge is more about art direction than byte-counting. At 16K, you are essentially making a short demo that happens to fit in 16K – the size constraint shapes your ambition but does not dictate your instruction choices. The size-coding techniques from this chapter still help at these larger budgets, but their impact is proportionally smaller.

13.7 Practical: Writing a 256-Byte Intro Step by Step

Start with a working attribute plasma (~400 bytes) and optimise it to 256.

Step 1: The Unoptimised Version

A simple attribute plasma: fill 768 bytes of attribute memory with values from sine sums, offset by a frame counter. Sound: a cycling melody on AY channel A. This version is clean, readable, and roughly 400 bytes – the sine table (32 bytes), note table (16 bytes), inline AY writes, and the plasma loop with table lookups.

Step 2: Replace CALL with RST

Any call to a ROM address matching an RST vector saves 2 bytes per invocation. For AY output, replace the six verbose inline register writes (~60 bytes) with a small subroutine:

```

ay_write:                                ; register in A, value in E
    ld    bc, $FFFD
    out   (c), a
    ld    b, $BF
    out   (c), e
    ret                                ; 8 bytes total

```

Six calls (5 bytes each: load A + load E + CALL) = 30 + 8 = 38 bytes. Savings: ~22 bytes.

Step 3: Overlap Data with Code

The 32-byte sine table at the entry point decodes as mostly harmless Z80 instructions (\$00=NOP, \$06=LD B,n, \$0C=INC C...). Place it at the program start. On first execution, the CPU stumbles through these “instructions,” scrambling some registers. The main loop then jumps past the table and never executes it again – but the data remains for lookups. The table bytes serve double duty.

Step 4: Exploit Register State

After the plasma loop writes 768 attributes, HL = \$5B00 and BC = 0 (from any LDIR used in initialisation). If the next operation needs these values, skip the explicit loads. Art-Top’s discovery in NHBF was exactly this: register values from screen clearing matched the text string length. Not planned. Noticed.

After each optimisation pass, annotate what every register contains at every point. Register state is a shared resource – the fundamental currency of size-coding.

Step 5: Smaller Encodings Everywhere

- LD A, 0 -> XOR A (save 1 byte)
- LD HL, nn + LD A, (HL) -> LD A, (nn) (save 1 byte if HL not needed)
- JP -> JR everywhere (save 1 byte each)
- CALL sub : ... : RET -> fall through directly (save 4 bytes)
- PUSH AF for temporary saves vs LD (var), A (save 2 bytes)

Step 6: Counting Bytes Precisely

Intuition about “how big is this” is unreliable. You need to count. There are three methods, and serious size-coders use all three.

Assembler output. sjasmplus can report the assembled size. The DISPLAY directive prints to the console during assembly, and ASSERT enforces the limit:

```

intro_end:
    ASSERT intro_end - init <= 256, "Intro exceeds 256 bytes!"
    DISPLAY "Intro size: ", /D, intro_end - init, " bytes"

```

Run the assembler after every change. The DISPLAY line tells you where you stand; the ASSERT catches overflows before you waste time testing a broken binary.

Symbol file analysis. Assemble with --sym=build/intro.sym to get a symbol table. Compare label addresses to find exactly how many bytes each section occupies. When your intro is 262 bytes and you need to cut 6, the symbol file tells you that

the AY init is 22 bytes (can you cut 2?), the effect loop is 38 bytes (can you merge the row and column counters?), the frame counter writeback is 8 bytes (can you restructure to make it 5?). Without this breakdown, you are guessing.

Hex dump inspection. After assembling, examine the raw binary in a hex editor (or `xxd build/intro.bin`). The hex dump shows you the actual bytes the CPU will execute. You will spot redundancies invisible in the source: two consecutive loads that could be one, an opcode whose value happens to match data you need elsewhere, a sequence of NOPs left by an accidental alignment. The hex dump is the ground truth. The source is an abstraction over it.

The Final Push

The last 10-20 bytes are the hardest. Structural rearrangement: reorder code so fall-throughs eliminate JR instructions. Merge the sound and visual loops. Embed data bytes in the instruction stream – if you need `$07` as data and also need an RLCA (opcode `$07`), arrange for one to serve as both.

At this stage, keep a log. Write down every change you try: “moved AY init before pixel fill: saved 2 bytes (C register reuse), lost 1 byte (need extra LD B). Net: +1 byte.” Many changes do not help. Some make things worse. Without a log, you will try the same dead-end twice. With a log, you build a map of the solution space.

Try radical restructuring. Can the visual effect loop also update the AY? If the inner loop iterates 768 times (once per attribute cell), and you write a new tone value every 32 iterations (once per row), the sound update happens inside the visual loop at the cost of one BIT 4, E / JR NZ check – 4 bytes to merge two routines that previously needed separate framework code. Sometimes merging saves 10 bytes; sometimes it costs 5. You will not know until you try.

The escape hatch: choose a different effect. If your plasma needs a sine table and you are 30 bytes over, no amount of micro-optimisation will save you. Switch to an effect that generates its visual from pure register arithmetic: XOR patterns, modular arithmetic, bit manipulation. An XOR interference pattern like the one in `intro256.a80` needs zero data bytes. The visual is less smooth than a sine plasma, but it fits. At 256 bytes, “fits” is the only criterion that matters.

You stare at the hex dump. You try moving the sound routine before the visual routine. You try replacing the sine table with a runtime generator. Each attempt reshuffles the bytes. Sometimes everything lines up.

The satisfaction of fitting a coherent audiovisual experience into 256 bytes – of solving the puzzle – is real and specific and unlike any other feeling in programming.

13.8 Size-Coding Music: Bytebeat on AY

In the PC demoscene, **bytebeat** is a formula-driven approach to sound: a single expression like `t*( (t>>12|t>>8)&63&t>>4)` generates PCM samples, producing surprisingly complex music from a few bytes of code. The concept was popularised by Viznut (Ville-Matias Heikkilä) in 2011, and 256-byte PC intros routinely use bytebeat for their soundtracks.

On the ZX Spectrum, the situation is different. The AY-3-8910 is not a DAC – it is a tone and noise generator with per-channel period and volume registers. You cannot feed it PCM samples in the traditional sense (volume-register sample playback exists but costs too many cycles for a size-coded intro). Instead, “AY bytebeat” means computing **tone periods and volume envelopes from mathematical formulas** driven by a frame counter.

The principle is the same as PC bytebeat: replace stored music data with a formula. The output target is different.

The Minimal AY Formula Engine

A typical approach in a 256-byte intro:

```
; Frame-driven AY "bytebeat" – ~20 bytes
; A = frame counter (incremented each HALT)
    ld    e, a
    and   $1F                ; period = low 5 bits of frame
    ld    d, a                ; D = tone period low
    ld    a, e
    rrca
    rrca
    rrca
    and   $0F                ; volume = bits 5-7 of frame, shifted
    ; Write to AY: register 0 = tone period low, register 8 = volume
```

This produces a cycling tone that sweeps through periods and fades in/out – not music in any traditional sense, but recognisably structured sound. The trick is choosing formulas that produce **musically interesting patterns** from simple bitwise operations.

Techniques for Better-Sounding Formulas

Pentatonic masking. Raw bitwise formulas produce chromatic noise. Mask the period value through a pentatonic lookup (5 bytes: the note intervals) to constrain output to a pleasant scale. Five bytes of data buys musically coherent sound.

Multi-channel formulas. The AY has three tone channels. Use different bit rotations of the same frame counter for each channel – they will produce related but distinct patterns, creating an impression of harmony:

```
    ld    a, (frame)
    call  .write_ch_a        ; channel A: raw formula
    ld    a, (frame)
    rrca
    rrca                    ; channel B: frame >> 2
    call  .write_ch_b
    ld    a, (frame)
    add   a, a               ; channel C: frame << 1
    call  .write_ch_c
```

Noise percussion. Toggle the noise generator on specific frame intervals (every 8th or 16th frame) for a rhythmic pulse. Cost: one AND + one OUT — about 6 bytes for a basic kick pattern.

LD A,R as entropy. The R register (memory refresh counter) is effectively random from a musical perspective. Mix it with the frame counter: `ld a,r : xor (frame)` produces evolving textures that never quite repeat. Useful for ambient or experimental soundscapes.

Bytebeat vs. Sequenced Music

	Bytebeat (formula)	Sequenced (pattern data)
Bytes	10-30 (code only)	200-400 (player) + 500+ (patterns)
Musical quality	Abstract, generative, alien	Melodic, structured, human
Best at	256b, 512b	1K, 4K
Sound	Rhythmic noise, sweeps, drones	Actual tunes

At 256 bytes, bytebeat is your only realistic option – there is no room for a pattern player. At 512, you can afford a tiny sequencer with 4-8 notes. At 4K, use a real player. The bytebeat approach is not inferior – it produces a *different kind* of sound that fits the aesthetic of tiny programs. Some of the most memorable 256-byte intros are memorable precisely because their sound is alien and generative, not because it imitates conventional music.

13.9 Procedural Graphics: The Rendered GFX Compo

Some demoscene parties run a **rendered graphics** (or **procedural graphics**) competition: submit a program that generates a static image. No pre-drawn bitmaps, no loaded data – every pixel must be computed. The visual output is judged as artwork, but it must be born from code.

On the Spectrum, this means your program fills the 6,912-byte screen area (bitmap + attributes) algorithmically, then halts. The image stays on screen for judging. File size limits vary – some compos allow any size, others impose 256-byte or 4K limits, turning it into a hybrid of size-coding and digital art.

Why the Spectrum Is Interesting for This

The Spectrum's display constraints – 1-bit pixels with 8×8 attribute colour – make procedural graphics a genuinely different challenge from doing it on a 256-colour VGA or 24-bit framebuffer. You cannot just compute RGB values per pixel. You must think in terms of:

- **Pixel patterns** within 8×8 character cells (dithering, halftone)
- **Attribute colour** per cell (2 colours from a palette of 15)
- **The interaction** between pixel pattern and attribute – a gradient needs both smooth dithering AND smooth attribute transitions

This constraint creates a distinctive visual style. Procedural Spectrum graphics look like nothing else – the colour grid gives them a mosaic quality that is part of the aesthetic, not a limitation to hide.

Common Approaches

Mandelbrot and Julia sets. The classic choice. The iteration loop is compact (~30-50 bytes for the core), and the fractal detail is infinite – zoom coordinates and iteration count are the only parameters. Map iteration count to dither pattern for pixel data, map to palette index for attributes. A Mandelbrot renderer fits comfortably in 256 bytes and produces images that look hand-crafted.

Interference patterns. Multiple overlapping sine or cosine waves, sampled at each pixel position. $\text{pixel} = \sin(x * \text{freq1} + \text{phase1}) + \sin(y * \text{freq2} + \text{phase2}) > \text{threshold}$. Produces organic, flowing shapes. On the Spectrum, threshold the sum to get the pixel bit, quantise to get the attribute colour.

Distance fields. Compute the distance from each pixel to a set of shapes (circles, lines, Bézier curves). Threshold the distance for pixel data, map it to colour for attributes. A few shapes can produce surprisingly complex images – overlapping circles alone can create intricate patterns.

L-systems and fractals. Recursive branching structures (trees, ferns, Sierpinski triangles). The recursion maps naturally to stack-based Z80 code, and the visual output has organic complexity from minimal code. A Sierpinski triangle renderer is about 20 bytes; a branching tree with random angles is perhaps 80.

The Byte Budget for Art

In a size-limited rendered GFX compo, every byte goes toward visual complexity. There is no frame loop, no sound, no animation – just a straight-line program that fills the screen and stops. This means your full budget goes to rendering code and coordinate generation. At 256 bytes, you can produce a detailed fractal. At 4K (compressed), you can generate images with multiple layers, computed textures, and careful dithering that approach hand-drawn quality.

The judging criterion is purely visual – the audience votes on the image, not the code. But the code constraint shapes the aesthetic. Procedural Spectrum graphics have a recognisable look: mathematical precision, fractal detail, and the characteristic colour grid of attribute-based rendering. The best entries embrace these constraints as style rather than fighting them.

13.10 Size-Coding as Art

Size-coding teaches you things that improve all your coding: the discipline of questioning every byte sharpens instruction-encoding awareness, the habit of looking for overlaps transfers to any optimisation work, and the practice of exploiting initial state and side effects makes you a better systems programmer.

Summary

- **Size-coding** competitions require complete programs in 256, 512, 1K, 4K, or 8K bytes – strict limits that demand a fundamentally different approach to programming.
 - **The size-coder's toolkit** includes register initialization assumptions, DJNZ as a combined decrement-and-branch, RST as a 1-byte CALL, overlapping instructions, and flag abuse via SBC A,A – tricks that save 1-5 bytes each but compound across a program.
 - **NHBF** (UriS, CC 2025) demonstrates the 256-byte mindset: every byte does double duty, register states from one routine feed into the next, instruction choice is driven purely by encoding size.
 - **The byte budget** for a typical 256-byte intro allocates ~90-130 bytes to framework (screen fill, AY init, frame sync, loop structure), leaving 120-160 bytes for the actual creative effect.
 - **Choosing the right effect** matters more than micro-optimisation: attribute-based visuals with arithmetic formulas (XOR, modular math) encode cheaply; pixel-level effects and data tables consume too many bytes at 256.
 - **The LPRINT trick** (diver4d, 2015) redirects BASIC's printer output to screen memory via address 23681, producing complex visual patterns in a handful of bytes – from pirated cassette loaders to demo art.
 - **Each size tier is qualitatively different:** 256 bytes allows one effect with minimal sound; 512 adds sine tables and two-channel music; 1K enables pixel-level effects, tracker music, and multiple parts; 4K crosses the threshold into mini-demo territory with compression, full soundtracks, and multi-effect compositions.
 - **4K intros** are where compression becomes viable: a ~200-byte decompressor unlocks 6-8K of working space, music players with pattern data fit comfortably, and scene tables enable 2-4 distinct effects with transitions. The optimisation target shifts from raw assembled size to compressed packed size.
 - **AY bytebeat** replaces stored music data with formulas: compute tone periods and volumes from the frame counter using bitwise arithmetic. At 256 bytes, formula-driven sound (10-30 bytes) is the only option; at 4K, switch to a real pattern player. Pentatonic masking, multi-channel bit rotation, and noise percussion add musicality for minimal bytes.
 - **Procedural graphics** (rendered GFX) competitions require every pixel to be computed, not loaded. The Spectrum's 1-bit pixels with 8×8 attribute colour make this a unique challenge – Mandelbrot sets, interference patterns, distance fields, and L-systems all produce distinctive results within the attribute grid aesthetic.
 - **The optimisation process** moves from structural changes (eliminating tables, merging loops) to encoding choices (RST for CALL, JR for JP, XOR A for LD A,0) to serendipitous discoveries (register states aligning with data needs).
 - **Counting bytes precisely** – via assembler DISPLAY/ASSERT, symbol file analysis, and hex dump inspection – is essential. Intuition about code size is unreliable.
 - **The ORG trick** – choosing your load address so that address bytes double as useful data – represents the deepest level of the puzzle.
-

Try It Yourself

1. **Start large, shrink small.** Write an attribute plasma with a frame counter. Get it working at any size. Then optimise to 512 bytes, tracking every byte saved and how.
2. **Explore LPRINT.** In BASIC, try `POKE 23681,64 : FOR i=1 TO 500 : LPRINT CHR$(RND*96+32); : NEXT i`. Watch the transposed data fill the screen. Experiment with different character ranges.
3. **Map your register state.** Write a small program and annotate what every register contains at every point. Look for places where one routine's output matches another's needed input.
4. **Study the RST vectors.** Disassemble the Spectrum ROM at \$0000, \$0008, \$0010, \$0018, \$0020, \$0028, \$0030, \$0038. These are your "free" subroutines.
5. **The 256-byte challenge.** Push the practical from this chapter to 256 bytes. You will need to make hard choices about what to keep and what to cut. That is the point.

Next: Chapter 14 – Compression: More Data in Less Space. We move from programs that fit in 256 bytes to the problem of fitting kilobytes of data into kilobytes of storage, with Introspec's comprehensive benchmark of 10 compressors as our guide.

Sources: UriS "NHBF Making-of" (Hype, 2025); diver4d "LPRINT Secrets" (Hype, 2015)

Chapter 14: Compression — More Data in Less Space

The ZX Spectrum 128K has 128 kilobytes of RAM. That sounds generous until you start subtracting: the screen takes 6,912 bytes (6,144 pixels + 768 attributes), the system variables claim their share, the AY music player and its pattern data want a bank or two, your code occupies another few thousand bytes, and the stack needs room to breathe. By the time you sit down to store the actual content of your demo — the graphics, the animation frames, the precalculated lookup tables — you are fighting for every byte.

A single full-screen image on the Spectrum is 6,912 bytes. A 4K intro can fit roughly 0.6 of one. A 48K demo could theoretically hold seven screens with nothing else. But demos are not slideshows. They have music. They have code. They have effects that demand tables of precalculated data. The question is not whether to compress — it is which compressor to use, and when.

This chapter is built around a benchmark. In 2017, Introspec (spke, Life on Mars) published “Data Compression for Modern Z80 Coding” on Hype — a meticulous comparison of ten compression tools tested against a carefully designed corpus. That article, with its 22,000 views and hundreds of comments, became the reference that ZX coders consult when choosing a compressor. We will walk through his results, understand the tradeoffs, and learn to pick the right tool for each job.

The Memory Problem

Let us be concrete about the constraints. Consider Break Space by Thesuper (Chaos Constructions 2016, 2nd place) — a demo with 19 scenes running on a ZX Spectrum 128K. One of those scenes, the Magen Fractal by psndcj, displays 122 frames of animation. Each frame is a full 6,912-byte screen. Uncompressed, that is 843,264 bytes — over six times the machine’s total RAM.

psndcj compressed all 122 frames into 10,512 bytes. That is 1.25% of the original size. The entire animation, every frame of it, fits in less space than two uncompressed screens.

Another scene in Break Space, the Mondrian animation, packs 256 hand-drawn frames — every square cut separately, individually compressed — into 3 kilobytes.

These are not theoretical exercises. They are production techniques from a demo that competed at one of the scene’s most prestigious parties. Compression is not

an optimisation you apply at the end. It is a fundamental architectural decision that determines what your demo can contain.

Compression as bandwidth amplifier

Introspec articulated the insight that elevates compression from a storage trick to a performance technique: **compression acts as a method to increase effective memory bandwidth.**

Suppose an effect needs 2 KB of data per frame. Store it compressed to 800 bytes and decompress with LZ4 at 34 T-states per output byte. The decompression costs 69,632 T-states — almost exactly one frame. But you can overlap it with border time, double-buffer a frame ahead, and interleave with effect rendering. The result: more data flowing through the system than the bus could deliver from uncompressed storage. The decompressor is a data amplifier.

The Benchmark

Introspec did not simply run each compressor on a few files and eyeball the results. He designed a corpus and measured systematically.

The corpus

The test data totalled 1,233,995 bytes across five categories:

- **Calgary corpus** — the standard academic compression benchmark (text, binary, mixed)
- **Canterbury corpus** — a more modern academic standard
- **30 ZX Spectrum graphics** — loading screens, multicolour images, game screens
- **24 music files** — PT3 patterns, AY register dumps, sample data
- **Miscellaneous ZX data** — tilemaps, lookup tables, mixed demo data

This mix matters. A compressor that excels on English text may struggle with ZX graphics, where long runs of zeros in the pixel area alternate with near-random attribute data. Testing on real Spectrum data — the data you will actually compress — is essential.

The results

Ten tools. Measured on total compressed size (lower is better), decompression speed in T-states per output byte (lower is faster), and decompressor code size in bytes (smaller is better for size-coded productions).

Tool	Compressed (bytes)	Ratio	Speed (T/byte)	Decompressor Size	Notes
Ex-omizer	596,161	48.3%	~250	~170 bytes	Best compression ratio
ApLib	606,833	49.2%	~105	199 bytes	Good balance
Pu-Crunch	616,855	50.0%	—	—	Complex LZ alternative
Hrust 1	613,602	49.7%	—	—	Relocatable stack de-packer
Pletter 5	635,797	51.5%	~69	~120 bytes	Fast + decent compression
MegaLZ	636,910	51.6%	~130	~110 bytes	Revived by Introspec in 2019
ZX7	653,879	53.0%	~107	69 bytes	Tiny decompressor
ZX0	—	~52%	~100	~70 bytes	Successor to ZX7
LZ4	722,522	58.6%	~34	~100 bytes	Fastest de-compression

Hrum	—	~52%	—	—	De- clared obso- lete
------	---	------	---	---	--------------------------------

Only Exomizer broke the 600,000-byte barrier across the full corpus. But Exomizer's decompression speed — roughly 250 T-states per output byte — makes it impractical for anything that needs to decompress during playback.

The tradeoff triangle

Every compressor makes a tradeoff between three qualities:

1. **Compression ratio** — how small the compressed data gets
2. **Decompression speed** — how many T-states per output byte
3. **Decompressor code size** — how many bytes the decompression routine occupies

You cannot have all three. Exomizer wins on ratio but is slow to decompress and has a large decompressor. LZ4 is the fastest to decompress but wastes 10 percentage points of compression ratio. ZX7 has a 69-byte decompressor but compresses less aggressively than Exomizer.

Introspec's genius was to map these tradeoffs on a Pareto frontier — a curve where no tool can improve on one dimension without losing on another. If a compressor is dominated on all three axes by another tool, it is obsolete. If it sits on the frontier, it is the right choice for some use case.

graph LR

```

SRC["Source Data<br>(raw bytes)"] --> EX0["Exomizer<br>48.3% ratio<br>~250
↳ T/byte<br>170B decompressor"]
SRC --> APL["ApLib<br>49.2% ratio<br>~105 T/byte<br>199B decompressor"]
SRC --> PLT["Pletter 5<br>51.5% ratio<br>~69 T/byte<br>~120B decompressor"]
SRC --> ZX0["ZX0<br>~52% ratio<br>~100 T/byte<br>~70B decompressor"]
SRC --> LZ4["LZ4<br>58.6% ratio<br>~34 T/byte<br>~100B decompressor"]

EX0 --> T1["Best ratio<br>Slowest decompression"]
LZ4 --> T2["Worst ratio<br>Fastest decompression"]
ZX0 --> T3["Smallest decompressor<br>Good all-around"]

style EX0 fill:#fdd,stroke:#333
style LZ4 fill:#ddf,stroke:#333
style ZX0 fill:#dfd,stroke:#333
style T1 fill:#fdd,stroke:#933
style T2 fill:#ddf,stroke:#339
style T3 fill:#dfd,stroke:#393

```

The tradeoff: Smaller compressed size = slower decompression. No compressor wins on all three axes (ratio, speed, decompressor size). Choose based on your use case: Exomizer for one-time loads, LZ4 for real-time streaming, ZX0 for size-coded intros.

His practical recommendations are crisp:

- **Maximum compression, speed irrelevant:** Exomizer. Use for one-time decompression at load time — loading screens, level data, anything you unpack once into a buffer and use repeatedly.
- **Good compression, moderate speed (~105 T/byte):** ApLib. A solid general-purpose choice when you need decent ratio and can afford ~105 T-states per byte.
- **Fast decompression (~69 T/byte):** Pletter 5. When you need to decompress during gameplay or between scenes and cannot afford Exomizer's sluggish unpacking.
- **Fastest decompression (~34 T/byte):** LZ4. The only choice for real-time streaming — decompressing data as you play it back. At 34 T-states per output byte, LZ4 can decompress over 2,000 bytes per frame. That is a 2 KB/frame data pipe.
- **Smallest decompressor (69-70 bytes):** ZX7 or ZX0. When the decompressor itself must be tiny — in 256-byte, 512-byte, or 1K intros where every byte of code counts.

Let these numbers guide your decisions. There is no universally “best” compressor. There is only the best compressor for your specific constraints.

How LZ Compression Works

All of the compressors in the table above belong to the Lempel-Ziv family. Understanding the core idea helps you predict which data compresses well and which does not.

LZ compression replaces repeated byte sequences with back-references. A match says: “copy N bytes from position P bytes back in the already-decoded stream.” The compressed stream alternates between **literals** (raw bytes with no useful match) and **matches** (offset + length pairs that reference earlier output).

LZ compression: a byte stream with literals and back-references. Repeated sequences are replaced by (offset, length) tokens, reducing the stream from 16 bytes to 10 tokens.

Figure 31: LZ compression: a byte stream with literals and back-references. Repeated sequences are replaced by (offset, length) tokens, reducing the stream from 16 bytes to 10 tokens.

The differences between compressors come down to encoding: how many bits for the offset, how many for the length, how to signal literal vs match. Exomizer uses sophisticated variable-length bit-level codes that compress tightly but require careful bit-extraction to decode — hence the ~250 T-states per byte. LZ4 uses byte-aligned tokens that the Z80 processes with simple shifts and masks — hence ~34 T-states per byte at the cost of 10 percentage points of ratio. ZX0 uses single-bit flags (0 = literal, 1 = match) with Elias interlaced codes for lengths, hitting a sweet spot between size and speed.

ZX Spectrum data compresses well because it has structure: large areas of identical bytes (black backgrounds, blank attributes), repeated patterns (tiles, fonts, UI), and correlated pixel data at regular offsets. Music compresses well too — PT3

patterns are full of repeated note sequences and empty rows. What compresses poorly: random data, already-compressed data, and very short files where encoding overhead exceeds savings.

ZX0 — The Size Coder’s Choice

ZX0, created by Einar Saukas, is the spiritual successor to ZX7 and has become the default compressor for modern ZX Spectrum development. It deserves special attention.

Why ZX0 exists

ZX7 was already remarkable: a 69-byte decompressor that achieved respectable compression ratios. But Saukas saw room for improvement. ZX0 uses an optimal parsing algorithm — it does not just find good matches, it finds the *best possible sequence* of matches and literals for the entire file. The result is compression ratios close to much larger compressors, with a decompressor that stays in the 70-byte range.

The decompressor

The Z80 decompressor for ZX0 is hand-optimised assembly, designed specifically for the Z80’s instruction set. It exploits the Z80’s flag register, its block transfer instructions, and the exact timing of conditional jumps to squeeze maximum functionality into minimum bytes. Here is the kind of code we are talking about:

```
; ZX0 decompressor – standard version
; HL = source (compressed data)
; DE = destination (output buffer)
; Uses: AF, BC, DE, HL
dxx0_standard:
    ld     bc, $ffff      ; initial offset = -1
    push  bc
    inc   bc              ; BC = 0 (literal length counter)
    ld    a, $80          ; bit buffer: only flag bit set
dxx0s_literals:
    call  dxx0s_elias      ; read literal length
    ldir                     ; copy literals
    add   a, a            ; read flag bit
    jr    c, dxx0s_new_offset
    call  dxx0s_elias      ; read match length
    ex    (sp), hl        ; retrieve offset from stack
    push  hl              ; put it back
    add   hl, de          ; calculate match address
    ldir                     ; copy match
    add   a, a            ; read flag bit
    jr    nc, dxx0s_literals
dxx0s_new_offset:
    ; ... offset decoding continues ...
```

Every instruction pulls double duty. The accumulator serves as both a bit buffer and a working register. The stack holds the last-used offset for repeated matches. The LDIR instruction handles both literal copying and match copying, keeping the code small.

At roughly 70 bytes, the entire decompressor fits in less space than a single ZX Spectrum character row. For a 256-byte intro, that leaves 186 bytes for everything else — the effect, the animation, the music. For a 4K intro, 70 bytes is negligible overhead. This is why ZX0 has become ubiquitous.

When to use ZX0

- **256-byte to 1K intros:** The tiny decompressor is essential. Every byte saved on the decompressor is a byte available for content.
- **4K intros:** ZX0 can decompress 4,096 bytes into 15–30 KB of code and data. SerzhSoft’s Megademica (1st place, Revision 2019) used this exact strategy to fit what reviewers called “a full new-school demo” into a 4K intro.
- **General demo and game development:** When you need a solid all-around compressor with a small footprint. ZX0 is not the fastest decompressor, but it is fast enough for one-time decompression at load time, and its ratio is competitive with tools that have much larger decompressors.
- **RED REDUX** (2025) used the newer ZX2 variant (also by Saukas) to achieve the remarkable feat of including Protracker music in a 256-byte intro.

ZX0 is not the right choice for real-time streaming (use LZ4) or for maximum compression at any cost (use Exomizer). But for the vast majority of ZX Spectrum projects, it is the correct default.

RLE and Delta Coding

Not everything needs a full LZ compressor. Two simpler techniques handle specific data types more effectively.

RLE: Run-Length Encoding

The simplest scheme: replace a run of identical bytes with a count and a value. The decompressor is trivial:

```
; Minimal RLE decompressor – HL = source, DE = destination
rle_decompress:
    ld     a, (hl)          ; read count
    inc    hl
    or     a
    ret    z                ; count = 0 means end
    ld     b, a
    ld     a, (hl)          ; read value
    inc    hl
.fill:   ld     (de), a
    inc    de
    djnz   .fill
    jr     rle_decompress
```

Only 12 bytes of decompressor code. RLE compresses beautifully when data contains long runs — blank screens, solid-colour backgrounds, attribute fills. It compresses terribly on complex pixel art. The advantage over LZ: for size-coded intros where even ZX0's 70 bytes feel expensive, a 12-byte RLE scheme frees precious space.

RLE also benefits from **data transposition**: if your data is a 2D block (e.g., 32×24 attributes) where columns are more uniform than rows, transposing to column-major order creates longer runs. The cost is an un-transpose pass after decompression (~13 T-states/byte). Whether the total (12-byte decompressor + un-transpose code + compressed data) beats ZX0 (70-byte decompressor + compressed data) depends on your data — measure both.

Sidebar: Ped7g's Self-Modifying RLE — 9 Bytes That Rewrite Themselves

For 256-byte intros, even 12 bytes feels expensive. Ped7g (Peter Helcmanovsky, sjasmpus maintainer) contributed a self-modifying RLE depacker that compresses the decoder itself to **9 bytes of core code** — and the exit mechanism is built into the data stream.

The trick: the RLE data lives in memory *before* the depacker code. The data stream ends with the bytes \$18, \$00, which the depacker writes into the target buffer at a calculated position so that the bytes overwrite the `ld (hl),c` instruction. The byte sequence \$18, \$23 assembles to `jr +$23`, which jumps forward past the depacker into the intro's main code. The data literally rewrites the code to terminate itself.

Here is the complete working mini-intro — a 120-byte binary that fills the screen with coloured stripes using only the self-modifying RLE:

```
““z80 id:ch14_ped7g_rle_mini_intro ; Ped7g's self-modifying RLE mini-
intro ; Assemble with sjasmpus: sjasmpus rle_intro.a80 ; ; The RLE data
is a stream of (value, count) pairs read via POP BC. ; SP walks through
the data as a read pointer. ; The db $18,$00 at the end of the data stream
overwrites ld (hl),c ; to become jr +$23, exiting the depack loop into
intro_start. ; ; Contributed by Ped7g (Peter Helcmanovsky) — sjasmpus
maintainer ; and ZX Spectrum Next contributor. Used with permission.
```

```
DEVICE ZXSPECTRUM48, $8000
```

```
target EQU $4000 ORG $5B00 ; loading address → print buffer
```

```
intro_data: dw target ; initial HL value (POP HL) ; RLE pairs: value, count
(count=0 means 256 iterations) .(43) db $AA, 0, $00, 0 ; alternating stripe
pattern db $43, 322, $44, 324, $45, 323, $46, 322, $47, 322 db $46, 322,
$45, 323, $44, 324, $43, 322 db $18, $00 ; data that will overwrite ld (hl),c
; creating jr rle_loop_inner+$25 rle_start: ei ; simulate post-LOAD BASIC
environment ld sp, intro_data pop hl ; HL = target address rle_loop_outer:
pop bc ; C = value, B = repeat count rle_loop_inner: ld (hl), c ; ← THIS
instruction gets overwritten inc hl ; by the $18,$00 data to become djnz
rle_loop_inner ; jr +$23, jumping to intro_start jr rle_loop_outer ; 31
bytes of space — fill with helper code ds $1F intro_start: assert $ ==
rle_loop_inner + 2 + $23 inc a and 7 out (254), a ; cycle border colours jr
intro_start
```

```

SAVESNA "rle_intro.sna", rle_start
SAVEBIN "rle_intro.bin", intro_data, $ - intro_data
""

```

Ped7g's 120-byte self-modifying RLE intro running on ZX Spectrum — coloured attribute stripes unpacked by a 9-byte depack loop, rainbow border from the cycling OUT (254),A

Figure 32: Ped7g's 120-byte self-modifying RLE intro running on ZX Spectrum — coloured attribute stripes unpacked by a 9-byte depack loop, rainbow border from the cycling OUT (254),A

Byte count analysis. The depack loop is 9 bytes: `pop bc (1) + ld (hl),c (1) + inc hl (1) + djnz (2) + jr (2) + pop hl (1) + ld sp,nn (3) = 9 core + 6 setup = 15 bytes total` for a self-contained RLE decoder with built-in exit. Compare to the 12-byte minimal RLE from the previous section, which still needs external setup and a termination check.

Interrupt safety. SP is used as a data pointer, so interrupts will corrupt the stack. The `ei` at the start is intentional — in a 256-byte intro loaded from BASIC, interrupts are already enabled. The occasional interrupt writes to already-consumed data behind the SP pointer, so the depack completes correctly. For the intro code itself, SP has moved past the data and the stack works normally. But do not combine this technique with IM2 or interrupt-driven music.

Advanced variants. Ped7g notes several alternative exit strategies: (1) if the target area extends behind the depack code, the RLE data can overwrite the `jr rle_loop_outer` offset to jump further; (2) the `jp $C3C3` trick — place `$C3` values in the data with exact counts so that `DJNZ` terminates when `jp $C3C3` assembles in memory, and align the intro so address `$C3C3` is the continuation code. As Ped7g says: “you can invent many such things — it always depends on the specific situation.”

Credit: Contributed by Ped7g (Peter Helcmanovsky) — `sjasmplus` maintainer and ZX Spectrum Next contributor. Used with permission.

Delta coding: store what changed

Delta coding stores differences between consecutive values rather than absolute values. Two animation frames that are 90% identical? Store only the changed bytes — a list of (position, new_value) pairs. If only 691 bytes differ out of 6,912, the delta is 2,073 bytes (3 bytes per change) instead of a full frame. Apply LZ on top of the delta stream and it compresses further — the difference stream has more zeros and repeated small values than raw frame data.

The Break Space Magen Fractal exploits this: 122 frames at 6,912 bytes each, compressed to 10,512 bytes total, because each frame differs from the previous by a small amount. Delta + LZ is the standard pipeline for multi-frame animations, scrolling tilemaps, and sprite animations where the figure changes pose but the background stays fixed.

Pre-Compression Data Preparation

Delta coding is not the only trick. The compressor only sees the byte stream you feed it. If you restructure the data before compression, the same LZ algorithm can achieve dramatically different ratios. This is the art of pre-compression preparation — and it is often more valuable than switching packers.

Entropy: the theoretical floor

Shannon entropy measures the minimum bits per byte needed to represent your data, assuming an ideal encoder. A completely random byte stream has entropy 8.0 bits/byte — incompressible. A file of identical bytes has entropy 0.0. Real Spectrum data falls somewhere between. A raw sine table might have entropy 6.75 bits/byte. Apply delta encoding, and it drops to 2.85. Apply the second derivative, and it falls to 1.49 — an 78% reduction. That is the theoretical room the compressor has to work with.

You do not need to compute entropy by hand. The formula is simple enough for a Python script:

```
from collections import Counter

def entropy(data: bytes) -> float:
    """Shannon entropy in bits per byte. Lower = more compressible."""
    counts = Counter(data)
    n = len(data)
    return -sum(c/n * math.log2(c/n) for c in counts.values())
```

Run this on your raw data, then on delta-encoded data, then on transposed data. The transform that gives the lowest entropy will compress best, regardless of which packer you use.

The second derivative: sinusoidal and quadratic data

Delta encoding stores first differences: $d[i] = \text{data}[i] - \text{data}[i-1]$. For a linear ramp (0, 3, 6, 9...), the delta stream is constant (3, 3, 3...) — perfect for compression. But sine waves and smooth curves produce a delta stream that itself varies smoothly. The second derivative (delta of delta) catches this:

Data type	Raw entropy	1st derivative	2nd derivative
Sine table (256B)	6.75	2.85	1.49
Linear ramp	7.00	0.00	0.00
Quadratic curve	6.80	3.20	0.00
Random bytes	8.00	8.00	8.00

The second derivative of a quadratic function is a constant. This is not abstract calculus — it is the difference between 6.80 and 0.00 bits per byte. A 256-byte quadratic lookup table, second-derivative encoded, compresses to almost nothing.

Here is the creative insight: sinusoidal decay and quadratic decay are often visually indistinguishable in a demo effect. If you are animating a particle that slows down, the audience cannot tell whether you used $\sin(t)$ or $at^2 + bt + c$. But

the compressor can: the quadratic version has a perfectly linear first derivative and a constant second derivative. If your animation can tolerate a quadratic approximation, you save bytes not by switching compressors, but by switching curves.

Transposition: column-major for tabular data

Demoscene data is often tabular — 3D vertex tables (X, Y, Z per vertex), animation keyframes (angle, radius, speed per frame), colour palettes (R, G, B per entry). When stored row-major ($X_0 Y_0 Z_0 X_1 Y_1 Z_1 \dots$), consecutive bytes are from different columns with different statistical properties. Delta encoding makes this *worse*:

Row-major: 128 64 200 129 63 201 130 62 202 ...

Delta: 64 136 57 190 138 57 190 138 ... (wild jumps between columns)

Transpose to column-major ($X_0 X_1 X_2 \dots Y_0 Y_1 Y_2 \dots Z_0 Z_1 Z_2 \dots$) and now consecutive bytes are from the same column. Delta encoding now sees smooth progressions:

Column-major: 128 129 130 131 ... 64 63 62 61 ... 200 201 202 203 ...

Delta: 1 1 1 1 ... -1 -1 -1 ... 1 1 1 1 ... (trivial)

The numbers are striking. A 768-byte vertex table (256 vertices $\times$ 3 columns):

Layout	Entropy (raw)	Entropy (delta)
Row-major (X,Y,Z interleaved)	7.52	7.66 (worse!)
Column-major, stride 3	7.52	2.58

Delta encoding on row-major data *increased* entropy. The same delta on transposed data reduced it by 65%. The compressor does not know your data is tabular — you have to tell it, by reordering.

The rule: if your data has columns with different patterns, **always transpose before compressing**. The stride (number of columns) does not need to be guessed — try a few divisors of the data length and pick whichever gives the lowest delta entropy.

On the Spectrum, the decompressor just writes bytes sequentially. The transposition happens in your build tools, not at runtime. Zero runtime cost.

Interleaving planes: masks and pixels

Sprites with masks are a special case of transposition. Stored as mask-pixel-mask-pixel per row, consecutive bytes alternate between two completely different distributions (masks are mostly \$FF or \$00; pixels have diverse values). Separate all mask bytes from all pixel bytes:

Before: FF 3C FF 18 FF 00 ... (mask, pixel, mask, pixel)

After: FF FF FF ... 3C 18 00 ... (all masks, then all pixels)

The mask block compresses to nearly nothing (long runs of \$FF). The pixel block compresses normally. Combined ratio improves 10–20% over interleaved storage, depending on sprite complexity.

Pattern detection: when not to compress

Sometimes data has structure that a generator can reproduce more cheaply than a decompressor. If your data is periodic with period P , storing one period plus a tiny replay loop takes $P + \sim 10$ bytes. If P is small relative to the total data, this beats any compressor.

Sine tables are the canonical case. A 256-byte sine table compresses to ~ 140 bytes with ZX0. But a Spectrum-friendly sine generator (using the ROM calculator or a CORDIC kernel) produces the same 256 bytes from under 30 bytes of code. For demo-quality accuracy, even a simple quadratic approximation per quarter-wave suffices.

The decision tree: (1) Can you generate it from a formula in fewer bytes than compressed size? Generate. (2) Is the data periodic? Store one period + loop. (3) Is the data tabular? Transpose + delta + LZ. (4) Is the data sequential frames? Delta + LZ. (5) None of the above? Just compress it.

Practical transforms for common demo data

Data type	Best pre-transform	Why
Sine/cosine tables	2nd derivative, or generate at runtime	Smooth acceleration $\rightarrow$ constant 2nd deriv
3D vertex tables	Transpose (stride = fields per vertex) + delta	Separates axes; smooth trajectories per axis
Precalculated animation	Delta between frames + LZ	High inter-frame redundancy
AY register dumps	Transpose (stride = 14, one per register) + delta	Each register varies smoothly between frames
Colour ramps / gradients	1st derivative	Linear or near-linear progression
Tile maps	Transpose (stride = map width) + delta	Spatial locality: adjacent tiles are similar
Bitmap font data	Separate bit-planes, or store as 1-bit + RLE	Lots of zero bytes in descenders
Particle positions	Sort by one axis, then delta-encode each axis	Sorted order maximises delta compression

The key insight: **every byte you save with a free pre-transform is a byte you do not need a more expensive packer to save.** Transpose + delta + Pletter 5 (fast decompressor) often beats raw Exomizer (slow decompressor) on structured data. You get a better ratio *and* faster decompression.

The Practical Pipeline

Understanding compression algorithms is useful. Integrating them into your build pipeline is essential.

From asset to binary

The pipeline: source asset (PNG) -> converter (png2scr) -> compressor (zx0) -> assembler (sjasmpplus) -> .tap file. The compressor runs on your development machine, not the Spectrum. For ZX0: zx0 screen.scr screen.zx0. Include the result with sjasmpplus's INCBIN directive:

```
compressed_screen:
    incbin "assets/screen.zx0"
```

At runtime, decompress with a simple call:

```
ld    hl, compressed_screen    ; source: compressed data
ld    de, $4000                ; destination: screen memory
call  dzx0_standard            ; decompress
```

Makefile integration

The compression step belongs in your Makefile, not in your head:

```
zx0 $< $@

demo.tap: main.asm assets/screen.zx0
    sjasmpplus main.asm --raw=demo.bin
    bin2tap demo.bin demo.tap
```

Change a source PNG, run make, and the compressed binary is regenerated automatically. No manual steps, no forgotten recompression.

Example: loading screen with ZX0

A complete minimal example — decompress a loading screen to video memory and wait for a keypress:

```
; loading_screen.asm — assemble with sjasmpplus
    org $8000
start:
    ld    hl, compressed_screen
    ld    de, $4000
    call  dzx0_standard

.wait: xor    a
        in    a, ($fe)
        cpl
        and   $1f
        jr    z, .wait
        ret

    include "dzx0_standard.asm"

compressed_screen:
    incbin "screen.zx0"

    display "Total: ", /d, $ - start, " bytes"
```

ZX0 decompression demo – a compressed loading screen unpacked to video memory in real time

Figure 33: ZX0 decompression demo – a compressed loading screen unpacked to video memory in real time

Use sjasmpplus’s `DISPLAY` directive to print size information during assembly. Always know exactly how large your compressed data is — the difference between ZX0 and Exomizer on a single loading screen can be 400 bytes, and over 8 scenes that adds up.

Choosing the right compressor

Ask in order: (1) Size-coded intro? ZX0/ZX7 — the 69–70 byte decompressor is non-negotiable. (2) Real-time streaming? LZ4 — nothing else is fast enough. (3) One-time load? Exomizer — maximum ratio, speed irrelevant. (4) Need a balance? ApLib or Pletter 5, both on the Pareto frontier. (5) Data full of identical runs? Custom RLE. (6) Sequential animation frames? Delta coding first, then LZ.

The MegaLZ Revival

In 2017, Introspec declared MegaLZ “morally obsolete.” Two years later, he resurrected it himself.

The insight: the compression *format* and the *decompressor implementation* are separable problems. MegaLZ’s format was good — the first Spectrum compressor to use an optimal parser (LVD, 2005), with Elias gamma codes and a slightly larger window than Pletter 5. What was bad was the Z80 decompressor. Introspec wrote two new ones:

- **Compact:** 92 bytes, ~98 T-states per byte
- **Fast:** 234 bytes, ~63 T-states per byte — faster than three consecutive LDIRs

With these decompressors, MegaLZ “handily beats Pletter 5 and ZX7” on the combined ratio-plus-speed metric. The lesson: do not assume a compressor is dead. The format is the hard part. The decompressor is Z80 code, and Z80 code can always be rewritten.

What the Numbers Mean in Practice

4K intro: 4,096 bytes total. ZX0 decompressor: ~70 bytes. Engine + music + effects: ~2,400 bytes. That leaves ~1,626 bytes for compressed data, which decompress to ~3,127 bytes of raw assets. SerzhSoft’s Megadematica (1st place, Revision 2019) compressed tunnel effects, transitions, AY music, and rapid scene changes into exactly 4,096 bytes. It was nominated for Outstanding Technical Achievement at the Meteoriks.

Real-time streaming: You need 2 KB of data per frame at 50 fps. LZ4 at 34 T/byte decompresses 2,048 bytes in 69,632 T-states — almost exactly one frame (69,888

T-states on 48K). Tight but feasible with overlapped border-time decompression. ApLib would need 215,040 T-states for the same data — over three frames. Exomizer over seven. For streaming, LZ4 is the only option.

128K multi-scene demo: Eight scenes, each with a 6,912-byte loading screen. Exomizer compresses each to ~3,338 bytes; ZX0 to ~3,594 bytes. The difference: 256 bytes per screen, 2,048 bytes across 8 scenes. When decompression happens during scene transitions, Exomizer's slow decompression is invisible. The 2 KB savings is not.

256-byte intro: ZX0's 70-byte decompressor leaves 186 bytes for everything. More commonly at this size, you skip LZ and generate data procedurally with LFSR generators and ROM calculator calls. But when you need specific non-algorithmic data — a colour ramp, a bitmap fragment — ZX0 remains the tool.

Summary: The Compressor Cheat Sheet

Your situation	Use this	Why
One-time load, maximum ratio	Exomizer	48.3% ratio, speed irrelevant
General purpose, good balance	ApLib	49.2% ratio, ~105 T/byte
Need speed + decent ratio	Pletter 5	51.5% ratio, ~69 T/byte
Real-time streaming	LZ4	~34 T/byte, 2+ KB per frame
Size-coded intro (256b-1K)	ZX0 / ZX7	69-70 byte decompressor
4K intro	ZX0	Tiny decompressor + good ratio
Runs of identical bytes	RLE (custom)	Decompressor under 30 bytes
Sequential animation frames	Delta + LZ	Exploit inter-frame redundancy

The numbers are the answer. Not opinions, not folklore, not "I heard Exomizer is the best." Introspec tested ten compressors across 1.2 megabytes of real Spectrum data and published the results. Use his numbers. Choose the compressor that fits your constraints. Then move on to the hard part — making something worth compressing.

Try It Yourself

1. **Compress a loading screen.** Take any ZX Spectrum .scr file (grab one from zxart.ee or create your own in Multipaint). Compress it with ZX0 and Exomizer. Compare the sizes. Then write the minimal loader shown in this chapter to

decompress and display it. Time the decompression using border-colour timing from Chapter 1.

2. **Measure the streaming limit.** Write a tight loop that decompresses data with the ZX0 standard decompressor and measures how many bytes it can decompress per frame. Compare with the LZ4 decompressor. Verify the numbers from the benchmark table against your own measurements.
3. **Build a delta compressor.** Take two ZX Spectrum screens that differ slightly (save a game screen, move a sprite, save again). Write a simple tool (in Python or your language of choice) that produces a delta stream: a list of (offset, new_value) pairs for bytes that differ. Compare the size of the delta stream to the size of the full second screen. Then compress the delta stream with ZX0 and compare again.
4. **Integrate compression into a Makefile.** Set up a project with a Makefile that automatically compresses assets as a build step. Change a source PNG, run make, and verify that the compressed binary is regenerated and the final .tap file is updated. This is the workflow you will use for every project from now on.
5. **Transpose and measure.** Create a 768-byte file of 256 (X, Y, Z) triplets where X is a sine wave, Y is a cosine, and Z is a linear ramp. Measure the entropy of the raw file. Then transpose it (all X values, then all Y, then all Z) and measure again. Apply delta encoding to both versions and compare. You should see the transposed+delta version drop below 3 bits/byte, while the raw+delta version stays above 7. Compress both with ZX0 and compare the actual sizes — the entropy numbers predict the winner.
6. **The quadratic substitution.** Generate a 256-byte sine table and a 256-byte quadratic approximation (fit $ax^2 + bx + c$ to one quarter-wave, mirror for the full cycle). Plot both — they should be visually identical. Now compute the second derivative of each. The sine's second derivative has entropy ~ 1.5 bits/byte; the quadratic's is exactly 0. Compress both with ZX0. The quadratic version is smaller, and the animation looks the same.

Sources: Introspec “Data Compression for Modern Z80 Coding” (Hype, 2017); Introspec “Compression on the Spectrum: MegaLZ” (Hype, 2019); Break Space NFO (Thesuper, 2016); Einar Saukas, ZX0 (github.com/einar-saukas/ZX0); Ped7g (Peter Helcmanovsky), self-modifying RLE depacker (contributed with permission, 2026)

Chapter 15: Anatomy of Two Machines

“Design characterizes realizational, stylistic, ideological integrity.” – Introspec (spke), “For Design” (Hype 2015)

Welcome to Part V. We are building a game.

Parts I through IV gave you the demoscener’s toolkit: cycle counting, screen tricks, inner loop optimisation, sound architecture, compression. A game makes different demands. You need the full memory map, not just the screen region. You need to understand banking, because your levels, music, and sprite data will not fit in one contiguous block. You need to know how the Agon Light 2’s two processors talk to each other, because your game loop straddles that boundary.

This chapter is the hardware reference for everything that follows. Where Chapter 1 gave you the frame budget and Chapter 2 gave you the screen layout, this chapter gives you everything else. Keep it bookmarked.

15.1 ZX Spectrum 128K: Memory Map

The original 48K Spectrum had a simple memory model: 16 KB ROM at \$0000-\$3FFF, 48 KB RAM at \$4000-\$FFFF. The 128K model keeps this layout visible to the CPU but hides a banking system underneath.

The 128K has eight 16 KB pages of RAM (pages 0-7, totalling 128 KB) and two 16 KB ROMs (ROM 0: 128K editor, ROM 1: 48K BASIC). At any moment, the Z80 sees a 64 KB address space divided into four 16 KB slots:

Address Range	Contents	Notes
\$0000-\$3FFF	ROM (0 or 1)	Selected by bit 4 of \$7FFD
\$4000-\$7FFF	RAM page 5	Always page 5. Screen memory lives here.
\$8000-\$BFFF	RAM page 2	Always page 2.
\$C000-\$FFFF	RAM page N	Switchable: any page 0-7 via \$7FFD

Pages 5 and 2 are hardwired into their slots. You cannot swap them out. This means the screen (\$4000-\$57FF) is always accessible, and your main code (typically ORG'd at \$8000) sits in page 2 where it will not disappear when you switch banks.

The top 16 KB slot at \$C000-\$FFFF is the flexible one. Write to port \$7FFD and the page mapped there changes.

The \$7FFD Port: Bank Switching

Port \$7FFD controls memory configuration on the 128K. It is write-only – you cannot read it back. This means you must shadow its value in a RAM variable if you need to know the current state.

Port \$7FFD bit layout:

- Bit 0-2: RAM page mapped at \$C000 (0-7)
- Bit 3: Screen select (0 = normal screen at page 5,
1 = shadow screen at page 7)
- Bit 4: ROM select (0 = 128K ROM, 1 = 48K BASIC ROM)
- Bit 5: Disable paging (set this and banking is locked
until next reset -- used by 48K BASIC)
- Bits 6-7: Unused

A typical bank-switch routine:

```
; Switch RAM page at $C000 to page number in A (0-7)
; Preserves other $7FFD bits from shadow variable
bank_switch:
    ld    b, a                ; 4T    save desired page
    ld    a, (bank_shadow)    ; 13T   load current $7FFD state
    and   %11111000          ; 7T    clear page bits (0-2)
    or    b                   ; 4T    insert new page number
    ld    (bank_shadow), a    ; 13T   update shadow
    ld    bc, $7FFD           ; 10T
    out   (c), a              ; 12T   do the switch
    ret                        ; 10T
                        ; --- 73T total
```

Those 73 T-states are not free, but they are negligible compared to a frame budget of 70,000+. The real cost of banking is architectural, not temporal: you must design your data layout so that you never need to access two different banked pages simultaneously. Music data in page 4, level data in page 3, sprite graphics in page 6 – but your music player and your renderer cannot both be running from \$C000 at the same time.

The shadow screen. Bit 3 of \$7FFD selects which RAM page the ULA reads for display: page 5 (normal) or page 7 (shadow). This gives you hardware double-buffering – draw into the shadow screen while the ULA displays the normal one, then flip by toggling bit 3. We will use this heavily in Chapter 17 (Scrolling) and Chapter 21 (Full Game).

A Practical Memory Map for a 128K Game

Here is how a real game might lay out its 128 KB across the eight pages:

Page	Slot	Usage
0	\$C000 (banked)	Level data (maps, tile definitions)
1	\$C000 (banked)	Sprite graphics set 1
2	\$8000-\$BFFF (fixed)	Main game code, entity system, interrupt handler
3	\$C000 (banked)	Sprite graphics set 2, lookup tables
4	\$C000 (banked)	Music data (.pt3 patterns, instruments)
5	\$4000-\$7FFF (fixed)	Primary screen, attribute memory, system vars
6	\$C000 (banked)	Sound effects, additional level data
7	\$C000 (banked)	Shadow screen (double-buffer target)

Notice: page 7 serves double duty. The ULA can display it as the shadow screen, but you can also bank it into \$C000 and use it as a 16 KB data page when you are not double-buffering. Many demos exploit this.

The critical constraint: **your interrupt handler and main loop must live in pages 2 or 5**, because those are the only pages guaranteed to be mapped at all times. If an interrupt fires while page 4 is banked in at \$C000, and your interrupt handler lives at \$C000, the CPU jumps into your music data instead of your code. The result is a crash, usually a spectacular one.

Rule: never put time-critical code in a banked page unless you are absolutely certain which page is active when that code runs.

15.2 Contended Memory: The Practical Truth

In Chapter 1, we established that Pentagon clones have no contended memory and that cycle counts are reliable everywhere. That is true, and the Pentagon remains the standard for cycle-counted demo work. But if you are writing a game for release, your players will include people running original Sinclair hardware, Amstrad +2A/+3 models, and modern FPGA clones that emulate the original timing. You need to know what contended memory does and how to avoid its worst effects.

Introspec covered this exhaustively in his “GO WEST” articles on Hype (2015). Here is the practical summary.

What Gets Contended

On the original Sinclair machines, the ULA and the CPU share a memory bus. When the ULA is reading screen data to paint the display (during the 192 active scanlines), any CPU access to certain RAM pages gets delayed. The CPU is literally halted for extra T-states while the ULA finishes its read.

The contended pages differ between models:

ZX Spectrum 128K Memory Map




Figure 34: ZX Spectrum 128K memory map

Model	Contended Pages	Always Fast
48K	Page 5 only (\$4000-\$7FFF)	\$8000-\$FFFF (uncontended)
128K / +2	Pages 1, 3, 5, 7	Pages 0, 2, 4, 6
+2A / +2B / +3	Pages 4, 5, 6, 7	Pages 0, 1, 2, 3

The 48K contends only the lower 16 KB of RAM (\$4000-\$7FFF, page 5) – the upper 32 KB (\$8000-\$FFFF) is uncontended. On the 128K, the pattern is every odd page. On the +2A/+3, it flips: the high pages are contended.

This has immediate practical consequences. On a 128K, your main code at \$8000 (page 2) is in uncontended memory – fast. The screen at \$4000 (page 5) is contended – writes to screen memory are slower during active display. And page 7 (the shadow screen) is also contended, which means double-buffer fills to the shadow screen are slower than you might expect on original hardware.

How Much Slower?

Introspec measured the actual penalties:

- **Random byte access to contended memory:** approximately **0.92 extra T-states per byte** on average during active display
- **Stack operations (PUSH/POP) to contended memory:** approximately **1.3 extra T-states per byte** on average
- **During border time: zero penalty** – contention only occurs while the ULA is actively painting scanlines

That 0.92 figure means a LD A, (HL) that should cost 7 T-states will cost, on average, about 7.92 T-states when HL points into contended memory during active display. A PUSH that writes two bytes to contended memory at 11 T-states will cost about 13.6 T-states instead.

These averages hide a messy reality: the actual penalty depends on where in the ULA's 8-T-state read cycle your CPU access lands. The pattern repeats every 8 T-states: penalties of 6, 5, 4, 3, 2, 1, 0, 0 extra states. You can land anywhere in this cycle, and the penalty compounds with each memory access within an instruction. This makes precise cycle counting on contended machines genuinely difficult.

The Practical Response

For game development, not demo effects, the approach is straightforward:

1. **Put your code in uncontended memory.** On the 128K, ORG at \$8000 (page 2) – always fast.
2. **Write to the screen during border time when possible.** The top and bottom borders give you contention-free access to screen memory. So does the left/right border of each scanline.
3. **Do not worry about precise contention modelling.** Budget a 15-20% slowdown for code that touches screen memory during active display, and design your frame budget with that margin. This is not cycle-counted demo work; this is game development.
4. **Test on real hardware or accurate emulators.** Fuse emulates contention correctly. Unreal Speccy (Pentagon mode) does not, by design. ZEsarUX can emulate multiple models.

Introspec's advice from GO WEST boils down to this: **contended memory is a portability issue, not a drama**. If your code runs on Pentagon, it will almost certainly run on original hardware too – just a little slower during screen writes. The places where contention actually breaks things are cycle-exact raster effects (multicolour, floating bus sync), and those are demo techniques, not game techniques.

15.3 ULA Timing

The ULA generates the video signal and the CPU interrupt. Understanding its timing is essential for border effects, interrupt-driven music, and screen synchronisation.

Frame Structure

A complete frame consists of scanlines. The scanline width and total scanline count differ between models:

Machine	T-states/line	Scanlines	T-states/frame
ZX Spectrum 48K	224	312	69,888
ZX Spectrum 128K	228	311	70,908
Pentagon 128	224	320	71,680

Note the 128K's wider scanline (228 vs 224 T-states). The extra 4 T-states per line are in the border/sync portion, not the active display.

Tact-Maps: Frame Regions

The frame divides into three regions. The interrupt fires at the start of vertical blank, before the top border. Here is the timing map for each model:

Pentagon 128 (71,680 T-states)

Top border	80 lines × 224T = 17,920T	No screen reads. No contention.
Active display	192 lines × 224T = 43,008T	ULA reads screen memory. No contention on Pentagon.
Bottom border	48 lines × 224T = 10,752T	No screen reads. No contention.
Total: 71,680T		

ZX Spectrum 128K (70,908 T-states)

Top border	63 lines × 228T = 14,364T	No screen reads. No contention.
Active display	192 lines × 228T = 43,776T	ULA reads screen memory. Contention on pages 1,3,5,7.

Bottom border	56 lines × 228T = 12,768T	No screen reads. No contention.
	Total: 70,908T	

ZX Spectrum 48K (69,888 T-states)

Top border	64 lines × 224T = 14,336T	No screen reads. No contention.
Active display	192 lines × 224T = 43,008T	ULA reads screen memory. Contention on \$4000-\$7FFF.
Bottom border	56 lines × 224T = 12,544T	No screen reads. No contention.
	Total: 69,888T	

After a HALT, you have the entire top border period – 17,920 T-states on Pentagon, 14,364 on 128K – to do work before the beam enters the active display area and contention begins. This is why well-structured Spectrum code does screen writes at the top of the frame: you get contention-free access to screen memory during the border period.

Scanline Timing

Each scanline breaks down into an active portion (where the ULA reads screen data) and border/sync portions:

48K and Pentagon (224 T-states per line):

24T right border
48T horizontal sync + retrace
24T left border

128K (228 T-states per line):

24T right border
52T horizontal sync + retrace
24T left border

During the 128 active T-states, memory access to contended pages gets delayed (on non-Pentagon machines). During the remaining 96 T-states (or 100 on 128K), no contention. Even during active display, roughly half of each scanline is contention-free.

Total vs Practical Budget

The frame totals above are the time between interrupts. The *practical* budget – T-states available for your code – is less:

Overhead	Cost
HALT + interrupt acknowledge (IM1)	~30 T-states
Minimal ISR (EI + RET)	~14 T-states
Typical PT3 music player (in ISR)	~3,000–5,000 T-states
Main loop housekeeping (frame counter, HALT jump)	~20–50 T-states

Practical budgets with a music player running:

Machine	Total	After PT3 player	After player + contention margin
Pentagon	71,680	~66,000–68,000	~66,000–68,000 (no contention)
128K	70,908	~65,000–67,000	~55,000–60,000 (screen writes during active display)
48K	69,888	~64,000–66,000	~55,000–60,000 (\$4000–\$7FFF contended)

When this book says “frame budget of ~70,000 T-states,” it means the total. When planning your inner loops, budget for the practical figure – typically 65,000–68,000 on Pentagon with music.

15.4 Floating Bus, ULA Snow, and the \$7FFD Bug

These are three hardware quirks that appear on original Sinclair machines but not on most clones. You may never encounter them in game development, but they can cause mysterious bugs if you do not know they exist.

Floating Bus

On original Spectrum hardware, reading from an unattached port returns whatever data the ULA happens to be putting on the data bus at that moment. During active display, the ULA is reading screen memory, so a read from port \$FF returns the byte the ULA is currently reading.

Demo coders exploit this for beam synchronisation: read the floating bus in a tight loop until you see a known value from screen memory, and you know exactly where the beam is. This is the cheapest sync method – no interrupt timing required.

Games rarely need this, but be aware: if your code reads from a port that does not exist on the hardware, the return value is unpredictable and varies between models. The floating bus is *not* emulated on Pentagon, Scorpion, or ZX Next.

The \$7FFD Read Bug

Port \$7FFD is write-only. But on some Spectrum models, reading from port \$7FFD (even unintentionally, through an instruction that happens to put \$7FFD on the address bus) causes the floating bus value to be written into the port. This triggers a spurious page switch.

The practical danger: the Z80 instruction LD A, (nn) puts the address nn on the bus during execution. If nn happens to be \$7FFD and you are reading data stored at address \$7FFD, the memory read can trigger a port write on original hardware. This is an obscure bug but a real one. Avoid storing data at address \$7FFD.

ULA Snow

If the Z80’s I register (used for IM2 interrupt vector table base) is set to a value in the range \$40–\$7F, the DRAM refresh cycle during every M1 opcode fetch puts

an address in the \$4000-\$7FFF range onto the address bus. This conflicts with the ULA's screen reads and produces visual "snow" – random noise on the display.

The fix is simple: **never set I to a value between \$40 and \$7F**. The typical IM2 setup uses I = \$FE with a 257-byte table of identical vectors at \$FE00-\$FF00. This keeps I well above the danger zone.

15.5 Clone Differences

The ZX Spectrum ecosystem includes dozens of clones, but three matter most for modern development: the Pentagon 128, the Scorpion ZS-256, and the ZX Spectrum Next.

Pentagon 128

The Pentagon is the standard platform for the Russian demoscene and the primary target for this book's demoscene chapters.

Parameter	Pentagon 128	Original 128K
CPU clock	3.5 MHz	3.5 MHz
T-states per frame	71,680	70,908
Scanlines per frame	320	311
Contended memory	None	Pages 1, 3, 5, 7
Border lines (top)	80	63
Border lines (bottom)	48	56

The extra 772 T-states per frame (71,680 vs 70,908) come from the additional scanlines. The border is distributed differently: a taller top border and a shorter bottom border. This affects border effects – demo code that produces a symmetrical border pattern on the 128K will be slightly asymmetric on the Pentagon.

The absence of contended memory is the Pentagon's defining feature for programmers. Every instruction costs exactly what the datasheet says. This is why we use Pentagon timing throughout this book.

7 MHz Turbo mode. Many Pentagon-compatible machines (Pentagon 512, Pentagon 1024, ATM Turbo 2+) offer a 7 MHz turbo mode. The CPU runs at double speed, but the ULA timing stays the same. This means the frame budget doubles to approximately 143,360 T-states in turbo mode. The catch: turbo mode is not standard across all machines, and code that relies on it will not run on a stock Pentagon 128 or any Sinclair hardware.

For games, turbo mode is a luxury that lets you run more complex logic or more sprites per frame. For demos targeting compo rules, it is usually forbidden – competitions specify "Pentagon 128K, 3.5 MHz."

Scorpion ZS-256

The Scorpion is a Ukrainian clone with 256 KB of RAM (16 pages of 16 KB) and several hardware extensions.

Feature	Scorpion ZS-256
RAM	256 KB (16 pages)
Banking	Extended \$1FFD port for pages 8-15
Graphics	GMX mode: 320x200, 16 colours from 256
Contended memory	None
Frame timing	Pentagon-compatible (71,680 T-states)

The doubled RAM is useful for games: you get 16 data pages instead of 8. The extra pages are accessed via port \$1FFD, which uses a similar scheme to \$7FFD but controls the additional RAM.

GMX (Graphics Mode Extended) is the Scorpion's party trick: a 320x200 display with 16 colours chosen from a 256-colour palette. This breaks completely with the Spectrum's attribute-based display, offering a linear framebuffer closer to what you might see on an Amiga or a PC VGA. The GMX framebuffer is large (32,000 bytes for 4-bit colour) and lives in the extended RAM pages.

Few games target GMX because it limits your audience to Scorpion owners. But it demonstrates what Z80 hardware can do when freed from the ULA's attribute grid.

ZX Spectrum Next

The ZX Spectrum Next is the modern flagship of the platform: an FPGA-based machine that is backward-compatible with the original Spectrum but adds substantial new hardware.

Feature	ZX Spectrum Next
CPU	Z80N (Z80 + new instructions) at 3.5 / 7 / 14 / 28 MHz
RAM	1 MB (extendable to 2 MB), 8 KB MMU pages
MMU	8 slots x 8 KB = fine-grained memory mapping
Layer 2	256x192 or 320x256, 8-bit colour (256 colours)
Tilemap	Hardware tilemap layer, 40x32 or 80x32 tiles
Sprites	128 hardware sprites, 16x16, up to 12 per scanline
Copper	Co-processor for per-scanline register changes
DMA	zxndMA for fast block transfers
AY sound	3 x AY-3-8910 (9 channels) with per-channel stereo panning

The Next's **MMU** is fundamentally different from the 128K's banking. Instead of one switchable 16 KB slot, the Next divides the entire 64 KB address space into eight 8 KB slots. Each slot can be independently mapped to any 8 KB page from the 1-2 MB RAM pool. This means you can have fine-grained control:

```
; Map 8KB page $0A into slot 3 ($6000-$7FFF)
ld  a, $0A
ld  bc, $243B      ; Next register select port
ld  a, $53         ; Register $53 = MMU slot 3
out (c), a
ld  bc, $253B      ; Next register data port
```

```
ld    a, $0A          ; Page $0A
out   (c), a
```

This is much more flexible than the 128K's single switchable slot. You can map sprite data into one 8 KB window, level data into another, and music data into a third – all simultaneously visible.

Layer 2 gives you a 256-colour bitmap display without attribute clash. This is the single biggest quality-of-life improvement for game developers: no more careful attribute planning, no more colour clash workarounds. Just a framebuffer where each byte is one pixel. The cost is memory: a 256x192 Layer 2 screen is 49,152 bytes.

Hardware sprites on the Next provide 128 sprite slots, each 16x16 pixels with 8-bit colour, up to 12 per scanline. Sprite attributes (position, pattern, rotation) are set through Next registers and port \$57. No software rendering needed.

The Copper is a co-processor that executes a simple program synchronised to the beam position. It can write to any Next register at any scanline, enabling per-line palette changes, scroll offsets, and raster effects without consuming Z80 T-states – a deliberate homage to the Amiga Copper.

zxndMA provides hardware-accelerated block transfers at approximately 2 T-states per byte – about 10 times faster than LDIR. For filling the Layer 2 framebuffer or transferring sprite data, DMA is transformative.

The Next is essentially a different machine that happens to be backward-compatible. The interesting constraints shift from “can I fit this in the frame budget” to “how do I best use multiple hardware layers.”

15.6 Agon Light 2: A Different Beast

The Agon Light 2 is the second platform for our game development chapters. It runs a Zilog eZ80 – a direct descendant of the Z80 – at 18.432 MHz, with 512 KB of flat RAM and a separate ESP32 co-processor handling video and audio. The architecture is fundamentally different from the Spectrum: instead of a CPU that shares a bus with a fixed video output chip, the Agon uses two independent processors communicating over a serial link.

Dual-Processor Architecture

The Agon's defining characteristic is the split between the **eZ80** (your CPU) and the **ESP32** (the VDP, Video Display Processor):

```

|   eZ80   | 18.432 MHz
|  512 KB  | Your code runs here
| MOS API  | 24-bit addressing
+-----+
|
|   UART serial
| (1,152 Kbaud)
|
+-----+
```

	ESP32		240 MHz dual-core
	FabGL		Video: up to 640x480
	VDP		Audio: waveforms, samples
+-----+			

This split has important consequences:

1. **No shared video memory.** You cannot write directly to a framebuffer. Every pixel, every sprite, every tile operation is a *command* sent over the serial link from the eZ80 to the ESP32.
2. **Latency.** The serial link runs at 1,152,000 baud (since MOS 1.03). A single command byte takes about 8.7 microseconds to transmit. Complex drawing operations (fill rectangle, draw bitmap) require multiple bytes and the VDP needs time to execute them.
3. **Asynchronous rendering.** The VDP processes commands from a buffer. Your eZ80 code sends commands and continues running. The VDP catches up independently. This means you do not have the Spectrum's tight coupling between CPU work and screen output – but you also cannot precisely control when pixels appear.
4. **Independent frame rate.** The VDP renders at its own rate (typically 60 Hz). Your eZ80 game loop can run at whatever rate it wants; the VDP will display whatever it has most recently drawn.

For Spectrum programmers, this is a different approach entirely. You go from “I write bytes to video memory and they appear on the next scanline” to “I send drawing commands and trust the VDP to render them eventually.” The upside is enormously reduced CPU overhead for graphics. The downside is less control.

eZ80 Memory Model: 24-Bit Flat

The eZ80 has a 24-bit address bus, giving it a theoretical 16 MB address space. The Agon Light 2 maps 512 KB of RAM into the bottom of this space:

Address Range	Size	Contents
\$000000-\$07FFFF	512 KB	RAM
\$080000-\$0FFFFFF	512 KB	RAM (mirror, on some boards)
\$A00000-\$FFFFFF	varies	I/O, on-chip peripherals

No banking. No page switching. No contended memory. Your code, your data, your buffers, your lookup tables – everything lives in one flat, linearly addressable space. After the Spectrum's 8-page juggling act, the simplification is immediate.

The eZ80 supports two operating modes that determine how it uses this address space.

ADL Mode vs Z80 Mode

This is the most important architectural distinction on the Agon, and it trips up newcomers constantly.

Z80 mode (also called Z80-compatible mode) makes the eZ80 behave like a classic Z80: 16-bit registers, 16-bit addresses, 64 KB address space. All standard Z80

code runs unmodified. The upper 8 bits of the address come from the MBASE register, creating a 64 KB “window” into the 24-bit address space. This is what you use when porting existing Z80 code.

ADL mode (Address Data Long) is the eZ80’s native mode: 24-bit registers, 24-bit addresses, full 16 MB address space. HL, BC, DE, SP, and IX/IY are all 24 bits wide. LD HL,\$123456 loads a 3-byte value. PUSH HL pushes 3 bytes onto the stack (not 2). Every pointer is 3 bytes.

```
; ADL mode: 24-bit addressing, full 512KB accessible
ld    hl, $040000      ; point to a buffer 256KB into RAM
ld    (hl), $FF        ; write directly -- no banking needed
ld    bc, 1024
ld    de, $040001
ldir                                ; fill 1KB in one shot
```

MOS (the Agon’s operating system) boots the eZ80 in ADL mode, and most Agon software stays in ADL mode. The key differences from Z80 mode:

Feature	Z80 Mode	ADL Mode
Register width	16 bits	24 bits
Address space	64 KB (via MBASE)	16 MB (24-bit)
PUSH/POP size	2 bytes	3 bytes
JP/CALL addresses	16-bit	24-bit
Stack frame size	2 bytes per entry	3 bytes per entry
Instruction encoding	Z80-compatible	Extended (3-byte addresses)

The trap: if you write code assuming 16-bit values and run it in ADL mode, things break in subtle ways. A PUSH HL pushes 3 bytes, not 2, so your stack-based data structures have a different size. A JP (HL) jumps to a 24-bit address, so lookup tables of 16-bit addresses will not work. The eZ80 provides LD.S and LD.L suffixed instructions to explicitly control the data width, and you can switch between modes with JP.LIL / JP.SIS prefixes, but this gets complicated fast.

The practical rule for games: stay in ADL mode. Use 24-bit addresses everywhere. Do not try to share code between a Spectrum build and an Agon build at the source level – the addressing is too different. Instead, share *algorithms* and *data formats*, with platform-specific implementations for memory access, I/O, and graphics.

MOS API: The Operating System

MOS (Machine Operating System) provides system services on the Agon: file I/O, keyboard input, timer access, and VDP communication. MOS calls are made through RST \$08 with a function number in the A register:

```
; MOS API: open a file
ld    hl, filename      ; pointer to null-terminated filename
ld    c, $01            ; mode: read
rst    $08              ; MOS call
db    $0A               ; function $0A: ffs_fopen
; Returns file handle in A
```

```
filename:
    db    "level1.dat", 0
```

Key MOS functions for game development:

Function	Code	Description
mos_getkey	\$00	Read keyboard (non-blocking)
mos_load	\$01	Load file from SD card
mos_save	\$02	Save file to SD card
mos_sysvars	\$08	Get pointer to system variables (vsync counter, etc.)
ffs_fopen	\$0A	Open file
ffs_fclose	\$0B	Close file
ffs_fread	\$0C	Read from file
mos_getrtc	\$12	Get real-time clock

File I/O on the Agon is trivially easy compared to the Spectrum. No tape loading, no esxDOS wrappers, no TR-DOS: just open a file from the SD card and read it into memory. Level data, sprite sheets, music – load them on demand, no banking gymnastics.

VDP Commands: Talking to the Screen

All graphics go through VDU commands sent to the ESP32 VDP. The eZ80 sends bytes to a VDU output stream; the VDP interprets them as drawing instructions:

```
; VDP: draw a filled rectangle at (10, 10)
    rst $10 : db 25          ; PLOT command
    rst $10 : db 85          ; mode: filled rectangle
    rst $10 : db 10          ; x low
    rst $10 : db 0           ; x high
    rst $10 : db 10          ; y low
    rst $10 : db 0           ; y high
```

Verbose compared to LD (HL),A, but the VDP does the rendering on the ESP32. The VDP supports bitmap modes (up to 640x480), up to 256 hardware sprites (each up to 64x64), hardware tilemaps with scrolling, and audio (waveforms, ADSR, samples).

The bottleneck is the serial link, not the CPU. A complex scene with many sprite updates can saturate the UART, causing visual lag. Minimise VDP commands per frame: batch updates, use hardware scrolling instead of redrawing tiles, and let the sprite engine do the heavy lifting.

15.7 Comparing the Platforms

Let us lay out the two machines side by side, focusing on what matters for the game engine we will build in Chapters 16-19.

Feature	ZX Spectrum 128K	Agon Light 2
CPU	Z80A @ 3.5 MHz	eZ80 @ 18.432 MHz
T-states per frame	~70,908 (128K, 50 Hz) / 71,680 (Pentagon, 50 Hz)	~307,200 (60 Hz)
RAM	128 KB (8 x 16 KB pages)	512 KB (flat)
Address space	64 KB (banked)	16 MB (24-bit)
Screen memory	Shared bus, direct write	Separate VDP, command-based
Colours	15 (8 base x bright, minus overlap)	Up to 64 in standard modes
Resolution	256x192 (attribute colour per 8x8)	Configurable, up to 640x480
Sprites	Software only	Up to 256 hardware sprites
Scrolling	Software only (manual shift/copy)	Hardware scroll offsets
Sound	AY-3-8910 (3 channels)	ESP32 audio (multi-channel, waveforms)
Storage	Tape / DivMMC (esxDOS)	SD card (FAT32)
Double buffering	Shadow screen (page 7)	VDP-managed

The frame budget ratio is approximately 4:1 in the Agon's favour. But the Agon's graphics go through a serial bottleneck, so raw CPU speed does not translate directly to rendering speed. On the Spectrum, PUSH HL writes two bytes to the screen in 11 T-states. On the Agon, updating a sprite position requires 6+ bytes over a 1,152 Kbaud link, taking hundreds of microseconds regardless of CPU speed.

The Spectrum rewards byte-level optimisation. The Agon rewards architectural decisions. Both reward careful thinking about frame budgets.

15.8 Practical: Memory Inspector Utility

Let us build a simple memory inspector for both platforms. This utility displays a region of RAM as hex bytes on screen, and lets you navigate through memory with the keyboard. It is the kind of tool you will use constantly during development.

Spectrum Version

The Spectrum version writes directly to screen memory. We display 16 rows of 16 bytes (256 bytes per page) with the start address shown at the left.

```
; Memory Inspector - ZX Spectrum 128K
; Displays 256 bytes of memory as hex, navigable with keys
; ORG $8000 (page 2, uncontended)

    ORG  $8000

SCREEN_ATTR EQU $5800
START_ADDR  EQU inspect_addr      ; address to inspect (self-mod)
```

```

start:
    call clear_screen

main_loop:
    halt                ; sync to frame

    ; Read keyboard
    call read_keys      ; returns: A = action
    cp    1
    jr    z, .page_up   ; Q = previous page
    cp    2
    jr    z, .page_down ; A = next page
    cp    3
    jr    z, .bank_up   ; P = next bank
    cp    4
    jr    z, .bank_down ; 0 = previous bank
    jr    .draw

.page_up:
    ld    hl, (inspect_addr)
    ld    de, -256
    add   hl, de
    ld    (inspect_addr), hl
    jr    .draw

.page_down:
    ld    hl, (inspect_addr)
    ld    de, 256
    add   hl, de
    ld    (inspect_addr), hl
    jr    .draw

.bank_up:
    ld    a, (current_bank)
    inc   a
    and   7                ; wrap 0-7
    ld    (current_bank), a
    call  bank_switch
    jr    .draw

.bank_down:
    ld    a, (current_bank)
    dec   a
    and   7
    ld    (current_bank), a
    call  bank_switch

.draw:
    ; Display current bank and address
    call draw_header

    ; Display 16 rows x 16 bytes
    ld    hl, (inspect_addr)
    ld    b, 16            ; 16 rows

```

```

        ld    de, $4060                ; screen position (row 3, col 0)
.row_loop:
    push bc
    push hl

    ; Print address
    ld    a, h
    call print_hex                    ; print high byte of address
    ld    a, l
    call print_hex                    ; print low byte
    ld    a, ':'
    call print_char

    ; Print 16 hex bytes
    pop   hl
    push  hl
    ld    b, 16
.byte_loop:
    ld    a, (hl)
    call print_hex                    ; 17T call + print routine
    inc   hl
    ld    a, ' '
    call print_char
    djnz .byte_loop

    pop   hl
    ld    de, 16
    add   hl, de                      ; advance to next row
    pop   bc

    ; Move screen pointer down one character row
    call next_char_row

    djnz .row_loop

    jr    main_loop

; --- Data ---
inspect_addr:  dw $C000                ; start address to inspect
current_bank:  db 0                    ; current bank at $C000
bank_shadow:   db 0                    ; shadow of port $7FFD

; read_keys, print_hex, print_char, clear_screen,
; draw_header, next_char_row, bank_switch: implementations
; omitted for brevity -- see examples/mem_inspect.a80
; for the complete compilable source.

```

The key architectural point: we inspect \$C000 because that is the banked slot. By changing `current_bank`, we can page through all 8 RAM pages using the `bank_switch` routine from section 15.1. The inspector itself lives at \$8000 (page 2), safe from banking changes.

Agon Version

The Agon version uses MOS system calls for keyboard input and VDP text output. No screen memory calculation, no attribute handling – just send text to the VDP.

```
; Memory Inspector - Agon Light 2 (ADL mode)
    .ASSUME ADL=1
    ORG $040000

main_loop:
    ; Wait for vsync via MOS sysvar
    rst $08
    db $08                      ; mos_sysvars
    ld a, (ix+$00)              ; sysvar_time (low byte)
.wait_vsync:
    cp (ix+$00)
    jr z, .wait_vsync          ; spin until counter changes

    ; Check keyboard (Q = up, A = down)
    rst $08
    db $00                      ; mos_getkey
    ; ... navigation same as Spectrum version ...

.draw:
    rst $10
    db 30                      ; VDU 30 = cursor home

    ld hl, (inspect_addr)      ; 24-bit load!
    ld b, 16
.row_loop:
    push bc
    push hl
    call print_hex24           ; print full 24-bit address
    ld a, ':'
    rst $10
    pop hl
    push hl
    ld b, 16
.byte_loop:
    ld a, (hl)                 ; direct 24-bit access, no banking
    call print_hex8
    inc hl
    djnz .byte_loop
    pop hl
    ld de, 16
    add hl, de
    pop bc
    djnz .row_loop
    jr main_loop

inspect_addr: dl $000000      ; 24-bit address (dl, not dw)
; Full source: examples/mem_inspect_agon.a80
```

Notice the contrast:

- **No bank switching.** The Agon inspector can look at any address in 512 KB directly. LD HL, \$070000 and you are inspecting 448 KB into RAM. No ports, no shadow variables, no risk of banking in the wrong page.
- **No screen address calculation.** Text output goes through RST \$10, and the VDP handles cursor positioning, character rendering, and scrolling.
- **24-bit data directives.** We use dl (define long) for 3-byte pointers instead of dw (define word).
- **VSynC through system variables.** MOS provides a sysvar_time counter that increments each frame. We spin-wait on it for frame synchronisation – cruder than the Spectrum’s HALT, but functional.

Both inspectors do the same job. The Spectrum version is more code (you must handle everything yourself) but gives you total control. The Agon version is less code (the OS and VDP handle display) but gives you less control over exactly how the output looks.

This mirrors the broader development experience on both platforms. The Spectrum demands more effort for less visual richness. The Agon demands less effort for more visual richness. Both reward understanding the hardware.

Summary

- The **ZX Spectrum 128K** has 128 KB of RAM in 8 pages of 16 KB. Pages 2 and 5 are fixed in the address space; the top 16 KB slot at \$C000 is switchable via port \$7FFD. Keep your main code in page 2 and your interrupt handler away from banked memory.
- **Contended memory** slows CPU access to certain RAM pages during active display on original Sinclair hardware. Average penalty: ~0.92 extra T-states per byte. Pentagon clones have no contention. For game development, budget a 15-20% overhead on screen writes and keep time-critical code in uncontended pages.
- **ULA timing:** the interrupt fires at the top of the frame. You get ~14,000 T-states of contention-free time before the beam enters the active display area. Use this window for screen writes.
- The **\$7FFD port** is write-only. Shadow its value in RAM. Bit 3 selects the shadow screen (page 7) for double buffering. Bit 5 disables paging permanently until reset.
- **Floating bus, ULA snow,** and the **\$7FFD read bug** are quirks of original Sinclair hardware. Avoid I register values \$40-\$7F. Do not store data at address \$7FFD. The floating bus is not present on clones.
- **Pentagon 128:** no contended memory, 71,680 T-states per frame, 320 scan-lines. The demoscene standard. 7 MHz turbo mode doubles the frame budget on some variants.
- **Scorpion ZS-256:** 256 KB RAM (16 pages), GMX 320x200x16 colour mode.

- **ZX Spectrum Next:** 1-2 MB RAM with 8 KB MMU pages, Layer 2 (256-colour bitmap), 128 hardware sprites, Copper co-processor, zxnDMA, triple AY sound.
- The **Agon Light 2** uses a dual-processor architecture: eZ80 @ 18.432 MHz for logic, ESP32 for video/audio. 512 KB flat RAM, 24-bit addressing (ADL mode), MOS API for system services, VDP commands for all graphics.
- **ADL mode vs Z80 mode:** ADL mode uses 24-bit registers and addresses. Z80 mode emulates classic Z80 with 16-bit addresses via MBASE. Stay in ADL mode for new Agon code.
- The **serial link** between eZ80 and ESP32 is the Agon's bottleneck. Minimise VDP command traffic per frame. Use hardware sprites and tilemaps to reduce the number of drawing commands.
- Both platforms reward careful frame budget management. The Spectrum gives you ~70,000 T-states and demands byte-level optimisation. The Agon gives you ~307,000 T-states (at 60 Hz) but throttles graphics through a serial link. Different constraints, same discipline.

Sources: Introspec “GO WEST Parts 1-2” (Hype 2015); ZX Spectrum 128K Service Manual; Zilog eZ80 CPU User Manual; Agon Light 2 Documentation (Bernardo Kastrup); ZX Spectrum Next User Manual (2nd Edition)

Chapter 16: Fast Sprites

“Two colours per cell? Fine. But those two colours are going to move.”

Every game needs things that move. Bullets, enemies, the player character, explosions. On any hardware with a blitter or a GPU, the mechanics of putting a small image at an arbitrary screen position are handled for you. On the ZX Spectrum, they are your problem.

The Spectrum has no hardware sprites, no blitter, no co-processor. Every pixel of every sprite is placed by your Z80 code, one instruction at a time, into the same video memory the ULA is reading 50 times per second. And because the screen memory layout is interleaved (Chapter 2), “one row down” means `INC H` — unless you are crossing a character boundary, in which case it means something considerably uglier.

This chapter presents six methods for drawing sprites on the Spectrum, from the simplest XOR routine to compiled sprites that execute at the theoretical maximum speed of the hardware. Each method has trade-offs. We will also look at the Agon Light 2, where the VDP provides hardware sprites and the entire problem collapses to a handful of API calls.

Method 1: XOR Sprites

The simplest approach

XOR drawing is the minimum viable sprite. It requires no mask data, no background save buffer, and no erase step. You draw the sprite by XORing its pixel data with the screen, and you erase it by XORing the same data again — the property of XOR that $A \oplus B \oplus B = A$ guarantees that the background is restored.

Here is a complete 16x16 XOR sprite routine:

```
; Draw a 16x16 XOR sprite
; Input:  HL = screen address (top-left byte of sprite position)
;         IX = pointer to sprite data (32 bytes: 2 bytes x 16 rows)
;
xor_sprite_16x16:
    ld    b, 16                ; 7 T   16 rows

    .row:
        ld    a, (ix+0)        ; 19 T   left byte of sprite row
```

```

xor  (hl)          ; 7 T   combine with screen
ld   (hl), a       ; 7 T   write back
inc  l             ; 4 T   move right one byte

ld   a, (ix+1)     ; 19 T   right byte of sprite row
xor  (hl)          ; 7 T
ld   (hl), a       ; 7 T   write back
dec  l             ; 4 T   restore column

inc  ix            ; 10 T   \  advance sprite
inc  ix            ; 10 T   /  data pointer

inc  h             ; 4 T   move down one pixel row
ld   a, h          ; 4 T   \
and  7             ; 7 T   | check character
jr   nz, .no_boundary ; 12/7T /  boundary crossing

; Character boundary: adjust HL
ld   a, l          ; 4 T
add  a, 32         ; 7 T
ld   l, a          ; 4 T
jr   c, .no_fix    ; 12/7T
ld   a, h          ; 4 T
sub  8             ; 7 T
ld   h, a          ; 4 T
.no_fix:

.no_boundary:
djnz .row          ; 13 T   (8 on last)
ret               ; 10 T

```

The inner loop costs 134 T-states per row in the common case (no boundary crossing): two IX loads at 19T each, two XOR-and-store sequences at 18T each, two INC IX at 10T, row advance (INC H + check) at 27T, and DJNZ at 13T. At character boundaries, the cost rises to ~164 T-states (the extra adjustment instructions replace the taken JR). For 16 rows: approximately 2,200 T-states per draw. To erase the sprite, call the same routine again with the same screen address — the XOR undoes itself.

Total cost to animate one XOR sprite: ~4,400 T-states per frame (draw + erase).

When XOR works

XOR sprites are perfect for:

- **Cursors.** A blinking text cursor, a crosshair, a selection highlight. Anything that sits on top of a static background and does not need to look pretty.
- **Bullets.** A 2x2 or 4x4 projectile that moves fast enough that visual glitches are invisible.
- **Debug markers.** Plotting collision boxes, entity positions, path nodes.

When XOR fails

XOR has two serious problems. First, the visual quality is poor. Where the sprite overlaps with existing screen data, the pixels invert rather than replacing. A white sprite passing over white text turns invisible. A carefully drawn character becomes a mess of inverted pixels against a detailed background.

Second, XOR gives you no control over attributes. The sprite's colour is whatever ink/paper combination happens to be in the cells it overlaps. For a bullet or cursor this is fine. For a character sprite, it is not.

Despite its limitations, XOR is useful enough that every game programmer should have it in their toolkit. Twenty lines of code, zero extra memory, and it just works.

Method 2: OR+AND Masked Sprites

The industry standard

Almost every commercial Spectrum game released after 1984 used this technique or a close variant. A masked sprite carries two pieces of data for each row: a *mask* and the *graphic*. The mask defines the sprite's shape — which pixels belong to the sprite and which are transparent. The graphic defines the sprite's appearance — which of the shaped pixels are set.

The drawing algorithm for each byte is:

1. Read the screen byte.
2. AND it with the mask. This clears the pixels where the sprite will appear, leaving the rest of the background intact.
3. OR it with the graphic. This stamps the sprite's pixels into the cleared area.

The result: the sprite appears on screen with transparent areas showing the background through. No XOR artifacts. No inverted pixels. Clean, professional-looking sprites.

AND+OR sprite compositing — the mask punches a transparent hole in the background (AND), then sprite data fills it (OR), preserving surrounding pixels at bit level

Figure 35: AND+OR sprite compositing — the mask punches a transparent hole in the background (AND), then sprite data fills it (OR), preserving surrounding pixels at bit level

Data format

For a 16x16 sprite, each row contains 4 bytes: mask-left, graphic-left, mask-right, graphic-right. The mask byte has 1 for transparent pixels and 0 for opaque pixels (because ANDing with 1 preserves the background, ANDing with 0 clears it). Total data per sprite: 16 rows x 4 bytes = 64 bytes.

The inner loop

```

; Draw a 16x16 masked sprite (byte-aligned)
; Input:  HL = screen address
;         DE = pointer to sprite data
;         Format per row: mask_L, gfx_L, mask_R, gfx_R
;
masked_sprite_16x16:
    ld    b, 16                ; 7 T

.row:
    ; --- Left byte ---
    ld    a, (de)              ; 7 T    load mask
    and   (hl)                 ; 7 T    clear sprite-shaped hole in background
    inc   de                   ; 6 T
    ld    c, a                 ; 4 T    save masked background

    ld    a, (de)              ; 7 T    load graphic
    or    c                   ; 4 T    stamp sprite into hole
    ld    (hl), a              ; 7 T    write to screen
    inc   de                   ; 6 T
    inc   l                   ; 4 T    move right

    ; --- Right byte ---
    ld    a, (de)              ; 7 T    load mask
    and   (hl)                 ; 7 T
    inc   de                   ; 6 T
    ld    c, a                 ; 4 T

    ld    a, (de)              ; 7 T    load graphic
    or    c                   ; 4 T
    ld    (hl), a              ; 7 T
    inc   de                   ; 6 T
    dec   l                   ; 4 T    restore column

    ; --- Next row (DOWN_HL) ---
    inc   h                   ; 4 T
    ld    a, h                 ; 4 T
    and   7                   ; 7 T
    jr    nz, .no_boundary    ; 12/7T

    ld    a, l                 ; 4 T
    add   a, 32                ; 7 T
    ld    l, a                 ; 4 T
    jr    c, .no_fix           ; 12/7T
    ld    a, h                 ; 4 T
    sub   8                   ; 7 T
    ld    h, a                 ; 4 T

.no_fix:
.no_boundary:
    djnz .row                 ; 13 T
    ret                      ; 10 T

```

Cycle counting


Let us count the common case (no character boundary crossing). Note that JR NZ is taken (12T) in the common case because the boundary crossing is rare — only 1 in 8 rows crosses a character boundary.

Section	Instructions	T-states
Left byte: mask+draw	ld a,(de) + and (hl) + inc de + ld c,a + ld a,(de) + or c + ld (hl),a + inc de + inc l	52
Right byte: mask+draw	Same sequence + dec l	52
Row advance	inc h + ld a,h + and 7 + jr nz (taken)	27
Loop	djnz	13
Total per row		144

For 16 rows: $16 \times 144 = \mathbf{2,304 \text{ T-states}}$ (common case). Add boundary-crossing overhead for ~2 boundaries in a 16-pixel sprite: roughly **2,400 T-states** total.

But this only draws the sprite. You also need to erase the previous frame's sprite, which means restoring the background — we will address this in Method 6 (Dirty Rectangles). For now, note that the draw alone is about 35% more expensive than XOR, but the visual quality is incomparably better.

Masked Sprite Rendering: AND Clears, OR Stamps



Mask has 0s where sprite pixels go (AND clears them), 1s elsewhere (AND preserves background)

Figure 36: AND mask clears the sprite area on a checkered background, then OR stamps the graphic. The mask has 0-bits where sprite pixels go (AND clears them) and 1-bits elsewhere (AND preserves the background).

Byte alignment and the shift problem

The routine above assumes the sprite starts at a byte boundary — that is, the x coordinate is a multiple of 8. In practice, game characters move pixel by pixel, not in 8-pixel jumps. If your sprite's x position is 53, it starts at byte column 6, pixel 5 within that byte. The sprite data needs to be shifted right by 5 bits.

You can shift at draw time:

```

; Shift mask and graphic right by A bits
; This adds significant cost per byte
    ld    a, (de)          ; 7 T    load mask byte
    ld    c, a              ; 4 T
    ld    a, b              ; 4 T    shift count
.shift:
    srl    c                ; 8 T    \
    dec    a                ; 4 T    | per-bit shift loop
    jr     nz, .shift       ; 12 T    /

```

Each bit of shift costs 24 T-states per byte in this naive loop. For a 5-bit shift on a 16-wide sprite (3 bytes per row after shift, since the sprite spills into a third byte), you are looking at an extra $5 \times 24 \times 3 = 360$ T-states per row — doubling the draw cost. For 8 sprites at 25 fps, this shift overhead alone would consume roughly 46,000 T-states per frame — over 60% of your budget.

This is why pre-shifted sprites exist.

Method 3: Pre-Shifted Sprites

The memory-for-speed trade-off

The idea is simple: instead of shifting the sprite data at draw time, pre-compute shifted versions of the sprite at load time (or at assembly time) and store them alongside the original. When you need to draw the sprite at pixel offset 3 within a byte, you use the version that was pre-shifted by 3 pixels.

There are two common configurations:

4 shifted copies (shift by 0, 2, 4, 6 pixels). This gives 2-pixel horizontal resolution. The sprite snaps to even pixel positions, which is often acceptable for game characters. Memory cost: 4x the unshifted data.

8 shifted copies (shift by 0, 1, 2, 3, 4, 5, 6, 7 pixels). Full pixel-level horizontal positioning. Memory cost: 8x the unshifted data. But each shifted version is also wider: a 16-pixel-wide sprite shifted by 1–7 bits spills into a third byte column, so each shifted copy is 3 bytes wide instead of 2.

Memory calculation

For a 16x16 masked sprite:

Configuration	Bytes per row	Rows	Copies	Total
Unshifted only	4 (mask+gfx x 2 bytes)	16	1	64
4 shifts	4	16	4	256
8 shifts (3 bytes wide)	6 (mask+gfx x 3 bytes)	16	8	768

For a sprite with 4 animation frames, multiply by 4:

Configuration	Per frame	4 frames	8 sprites
Unshifted + runtime shift	64	256	2,048
4 pre-shifts	256	1,024	8,192
8 pre-shifts	768	3,072	24,576

24 KB for 8 sprites with full pre-shifting. On a 128K Spectrum, that is one and a half memory banks just for sprite data. On a 48K machine, it is nearly half of available RAM. The trade-off is stark.

Practical compromise

Most games use 4 pre-shifted copies. The 2-pixel horizontal resolution is barely noticeable in gameplay. Some games use 8 copies for the player character (where smooth movement matters most) and 4 copies or even runtime shifting for less important sprites.

The draw routine for pre-shifted sprites is identical to the byte-aligned masked sprite routine — you just select the correct pre-shifted data set before calling it:

```
; Select pre-shifted sprite data
; Input:  A = x coordinate (0-255)
;         IY = base of pre-shift table (4 entries, each pointing to 16-row data)
; Output: DE = pointer to correct shifted sprite data
;
select_preshift:
    and $06                ; 7 T    mask to shifts 0,2,4,6 (4 copies)
    ld  c, a                ; 4 T
    ld  b, 0                ; 7 T
    add iy, bc              ; 15 T
    ld  e, (iy+0)           ; 19 T
    ld  d, (iy+1)           ; 19 T    DE = pointer to shifted data
    ret
```

The draw time is the same as Method 2: ~2,300 T-states. But you have eliminated the per-pixel shift cost entirely. The price is paid in memory, not T-states.

Method 4: Stack Sprites (The PUSH Method)

The fastest output on the Z80

We saw in Chapter 3 that PUSH writes 2 bytes in 11 T-states — 5.5 T-states per byte, the fastest write operation on the Z80. Stack sprites exploit this for sprite output: set SP to the bottom of the sprite's screen area, load register pairs with sprite data, and PUSH them onto the screen.

The technique requires a critical setup:

1. **DI** — disable interrupts. If an interrupt fires while SP points to the screen, the CPU pushes its return address into your pixel data, corrupting the display and probably crashing.

2. **Save SP** — stash the real stack pointer using self-modifying code.
3. **Set SP** to the bottom-right of the sprite area on screen (PUSH works downward).
4. **Load and PUSH** — load sprite data into register pairs and PUSH them in sequence.
5. **Restore SP and EI** — put the stack back and re-enable interrupts.

The inner loop

For a 16x16 sprite (2 bytes wide), each row is a single PUSH:

```
; Stack sprite: 16x16, writes 2 bytes per row via PUSH
; Input:  screen_addr = pre-calculated bottom-right screen address
;         sprite_data = 32 bytes of pixel data (2 bytes x 16 rows,
;         stored bottom-to-top because PUSH goes downward)
;
stack_sprite_16x16:
    di                                ; 4 T

    ld  (restore_sp + 1), sp          ; 20 T  save SP (self-mod)

    ld  sp, (screen_addr)              ; 20 T  SP = bottom of sprite on screen
    ld  ix, sprite_data                ; 14 T

    ; Row 15 (bottom) - each PUSH writes 2 bytes and decrements SP
    ld  h, (ix+31)                     ; 19 T  \
    ld  l, (ix+30)                     ; 19 T  | load bottom row
    push hl                            ; 11 T  / write to screen

    ; But wait --- SP just decremented by 2, and the next row UP
    ; on the Spectrum screen is NOT at SP-2. The interleaved layout
    ; means "one row up" is at a completely different address.
    ;
    ; This is the fundamental problem with stack sprites on the
    ; Spectrum: the screen is not contiguous in memory.
```

And here is the fundamental difficulty. The PUSH method is the fastest possible write, but the Spectrum's interleaved screen layout means that consecutive screen rows are not at consecutive addresses. SP decrements linearly, but screen rows follow the 010TTSSS LLLCCCCC pattern from Chapter 2.

Making it work: pre-calculated SP chain

The solution is to not rely on SP auto-decrement for row navigation. Instead, you explicitly set SP for each row:

```
; Stack sprite: 16x16 with explicit SP per row
; This is the practical version --- each row gets SP set independently
;
stack_sprite_16x16:
    di                                ; 4 T
    ld  (restore_sp + 1), sp          ; 20 T
```



```

ld    hl, (sprite_data + 0)    ; 16 T    row 0 data
ld    sp, (row_addrs + 0)      ; 20 T    SP = screen addr for row 0 + 2
push hl                        ; 11 T    write row 0
                                   ;      total per row: 47 T

ld    hl, (sprite_data + 2)    ; 16 T    row 1 data
ld    sp, (row_addrs + 2)      ; 20 T
push hl                        ; 11 T
; ... repeated for all 16 rows ...

restore_sp:
ld    sp, $0000                ; 10 T    self-modified
ei                                ; 4 T
ret                                ; 10 T

```

Each row costs 47 T-states. For 16 rows: $16 \times 47 = 752$ T-states, plus setup and teardown (~60 T-states). Total: approximately **810 T-states**.

Compare this to Method 2's ~2,300 T-states. The stack sprite is nearly 3x faster — but it comes with constraints.

The costs

No masking. The PUSH writes 2 bytes unconditionally. It overwrites whatever was on screen. There is no AND-with-mask step. The sprite is always a solid rectangle — any “transparent” pixels will show the sprite’s background colour, not the game background. For sprites on a solid-colour background (many classic Spectrum games used black), this is fine. For sprites over detailed backgrounds, it is not.

Pre-calculated row addresses. You need a table of screen addresses for all 16 rows of the sprite, updated whenever the sprite moves. This is a per-frame setup cost — not enormous, but not free.

Interrupts are off. For 8 sprites, roughly 6,500 T-states with interrupts disabled. If your music runs from IM2, schedule sprite drawing immediately after HALT.

Data must be stored in PUSH order. Since PUSH writes the high byte at (SP-1) and the low byte at (SP-2), and SP decrements *before* writing, the data layout requires careful attention. The sprite data is stored reversed: the rightmost byte of a row becomes the low byte loaded into the register, the leftmost becomes the high byte.

When to use stack sprites

Stack sprites are the weapon of choice when you need raw speed and your background is simple enough that full-rectangle overwrites are acceptable. Arcade-style games with black backgrounds, score overlays, and fast-moving objects are the natural fit. The PUSH method is also used for screen clearing and bulk data output (Chapter 3), where the “no masking” limitation is irrelevant.

Method 5: Compiled Sprites

The sprite is the code

A compiled sprite takes the code-generation philosophy from Chapter 3 to its logical conclusion. Instead of a data table interpreted by a draw routine, the sprite *is* an executable routine. Each visible pixel byte in the sprite becomes a LD (HL), n instruction. Transparent runs become INC L or INC H instructions to skip over them. The entire sprite is rendered by CALL-ing it.

How it works

Consider a simple 8x8 sprite with some transparent pixels. In a masked sprite, you would store mask+graphic pairs and run the AND/OR loop. In a compiled sprite, you generate Z80 instructions at assembly time (or at load time):

```
; Compiled sprite for a small arrow shape
; Input: HL = screen address of top-left byte
; The sprite draws itself.
;
compiled_arrow:
    ; Row 0: one pixel byte
    ld  (hl), $18          ; 10 T  ..##....
    inc h                  ; 4 T   next row

    ; Row 1: one pixel byte
    ld  (hl), $3C          ; 10 T  ..####..
    inc h                  ; 4 T

    ; Row 2: one pixel byte
    ld  (hl), $7E          ; 10 T  .#####.
    inc h                  ; 4 T

    ; Row 3: one pixel byte
    ld  (hl), $FF          ; 10 T  #####
    inc h                  ; 4 T

    ; Row 4: one pixel byte
    ld  (hl), $3C          ; 10 T  ..####..
    inc h                  ; 4 T

    ; Row 5: one pixel byte
    ld  (hl), $3C          ; 10 T  ..####..
    inc h                  ; 4 T

    ; Row 6: one pixel byte
    ld  (hl), $3C          ; 10 T  ..####..
    inc h                  ; 4 T

    ; Row 7: one pixel byte
    ld  (hl), $3C          ; 10 T  ..####..
```

```

ret                                ; 10 T

; Total: 8 x (10 + 4) + 10 = 122 T-states
; Compare: masked routine for 8x8 = ~600 T-states

```

That is 122 T-states for an 8x8 sprite. The masked approach takes roughly 5x longer.

16x16 compiled sprite

For a wider sprite, each row may have multiple LD (HL), n instructions separated by INC L:

```

; Compiled sprite: 16x16 (2 bytes wide)
; Input: HL = screen address of top-left
;
compiled_sprite_16x16:
    ; Row 0
    ld  (hl), $3C          ; 10 T   left byte
    inc l                  ; 4 T
    ld  (hl), $0F          ; 10 T   right byte
    dec l                  ; 4 T   restore column
    inc h                  ; 4 T   next row
                                ;      row cost: 32 T

    ; Row 1
    ld  (hl), $7E          ; 10 T
    inc l                  ; 4 T
    ld  (hl), $1F          ; 10 T
    dec l                  ; 4 T
    inc h                  ; 4 T
                                ;      row cost: 32 T

    ; ... rows 2-6 similar ...

    ; Row 7 (character boundary)
    ld  (hl), $FF          ; 10 T
    inc l                  ; 4 T
    ld  (hl), $FF          ; 10 T
    dec l                  ; 4 T
    ; Character boundary crossing:
    ld  a, l               ; 4 T
    add a, 32              ; 7 T
    ld  l, a               ; 4 T
    ld  a, h               ; 4 T
    sub 8                  ; 7 T
    ld  h, a               ; 4 T   boundary cost: 30 T
    inc h                  ; 4 T
                                ;      row cost: 62 T

    ; Rows 8-15 similar to 0-6, with another boundary at row 15
    ; ...

```

```
ret                                ; 10 T
```

Per row (common case): 32 T-states. For 16 rows with 1-2 boundary crossings: roughly **570 T-states**.

The trade-offs

Strengths: - The fastest masked-capable sprite method. You can incorporate AND masking into compiled sprites — each byte becomes `ld a,(hl) / and mask / or graphic / ld (hl),a` instead of a simple `ld (hl),n`. Even with masking, the compiled approach avoids loop overhead, data pointer management, and row counting. - No loop overhead at all. The code is completely straight-line. - Transparent regions cost nothing if they span entire bytes — you just skip them with `INC L` or `INC H`.

Weaknesses: - **Code size.** Each visible byte takes 2 bytes of code (`LD (HL), n`). With masking (4 instructions per byte), code size roughly triples. A full set of 8 pre-shifted compiled sprites with 4 animation frames can reach several kilobytes per sprite. - **No runtime data changes.** Pixel values are baked into instruction operands. Animation requires a separate compiled routine per frame. - **Boundary handling is baked in.** Character boundary crossings sit at fixed positions, so the sprite must maintain consistent vertical alignment or you need multiple compiled versions.

Compiled sprites with masking

For sprites that need to appear over a detailed background, you compile the mask into the code:

```
; Compiled sprite with masking: one byte
; Instead of ld (hl),n, we do:
ld  a, (hl)           ; 7 T  read screen
and $C3               ; 7 T  mask: clear sprite pixels
or  $3C               ; 7 T  graphic: stamp sprite
ld  (hl), a           ; 7 T  write back
                        ;      per-byte cost: 28 T
```

28 T-states per byte, versus 52 T-states per byte in the general-purpose masked routine (Method 2). The saving comes from eliminating pointer management, loop counting, and data loading — the mask and graphic values are immediate operands.

For 16 rows x 2 bytes: $16 \times (28 + 28 + 4 + 4 + 4) = 16 \times 68 =$ **1,088 T-states**. This is about half the cost of the generic masked routine, with full transparency support.

Method	16x16 draw cost	Masking	Notes
XOR sprite	~2,200 T	No	Draw + erase = ~4,400 T
OR+AND masked	~2,400 T	Yes	Standard approach
Pre-shifted masked	~2,400 T	Yes	No shift cost; 4-8x memory
Stack sprite (PUSH)	~810 T	No	DI required; solid rectangle
Compiled (no mask)	~570 T	No	Code = sprite; large footprint
Compiled (masked)	~1,088 T	Yes	Best of both; largest footprint




Figure 37: Sprite rendering methods comparison

Method 6: Dirty Rectangles

The background problem

Methods 1–5 all address the question of *putting* pixels on screen. But sprites move. Every frame, the sprite is in a new position. Before drawing the sprite at its new position, you must erase it from the old position — or the screen fills up with ghostly after-images.

The XOR method handles this implicitly: XOR the old position to erase, XOR the new position to draw. But for all other methods, you need a way to restore the background.

There are three common approaches:

Full screen clear. Wipe the pixel area every frame (~36,000 T-states with PUSH from Chapter 3), then redraw everything. Feasible but expensive.

Background save/restore. Before drawing each sprite, save the screen behind it. To erase, copy the saved buffer back. Cost is $O(\text{sprite_size})$ per sprite, not $O(\text{screen_size})$.

Dirty rectangle tracking. A refinement: track which rectangles were modified, restore only those, then draw new sprites (saving new background as you go).

The save/restore cycle

The practical approach for most Spectrum games is per-sprite background save/restore. Here is the cycle for one sprite per frame:

1. **Restore** the background saved last frame (copy saved buffer to old screen position).
2. **Save** the background at the new screen position (copy screen to save buffer).
3. **Draw** the sprite at the new position.

The order matters. You restore before saving to avoid overwriting the new background save with stale data if sprites overlap.

Save/restore routine

For a 16x16 sprite (2 bytes wide, 16 rows), the background buffer is 32 bytes:

```
; Save background behind a 16x16 sprite
; Input:  HL = screen address (top-left)
;         DE = pointer to save buffer (32 bytes)
;
save_background_16x16:
    ld    b, 16                ; 7 T

.row:
    ld    a, (hl)              ; 7 T    read left byte
    ld    (de), a              ; 7 T    save it
    inc   de                   ; 6 T
    inc   l                    ; 4 T

    ld    a, (hl)              ; 7 T    read right byte
    ld    (de), a              ; 7 T    save it
    inc   de                   ; 6 T
    dec   l                    ; 4 T

    ; DOWN_HL (same as sprite routines)
    inc   h                    ; 4 T
    ld    a, h                 ; 4 T
    and   7                    ; 7 T
    jr    nz, .no_boundary    ; 12/7T
    ld    a, l                 ; 4 T
    add   a, 32                ; 7 T
    ld    l, a                 ; 4 T
    jr    c, .no_fix           ; 12/7T
    ld    a, h                 ; 4 T
    sub   8                    ; 7 T
    ld    h, a                 ; 4 T
.no_fix:
.no_boundary:
    djnz  .row                 ; 13 T
    ret                        ; 10 T
```

The restore routine is identical with source and destination swapped: read from the buffer, write to the screen. Each routine takes approximately **1,500 T-states** for 16 rows.

Full frame budget

Let us calculate the per-frame cost for 8 animated 16x16 sprites using OR+AND masking with background save/restore:

Operation	Per sprite	8 sprites
Restore background	~1,500 T	12,000 T
Save new background	~1,500 T	12,000 T
Draw sprite (masked)	~2,400 T	19,200 T
Total	~5,400 T	~43,200 T

On a Pentagon (71,680 T-states per frame): 43,200 T leaves **28,480 T** for game logic, input processing, music, and everything else. At 25 fps you have twice the budget (two frames per update), giving ~100,000 T-states for non-sprite work. This is comfortable for a game.

If you use compiled masked sprites instead:

Operation	Per sprite	8 sprites
Restore background	~1,500 T	12,000 T
Save new background	~1,500 T	12,000 T
Draw sprite (compiled, masked)	~1,088 T	8,704 T
Total	~4,088 T	~32,704 T

That saves over 10,000 T-states per frame — a meaningful improvement that buys you more room for game logic or more sprites.

Draw order and overlap

When multiple sprites overlap, draw order matters. The simplest correct approach:

1. Restore all backgrounds (in reverse draw order, to handle overlaps correctly).
2. Save all new backgrounds.
3. Draw all sprites.

Reverse-order restoration ensures that when two sprites overlapped last frame, the earlier sprite's save buffer (which captured the clean background) is restored last, correctly cleaning up the overlap area.

The reasoning: if sprite A was drawn on top of sprite B, A's save buffer contains B's pixels. Restoring A first exposes B, then restoring B reveals the clean background. Forward-order restoration leaves artifacts. Many games sidestep this by preventing overlap or accepting minor glitches.

Optimising the Inner Loops

Eliminating pointer management

The routines above spend significant time on pointer management: `inc de`, `inc l`, `dec l`, and the `DOWN_HL` boundary logic. Several optimisations can reduce this

overhead.

Use LDI instead of manual copy. For save/restore operations, an LDI chain (Chapter 3) copies a byte from (HL) to (DE), increments both, and decrements BC — all in 16 T-states. Compared to our manual `ld a,(hl) + ld (de),a + inc de + inc l` at 24 T-states, LDI saves 8 T-states per byte. For a 16x16 sprite (32 bytes), that is 256 T-states saved per save or restore operation.

```
; Save background using LDI (partial unroll, 2 bytes per row)
; HL = screen address, DE = save buffer
;
save_bg_ldi:
    ld    b, 16                ; 7 T
.row:
    ldi                    ; 16 T    copy left byte
    ldi                    ; 16 T    copy right byte
    dec    l                  ; 4 T    \
    dec    l                  ; 4 T    /  LDI advanced L by 2, undo it

    ; DOWN_HL
    inc    h                  ; 4 T
    ld     a, h               ; 4 T
    and    7                  ; 7 T
    jr     nz, .no_boundary   ; 12/7T
    ld     a, l               ; 4 T
    add    a, 32              ; 7 T
    ld     l, a               ; 4 T
    jr     c, .no_fix         ; 12/7T
    ld     a, h               ; 4 T
    sub    8                  ; 7 T
    ld     h, a               ; 4 T
.no_fix:
.no_boundary:
    djnz   .row               ; 13 T
    ret                                ; 10 T
```

Common-case row cost: $16 + 16 + 4 + 4 + 4 + 4 + 7 + 12 + 13 = \mathbf{80 \text{ T-states}}$ (JR NZ is taken at 12T in the common case — no boundary crossing). For 16 rows: approximately **1,280 T-states** — a worthwhile improvement over the 1,500 T-states of the manual approach.

Combine save and draw. Instead of save-then-draw as two separate passes over the screen area, combine them into one pass: for each byte, read the screen (save it), then write the sprite data. This halves the number of row-advance operations and eliminates one complete DOWN_HL traversal:

```
; Combined save-and-draw for masked sprite
; HL = screen address, DE = sprite data (mask, gfx pairs)
; IX = save buffer
;
save_and_draw_16x16:
    ld    b, 16                ; 7 T
.row:
    ; Left byte
```



```

ld  a, (hl)          ; 7 T   read screen (for saving)
ld  (ix+0), a        ; 19 T   save to buffer
ld  c, a             ; 4 T

ld  a, (de)          ; 7 T   load mask
and c                ; 4 T   mask background
inc de               ; 6 T
ld  c, a             ; 4 T

ld  a, (de)          ; 7 T   load graphic
or  c                ; 4 T   stamp sprite
ld  (hl), a          ; 7 T   write to screen
inc de               ; 6 T
inc l                ; 4 T

; Right byte (similar)
ld  a, (hl)          ; 7 T
ld  (ix+1), a        ; 19 T
ld  c, a             ; 4 T

ld  a, (de)          ; 7 T
and c                ; 4 T
inc de               ; 6 T
ld  c, a             ; 4 T

ld  a, (de)          ; 7 T
or  c                ; 4 T
ld  (hl), a          ; 7 T
inc de               ; 6 T
dec l                ; 4 T

; Advance IX and HL
inc ix               ; 10 T
inc ix               ; 10 T

inc h                ; 4 T
ld  a, h             ; 4 T
and 7                ; 7 T
jr  nz, .no_boundary ; 12/7T
ld  a, l             ; 4 T
add a, 32            ; 7 T
ld  l, a             ; 4 T
jr  c, .no_fix       ; 12/7T
ld  a, h             ; 4 T
sub 8                ; 7 T
ld  h, a             ; 4 T
.no_fix:
.no_boundary:
djnz .row            ; 13 T
ret                  ; 10 T

```

This combines save and draw into a single pass. The cost per row (common case):

roughly **205 T-states** (JR NZ taken at 12T). For 16 rows: approximately **3,400 T-states** — compared to separate save (~1,280 T) + draw (~2,400 T) = 3,680 T-states. The saving is modest (~280 T per sprite), but it adds up across 8 sprites.

For maximum performance, unroll the entire routine: no DJNZ loop, explicit per-row code with boundary crossings baked in at rows 7 and 15. This eliminates loop and boundary-test overhead, bringing the total to roughly **2,780 T-states** at the cost of ~300 bytes of code per sprite routine.

Agon Light 2: Hardware VDP Sprites

The Agon Light 2 takes a fundamentally different approach. The eZ80 communicates with a VDP (Video Display Processor) on an ESP32, which handles all sprite rendering in hardware. The CPU uploads bitmaps, then issues commands to position, show, hide, and animate sprites. The VDP composites sprites during its own rendering pass, with no CPU cost per pixel. Up to 256 sprite slots are supported.

The VDU command sequence for defining and activating a sprite:

```
VDU 23, 27, 1, w, h, format          ; Create sprite: w x h pixels
; ... upload bitmap data ...
VDU 23, 27, 4, x_lo, x_hi, y_lo, y_hi ; Set position
VDU 23, 27, 11                        ; Show sprite
```

In eZ80 assembly, these commands are sent as byte sequences to the VDP via RST \$10 (MOS character output). Each command is a sequence of ld a, byte : rst \$10 pairs.

Moving a sprite

Once defined, moving a sprite is just the position command:

```
; Agon: Move sprite 0 to (x, y)
; Input: BC = x, DE = y
;
move_sprite:
    ld  a, 23 : rst $10
    ld  a, 27 : rst $10
    ld  a, 4  : rst $10    ; move command
    ld  a, 0  : rst $10    ; sprite number

    ld  a, c  : rst $10    ; x low
    ld  a, b  : rst $10    ; x high
    ld  a, e  : rst $10    ; y low
    ld  a, d  : rst $10    ; y high
    ret
```

The CPU cost for moving a sprite is just the cost of sending ~10 bytes over the serial interface. At 1,152,000 baud, each byte takes roughly 9 microseconds, so moving one sprite takes approximately 90 microseconds — about 1,660 T-states at 18.432 MHz. Moving 8 sprites: ~13,000 T-states. The VDP handles all pixel compositing, transparency, and background management in hardware.

Scanline limits

The VDP has a practical limit on sprite pixels per horizontal scanline. When too many sprites overlap on the same line, some may flicker — the same phenomenon seen on the NES and Master System. A reasonable guideline is 8 to 12 16x16 sprites per scanline. For 8 sprites distributed across the screen, you are unlikely to hit this limit.

The trade-off

The Agon eliminates the entire sprite-drawing problem. No save/restore, no masking, no interleaved screen navigation. The cost is abstraction: no per-pixel tricks, no creative data manipulation, and dependency on VDP firmware capabilities. The Spectrum forces you to build everything from nothing. The Agon frees you to spend that effort on game design.

Practical: 8 Animated Sprites at 25 fps

Spectrum implementation

Our target: 8 animated 16x16 sprites with background save/restore, running at 25 fps (updating every 2 frames) on a ZX Spectrum 128K.

Architecture:

Each sprite has a data structure:

```
; Sprite structure (12 bytes per sprite)
;
SPRITE_X      EQU 0      ; x coordinate (0-255)
SPRITE_Y      EQU 1      ; y coordinate (0-191)
SPRITE_OLD_X  EQU 2      ; previous x (for erase)
SPRITE_OLD_Y  EQU 3      ; previous y
SPRITE_FRAME  EQU 4      ; current animation frame (0-3)
SPRITE_DIR    EQU 5      ; direction / flags
SPRITE_DX     EQU 6      ; x velocity (signed)
SPRITE_DY     EQU 7      ; y velocity (signed)
SPRITE_GFX    EQU 8      ; pointer to sprite graphic data (2 bytes)
SPRITE_SAVE   EQU 10     ; pointer to background save buffer (2 bytes)

SPRITE_SIZE   EQU 12
NUM_SPRITES   EQU 8
```

Frame cycle (every 2 VBLANKs):

```
main_loop:
    halt                ; wait for VBLANK
    halt                ; wait again (25 fps = every 2nd frame)

    ; Phase 1: Restore all backgrounds (reverse order)
    ld    ix, sprites + (NUM_SPRITES - 1) * SPRITE_SIZE
    ld    b, NUM_SPRITES
```

```

.restore_loop:
    call restore_sprite_bg
    ld    de, -SPRITE_SIZE
    add   ix, de
    djnz  .restore_loop

    ; Phase 2: Update positions
    ld    ix, sprites
    ld    b, NUM_SPRITES
.update_loop:
    call update_sprite_position
    ld    de, SPRITE_SIZE
    add   ix, de
    djnz  .update_loop

    ; Phase 3: Save backgrounds and draw (forward order)
    ld    ix, sprites
    ld    b, NUM_SPRITES
.draw_loop:
    call save_and_draw_sprite
    ld    de, SPRITE_SIZE
    add   ix, de
    djnz  .draw_loop

    ; Phase 4: Game logic, input, sound
    call process_input
    call update_game_logic
    call update_sound

    jr    main_loop

```

OR+AND masked sprite with movement, showing eight animated sprites over a patterned background

Figure 38: OR+AND masked sprite with movement, showing eight animated sprites over a patterned background

Cycle budget:

Phase	Cost
2 x HALT	0 T (waiting)
Restore 8 backgrounds	8 x 1,280 = 10,240 T
Update 8 positions	8 x 200 = 1,600 T
Save + Draw 8 sprites	8 x 3,400 = 27,200 T
Loop overhead	~2,000 T
Total sprite work	~41,040 T
Available for game logic	~102,000 T

With a 2-frame budget of 143,360 T-states (2 x 71,680 on Pentagon), we have roughly 102,000 T-states for game logic, input, and sound. This is generous —

enough for entity AI (Chapter 19), tile collision detection, music playback, and input processing.

Contention and sprite rendering (48K/128K Spectrum). The budget above assumes Pentagon timing with no wait states. On a standard 48K or 128K Spectrum, writes to video RAM (\$4000-\$5AFF for pixels, \$5800-\$5AFF for attributes) are contended — the ULA steals cycles during the active display period. Each PUSH in a stack sprite or each LD (HL),A in a masked sprite costs roughly 1 extra T-state per access to this region, adding up to ~1.3 T per 2-byte PUSH. For an 8-sprite system, this can add 3,000–5,000 T-states of overhead. **Practical rule:** on non-Pentagon machines, schedule sprite drawing during the border period (top/bottom border give ~14,000 T-states of contention-free time) or accept a 2-frame budget with wider margins. See Chapter 15.2 for the full anatomy of ULA contention patterns and mitigation strategies.

Before drawing each sprite, calculate the screen address from (x, y) using the routine from Chapter 2, and select the correct pre-shifted data based on x AND \$06 (for 4 shift levels). The pre-shift selection logic from Method 3 applies directly.

Agon implementation

On the Agon, the main loop becomes trivially simple: wait for VSync, update positions, send VDU 23,27,4 move commands for each sprite, and proceed to game logic. No save/restore, no masking, no screen address calculation, no interleaved row navigation. The VDP handles everything.

The contrast is instructive. On the Spectrum, sprite rendering is the dominant cost — over 40,000 T-states per frame, where every cycle saved in the inner loop translates directly to more sprites or more game logic. On the Agon, sprite rendering is effectively free from the CPU's perspective, and your engineering effort goes into game design rather than pixel-pushing. Both approaches have their satisfactions.

Summary

- **XOR sprites** are the simplest method: XOR to draw, XOR again to erase. ~2,200 T-states to draw a 16x16 sprite. No mask, no background save needed. Visual quality is poor (inverted pixels over background detail). Good for cursors, bullets, and debug markers.
- **OR+AND masked sprites** are the industry standard. Each byte goes through an AND-with-mask, OR-with-graphic sequence that produces clean transparency. ~2,400 T-states for a 16x16 sprite. This is what most commercial Spectrum games use.
- **Pre-shifted sprites** eliminate the per-pixel shift cost by storing 4 or 8 pre-computed shifted copies of the sprite data. Draw time is the same as the masked routine. Memory cost scales from 4x (4 shifts, 2-pixel resolution) to 8x (8 shifts, full pixel resolution). The standard memory-vs-speed trade-off.
- **Stack sprites (PUSH method)** are the fastest raw output: ~810 T-states for a 16x16 sprite. Requires DI/EI, explicit SP management per row, and produces

solid rectangles (no masking). Best for games with simple backgrounds.

- **Compiled sprites** turn the sprite into executable code. Each pixel byte becomes a LD (HL),n instruction. ~570 T-states without masking, ~1,088 T-states with compiled masking. The fastest masked method, at the cost of large code size. Animation requires separate compiled routines per frame.
 - **Dirty rectangles** with background save/restore are the standard technique for sprite animation. Save the background before drawing, restore it before drawing the next frame. Restore in reverse draw order to handle overlapping sprites correctly. The combined save-and-draw approach reduces per-sprite cost to ~3,400 T-states.
 - **8 sprites at 25 fps** on a Spectrum 128K costs approximately 41,000 T-states per update cycle (every 2 frames), leaving ~102,000 T-states for game logic — a comfortable budget for a real game.
 - **Agon Light 2 hardware sprites** eliminate the entire rendering problem. Define sprites once, move them with VDU commands. CPU cost is negligible. The trade-off is abstraction: you gain performance but lose the ability to do per-pixel tricks with sprite data.
 - The choice of sprite method depends on your game's requirements: background complexity, sprite count, animation needs, available memory, and target frame rate. Most Spectrum games use OR+AND masking with pre-shifting and dirty rectangles. Demoscene productions and performance-critical games reach for compiled sprites or the PUSH method.
-

Try It Yourself

1. **Implement all six methods.** Take a simple 8x8 sprite design and implement XOR, masked, and compiled versions. Use the border-colour timing harness from Chapter 1 to compare their draw costs. The difference should be clearly visible.
2. **Pre-shift generator.** Write a utility (in Python, Processing, or Z80 assembly) that takes a masked sprite and generates 4 pre-shifted versions. Store them in memory and write a draw routine that selects the correct version based on x-coordinate.
3. **Background save/restore demo.** Put a masked sprite on a patterned background (the checkerboard from Chapter 2's practical). Move the sprite with keyboard input. Verify that the background is correctly restored as the sprite moves. Then add a second sprite and verify that overlapping areas are handled correctly.
4. **The 8-sprite challenge.** Implement the full 8-sprite system with background save/restore and animation. Start with the OR+AND masked approach. Measure the cycle budget with border colours. If you have room, switch to compiled masked sprites and measure the improvement.
5. **Agon comparison.** If you have an Agon Light 2, implement the same 8-sprite animation using VDP hardware sprites. Compare the code complexity and the CPU budget available for game logic.

Sources: Spectrum graphics programming folklore, widely documented across the ZX community; Chapter 3 of this book for PUSH tricks and self-modifying code; Chapter 2 for screen layout and DOWN_HL; Agon Light 2 VDP documentation (Quark firmware, FabGL sprite API); the game-dev chapters of *book-plan.md* for the six-method framework and practical targets.

Chapter 17: Scrolling

“The screen is 256 pixels wide. The level is 8,000. Somehow, the player must move through it.”

Every side-scrolling game needs to move the world. The player runs right, the background shifts left. It looks simple. On hardware with a scroll register – the NES, the Mega Drive, the Agon Light 2 – it *is* simple: write an offset, and the hardware does the rest. On the ZX Spectrum, there is no scroll register. There is no hardware assistance of any kind. To scroll the screen, you move the bytes yourself. All 6,144 of them.

This chapter works through every practical scrolling method on the Spectrum, from the cheapest to the most expensive: attribute scrolling (768 bytes, trivial), vertical pixel scrolling (tricky because of the interleaved memory layout from Chapter 2), horizontal pixel scrolling (expensive – every byte in every row must be shifted), and the combined method that real games use to get smooth horizontal scrolling within a manageable budget. We will count every T-state, build comparison tables, and show how the shadow screen trick on the 128K makes the whole thing tearfree.

Then we will look at how the Agon Light 2 handles the same problem with hardware scroll offsets and tilemap support – a useful contrast that shows what “the same ISA with different hardware” really means in practice.

The Budget

Before we write a single instruction, let us establish what we are working with.

On a Pentagon (the timing model most Spectrum demos and games target), one frame is **71,680 T-states**. On a standard 48K/128K Spectrum, it is 69,888. We will use the Pentagon figure throughout, but the analysis applies to both – the difference is about 2.5%.

A full-screen scroll means moving data across all 6,144 bytes of pixel memory (and possibly the 768 bytes of attribute memory). The question is always the same: can we do it in one frame, and if so, how many T-states are left for everything else – game logic, sprite drawing, music, input?

Here is the raw cost of just *touching* every byte in the pixel area, using different methods:

Method	Per byte	6,144 bytes	% of frame
ldir	21 T	129,019 T	180%
ldi chain	16 T	98,304 T	137%
ld a,(hl) + ld (de),a + inc hl + inc de push (2 bytes)	24 T 5.5 T/byte	147,456 T 33,792 T	206% 47%

The first three methods cannot move the full pixel area in a single frame. Even ldi chains, the fastest copying method after PUSH, blow the budget by 37%. And scrolling is not just copying – horizontal scrolling requires a *shift* operation on every byte, which adds to the per-byte cost.

This is why scrolling on the Spectrum is a design problem, not just a coding problem. You cannot brute-force a full-screen pixel scroll at 50fps. You must choose your method based on what your game can afford.

Vertical Pixel Scrolling

Vertical scrolling moves the screen contents up or down by one or more pixel rows. Conceptually it is simple: copy each row to the position of the row above (for scrolling up) or below (for scrolling down). On a linear framebuffer, this would be a single block copy. On the Spectrum, the interleaved memory layout (Chapter 2) makes it considerably more interesting.

The Interleave Problem

Recall the screen address structure from Chapter 2:

Low byte: L L L C C C C C

Where TT = third (0-2), SSS = scanline within character cell (0-7), LLL = character row within third (0-7), CCCCC = column byte (0-31).

To scroll up by one pixel, you need to copy the contents of row N to row N-1, for every row from 1 to 191. The source and destination addresses for adjacent pixel rows are *not* separated by a constant offset. Within a character cell, consecutive rows differ by \$0100 in the high byte (just INC H / DEC H). But at character cell boundaries – every 8th row – the relationship changes: you must adjust L by 32 and reset the scanline bits in H. At third boundaries (every 64th row), the adjustment is different again.

The Algorithm: Scroll Up by One Pixel

The approach that works with the interleave, rather than against it, uses the split-counter structure from Chapter 2. Maintain two pointers (source and destination) and advance both using the screen's natural hierarchy: 3 thirds, 8 character rows per third, 8 scanlines per character row. Within each character cell, moving between scanlines is just INC H / INC D on the source and destination. At character

row boundaries, reset the scanline bits and add 32 to L. At third boundaries, add 8 to H and reset L. The inner loop copies 32 bytes per row with LDIR or an LDI chain, and the pointer advancement is folded into the outer loop structure.

Cost Analysis

For each of the 191 row copies, we must copy 32 bytes from source to destination. Using LDIR:

- Per row: $32 \text{ bytes} \times 21 \text{ T-states} - 5 = 667 \text{ T-states}$ for the LDIR, plus pointer management overhead.
- Pointer management (save/restore source and dest, advance scanline): approximately 60 T-states per row within a character cell, more at boundaries.

Total with LDIR: approximately 143,000 T-states. That is roughly **two full frames**. A vertical pixel scroll by one row using LDIR does not fit in a single frame.

We can do better. Replace the LDIR with an LDI chain – 32 LDI instructions per row:

- Per row: $32 \times 16 = 512 \text{ T-states}$ for the LDIs, plus $\sim 50 \text{ T-states}$ of pointer management.
- Total: $191 \times 562 = \mathbf{107,342 \text{ T-states}}$. Still over budget by about 50%.

The PUSH trick is awkward here because we need to copy *between* two non-contiguous areas with a non-constant relationship. PUSH writes to contiguous descending addresses, which does not match the interleaved source/destination pattern.

Partial Scrolling: The Practical Approach

The reality is that most games do not scroll the entire 192-line display. A typical game reserves:

- Top 2 character rows (16 pixels) for a status bar – not scrolled.
- Bottom 1 character row (8 pixels) for a score line – not scrolled.
- Middle: 21 character rows = 168 pixel rows = the scrolling play area.

168 rows of vertical pixel scrolling with LDI chains: $168 \times 562 = \mathbf{94,416 \text{ T-states}}$, or 132% of a frame. Still too much for a single frame if you want time for anything else.

This is why pure vertical pixel scrolling at 1px/frame is rare in Spectrum games. The common approaches are:

1. **Scroll by 8 pixels (one character row):** Move the attributes and the character-aligned pixel data. This is much cheaper because you only copy 21 character rows $\times$ 8 scanlines = 168 rows, but you can use a block copy trick: within each third, the character rows are contiguous in blocks. Cost: around 40,000–50,000 T-states with LDIR. Feasible.
2. **Scroll by 1 pixel using a counter:** Scroll 1px per frame visually by combining a character-level scroll (cheap, every 8 frames) with a pixel offset counter (draw new content at an offset within the 8px character cell). We will cover this combined approach in the horizontal scrolling section below, because it is far more commonly needed there.

3. **Use the shadow screen (128K only):** Draw the scrolled content into a back buffer, then flip. This eliminates tearing and lets you spread the work across frames. We cover this later in the chapter.

Scrolling by 8 Pixels (One Character Row)

Scrolling by a full character row is dramatically cheaper because the source and destination are related by a simple offset within each third. The character rows within a third are spaced 32 bytes apart in L. So scrolling one character row up means copying from L+32 to L, for each scanline and each third.

For scrolling the play area by one character row, the key insight is that within a single scanline, character rows are stored contiguously (32 bytes apart). Scanline 0 of char rows 0–7 in a third lives at \$xx00, \$xx20, \$xx40, ..., \$xxE0. Scrolling N character rows up by one position within a single scanline is therefore a single block copy of (N-1) x 32 bytes.

For a 20-character-row play area, one scanline's worth of data is $20 \times 32 = 640$ bytes. Scrolling that scanline means copying $19 \times 32 = 608$ bytes forward by 32. We do this for each of the 8 scanlines, handling third boundaries separately.

Estimated cost: 8 scanlines x ~12,700 T-states per scanline (608 bytes via LDIR) + third-boundary handling = approximately **105,000 T-states**. That is 146% of a frame.

Even character-row scrolling the full play area in one frame is tight. Games handle this by either:

- **Scrolling during the blank (border) period.** The top and bottom border on a Pentagon give approximately 14,000 T-states of free time where no contention occurs.
- **Splitting across two frames.** Scroll the top half one frame, the bottom half the next. The visual effect is a 25fps scroll at 8-pixel jumps.
- **Using the shadow screen** (see below).

Horizontal Pixel Scrolling

Horizontal scrolling is the bread and butter of side-scrolling games: the world moves left or right as the player walks. And it is the most expensive type of scrolling on the Spectrum, because it requires not just copying bytes but *shifting* them.

Why Horizontal Scrolling Is Expensive

When you scroll the screen left by one pixel, every byte in every row must shift its bits left by one position, and the bit that falls off the left edge of one byte must become the rightmost bit of its left neighbour. This is a rotate-and-carry chain across all 32 bytes of each row.

The Z80's RL (rotate left through carry) instruction is the tool for this. For a leftward scroll, each pixel moves one position left. Bit 7 is the leftmost pixel in a byte, bit 0 the rightmost. Shifting left means each byte's bit 7 exits and must enter bit 0 of

the byte to its left. The carry flag bridges adjacent bytes, so we process the row from **right to left**:

```
; Scroll one pixel row left by 1 pixel
; HL points to byte 31 (rightmost) of the row
;
; Process right to left. Each byte rotates left; carry propagates.
;
    or    a                ; 4 T    clear carry (no pixel entering from right)

    ; Byte 31 (rightmost)
    rl    (hl)              ; 15 T    shift left, bit 7 -> carry, carry -> bit 0
    dec    hl                ; 6 T
    ; Byte 30
    rl    (hl)              ; 15 T
    dec    hl                ; 6 T
    ; ...repeat for bytes 29 down to 0...
    ; Byte 0 (leftmost)
    rl    (hl)              ; 15 T    bit 7 of byte 0 is lost (scrolled off
    ⇨    screen)
```

Each byte costs: 15 (RL (HL)) + 6 (DEC HL) = **21 T-states** per byte. For 32 bytes per row: $32 \times 21 - 6 = \mathbf{666 \text{ T-states}}$ per row (we do not need the final DEC HL).

Actually, the first byte needs OR A (4 T) to clear carry. So one row costs: $4 + 32 \times 15 + 31 \times 6 = 4 + 480 + 186 = \mathbf{670 \text{ T-states}}$.

For 192 rows: $192 \times 670 = \mathbf{128,640 \text{ T-states}}$. That is **179% of a frame**.

A full-screen horizontal pixel scroll by one pixel does not fit in a single frame using RL chains. And this is *just the shifting* – we have not drawn any new content on the right edge.

Horizontal scrolling prototype — tiled play area with byte-level shift visualisation showing how RL chain propagates the carry bit across adjacent bytes

Figure 39: Horizontal scrolling prototype — tiled play area with byte-level shift visualisation showing how RL chain propagates the carry bit across adjacent bytes

The Full Budget Calculation

Let us lay out the complete per-row cost with all the overhead of navigating the interleaved screen:

Operation	T-states per row
Set HL to row start (or advance from previous)	~15
Set HL to rightmost byte: <code>ld a, l : or \$1F : ld l, a</code>	15
Clear carry: <code>or a</code>	4
$32 \times \text{rl (hl)}$	480
$31 \times \text{dec hl (between bytes)}$	186
Advance to next row (<code>inc h</code> or boundary cross)	4-77
Total per row (typical)	~704

For 192 rows: $192 \times 704 = \mathbf{135,168 \text{ T-states} = 189\% \text{ of one frame.}}$

For a 168-row play area: $168 \times 704 = \mathbf{118,272 \text{ T-states} = 165\% \text{ of one frame.}}$

There is no way to do a full-screen single-pixel horizontal scroll in one frame with standard methods on a 3.5 MHz Z80. This is the fundamental constraint that drives every scrolling technique in this chapter.

Scrolling on non-Pentagon machines. The budget above assumes Pentagon timing. On a standard 48K or 128K Spectrum, the RL chain writes to video RAM in contended memory (\$4000-\$7FFF), and every access during the active display period incurs wait states. Expect the per-row cost to rise from ~704 T to ~850 T — roughly a **20% overhead**. The shadow screen (page 7 on 128K) is also contended, so double-buffered scrolling does not escape the penalty. Buffer-to-screen transfers (LDIR or PUSH-based) hit the same wall: the *source* can be in uncontended RAM above \$8000, but the *destination* is always video memory. **Mitigation:** reduce the scrolling area, use character-row scrolling where possible, or synchronise the transfer to the border period. See Chapter 15.2 for contention timing tables and scheduling strategies.

Can We Do Better?

You might think unrolling or alternate addressing modes would help. They do not. RL (IX+d) costs 23 T-states — *more* than RL (HL) at 15 T. A load-rotate-store sequence (LD A, (HL) : RLA : LD (HL), A at 18 T per byte, plus 6 T for DEC HL = 24 T) is also slower. The RL (HL) : DEC HL chain at 21 T/byte is essentially optimal for horizontal pixel scrolling on the Z80.

Bottom line: The only way to make horizontal scrolling affordable is to reduce the number of rows or bytes you scroll.

Attribute (Character) Scrolling

If pixel scrolling is expensive, attribute scrolling is almost free by comparison. Attribute scrolling moves the display in 8-pixel (one character cell) jumps. You only move the 768 bytes of attribute memory and the corresponding character-aligned pixel blocks — or, more commonly, you move just the attributes and redraw the play area from a tilemap.

Scrolling Attributes with LDIR

The attribute area is linear: 32 bytes per row, 24 rows, sequential from \$5800 to \$5AFF. Scrolling left by one character column means copying bytes 1–31 to positions 0–30 in each row, then writing the new column at position 31.

For the entire 24-row attribute area:

```
; Scroll all attributes left by 1 character column
; New column data in a 24-byte table at new_col_data
;
scroll_attrs_left:
```

```

ld    hl, $5801          ; 10 T  source: column 1
ld    de, $5800          ; 10 T  dest: column 0
ld    bc, 767            ; 10 T  768 - 1 bytes
ldir                      ; 767*21 + 16 = 16,123 T

; Now fill the rightmost column with new data
ld    hl, new_col_data   ; 10 T
ld    de, $581F          ; 10 T  column 31 of row 0
ld    b, 24              ; 7 T
.fill_col:
ld    a, (hl)            ; 7 T
ld    (de), a            ; 7 T
inc   hl                 ; 6 T
; advance DE by 32 (next attribute row)
ld    a, e               ; 4 T
add   a, 32              ; 7 T
ld    e, a               ; 4 T
jr    nc, .no_carry      ; 12/7 T
inc   d                  ; 4 T
.no_carry:
djnz  .fill_col          ; 13 T
ret

; Total LDIR: ~16,123 T
; Total column fill: ~24 * 50 = ~1,200 T
; Grand total: ~17,323 T = 24.2% of frame

```

17,323 T-states for a full-screen attribute scroll. That is about 24% of a frame. Compare this to the 135,000+ T-states for a pixel scroll. Attribute scrolling is nearly 8x cheaper.

The catch: the scroll jumps by 8 pixels at a time. The visual result is coarse and jerky. For text scrollers in demos, this is often acceptable – the viewer reads the text, not the smoothness. For a game, 8-pixel jumps feel terrible. That is where the combined method comes in.

The Combined Method: Character Scroll + Pixel Offset

This is the technique that most Spectrum side-scrolling games actually use. The idea is simple and powerful:

1. Maintain a **pixel offset** counter from 0 to 7. Each frame, increment the offset.
2. When the offset reaches 8, reset it to 0 and perform an **attribute/character scroll** – the cheap operation.
3. On every frame, render the play area with the current pixel offset applied. This offset shifts the entire display by 0–7 pixels within the current character column positions.

The pixel offset can be applied in two ways:

Method A: Shift the new column. Only shift the one column of pixel data (the column being scrolled into view) by the current offset. The rest of the screen is

drawn from tiles at character alignment. This works when you have a tile-based renderer that redraws from a map.

Method B: Hardware-style virtual offset. Maintain a rendering offset that controls where within each character cell the tile data begins. This is conceptually similar to a hardware scroll register but implemented in software.

Method A is more common in practice. Let us walk through it.

How It Works

Imagine the play area is 20 characters (160 pixels) wide and 20 characters tall. The level data is a tilemap where each tile is 8x8 pixels (one character cell).

The scroll state consists of: - `scroll_tile_x`: which tile column is at the left edge of the screen (integer, advances by 1 every 8 frames). - `scroll_pixel_x`: pixel offset within the current tile (0-7, advances by 1 each frame).

Each frame:

1. **If `scroll_pixel_x` is 0:** Redraw the entire play area from the tilemap at character alignment. This is a tile renderer, which we can make fast using LDIR or LDI chains (each tile row is 1 byte or a few bytes of data copied to the right screen address). Cost: 20 columns x 20 rows x ~100 T per tile = ~40,000 T. Affordable.
2. **If `scroll_pixel_x` is 1-7:** Redraw the play area shifted by `scroll_pixel_x` pixels. For most of the play area, the tiles are character-aligned and can be drawn normally – the pixel offset only affects the **leftmost and rightmost visible columns**, where a tile is partially visible.

Wait – that is the efficient interpretation, but it requires a tile renderer that clips at sub-character boundaries. The simpler (and more common) approach is:

The Simple Combined Method

1. Every 8 frames, perform a character-level scroll (LDIR the attribute and pixel data left by one column). Cost: ~17,000 T for attributes + ~40,000 T for pixel data = ~57,000 T. Done once every 8 frames.
2. Every frame, shift a **narrow window** by 1 pixel. This window is only 1 column (32 bytes) or 2 columns (64 bytes) wide – the seam between the old data and the newly arriving column.
3. **Between character scrolls**, the display shows the last character-scrolled position with a 0-7 pixel offset applied to the edge column. The player perceives smooth 1-pixel-per-frame scrolling.

Here is the per-frame cost breakdown:

Operation	T-states	Frequency
Character scroll (full play area)	~57,000	Every 8th frame
Pixel-shift 1-2 edge columns (20 rows x 2 cols x 21 T/byte x 8 scanlines)	~6,720	Every frame

Operation	T-states	Frequency
Draw new tile column at right edge	~5,000	Every 8th frame
Attribute column update	~1,200	Every 8th frame

On 7 out of 8 frames: ~6,720 T-states for the edge pixel shift. That is under 10% of the frame budget. Plenty of room for game logic, sprites, and music.

On every 8th frame: ~6,720 + 57,000 + 5,000 + 1,200 = ~69,920 T-states. That is 97.5% of the frame budget. Tight, but doable - especially if you split the character scroll across two frames or use the shadow screen.

Implementation: The Edge-Column Pixel Shift

The key inner routine shifts 1 or 2 columns of pixel data by 1 pixel. For a 2-column (16-pixel) window, each row has 2 bytes to shift:

```
; Shift 2 bytes left by 1 pixel with carry propagation
; HL points to the right byte of the pair
;
    or    a                ; 4 T    clear carry
    rl    (hl)             ; 15 T   right byte: shift left, bit 7 -> carry
    dec   hl               ; 6 T
    rl    (hl)             ; 15 T   left byte: carry -> bit 0, bit 7 lost
                                ; total: 40 T per row (for 2-byte window)
```

For 160 rows (20 char rows x 8 scanlines): 160 x 40 = **6,400 T-states**. With pointer advancement overhead (~20 T per row), the total is about **9,600 T-states** per frame. Very affordable.

The Rendering Pipeline

Here is the complete per-frame sequence for a combined horizontal scroller:

```
frame_loop:
    halt                ; wait for interrupt

    ; --- Always: advance pixel offset ---
    ld    a, (scroll_pixel_x)
    inc   a
    cp    8
    jr    nz, .no_char_scroll

    ; --- Every 8th frame: character scroll ---
    xor   a                ; reset pixel offset to 0
    ld    (scroll_pixel_x), a

    ; Advance tile position
    ld    hl, (scroll_tile_x)
    inc   hl
    ld    (scroll_tile_x), hl
```



```

    ; Scroll pixel data left by 1 column (8 pixels)
    call scroll_pixels_left_char

    ; Scroll attributes left by 1 column
    call scroll_attrs_left

    ; Draw new tile column on right edge
    call draw_right_column

    jr    .scroll_done

.no_char_scroll:
    ld    (scroll_pixel_x), a

    ; Shift the edge columns by 1 pixel
    call shift_edge_columns

.scroll_done:
    ; --- Game logic, sprites, music ---
    call update_entities
    call draw_sprites
    call play_music

    jr    frame_loop

```

This is the skeleton of a real Spectrum side-scroller. The key insight is that smooth 1-pixel scrolling is achieved *without* shifting the entire screen every frame. The expensive character-level scroll happens only once every 8 frames, and the per-frame work is minimal.

Scrolling the Pixel Data by One Character Column

The character-level pixel scroll (step 2 in the pipeline above) shifts 8 pixels' worth of data leftward for every row. Because 8 pixels = 1 byte, this is a *byte-level* copy, not a bit-level rotate. Each row's 32 bytes shift left by 1 byte: byte[1] goes to byte[0], byte[2] goes to byte[1], ..., byte[31] goes to byte[30], and byte[31] is cleared or filled with new data.

For a single row, this is a 31-byte LDIR:

```

; Shift one pixel row left by 8 pixels (1 byte)
; HL = address of byte 1 (source), DE = address of byte 0 (dest)
; BC = 31
;
    ldir                                ; 31*21 - 5 = 646 T per row... wait.
                                         ; Actually: 30*21 + 16 = 646 T. Yes.

```

For the full play area (168 rows): 168 x 646 = 108,528 T-states + row navigation overhead.

A better approach leverages the fact that within each scanline of a character row, the bytes are contiguous. For 20 character columns, one scanline's data is 20

contiguous bytes. Scrolling that scanline left by 1 byte means LDIR of 19 bytes:

```
; Scroll one scan line of the play area left by 1 character column
; Play area is 20 columns wide (columns 2-21, for example)
; Source: column 3, Dest: column 2, count: 19
;
ld    hl, row_addr + 3    ; source = byte 3 of this scan line
ld    de, row_addr + 2    ; dest  = byte 2
ld    bc, 19              ; 19 bytes to copy
ldir                      ; 18*21 + 16 = 394 T
```

For 160 rows: $160 \times 394 = 63,040$ T-states. Add ~ 20 T per row for pointer navigation: $160 \times 414 = \mathbf{66,240}$ T-states. That is 92% of a frame. Doable but tight on the “every 8th frame” budget.

With LDI chains (19 LDIs per row): $19 \times 16 = 304$ T per row. For 160 rows: $160 \times 324 = \mathbf{51,840}$ T-states = 72% of a frame. Now we have 28% left for drawing the new column and updating attributes.

The Shadow Screen Trick

The ZX Spectrum 128K has a feature that transforms the scrolling problem: **two screen buffers**. The standard screen lives at \$4000 in page 5 (always mapped at \$4000-\$7FFF). The shadow screen lives at \$C000 in page 7 (mapped at \$C000-\$FFFF when page 7 is banked in).

Port \$7FFD controls which screen is displayed:

```
; Bit 3 of port $7FFD selects the display screen:
; Bit 3 = 0: display page 5 (standard screen at $4000)
; Bit 3 = 1: display page 7 (shadow screen at $C000)

ld    a, (current_bank)
or    %00001000          ; set bit 3: display shadow screen
ld    bc, $7FFD
out   (c), a
```

The trick for scrolling:

1. **Frame N:** The player sees the standard screen (page 5). Meanwhile, you draw the *next* frame’s scrolled content into the shadow screen (page 7, at \$C000).
2. **Frame N+1:** Flip the display to the shadow screen. The player now sees the freshly drawn frame with no tearing. Meanwhile, you start drawing frame N+2 into the now-hidden standard screen.

This double-buffering approach eliminates tearing completely and gives you a full frame (or more) to prepare each scrolled frame. The cost is that you need to maintain two complete screen states, and each “scroll” is actually a full redraw of the play area into the back buffer.

```
; Flip displayed screen and return back buffer address in HL
;
; screen_flag: 0 = showing page 5, drawing to page 7
```

```

;           1 = showing page 7, drawing to page 5
;
flip_screens:
    ld  a, (screen_flag)
    xor 1           ; 7 T toggle (XOR with immediate)
    ld  (screen_flag), a

    ld  hl, $C000      ; assume drawing to page 7
    or  a
    jr  z, .show_page5

    ; Now showing page 7, draw to page 5
    ld  hl, $4000
    ld  a, (current_bank)
    or  %00001000      ; bit 3 set: display page 7
    jr  .do_flip

.show_page5:
    ld  a, (current_bank)
    and %11110111      ; bit 3 clear: display page 5

.do_flip:
    ld  bc, $7FFD
    out (c), a
    ld  (current_bank), a
    ret                ; HL = back buffer address

```

Shadow Screen Scrolling Strategy

With double-buffering, the scrolling approach changes:

Instead of scrolling the live screen in place (which causes tearing and must complete within one frame), you **redraw the play area from the tilemap** into the back buffer at the new scroll position. This is fundamentally different. You are not *moving* existing screen data – you are *rendering fresh* from the map.

This is more work per frame (you redraw the whole play area, not just shift it), but it has significant advantages:

1. **No tearing.** The player never sees a half-scrolled screen.
2. **No edge-column shifting.** You render each tile at its correct sub-character offset directly.
3. **Flexible scroll speed.** You can scroll 1, 2, or 3 pixels per frame without changing the rendering logic.
4. **Simpler code.** A tile renderer is simpler than a combined shift-and-copy scroller.

The cost of a full play-area redraw from tiles depends on your tile renderer. With 20 x 20 tiles, each tile being 8 bytes (8 scanlines x 1 byte), and using LDI chains:

- 400 tiles x 8 bytes x 16 T per LDI = 51,200 T-states for data output.
- Plus tile address lookups and screen address calculations: ~20 T per tile x 400 = 8,000 T.
- **Total: ~59,200 T-states** = 82% of a frame.

This leaves 18% (~12,900 T-states) for sprites, game logic, and music. Tight but workable.

Comparison: Scrolling Methods on ZX Spectrum

Method	T-states/frame	% of frame	Visual quality	Notes
Full pixel scroll (horizontal, 1px)	~135,000	189%	Smooth	Impossible at 50fps
Full pixel scroll (vertical, 1px)	~107,000	149%	Smooth	Impossible at 50fps
Attribute scroll only	~17,000	24%	Jerky (8px jumps)	Very cheap
Combined (char + pixel edge)	~10,000 avg, ~70,000 peak	14%/98%	Smooth	Best single-buffer method
Shadow screen + tile redraw	~59,000	82%	Smooth, tear-free	Requires 128K
Character scroll (8px jumps)	~52,000–66,000	73–92%	Jerky	For scrolling text/status

Scrolling Right (and the Direction Problem)

Everything above describes a leftward scroll (the player moves right, the world shifts left). What about scrolling right?

For attribute scrolling, reverse the LDIR direction. Copy bytes 0–30 to positions 1–31, right to left. LDIR copies forward (low to high addresses), so for a rightward scroll you need LDDR (copy backward):

```
; Scroll attributes right by 1 character column
;
scroll_attrs_right:
    ld    hl, $5ADE          ; source: last row, column 30
    ld    de, $5ADF          ; dest: last row, column 31
    ld    bc, 767            ; 768 - 1 bytes
    lddr                     ; 767*21 + 16 = 16,123 T
    ret
```




Figure 40: Scrolling technique cost comparison

For pixel bit-shifting, a rightward scroll uses RR (HL) instead of RL (HL), processing left to right:

```

; Scroll one pixel row RIGHT by 1 pixel
; HL points to byte 0 (leftmost)
;
    or    a            ; 4 T    clear carry
    rr    (hl)         ; 15 T   shift right, bit 0 -> carry
    inc   hl           ; 6 T
    rr    (hl)         ; 15 T   carry -> bit 7
    inc   hl           ; 6 T
; ... 32 bytes total ...

```

The per-byte cost is identical: 21 T-states. A rightward scroll costs the same as a leftward scroll. The combined method works in both directions with the same budget.

For bidirectional scrolling (the player can go left or right), you need two versions of the character-scroll and edge-shift routines, switched based on direction. Self-modifying code is useful here: before the scroll, patch the RL/RR opcode and the INC/DEC direction in the shift routine. This avoids a branch inside the inner loop (see Chapter 3 for the SMC pattern).

Agon Light 2: Hardware Scrolling

The Agon Light 2's VDP (Video Display Processor) handles scrolling entirely differently from the Spectrum. Where the Spectrum programmer must move bytes manually, the Agon provides hardware-level support for scroll offsets and tilemaps.

Hardware Scroll Offsets

The VDP supports a viewport offset for bitmap modes. By setting the scroll offset registers, you shift the entire displayed image without moving any pixel data. The eZ80 sends a VDP command via the serial link:

```
; Agon: set horizontal scroll offset
; VDU 23, 0, &C3, x_low, x_high
;
    ld    a, 23
    call  vdu_write      ; VDU command prefix
    ld    a, 0
    call  vdu_write
    ld    a, $C3         ; set scroll offset command
    call  vdu_write
    ld    a, (scroll_x)
    call  vdu_write      ; x offset low byte
    ld    a, (scroll_x+1)
    call  vdu_write      ; x offset high byte
```

The hardware applies this offset when reading the framebuffer for display. No pixel data is moved, no CPU T-states are spent on shifting bytes, and the scroll is perfectly smooth at any speed. The CPU cost is just the serial communication overhead (a few hundred T-states for the VDU command sequence).

Tilemap Scrolling

The VDP's tilemap mode provides native tile-based rendering. You define a set of tiles (8x8 or 16x16 pixel patterns), build a map array that references tile indices, and the hardware renders the map at display time. Scrolling is achieved by changing the tilemap's viewport offset:

```
; Agon: set tilemap scroll offset
; VDU 23, 27, <tilemap_scroll_command>, offset_x, offset_y
;
    ld    a, 23
    call  vdu_write
    ld    a, 27
    call  vdu_write
    ld    a, 14          ; set tilemap scroll offset
    call  vdu_write
    ; ... send x and y offsets ...
```

The tilemap wraps around automatically. As the viewport scrolls past the edge of the map, the hardware wraps to the beginning (or you can update the edge column with new tile indices – the ring-buffer column loading technique).

Ring-Buffer Column Loading

For an infinitely scrolling level, the tilemap acts as a ring buffer. The map is wider than the screen by at least one column. As the player scrolls right:

1. The hardware scroll offset advances by 1 pixel per frame (or whatever speed you want).
2. When a new tile column is about to scroll into view, the eZ80 writes new tile indices into the column that just scrolled off the left edge.
3. The tilemap wraps, and the newly written column appears on the right.

```
; Ring-buffer column loading (Agon, conceptual)
;
; tilemap is 40 columns wide, screen shows 32
; scroll_col tracks which column is at the left edge
;
ring_buffer_load:
    ld  a, (scroll_col)
    add a, 32           ; column about to appear on right
    and 39             ; wrap to tilemap width (mod 40)
    ld  c, a           ; C = column index to update

    ; Load new tile data for this column from the level map
    ; (level_map is a wider array of tile indices)
    ld  hl, (level_ptr) ; pointer into the level data
    ld  b, 20           ; 20 rows
.load_col:
    ld  a, (hl)         ; read tile index from level
    inc hl
    ; Write tile index to tilemap at (C, row)
    call set_tilemap_cell ; VDP command to set one cell
    djnz .load_col

    ld  (level_ptr), hl
    ret
```

The CPU work per frame is minimal: writing 20 tile indices via VDP commands, perhaps 2,000–3,000 T-states total. The rest of the frame is available for game logic. Compare this to the Spectrum’s 59,000+ T-states for a tile-based scroll redraw. The Agon’s hardware tilemap buys you roughly a 20x reduction in CPU cost for scrolling.

Comparison: Spectrum vs. Agon Scrolling

Aspect	ZX Spectrum	Agon Light 2
Scroll granularity	Software-limited; 1px possible but expensive	1px native, zero CPU cost
CPU cost per frame	10,000–135,000 T	500–3,000 T
Tearing	Visible without double-buffering	None (VDP handles sync)

Aspect	ZX Spectrum	Agon Light 2
Direction change	Requires alternate routines or SMC	Change offset sign
Map size limit	Limited by RAM, no hardware support	Tilemap size limited by VDP memory
Colour per tile	2 colours per 8x8 cell (attribute)	Full colour per pixel

The contrast is stark. What the Spectrum programmer spends most of their frame budget on – moving pixel data across a scrambled memory layout – the Agon handles with a register write. The hardware design choices propagate through every level of the software. The Spectrum’s constraints forced the development of the combined scroll method, tile engines, and shadow-screen tricks. The Agon’s constraints are elsewhere (serial VDP latency, command overhead for complex scenes).

Practical: Horizontal Side-Scrolling Level

Spectrum Version: Combined Character + Pixel Scroll

Build a horizontal side-scroller with a 20x20-character play area that scrolls smoothly at 1 pixel per frame. The level data is a tilemap stored in a memory bank.

Here is the complete structure:

```
; Side-scroller engine – ZX Spectrum 128K
; Uses combined character + pixel method with shadow screen.
;
    ORG $8000

PLAY_X      EQU 2           ; play area starts at column 2
PLAY_Y      EQU 2           ; play area starts at char row 2
PLAY_W      EQU 20          ; play area width in characters
PLAY_H      EQU 20          ; play area height in characters

scroll_pixel_x:  DB 0        ; pixel offset 0-7
scroll_tile_x:   DW 0        ; tile column at left edge
screen_flag:     DB 0        ; which screen is visible
current_bank:    DB 0        ; current $7FFD value

; --- Main loop ---
main:
    halt                ; 4 T    sync to frame

    ; Advance scroll
    ld  a, (scroll_pixel_x) ; 13 T
    inc a                  ; 4 T
    cp  8                  ; 7 T
    jr  c, .pixel_only     ; 12/7 T
```



```

; Character scroll frame
xor  a
ld   (scroll_pixel_x), a

; Advance tile position
ld   hl, (scroll_tile_x)
inc  hl
ld   (scroll_tile_x), hl

; Get back buffer address
call get_back_buffer      ; HL = $4000 or $C000

; Redraw full play area from tilemap into back buffer
call render_play_area     ; ~50,000 T

; Flip screens
call flip_screens         ; ~30 T

jr   .frame_done

.pixel_only:
ld   (scroll_pixel_x), a

; Shift edge columns in current (non-displayed) buffer
call get_back_buffer
call shift_edge_columns   ; ~9,600 T

call flip_screens

.frame_done:
call update_player        ; ~2,000 T
call draw_sprites         ; ~5,000 T
call play_music           ; ~3,000 T (IM2 handler)

jr   main

; --- Render full play area from tilemap ---
; Input: HL = base address of target screen ($4000 or $C000)
;
render_play_area:
; For each tile in the play area:
;   Look up tile index from tilemap
;   Copy 8 bytes of tile data to screen, navigating interleave
;
; 20 columns x 20 rows = 400 tiles
; Each tile: 8 scan lines x 1 byte = 8 LDI operations
; Per tile: lookup (20 T) + 8 x (LDI 16 T + INC H 4 T) = 180 T
; Total: 400 x 180 = 72,000 T
;
; (Actual implementation uses PUSH tricks and

```

```

; pre-computed screen address tables for ~55,000 T)
ret

; --- Shift edge columns by 1 pixel ---
; Shifts the 2 rightmost columns of the play area left by 1 pixel
;
shift_edge_columns:
; For each of 160 pixel rows in the play area:
;   Navigate to the correct screen address
;   RL (HL) on the 2 edge bytes, right to left
;
; Per row: 40 T (2 bytes shifted) + 20 T (navigation)
; Total: 160 x 60 = 9,600 T
ret

```

Horizontal pixel scroller showing smooth combined character-plus-pixel scrolling over a tiled play area

Figure 41: Horizontal pixel scroller showing smooth combined character-plus-pixel scrolling over a tiled play area

Agon Version: Hardware Tilemap Scrolling

The Agon version is dramatically simpler. The main loop calls `vsync`, increments a 16-bit scroll offset, sends it to the VDP via the `set_scroll_offset` routine (a handful of `vdu_write` calls), and every 8 pixels calls `ring_buffer_load` to update one column of tile indices. The entire scroll costs under 3,000 T-states per frame, leaving 365,000+ T-states for game logic, AI, physics, and rendering. The Spectrum version is a careful exercise in cycle-counting where every technique from Chapters 2 and 3 comes together to achieve what the Agon does with a hardware register.

Vertical + Horizontal: Combined Scrolling

Some games scroll in both directions simultaneously. On the Spectrum, apply the combined method to both axes: character scroll + pixel offset (0-7) for each. The character scroll in each direction happens once every 8 frames. Both coinciding on the same frame is a 1/64 probability (about every 1.3 seconds) – either accept one dropped frame or split the work. The per-frame edge-shifting cost for both axes: horizontal edge columns (~9,600 T) + vertical edge rows (~6,400 T) = ~16,000 T = 22% of the frame. Manageable.

Optimisation Tips

1. Use a Screen Address Lookup Table

Pre-compute a table of 192 screen addresses (one per pixel row) in RAM. Cost: 384 bytes. Benefit: a 16-bit table lookup (about 30 T-states) replaces the bit-shuffling address calculation (91 T-states).

2. Scroll Only What Is Visible

If sprites cover part of the play area, you can skip scrolling the rows behind opaque sprites. Track which rows need scrolling with a dirty-row bitmap. This optimisation pays off when sprites cover a significant fraction of the play area.

3. Use PUSH for the Character Scroll

For the character-level pixel data scroll (copying 19 bytes left per scanline), the PUSH trick works well. Set SP to the end of each scanline's play area, POP 10 bytes, shift the register contents, and PUSH them back one byte offset. This is complex to set up but reduces the per-scanline cost by 30-40%.

4. Split the Character Scroll Across Frames

If the character scroll (every 8th frame) is too expensive for one frame, split it: scroll the top half of the play area on frame N and the bottom half on frame N+1. The visual artefact (the top half shifts 1 frame before the bottom) is barely noticeable at 50fps.

5. Palette and Attribute Tricks

For attribute-only scrolling (no pixel data involved), consider using FLASH or BRIGHT changes to create the illusion of motion within a static pixel grid. A rotating set of attribute colours on character-aligned tiles can simulate flow, water, or conveyor belts without moving any pixel data at all.

Summary

- **Full-screen pixel scrolling on the ZX Spectrum is impossible at 50fps.** Horizontal pixel scrolling costs ~135,000 T-states for 192 rows (189% of the frame budget). Vertical costs ~107,000 T-states (149%). The interleaved memory layout adds complexity to vertical scrolling, and the lack of a barrel shifter makes horizontal scrolling inherently expensive.
- **Attribute scrolling is cheap** at ~17,000 T-states (24% of frame), but moves in coarse 8-pixel jumps.
- **The combined method** is what real games use: character-level scrolls (every 8 frames) plus per-frame pixel-shifting of the 1-2 edge columns. Average per-frame cost is under 10,000 T-states. The spike on character-scroll frames (~70,000 T) can be managed with the shadow screen or by splitting across frames.
- **The shadow screen** (128K, page 7) provides tear-free double buffering. Draw the next frame into the back buffer, then flip the display. This changes the scrolling strategy from "shift existing data" to "redraw from tilemap," which is conceptually simpler and eliminates tearing.
- **Horizontal scrolling direction** does not change the cost. Rightward scrolling uses RR (HL) instead of RL (HL), left-to-right instead of right-to-left, at the same 21 T-states per byte.

- **Vertical pixel scrolling** is complicated by the Spectrum's interleaved screen layout. Moving one pixel row down means navigating the 010TTSSS LLLCCCC address structure, with different pointer adjustments at character-cell and third boundaries. The split-counter approach from Chapter 2 is essential.
- **The Agon Light 2** provides hardware scroll offsets and tilemap rendering that reduce the CPU cost of scrolling to a few VDP commands per frame (~500–3,000 T-states). Ring-buffer column loading keeps the tilemap updated as new terrain scrolls into view. What the Spectrum programmer builds in 70,000 T-states, the Agon handles with a register write.
- **Key techniques from earlier chapters** are essential here: the interleaved screen layout and split-counter navigation (Chapter 2), LDI chains and PUSH tricks for fast data movement (Chapter 3), and self-modifying code for direction-switchable scroll routines (Chapter 3).

Sources: Introspec “Eshchy raz pro DOWN_HL” (Hype, 2020) for interleaved screen navigation; Introspec “GO WEST Part 1” (Hype, 2015) for contended memory costs; DenisGrachev “Ringo Render 64x48” (Hype, 2022) for half-character displacement scrolling; ZX Spectrum 128K Technical Manual for port \$7FFD and shadow screen; Agon Light 2 VDP documentation for tilemap and scroll offset commands.

Chapter 18: Game Loop and Entity System

“A game is a demo that listens.”

Every demo effect we have built so far runs in a closed loop: calculate, render, repeat. The viewer watches. The code does not care whether anyone is in the room. A game breaks this contract. A game *responds*. The player presses a key and something must change – immediately, reliably, within the same frame budget we have been counting since Chapter 1.

This chapter is about building the architecture that makes a game possible on the ZX Spectrum and Agon Light 2. Not the rendering (Chapter 16 covered sprites, Chapter 17 covered scrolling) and not the physics (Chapter 19 will handle collisions and AI). This chapter is the skeleton: the main loop that drives everything, the state machine that organises the flow from title screen to gameplay to game over, the input system that reads the player’s intentions, and the entity system that manages every object in the game world.

By the end, you will have a working game skeleton with 16 active entities – a player, eight enemies, and seven bullets – running within the frame budget on both platforms.

18.1 The Main Loop

Every game on the ZX Spectrum follows the same fundamental rhythm:

2. Read input -- what does the player want?
3. Update state -- move entities, run AI, check collisions
4. Render -- draw the frame
5. Go to 1

This is the game loop. It is not complicated. Its power comes from the fact that it runs fifty times per second, every second, and everything the player experiences emerges from this cycle.

Here is the minimal implementation:

```
ORG  $8000

; Install IM1 interrupt handler (standard for games)
```

```

im    1
ei

main_loop:
    halt                ; 4T + wait -- sync to frame interrupt

    call read_input      ; poll keyboard/joystick
    call update_entities ; move everything, run logic
    call render_frame     ; draw to screen

    jr    main_loop      ; 12T -- loop forever

```

The HALT instruction is the heartbeat. When the CPU executes HALT, it stops and waits for the next maskable interrupt. On the Spectrum, the ULA fires this interrupt at the start of each frame – once every 1/50th of a second. The CPU wakes up, the IM1 handler at address \$0038 runs (on a stock ROM this simply increments the frame counter), and then execution resumes at the instruction after HALT. Your main loop code runs, does its work, and hits HALT again to wait for the next frame.

This gives you exactly one frame’s worth of T-states to do everything. If your work finishes early, the CPU idles inside HALT until the next interrupt – no wasted power, no drift, perfect sync. If your work takes too long and the interrupt fires before you reach HALT, you miss a frame. The loop still works (the next HALT will catch the following interrupt), but the game drops to 25 fps for that frame. Miss consistently and you are at 25 fps permanently. Miss badly and you are at 16.7 fps (every third frame). The player notices.

The Frame Budget, Revisited

We established the numbers in Chapter 1, but they bear repeating in the context of a game:

Machine	T-states per frame	Practical budget
ZX Spectrum 48K	69,888	~62,000 (after interrupt overhead)
ZX Spectrum 128K	70,908	~63,000
Pentagon 128	71,680	~64,000
Agon Light 2	~368,640	~360,000

The “practical budget” accounts for the interrupt handler, the HALT instruction itself, and the border timing overhead. On the Spectrum, you have roughly 64,000 T-states of usable time per frame. On the Agon, you have over five times that.

How does a typical game spend those 64,000 T-states? Here is a realistic breakdown for a Spectrum platformer:

Subsystem	T-states	% of budget
Input reading	~500	0.8%
Entity update (16 entities)	~8,000	12.5%
Collision detection	~4,000	6.3%
Music player (PT3)	~5,000	7.8%
Sprite rendering (8 visible)	~24,000	37.5%

Subsystem	T-states	% of budget
Background/scroll update	~12,000	18.8%
Miscellaneous (HUD, state)	~3,000	4.7%
Remaining headroom	~7,500	11.7%

That 11.7% headroom is your safety margin. Eat into it and you start dropping frames on complex scenes. The border-colour profiling technique from Chapter 1 – red for sprites, blue for music, green for logic – is how you monitor this budget during development. Use it constantly.

On the Agon, the same game logic runs in a fraction of the budget. The entity update, collision detection, and input reading might consume 15,000 T-states total – about 4% of the Agon’s frame. The VDP handles sprite rendering on the ESP32 coprocessor, so CPU-side sprite cost drops to the VDU command overhead. You have enormous room for more complex AI, more entities, or simply less stress.

18.2 The Game State Machine

A game is not one loop – it is several. The title screen has its own loop (animate logo, wait for keypress). The menu has its own loop (highlight options, read input). The gameplay loop is what we described above. The pause screen freezes the gameplay loop and runs a simpler one. The game-over screen has yet another.

The cleanest way to organise these is a **state machine**: a variable that tracks which state the game is in, and a table of handler addresses – one per state.

State Definitions

```
; Game states (byte values, used as table offsets)
STATE_TITLE      EQU 0
STATE_MENU       EQU 2      ; x2 because each table entry is 2 bytes
STATE_GAME       EQU 4
STATE_PAUSE      EQU 6
STATE_GAMEOVER   EQU 8

; Current state variable
game_state:      DB  STATE_TITLE
```

The Jump Table

```
; Table of handler addresses, indexed by state
state_table:
    DW  state_title      ; STATE_TITLE  = 0
    DW  state_menu       ; STATE_MENU   = 2
    DW  state_game       ; STATE_GAME  = 4
    DW  state_pause      ; STATE_PAUSE = 6
    DW  state_gameover    ; STATE_GAMEOVER = 8
```




Figure 42: Game loop architecture




Figure 43: Game state machine: Title, Menu, Game, Pause, and Game Over states connected by labelled transitions. Each state runs its own loop; transitions happen via a jump table.

The Dispatcher

The main loop becomes a dispatcher that reads the current state and jumps to the appropriate handler:

```

main_loop:
    halt                ; sync to frame

    ; Dispatch to current state handler
    ld  a, (game_state) ; 13T  load state index
    ld  l, a             ; 4T
    ld  h, 0             ; 7T
    ld  de, state_table ; 10T
    add hl, de           ; 11T  HL = state_table + offset
    ld  e, (hl)          ; 7T  low byte of handler address
    inc hl               ; 6T
    ld  d, (hl)          ; 7T  high byte of handler address
    ex  de, hl           ; 4T  HL = handler address
    jp  (hl)             ; 4T  jump to handler
                        ; --- 73T total dispatch overhead
  
```

The JP (HL) instruction is the key. It does not jump to the address stored *at* HL – it jumps to the address *in* HL. This is Z80’s indirect jump, and it costs only 4 T-states. The entire dispatch – loading the state variable, computing the table offset, reading the handler address, and jumping – takes 73 T-states. That is negligible: about 0.1% of the frame budget.

Each handler runs its own logic and then jumps back to `main_loop`:

```

state_title:
    call draw_title_screen
    call read_input
    ; Check for start key (SPACE or ENTER)
    ld  a, (input_flags)
    bit  BUTTON_FIRE, a
    jr  z, .no_start
    ; Transition to menu
    ld  a, STATE_MENU
    ld  (game_state), a
    call init_menu          ; set up menu screen
.no_start:
    jp  main_loop

state_game:
    call read_input
    ; Check for pause
    ld  a, (input_keys)
    bit  KEY_P, a
    jr  z, .not_paused
    ld  a, STATE_PAUSE
    ld  (game_state), a
    jp  main_loop
.not_paused:
    call update_entities
    call check_collisions
    call render_frame
    call update_music      ; AY player -- see Chapter 11
    jp  main_loop

state_pause:
    ; Game is frozen -- only check for unpause
    call read_input
    ld  a, (input_keys)
    bit  KEY_P, a
    jr  z, .still_paused
    ld  a, STATE_GAME
    ld  (game_state), a
.still_paused:
    ; Optionally blink "PAUSED" text
    call blink_pause_text
    jp  main_loop

state_gameover:
    call draw_gameover_screen
    call read_input
    ld  a, (input_flags)
    bit  BUTTON_FIRE, a
    jr  z, .wait
    ld  a, STATE_TITLE

```

```

        ld    (game_state), a
        call init_title
.wait:
        jp    main_loop

```

Why Not a Chain of Comparisons?

You might be tempted to write the dispatcher as:

```

ld    a, (game_state)
cp    STATE_TITLE
jp    z, state_title
cp    STATE_MENU
jp    z, state_menu
cp    STATE_GAME
jp    z, state_game
; ...

```

This works, but it has two problems. First, the cost grows linearly: each additional state adds a CP (7T) and a JP Z (10T), so the worst case is 17T per state. With 5 states, the game state (the most common case) might take 51T to reach if it is the third comparison. The jump table takes 73T regardless of which state is active – it is $O(1)$, not $O(n)$.

Second, and more importantly, the jump table scales cleanly. Adding a sixth state (say, STATE_SHOP) means adding one DW entry to the table and one constant definition. The dispatcher code does not change at all. With comparison chains, you add more instructions to the dispatcher itself, and the ordering starts to matter for performance. The table approach is both faster in the common case and cleaner to maintain.

State Transitions

State transitions happen by writing a new value to `game_state`. Typically you also call an initialisation routine for the new state:

```

; Transition: Game -> Game Over
game_over_transition:
    ld    a, STATE_GAMEOVER
    ld    (game_state), a
    call init_gameover      ; set up game over screen, save score
    ret

```

Keep transitions explicit and centralised. A common bug in Z80 games is a state transition that forgets to initialise the new state's data – the game-over screen shows garbage because nobody cleared the screen or reset the animation counter. Every state should have an `init_` routine that the transition calls.

18.3 Input: Reading the Player

ZX Spectrum Keyboard

The Spectrum keyboard is read through port \$FE. The keyboard is wired as a matrix of 8 half-rows, each selected by setting a bit low in the high byte of the port address. Reading port \$FE with a specific high byte returns the state of that half-row: 5 bits, one per key, where 0 means pressed and 1 means not pressed.

The half-row map:

High byte	Keys (bit 0 to bit 4)
\$FE (bit 0 low)	SHIFT, Z, X, C, V
\$FD (bit 1 low)	A, S, D, F, G
\$FB (bit 2 low)	Q, W, E, R, T
\$F7 (bit 3 low)	1, 2, 3, 4, 5
\$EF (bit 4 low)	0, 9, 8, 7, 6
\$DF (bit 5 low)	P, O, I, U, Y
\$BF (bit 6 low)	ENTER, L, K, J, H
\$7F (bit 7 low)	SPACE, SYMSHIFT, M, N, B

The standard game controls – Q/A/O/P for up/down/left/right and SPACE for fire – span three half-rows. Here is a routine that reads them and packs the result into a single byte:

```
; Input flag bits
INPUT_RIGHT EQU 0
INPUT_LEFT  EQU 1
INPUT_DOWN  EQU 2
INPUT_UP    EQU 3
INPUT_FIRE  EQU 4

; Read QAOP+SPACE into input_flags
; Returns: A = input_flags byte, also stored at (input_flags)
read_keyboard:
    ld    d, 0                ; 7T    accumulate result in D

    ; Read O and P: half-row $DF (P=bit0, O=bit1)
    ld    bc, $DFFE           ; 10T
    in    a, (c)              ; 12T
    bit   0, a                 ; 8T    P key
    jr    nz, .no_right       ; 12/7T
    set   INPUT_RIGHT, d      ; 8T

.no_right:
    bit   1, a                 ; 8T    O key
    jr    nz, .no_left        ; 12/7T
    set   INPUT_LEFT, d       ; 8T

.no_left:

    ; Read Q and A: half-rows $FB (Q=bit0) and $FD (A=bit0... wait)
    ; Q is in half-row $FB at bit 0
```

```

    ld    b, $FB                ; 7T
    in     a, (c)                ; 12T
    bit    0, a                  ; 8T    Q key
    jr     nz, .no_up            ; 12/7T
    set    INPUT_UP, d           ; 8T
.no_up:

    ; A is in half-row $FD at bit 0
    ld     b, $FD                ; 7T
    in     a, (c)                ; 12T
    bit    0, a                  ; 8T    A key
    jr     nz, .no_down         ; 12/7T
    set    INPUT_DOWN, d         ; 8T
.no_down:

    ; SPACE: half-row $7F at bit 0
    ld     b, $7F                ; 7T
    in     a, (c)                ; 12T
    bit    0, a                  ; 8T
    jr     nz, .no_fire         ; 12/7T
    set    INPUT_FIRE, d         ; 8T
.no_fire:

    ld     a, d                  ; 4T
    ld     (input_flags), a      ; 13T
    ret                                ; 10T
    ; Total: ~220T worst case (all keys pressed)

```

At roughly 220 T-states worst case, input reading is trivial in the frame budget. Even on the Spectrum you could afford to read the keyboard ten times per frame and barely notice.

Kempston Joystick

The Kempston interface is even simpler. One port read returns all five directions plus fire:

```

; Kempston joystick port
KEMPSTON_PORT EQU $1F

; Read Kempston joystick
; Returns: A = joystick state
; bit 0 = right, bit 1 = left, bit 2 = down, bit 3 = up, bit 4 = fire
read_kempston:
    in     a, (KEMPSTON_PORT)    ; 11T
    and    %00011111             ; 7T    mask to 5 bits
    ld     (input_flags), a      ; 13T
    ret                                ; 10T
    ; Total: 41T

```

Notice something convenient: the Kempston bit layout matches our INPUT_* flag definitions exactly. This is not a coincidence – the Kempston interface was designed with this standard in mind, and most Spectrum games adopt the same bit ordering.

If you support both keyboard and joystick, you can OR the results together:

```
read_input:
    call read_keyboard      ; D = keyboard flags
    push de
    call read_kempston     ; A = joystick flags
    pop de
    or d                    ; combine both sources
    ld (input_flags), a
    ret
```

Now the rest of your code only checks `input_flags` and does not care whether the input came from the keyboard or a joystick.

Edge Detection: Press vs Hold

For some actions – firing a bullet, opening a menu – you want to respond to the *press* event, not the held state. If you check bit `INPUT_FIRE`, a every frame, the player fires a bullet every 1/50th of a second while holding the button. That might be intentional for rapid-fire, but for a single-shot weapon or a menu selection, you need edge detection.

The technique: store the previous frame’s input alongside the current frame’s, and XOR them to find the bits that changed:

```
input_flags:      DB 0      ; current frame
input_prev:       DB 0      ; previous frame
input_pressed:    DB 0      ; newly pressed this frame (edges)

read_input_with_edges:
    ; Save previous state
    ld a, (input_flags)
    ld (input_prev), a

    ; Read current state
    call read_input      ; updates input_flags

    ; Compute edges: pressed = current AND NOT previous
    ld a, (input_prev)
    cpl                ; 4T  invert previous
    ld b, a             ; 4T
    ld a, (input_flags) ; 13T
    and b               ; 4T  current AND NOT previous
    ld (input_pressed), a ; 13T = newly pressed this frame
    ret
```

Now `input_pressed` has a 1 bit only for buttons that were *not* pressed last frame but *are* pressed this frame. Use `input_flags` for continuous actions (movement) and `input_pressed` for one-shot actions (fire, jump, menu select).

Agon Light 2: PS/2 Keyboard via MOS

The Agon reads its PS/2 keyboard through the MOS (Machine Operating System) API. The eZ80 does not directly scan a keyboard matrix – instead, the ESP32 VDP

coprocessor handles the keyboard hardware and passes keypress events to the eZ80 via a shared buffer.

The MOS system variable `sysvar_keyascii` (at address `$0800 + offset`) holds the ASCII code of the most recently pressed key, or 0 if no key is down. For game controls, you typically poll this variable or use the MOS `waitvblank` / keyboard API calls:

```
; Agon: Read keyboard via MOS sysvar
; MOS sysvar_keyascii at IX+$05
read_input_agon:
    ld    a, (ix + $05)      ; read last key from MOS sysvars
    ; Map ASCII to input_flags
    cp    'o'
    jr    nz, .not_left
    set   INPUT_LEFT, d
.not_left:
    cp    'p'
    jr    nz, .not_right
    set   INPUT_RIGHT, d
.not_right:
    ; ... etc for Q, A, SPACE
    ld    a, d
    ld    (input_flags), a
    ret
```

The Agon also supports reading individual key states via VDU commands (VDU 23,0,\$01,keycode), which return whether a specific key is currently held. This is closer to the Spectrum's half-row approach and better suited for games that need simultaneous key detection. The MOS API handles the PS/2 protocol, scan code translation, and auto-repeat – none of which you need to worry about.

18.4 The Entity Structure

A game entity is anything that moves, animates, interacts, or needs per-frame updating: the player character, enemies, bullets, explosions, floating score numbers, power-ups. On the Z80, we represent each entity as a fixed-size block of bytes in memory.

Structure Layout

Here is the entity structure we will use throughout the game-dev chapters:

```
-----
+0    1    x_frac    X sub-pixel (low byte of 8.8 fixed-point)
+1    1    x_int     X pixel column (high byte, 0-255)
+2    1    y_frac    Y sub-pixel (low byte of 8.8 fixed-point)
+3    1    y_int     Y pixel row (high byte, 0-191)
+4    1    type      Entity type (0=inactive, 1=player, 2=enemy, 3=bullet,
↪ ...)
+5    1    state     Entity state (0=idle, 1=active, 2=dying, 3=dead, ...)
```

+6	1	anim_frame	Current animation frame index
+7	1	dx_frac	Horizontal velocity, fractional part
+8	1	dx_int	Horizontal velocity, integer part (signed)
+9	1	dy_frac	Vertical velocity, fractional part
+10	1	dy_int	Vertical velocity, integer part (signed)
+11	1	health	Hit points remaining
+12	1	flags	Bit flags (see below)
+13	1	width	Bounding box width in pixels (for collision, Ch.19)
+14	1	height	Bounding box height in pixels (for collision, Ch.19)
+15	1	(reserved)	

16 bytes total per entity

Flag bits in the flags byte:

Bit 7: ACTIVE -- entity is alive and should be updated/rendered
 Bit 4: VISIBLE -- entity should be rendered (active but invisible = logic
 ⇨ only)
 Bit 2: COLLIDABLE -- entity participates in collision detection
 Bit 1: FACING_LEFT -- horizontal facing direction
 Bit 3: INVINCIBLE -- temporary invulnerability (player after being hit)
 Bit 0: ON_GROUND -- entity is standing on solid ground (set by physics)
 Bit 5-6: reserved

ACTIVE is at bit 7 so that bit 7, (ix+12) gives a fast alive check, and `ld a, (ix+12)` / and \$80 tests it with one logical instruction. ON_GROUND at bit 0 allows bit 0, (ix+12) for the grounded test that Chapter 19's physics will use.

Why 16 Bytes?

Sixteen bytes is a deliberate choice. It is a power of two, which makes index-to-address conversion trivial – four left shifts instead of the multiply-by-constant decomposition that odd sizes require. Sixteen entities occupy 256 bytes, exactly one page, which simplifies certain address calculations. And all 16 bytes are used: positions, velocities, type, state, animation, health, flags, bounding box dimensions, and one reserved byte for future use. Nothing is wasted.

```
; Calculate entity address from index in A
; Input: A = entity index (0-15)
; Output: HL = address of entity structure
; Destroys: nothing extra
get_entity_addr:
    ld    l, a                ; 4T
    ld    h, 0                ; 7T
    add   hl, hl               ; 11T  x2
    add   hl, hl               ; 11T  x4
    add   hl, hl               ; 11T  x8
    add   hl, hl               ; 11T  x16
    ld    de, entity_array    ; 10T
    add   hl, de               ; 11T  HL = entity_array + index * 16
    ret                                ; 10T
; Total: 86T
```

Four ADD HL,HL instructions multiply by 16. No saved copy, no two-part decomposi-

tion, no extra register needed. Compare this to multiplying by 10 (which requires saving an intermediate result in DE for the 8+2 decomposition, 94 T-states) or by 12 (8+4, same overhead). Sixteen is the sweet spot: enough room for everything Chapter 19's physics and collision code will need, with the simplest possible index arithmetic.

Both Axes in 8.8 Fixed-Point

Both X and Y use 16-bit 8.8 fixed-point: the high byte is the pixel coordinate and the low byte is a sub-pixel fraction for smooth movement. This is essential for any game where objects move at fractional-pixel speeds. A character moving at 1.5 pixels per frame with only integer coordinates would alternate between 1-pixel and 2-pixel steps, producing visible judder. With 8.8 fixed-point, the movement is smooth: add \$0180 each frame and the pixel position advances 1, 2, 1, 2... in a pattern the eye perceives as steady 1.5 pixels per frame.

Y needs the same precision as X. A platformer with gravity applies a small downward acceleration each frame – without sub-pixel Y, the jump arc becomes a staircase of integer steps instead of a smooth parabola. The Spectrum's screen is 192 pixels tall, so a single byte covers the pixel range, but the sub-pixel byte at +2 (y_frac) gives gravity and jumping the resolution they need. Chapter 19's physics code relies on this.

The 8.8 Fixed-Point System

Fixed-point arithmetic was introduced in Chapter 4. Here is a quick recap of how it applies to entity movement. Both X and Y work the same way – a 16-bit add of position and velocity, with carry propagating from the fractional byte to the pixel byte:

```
; Move entity: add velocity to position (one axis)
; IX = entity pointer
; Example: X axis (same pattern for Y at offsets +2/+3 with +9/+10)
move_entity_x:
    ld    a, (ix+0)          ; 19T  x_frac
    add   a, (ix+7)          ; 19T  + dx_frac
    ld    (ix+0), a          ; 19T  store x_frac
    ld    a, (ix+1)          ; 19T  x_int
    adc   a, (ix+8)          ; 19T  + dx_int (with carry from frac)
    ld    (ix+1), a          ; 19T  store x_int
    ret
; Total: 114T per axis
```

The beauty of fixed-point: the 8-bit ADD on the fractional bytes produces a carry that ADC propagates into the integer byte. No 16-bit register juggling needed – just two adds and two stores. A velocity of (\$80, \$00) moves 0.5 pixels per frame; (\$80, \$01) moves 1.5 pixels per frame. The same pattern applies to both axes, and Chapter 19's move_entity routine does exactly this for X and Y in sequence.

18.5 The Entity Array

Entities live in a statically allocated array. No dynamic memory allocation, no linked lists, no heap. Static arrays are the standard approach on the Z80 for good reason: they are fast, predictable, and cannot fragment.

```
; Entity array: 16 entities, 16 bytes each = 256 bytes (one page)
MAX_ENTITIES    EQU 16
ENTITY_SIZE     EQU 16

entity_array:
    DS    MAX_ENTITIES * ENTITY_SIZE    ; 256 bytes, zeroed at init
```

Entity Slot Allocation

Slot 0 is always the player. Slots 1-8 are enemies. Slots 9-15 are projectiles and effects (bullets, explosions, score popups). This fixed partitioning simplifies the code: when you need to iterate over enemies for AI, you iterate slots 1-8. When a bullet needs spawning, you search slots 9-15. The player is always at a known address.

```
; Fixed slot assignments
SLOT_PLAYER      EQU 0
SLOT_ENEMY_FIRST EQU 1
SLOT_ENEMY_LAST  EQU 8
SLOT_PROJ_FIRST  EQU 9
SLOT_PROJ_LAST   EQU 15
```

Iterating Entities

The core update loop walks through every entity slot, checks the ACTIVE flag, and calls the appropriate update handler:

```
; Update all active entities
; Total cost: ~2,500T for 16 entities (most inactive), up to ~8,000T (all
;   ↪ active)
update_entities:
    ld    ix, entity_array    ; 14T  IX points to first entity
    ld    b, MAX_ENTITIES     ; 7T   loop counter

.loop:
    ; Check if entity is active
    ld    a, (ix + 12)        ; 19T  load flags byte (offset +12)
    bit   7, a                ; 8T   test ACTIVE flag (bit 7)
    jr    z, .skip            ; 12/7T skip if inactive

    ; Entity is active -- dispatch by type
    ld    a, (ix + 4)          ; 19T  load type byte (offset +4)
    ; Jump table dispatch based on type
    call  update_by_type      ; ~200-500T depending on type

.skip:
    ; Advance IX to next entity
```

```

ld    de, ENTITY_SIZE    ; 10T
add   ix, de              ; 15T  IX += 16
djnz  .loop               ; 13/8T
ret

```

This uses IX as the entity pointer, which is convenient because IX-indexed addressing lets you access any field by its offset: (IX+0) is x_frac, (IX+1) is x_int, (IX+4) is type, (IX+12) is flags, and so on. The downside of IX is cost: every LD A, (IX+n) takes 19 T-states versus 7 for LD A, (HL). For the entity update loop, which runs 16 times per frame, this overhead is acceptable. For the inner rendering loop where you touch entity data thousands of times per frame, you would copy the relevant fields into registers first.

Update Dispatch by Type

Each entity type has its own update handler. We use the same jump-table technique as the game state machine:

```

; Entity type constants
TYPE_INACTIVE EQU 0
TYPE_PLAYER   EQU 1
TYPE_ENEMY    EQU 2
TYPE_BULLET  EQU 3
TYPE_EXPLOSION EQU 4

; Handler table (2 bytes per entry)
type_handlers:
    DW  update_inactive    ; type 0: no-op, should not be called
    DW  update_player      ; type 1
    DW  update_enemy       ; type 2
    DW  update_bullet      ; type 3
    DW  update_explosion    ; type 4

; Dispatch to type handler
; Input: A = entity type, IX = entity pointer
update_by_type:
    add  a, a                ; 4T   type * 2 (table entries are 2 bytes)
    ld   l, a                ; 4T
    ld   h, 0                ; 7T
    ld   de, type_handlers   ; 10T
    add  hl, de              ; 11T
    ld   e, (hl)             ; 7T
    inc  hl                  ; 6T
    ld   d, (hl)             ; 7T
    ex   de, hl              ; 4T
    jp   (hl)                ; 4T   jump to handler (RET will return to caller)
                                ; --- 64T dispatch overhead

```

Each handler receives IX pointing to the entity and can access all fields via indexed addressing. When the handler executes RET, it returns to the entity update loop, which advances to the next slot.

The Player Update Handler

Here is a typical player update – read input flags, apply movement, update animation:

```
; Update player entity
; IX = entity pointer (slot 0)
update_player:
    ; Read horizontal input
    ld  a, (input_flags)    ; 13T
    bit INPUT_RIGHT, a      ; 8T
    jr  z, .not_right       ; 12/7T
    ; Move right: add dx to X (8.8 fixed-point)
    ld  a, 2                ; 7T  dx_frac = 2 subpixels per frame
    add a, (ix + 0)         ; 19T  add to x_frac
    ld  (ix + 0), a         ; 19T
    jr  nc, .no_carry_r     ; 12/7T
    inc (ix + 1)            ; 23T  carry into x_int
.no_carry_r:
    res FLAG_FACING_L, (ix + 12) ; 23T  clear FACING_LEFT flag
    jr  .horiz_done         ; 12T
.not_right:
    bit INPUT_LEFT, a       ; 8T
    jr  z, .horiz_done       ; 12/7T
    ; Move left: subtract dx from X
    ld  a, (ix + 0)         ; 19T  load x_frac
    sub 2                   ; 7T  subtract dx_frac
    ld  (ix + 0), a         ; 19T
    jr  nc, .no_borrow_l    ; 12/7T
    dec (ix + 1)            ; 23T  borrow from x_int
.no_borrow_l:
    set FLAG_FACING_L, (ix + 12) ; 23T  set FACING_LEFT flag
.horiz_done:

    ; Update animation frame (cycle every 8 frames)
    ld  a, (ix + 6)         ; 19T  anim_frame (offset +6)
    inc a                   ; 4T
    and 7                   ; 7T  wrap 0-7
    ld  (ix + 6), a         ; 19T
    ret
; Total: ~250-350T depending on input
```

This is deliberately simple. Chapter 19 will add gravity, jumping, and collision response. For now, the point is the *structure*: entity pointer in IX, fields accessed by offset, input flags driving state changes, animation counter ticking.

18.6 The Object Pool

Bullets, explosions, and particle effects are transient. A bullet exists for a fraction of a second before it hits something or leaves the screen. An explosion animates

for 8-16 frames and vanishes. You could spawn these dynamically, but on the Z80, “dynamic” means searching for free memory, managing allocation, and risking fragmentation. Instead, we use an **object pool**: a fixed set of slots that entities are activated into and deactivated from.

We already have the pool – it is the entity array. Slots 9-15 are the projectile/effect pool. Spawning a bullet means finding an inactive slot in that range and filling it in. Destroying a bullet means clearing its ACTIVE flag.

Spawning a Bullet

```
; Spawn a bullet at position (B=x_pixel, C=y_pixel)
; moving in direction determined by player facing
; Returns: carry set if no free slot available
spawn_bullet:
    ld    ix, entity_array + (SLOT_PROJ_FIRST * ENTITY_SIZE)
    ld    d, SLOT_PROJ_LAST - SLOT_PROJ_FIRST + 1 ; 7 slots to check

.find_slot:
    ld    a, (ix + 12)      ; 19T flags
    bit   7, a              ; 8T  ACTIVE?
    jr    z, .found         ; 12/7T found an inactive slot

    push  de                ; 11T save loop counter (D)
    ld    de, ENTITY_SIZE   ; 10T
    add   ix, de            ; 15T next slot
    pop   de                ; 10T restore loop counter
    dec   d                 ; 4T
    jr    nz, .find_slot    ; 12T

    ; No free slot -- set carry and return
    scf                      ; 4T
    ret

.found:
    ; Fill in the bullet entity
    ld    (ix + 0), 0        ; x_frac = 0
    ld    (ix + 1), b        ; x_int = B
    ld    (ix + 2), 0        ; y_frac = 0
    ld    (ix + 3), c        ; y_int = C
    ld    (ix + 4), TYPE_BULLET ; type
    ld    (ix + 5), 1        ; state = active
    ld    (ix + 6), 0        ; anim_frame = 0
    ld    (ix + 11), 1       ; health = 1 (dies on first collision)

    ; Set velocity based on player facing
    ld    a, (entity_array + 12) ; player flags
    bit   FLAG_FACING_L, a      ; FACING_LEFT?
    jr    z, .fire_right
    ld    (ix + 7), 0           ; dx_frac = 0
    ld    (ix + 8), -4         ; dx_int = -4 (fast, leftward)
    jr    .set_flags
```

```
.fire_right:
    ld    (ix + 7), 0          ; dx_frac = 0
    ld    (ix + 8), 4          ; dx_int = +4 (fast, rightward)
.set_flags:
    ld    (ix + 9), 0          ; dy_frac = 0
    ld    (ix + 10), 0         ; dy_int = 0 (horizontal bullet)
    ld    (ix + 12), %10010100 ; flags: ACTIVE + VISIBLE + COLLIDABLE
    or    a                    ; clear carry (success)
    ret
```

Deactivating an Entity

When a bullet leaves the screen or an explosion finishes its animation, deactivation is a single instruction:

```
; Deactivate entity at IX
deactivate_entity:
    ld    (ix + 12), 0         ; 19T clear all flags (ACTIVE=0)
    ret
```

That is it. Next frame, the update loop sees bit 7 of flags clear and skips the slot. The slot is now available for the next `spawn_bullet` call to reuse.

Bullet Update Handler

```
; Update a bullet entity
; IX = entity pointer
update_bullet:
    ; Move horizontally (8.8 fixed-point add)
    ld    a, (ix + 0)          ; 19T x_frac
    add    a, (ix + 7)          ; 19T + dx_frac
    ld    (ix + 0), a          ; 19T
    ld    a, (ix + 1)          ; 19T x_int
    adc    a, (ix + 8)          ; 19T + dx_int (with carry)
    ld    (ix + 1), a          ; 19T

    ; Check screen bounds (0-255 pixel range)
    or    a                    ; 4T
    jr    z, .off_screen       ; if x_int=0, leftward bullet exited
    cp    248                  ; 7T near right edge?
    jr    nc, .off_screen      ; past right boundary

    ; Still alive -- return
    ret

.off_screen:
    ; Deactivate
    ld    (ix + 12), 0         ; clear flags
    ret
; Total: ~170T active, ~190T when deactivating
```

Pool Sizing

Seven projectile slots (indices 9-15) might sound limited. In practice, it is more than enough for most Spectrum games. Consider: a bullet that crosses the full screen width (256 pixels) at 4 pixels per frame takes 64 frames – over a second. If the player fires once every 8 frames (a rapid fire rate), at most 8 bullets can exist simultaneously. Seven slots with occasional spawn failures (the bullet simply does not fire that frame) feels natural, not buggy. The player is unlikely to notice a missed bullet at the edge of their fire rate.

If you need more, expand the entity array. But be aware of the cost: each additional entity adds ~160 T-states to the worst-case update loop (when active) and ~50 T-states even when inactive (the ACTIVE flag check and the IX advance still run). Thirty-two entities with all active would consume roughly 16,000 T-states in the update loop alone – a quarter of the frame budget before you have rendered a single pixel. Memory is not an issue – 32 entities at 16 bytes each is 512 bytes, still trivial.

On the Agon, you can afford larger pools. With 360,000 T-states per frame and hardware sprite rendering, 64 or even 128 entities are feasible.

18.7 Explosion and Effect Entities

Explosions, score popups, and particle effects use the same entity slots as bullets. The difference is in their update handlers: they animate through a sequence of frames and then self-destruct.

```
; Update an explosion entity
; IX = entity pointer
update_explosion:
    ; Advance animation frame
    ld  a, (ix + 6)          ; 19T  anim_frame (offset +6)
    inc a                    ; 4T
    cp  8                    ; 7T  8 frames of animation
    jr  nc, .done            ; 12/7T animation complete

    ld  (ix + 6), a          ; 19T  store new frame
    ret

.done:
    ; Animation complete -- deactivate
    ld  (ix + 12), 0         ; 19T  clear flags
    ret
```

To spawn an explosion when an enemy dies:

```
; Spawn explosion at the enemy's position
; IX currently points to the dying enemy
spawn_explosion_at_entity:
    ld  b, (ix + 1)          ; enemy's x_int
    ld  c, (ix + 3)          ; enemy's y_int
```

```

; Find a free projectile/effect slot
push ix
ld  ix, entity_array + (SLOT_PROJ_FIRST * ENTITY_SIZE)
ld  d, SLOT_PROJ_LAST - SLOT_PROJ_FIRST + 1

.find:
ld  a, (ix + 12)
bit 7, a
jr  z, .got_slot
ld  e, ENTITY_SIZE
add ix, de
dec d
jr  nz, .find
pop ix
ret                                ; no free slot -- skip explosion

.got_slot:
ld  (ix + 0), 0                    ; x_frac
ld  (ix + 1), b                    ; x_int
ld  (ix + 2), 0                    ; y_frac
ld  (ix + 3), c                    ; y_int
ld  (ix + 4), TYPE_EXPLOSION
ld  (ix + 5), 1                    ; state = active
ld  (ix + 6), 0                    ; anim_frame = 0
ld  (ix + 7), 0                    ; dx_frac = 0 (stationary)
ld  (ix + 8), 0                    ; dx_int = 0
ld  (ix + 9), 0                    ; dy_frac = 0
ld  (ix + 10), 0                   ; dy_int = 0
ld  (ix + 11), 0                   ; health = 0
ld  (ix + 12), %10010000 ; ACTIVE + VISIBLE, not COLLIDABLE
pop ix
ret

```

The pattern is always the same: find a free slot, fill in the structure, set the flags. The update handler does type-specific work. The deactivation clears the flags. The slot is reused next time something needs to spawn. This is the entire dynamic object lifecycle on the Z80 – no allocator, no garbage collector, no free list. Just an array and a flag.

18.8 Putting It All Together: The Game Skeleton

Here is the complete game skeleton that ties everything together. This is a compilable framework with all the pieces wired up: state machine, input, entity system, and the main loop.

```

ORG  $8000

; =====
; Constants
; =====

```



```

MAX_ENTITIES      EQU 16
ENTITY_SIZE       EQU 16

STATE_TITLE       EQU 0
STATE_MENU        EQU 2
STATE_GAME        EQU 4
STATE_PAUSE       EQU 6
STATE_GAMEOVER    EQU 8

TYPE_INACTIVE     EQU 0
TYPE_PLAYER       EQU 1
TYPE_ENEMY        EQU 2
TYPE_BULLET      EQU 3
TYPE_EXPLOSION    EQU 4

INPUT_RIGHT       EQU 0
INPUT_LEFT        EQU 1
INPUT_DOWN        EQU 2
INPUT_UP          EQU 3
INPUT_FIRE        EQU 4

FLAG_ACTIVE       EQU 7
FLAG_VISIBLE      EQU 4
FLAG_COLLIDABLE   EQU 2
FLAG_FACING_L     EQU 1

; =====
; Entry point
; =====
entry:
    di
    ld    sp, $C000          ; set stack (below banked memory on 128K)
                                ; NOTE: $FFFF is in banked page on 128K Spectrum,
                                ; which causes stack corruption during bank
    ↪ switches.              ; Use $C000 (or $BFFF) for 128K compatibility.

    im    1
    ei

    ; Clear entity array
    ld    hl, entity_array
    ld    de, entity_array + 1
    ld    bc, MAX_ENTITIES * ENTITY_SIZE - 1
    ld    (hl), 0
    ldir

    ; Start in title state
    ld    a, STATE_TITLE
    ld    (game_state), a

; =====

```

```

; Main loop with state dispatch
; =====
main_loop:
    halt                ; sync to frame interrupt

    ; --- Border profiling: red = active processing ---
    ld    a, 2
    out   ($FE), a

    ; Dispatch to current state
    ld    a, (game_state)
    ld    l, a
    ld    h, 0
    ld    de, state_table
    add   hl, de
    ld    e, (hl)
    inc   hl
    ld    d, (hl)
    ex    de, hl
    jp    (hl)

; Called by each state handler when done
return_to_loop:
    ; --- Border black: idle ---
    xor   a
    out   ($FE), a
    jr    main_loop

; =====
; State table
; =====
state_table:
    DW    state_title
    DW    state_menu
    DW    state_game
    DW    state_pause
    DW    state_gameover

; =====
; State: Title screen
; =====
state_title:
    call  read_input_with_edges
    ld    a, (input_pressed)
    bit   INPUT_FIRE, a
    jr    z, .wait
    ; Transition to game
    ld    a, STATE_GAME
    ld    (game_state), a
    call  init_game
.wait:

```

```

        jp    return_to_loop

; =====
; State: Game
; =====
state_game:
    call read_input_with_edges

    ; Check pause
    ; (using 'P' key -- half-row $DF, bit 0 is P, but for simplicity
    ; we check input_pressed bit 4 / FIRE as a toggle here)

    ; Update all entities
    call update_entities

    ; Render all visible entities
    call render_entities

    ; Update music
    ; call music_play          ; PT3 player -- see Chapter 11

    jp    return_to_loop

; =====
; State: Pause (minimal)
; =====
state_pause:
    call read_input_with_edges
    ld    a, (input_pressed)
    bit   INPUT_FIRE, a
    jr    z, .still_paused
    ld    a, STATE_GAME
    ld    (game_state), a
.still_paused:
    jp    return_to_loop

; =====
; State: Game Over (minimal)
; =====
state_gameover:
    call read_input_with_edges
    ld    a, (input_pressed)
    bit   INPUT_FIRE, a
    jr    z, .wait
    ld    a, STATE_TITLE
    ld    (game_state), a
.wait:
    jp    return_to_loop

; =====
; State: Menu (minimal -- expand for your game)

```

```

; =====
state_menu:
    ; A full menu would display options and handle UP/DOWN/FIRE.
    ; For this skeleton, the menu simply transitions to the title.
    ; See Exercise 2 below for adding a real menu with item selection.
    jp    state_title

; =====
; Init game: set up player and enemies
; =====
init_game:
    ; Clear entity array
    ld    hl, entity_array
    ld    de, entity_array + 1
    ld    bc, MAX_ENTITIES * ENTITY_SIZE - 1
    ld    (hl), 0
    ldir

    ; Set up player (slot 0)
    ld    ix, entity_array
    ld    (ix + 0), 0          ; x_frac = 0
    ld    (ix + 1), 128       ; x_int = 128 (centre)
    ld    (ix + 2), 0          ; y_frac = 0
    ld    (ix + 3), 160       ; y_int = 160 (near bottom)
    ld    (ix + 4), TYPE_PLAYER
    ld    (ix + 5), 1          ; state = active
    ld    (ix + 6), 0          ; anim_frame
    ld    (ix + 7), 0          ; dx_frac
    ld    (ix + 8), 0          ; dx_int
    ld    (ix + 9), 0          ; dy_frac
    ld    (ix + 10), 0         ; dy_int
    ld    (ix + 11), 3         ; health = 3
    ld    (ix + 12), %10010100 ; ACTIVE + VISIBLE + COLLIDABLE

    ; Set up 8 enemies (slots 1-8) in a formation
    ld    ix, entity_array + ENTITY_SIZE ; slot 1
    ld    b, 8                  ; 8 enemies
    ld    c, 24                 ; starting X pixel

.enemy_loop:
    ld    (ix + 0), 0          ; x_frac
    ld    (ix + 1), c          ; x_int
    ld    (ix + 2), 0          ; y_frac
    ld    (ix + 3), 32         ; y_int = 32 (near top)
    ld    (ix + 4), TYPE_ENEMY
    ld    (ix + 5), 1          ; state = active
    ld    (ix + 6), 0          ; anim_frame
    ld    (ix + 7), 0          ; dx_frac
    ld    (ix + 8), 1          ; dx_int = 1 (moving right slowly)
    ld    (ix + 9), 0          ; dy_frac
    ld    (ix + 10), 0         ; dy_int = 0

```

```

    ld    (ix + 11), 1        ; health = 1
    ld    (ix + 12), %10010100 ; ACTIVE + VISIBLE + COLLIDABLE

    ; Advance to next slot and X position
    ld    de, ENTITY_SIZE
    add   ix, de
    ld    a, c
    add   a, 28                ; 28 pixels apart
    ld    c, a
    djnz  .enemy_loop

    ret

; =====
; Input system
; =====
input_flags:      DB  0
input_prev:       DB  0
input_pressed:    DB  0

read_input_with_edges:
    ; Save previous
    ld    a, (input_flags)
    ld    (input_prev), a

    ; Read keyboard (QAOP + SPACE)
    ld    d, 0

    ; P key: half-row $DF, bit 0
    ld    bc, $DFFE
    in    a, (c)
    bit   0, a
    jr    nz, .no_right
    set   INPUT_RIGHT, d
.no_right:
    ; O key: half-row $DF, bit 1
    bit   1, a
    jr    nz, .no_left
    set   INPUT_LEFT, d
.no_left:

    ; Q key: half-row $FB, bit 0
    ld    b, $FB
    in    a, (c)
    bit   0, a
    jr    nz, .no_up
    set   INPUT_UP, d
.no_up:

    ; A key: half-row $FD, bit 0
    ld    b, $FD

```

```

    in    a, (c)
    bit   0, a
    jr    nz, .no_down
    set   INPUT_DOWN, d
.no_down:

    ; SPACE: half-row $7F, bit 0
    ld    b, $7F
    in    a, (c)
    bit   0, a
    jr    nz, .no_fire
    set   INPUT_FIRE, d
.no_fire:

    ld    a, d
    ld    (input_flags), a

    ; Compute edges
    ld    a, (input_prev)
    cpl
    ld    b, a
    ld    a, (input_flags)
    and   b
    ld    (input_pressed), a
    ret

; =====
; Entity update loop
; =====
update_entities:
    ld    ix, entity_array
    ld    b, MAX_ENTITIES

.loop:
    push  bc
    ld    a, (ix + 12)      ; flags
    bit   FLAG_ACTIVE, a
    jr    z, .skip

    ld    a, (ix + 4)       ; type
    call  update_by_type

.skip:
    ld    de, ENTITY_SIZE
    add   ix, de
    pop   bc
    djnz  .loop
    ret

; =====
; Type dispatch

```

```

; =====
type_handlers:
    DW  .nop_handler      ; TYPE_INACTIVE
    DW  update_player
    DW  update_enemy
    DW  update_bullet
    DW  update_explosion

update_by_type:
    add  a, a
    ld   l, a
    ld   h, 0
    ld   de, type_handlers
    add  hl, de
    ld   e, (hl)
    inc  hl
    ld   d, (hl)
    ex   de, hl
    jp   (hl)

.nop_handler:
    ret

; =====
; Player update
; =====
update_player:
    ld   a, (input_flags)
    bit  INPUT_RIGHT, a
    jr   z, .not_right
    ld   a, (ix + 0)
    add  a, 2                ; move right (subpixel)
    ld   (ix + 0), a
    jr   nc, .x_done_r
    inc  (ix + 1)
.x_done_r:
    res  FLAG_FACING_L, (ix + 12)
    jr   .horiz_done
.not_right:
    bit  INPUT_LEFT, a
    jr   z, .horiz_done
    ld   a, (ix + 0)
    sub  2                ; move left (subpixel)
    ld   (ix + 0), a
    jr   nc, .x_done_l
    dec  (ix + 1)
.x_done_l:
    set  FLAG_FACING_L, (ix + 12)
.horiz_done:

    ; Fire bullet on press (edge-detected)

```

```

    ld    a, (input_pressed)
    bit   INPUT_FIRE, a
    jr    z, .no_fire
    ld    b, (ix + 1)          ; player x_int
    ld    c, (ix + 3)          ; player y_int
    call  spawn_bullet
.no_fire:

    ; Animate
    ld    a, (ix + 6)
    inc   a
    and   7
    ld    (ix + 6), a
    ret

; =====
; Enemy update (simple patrol)
; =====
update_enemy:
    ; Move by dx_int (integer-only for simple patrol)
    ld    a, (ix + 8)          ; dx_int
    add   a, (ix + 1)          ; add to x_int
    ld    (ix + 1), a

    ; Bounce at screen edges
    cp    240
    jr    c, .no_bounce_r
    ld    (ix + 8), -1         ; reverse direction
    jr    .bounce_done
.no_bounce_r:
    cp    8
    jr    nc, .bounce_done
    ld    (ix + 8), 1          ; reverse direction
.bounce_done:

    ; Animate
    ld    a, (ix + 6)
    inc   a
    and   3                    ; 4-frame animation cycle
    ld    (ix + 6), a
    ret

; =====
; Bullet update
; =====
update_bullet:
    ld    a, (ix + 8)          ; dx_int
    add   a, (ix + 1)          ; add to x_int (simplified: integer movement)
    ld    (ix + 1), a

    ; Off screen?

```



```

    cp    248
    jr    nc, .deactivate
    or    a
    jr    z, .deactivate
    ret

.deactivate:
    ld    (ix + 12), 0        ; clear all flags
    ret

; =====
; Explosion update
; =====
update_explosion:
    ld    a, (ix + 6)        ; anim_frame
    inc    a
    cp    8                  ; 8 frames
    jr    nc, .done
    ld    (ix + 6), a
    ret
.done:
    ld    (ix + 12), 0
    ret

; =====
; Spawn bullet
; =====
spawn_bullet:
    ; B = x_int, C = y_int
    push    ix
    ld    ix, entity_array + (9 * ENTITY_SIZE) ; first projectile slot
    ld    d, 7                ; 7 slots to search

.find:
    ld    a, (ix + 12)
    bit    FLAG_ACTIVE, a
    jr    z, .found
    push    de                ; save loop counter in D
    ld    de, ENTITY_SIZE
    add    ix, de
    pop    de                ; restore loop counter
    dec    d
    jr    nz, .find
    ; No free slot
    pop    ix
    scf
    ret

.found:
    ld    (ix + 0), 0        ; x_frac
    ld    (ix + 1), b        ; x_int

```

```

    ld    (ix + 2), 0          ; y_frac
    ld    (ix + 3), c          ; y_int
    ld    (ix + 4), TYPE_BULLET
    ld    (ix + 5), 1          ; state
    ld    (ix + 6), 0          ; anim_frame
    ld    (ix + 9), 0          ; dy_frac = 0
    ld    (ix + 10), 0         ; dy_int = 0

    ; Direction from player facing
    ld    a, (entity_array + 12) ; player flags
    bit   FLAG_FACING_L, a
    jr    z, .right
    ld    (ix + 7), 0          ; dx_frac
    ld    (ix + 8), -4         ; dx_int = -4
    jr    .dir_done
.right:
    ld    (ix + 7), 0          ; dx_frac
    ld    (ix + 8), 4          ; dx_int = +4
.dir_done:
    ld    (ix + 11), 1         ; health = 1
    ld    (ix + 12), %10010100 ; ACTIVE + VISIBLE + COLLIDABLE

    pop   ix
    or    a                    ; clear carry
    ret

; =====
; Render entities (stub -- see Chapter 16 for sprite rendering)
; =====
render_entities:
    ; In a real game, this iterates visible entities and draws sprites.
    ; See Chapter 16 for OR+AND masked sprites, pre-shifted sprites,
    ; and the dirty-rectangle system.
    ; For this skeleton, we use a minimal attribute-block renderer.
    ld    ix, entity_array
    ld    b, MAX_ENTITIES

.loop:
    push  bc
    ld    a, (ix + 12)
    bit   FLAG_VISIBLE, a
    jr    z, .skip

    ; Draw a 1-character coloured block at entity position.
    ; For real sprite rendering, see Chapter 16 (OR+AND masks,
    ; pre-shifted sprites, compiled sprites, dirty rectangles).
    ld    a, (ix + 1)          ; x_int
    rrca                      ; /2
    rrca                      ; /4
    rrca                      ; /8 = character column
    and   $1F                 ; mask to 0-31

```

```

    ld    e, a
    ld    a, (ix + 3)      ; y_int
    rrca
    rrca
    rrca
    and   $1F              ; character row (0-23)
    ; Compute attribute address: $5800 + row*32 + col
    ld    l, a
    ld    h, 0
    add   hl, hl           ; row * 2
    add   hl, hl           ; row * 4
    add   hl, hl           ; row * 8
    add   hl, hl           ; row * 16
    add   hl, hl           ; row * 32
    ld    d, 0
    add   hl, de           ; + column
    ld    de, $5800
    add   hl, de           ; HL = attribute address

    ; Colour by type
    ld    a, (ix + 4)      ; type
    add   a, a             ; type * 2 (crude colour mapping)
    or    %01000000       ; BRIGHT bit
    ld    (hl), a          ; write attribute

.skip:
    ld    de, ENTITY_SIZE
    add   ix, de
    pop   bc
    djnz .loop
    ret

; =====
; Data
; =====
game_state:    DB    STATE_TITLE

entity_array:
    DS    MAX_ENTITIES * ENTITY_SIZE, 0 ; 256 bytes

```

This skeleton compiles, runs, and does something visible: coloured blocks move across the attribute grid. The player block responds to QAOP controls. Pressing SPACE spawns bullets that fly across the screen. Enemies bounce between the screen edges. When a bullet exits the screen, its slot frees up for the next shot.

It is ugly – attribute blocks instead of sprites, no scrolling, no sound. But the architecture is complete. Every piece from this chapter is present and wired together: the HALT-driven main loop, the state machine dispatch, the input reader with edge detection, the entity array with per-type update handlers, and the object pool for projectiles. Chapters 16 and 17 provide the rendering. Chapter 19 provides the physics and collisions. Chapter 11 provides the music. This skeleton is where they all plug in.

18.9 Agon Light 2: The Same Architecture, More Room

The Agon Light 2 uses the same fundamental game loop structure. The eZ80 runs Z80 code natively (in Z80-compatible mode or ADL mode), so the HALT-based main loop, the state machine, the entity system, and the input logic all translate directly.

The key differences:

Frame sync. The Agon uses MOS's waitvblank call (RST \$08, function \$1E) instead of HALT for frame synchronisation. The VDP generates the vertical blank signal and the MOS API exposes it.

Input. Keyboard reading goes through MOS system variables rather than direct port I/O. The half-row matrix does not exist – the PS/2 keyboard is handled by the ESP32 VDP. The input abstraction layer we built (everything funnels into `input_flags`) means the rest of the game code does not care about the difference.

Entity budget. With ~360,000 T-states per frame and hardware sprite rendering, the entity update loop is no longer a bottleneck. You could update 64 entities with complex AI and still use under 10% of the frame budget. The limiting factor on the Agon is VDP sprite count per scanline (typically 16-32 hardware sprites visible on the same line) rather than CPU time.

Rendering. The Agon's VDP handles sprite rendering. Instead of manually blitting pixels into screen memory (Chapter 16's six methods), you issue VDU commands to position hardware sprites. The CPU cost per sprite drops from ~2,400 T-states (OR+AND blit on Spectrum) to ~50-100 T-states (sending a VDU position command). This frees enormous CPU time for game logic.

Memory. The Agon has 512KB of flat memory – no banking, no contended regions. Your entity array, lookup tables, sprite data, level maps, and music can all coexist without the bank-switching gymnastics that Chapter 15 describes for the Spectrum 128K.

The practical takeaway: on the Agon, this chapter's architecture scales effortlessly. More entities, more complex state machines, more AI logic – none of it threatens the frame budget. The discipline of counting every T-state still matters (it is good engineering), but the constraints that force agonising trade-offs on the Spectrum simply do not apply.

18.10 Design Decisions and Trade-Offs

Fixed vs Variable Frame Rate

This chapter's main loop assumes a fixed 50 fps frame rate: do everything in one frame, or drop. The alternative is a variable time step: measure how long the frame took and scale all movement by delta-time. Variable time steps are standard in modern game engines but add complexity on the Z80 – you need a frame timer, multiplication by delta in every movement calculation, and careful handling of physics stability at variable rates.

For Spectrum games, fixed 50 fps is almost universally the right choice. The hardware is deterministic, the frame budget is predictable, and the simplicity of fixed-step physics (everything moves by constant amounts each frame) eliminates an entire category of bugs. If your game drops below 50 fps, the answer is to optimise until it does not – not to add a variable time step.

On the Agon, with its larger budget, you are even less likely to need variable timing. Fix the frame rate at 50 or 60 fps and keep life simple.

Entity Size: 16 Bytes vs Larger

Our 16-byte entity structure is a good fit for the data we need: positions, velocities, type, state, animation, health, flags, and bounding box dimensions for collision detection. Some commercial Spectrum games used 32 bytes or more per entity, storing additional fields like previous position (for dirty-rectangle erasure), sprite address, AI timer, and more.

The trade-off is iteration speed versus field access. Our 16-entity array takes 256 bytes (one page) and the full update loop runs in ~8,000 T-states. A 32-byte structure with 16 entities takes 512 bytes (still small) but the iteration overhead grows because IX advances by 32 each step, and indexed accesses to fields at high offsets like (IX+28) take the same 19 T-states but are harder to keep track of.

If you need more per-entity data, consider splitting the structure: a compact “hot” array (position, type, flags – the fields touched every frame) and a parallel “cold” array (sprite address, AI state, score value – fields accessed only when needed). This is the same structure-of-arrays versus array-of-structures trade-off that modern game engines face, applied at the Z80 scale.

When to Use HL Instead of IX

IX-indexed addressing is convenient but expensive: 19 T-states per access versus 7 for (HL). In the update loop (called 16 times per frame), the IX overhead is acceptable – 16 x 12 extra T-states per access = 192 T-states, negligible.

But in the rendering loop, where you might touch 4-6 entity fields for each of 8 visible sprites, the cost adds up. The technique: at the start of the render pass for each entity, copy the fields you need into registers:

```
; Copy entity fields to registers for fast rendering
ld  l, (ix + 0)      ; 19T  x_frac
ld  h, (ix + 1)      ; 19T  x_int
ld  c, (ix + 3)      ; 19T  y_int
ld  a, (ix + 6)      ; 19T  anim_frame
; Now render using H (X pixel), C (Y), A (frame)
; All subsequent accesses are register-to-register: 4T each
```

Four IX accesses at 19T = 76T up front, then the entire render routine uses 4T register accesses instead of 19T IX accesses. If the render routine touches those fields 10 times, you save (19-4) x 10 - 76 = 74 T-states per entity. Small, but over 8 entities per frame, that is 592 T-states – enough to draw another half-sprite.

Summary

- The **game loop** is HALT -> Input -> Update -> Render -> repeat. The HALT instruction synchronises to the frame interrupt, giving you exactly one frame's worth of T-states (roughly 64,000 on a Pentagon, 360,000 on the Agon).
 - A **state machine** with a jump table of handler pointers (DW state_title, DW state_game, etc.) organises the flow from title screen through gameplay to game over. Dispatch costs 73 T-states regardless of state count – constant time, clean scaling.
 - **Input reading** on the Spectrum uses IN A, (C) to poll keyboard half-rows through port \$FE. Five keys (QAOP + SPACE) cost roughly 220 T-states to read. The Kempston joystick is a single 11T port read. Edge detection (press vs hold) uses XOR between current and previous frames.
 - The **entity structure** is 16 bytes: X and Y (both 8.8 fixed-point), type, state, anim_frame, dx and dy (both int+frac), health, flags, width, height, and one reserved byte. Sixteen entities occupy 256 bytes (one page). Multiplication by 16 for index-to-address conversion is four left shifts – the simplest possible power-of-two arithmetic.
 - The **entity array** is statically allocated with fixed slot assignments: slot 0 for the player, slots 1-8 for enemies, slots 9-15 for projectiles and effects. Iteration checks the ACTIVE flag and dispatches to per-type handlers via a second jump table.
 - The **object pool** is the entity array itself. Spawning sets fields and the ACTIVE flag. Deactivation clears the flags. Free-slot search is a linear scan of the relevant slot range. Seven projectile slots handle typical fire rates without the player noticing missed spawns.
 - **IX-indexed addressing** is convenient for entity field access (19T per access) but expensive in inner loops. Copy fields to registers at the start of rendering for 4T access throughout.
 - The Agon Light 2 uses the same architecture with more headroom. MOS waitvblank replaces HALT, PS/2 keyboard replaces half-row scanning, hardware sprites replace CPU blitting. The entity update loop is no longer the bottleneck.
 - The practical skeleton in this chapter runs a state machine, 16 entities (1 player + 8 enemies + 7 bullet/effect slots), input with edge detection, per-type update handlers, and a minimal attribute-block renderer. The 16-byte entity structure carries forward directly into Chapter 19's collision detection and physics – same offsets, same flags, no refactoring needed. This skeleton is the chassis that Chapters 16 (sprites), 17 (scrolling), 19 (collisions), and 11 (sound) plug into.
-

Try It Yourself

1. **Build the skeleton.** Compile the game skeleton from section 18.8 and run it in an emulator. Use QAOP to move the player block and SPACE to fire. Watch

the coloured blocks move. Add border-colour profiling (Chapter 1) to see how much frame budget is used.

2. **Add a sixth state.** Implement STATE_SHOP between Menu and Game. The shop screen should display three items and let the player choose one with UP/DOWN and FIRE. This exercises the state machine – add the constant, the table entry, the handler, and the transition logic.
3. **Expand the entity count.** Increase MAX_ENTITIES to 32, add 16 more enemies, and measure the frame budget impact with border profiling. At what entity count does the update loop start threatening 50 fps?
4. **Implement Kempston support.** Add the Kempston joystick reader and combine it with keyboard input using OR. Test in an emulator that supports Kempston emulation (Fuse: Options -> Joysticks -> Kempston).
5. **Split hot and cold data.** Create a second “cold” array with 4 bytes per entity (sprite address, AI timer, AI state, score value). Modify the update loop to access cold data only when the entity type requires it (enemies for AI, not bullets). Measure the T-state savings.

Next: Chapter 19 – Collisions, Physics, and Enemy AI. We will add AABB collision detection, gravity, jumping, and four enemy behaviour patterns to the skeleton from this chapter.

Sources: ZX Spectrum keyboard matrix and port layout from Sinclair Research technical documentation; Kempston joystick interface specification; fixed-point arithmetic techniques from Chapter 4 (Dark/X-Trade, Spectrum Expert #01, 1997); border-colour profiling from Chapter 1; sprite rendering methods referenced from Chapter 16; AY music integration referenced from Chapter 11; Agon Light 2 MOS API documentation (Bernardo Kastrup, agon-light.com)

Chapter 19: Collisions, Physics, and Enemy AI

“Every game is a lie. Physics is faked. Intelligence is a table lookup. The player never notices because the lies are told at 50 frames per second.”

In Chapter 18, we built a game loop, an entity system that tracks sixteen objects, and an input handler. But right now our player walks through walls, floats above the ground, and the enemies stand still. A game with no collisions is a screensaver. A game with no physics is a slide puzzle. A game with no AI is a sandbox with nothing to push back.

This chapter adds the three systems that turn a tech demo into a game: collision detection, physics, and enemy AI. All three share a design philosophy: fake it well enough, fast enough, and nobody will know the difference. We build on the 16-byte entity structure introduced in Chapter 18 – X/Y positions in 8.8 fixed-point, velocity in dx/dy (each split into int+frac), type, state, flags, and bounding box dimensions.

Part 1: Collision Detection

AABB: The Only Shape You Need

Axis-Aligned Bounding Boxes. Every entity gets a rectangle defined by its position and dimensions: left edge, right edge, top edge, bottom edge. Two rectangles overlap if and only if all four of these conditions are true:

1. A's left edge is less than B's right edge
2. A's right edge is greater than B's left edge
3. A's top edge is less than B's bottom edge
4. A's bottom edge is greater than B's top edge

If any one of these conditions fails, the boxes do not overlap. This is the **early exit** that makes AABB fast: on average, most entity pairs are *not* colliding, so most checks bail out after one or two comparisons rather than doing all four.

graph TD

```
START["Check collision<br>between Entity A and Entity B"] --> X1{"A.left <  
↪ B.right?<br>(A.x < B.x + B.width)"}  
X1 -- No --> MISS["No collision<br>(clear carry, return)"]  
X1 -- Yes --> X2{"A.right > B.left?<br>(A.x + A.width > B.x)"}  
X2 -- No --> MISS  
X2 -- Yes --> Y1{"A.top < B.bottom?<br>(A.y < B.y + B.height)"}  
Y1 -- No --> MISS  
Y1 -- Yes --> COLLIDING
```



```

Y1 -- No --> MISS
Y1 -- Yes --> Y2{"A.bottom > B.top?<br>(A.y + A.height > B.y)"}
Y2 -- No --> MISS
Y2 -- Yes --> HIT["COLLISION!<br>(set carry, return)"]

```

```

style MISS fill:#dfd,stroke:#393
style HIT fill:#fdd,stroke:#933
style X1 fill:#eef,stroke:#339
style X2 fill:#eef,stroke:#339
style Y1 fill:#fee,stroke:#933
style Y2 fill:#fee,stroke:#933

```

Early exit saves cycles: Most entity pairs are far apart. The first X-overlap test rejects them in ~91 T-states. Only pairs that pass all four tests (worst case: ~270 T-states) are actual collisions. Test horizontal overlap first in side-scrollers – entities are more spread out on X than Y.




Figure 44: AABB collision detection: two bounding boxes with axis projections. Left panel shows overlapping boxes (collision), right panel shows separated boxes (no collision). Overlap is tested on both axes independently.

On the Z80, we store entity positions as 8.8 fixed-point values, but for collision detection we only need the integer part – the high byte of each coordinate. Pixel-level precision is more than enough. Here is a complete AABB collision routine:

```

; check_aabb -- Test whether two entities overlap
;
; Input:  IX = pointer to entity A
;         IY = pointer to entity B
; Output: Carry set if collision, clear if no collision
;
; Entity structure offsets (from Chapter 18):
; +0  x_frac    (low byte of 8.8 X position)

```

```

; +1 x_int      (high byte -- the pixel X coordinate)
; +2 y_frac
; +3 y_int
; +4 type
; +5 state
; +6 anim_frame
; +7 dx_frac
; +8 dx_int
; +9 dy_frac
; +10 dy_int
; +11 health
; +12 flags
; +13 width      (bounding box width in pixels)
; +14 height     (bounding box height in pixels)
; +15 (reserved)
;
; Cost: 91-270 T-states (Pentagon), depending on early exit
; Average case (no collision): ~120 T-states

```

check_aabb:

```

; --- Test 1: A.left < B.right ---
; A.left = A.x_int
; B.right = B.x_int + B.width
ld  a, (iy+1)      ; 19T B.x_int
add a, (iy+13)     ; 19T + B.width = B.right
ld  b, a           ; 4T B = B.right (save for test 2)
ld  a, (ix+1)      ; 19T A.x_int = A.left
cp  b             ; 4T A.left - B.right
jr  nc, .no_collision ; 12/7T if A.left >= B.right, no collision
                          ; --- early exit: 91T (taken, incl. .no_collision)
↪ ---

; --- Test 2: A.right > B.left ---
; A.right = A.x_int + A.width
; B.left = B.x_int
add a, (ix+13)     ; 19T A.x_int + A.width = A.right
ld  b, (iy+1)      ; 19T B.x_int = B.left
cp  b             ; 4T A.right - B.left (we need A.right > B.left)
jr  c, .no_collision ; 12/7T if A.right < B.left, no collision
jr  z, .no_collision ; 12/7T if A.right = B.left, touching but not
↪ overlapping

; --- Test 3: A.top < B.bottom ---
ld  a, (iy+3)      ; 19T B.y_int
add a, (iy+14)     ; 19T + B.height = B.bottom
ld  b, a           ; 4T
ld  a, (ix+3)      ; 19T A.y_int = A.top
cp  b             ; 4T A.top - B.bottom
jr  nc, .no_collision ; 12/7T if A.top >= B.bottom, no collision

; --- Test 4: A.bottom > B.top ---

```

```

    add a, (ix+14)      ; 19T  A.y_int + A.height = A.bottom
    ld  b, (iy+3)       ; 19T  B.y_int = B.top
    cp  b               ; 4T   A.bottom - B.top
    jr  c, .no_collision ; 12/7T
    jr  z, .no_collision ; 12/7T

    ; All four tests passed -- collision detected
    scf                ; 4T   set carry flag
    ret                ; 10T

.no_collision:
    or  a               ; 4T   clear carry flag
    ret                ; 10T

```

AABB collision detection test with two entities, border colour changes on overlap

Figure 45: AABB collision detection test with two entities, border colour changes on overlap

The IX/IY indexed addressing is convenient but expensive – 19 T-states per access versus 7 for `ld a, (hl)`. For a game with 8 enemies and 7 bullets, it is acceptable. Worst case (all four tests pass, collision detected): approximately 270 T-states. Best case (first test fails): approximately 91 T-states. For 8 enemies checked against the player, the average case is about $8 \times 120 = 960$ T-states – 1.3% of the Pentagon frame budget. Collisions are cheap.

Overflow warning: The `ADD A, (ix+13)` instructions compute $x + \text{width}$ in an 8-bit register. If an entity is positioned at $X=240$ with $\text{width}=24$, the result wraps around to 8, producing incorrect comparisons. Ensure that entity positions are clamped so that $x + \text{width}$ and $y + \text{height}$ never exceed 255 – typically by limiting the play area to leave a margin at the right and bottom edges. Alternatively, promote the comparison to 16-bit arithmetic at the cost of additional instructions.

Ordering the Tests for Fastest Rejection

The order matters. In a side-scroller, entities far apart horizontally are the common case. Testing horizontal overlap first rejects these after two comparisons. You can go further with a quick pre-rejection:

```

; Quick X-distance rejection before calling check_aabb
; If the horizontal distance between entities exceeds
; MAX_WIDTH (the widest entity), they cannot collide.

    ld  a, (ix+1)      ; 19T  A.x_int
    sub (iy+1)         ; 19T  - B.x_int
    jr  nc, .pos_dx    ; 12/7T
    neg                     ; 8T  absolute value
.pos_dx:
    cp  MAX_WIDTH      ; 7T   widest possible entity
    jr  nc, .skip      ; 12/7T too far apart, skip AABB check
    call check_aabb    ; only test close pairs
.skip:

```

This pre-rejection costs about 60 T-states, saving the 82+ T-states of the full AABB

check. On a scrolling level, typically only 2-3 enemies are close enough to need the full test.

Tile Collisions: The Tilemap as a Collision Surface

In a platformer, the player collides with the world – floors, walls, ceilings, spikes. We use the tilemap itself as a lookup table: convert pixel position to a tile index, look up the tile type, branch on the result. One array lookup replaces dozens of rectangle checks.

Assume a 32x24 tilemap with 8x8 pixel tiles (the natural Spectrum character grid):

```
; tile_at -- Look up the tile type at a pixel position
;
; Input:  B = pixel X, C = pixel Y
; Output: A = tile type (0=empty, 1=solid, 2=hazard, 3=ladder, etc.)
;
; Map is 32 columns wide, stored row-major at 'tilemap'
; Cost: ~182 T-states (Pentagon)
```

```
tile_at:
    ld    a, c                ; 4T    pixel Y
    srl   a                    ; 8T    /2
    srl   a                    ; 8T    /4
    srl   a                    ; 8T    /8 = tile row
    ld    l, a                ; 4T

    ; Multiply row by 32 (shift left 5)
    ld    h, 0                ; 7T
    add   hl, hl               ; 11T   *2
    add   hl, hl               ; 11T   *4
    add   hl, hl               ; 11T   *8
    add   hl, hl               ; 11T  *16
    add   hl, hl               ; 11T  *32

    ld    a, b                ; 4T    pixel X
    srl   a                    ; 8T    /2
    srl   a                    ; 8T    /4
    srl   a                    ; 8T    /8 = tile column
    ld    e, a                ; 4T
    ld    d, 0                ; 7T
    add   hl, de               ; 11T   row*32 + column = tile index

    ld    de, tilemap         ; 10T
    add   hl, de               ; 11T   absolute address

    ld    a, (hl)             ; 7T    tile type
    ret                                ; 10T
```

Now check the corners and edges of the entity against the tilemap:

```
; check_player_tiles -- Check player against tilemap
;
```

```
; Input: IX = player entity
; Output: Updates player position/velocity based on tile collisions
;
; We check up to 6 points around the player's bounding box,
; but bail out as soon as we find a solid tile.
```

```
check_player_tiles:
```

```
; --- Check below (feet) ---
; Bottom-left corner of player
ld  b, (ix+1)      ; 19T x_int
ld  a, (ix+3)      ; 19T y_int
add a, (ix+14)     ; 19T + height = bottom edge
ld  c, a           ; 4T
call tile_at       ; 17T+body
cp  TILE_SOLID     ; 7T
jr  z, .on_ground  ; 12/7T

; Bottom-right corner
ld  a, (ix+1)      ; 19T x_int
add a, (ix+13)     ; 19T + width
dec a              ; 4T -1 (rightmost pixel of entity)
ld  b, a
ld  a, (ix+3)
add a, (ix+14)
ld  c, a
call tile_at
cp  TILE_SOLID
jr  z, .on_ground

; Not standing on solid ground -- apply gravity
jr  .in_air
```

```
.on_ground:
```

```
; Snap Y to top of tile, clear vertical velocity
ld  a, c           ; bottom edge Y
and  %11111000     ; align to tile boundary (clear low 3 bits)
sub  (ix+14)        ; subtract height to get top-left Y
ld  (ix+3), a       ; snap y_int
xor  a
ld  (ix+9), a       ; dy_frac = 0
ld  (ix+10), a      ; dy_int = 0
set  0, (ix+12)     ; set "on_ground" flag in flags byte
jr  .check_walls
```

```
.in_air:
```

```
res  0, (ix+12)     ; clear "on_ground" flag
```

```
.check_walls:
```

```
; --- Check right (wall) ---
ld  a, (ix+1)
add a, (ix+13)      ; right edge
```

```

ld    b, a
ld    a, (ix+3)
add   a, 4           ; check midpoint vertically
ld    c, a
call  tile_at
cp    TILE_SOLID
jr    nz, .check_left

; Push out left: snap X to left edge of tile
ld    a, b
and   %11111000
dec   a
sub   (ix+13)
inc   a
ld    (ix+1), a
xor   a
ld    (ix+7), a      ; dx_frac = 0
ld    (ix+8), a      ; dx_int = 0

.check_left:
; --- Check left (wall) ---
ld    b, (ix+1)      ; left edge
ld    a, (ix+3)
add   a, 4
ld    c, a
call  tile_at
cp    TILE_SOLID
jr    nz, .check_ceiling

; Push out right: snap X to right edge of tile + 1
ld    a, b
and   %11111000
add   a, 8
ld    (ix+1), a
xor   a
ld    (ix+7), a
ld    (ix+8), a

.check_ceiling:
; --- Check above (head) ---
ld    b, (ix+1)
ld    c, (ix+3)      ; top edge
call  tile_at
cp    TILE_SOLID
ret   nz

; Hit ceiling: push down, zero vertical velocity
ld    a, c
and   %11111000
add   a, 8           ; bottom of ceiling tile
ld    (ix+3), a

```

```

xor  a
ld   (ix+9), a
ld   (ix+10), a
ret

```

The critical insight: point-in-tile lookups are $O(1)$ array accesses. Each `tile_at` call costs ~ 182 T-states. The entire tile collision system (checking feet, head, left, and right) costs roughly 800-1,200 T-states per entity, regardless of map size.

Sliding Collision Response

When the player hits a wall while moving diagonally, they should *slide*, not stop dead. Resolve collisions on each axis independently:

1. Apply horizontal velocity. Check horizontal tile collisions. If blocked, push out and zero horizontal velocity.
2. Apply vertical velocity. Check vertical tile collisions. If blocked, push out and zero vertical velocity.

This is exactly what `check_player_tiles` does – each axis is handled separately. Diagonal motion against a wall naturally becomes sliding. Most platformers apply X first (player-controlled), then Y (gravity). Experiment with both orders and feel the difference.

Part 2: Physics

What we are building is not a rigid-body simulation – it is a small set of rules that produce the *feeling* of weight and momentum. Three operations cover 90% of what a platformer needs: gravity, jumping, and friction.

Gravity: Falling Convincingly

Every frame, add a constant to the entity's vertical velocity:

```

; apply_gravity -- Add gravity to an entity's vertical velocity
;
; Input:  IX = entity pointer
; Output: dy updated (8.8 fixed-point, positive = downward)
;
; GRAVITY_FRAC and GRAVITY_INT define the gravity constant
; in 8.8 fixed-point. A good starting value: 0.25 per frame
; = $0040 (INT=0, FRAC=64, i.e. 64/256 = 0.25 pixels/frame^2)
;
; Cost: ~50 T-states (Pentagon)

```

```

GRAVITY_FRAC equ 40h      ; 0.25 pixels/frame^2 (fractional part)
GRAVITY_INT  equ 00h      ; (integer part)
MAX_FALL_INT equ 04h      ; terminal velocity: 4 pixels/frame

```

```

apply_gravity:
    ; Skip if entity is on the ground

```

```

bit  0, (ix+12)      ; 20T  check on_ground flag
ret  nz              ; 11/5T  on ground -- no gravity

; dy += gravity (16-bit fixed-point add)
ld   a, (ix+9)       ; 19T  dy_frac
add  a, GRAVITY_FRAC ; 7T
ld   (ix+9), a       ; 19T

ld   a, (ix+10)      ; 19T  dy_int
adc  a, GRAVITY_INT  ; 7T  add with carry from frac
ld   (ix+10), a      ; 19T

; Clamp to terminal velocity
cp   MAX_FALL_INT    ; 7T
ret  c               ; 11/5T  below terminal velocity, done
ld   (ix+10), MAX_FALL_INT ; 19T  clamp integer part
xor  a               ; 4T
ld   (ix+9), a       ; 19T  zero fractional part (exact clamp)
ret                      ; 10T

```

The fixed-point representation from Chapter 4 is doing the heavy lifting here. Gravity is 0.25 pixels per frame squared – a value that would be impossible to represent with integer arithmetic. In 8.8 fixed-point, it is simply \$0040. Each frame, *dy* grows by 0.25. After 4 frames, the entity is falling at 1 pixel per frame. After 16 frames, it is falling at 4 pixels per frame (terminal velocity). The acceleration curve feels natural because it *is* natural – constant acceleration is just velocity incrementing linearly.

The terminal velocity clamp prevents entities from falling so fast they skip through floors (the “tunnelling” problem). A maximum fall speed of 4 pixels per frame means the entity can never move more than half the tile height in one frame, so tile collision checks will always catch it.

Why Fixed-Point Matters Here

Without fixed-point, gravity is either 0 or 1 pixel per frame – float or stone, nothing in between. 8.8 fixed-point gives you 256 values between each integer. \$0040 (0.25) produces a gentle arc. \$0080 (0.5) feels heavy. \$0020 (0.125) feels like a moon jump. Tuning these constants is where your game finds its character. If the fixed-point fundamentals are hazy, revisit Chapter 4.

Jump: The Anti-Gravity Impulse

Jumping is the simplest physics operation in the game: set the vertical velocity to a large negative value (upward). Gravity will decelerate it, bring it to zero at the apex, and pull it back down. The jump arc is a natural parabola – no explicit arc calculation needed.

```

; try_jump -- Initiate a jump if the player is on the ground
;
; Input:  IX = player entity
; Output: dy set to -jump_force if on ground
;

```



```
; JUMP_FORCE defines the initial upward velocity in 8.8 fixed-point.
; A good starting value: -3.5 pixels/frame = $FC80
; (INT = $FC = -4 signed, FRAC = $80 = +0.5, so -4 + 0.5 = -3.5)
;
; Cost: ~50 T-states (Pentagon)
```

```
JUMP_FRAC equ 80h          ; fractional part of jump force
JUMP_INT  equ 0FCh         ; integer part (-4 signed + 0.5 frac = -3.5)
```

```
try_jump:
    ; Must be on ground to jump
    bit 0, (ix+12)          ; 20T on_ground flag
    ret z                   ; 11/5T in air -- cannot jump

    ; Set upward velocity
    ld (ix+9), JUMP_FRAC    ; 19T dy_frac
    ld (ix+10), JUMP_INT    ; 19T dy_int = -3.5 (upward)

    ; Clear on_ground flag
    res 0, (ix+12)          ; 23T

    ; (Optional: play jump sound effect here)
    ret                     ; 10T
```

With gravity at $0.25/\text{frame}^2$ and jump force at $-3.5/\text{frame}$, the player rises for 14 frames to a peak of about 24 pixels (~3 tiles), then falls for another 14 frames. Total airtime: 28 frames, just over half a second. Responsive but not twitchy.

Variable-Height Jumps

If the player releases the jump button while ascending, cut the upward velocity in half. A tap produces a short hop, a hold produces a full jump.

```
; check_jump_release -- Cut jump short if button released
;
; Input:  IX = player entity
; Output: dy halved if ascending and jump button not held
;
; Cost: ~40 T-states (Pentagon)

check_jump_release:
    ; Only relevant while ascending
    bit 7, (ix+10)          ; 20T check sign of dy_int
    ret z                   ; 11/5T not ascending (dy >= 0), skip

    ; Check if jump button is still held
    ; (assume A contains current input state from input handler)
    bit 4, a                 ; 8T bit 4 = fire/jump
    ret nz                   ; 11/5T still held, do nothing

    ; Button released -- halve upward velocity
    ; Arithmetic right shift of 16-bit dy (preserves sign)
```

```

    ld  a, (ix+10)      ; 19T  dy_int
    sra a               ; 8T   shift right arithmetic (sign-extending)
    ld  (ix+10), a      ; 19T
    ld  a, (ix+9)       ; 19T  dy_frac
    rra                ; 4T   rotate right through carry (carry from SRA
↪  above)
    ld  (ix+9), a       ; 19T
    ret                 ; 10T

```

This is a 16-bit arithmetic right shift: SRA preserves the sign on the high byte, RRA picks up the carry on the low byte. Upward velocity halves, the arc flattens. Forty T-states for a vastly better-feeling jump.

Friction: Slowing Down on the Ground

When the player releases the direction keys, they should decelerate, not stop dead. The operation is a single right-shift of horizontal velocity.

```

; apply_friction -- Decelerate horizontal movement
;
; Input:  IX = entity pointer
; Output: dx decayed toward zero
;
; Friction is applied as a right shift (divide by power of 2).
; SRA by 1 = divide by 2 (heavy friction, like rough ground)
; SRA by 1 every other frame = divide by ~1.4 (lighter friction)
;
; Cost: ~55 T-states (Pentagon)

```

```

apply_friction:
    ; Only apply friction on the ground
    bit 0, (ix+12)      ; 20T  on_ground flag
    ret z               ; 11/5T in air -- no ground friction

    ; 16-bit arithmetic right shift of dx (signed)
    ld  a, (ix+8)       ; 19T  dx_int
    sra a               ; 8T   shift right, sign-extending
    ld  (ix+8), a       ; 19T
    ld  a, (ix+7)       ; 19T  dx_frac
    rra                ; 4T   rotate right through carry
    ld  (ix+7), a       ; 19T
    ret                 ; 10T

```

Shifting right by 1 divides velocity by 2 every frame – the player stops within a few frames. For ice, apply friction less frequently:

```

; apply_friction_ice -- Light friction, every other frame
;
    ld  a, (frame_counter)
    and 1
    ret nz               ; skip odd frames
    jr  apply_friction  ; apply on even frames only

```

Vary friction by surface type – look up the tile under the entity’s feet and branch:

```

; Determine surface type
ld  b, (ix+1)      ; player X
ld  a, (ix+3)      ; player Y
add a, (ix+14)     ; + height (feet position)
inc a              ; one pixel below feet
ld  c, a
call tile_at
cp  TILE_ICE
jr  z, .ice_friction
; Default: heavy friction
call apply_friction
jr  .done
.ice_friction:
    call apply_friction_ice
.done:

```

Applying Velocity to Position

The final step: move the entity by its velocity via 16-bit fixed-point addition on each axis:

```

; move_entity -- Apply velocity to position
;
; Input:  IX = entity pointer
; Output: X and Y positions updated by dx and dy
;
; Cost: ~80 T-states (Pentagon)

move_entity:
    ; X position += dx (16-bit fixed-point add)
    ld  a, (ix+0)      ; 19T  x_frac
    add a, (ix+7)      ; 19T  + dx_frac
    ld  (ix+0), a      ; 19T

    ld  a, (ix+1)      ; 19T  x_int
    adc a, (ix+8)      ; 19T  + dx_int (with carry)
    ld  (ix+1), a      ; 19T

    ; Y position += dy (16-bit fixed-point add)
    ld  a, (ix+2)      ; 19T  y_frac
    add a, (ix+9)      ; 19T  + dy_frac
    ld  (ix+2), a      ; 19T

    ld  a, (ix+3)      ; 19T  y_int
    adc a, (ix+10)     ; 19T  + dy_int (with carry)
    ld  (ix+3), a      ; 19T
    ret                ; 10T

```

The Physics Loop

Putting it all together, the per-frame physics update for one entity looks like this:

```

; update_physics -- Full physics update for one entity
;
; Input:  IX = entity pointer
; Call order matters: gravity first, then move, then collide

update_entity_physics:
    call apply_gravity      ; accumulate downward velocity
    call apply_friction    ; decay horizontal velocity
    call move_entity       ; apply velocity to position
    call check_player_tiles ; resolve tile collisions
    ret

```

The order is deliberate: forces first, then move, then collide. This is the standard for platformers. Total cost per entity: approximately 1,000-1,500 T-states (dominated by tile collision lookups at ~182T each). For 16 entities: 16,000-24,000 T-states, about 25-33% of the Pentagon frame budget. In practice, only the player and gravity-affected enemies need full tile collision checks - bullets and effects can use simpler bounds tests.

Part 3: Enemy AI

You do not have the T-states for pathfinding or decision trees. What you have is a jump table and a state byte. That is enough.

The Finite State Machine

Every enemy has a state byte (offset +5 in our entity structure) that selects which behaviour routine runs this frame:

State	Name	Behaviour
0	PATROL	Walk back and forth between two points
1	CHASE	Move toward the player
2	ATTACK	Fire a projectile or charge
3	RETREAT	Move away from the player
4	DEATH	Play death animation, then deactivate

Transitions are simple conditions: proximity checks, cooldown timers, health thresholds. Each is a comparison or bit test - never anything expensive.

The JP Table

The core of the AI dispatcher is a **jump table** indexed by the state byte. O(1) dispatch regardless of how many states you have:

```

; ai_dispatch -- Run the AI for one enemy entity
;
; Input:  IX = enemy entity pointer
; Output: Entity state/position/velocity updated
;

```

```

; The state byte (ix+5) indexes into a table of handler addresses.
; Each handler is responsible for:
;   1. Executing this frame's behaviour
;   2. Checking transition conditions
;   3. Setting (ix+5) to the new state if transitioning
;
; Cost: ~45 T-states dispatch overhead + handler cost

; State constants
ST_PATROL equ 0
ST_CHASE  equ 1
ST_ATTACK equ 2
ST_RETREAT equ 3
ST_DEATH  equ 4

```

```

ai_dispatch:
    ld  a, (ix+5)          ; 19T  load state byte
    add a, a               ; 4T   *2 (each table entry is 2 bytes)
    ld  e, a               ; 4T
    ld  d, 0               ; 7T
    ld  hl, ai_state_table ; 10T
    add hl, de             ; 11T  HL = table + state*2
    ld  e, (hl)            ; 7T   low byte of handler address
    inc hl                 ; 6T
    ld  d, (hl)            ; 7T   high byte of handler address
    ex  de, hl             ; 4T   HL = handler address
    jp  (hl)               ; 4T   jump to handler
                                ; (handler returns via RET)

```

```

ai_state_table:
    dw  ai_patrol          ; state 0
    dw  ai_chase           ; state 1
    dw  ai_attack          ; state 2
    dw  ai_retreat         ; state 3
    dw  ai_death           ; state 4

```

The `jp (hl)` instruction costs only 4 T-states – the entire dispatch overhead is about 45 T-states regardless of state count. Note: `jp (hl)` jumps to the address *in* HL, not the address pointed to by HL. The parentheses are a Zilog notation quirk.

Patrol: The Dumb Walk

The simplest AI behaviour: walk in one direction until you reach a boundary, then turn around.

```

; ai_patrol -- Walk back and forth between two points
;
; Input:  IX = enemy entity
; Output: Position updated, state may transition to CHASE
;
; The enemy walks at a constant speed (PATROL_SPEED).
; Direction is stored in bit 1 of the flags byte:

```

```

; bit 1 = 0: moving right
; bit 1 = 1: moving left
;
; Patrol boundaries are defined per-enemy type (hardcoded
; or stored in a level table). Here we use a simple range
; check against the initial spawn position +/- PATROL_RANGE.
;
; Cost: ~120 T-states (Pentagon)

PATROL_SPEED equ 1      ; 1 pixel per frame
PATROL_RANGE equ 32     ; 32 pixels from centre point

ai_patrol:
    ; --- Move in current direction ---
    bit 1, (ix+12)      ; 20T check direction flag
    jr  nz, .move_left  ; 12/7T

.move_right:
    ld  a, (ix+1)        ; 19T x_int
    add a, PATROL_SPEED  ; 7T
    ld  (ix+1), a        ; 19T
    ; Check right boundary
    cp  PATROL_RIGHT_LIMIT ; 7T (or use spawn_x + PATROL_RANGE)
    jr  c, .check_player ; 12/7T not at edge yet
    ; Hit right edge -- turn left
    set 1, (ix+12)       ; 23T set direction = left
    jr  .check_player    ; 12T

.move_left:
    ld  a, (ix+1)        ; 19T x_int
    sub PATROL_SPEED     ; 7T
    ld  (ix+1), a        ; 19T
    ; Check left boundary
    cp  PATROL_LEFT_LIMIT ; 7T
    jr  nc, .check_player ; 12/7T
    ; Hit left edge -- turn right
    res 1, (ix+12)       ; 23T

.check_player:
    ; --- Detection: is the player nearby? ---
    ; Simple range check: |player.x - enemy.x| < DETECT_RANGE
    ld  a, (player_x)    ; 13T player x_int (cached in RAM)
    sub (ix+1)           ; 19T delta X
    jr  nc, .pos_dx      ; 12/7T
    neg                             ; 8T absolute value

.pos_dx:
    cp  DETECT_RANGE     ; 7T e.g., 48 pixels
    ret nc               ; 11/5T too far -- stay in PATROL

    ; Player detected -- transition to CHASE
    ld  (ix+5), ST_CHASE ; 19T set state = CHASE

```

```
ret                                ; 10T
```

A patrol enemy costs about 120 T-states per frame. That is trivial. Eight patrolling enemies cost under 1,000 T-states – barely a blip in the frame budget.

Chase: The Relentless Follower

The chase behaviour is simple: compute the sign of the horizontal distance between the enemy and the player, and move in that direction.

```
; ai_chase -- Move toward the player
;
; Input:  IX = enemy entity
; Output: Position updated, state may transition to ATTACK or RETREAT
;
; Cost: ~100 T-states (Pentagon)
```

```
CHASE_SPEED equ 2                ; faster than patrol
```

```
ai_chase:
; --- Move toward player ---
ld  a, (player_x)                ; 13T
sub (ix+1)                       ; 19T  dx = player.x - enemy.x
jr  z, .vertical                 ; 12/7T  same column -- skip horizontal

; Sign of dx determines direction
jr  c, .chase_left               ; 12/7T  player is to the left (dx negative)
```

```
.chase_right:
ld  a, (ix+1)                    ; 19T
add a, CHASE_SPEED               ; 7T
ld  (ix+1), a                    ; 19T
res 1, (ix+12)                   ; 23T  face right
jr  .check_attack                ; 12T
```

```
.chase_left:
ld  a, (ix+1)                    ; 19T
sub CHASE_SPEED                  ; 7T
ld  (ix+1), a                    ; 19T
set 1, (ix+12)                   ; 23T  face left
```

```
.vertical:
```

```
.check_attack:
; --- Close enough to attack? ---
ld  a, (player_x)
sub (ix+1)
jr  nc, .pos_atk
neg
```

```
.pos_atk:
cp  ATTACK_RANGE                 ; 7T  e.g., 16 pixels
jr  nc, .check_retreat           ; 12/7T  not close enough
```

```

; In attack range -- transition to ATTACK
ld  (ix+5), ST_ATTACK
ret

.check_retreat:
; --- Low health? Retreat. ---
ld  a, (ix+11)      ; 19T health
cp  RETREAT_THRESHOLD ; 7T e.g., 2 out of 8
ret  nc              ; 11/5T health OK -- stay in CHASE

; Health critical -- retreat
ld  (ix+5), ST_RETREAT
ret

```

The sign-of-dx technique: subtract, check carry. Carry set means player is to the left. Two instructions, no trigonometry, no pathfinding.

Attack: Fire and Cooldown

The ATTACK state fires a projectile, then waits for a cooldown timer. We reuse the `anim_frame` field (offset +6) as a countdown.

```

; ai_attack -- Fire projectile, then cool down
;
; Input:  IX = enemy entity
; Output: May spawn a bullet, transitions back to CHASE when ready
;
; Cost: ~60 T-states (cooldown tick) or ~150 T-states (fire + spawn)

```

```

ATTACK_COOLDOWN equ 30      ; 30 frames between shots (0.6 seconds)

```

```

ai_attack:
; --- Cooldown timer ---
ld  a, (ix+6)      ; 19T anim_frame used as cooldown
or  a              ; 4T
jr  z, .fire       ; 12/7T timer expired -- fire

; Decrement cooldown
dec (ix+6)         ; 23T
ret                ; 10T wait

```

```

.fire:
; --- Spawn a bullet ---
; Find a free slot in the entity pool (bullet type)
call find_free_entity ; returns IY = free entity, or Z flag if none
ret  z                ; no free slots -- skip this shot

; Configure the bullet entity
ld  a, (ix+1)
ld  (iy+1), a         ; bullet X = enemy X
ld  a, (ix+3)
add a, 4

```



```

    ld    (iy+3), a          ; bullet Y = enemy Y + 4 (mid-body)
    ld    (iy+4), TYPE_BULLET ; entity type
    ld    (iy+5), 0          ; state = 0 (active)

    ; Bullet direction: toward the player
    ld    a, (player_x)
    sub   (ix+1)
    jr    c, .bullet_left

.bullet_right:
    ld    (iy+8), BULLET_SPEED ; dx_int = positive
    jr    .fire_done

.bullet_left:
    ld    a, 0
    sub   BULLET_SPEED
    ld    (iy+8), a          ; dx_int = negative (two's complement)

.fire_done:
    ; Set cooldown and return to CHASE
    ld    (ix+6), ATTACK_COOLDOWN ; reset cooldown timer
    ld    (ix+5), ST_CHASE        ; back to chase state
    ret

```

The `find_free_entity` routine scans for an inactive slot using the same linear search pattern as Chapter 18's `spawn_bullet`. If the pool is full, the shot is dropped.

Retreat: The Reverse Chase

The mirror of chase – compute sign of dx, move the other way:

```

; ai_retreat -- Move away from the player
;
; Input:  IX = enemy entity
; Output: Position updated, transitions to PATROL if far enough away
;
; Cost: ~100 T-states (Pentagon)

```

```

RETREAT_DISTANCE equ 64    ; flee until 64 pixels away

```

```

ai_retreat:
    ; --- Move away from player ---
    ld    a, (player_x)
    sub   (ix+1)          ; dx = player.x - enemy.x
    jr    c, .flee_right  ; player is left, so flee right

.flee_left:
    ld    a, (ix+1)
    sub   CHASE_SPEED
    ld    (ix+1), a
    set   1, (ix+12)      ; face left (fleeing)
    jr    .check_safe

```

```

.flee_right:
    ld    a, (ix+1)
    add   a, CHASE_SPEED
    ld    (ix+1), a
    res   1, (ix+12)      ; face right (fleeing)

.check_safe:
    ; --- Far enough away? Return to patrol ---
    ld    a, (player_x)
    sub   (ix+1)
    jr    nc, .pos_ret
    neg
.pos_ret:
    cp    RETREAT_DISTANCE
    ret   c                ; not far enough -- keep fleeing

    ; Safe distance reached -- return to PATROL
    ld    (ix+5), ST_PATROL
    ret

```

Death: Animate and Remove

Health reaches zero, state becomes DEATH. The handler plays an animation, then deactivates the entity.

```

; ai_death -- Play death animation, then deactivate
;
; Input:  IX = enemy entity
; Output: Entity deactivated after animation completes
;
; Uses anim_frame as a countdown. When it reaches 0,
; the entity is marked inactive.
;
; Cost: ~40 T-states per frame

```

```

DEATH_FRAMES equ 8          ; 8 frames of death animation

```

```

ai_death:
    ld    a, (ix+6)          ; 19T  anim_frame (countdown)
    or    a                   ; 4T
    jr    z, .deactivate     ; 12/7T

    dec   (ix+6)             ; 23T  count down
    ret                                ; 10T

.deactivate:
    res   7, (ix+12)         ; 23T  clear "active" flag (bit 7 of flags)
    ret                                ; 10T

```

Once bit 7 is cleared, the entity vanishes from rendering and its slot becomes available for reuse.

Optimisation: Update AI Every 2nd or 3rd Frame

Players cannot tell the difference between 50 Hz AI and 25 Hz AI. The screen and player input run at 50 fps, but enemy decisions at 25 fps (every 2nd frame) or 16.7 fps (every 3rd) are indistinguishable. Velocity carries the entity smoothly between AI ticks.

```
; update_all_ai -- Update enemy AI on alternate frames
;
; Input:  frame_counter = current frame number
; Output: All enemies updated (on even frames only)

update_all_ai:
    ld  a, (frame_counter) ; 13T
    and 1                  ; 7T  check bit 0
    ret nz                  ; 11/5T  odd frame -- skip AI entirely

    ; Even frame -- run AI for all active enemies
    ld  ix, entity_array + ENTITY_SIZE ; skip player (entity 0)
    ld  b, MAX_ENEMIES      ; 8 enemies
.loop:
    push bc                ; 11T  save counter

    ; Check if entity is active
    bit 7, (ix+12)         ; 20T
    call nz, ai_dispatch   ; 17T + handler (only if active)

    ; Advance to next entity
    ld  de, ENTITY_SIZE    ; 10T  advance to next entity
    add ix, de             ; 15T

    pop  bc                ; 10T
    djnz .loop             ; 13/8T
    ret
```

This halves the AI cost. For 3rd-frame updating, use a modulo-3 check:

```
    ld  a, (frame_counter)
    ld  b, 3
    ; A mod 3: subtract 3 repeatedly
.mod3:
    sub b
    jr  nc, .mod3
    add a, b                ; restore: A = frame_counter mod 3
    or  a
    ret nz                  ; skip unless remainder is 0
```

The key insight: physics runs every frame for smooth movement. AI runs every 2nd or 3rd frame for decisions. The player sees fluid motion with slightly delayed reactions, and the result feels natural.

Part 4: Practical - Four Enemy Types

Four enemy types, each with distinct behaviour, wired into the entity system from Chapter 18.

1. The Walker – patrols a platform, reverses at edges. Detects the player by proximity. Chase behaviour: follow at ground level. Damage: contact only (no projectiles). Health: 1 hit.

State	Behaviour	Transition
PATROL	Walk left/right within range	Player within 48px: CHASE
CHASE	Move toward player at 2x speed	Within 16px: ATTACK
ATTACK	Pause, lunge forward	Cooldown expires: CHASE
DEATH	Flash 8 frames, deactivate	–

2. The Shooter – stands still (or patrols slowly), fires projectiles when the player is in range. Maintains distance.

State	Behaviour	Transition
PATROL	Slow walk or stationary	Player within 64px: ATTACK
ATTACK	Fire bullet, cooldown 30 frames	Player out of range: PATROL
RETREAT	Move away if player too close	Distance > 32px: ATTACK
DEATH	Explode animation, deactivate	–

3. The Swooper – moves vertically in a sine pattern (or simple up/down), dives toward the player when aligned.

```
; ai_patrol_swooper -- Vertical sine wave patrol
;
; Input:  IX = swooper entity
; Output: Position updated with vertical oscillation
;
; Uses anim_frame as the sine table index, incrementing each AI tick
;
; Cost: ~80 T-states (Pentagon)

ai_patrol_swooper:
    ; Vertical oscillation
    ld    a, (ix+6)          ; 19T  anim_frame = sine index
    inc   (ix+6)             ; 23T  advance for next frame
    ld    h, sine_table >> 8 ; 7T   sine table base (page-aligned, per Ch.4)
    ld    l, a               ; 4T   index
    ld    a, (hl)            ; 7T   signed sine value (-128..+127)
    sra   a                  ; 8T   /2 (reduce amplitude)
    sra   a                  ; 8T   /4
    add   a, (ix+3)          ; 19T   base Y + oscillation
    ld    (ix+3), a          ; 19T
```

```

; Check for dive: is player directly below?
ld  a, (player_x)
sub (ix+1)
jr  nc, .pos
neg
.pos:
cp  8                ; within 8 pixels horizontally?
ret nc              ; not aligned -- stay patrolling

; Player below and aligned -- switch to dive (CHASE)
ld  (ix+5), ST_CHASE
ret

```

The Swooper uses the sine table from Chapter 4 for vertical oscillation. When the player passes underneath, it dives.

4. The Ambusher - sits dormant until the player is very close, then activates aggressively.

```

; ai_patrol_ambusher -- Dormant until player is adjacent
;
; Input:  IX = ambusher entity
; Output: Activates if player within 16 pixels
;
; Cost: ~50 T-states (Pentagon)

```

```
AMBUSH_RANGE equ 16
```

```

ai_patrol_ambusher:
; Check proximity (Manhattan distance for cheapness)
ld  a, (player_x)
sub (ix+1)
jr  nc, .px
neg
.px:
ld  b, a                ; |dx|

ld  a, (player_y)
sub (ix+3)
jr  nc, .py
neg
.py:
add a, b                ; Manhattan distance = |dx| + |dy|
cp  AMBUSH_RANGE
ret nc                  ; too far -- stay dormant

; Player is close -- activate!
ld  (ix+5), ST_CHASE ; go straight to aggressive chase
; (Could also play an activation sound/animation here)
ret

```

Manhattan distance ($|dx| + |dy|$) costs about 30 T-states versus ~200 for Euclidean.

For proximity checks, it is good enough.

Wiring It Into the Game Loop

The complete per-frame update, building on Chapter 18:

```
game_frame:
    halt                ; wait for VBlank

    ; --- Input ---
    call read_input     ; Chapter 18

    ; --- Player physics ---
    ld    ix, entity_array    ; player is entity 0
    call handle_player_input  ; set dx from keys, try_jump from fire
    call update_entity_physics ; gravity + friction + move + tile collide

    ; --- Enemy AI (every 2nd frame) ---
    call update_all_ai

    ; --- Enemy physics (every frame) ---
    call update_all_enemy_physics

    ; --- Entity-vs-entity collisions ---
    call check_all_collisions

    ; --- Render ---
    call render_entities    ; Chapter 16 sprites
    call update_music       ; Chapter 11 AY

    jr    game_frame
```

The `check_all_collisions` routine tests player vs enemies and bullets vs entities:

```
; check_all_collisions -- Test player vs enemies, bullets vs enemies
;
; Cost: ~2,000-3,000 T-states depending on active entity count
```

```
check_all_collisions:
    ld    ix, entity_array    ; player entity
    ld    iy, entity_array + ENTITY_SIZE
    ld    b, MAX_ENEMIES + MAX_BULLETS ; 8 enemies + 7 bullets

.loop:
    push bc

    ; Skip inactive entities
    bit  7, (iy+12)
    jr   z, .next

    ; Is this an enemy? Check player vs enemy
    ld    a, (iy+4)            ; entity type
    cp    TYPE_BULLET
```

```

jr    z, .check_bullet

; Enemy: test against player
call check_aabb
jr    nc, .next          ; no collision
call handle_player_hit    ; damage player, knockback, etc.
jr    .next

.check_bullet:
; Bullet: check against all enemies (or just nearby ones)
; For simplicity, check bullet source -- don't hit the shooter
; This is handled by a "source" field or by checking type
call check_bullet_collisions

.next:
ld    de, ENTITY_SIZE
add   iy, de
pop   bc
djnz .loop
ret

```

Agon Light 2 Notes

The same physics and AI code runs unchanged on the Agon – pure Z80 arithmetic with no hardware dependencies. The Agon’s ~368,000 T-state budget means you can afford more entities (32 or 64), per-frame AI (no 2nd-frame skip needed), more collision check points, and richer state machines. Keep the physics constants identical between platforms so the game *feels* the same. The Agon VDP provides hardware sprite collision for bullet-vs-enemy checks, but tile collision remains a Z80 tilemap lookup.

Tuning Guide

The numbers in this chapter are starting points, not commandments. Here is a reference table for tuning the feel of your platformer:

Parameter	Value	Effect
GRAVITY_FRAC	\$20 (0.125)	Floaty, moon-like
GRAVITY_FRAC	\$40 (0.25)	Standard platformer feel
GRAVITY_FRAC	\$60 (0.375)	Heavy, fast-falling
JUMP_INT	\$FD (-3)	Low jump (~2 tiles)
JUMP_INT:FRAC	\$FC:\$80 (-3.5)	Medium jump (~3 tiles)
JUMP_INT	\$FB (-5)	High jump (~5 tiles)
PATROL_SPEED	1	Slow, predictable
CHASE_SPEED	2	Matches player walk speed
CHASE_SPEED	3	Faster than player – forces jumping
DETECT_RANGE	32	Short range, enemy is “dumb”
DETECT_RANGE	64	Medium range, balanced

Parameter	Value	Effect
DETECT_RANGE	128	Long range, enemy is aggressive
ATTACK_COOLDOWN	15	Rapid fire (2 shots/second at 25 Hz AI)
ATTACK_COOLDOWN	30	Moderate fire rate
ATTACK_COOLDOWN	60	Slow, deliberate
Friction shift	»1 every frame	Stops in ~3 frames (sticky)
Friction shift	»1 every 2 frames	Stops in ~6 frames (smooth)
Friction shift	»1 every 4 frames	Stops in ~12 frames (ice)

Play-test constantly. Change one number, play for thirty seconds, feel the difference. Physics tuning is not engineering – it is craft. The numbers should be in a constants block at the top of your source file, clearly labelled, easy to modify.

Summary

- **AABB collision** uses four comparisons with early exit. Most pairs are rejected after one or two tests. Cost: 91-270 T-states per pair on the Z80 (IX/IY indexed addressing dominates). Order the tests to reject the most common non-collision case first (usually horizontal). Watch for 8-bit overflow when computing $x + \text{width}$ near screen edges.
- **Tile collision** converts pixel coordinates to a tile index via right-shift and lookup. $O(1)$ per point checked, regardless of map size. Check the four corners and edge midpoints of the entity's bounding box.
- **Sliding collision response** resolves collisions on each axis independently. Apply X velocity then check X collisions; apply Y velocity then check Y collisions. Diagonal motion against a wall naturally becomes sliding.
- **Gravity** is a fixed-point addition to vertical velocity every frame: $dy += \text{gravity}$. With 8.8 format, sub-pixel values like 0.25 pixels/frame² produce smooth, natural-feeling acceleration curves.
- **Jumping** sets vertical velocity to a negative value. Gravity decelerates it, producing a parabolic arc with no explicit curve calculation. Variable-height jumps cut the velocity in half when the button is released.
- **Friction** is a right-shift of horizontal velocity: $dx \gg= 1$. Vary the frequency of application for different surface types (every frame = rough ground, every 4th frame = ice).
- **Enemy AI** uses a finite state machine with JP-table dispatch. Five states (Patrol, Chase, Attack, Retreat, Death) cover most platformer enemy behaviours. Dispatch cost: ~45 T-states regardless of state count.
- **Chase** uses the sign of $\text{player.x} - \text{enemy.x}$ for direction. Two instructions, zero trigonometry.
- **Update AI every 2nd or 3rd frame** to halve or third the CPU cost. Physics runs every frame for smooth movement; AI decisions can lag by 1-2 frames without the player noticing.
- **Four enemy types** (Walker, Shooter, Swooper, Ambusher) demonstrate how the same state machine framework produces varied behaviours by changing a few constants and one or two state handlers.
- **Total cost** for a 16-entity game (physics + collisions + AI): approximately 15,000-20,000 T-states per frame on the Spectrum (about 25-28% of the

Pentagon budget), leaving room for rendering and sound.

Try It Yourself

1. **Build the AABB test.** Place two entities on screen. Move one with the keyboard. Change the border colour when they collide. Verify the early-exit behaviour by placing entities far apart and measuring T-states with the border-colour harness from Chapter 1.
2. **Implement tile collision.** Create a simple tilemap with solid blocks and empty space. Move the player with keyboard input and gravity. Verify that the player lands on platforms, cannot walk through walls, and slides along surfaces when moving diagonally.
3. **Tune the physics.** Using the tuning guide above, adjust gravity and jump force to create three different feels: floaty (moon), standard (Mario-like), and heavy (Castlevania-like). Play each for a minute and note how the constants change the experience.
4. **Build all four enemy types.** Start with the Walker (patrol + chase only), then add the Shooter (projectiles), the Swooper (sine-wave motion), and the Ambusher (dormant activation). Test each one individually before combining them in one level.
5. **Profile the frame budget.** With all 16 entities active, use the multi-colour border profiler (Chapter 1) to visualise how much of the frame is spent on physics (red), AI (blue), collisions (green), and rendering (yellow). Adjust the AI update frequency and measure the difference.

Sources: Dark “Programming Algorithms” (Spectrum Expert #01, 1997)
– fixed-point arithmetic foundations; Game development folklore and Z80 platform knowledge; Jump table technique is standard Z80 practice documented across the ZX Spectrum development community

Chapter 20: Demo Workflow — From Idea to Compo

“Design is the complete aggregate of all demo components, both visible and concealed. Design characterises realizational, stylistic, ideological integrity.” – Introspec, “For Design,” Hype, 2015

A demo is not built in a single session of inspired coding. It is a project – one with deadlines, dependencies, creative decisions that must be locked down weeks before the party, technical gambles that pay off or do not, and a final submission that either works on the compo machine or fails in front of an audience. The distance between “I have an idea for a demo” and “it placed third at DiHalt” is measured not in lines of code but in workflow: how you organise effects, how you schedule your time, how you build and test, how you handle the inevitable moment when the music is not ready and the party is in four days.

This chapter is about that workflow. We have spent nineteen chapters on techniques – inner loops, cycle counting, compression, sound, sync. Now we step back and ask: how does a real demo come together? How do you go from a blank screen to a two-minute production that runs reliably, looks intentional, and lands at the right party on the right day?

The answers come from three sources. restorer’s making-of article for Lo-Fi Motion (Hype, 2020) provides a detailed case study of a working production pipeline – fourteen effects built in two weeks of evening coding, with a scene table system and a toolchain that any reader can replicate. Introspec’s philosophical essays on Hype – “For Design” and “MORE” (both 2015) – articulate the design thinking that separates a collection of effects from a coherent demo. And the broader making-of culture of the ZX Spectrum scene – from Eager’s detailed NFO to GABBA’s iOS video editor workflow to NHBF’s 256-byte puzzle – gives us a gallery of approaches to learn from.

20.1 What “Design” Means in a Demo

When demosceners say “design,” they do not mean graphic design. They do not mean UI layout or colour theory, although those matter. Introspec’s definition, published on Hype in January 2015, is broader and more demanding:

Design is the complete aggregate of all demo components, both visible and concealed.

This definition includes the effects the audience sees, the transitions between them, the music choice and how it synchronises with visuals, the colour palette, the pacing, the emotional arc – but also the code architecture, the memory layout, the compression strategy, the build pipeline, and the decisions about what to leave out. A demo with beautiful pixel art and terrible pacing has poor design. A demo with crude visuals but perfect musical sync and a clear emotional arc might have excellent design. A deliberately ugly demo, one that chooses its aesthetic with intention, can have outstanding design.

The implications for workflow are immediate. If design encompasses everything, then design decisions happen at every stage. The choice of assembler constrains what your build pipeline can do. The memory map determines which effects can coexist. The order in which you build effects determines what you can cut if time runs out. Every technical choice has an aesthetic consequence, and every aesthetic choice has a technical cost.

Demo production must be simultaneously bottom-up and top-down, with constant feedback between creative vision and technical reality. The workflow must support that feedback loop.

20.2 Lo-Fi Motion: A Complete Case Study

In September 2020, restorer published a making-of article for Lo-Fi Motion, a ZX Spectrum demo released at DiHalt 2020. The article is valuable not for any single technical insight but because it documents an *entire production pipeline* – from initial concept to finished binary – in enough detail to reproduce.

The Concept: “Belarusian Pixel”

Lo-Fi Motion uses attribute-resolution graphics – what restorer calls “lo-fi” rendering. Most effects work at the 32x24 attribute grid or at a doubled 32x48 resolution using half-character rows. No pixel-level rendering. The aesthetic is deliberately blocky, embracing the ZX Spectrum’s attribute grid rather than fighting it. The name says it: this is lo-fi, and the motion is the point.

This is a design decision with cascading technical benefits. Attribute-resolution effects are cheap to calculate (192 or 384 bytes per frame instead of 6,144), cheap to store (small frame buffers mean more room for compressed data), and fast to display (writing 768 bytes to attribute RAM fits easily within one frame). The lo-fi aesthetic is not a compromise – it is a choice that unlocks fourteen effects in two weeks of evening work.

The Scene Table

At the centre of Lo-Fi Motion’s architecture is the **scene table** – a data structure that drives the entire demo. Each entry in the table describes one scene:

```
; Scene table entry (conceptual structure)
```

```
scene_entry:
```

```
    DB  bank_number          ; which 16K memory bank holds this effect's code
    DW  entry_address        ; start address of the effect routine
```

```
DW  frame_duration      ; how many frames this scene runs
DB  param_byte_1        ; effect-specific parameter
DB  param_byte_2        ; effect-specific parameter
; ... additional parameters as needed
```

The demo engine reads the scene table linearly. For each entry, it pages in the specified memory bank, jumps to the entry address, and runs the effect for the specified number of frames. When the duration expires, it advances to the next entry. The entire demo – all fourteen effects, all transitions, all timing – is encoded in this one table.

This is the same architectural pattern we saw in Chapter 12’s scripting engine, stripped to its essentials. The Eager scripting engine had two levels (outer script for effects, inner script for parameter variations) and the kWORK command for asynchronous frame generation. Lo-Fi Motion’s scene table is simpler: one level, synchronous generation, no asynchronous buffering. The simplicity is the point. It works. It was built in two weeks.

The scene table pattern has a critical workflow advantage: it separates content from engine. Adding a new effect means writing the effect routine and adding one entry to the table. Reordering the demo means rearranging table entries. Adjusting timing means changing duration values. The engine code does not change. This separation means you can iterate on the demo’s structure – its pacing, its order, its timing – without touching the engine, and iterate on individual effects without touching the structure.

Fourteen Effects

Lo-Fi Motion contains approximately fourteen distinct visual effects. restorer lists them by their working names: raskolbas, slime, fire, interp, plasma, rain, dina, rtzoomer, rbars, bigpic, and several others. Each effect is a self-contained routine that renders into a virtual buffer.

The virtual buffer is a key architectural choice. Most effects do not write directly to screen memory. Instead, they render to a **1-byte-per-pixel buffer** – a block of RAM where each byte represents one attribute cell’s colour value. The buffer is typically 32 bytes wide and 24 or 48 bytes tall (for half-character resolution). After the effect renders into the buffer, a separate output routine copies the buffer to attribute RAM, performing any necessary format conversion.

This indirection costs a few hundred T-states per frame but provides two benefits. First, effects are isolated from the screen’s physical layout. An effect that renders into a linear buffer does not need to know about the attribute memory’s address structure. Second, effects can be composed: two effects can render into separate buffers, and a mixing routine can blend them before output. Lo-Fi Motion uses this for transitions – crossfading between two effects by interpolating their buffer values.

The buffer also enables the half-character resolution mode. A 32x48 buffer maps to the screen by using two attribute writes per character cell (one for the “top half” and one for the “bottom half”), exploiting the timing trick of rewriting attributes mid-scanline. This doubles the vertical resolution at the cost of more complex output code and tighter timing constraints.

The Effects Themselves

Each effect is a variation on themes from earlier chapters: **plasma** (sum-of-sines from Chapter 9), **rotozoomer** (texture-walking from Chapter 7), **fire** (cellular automaton averaging neighbours), **rain** (particle system), and **bigpic** (pre-compressed bitmap animation decompressed frame-by-frame, using techniques from Chapter 14). None of these effects is novel. The point is that at attribute resolution, every one of them is cheap enough that fourteen fit in one demo built in two weeks. The lo-fi decision is a force multiplier.

The Toolchain

restorer's toolchain for Lo-Fi Motion is a concrete answer to the question "what tools do I need to make a demo?"

Assembler: sjasmpplus. The standard Z80 macro assembler for the modern ZX scene. Memory banking (SLOT/PAGE directives), conditional assembly, macros, INCBIN for embedded data, DISPLAY for build-time diagnostics, output to .tap/.sna/.trd. The scene table, effect code, compressed data, and engine all compile in one sjasmpplus invocation.

Emulator: zemu. restorer's emulator of choice for Lo-Fi Motion. Unreal Speccy and Fuse are equally common. What matters is accurate timing and fast reload – you need to test a new build every few minutes.

Graphics: BGE 3.05 + Photoshop. BGE (Burial Graphics Editor, by Sinn/Delirium Tremens) is a ZX Spectrum-native graphics editor, widely used in the Russian scene for creating attribute-level artwork directly on the target platform. Pre-rendered PC images go through Photoshop (or Multipaint, GIMP) and custom scripts.

Scripts: Ruby. Conversion pipeline automation: images to attribute data, sine tables to binary includes, animation sequences to delta-compressed streams. Python, Perl, and Processing are equally common. What matters is that conversion is automated and repeatable.

Compression: hrust1opt. Hrust 1 with optimal parsing. The Z80 decompressor is relocatable (uses the stack for its working buffer), convenient for demos that page data in and out of banked memory.

The practical lesson: there is no single "right" toolchain. The right one is where every step from source asset to final binary is automated, changing one input regenerates all dependent outputs, and the whole build completes in seconds. Any manual step is a bug waiting to happen at 2 AM before the compo deadline.

The Build Pipeline

The tools chain together through a **Makefile** (or equivalent build script). The pipeline for Lo-Fi Motion looks approximately like this:

```

    |
    v
Ruby conversion scripts
    |
    v
Binary includes (.bin, .hru)

```

```

    |
    v
sjasmpplus assembly
    |
    v
Output binary (.trd or .sna)
    |
    v
Test in emulator (zemu)

```

Each arrow is a Makefile rule. Change a PNG, run make, and the entire chain re-executes – conversion, compression, assembly – producing a fresh binary in seconds. Load it in the emulator, watch the result, decide what to change, edit the source, run make again. This edit-build-test loop, measured in seconds, is what makes it possible to build fourteen effects in two weeks.

The Makefile also serves as documentation. Reading the build rules tells you exactly which scripts produce which outputs, which effects depend on which data files, and what the complete dependency graph looks like. When you return to the project after a six-month break, the Makefile tells you how everything fits together.

The Timeline: Two Weeks of Evenings

Lo-Fi Motion was built in approximately two weeks of evening coding sessions. restorer was working a day job. The evenings were the only available time.

This timeline is realistic for a lo-fi attribute demo, and instructive for anyone planning their first production. The breakdown looks roughly like this:

- **Days 1-2:** Engine architecture. Scene table system, virtual buffer, output routine, basic framework. Get one effect (plasma) running through the full pipeline.
- **Days 3-7:** Effects. Two to three effects per evening once the framework is solid. Each effect is 100-300 lines of assembly, rendering into the virtual buffer. Test each one individually.
- **Days 8-10:** Content. Pre-rendered images, font data, conversion scripts. This is where the Ruby scripts earn their keep.
- **Days 11-12:** Integration. All effects in the scene table, timing adjusted to the music, transitions tuned. This is where the scene table's edit-and-rebuild workflow pays off.
- **Days 13-14:** Polish and debugging. Border colours for timing visualisation (Chapter 1), fixing effects that break on edge cases, final compression pass to fit everything in memory.

The critical observation: the engine and pipeline consume the first two days. Every subsequent day benefits from that investment. If you skip the pipeline work and hardcode your first effect directly into screen memory, you save a day up front and lose a week later when you try to add a second effect and discover that nothing is modular.

20.3 Making-of Culture

The ZX Spectrum demoscene has a strong culture of documenting how demos are made. This is not universal in the broader demoscene – on many platforms, demos ship with no documentation beyond credits. On the Spectrum scene, detailed making-of articles are a tradition, and Hype (hype.retroscene.org) is the primary venue for publishing them.

Eager: The Technical NFO

When Introspec released Eager (to live) at 3BM Open Air 2015, the ZIP file included a file `_id.diz` – the traditional demo information file – that went far beyond credits and greetings. It was a technical writeup: the attribute tunnel approach, the four-fold symmetry optimisation, the digital drum hybrid technique, the asynchronous frame generation architecture. Kylearan, reviewing the demo on Pouet, wrote: “Big thanks for the nfo file alone, I love reading technical write-ups! Helps in understanding what I’m seeing/hearing, too.”

Introspec then published an even more detailed making-of article on Hype, which became the primary source for Chapters 9 and 12 of this book. The article explained not just *what* the demo does but *why* – the reasoning behind each technical decision, the constraints that drove the architecture, the creative goals that shaped the visual design.

This level of documentation serves multiple purposes. For the audience, it deepens appreciation – understanding how an effect works makes watching it more rewarding, not less. For other coders, it is education – the making-of articles on Hype are the closest thing the ZX scene has to a technical curriculum. For the author, it is a form of closure – articulating the decisions forces you to understand your own work, and the community feedback (Hype comments can be hundreds of posts deep) stress-tests your reasoning.

GABBA: A Different Workflow

diver4d’s making-of article for GABBA (2019) documents a radically different workflow from Eager’s. Where Introspec spent weeks on a scripting engine and asynchronous frame buffer, diver4d used Luma Fusion – an iOS video editor – as his synchronisation tool.

We covered the technical details in Chapter 12. The workflow insight is what matters here: diver4d recognised that frame-level audio-visual synchronisation is a *video editing* problem, not a *programming* problem. By doing the sync work in a tool designed for it, he could iterate on timing in seconds instead of minutes. The Z80 code was the implementation layer; the creative decisions happened in the video editor.

This is a general principle. Demo workflow is not about doing everything in assembly. It is about using the right tool for each task. Assembly for inner loops. Processing or Ruby for code generation. Photoshop or Multipaint for graphics. A video editor for timing. A Makefile to tie them together. The demo is the output; the tools are whatever gets you there fastest.

NHBF: The Puzzle

UriS’s making-of for NHBF (2025) documents a workflow at the opposite extreme from Lo-Fi Motion’s fourteen-effect pipeline. NHBF is a 256-byte intro – the entire program, code and data, fits in less space than a single attribute frame. The “workflow” is one person staring at a hex dump, constantly reshuffling instructions to find shorter encodings, discovering that register values from one routine happen to match the data needs of another.

We covered the specific techniques in Chapter 13. The workflow lesson is about constraint-driven creativity. UriS describes the process as “playing puzzle-like games” – an apt metaphor because the optimisation space in 256-byte coding is combinatorial. You cannot plan a path to the solution. You can only keep rearranging pieces and remain alert for serendipitous alignments. Art-Top’s discovery that register values from the screen-clearing routine matched the text string length was not planned. It was noticed.

This matters for demo workflow at any scale. Even in a full-sized demo with a proper engine and a Makefile and a scene table, there are moments when the best solution comes from stepping back and looking at the whole picture, noticing an accidental alignment between two systems that were designed independently. The puzzle-solving mindset is not exclusive to size-coding. It is a mode of thinking that improves all demo work.

20.4 The Toolchain in Detail

The ZX Spectrum demo toolchain has converged on a standard set. Here is a typical project layout:

```
main.asm           ; entry point, scene table, engine loop
engine.asm         ; scene table interpreter, buffer management
effects/
  plasma.asm       ; individual effect routines
  fire.asm
  rotozoomer.asm
sound/
  player.asm       ; music player (PT3 or custom)
  drums.asm        ; digital drum sample playback
data/
  music.pt3        ; music file (INCBIN)
  screens.zx0      ; compressed graphics (INCBIN)
  sinetable.bin    ; pre-generated lookup table (INCBIN)
Makefile
tools/
  gen_sinetable.rb  ; Ruby script: generate sine table
  convert_gfx.rb    ; Ruby script: PNG to attribute data
```

Assembler: sjasmplus

The workhorse. Memory banking via SLOT/PAGE directives, conditional assembly, macros, INCBIN for embedded data, DISPLAY for build-time diagnostics, and output

to .tap/.sna/.trd. A typical demo compiles in a single sjasmpplus invocation.

Emulators

Unreal Speccy is preferred by many Russian-scene demosceners for its deterministic timing and accurate Pentagon emulation, with TR-DOS, TurboSound, and multiple clone model support. **Fuse** is widely available on Linux and macOS. **zemu** is another option, used by restorer for Lo-Fi Motion. For source-level debugging, **DeZog** in VS Code connects to ZEsarUX and provides breakpoints, register inspection, and memory views.

Pick one emulator for primary development. Test on others before release. Demos that work in one emulator and crash in another are a party tradition best avoided.

Graphics and Code Generation

Multipaint enforces attribute constraints in real time – purpose-built for 8-bit pixel art. **Photoshop, GIMP, or Aseprite** offer creative freedom but require conversion scripts (Python, Ruby, Processing) to quantise and export. **Processing** handles generative graphics and code generation – Introspec used it to generate the chaos zoomer’s unrolled code sequences (Chapter 9).

Build Automation and CI

Your Makefile must automate the full pipeline: source assets to conversion scripts to compression to assembly. If any step requires manual intervention, it will fail at 2 AM before the deadline.

CI via GitHub Actions is increasingly common. A workflow that builds on every push catches implicit dependencies – the demo assembles on your machine but fails on a clean environment because of an undeclared tool version. Lo-Fi Motion’s source is on GitHub, published as a reference implementation: clone it, run make, get a working binary. This openness is unusual in the demoscene and valuable for learning.

Synchronisation and Compositing

The hardest part of a demo is not the effects — it is the *timing*. When to start the plasma. When to cut to the scroller. Which beat triggers the colour flash. This is synchronisation, and the ZX Spectrum scene has evolved a layered approach that combines demoscene-specific tools with general-purpose video editing.

The sync table. At the Z80 level, synchronisation is a data table:

```
dw 0,      effect_logo      ; frame 0: show logo
dw 150,    effect_plasma    ; frame 150: start plasma
dw 312,    flash_border     ; frame 312: beat hit, flash
dw 500,    effect_scroll    ; frame 500: start scroller
dw 0
```

The engine increments a frame counter each VBlank, compares it against the next entry in the table, and dispatches when the frame arrives. This is the simplest possible sync mechanism. It is also what every ZX Spectrum demo ultimately runs — regardless of how those frame numbers were determined.

The question is: how do you *find* the right frame numbers? Five approaches exist, from the simplest to the most sophisticated. (Appendix J covers each tool’s full workflow, export pipelines, and step-by-step recipes.)

Approach 1: Vortex Tracker + manual timing. Open your .pt3 in Vortex Tracker II. The bottom-right corner shows the current position (pattern, row, frame). Play the tune, note the frame numbers where beats, accents, and phrase transitions occur. Write them into your sync table. Rebuild, test, adjust. This is the approach most ZX demosceners use, including Kolnogorov (Vein): “Vortex + video editor. In Vortex the frame is shown in the bottom-right corner — I looked at which frames to hook onto, created a table with dw frame, action entries, and synced from that.”

The advantage: you hear the music and see the numbers simultaneously. The disadvantage: iterating is slow — every change requires rebuilding the demo and watching it from the beginning.

Approach 2: Video editor as sync planner. diver4d’s GABBA workflow recognised that frame-level synchronisation is a video editing problem. Capture each effect as a video clip, import the clips and music into a video editor (DaVinci Resolve, Blender VSE), scrub to find the perfect cut points, and read off the frame numbers. Kolnogorov: “I exported effect clips to video, assembled them in a video editor, attached the music track, and looked at what order the effects work best in, noting the frames where events should happen.” The important word is *looked* — this is a visual, intuitive process. (Appendix J.2–J.3 covers Blender VSE, DaVinci Resolve, and the GABBA workflow in detail.)

Approach 3: GNU Rocket. The standard sync tool across the PC and Amiga demoscenes — a tracker-like editor where columns are named parameters and rows are time steps. You set keyframes with interpolation (step, linear, smooth, ramp) and edit live while the demo runs via TCP. A Z80 client is impractical, but the workflow transfers: design sync curves in Rocket, export keyframes, convert to Z80 dw/db tables with a Python script. (Appendix J.2 describes the full Rocket → Z80 pipeline; Appendix J.7 provides a step-by-step recipe.)

Approach 4: Blender for pre-visualisation. For complex demos, storyboard effects as colour-coded strips on the VSE timeline with the music track, animate placeholder parameters in the Graph Editor, then export frame numbers and keyframe values via Blender’s Python API directly as Z80-ready data. (Appendix J.2–J.3 covers both the VSE and Graph Editor workflows.)

Approach 5: Game engines as data generators. Unity and Unreal are overkill as *demo engines* but perfect as *data generators*: VR motion capture (draw trajectories with a controller), GPU particle simulation (export positions per frame), and shader prototyping (iterate an algorithm at full speed, then translate to Z80). Blender covers most of this for non-VR work. The export pipeline is always the same: float → 8-bit fixed-point → delta-encode → transpose → compress → INCBIN. (Appendix J.4 covers the full pipeline with comparison tables and a step-by-step VR capture recipe.)

The PC demoscene has a parallel ecosystem of demo-making tools built on the same philosophy of procedural generation and extreme compression: Farbrausch’s Werkkzeug/kkrunchy (open-sourced 2012), TiXL (node-based motion graphics, MIT), Bonzomatic (live shader coding), and music synths like Sointu and WaveSabre. None target Z80 directly, but the thinking is identical — the ZX Spectrum equivalent of Werkkzeug’s node

graph is your Python build script that generates lookup tables and emits INCBIN directives. Appendix J.5 covers the history and Appendix J.6 surveys the music tools, including Furnace — a modern tracker with direct AY-3-8910 support.

FIGURE: Vortex Tracker II – frame counter

VT2 main window with pattern editor visible.
Bottom-right: position display showing pattern:row and absolute frame number.
Highlight/circle the frame counter.

Caption: "The frame number in VT2's status bar maps directly to the PT3 player's interrupt counter on the Spectrum. What you see here is what your sync table references."

Screenshot needed: open any .pt3 in VTI fork, play to a mid-song position, capture with frame number visible.

See Appendix J for pseudo-screenshots of GNU Rocket, Blender VSE, Blender Graph Editor, and TiXL, plus detailed tool descriptions and five step-by-step export recipes.

The human touch. Kolnogorov articulates a principle that all experienced demosceners understand but rarely state explicitly: “Even if we know the snare hits every 16 notes, and we flash the border every 16 notes — it will look dead and robotic. The essence of sync is that it should be deliberately uneven and broken in places.”

Algorithmic sync — trigger on every beat, fade on every phrase boundary — feels mechanical. The best demo sync follows musical *phrases*, not individual beats. Some events fire slightly before the beat (building tension). Some fire after (surprise). Some phrases have no visual change at all (creating anticipation for the next hit). This is why manual sync tables, tediously assembled by a human watching and listening, consistently produce better results than any automated system.

The practical consequence: even if you use Rocket or Blender to plan your sync, the final pass is always manual. Watch the demo with the music. Adjust frame numbers by ear. Add the off-beat hits and the deliberate silences that make the sync feel alive.

20.5 Compo Culture

A demo without a compo is a video on YouTube. A demo at a compo is a performance — shown on a big screen, with an audience, with other entries to compare against, with prizes at stake. The compo is where the work meets its audience, and the culture around compos shapes the work.

The Major Parties

The ZX Spectrum demo scene is served by a handful of recurring parties, each with its own character.

Chaos Constructions (CC) is the largest and most prestigious ZX demo event, held in Saint Petersburg, Russia. The ZX demo compo at CC draws the strongest entries: Break Space (2016), Eager’s successors, and productions from groups like Thesuper, 4D+TBK, and Placeholders. CC is where you go to compete at the highest level. The audience is large, knowledgeable, and unforgiving.

DiHalt is held in Nizhny Novgorod, Russia, and has both a summer event and a “Lite” winter edition. DiHalt tends to be more experimental than CC – the audience is welcoming to first-time entrants, and the atmosphere encourages risk-taking. Lo-Fi Motion was released at DiHalt 2020. If you are entering your first compo, DiHalt Lite is a good choice.

Multimatograf is a smaller event with a tradition of encouraging new work. The compo categories are broad, the entry requirements are minimal, and the vibe is supportive. Introspec has reviewed Multimatograf compos on Hype, sometimes critically – he holds every party to the same standard – but the event itself is welcoming to beginners.

CAFe (Creative Art Festival) is a demoscene event with a broader scope (not exclusively ZX), but the ZX categories attract strong entries. GABBA took first place at CAFe 2019.

Revision is the world’s largest demoscene event, held annually in Saarbrücken, Germany. It is not ZX-specific, but the “8-bit demo” and “oldschool” categories welcome ZX Spectrum entries. Competing at Revision means showing your work to the global demoscene – an audience of thousands, most of whom have never seen a Spectrum demo. SerzhSoft’s Megademica won the 4K intro compo at Revision 2019, proving that ZX entries can compete on the global stage.

How to Enter Your First Compo

The process is less intimidating than it sounds.

- 1. Choose a party.** Start with a smaller event – DiHalt Lite, Multimatograf, or a local party if one exists in your area. Larger parties have higher expectations, and the pressure of competing against experienced groups at CC can be counterproductive for a first entry.

- 2. Know the rules.** Each party publishes compo rules specifying: platform requirements (which Spectrum model, which emulator configuration), file format (.tap, .trd, .sna), maximum file size, whether remote entries are accepted, and submission deadlines. Read the rules. Follow the rules. A technically impressive demo that ships as a .txz when the rules require .trd will be disqualified.

- 3. Test on the target platform.** If the party runs entries on real hardware (a physical Pentagon or Scorpion), test on that hardware or on an emulator configured to match. Demos that work perfectly on one machine model and crash on another are distressingly common. The differences are subtle: contended memory timings, bank switching delays, AY chip quirks. Chapter 15 covers the machine-specific details; Chapter 5 of Introspec’s GO WEST series covers the portability pitfalls.

4. Submit early. Most parties accept remote entries via email or a web form. Submit a day early if possible. Last-minute submissions are stressful and prone to errors (uploading the wrong file, forgetting to include a required metadata file). The party version can be imperfect – many demos are updated to “final” versions after the party, fixing bugs discovered during the compo showing.

5. Write a file_id.diz or NFO. Include a text file with credits (who did what), platform requirements (which model, which mode), and – if you are willing – a brief technical description. The audience appreciates knowing what they are looking at. The scene appreciates the documentation. And you will appreciate having written it when, three years later, you try to remember how the plasma table generation works.

6. Watch the compo. If you are at the party in person, watch your demo on the big screen with the audience. The experience of seeing your work displayed publicly, of hearing the audience react, of comparing your entry to the others – this is why compos exist. If you are submitting remotely, watch the stream if one is available. Some parties publish compo recordings on YouTube afterward.

7. Do not expect to win. Your first entry is a learning experience. The goal is to finish something, submit it, and see it shown. Placing is a bonus. The feedback you get – from the audience, from other sceners, from your own reaction to seeing it on a big screen – is worth more than any prize.

Remote entries are accepted at most ZX events. Lo-Fi Motion was a remote entry at DiHalt 2020. Some parties run online-only events streamed on YouTube or Twitch. If your nearest demoscene event is a 12-hour flight away, online compos are a viable starting point.

20.6 The Community

The ZX Spectrum demoscene is small enough that most active participants know each other, and large enough that it sustains multiple active communities.

Hype (hype.retroscore.org)

The primary Russian-language forum for ZX Spectrum demoscene discussion. Founded and moderated by Introspec, it hosts the making-of articles, technical tutorials, compo reviews, and design discussions that form the core source material for this book. Threads run to hundreds of comments, with experienced coders debating cycle counts in detail. For a non-Russian speaker, browser translation tools handle the prose well enough, and Z80 assembly reads the same in every alphabet.

The culture is direct. If you post a demo with a timing bug, someone will tell you exactly which T-state is wrong. This directness produces genuine technical discussion rather than polite but unhelpful encouragement.

ZXArt (zxart.ee)

The comprehensive archive of ZX Spectrum creative works – demos, music, graphics, games, magazines, and metadata. Every production in this book can be found on

ZXArt with screenshots, credits, party results, and downloads. ZXArt also hosts digitised ZX magazines in TRD format (Spectrum Expert, Born Dead, ZX Format), containing the original articles that established the techniques this book teaches.

Pouet (pouet.net)

The global demoscene production database. For the ZX scene, Pouet bridges to the wider community – ZX demos rated by people who primarily watch PC or Amiga productions. The perspective shift is valuable: a technically brilliant inner loop that impresses Hype readers might be invisible to a Pouet commenter who focuses on visual impact and music sync. Pouet also hosts NFO files – when you cannot find a making-of article on Hype, check the NFO on Pouet.

20.7 Project Management for Demo Makers

Demo making is project management. The project has a deadline (the party date), deliverables (the final binary), dependencies (music, graphics, effects, engine), and usually a team of contributors with competing priorities. Managing this is not glamorous, but it is what separates finished demos from abandoned prototypes.

The Minimum Viable Demo

Start with the simplest possible version of your demo that is complete – not polished, not impressive, but complete. One effect, one piece of music, a proper beginning and end. Get this running end-to-end through the full build pipeline within the first few days. This is your safety net. If everything goes wrong – if the complex effect you planned does not work, if the musician is late with the final track, if your hard drive dies a week before the party – you have something to submit.

Then iterate. Add effects one at a time. Replace the placeholder music when the final track arrives. Add transitions, polish timing, optimise memory usage. Each iteration produces a complete, submittable demo that is better than the last. At any point, you can stop and submit what you have.

This incremental approach is how Lo-Fi Motion was built. restorer did not write fourteen effects and then stitch them together. He built the engine and one effect, verified they worked, then added effects one by one. Each evening's work produced a slightly better demo. If he had run out of time at ten effects instead of fourteen, the demo would still have been complete and submittable.

Working with Collaborators

Most demos are collaborations. Three principles keep them on track:

Establish the data format early. The musician needs to know: PT3 or custom player? Single AY or TurboSound? How are drum triggers signalled? The artist needs to know: attribute resolution or pixel resolution? Colour constraints? Maximum file size? Receiving a TurboSound composition when your engine supports only single AY is a disaster, and it is your fault for not specifying constraints.

Communicate the timeline. If the party is in four weeks, tell the musician you need the track in two. The buffer is for integration, debugging, and surprises.

Provide placeholders. Use a placeholder .pt3 with the right tempo until the final track arrives. Use programmer art until the final graphics arrive. The engine should never depend on final assets. When real assets arrive, drop them into the pipeline and rebuild.

Debugging and Testing

Demo bugs are uniquely painful because they manifest in front of an audience. A crash during the compo showing is both a technical failure and a social embarrassment. Testing is not optional.

Test on multiple emulators. Each emulator has slightly different timing, memory initialisation, and AY behaviour. A demo that works in Unreal Speccy but crashes in Fuse probably has a timing or memory assumption that is valid on Pentagon but not on standard Spectrum.

Test from a cold start. Clear all memory before loading the demo. Do not assume any register values or memory contents from a previous program. If your demo works after running a previous demo but crashes from a fresh boot, you have an initialisation bug.

Test the compo file, not the development binary. The file you submit should be the exact file you tested. Not a "quickly recompiled" version with a last-minute fix. Last-minute fixes introduce last-minute bugs.

Use border colours for timing. The technique from Chapter 1: set the border to different colours at different points in the frame loop. If the border flash extends into the visible area, your code is too slow. If it does not, you have margin. This is the fastest way to verify that an effect fits within the frame budget.

20.8 Transcending the Platform: Introspec's "MORE"

In February 2015, Introspec published a short essay on Hype titled simply "MORE." It is not a technical article. It contains no code, no cycle counts, no inner loops. It is a challenge to the ZX Spectrum scene – and by extension, to everyone who works within hardware constraints.

The argument is that the best demos are not the ones that do the most impressive things *despite* the platform limitations. They are the ones that transcend the platform entirely – that create experiences which would be meaningful on any hardware. The platform's constraints shape the technique but should not limit the ambition.

Two pixels suffice to tell a story.

This is Introspec's most quoted line. It means: the artistic content of a demo is not determined by its resolution, its colour depth, its polygon count, or its sample rate. Two pixels – two attribute cells, two dots on a 32x24 grid – can tell a story if the timing is right, the context is clear, and the intention is genuine. The technology serves the art, not the other way around.

Introspec references James Houston’s “Big Ideas (Don’t Get Any)” – a video where a Sinclair ZX Spectrum, dot matrix printer, and other obsolete hardware perform a Radiohead song. The project is moving not because of the technical achievement but because the choice of hardware *means* something. The obsolescence is the point. The fragility is the beauty.

The practical implication: technique is necessary but not sufficient. You can master every effect in this book and still produce a demo nobody remembers. What makes a demo memorable is not what it does but what it says. Even an abstract demo has a personality: its pacing says something about tension and release; its colour palette evokes a mood; its music choice creates emotional context. The coder who treats these as afterthoughts produces a tech demo. The coder who treats them as design decisions produces a demo.

Lo-Fi Motion embraced its lo-fi aesthetic as identity. Eager turned the 32x24 grid from constraint into creative choice. NHBf found beauty in the puzzle of 256 bytes. In each case, the limitation became the medium.

This is what “MORE” demands. Not more polygons, not more colours, not more effects. More ambition. More intention. More willingness to treat a ZX Spectrum demo as an art form.

20.9 Your First Demo: A Practical Roadmap

For the reader who has followed this book from Chapter 1 and wants to make a demo, here is a concrete path.

```
graph TD
    IDEA["Idea<br>(visual concept, mood, music)"] -->
    ↪ PROTO["Prototype<br>(verify/ HTML/JS, or<br>quick Z80 test)"]
    PROTO --> IMPL["Z80 Implementation<br>(sjasmpplus, effect code)"]
    IMPL --> TIME{"Fits in<br>frame budget?"}
    TIME -- No --> OPT["Optimise<br>(unroll, precompute,<br>reduce scope)"]
    OPT --> IMPL
    TIME -- Yes --> POLISH["Polish<br>(transitions, sync to music,<br>colour palette)"]
    ↪ POLISH --> PARTY["Party Version<br>(submit to compo)"]
    PARTY --> FINAL["Final Version<br>(fix bugs, add credits,<br>test on hardware)"]

    IMPL --> |"border-colour timing"| TIME
    POLISH --> |"scene table reorder"| POLISH

    style IDEA fill:#ffd,stroke:#993
    style PROTO fill:#ddf,stroke:#339
    style IMPL fill:#dfd,stroke:#393
    style OPT fill:#fdd,stroke:#933
    style PARTY fill:#fdf,stroke:#939
    style FINAL fill:#dff,stroke:#399
```

The iterative loop: The path from implementation to timing check and back is where most development time is spent. The prototype stage

(HTML/JS or quick Z80 sketch) validates the visual concept before committing to full implementation. The scene table makes reordering effects trivial during the polish phase.

Demo framework with effect slots driven by a scene table, showing the engine cycling through multiple visual effects

Figure 46: Demo framework with effect slots driven by a scene table, showing the engine cycling through multiple visual effects

Week 1: Foundation

1. **Set up the toolchain.** Install sjasmpplus, choose an emulator (Unreal Speccy, Fuse, or ZEsarUX), set up a project directory with a Makefile. Verify that you can assemble a minimal program and run it in the emulator.
2. **Build the scene table engine.** Write a minimal engine that reads a scene table and calls effect routines for the specified duration. Start with Lo-Fi Motion's architecture: bank number, entry address, frame count. Get it working with a single dummy effect (fill the screen with a colour, increment the colour each frame).
3. **Add music.** Integrate a PT3 player into your IM2 interrupt handler (Chapter 11). Drop in any .pt3 file as a placeholder. Verify that music plays while the dummy effect runs.

Week 2: Effects

4. **Build your first real effect.** Attribute plasma is the natural starting point – it is cheap, visually rich, and well-understood (Chapter 9). Render into a virtual buffer and copy to attribute RAM.
5. **Build your second effect.** Fire, rotozoomer, rain, colour bars – pick one from the effects covered in Part II. Two effects and a transition between them constitute a minimal demo.
6. **Add a transition.** A simple crossfade between two attribute buffers: interpolate the colour values over 25-50 frames. Or a hard cut synchronised to a beat in the music.

Week 3: Polish

7. **Replace placeholder music.** If you have a musician collaborator, integrate the final track. If not, spend time choosing a .pt3 that fits your demo's mood and pace.
8. **Adjust timing.** This is where the scene table earns its value. Rearrange effects, adjust durations, align transitions to musical events. Rebuild and test repeatedly.
9. **Add a beginning and an end.** A loading screen (compressed with ZX0, Chapter 14), an opening title, a closing credits screen. First and last impressions matter.

Week 4: Release

10. **Test.** Multiple emulators. Cold boot. The exact file you will submit.
11. **Write the NFO.** Credits, platform requirements, greetings, and – if you are feeling generous – a technical description of how the effects work. Future you will be grateful.
12. **Submit.** Choose a party. Follow the rules. Upload the file. Then watch the compo and enjoy seeing your work on screen.

Your first entry is unlikely to place. Treat it as a learning exercise: the feedback from seeing your work on the big screen and comparing it to other entries is more valuable than any prize. Each subsequent demo will be better because you will know what to fix.

Summary

- **Design is everything.** Introspec defines demo design as “the complete aggregate of all demo components, both visible and concealed” – the code architecture, the memory layout, the pacing, and the emotional arc, not just the visual effects.
- **Lo-Fi Motion provides a replicable production template:** a scene table drives the demo structure, effects render into virtual 1-byte-per-pixel buffers, and the toolchain (sjasmpplus + zemu + Ruby scripts + hrust1opt) chains together through a Makefile. Fourteen effects were built in two weeks of evening work.
- **The scene table pattern** separates content from engine. Adding, removing, or reordering effects means editing a data table, not restructuring code. This supports rapid iteration on pacing and structure.
- **Making-of culture is a strength of the ZX scene.** Detailed technical writeups – from Eager’s NFO to GABBA’s video-editor workflow to NHBF’s 256-byte puzzle – serve as education, documentation, and community building.
- **The standard toolchain** converges on sjasmpplus (assembler), Unreal Speccy or Fuse (emulator), BGE or Multipaint (graphics), Ruby or Python scripts (conversion and code generation), ZX0 or hrust1opt (compression), and a Makefile (build automation). CI via GitHub Actions is increasingly common.
- **Synchronisation** is the hardest part of a demo. The layered approach: determine frame numbers in Vortex Tracker or a video editor (DaVinci Resolve, Blender VSE), optionally plan interpolated parameter curves in GNU Rocket, export to Z80 dw frame, action tables. The final pass is always manual — algorithmic sync feels robotic; human-placed sync follows phrases, not beats. (Appendix J covers all sync tools, data generation pipelines, and step-by-step export recipes.)
- **Compo culture** centres on events like Chaos Constructions, DiHalt, Multi-matograf, CAFe, and Revision. Entering your first compo requires choosing an appropriate event, following the rules, testing thoroughly, and submitting early.

- **The community** lives on Hype (technical discussion, making-of articles), ZXArt (production archive), and Pouet (global demoscene database).
- **Project management matters.** Build the minimum viable demo first, then iterate. Establish data formats with collaborators early. Test on multiple emulators from a cold start. Submit the exact file you tested.
- **Introspec's "MORE"** challenges demo makers to transcend platform limitations: "Two pixels suffice to tell a story." The technology serves the art, not the other way around. The best demos are not the most technically impressive – they are the ones where every component, visible and concealed, serves a coherent creative vision.

Next: Chapter 21 – Full Game: ZX Spectrum 128K. We move from demos to games, integrating everything from Parts I through V into a complete side-scrolling platformer.

Sources: restorer, "Making of Lo-Fi Motion," Hype, 2020 (hype.retroscore.org/blog/demo/1023.html); Introspec, "For Design," Hype, 2015 (hype.retroscore.org/blog/demo/64.html); Introspec, "MORE," Hype, 2015 (hype.retroscore.org/blog/demo/87.html); Introspec, "Making of Eager," Hype, 2015 (hype.retroscore.org/blog/demo/261.html); diver4d, "Making of GABBA," Hype, 2019 (hype.retroscore.org/blog/demo/948.html); UriS, "NHBF Making-of," Hype, 2025 (hype.retroscore.org/blog/dev/1120.html)

Chapter 21: Full Game - ZX Spectrum 128K

“The only way to know if your engine works is to ship a game.”

You have sprites (Chapter 16). You have scrolling (Chapter 17). You have a game loop and entity system (Chapter 18). You have collisions, physics, and enemy AI (Chapter 19). You have AY music and sound effects (Chapter 11). You have compression (Chapter 14). You have 128K of banked RAM and know how to address every byte of it (Chapter 15).

Now you must put all of it into a single binary that loads from tape, shows a loading screen, presents a menu, plays five levels of a side-scrolling platformer with four enemy types and a boss, tracks high scores, and fits on a .tap file.

This is the integration chapter. No new techniques appear here. Instead, we face the problems that only emerge when every subsystem must coexist: memory contention between graphics banks and code, frame budgets that overflow when scrolling and sprites and music and AI all demand their share, build systems that must coordinate a dozen data conversion steps, and the thousand small decisions about what goes where in 128K of banked memory.

The game we are building is called *Ironclaw* – a five-level side-scrolling platformer starring a mechanical cat navigating a series of increasingly hostile factory floors. The genre is deliberate: side-scrolling platformers demand every subsystem simultaneously and leave nowhere to hide. If the scrolling stutters, you see it. If the sprite rendering overflows the frame, you feel it. If the collision detection fails, the player falls through the floor. A platformer is the hardest integration test a Z80 game engine can face.

21.1 Project Architecture

Before writing a single line of Z80, you need a directory structure that scales. A 128K game with five levels, a tileset, sprite sheets, a music score, and sound effects generates dozens of data files. If you do not organise them from the start, you will drown.

Directory Layout

```

src/
  main.a80      -- entry point, bank switching, state machine
  render.a80    -- tile renderer, scroll engine
  sprites.a80   -- sprite drawing routines (OR+AND masked)
  entities.a80  -- entity update, spawning, despawning
  physics.a80   -- gravity, friction, jump, collision response
  collisions.a80 -- AABB and tile collision checks
  ai.a80        -- enemy FSM: patrol, chase, attack, retreat, death
  player.a80    -- player input, state, animation
  hud.a80       -- score, lives, status bar
  menu.a80      -- title screen, options, high scores
  loader.a80    -- loading screen, tape/esxDOS loader
  music_driver.a80 -- PT3 player, interrupt handler
  sfx.a80       -- sound effects engine, channel stealing
  esxdos.a80    -- DivMMC file I/O wrappers
  banks.a80     -- bank switching macros and utilities
  defs.a80      -- constants, memory map, entity structure
data/
  levels/       -- level tilemaps (compressed)
  tiles/        -- tileset graphics
  sprites/      -- sprite sheets (pre-shifted)
  music/        -- PT3 music files
  sfx/          -- SFX definition tables
  screens/      -- loading screen, title screen
tools/
  png2tiles.py  -- PNG tileset converter
  png2sprites.py -- PNG sprite sheet converter (generates shifts)
  map2bin.py    -- Tiled JSON/TMX to binary tilemap
  compress.py   -- wrapper around ZX0/Pletter compression
build/         -- compiled output (gitignored)
Makefile       -- the build system

```

Every source file focuses on one subsystem. Every data file passes through a conversion pipeline before it reaches the assembler. The `tools/` directory holds Python scripts that convert artist-friendly formats (PNG images, Tiled editor maps) into assembler-ready binary data.

The Build System

The Makefile is the spine of the project. It must:

1. Convert all graphics from PNG to binary tile/sprite data
2. Convert level maps from Tiled format to binary tilemaps
3. Compress level data, graphics banks, and music with the appropriate compressor
4. Assemble all source files into a single binary
5. Generate the final `.tap` file with the correct loader

```

ASM      = sjasmpplus
COMPRESS = zx0
PYTHON   = python3

```

```

# Data conversion
data/tiles/tileset.bin: data/tiles/tileset.png
$(PYTHON) tools/png2tiles.py $< $@

data/sprites/player.bin: data/sprites/player.png
$(PYTHON) tools/png2sprites.py --shifts 4 $< $@

data/levels/level%.bin: data/levels/level%.tmx
$(PYTHON) tools/map2bin.py $< $@

# Compression (ZX0 for level data -- good ratio, small decompressor)
data/levels/level%.bin.zx0: data/levels/level%.bin
$(COMPRESS) $< $@

# Compression (Pletter for graphics -- faster decompression)
data/tiles/tileset.bin.plt: data/tiles/tileset.bin
pletter5 $< $@

# Assembly
build/ironclaw.tap: src/*.a80 data/levels/*.zx0 data/tiles/*.plt \
                    data/sprites/*.bin data/music/*.pt3
$(ASM) --fullpath src/main.a80 --raw=build/ironclaw.tap

.PHONY: clean
clean:
    rm -rf build/ data/**/*.*.bin data/**/*.*.zx0 data/**/*.*.plt

```

The key insight is the data pipeline. The artist exports a PNG tileset from Aseprite. The png2tiles.py script slices it into 8x8 or 16x16 tiles, converts each to the Spectrum's interleaved pixel format, and writes a binary blob. The level designer exports a Tiled .tmx map. The map2bin.py script extracts the tile indices and writes a compact binary tilemap. The compressor squeezes each blob. Only then does the assembler INCBIN the result into the appropriate memory bank.

This pipeline means the game's content is always in editable form (PNG, TMX), and the build system handles every conversion automatically. Change a tile in the PNG, type make, and the new tile appears in the game.

21.2 Memory Map: 128K Bank Assignments

The ZX Spectrum 128K has eight 16KB RAM banks, numbered 0 through 7. At any moment, the CPU sees a 64KB address space:

```

$4000-$7FFF   Bank 5 (always) -- screen memory (normal screen)
$8000-$BFFF   Bank 2 (always) -- typically code
$C000-$FFFF   Switchable -- banks 0-7, selected via port $7FFD

```

Banks 5 and 2 are hardwired to \$4000 and \$8000 respectively. Only the top 16KB window (\$C000-\$FFFF) is switchable. The bank selection register at port \$7FFD also controls which screen is displayed (bank 5 or bank 7) and which ROM page is

active.

```
; Port $7FFD layout:
; Bit 0-2: Bank number for $C000-$FFFF (0-7)
; Bit 3:   Screen select (0 = bank 5 normal, 1 = bank 7 shadow)
; Bit 4:   ROM select (0 = 128K editor, 1 = 48K BASIC)
; Bit 5:   Disable paging (PERMANENT -- cannot be undone without reset)
; Bits 6-7: Unused

; Switch to bank N at $C000
; Input: A = bank number (0-7)
; Preserves: all registers except A
switch_bank:
    or    %00010000          ; ROM 1 (48K BASIC) -- keep this set
    ld    (last_bank_state), a
    ld    bc, $7FFD
    out   (c), a
    ret

last_bank_state:
    db    %00010000          ; default: bank 0, normal screen, ROM 1
```

The critical rule: **always store your last write to \$7FFD** in a shadow variable. Port \$7FFD is write-only – you cannot read back the current state. If you need to change one bit (say, switch the screen) without disturbing the bank selection, you must read your shadow variable, modify the bit, write the result to both the port and the shadow.

Ironclaw Bank Allocation

Here is how Ironclaw distributes its 128KB across the eight banks:

```
Tileset graphics (compressed)
Decompression buffer

Bank 1 ($C000) -- Level data: tilemaps for levels 3-5 (compressed)
                Boss level data and patterns
                Enemy spawn tables

Bank 2 ($8000) -- FIXED: Main game code
                Player logic, physics, collisions
                Sprite routines, entity system
                State machine, HUD
                ~ 14KB code, 2KB tables/buffers


Bank 3 ($C000) -- Sprite graphics (pre-shifted x4)
                Player: 6 frames x 4 shifts = 24 variants
                Enemies: 4 types x 4 frames x 4 shifts = 64 variants
                Projectiles, particles, pickups
                ~ 12KB total

Bank 4 ($C000) -- Music: PT3 song data (title, levels 1-3)
                PT3 player code (resident copy)
```

Bank 5 (\$4000) -- FIXED: Normal screen
 Pixel data \$4000-\$57FF (6,144 bytes)
 Attributes \$5800-\$5AFF (768 bytes)
 Remaining ~9KB: interrupt handler, screen buffers

Bank 6 (\$C000) -- Music: PT3 song data (levels 4-5, boss, game over)
 SFX definition tables
 SFX engine code

Bank 7 (\$4000) -- Shadow screen (used for double buffering)
 Also usable as 16KB data storage when
 not actively double-buffering



Key: Banks 2 and 5 are always visible (hardwired).
 Only \$C000–\$FFFF is switchable.
 Port \$7FFD is write-only – always shadow its state!

Several things to notice about this layout:

Code lives in bank 2 (fixed). Because bank 2 is always mapped at \$8000-\$BFFF, your main game code is always accessible. You never need to page in code – only data. This eliminates the most dangerous class of banking bug: calling a routine that has been paged out.

Sprite graphics in bank 3, separate from level data in banks 0-1. When rendering a frame, the renderer needs both tile graphics (for the scrolling background) and sprite graphics (for entities). If both were in the same switchable bank, you would need to page back and forth mid-render. By placing them in separate banks,

you can page in the tile data, render the background, then page in the sprite data and render all entities, with only two bank switches per frame.

Music is split across banks 4 and 6. The PT3 player runs during the interrupt handler, which fires once per frame. The interrupt handler must page in the music bank, update the AY registers, and page back to whatever bank the main loop was using. Splitting music across two banks means the interrupt handler must know which bank contains the current song. We handle this with a variable:

```
current_music_bank:
    db    4                ; bank 4 by default

im2_handler:
    push af
    push bc
    push de
    push hl
    push ix
    push iy                ; IY must be preserved -- BASIC uses it
                           ; for system variables, and PT3 players
                           ; typically use IY internally

    ; Save current bank state
    ld    a, (last_bank_state)
    push af

    ; Page in music bank
    ld    a, (current_music_bank)
    call switch_bank

    ; Update PT3 player -- writes AY registers
    call pt3_play

    ; Check for pending SFX
    call sfx_update

    ; Restore previous bank
    pop   af
    ld    (last_bank_state), a
    ld    bc, $7FFD
    out   (c), a

    pop   iy
    pop   ix
    pop   hl
    pop   de
    pop   bc
    pop   af
    ei
    reti
```

The shadow screen in bank 7 is available for double-buffering during scroll updates (as described in Chapter 17). When you are not actively double-buffering

- during the menu, between levels, during cutscenes - bank 7 is 16KB of free storage. Ironclaw uses it to hold the decompressed tilemap of the current level during gameplay, freeing the switchable banks for graphics and music.

The Stack

The stack lives at the top of bank 2's address space, growing downward from \$BFFF. With ~14KB of code starting at \$8000, the stack has approximately 2KB of room - more than enough for normal call depth, but you must be vigilant. Deep recursion is not an option. If you are using stack-based sprite output (Chapter 16's PUSH method), remember that you are borrowing the stack pointer and must save and restore it with interrupts disabled.

21.3 The State Machine

A game is not one program. It is a sequence of modes - title screen, menu, gameplay, pause, game over, high scores - each with different input handling, different rendering, and different update logic. Chapter 18 introduced the state machine pattern. Here is how Ironclaw implements it at the top level.

```
; Game states
STATE_LOADER      equ 0
STATE_TITLE       equ 1
STATE_MENU        equ 2
STATE_GAMEPLAY    equ 3
STATE_PAUSE       equ 4
STATE_GAMEOVER    equ 5
STATE_HISCORE     equ 6
STATE_LEVELWIN    equ 7

; State handler table -- each entry is a 2-byte address
state_table:
    dw  state_loader      ; 0: loading screen + init
    dw  state_title       ; 1: title screen with animation
    dw  state_menu        ; 2: main menu (start, options, hiscores)
    dw  state_gameplay    ; 3: in-game
    dw  state_pause       ; 4: paused
    dw  state_gameover    ; 5: game over sequence
    dw  state_hiscore     ; 6: high score entry
    dw  state_levelwin    ; 7: level complete, advance

current_state:
    db  STATE_LOADER

; Main loop -- called once per frame after HALT
main_loop:
    halt                ; wait for frame interrupt

    ; Dispatch to current state handler
    ld  a, (current_state)
```

```

add  a, a          ; x2 for word index
ld   l, a
ld   h, 0
ld   de, state_table
add  hl, de
ld   a, (hl)
inc  hl
ld   h, (hl)
ld   l, a          ; HL = handler address
jp   (hl)          ; jump to handler

```

; Each handler ends with: `jp main_loop`

Each state handler owns the frame completely. The gameplay handler runs input, physics, AI, rendering, and HUD updates. The menu handler reads input and draws the menu. The pause handler simply waits for the unpause key, displaying a “PAUSED” overlay.

State transitions happen by writing a new value to `current_state`. The transition from `STATE_GAMEPLAY` to `STATE_PAUSE` requires no cleanup – the game state is untouched, and returning to `STATE_GAMEPLAY` resumes exactly where you left off. But the transition from `STATE_GAMEOVER` to `STATE_HISCORE` requires checking whether the player’s score qualifies, and the transition from `STATE_LEVELWIN` to `STATE_GAMEPLAY` requires loading and decompressing the next level’s data.

21.4 The Gameplay Frame

This is where integration happens. During `STATE_GAMEPLAY`, every frame must execute the following, in order:

- | | |
|-----------------------------|--|
| 2. Update player physics | ~800 T-states |
| 3. Update player state | ~400 T-states |
| 4. Update enemies (AI+phys) | ~4,000 T-states (8 enemies) |
| 5. Check collisions | ~2,000 T-states |
| 6. Update projectiles | ~500 T-states |
| 7. Scroll the viewport | ~8,000-15,000 T-states (depends on method) |
| 8. Render background tiles | ~12,000 T-states (exposed column/row) |
| 9. Erase old sprites | ~3,000 T-states (background restore) |
| 10. Draw sprites | ~8,000 T-states (8 entities x ~1,000 each) |
| 11. Update HUD | ~1,500 T-states |
| 12. [Music plays in IM2] | ~3,000 T-states (interrupt handler) |

Total: ~43,400-50,400 T-states

On a Pentagon with 71,680 T-states per frame, that leaves 21,000-28,000 T-states of slack. That sounds comfortable, but it is deceptive. Those estimates are averages. When four enemies are on screen and the player is jumping across a gap with projectiles flying, the worst case can spike 20-30% above the average. Your slack is your safety margin.

The order matters. Input must come first – you need the player’s intent before simulating physics. Physics must precede collision detection – you need to know

where entities want to move before checking whether they can. Collision response must precede rendering – you need final positions before drawing anything. And sprites must be drawn after the background, because sprites overlay the tiles.

Reading Input

```
; Read keyboard and Kempston joystick
; Returns result in A: bit 0=right, 1=left, 2=down, 3=up, 4=fire
read_input:
    ld    d, 0                ; accumulate result

    ; Kempston joystick (active high)
    in    a, ($1F)            ; Kempston port
    and    %00011111          ; mask 5 bits: fire, up, down, left, right
    ld    d, a

    ; Keyboard: QAOP + space (merge with joystick)
    ; Q = up
    ld    bc, $FBFE           ; half-row Q-T
    in    a, (c)
    bit    0, a                ; Q key
    jr    nz, .not_q
    set    3, d                ; up
.not_q:
    ; O = left
    ld    b, $DF              ; half-row Y-P
    in    a, (c)
    bit    1, a                ; O key
    jr    nz, .not_o
    set    1, d                ; left
.not_o:
    ; P = right
    bit    0, a                ; P key (same half-row)
    jr    nz, .not_p
    set    0, d                ; right
.not_p:
    ; A = down
    ld    b, $FD              ; half-row A-G
    in    a, (c)
    bit    0, a                ; A key
    jr    nz, .not_a
    set    2, d                ; down
.not_a:
    ; Space = fire
    ld    b, $7F              ; half-row space-B
    in    a, (c)
    bit    0, a                ; space
    jr    nz, .not_fire
    set    4, d                ; fire
.not_fire:
```

```

ld    a, d
ld    (input_state), a
ret

```

Note that keyboard reads use IN A, (C) with the half-row address in B. Each key maps to a bit in the result byte. Merging keyboard and joystick into a single byte means the rest of the game logic does not care which input device the player uses.

The Scroll Engine

Scrolling is the most expensive operation in the frame. Chapter 17 covered the techniques in detail; here is how they integrate into the game.

Ironclaw uses the **combined scroll** method: character-granularity scrolling (8-pixel jumps) for the main viewport, with a pixel offset (0-7) within the current 8-pixel window for smooth visual movement. When the pixel offset reaches 8, the viewport shifts by one tile column and the offset resets to 0.

The viewport is 30 characters wide (240 pixels) and 20 characters tall (160 pixels), leaving room for a 2-character HUD at the top and bottom. The level tilemap is typically 256-512 tiles wide and 20 tiles tall.

When the viewport shifts by one tile column, the renderer must:

1. Copy 29 columns of the current screen one character left (or right)
2. Draw the newly exposed column of tiles from the tilemap

The column copy is an LDIR chain: 20 rows x 8 pixel lines x 29 bytes = 4,640 bytes at 21 T-states each = 97,440 T-states. That is more than an entire frame. This is why Chapter 17's shadow screen technique is essential.

```

; Shadow screen double-buffer scroll
; Frame N: display screen is bank 5, draw screen is bank 7
; 1. Draw the shifted background into bank 7
; 2. Flip: set bit 3 of $7FFD to display bank 7
; Frame N+1: display screen is bank 7, draw screen is bank 5
; 3. Draw the shifted background into bank 5
; 4. Flip: clear bit 3 of $7FFD to display bank 5

```

```

flip_screen:
    ld    a, (last_bank_state)
    xor   %00001000      ; toggle screen bit (bit 3)
    ld    (last_bank_state), a
    ld    bc, $7FFD
    out   (c), a
    ret

```

But even with double-buffering, the full column copy is expensive. Ironclaw optimises this by spreading the work: during smooth sub-tile scrolling (pixel offset 1-7), there is no column copy – only the offset changes. The expensive column copy happens only on tile boundaries, roughly every 4-8 frames depending on the player's speed. Between those spikes, scroll rendering is nearly free.

When a tile boundary is crossed, the column copy can be spread across two frames using the double buffer: Frame N draws the top half of the shifted screen into

the back buffer, Frame N+1 draws the bottom half and flips. The player sees a seamless scroll because the flip only happens when the back buffer is complete.

21.5 Sprite Integration

Ironclaw uses OR+AND masked sprites (Chapter 16, method 2) for all game entities. This is the standard technique: for each sprite pixel, AND with a mask byte to clear the background, then OR with the sprite data to set the pixels.

Each 16x16 sprite has four pre-shifted copies (Chapter 16, method 3), one for each 2-pixel horizontal alignment. This reduces per-pixel shifting from a runtime operation to a table lookup. The cost: each sprite frame requires 4 variants x 16 lines x 4 bytes/line (2 bytes data + 2 bytes mask, widened to handle shift overflow) = 256 bytes. But the rendering speed drops from ~1,500 T-states to ~1,000 T-states per sprite, and with 8-10 sprites on screen, that savings adds up.

Pre-shifted sprite data lives in bank 3. During the sprite rendering phase, the renderer pages in bank 3, iterates through all active entities, and draws each one:

```
; Draw all active entities
; Assumes bank 3 (sprite graphics) is paged in at $C000
render_entities:
    ld    ix, entity_array
    ld    b, MAX_ENTITIES

.loop:
    push bc

    ; Check if entity is active
    ld    a, (ix + ENT_FLAGS)
    bit   FLAG_ACTIVE, a
    jr    z, .skip

    ; Calculate screen position from world position and viewport
    ld    l, (ix + ENT_X)
    ld    h, (ix + ENT_X + 1)
    ld    de, (viewport_x)
    or    a                      ; clear carry
    sbc   hl, de                 ; screen_x = world_x - viewport_x
    ; Check if on screen (0-239)
    bit   7, h
    jr    nz, .skip              ; off-screen left (negative)
    ld    a, h
    or    a
    jr    nz, .skip              ; off-screen right (> 255)
    ld    a, l
    cp    240
    jr    nc, .skip              ; off-screen right (240-255)

    ; Store screen X for sprite routine
    ld    (sprite_screen_x), a
```

```

; Y position (already in screen coordinates for simplicity)
ld  a, (ix + ENT_Y)
ld  (sprite_screen_y), a

; Look up sprite graphic address from type + frame + shift
call get_sprite_address ; returns HL = address in bank 3

; Draw masked sprite at (sprite_screen_x, sprite_screen_y)
call draw_sprite_masked

.skip:
pop  bc
ld  de, ENT_SIZE
add ix, de          ; next entity
djnz .loop
ret

```

Background Restore (Dirty Rectangles)

Before drawing sprites at their new positions, you must erase them from their old positions. Ironclaw uses the dirty rectangle method from Chapter 16: before drawing a sprite, save the background beneath it to a buffer. Before the next frame's sprite rendering pass, restore those saved backgrounds.

```

; Dirty rectangle entry: 4 bytes
;  byte 0: screen address low
;  byte 1: screen address high
;  byte 2: width in bytes
;  byte 3: height in pixel lines

; Save background before drawing sprite
save_background:
; HL = screen address, B = height, C = width
ld  de, bg_save_buffer
ld  (bg_save_ptr), de
; ... copy rectangle from screen to buffer ...
ret

; Restore all saved backgrounds (called before new sprite render pass)
restore_backgrounds:
ld  hl, dirty_rect_list
ld  b, (hl)          ; count of dirty rectangles
inc hl
or  a
ret  z                ; no sprites last frame

.loop:
push bc
; Read rectangle descriptor
ld  e, (hl)
inc hl

```

```

ld    d, (hl)          ; DE = screen address
inc   hl
ld    b, (hl)          ; B = height
inc   hl
ld    c, (hl)          ; C = width
inc   hl
push  hl

; Copy saved background back to screen
; ... copy from bg_save_buffer to screen ...

pop   hl
pop   bc
djnz  .loop
ret

```

The cost of dirty rectangles is proportional to the number and size of sprites. For 8 entities of 16x16 pixels (3 bytes wide after shift), saving and restoring costs roughly $8 \times 16 \times 3 \times 2$ (save + restore) $\times \sim 10$ T-states/byte = $\sim 7,680$ T-states. Not cheap, but predictable.

21.6 Collisions, Physics, and AI in Context

Chapters 18 and 19 covered these systems in isolation. In the integrated game, the key challenge is ordering: which system runs first, and what data does each need from the others?

Physics-Collision Loop

The physics update must interleave with collision detection. The pattern is:

2. Apply input: if (input_right) velocity_x += ACCEL
3. Horizontal move:
 - a. new_x = x + velocity_x
 - b. Check tile collisions at (new_x, y)
 - c. If blocked: push back to tile boundary, velocity_x = 0
 - d. Else: x = new_x
4. Vertical move:
 - a. new_y = y + velocity_y
 - b. Check tile collisions at (x, new_y)
 - c. If blocked: push back, velocity_y = 0, set on_ground flag
 - d. Else: y = new_y, clear on_ground flag
5. If (on_ground AND input_jump): velocity_y = -JUMP_FORCE

The horizontal and vertical moves are separate because collision response must handle each axis independently. If you move diagonally and hit a corner, you want to slide along the wall on one axis while stopping on the other. Checking both axes simultaneously leads to “sticking” bugs where the player gets trapped on corners.

All positions use 8.8 fixed-point format (Chapter 4): the high byte is the pixel coordinate, the low byte is the fractional part. Velocity values are also 8.8. This

gives sub-pixel movement precision without requiring any multiplication in the core physics loop – addition and shifting are sufficient.

```
; Apply gravity to entity at IX
; velocity_y is 16-bit signed, 8.8 fixed-point
apply_gravity:
    ld    l, (ix + ENT_VY)
    ld    h, (ix + ENT_VY + 1)
    ld    de, GRAVITY          ; e.g., $0040 = 0.25 pixels/frame/frame
    add   hl, de
    ; Clamp to terminal velocity
    ld    a, h
    cp    MAX_FALL_SPEED      ; e.g., 4 pixels/frame
    jr    c, .no_clamp
    ld    hl, MAX_FALL_SPEED * 256
.no_clamp:
    ld    (ix + ENT_VY), l
    ld    (ix + ENT_VY + 1), h
    ret
```

Tile Collision

The tile collision check converts a pixel coordinate to a tile index, then looks up the tile type in the level's collision map:

```
; Check tile at pixel position (B=x, C=y)
; Returns: A = tile type (0=empty, 1=solid, 2=hazard, 3=platform)
check_tile:
    ; Convert pixel X to tile column: x / 8
    ld    a, b
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    ; A = column (0-31)
    ld    l, a

    ; Convert pixel Y to tile row: y / 8
    ld    a, c
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    srl   a
    ; A = row (0-23)

    ; Tile index = row * level_width + column
    ld    h, 0
    ld    d, h
    ld    e, a
    ; Multiply row by level_width (e.g., 256 = trivial: just use E as high byte)
    ; For level_width = 256: address = level_map + row * 256 + column
    ld    d, e
    ; D = row = high byte of offset
    ld    e, l
    ; E = column = low byte of offset
    ld    hl, level_collision_map
    add   hl, de
    ld    a, (hl)
    ; A = tile type
```

```
ret
```

For Ironclaw, level widths are set to 256 tiles. This is not a coincidence – it makes the row multiplication trivial (the row number becomes the high byte of the offset). A 256-tile-wide level at 8 pixels per tile is 2,048 pixels, roughly 8.5 screens wide. For longer levels, you can use 512-tile width (multiply row by 2 via `SLA E : RL D`), though this costs a few extra T-states per lookup.

Enemy AI

Each enemy type has a finite state machine (Chapter 19). The state is stored in the entity structure:

```
; Entity structure (16 bytes per entity)
ENT_X      equ 0      ; 16-bit, 8.8 fixed-point
ENT_Y      equ 2      ; 16-bit, 8.8 fixed-point
ENT_VX      equ 4      ; 16-bit, 8.8 fixed-point
ENT_VY      equ 6      ; 16-bit, 8.8 fixed-point
ENT_TYPE    equ 8      ; entity type (player, walker, flyer, shooter, boss)
ENT_STATE   equ 9      ; FSM state (idle, patrol, chase, attack, retreat, dying)
ENT_ANIM     equ 10     ; animation frame counter
ENT_HEALTH  equ 11     ; hit points
ENT_FLAGS   equ 12     ; bit flags: active, on_ground, facing_left, invuln, ...
ENT_TIMER    equ 13     ; general-purpose timer (attack cooldown, etc.)
ENT_AUX1     equ 14     ; type-specific data (patrol point, projectile type, etc.)
ENT_AUX2     equ 15     ; type-specific data
ENT_SIZE     equ 16
```

```
MAX_ENTITIES equ 16    ; player + 8 enemies + 7 projectiles
```

Ironclaw's four enemy types:

1. **Walker** – Patrols between two points. When the player is within 64 pixels horizontally, switches to Chase state (walks toward player). Switches to Attack (contact damage) on collision. Returns to Patrol when player moves away or the enemy reaches a ledge.
2. **Flyer** – Sine-wave vertical movement (using the sine table from Chapter 4). Ignores tile collisions. Chases the player horizontally when within range. Drops projectiles at intervals.
3. **Shooter** – Stationary. Fires a horizontal projectile every N frames when the player is within line-of-sight (same tile row, no solid tiles between). The projectile is a separate entity allocated from the entity pool.
4. **Boss** – Multi-phase FSM. Phase 1: patrol platform, fire spread shots. Phase 2 (below 50% health): faster movement, targeted shots, summons walkers. Phase 3 (below 25% health): enrage, continuous fire, screen shake.

The key optimisation from Chapter 19: AI does not run every frame. Enemy AI updates are distributed across frames using a simple round-robin:

```
; Update AI for subset of enemies each frame
; ai_frame_counter cycles 0, 1, 2, 0, 1, 2, ...
update_enemy_ai:
    ld    a, (ai_frame_counter)
```

```

    inc a
    cp 3
    jr c, .no_wrap
    xor a
.no_wrap:
    ld (ai_frame_counter), a

    ; Only update enemies where (entity_index % 3) == ai_frame_counter
    ld ix, entity_array + ENT_SIZE ; skip player (index 0)
    ld b, MAX_ENTITIES - 1
    ld c, 0 ; entity index counter

.loop:
    push bc
    ld a, (ix + ENT_FLAGS)
    bit FLAG_ACTIVE, a
    jr z, .next

    ; Check if this entity's turn
    ld a, c
    ld e, 3
    call mod_a_e ; A = entity_index % 3
    ld b, a
    ld a, (ai_frame_counter)
    cp b
    jr nz, .next

    ; Run AI for this entity
    call run_entity_ai ; dispatch based on ENT_TYPE and ENT_STATE

.next:
    pop bc
    inc c
    ld de, ENT_SIZE
    add ix, de
    djnz .loop
    ret

```

This means each enemy's AI runs once every 3 frames. At 50 fps, that is still ~17 AI updates per second per enemy – more than enough for responsive behaviour. The saving is significant: if AI costs ~500 T-states per enemy, running all 8 enemies every frame costs 4,000 T-states. Running 2-3 enemies per frame costs 1,000-1,500 T-states. Physics and collision detection still run every frame for smooth movement.

21.7 Sound Integration

Music

The PT3 player runs inside the IM2 interrupt handler, as shown in section 21.2. The player occupies approximately 1.5-2KB of code and executes once per frame,

taking ~2,500-3,500 T-states depending on the complexity of the current pattern row.

Each level has its own music track. When transitioning between levels, the game:

1. Fades out the current track (ramp AY volumes to 0 over 25 frames)
2. Pages in the appropriate music bank (bank 4 or 6)
3. Initialises the PT3 player with the new song's start address
4. Fades in

The PT3 data format is compact - a typical 2-3 minute game music loop compresses to 2-4KB with Pletter, which is why two music banks (4 and 6) can hold all six tracks (title, five levels, boss, game over).

Sound Effects

Sound effects use the priority-based channel stealing system from Chapter 11. When a sound effect triggers (player jumps, enemy dies, projectile fires), the SFX engine temporarily hijacks one AY channel, overriding whatever the music was doing on that channel. When the effect finishes, the channel returns to music control.

```
; SFX priority levels
SFX_JUMP      equ 1      ; low priority
SFX_PICKUP    equ 2
SFX_SHOOT     equ 3
SFX_HIT       equ 4
SFX_EXPLODE   equ 5      ; high priority
SFX_BOSS_DIE  equ 6      ; highest priority

; Trigger a sound effect
; A = SFX id
play_sfx:
    ; Check priority -- only play if higher than current SFX
    ld  hl, current_sfx_priority
    cp  (hl)
    ret c                ; current SFX has higher priority, ignore

    ; Set up SFX playback
    ld  (hl), a           ; update priority
    ; Look up SFX descriptor table
    add a, a              ; x2 for word index
    ld  l, a
    ld  h, 0
    ld  de, sfx_table
    add hl, de
    ld  a, (hl)
    inc hl
    ld  h, (hl)
    ld  l, a              ; HL = SFX descriptor address

    ; SFX descriptor: duration (byte), channel (byte),
    ;                      then per-frame: freq_lo, freq_hi, volume, noise
```

```

ld    a, (hl)
ld    (sfx_frames_left), a
inc   hl
ld    a, (hl)
ld    (sfx_channel), a
inc   hl
ld    (sfx_data_ptr), hl
ret

```

The SFX update runs inside the interrupt handler, after the PT3 player. If an SFX is active, it overwrites the AY register values that the PT3 player just set for the stolen channel. This means the music continues to play correctly on the other two channels, and the stolen channel produces the sound effect.

SFX definitions are procedural tables rather than sampled audio. Each entry is a sequence of per-frame register values:

```

; SFX: player jump -- ascending frequency sweep on channel C
sfx_jump_data:
    db    8                ; duration: 8 frames
    db    2                ; channel C (0=A, 1=B, 2=C)
    ; Per-frame: freq_lo, freq_hi, volume
    db    $80, $01, 15     ; frame 1: low pitch, full volume
    db    $60, $01, 14     ; frame 2: slightly higher
    db    $40, $01, 13
    db    $20, $01, 12
    db    $00, $01, 10
    db    $E0, $00, 8
    db    $C0, $00, 5
    db    $A0, $00, 2      ; frame 8: high pitch, fading out

```

This approach uses negligible memory (8-20 bytes per effect) and negligible CPU time (a few dozen T-states per frame to write 3-4 AY register values).

21.8 Loading: Tape and DivMMC

A ZX Spectrum game must load somehow. In the 1980s, that meant tape. Today, most users have a DivMMC (or similar) SD card interface running esxDOS. Ironclaw supports both.

The .tap File and BASIC Loader

The .tap file format is a sequence of data blocks, each preceded by a 2-byte length and a flag byte. A BASIC loader program (itself a block in the .tap) uses LOAD "" CODE commands to load each block to the correct address.

Ironclaw's .tap structure:

```

Block 1:  Loading screen (6912 bytes -> $4000)
Block 2:  Main code block (bank 2 content -> $8000)
Block 3:  Bank 0 data (level data + tiles, compressed)
Block 4:  Bank 1 data (more level data)

```

Block 5: Bank 3 data (sprite graphics)
 Block 6: Bank 4 data (music tracks 1-3)
 Block 7: Bank 6 data (music tracks 4-6, SFX)

The BASIC loader:

```
20 LOAD "" SCREEN$
30 LOAD "" CODE
40 BORDER 0: PAPER 0: INK 0: CLS
50 RANDOMIZE USR 32768
```

Line 10 sets RAMTOP below \$8000, protecting our code from BASIC's stack. Line 20 loads the loading screen directly into screen memory (the Spectrum's LOAD "" SCREEN\$ does this automatically). Line 30 loads the main code block. Line 40 clears the screen. Line 50 jumps to our code at \$8000.

But this only loads the main code block. The banked data (blocks 3-7) must be loaded by our own Z80 code, which pages in each bank and uses the ROM's tape loading routine:

```
; Load bank data from tape
; Called after main code is running
load_bank_data:
    ; Bank 0
    ld    a, 0
    call switch_bank
    ld    ix, $C000        ; load address
    ld    de, BANK0_SIZE   ; data length
    call load_tape_block

    ; Bank 1
    ld    a, 1
    call switch_bank
    ld    ix, $C000
    ld    de, BANK1_SIZE
    call load_tape_block

    ; ... repeat for banks 3, 4, 6 ...
    ret

; Load one tape block using ROM routine
; IX = address, DE = length
load_tape_block:
    ld    a, $FF           ; data block flag (not header)
    scf                    ; carry set = LOAD (not VERIFY)
    call $0556             ; ROM tape loading routine
    ret    nc              ; carry clear = load error
    ret
```

esxDOS Loading (DivMMC)

For users with a DivMMC or similar hardware, loading from an SD card is dramatically faster and more reliable. The esxDOS API provides file operations through RST \$08 followed by a function number:

```

; esxDOS function codes
F_OPEN      equ  $9A
F_CLOSE     equ  $9B
F_READ      equ  $9D
F_WRITE     equ  $9E
F_SEEK      equ  $9F
F_OPENDIR   equ  $A3
F_READDIR   equ  $A4

; esxDOS open modes
FA_READ     equ  $01
FA_WRITE    equ  $06
FA_CREATE   equ  $0E

; Open a file
; IX = pointer to null-terminated filename
; Returns: A = file handle (or carry set on error)
esx_open:
    ld  a, '*'           ; use default drive
    ld  b, FA_READ       ; open for reading
    rst $08
    db  F_OPEN
    ret

; Read bytes from file
; A = file handle, IX = destination address, BC = byte count
; Returns: BC = bytes actually read (or carry set on error)
esx_read:
    rst $08
    db  F_READ
    ret

; Close a file
; A = file handle
esx_close:
    rst $08
    db  F_CLOSE
    ret

Ironclaw detects whether esxDOS is present at startup by checking for the DivMMC
signature. If present, it loads all data from files on the SD card instead of tape:

; Load game data from esxDOS
; All bank data stored in separate files on SD card
load_from_esxdos:
    ; Load bank 0: levels + tiles
    ld  a, 0
    call switch_bank
    ld  ix, filename_bank0
    call esx_open
    ret  c                ; error -- fall back to tape
    push af               ; save file handle

```

```

    ld    ix, $C000
    ld    bc, BANK0_SIZE
    pop   af                      ; A = file handle (esxDOS preserves this)
    push  af
    call  esx_read
    pop   af
    call  esx_close

    ; Repeat for other banks...
    ; Bank 1
    ld    a, 1
    call  switch_bank
    ld    ix, filename_bank1
    call  esx_open
    ret   c
    ; ... (same pattern) ...

ret

filename_bank0:  db "IRONCLAW.B0", 0
filename_bank1:  db "IRONCLAW.B1", 0
filename_bank3:  db "IRONCLAW.B3", 0
filename_bank4:  db "IRONCLAW.B4", 0
filename_bank6:  db "IRONCLAW.B6", 0

The detection code:

; Detect esxDOS presence
; Sets carry if esxDOS is NOT available
detect_esxdos:
    ; Try to open a nonexistent file -- if RST $08 returns
    ; without crashing, esxDOS is present
    ld    a, '*'
    ld    b, FA_READ
    ld    ix, test_filename
    rst   $08
    db    F_OPEN
    jr    c, .not_present    ; carry set = open failed, but esxDOS handled it
    ; File actually opened -- close it and return success
    call  esx_close
    or    a                  ; clear carry
    ret

.not_present:
    ; esxDOS returned an error -- it IS present, just file not found
    ; Distinguish from "RST $08 went to ROM and crashed"
    ; by checking if we're still running. If we're here, esxDOS is present.
    or    a                  ; clear carry = esxDOS present
    ret

test_filename:  db "IRONCLAW.B0", 0

```

In practice, the safest detection method checks for the DivMMC's identification byte at a known trap address, or uses a known-safe RST \$08 call. The method

above works because if esxDOS is not present, RST \$08 jumps to the ROM's error handler, which for the 128K ROM at address \$0008 is a benign return that leaves carry clear. Production code should use a more robust check; the pattern above illustrates the concept.

21.9 Loading Screen, Menu, and High Scores

Loading Screen

The loading screen is the player's first impression. It loads as LOAD "" SCREEN\$ in the BASIC loader, meaning it appears while the remaining data blocks are loading from tape. On esxDOS, the loading is fast enough that you may want to display the screen for a minimum duration:

```
show_loading_screen:
    ; Loading screen is already in screen memory ($4000) from BASIC loader
    ; If loading from esxDOS, load it explicitly:
    ld  ix, filename_screen
    call esx_open
    ret  c
    push af
    ld  ix, $4000
    ld  bc, 6912
    pop af
    push af
    call esx_read
    pop af
    call esx_close

    ; Minimum display time: 100 frames (2 seconds)
    ld  b, 100
.wait:
    halt
    djnz .wait
    ret
```

```
filename_screen: db "IRONCLAW.SCR", 0
```

The loading screen itself is a standard Spectrum screen file: 6,144 bytes of pixel data followed by 768 bytes of attributes, totalling 6,912 bytes. Create it in any Spectrum-compatible art tool (ZX Paintbrush, SEViewer, or Multipaint) or convert a modern image with a dithering tool.

Title Screen and Menu

The title screen state displays the game logo and animated background, then transitions to the menu on any keypress:

```
state_title:
    ; Animate background (e.g., scrolling starfield, colour cycling)
    call title_animate
```

```

; Check for keypress
xor  a
in   a, ($FE)      ; read all keyboard half-rows at once
cpl                     ; invert (keys are active low)
and  $1F           ; mask 5 key bits
jr   z, .no_key
ld   a, STATE_MENU
ld   (current_state), a
.no_key:
jp   main_loop

```

The menu offers three options: Start Game, Options, High Scores. Navigation uses up/down keys, selection uses fire/enter. The menu is a simple state machine within the STATE_MENU handler:

```

menu_selection:
    db  0           ; 0=Start, 1=Options, 2=HiScores

state_menu:
    ; Draw menu (only redraw on selection change)
    call draw_menu

    ; Read input
    call read_input
    ld  a, (input_state)

    ; Up
    bit 3, a
    jr  z, .not_up
    ld  a, (menu_selection)
    or  a
    jr  z, .not_up
    dec a
    ld  (menu_selection), a
    call play_menu_beep
.not_up:

    ; Down
    ld  a, (input_state)
    bit 2, a
    jr  z, .not_down
    ld  a, (menu_selection)
    cp  2
    jr  z, .not_down
    inc a
    ld  (menu_selection), a
    call play_menu_beep
.not_down:

    ; Fire / Enter
    ld  a, (input_state)

```

```

    bit 4, a
    jr  z, .no_fire
    ld  a, (menu_selection)
    or  a
    jr  nz, .not_start
    ; Start game
    call init_game
    ld  a, STATE_GAMEPLAY
    ld  (current_state), a
    jp  main_loop
.not_start:
    cp  1
    jr  nz, .not_options
    ; Options (toggle sound, controls, etc.)
    call show_options
    jp  main_loop
.not_options:
    ; High scores
    ld  a, STATE_HIScore
    ld  (current_state), a
    jp  main_loop

.no_fire:
    jp  main_loop

```

High Scores

High scores are stored in a 10-entry table in bank 2's data area:

```

; High score entry: 3 bytes name + 3 bytes BCD score = 6 bytes
; 10 entries = 60 bytes
HIScore_COUNT equ 10
HIScore_SIZE  equ 6

```

```

hiscore_table:
    ; Pre-filled defaults
    db  "ACE"
    db  $00, $50, $00      ; 005000 BCD
    db  "BOB"
    db  $00, $40, $00      ; 004000
    db  "CAT"
    db  $00, $30, $00      ; 003000
    ; ... 7 more entries ...
    ds  7 * HIScore_SIZE, 0

```

Scores use BCD (Binary Coded Decimal) – two decimal digits per byte, three bytes per score, giving a maximum of 999,999 points. BCD is preferable to binary for display because converting a 24-bit binary number to decimal on a Z80 requires expensive division. With BCD, the DAA instruction handles carry between digits automatically, and printing requires only masking nibbles:

```

; Add points to score
; DE = points to add (BCD, 2 bytes, max 9999)

```

```

add_score:
    ld    hl, player_score
    ld    a, (hl)
    add   a, e
    daa                      ; adjust for BCD
    ld    (hl), a
    inc   hl
    ld    a, (hl)
    adc   a, d
    daa
    ld    (hl), a
    inc   hl
    ld    a, (hl)
    adc   a, 0
    daa
    ld    (hl), a
    ret

player_score:
    db    0, 0, 0            ; 3 bytes BCD, little-endian

```

When the game ends, the code scans the high score table to see if the player's score qualifies. If it does, the game enters STATE_HISCORE for name entry (three characters, selected with up/down/fire).

On esxDOS systems, the high score table can be saved to the SD card. On tape systems, high scores persist only for the current session.

21.10 Level Loading and Decompression

When the player starts a level or completes one, the game must:

1. Page in the bank containing the level data (bank 0 for levels 1-2, bank 1 for levels 3-5)
2. Decompress the tilemap into bank 7 (the shadow screen bank, repurposed as a data buffer during level transitions)
3. Decompress the tile graphics into a buffer in bank 2 or bank 0
4. Initialise the entity array from the level's spawn table
5. Reset the viewport to the level's start position
6. Reset the scroll engine state

```

; Load and initialise level
; A = level number (0-4)
load_level:
    push af

    ; Determine which bank holds this level
    cp    2
    jr    nc, .bank1
    ; Levels 0-1: bank 0
    ld    a, 0

```

```

    call switch_bank
    pop af
    push af
    ; Look up compressed data address
    add a, a
    ld l, a
    ld h, 0
    ld de, level_ptrs_bank0
    add hl, de
    jr .decompress
.bank1:
    ; Levels 2-4: bank 1
    ld a, 1
    call switch_bank
    pop af
    push af
    sub 2                ; offset within bank 1
    add a, a
    ld l, a
    ld h, 0
    ld de, level_ptrs_bank1
    add hl, de

.decompress:
    ; HL points to 2-byte address of compressed level data in current bank
    ld a, (hl)
    inc hl
    ld h, (hl)
    ld l, a                ; HL = compressed data source (in $C000-$FFFF)

    ; Decompress tilemap into bank 7
    ; First, save current bank and switch to bank 7
    ; BUT: bank 7 is at $4000 (shadow screen), not $C000
    ; We decompress to $C000 in a temporary bank, then copy
    ; OR: decompress directly into shadow screen at $4000

    ; Simpler approach: decompress into a buffer at $8000+ area
    ; (we have ~2KB free above our code in bank 2)
    ; For large levels, use bank 7 at $4000:
    ; Enable shadow screen banking, then write to $4000-$7FFF

    ld de, level_buffer    ; destination in bank 2 work area
    call zx0_decompress     ; ZX0 decompressor: HL=src, DE=dest

    ; Initialise entities from spawn table
    pop af                ; A = level number
    call init_level_entities

    ; Set viewport to level start
    ld hl, 0
    ld (viewport_x), hl

```

```
ld    hl, 0
ld    (viewport_y), hl

; Reset scroll state
xor    a
ld    (scroll_pixel_offset), a
ld    (scroll_dirty), a

ret
```

The choice of compressor matters here. Level data is loaded once per level (during a transition screen), so decompression speed is not critical – we can afford Exomizer’s ~250 T-states per byte for the best compression ratio. But tile graphics may need to be decompressed during gameplay (if tiles are banked), so Pletter’s ~69 T-states per byte is preferable.

As discussed in Chapter 14, the decompressor code itself occupies memory. ZX0 at ~70 bytes is ideal for projects where code space is tight. Ironclaw includes both a ZX0 decompressor (for level data at load time) and a Pletter decompressor (for streamed tile data during gameplay).

21.11 Profiling with DeZog

You have written all the code. It compiles. It runs. The player walks, enemies patrol, tiles scroll, music plays. But the frame budget overflows on level 3, where six enemies and three projectiles are on screen simultaneously. The border stripe shows a red band that extends past the visible screen area. You are dropping frames.

This is where DeZog earns its place in your toolchain.

What is DeZog?

DeZog is a VS Code extension that provides a full debugging environment for Z80 programs. It connects to emulators (ZEsarUX, CSpect, or its own internal simulator) and gives you:

- Breakpoints (address, conditional, logpoints)
- Step-through execution (step into, step over, step out)
- Register watches (all Z80 registers, updated in real time)
- Memory viewer (hex dump with live updates)
- Disassembly view
- Call stack
- **T-state counter** – the profiling tool we need

The Profiling Workflow

The border stripe tells you *that* you are over budget. DeZog tells you *where*.

Step 1: Isolate the slow frame. Set a conditional breakpoint at the start of the main loop that triggers only when a “frame overflow” flag is set. Add code to set this flag when the frame takes too long:

```

; At the end of the gameplay frame, before HALT:
; Check if we're still in the current frame
; (a simple approach: read the raster line via floating bus
; or use a frame counter incremented by IM2)
ld  a, (frame_overflow_flag)
or  a
jr  z, .ok
; Frame overflowed -- set debug breakpoint trigger
nop                ; <-- set DeZog breakpoint here
.ok:

```

Step 2: Measure subsystem costs. DeZog's T-state counter lets you measure the exact cost of any code section. Place the cursor at the start of `update_enemy_ai`, note the T-state counter, step over the call, and note the new counter value. The difference is the exact cost.

A systematic profiling pass measures each subsystem:

read_input	187	0.3%	
update_player_physics	743	1.0%	
update_player_state	412	0.6%	
update_enemy_ai	4,231	5.9%	<-- worst case
check_all_collisions	2,847	4.0%	
update_projectiles	523	0.7%	
scroll_viewport	12,456	17.4%	<-- expensive
render_exposed_tiles	11,892	16.6%	<-- expensive
restore_backgrounds	3,214	4.5%	
draw_sprites	10,156	14.2%	<-- expensive
update_hud	1,389	1.9%	
[IM2 music interrupt]	3,102	4.3%	
TOTAL	51,152	71.4%	
Slack	20,528	28.6%	

That is the average case. Now profile the worst case – level 3, six enemies on screen, player near the right edge triggering a scroll:

read_input	187	0.3%	
update_player_physics	743	1.0%	
update_player_state	412	0.6%	
update_enemy_ai	5,891	8.2%	<-- 6 enemies active
check_all_collisions	4,156	5.8%	<-- more pairs
update_projectiles	1,247	1.7%	<-- 3 projectiles
scroll_viewport	14,892	20.8%	<-- scroll + new column
render_exposed_tiles	14,456	20.2%	<-- full column render
restore_backgrounds	4,821	6.7%	
draw_sprites	13,892	19.4%	<-- 10 entities
update_hud	1,389	1.9%	
[IM2 music interrupt]	3,102	4.3%	
TOTAL	65,188	90.9%	
Slack	6,492	9.1%	

Only 9% slack in the worst case. That is dangerously thin. One more enemy or a complex music pattern could push you over.

Step 3: Find the bottleneck. The profiling table makes it obvious: scrolling + tile rendering consume 41% of the frame in the worst case. Sprite rendering takes 19%. Enemy AI takes 8%.

Step 4: Optimise the bottleneck. Options, roughly in order of impact:

1. **Spread the scroll cost.** Instead of rendering the full new column in one frame, render half on frame N and half on frame N+1 using the double buffer (discussed in section 21.4). This cuts the scroll spike from ~29,000 to ~15,000 T-states per frame.
2. **Use compiled sprites for the player.** The player sprite is always on screen and always rendered. Switching from OR+AND masked (Chapter 16, method 2) to compiled sprites (method 5) saves ~30% per sprite draw, but increases memory usage. For one frequently-drawn entity, the trade-off is worth it.
3. **Reduce sprite overdraw.** If two enemies overlap, you are drawing pixels that will be overwritten. Sort entities by Y coordinate (back-to-front) and skip drawing for fully occluded sprites. This helps in the worst case when entities cluster.
4. **Tighten the AI.** Profile `run_entity_ai` for each enemy type. The Shooter's line-of-sight check (scanning tile columns for occlusion) is often the most expensive AI operation. Cache the result: only re-check line-of-sight every 8 frames instead of every 3.

After optimisation, the worst case drops to ~58,000 T-states, leaving 19% slack. That is comfortable.

DeZog Configuration for Ironclaw

DeZog connects to an emulator that supports its debug protocol. For ZX Spectrum 128K development, ZEsarUX is the recommended choice:

```
{
  "version": "0.2.0",
  "configurations": [
    {
      "type": "dezog",
      "request": "launch",
      "name": "Ironclaw (ZEsarUX)",
      "remoteType": "zesarux",
      "zesarux": {
        "hostname": "localhost",
        "port": 10000
      },
      "sjasmpplus": [
        {
          "path": "src/main.a80"
        }
      ],
      "topOfStack": "0xBFFF",
    }
  ]
}
```



```

        "load": "build/ironclaw.sna",
        "startAutomatically": true,
        "history": {
            "reverseDebugInstructionCount": 100000
        }
    }
}

```

The history setting enables reverse debugging – you can step backwards to see how you arrived at a bug. This is invaluable for tracking down collision glitches where an entity teleported through a wall three frames ago.

21.12 The Data Pipeline in Detail

Getting data from artist tools into the game is often the most underestimated part of a project. Ironclaw’s pipeline converts four kinds of assets:

Tilesets (PNG to Spectrum pixel format)

The artist draws tiles in Aseprite, Photoshop, or any pixel art tool as an indexed-colour PNG. Tiles are arranged in a grid on a single sheet. The conversion script:

1. Reads the PNG, verifies it is 1-bit (black and white) or indexed with Spectrum-compatible colours
2. Slices into 8x8 or 16x16 tiles
3. Converts each tile to the Spectrum’s interleaved pixel format (where row 0 is at offset 0, row 1 at offset 256, not offset 1 – matching the screen layout)
4. Optionally deduplicates identical tiles
5. Writes a binary blob and a symbol table mapping tile IDs to offsets

For attributes, each tile also carries a colour byte (INK + PAPER + BRIGHT). The script extracts this from the PNG’s palette and writes a parallel attribute table.

Sprite Sheets (PNG to pre-shifted sprite data)

Sprites follow a similar pipeline, but with an additional step: pre-shifting. The conversion script:

1. Reads the PNG sprite sheet
2. Slices into individual frames
3. Generates a mask for each frame (any non-background pixel produces a 0 in the mask, background produces 1)
4. For each frame, generates 4 horizontally shifted variants (0, 2, 4, 6 pixels offset)
5. Each shifted variant is widened by one byte (a 2-byte-wide sprite becomes 3 bytes wide to hold shifted overflow)
6. Writes interleaved data+mask bytes for efficient rendering

Level Maps (Tiled JSON to binary tilemap)

Levels are designed in Tiled, a free cross-platform tilemap editor. The designer places tiles visually, adds object layers for entity spawn points and triggers, and exports as JSON or TMX.

The conversion script:


1. Reads the Tiled export
2. Extracts the tile layer as a flat array of tile indices
3. Extracts the object layer for spawn points (enemy positions, player start, item locations)
4. Generates a collision map: for each tile, looks up whether it is solid, a platform, a hazard, or empty (based on a tile properties file)
5. Writes the tilemap, collision map, and spawn table as separate binary files

Music (Vortex Tracker II to PT3)

Music is composed in Vortex Tracker II, which exports directly to .pt3 format. The PT3 file is embedded into the bank data with INCBIN. The PT3 player code (widely available as open-source Z80 assembly, typically 1.5-2KB) is included in the music bank alongside the song data.

Putting It Together

The complete conversion pipeline for a level:



Every step is automated by the Makefile. The artist changes a tile, types make, and sees the result in the emulator.

21.13 Release Format: Building the .tap

The final deliverable is a .tap file. sjasmpplus can generate .tap output directly using its SAVETAP directive:

```

; main.a80 -- top-level assembly file

; Define the BASIC loader
DEVICE ZXSPECTRUM128

```

```

; Page in bank 2 at $8000
ORG $8000

; Include all game code
INCLUDE "defs.a80"
INCLUDE "banks.a80"
INCLUDE "render.a80"
INCLUDE "sprites.a80"
INCLUDE "entities.a80"
INCLUDE "physics.a80"
INCLUDE "collisions.a80"
INCLUDE "ai.a80"
INCLUDE "player.a80"
INCLUDE "hud.a80"
INCLUDE "menu.a80"
INCLUDE "loader.a80"
INCLUDE "music_driver.a80"
INCLUDE "sfx.a80"
INCLUDE "esxdos.a80"

; Entry point
entry:
    di
    ld    sp, $BFFF
    call init_system
    call detect_esxdos
    jr    c, .tape_load
    call load_from_esxdos
    jr    .loaded
.tape_load:
    call load_bank_data
.loaded:
    call init_interrupts
    ei
    jp    main_loop

; Bank data sections
; Each SLOT/PAGE directive places data into the correct bank
SLOT 3                ; use $C000 slot
PAGE 0                ; bank 0
ORG $C000
INCLUDE "data/bank0_levels.a80" ; INCBIN compressed level data

PAGE 1                ; bank 1
ORG $C000
INCLUDE "data/bank1_levels.a80"

PAGE 3                ; bank 3
ORG $C000
INCLUDE "data/bank3_sprites.a80"

```

```

PAGE 4                ; bank 4
ORG $C000
INCLUDE "data/bank4_music.a80"

PAGE 6                ; bank 6
ORG $C000
INCLUDE "data/bank6_sfx.a80"

; Save as .tap with BASIC loader
SAVETAP "build/ironclaw.tap", BASIC, "Ironclaw", 10, 2
SAVETAP "build/ironclaw.tap", CODE, "Screen", $4000, 6912, $4000
SAVETAP "build/ironclaw.tap", CODE, "Code", $8000, $-$8000, $8000

; Save bank snapshots (for .sna or manual loading)
SAVESNA "build/ironclaw.sna", entry

```

The exact SAVETAP syntax varies by sjasmplus version. For 128K games with banked data, the cleanest approach is to generate a .sna snapshot (which captures all bank states) for emulator testing, and a .tap with a BASIC loader plus machine code blocks for distribution.

Testing the Release

Before publishing, test on at least three emulators:

1. **Fuse** - the reference Spectrum emulator, accurate timing for original hardware
2. **Unreal Speccy** - Pentagon timing, the demoscene standard, good debugger
3. **ZEsarUX** - supports 128K banking, esxDOS emulation, DeZog integration

And if possible, test on real hardware with a DivMMC. Emulators occasionally differ in timing edge cases, and a game that runs perfectly in Fuse may drop frames on a real Spectrum due to contended memory effects that the emulator models slightly differently.

21.14 Final Polish

The difference between a working game and a finished game is polish. Here is a checklist of small touches that matter:

Screen transitions. Do not jump between screens instantly. A simple fade-to-black (write decreasing brightness to all attributes over 8 frames) or a wipe (clear columns from left to right over 16 frames) gives the game a professional feel. Cost: negligible - transitions happen between gameplay frames.

Death animation. When the player dies, freeze gameplay for 15 frames, flash the player sprite by toggling its INK between frames, play the death SFX, then respawn. Do not just teleport the player back to the checkpoint.

Screen shake. When the boss hits the ground or an explosion goes off, shift the viewport by 1-2 pixels for 4-6 frames. On the Spectrum, you can fake this by adjusting the scroll offset without actually moving any tiles. It is almost free and adds enormous impact.

Attract mode. After 30 seconds on the title screen with no input, start a demo playback – record the player’s input during a test run and replay it. This is how arcade games hook passersby, and it works for Spectrum games too.

Colour cycling. Animate menu text or logo colours by cycling attributes through a palette table. A 4-byte attribute cycle costs essentially zero CPU and makes static screens feel alive.

Input debounce. Ignore key presses that last fewer than 2 frames. Without debounce, the menu cursor will skip past options because the key was held for multiple frames. A simple frame counter per key fixes this:

```
; Debounced fire button
fire_held_frames:
    db    0

check_fire:
    ld    a, (input_state)
    bit   4, a
    jr    z, .released
    ; Fire is held
    ld    a, (fire_held_frames)
    inc   a
    ld    (fire_held_frames), a
    cp    1                ; only trigger on first frame of press
    ret                    ; Z flag set if this is the first frame
.released:
    xor   a
    ld    (fire_held_frames), a
    ret                    ; Z flag clear (no fire)
```

Summary

- **Project structure matters.** Separate source files by subsystem, data files by type. Use a Makefile to automate the full pipeline from PNG/TMX to .tap.
- **Memory map carefully.** Code in bank 2 (fixed at \$8000), level data in banks 0-1, sprite graphics in bank 3, music in banks 4 and 6, shadow screen in bank 7. Keep a shadow copy of port \$7FFD – it is write-only.
- **The interrupt handler owns the music.** The IM2 handler pages in the music bank, runs the PT3 player, updates SFX, and restores the previous bank. Keep it lean – ~3,000 T-states maximum.
- **The gameplay frame budget on Pentagon is 71,680 T-states.** A typical frame with scrolling, 8 sprites, and AI costs ~50,000 T-states average, ~65,000 worst case. Profile and optimise the worst case, not the average.
- **Scrolling is the most expensive single operation.** Use the combined scroll method (character-level LDIR + pixel offset) with shadow screen double-buffering. Spread the column copy across two frames when possible.

- **Run enemy AI every 2nd or 3rd frame.** Physics and collision detection run every frame; AI decisions can be amortised. This saves 2,000-3,000 T-states per frame in the worst case.
- **Use esxDOS for modern hardware.** The RST \$08 / F_OPEN / F_READ / F_CLOSE API is simple and fast. Detect DivMMC at startup and fall back to tape loading if absent.
- **Profile with DeZog.** The border stripe tells you that you are over budget. DeZog tells you where. Measure each subsystem, find the bottleneck, optimise it, measure again.
- **Choose the right compressor for each job.** Exomizer or ZX0 for one-time level loading (best ratio). Pletter for tile streaming during gameplay (fast decompression). See Chapter 14 for the full tradeoff analysis.
- **Polish is not optional.** Screen transitions, death animations, screen shake, input debounce, and attract mode are the difference between a tech demo and a game.
- **Test on multiple emulators and real hardware.** Fuse, Unreal Speccy, and ZEsarUX each model timing differently. DivMMC behaviour on real hardware can differ from emulated esxDOS.

Sources: World of Spectrum (ZX Spectrum 128K memory map and port \$7FFD documentation); Introspec “Data Compression for Modern Z80 Coding” (Hype 2017); esxDOS API documentation (DivIDE/DivMMC wiki); DeZog VS Code Extension documentation (GitHub: maziac/DeZog); sjasmplus documentation (SAVETAP, DEVICE, SLOT, PAGE directives); Vortex Tracker II PT3 format specification; Chapters 11, 14, 15, 16, 17, 18, 19 of this book.

Chapter 22: Porting — Agon Light 2

“The same instruction set, a completely different machine.”

You have built the game. Five levels, four enemy types, a boss fight, AY music with sound effects, a loading screen, and a menu system — all running on a ZX Spectrum 128K at 3.5 MHz, in 128 kilobytes of banked memory, rendering through a ULA that has not changed since 1982. Every byte is accounted for. Every cycle is earned.

Now you are going to port it to a machine that has the same CPU instruction set, five times the clock speed, four times the memory, hardware sprites, hardware tilemap scrolling, an SD card for loading, and a 24-bit flat address space with no banking.

This should be easy.

It is not easy. It is *different* in ways that will surprise you, and the surprises teach you things about both machines that you would not learn any other way.

Same ISA, Different World

The Agon Light 2 runs a Zilog eZ80 at 18.432 MHz with 512 KB of flat RAM. The eZ80 is a direct descendant of the Z80 — it executes the entire Z80 instruction set, uses the same register names, the same flags, the same mnemonics. If you write `LD A, (HL)` on a Spectrum and `LD A, (HL)` on an Agon, the opcode is identical. The behaviour is identical. A Z80 programmer can sit down at an Agon and start writing code immediately.

But the Agon is not a fast Spectrum. It is a fundamentally different architecture wearing a familiar face. The differences fall into three categories:

What the eZ80 adds. 24-bit registers, 24-bit addressing, a 16 MB address space (of which 512 KB is populated), new instructions for 24-bit arithmetic, and a mode system (ADL vs Z80-compatible) that controls register width and address generation.

What the VDP replaces. The Spectrum’s ULA — the chip that reads video memory and paints the screen — is replaced by a completely separate processor. The Agon’s VDP is an ESP32 microcontroller running the FabGL graphics library. It handles display output, sprites, tilemaps, and audio. The eZ80 CPU communicates with

the VDP over a high-speed serial link, sending command sequences. There is no shared video memory. You do not write pixels to an address; you send commands to a coprocessor.

What disappears. Banked memory, contended memory, the attribute grid, the interleaved screen layout, the 6,912-byte framebuffer, the border as a timing tool, direct framebuffer access, cycle-level raster synchronization. All gone.

To port our Spectrum game, we need to understand what transfers directly, what needs rewriting, and what needs rethinking from scratch.

The Architecture at a Glance

Before diving into code, let us lay out the two machines side by side.

Feature	ZX Spectrum 128K	Agon Light 2
CPU	Z80A	eZ80 (Z80-compatible + ADL extensions)
Clock	3.5 MHz (7 MHz on turbo clones)	18.432 MHz
RAM	128 KB (8 x 16 KB banks, switched via port \$7FFD)	512 KB flat (24-bit addressing)
Address space	16-bit (64 KB visible at a time)	24-bit (16 MB, 512 KB populated)
Video	ULA: 256x192, 8x8 attribute colour, direct memory-mapped	VDP (ESP32 + FabGL): multiple modes, up to 640x480, sprites, tilemaps
Framebuffer access	Direct: write to \$4000-\$5AFF	Indirect: send VDP commands over serial
Sprites	Software only	Hardware: up to 256, managed by VDP
Scrolling	Software only (shift entire framebuffer)	Hardware tilemap scrolling via VDP
Sound	AY-3-8910 (3 channels + noise)	VDP audio (ESP32 synthesis, multiple waveforms, ADSR)
Storage	Tape / DivMMC (esxDOS)	SD card (FAT32, MOS file API)
OS	None (bare metal) / esxDOS for file I/O	MOS (Machine Operating System)
Frame budget	~71,680 T-states (Pentagon)	~368,640 T-states (at 50 Hz)

The frame budget ratio is roughly 5:1. But this understates the real difference, because many operations that consume CPU T-states on the Spectrum — sprite rendering, screen scrolling, framebuffer management — are offloaded to the VDP on the Agon. The eZ80 CPU spends its cycles on game logic, not pixel pushing.

graph TD

```

subgraph ZX["ZX Spectrum: Draw Sprite (~2400T CPU)"]
    direction TB

```



```

    ZA["Calculate screen address<br>from (x, y) coordinates"] --> ZB["Select
↪ pre-shifted variant<br>(x mod 8 → shift table)"]
    ZB --> ZC["For each pixel row:<br>AND mask with screen byte<br>OR sprite
↪ data<br>write back to framebuffer"]
    ZC --> ZD["Advance to next screen row<br>(DOWN_HL: handle interleave)"]
    ZD --> ZE{"16 rows<br>done?"}
    ZE -- No --> ZC
    ZE -- Yes --> ZF["Done – pixels in<br>video RAM at $4000"]
end

subgraph AGON["Agon Light 2: Draw Sprite (~50T CPU)"]
    direction TB
    AA["Send VDU 23,27,4,N<br>(select sprite N)"] --> AB["Send VDU
↪ 23,27,13,x,y<br>(set position)"]
    AB --> AC["Send VDU 23,27,15<br>(update display)"]
    AC --> AD["Done – VDP renders<br>sprite in hardware"]
end

style ZC fill:#fdd,stroke:#933
style AA fill:#dfd,stroke:#393
style AB fill:#dfd,stroke:#393
style AC fill:#dfd,stroke:#393

```

The architectural shift: On the Spectrum, the CPU is the rendering engine — every pixel is placed by Z80 instructions. On the Agon, the CPU is a *command sequencer* — it tells the VDP what to draw, and the ESP32 coprocessor handles the actual rendering. The CPU cost drops from ~2,400T to ~50T per sprite, but you now manage an asynchronous command pipeline with serial latency.

ADL Mode vs Z80-Compatible Mode

This is the single most important architectural concept for any Z80 programmer approaching the eZ80. Get it wrong and your code will crash in ways that are difficult to debug. Get it right and you unlock the full power of the chip.

The eZ80 has two operating modes:

Z80-compatible mode (Z80 mode). Registers are 16 bits wide. Addresses are 16 bits. The MBASE register provides the upper 8 bits of every address, effectively placing your 64 KB window somewhere in the 16 MB address space. Code behaves exactly like a standard Z80 — LD HL,\$4000 loads a 16-bit value, JP (HL) jumps to a 16-bit address (with MBASE prepended), PUSH HL pushes 2 bytes onto the stack.

ADL mode (Address Data Long). Registers are 24 bits wide. Addresses are 24 bits. LD HL,\$040000 loads a 24-bit value, JP (HL) jumps to a full 24-bit address, PUSH HL pushes 3 bytes onto the stack. This is the native mode of the eZ80.

MOS boots the Agon in ADL mode. Your application starts in ADL mode. Most Agon software runs entirely in ADL mode. But Z80-compatible mode exists, and understanding the interaction between the two is critical.

Why You Might Use Z80 Mode

If you are porting Z80 code from the Spectrum, you might think: “just switch to Z80 mode and run my existing code.” This works, up to a point. Your 16-bit address calculations, your DJNZ loops, your LDIR block copies — they all behave identically in Z80 mode. MBASE is set so that the 16-bit addresses map to the right region of the Agon’s memory.

The problem is *interacting with everything else*. MOS API calls expect ADL mode. VDP commands are sent through MOS routines that assume 24-bit stack frames. If you are in Z80 mode and call a MOS routine, the stack frame will be wrong — MOS pushes 3 bytes per return address, your Z80-mode code pushed 2. The result is stack corruption and a crash.

The Mode Switching Mechanism

The eZ80 provides special prefixes for switching modes within a single instruction:

Prefix	Effect
.SIS (suffix)	Execute the following instruction in Z80 mode (short registers, short addresses)
.LIS	Execute in: Long registers, Short addresses
.SIL	Execute in: Short registers, Long addresses
.LIL	Execute in ADL mode (long registers, long addresses)

And for calls and jumps:

Instruction	From mode	To mode
CALL.IS addr	ADL	Z80
CALL.IL addr	Z80	ADL

The .IS suffix means “Instruction Short” — the call instruction itself uses short (16-bit) conventions for the return address. .IL means “Instruction Long” — the call pushes a 24-bit return address.

Here is the practical pattern for calling MOS from Z80-mode code:

```
; In Z80-compatible mode, calling a MOS API function
; We need to switch to ADL mode for the call

; Method: use RST.LIL $08 (MOS API entry point)
; .LIL means "long instruction, long mode" ---
; pushes a 24-bit return address and enters ADL mode
RST.LIL $08      ; call MOS API in ADL mode
DB      mos_func ; MOS function number follows
; MOS returns to us in Z80 mode (matching our caller)
```

MOS provides RST \$08 as a unified API entry point. The `.LIL` suffix handles the mode transition cleanly. After the call, execution returns to your Z80-mode code with the correct stack state.

The Practical Rule

For porting, the cleanest approach is: **run your game logic in ADL mode and translate your Z80 code to use 24-bit conventions from the start.** Do not try to run in Z80 mode and switch back and forth for every MOS call. The mode-switching overhead and the risk of stack mismatches are not worth it.

This means your port will not be a byte-for-byte copy of the Spectrum code. It will be a *translation*. The algorithms are the same. The logic is the same. The register usage is mostly the same. But every address is 24 bits wide, every stack push is 3 bytes, and every immediate address load carries an extra byte.

The MBASE Trap

If you do use Z80 mode, MBASE determines the upper 8 bits of every memory address. On boot, MOS sets MBASE to \$00, meaning Z80-mode addresses 0000 — —FFFF map to physical addresses \$000000–\$00FFFF. If your code or data lives above \$00FFFF (above the first 64 KB), Z80-mode code cannot reach it without changing MBASE.

This is a trap for Spectrum porters who think “I have 512 KB, I will put my level data at \$080000.” In Z80 mode, that address does not exist. You must either use ADL mode to access it or set MBASE to \$08 (making addresses 0000 — —FFFF map to \$080000–\$08FFFF). But changing MBASE affects *all* memory accesses, including instruction fetches — so your code had better be in that region too, or you will jump into garbage.

The advice is simple: stay in ADL mode. Use the full 24-bit address space natively.

What Transfers Directly

Not everything changes. A surprising amount of your Spectrum game logic ports with minimal modification.

Game Logic and Entity System

The entity system from Chapter 18 — the structure arrays holding X, Y, type, state, animation frame, velocity, health, and flags — transfers almost verbatim. The main loop structure (HALT, input, update, render, repeat) is identical in concept, though the specific HALT and interrupt mechanism differs.

Here is the entity update loop on the Spectrum:

```
; Spectrum: Update all entities
; IX points to entity array, B = entity count
update_entities:
    ld    ix,entities
    ld    b,MAX_ENTITIES
```



```

cp    (iy+ENT_X)          ; entity1.x+width > entity2.x?
ret   c                   ; no overlap
; overlap on X axis confirmed

```

Fixed-Point Arithmetic

All fixed-point 8.8 calculations — gravity, velocity, friction, acceleration — port without changes. The shift-and-add patterns, the 16-bit additions, the right-shift friction:

```

; Apply gravity: velocity_y += gravity
; Works identically on both platforms
ld    a,(ix+ENT_VY_L0)
add   a,GRAVITY_L0
ld    (ix+ENT_VY_L0),a
ld    a,(ix+ENT_VY_HI)
adc   a,GRAVITY_HI
ld    (ix+ENT_VY_HI),a

```

Byte-level arithmetic does not care whether the registers are notionally 16 or 24 bits wide. The accumulator is always 8 bits. Carry propagation works the same way.

State Machine

The game state machine (title, menu, gameplay, pause, game over) uses a jump table indexed by state number. On the Spectrum:

```

; Spectrum: dispatch game state
ld    a,(game_state)
add   a,a                ; multiply by 2 (16-bit pointers)
ld    e,a
ld    d,0
ld    hl,state_table
add   hl,de
ld    a,(hl)
inc   hl
ld    h,(hl)
ld    l,a
jp    (hl)

```

```

state_table:
dw    state_title
dw    state_menu
dw    state_game
dw    state_pause
dw    state_gameover

```

On the Agon, the pointer table stores 24-bit addresses:

```

; Agon (ADL mode): dispatch game state
ld    a,(game_state)
ld    l,a
ld    h,0                ; HL = state index

```

```

    ld    e,l
    ld    d,h                ; DE = copy of state index
    add   hl,hl              ; HL = state * 2
    add   hl,de              ; HL = state * 3 (24-bit pointers)
    ld    de,state_table
    add   hl,de
    ld    hl,(hl)            ; load 24-bit pointer
    jp    (hl)

state_table:
    dl    state_title        ; DL = define long (24-bit)
    dl    state_menu
    dl    state_game
    dl    state_pause
    dl    state_gameover

```

The change: pointers are 3 bytes instead of 2, so the index multiplication changes from *2 to *3, and the table uses DL (define long) instead of DW (define word). The logic is otherwise identical.

What Needs Rewriting

Rendering: From Framebuffer to VDP Commands

This is the largest single change in the port. On the Spectrum, rendering means writing bytes to video memory addresses. The entire rendering pipeline — sprite drawing, screen clearing, tile painting, scrolling — is CPU code that manipulates memory at \$4000-\$5AFF.

On the Agon, rendering means sending VDP command sequences. The VDP understands a protocol based on VDU byte streams (the same VDU command system used by BBC BASIC, extended with Agon-specific commands). You send a sequence of bytes to the VDP through MOS, and the ESP32 processes them.

Sprites

On the Spectrum (from Chapter 16), drawing a 16x16 masked sprite costs roughly 1,200 T-states of CPU time — reading mask bytes, ANDing with the screen, ORing the sprite data, writing back. You do this for every sprite, every frame.

On the Agon, you upload the sprite bitmap *once*, and then move it by sending a position update:

```

; Agon: Create and position a hardware sprite
; Step 1: Upload sprite bitmap (done once at init)
;   VDU 23, 27, 4, spriteNum    ; select sprite
;   VDU 23, 27, 0, w, h        ; set dimensions
;   followed by pixel data

; Step 2: Move sprite (done every frame)
; VDU 23, 27, 4, spriteNum      ; select sprite

```

```

; VDU 23, 27, 13, x.lo, x.hi, y.lo, y.hi ; set position

move_sprite:
    ; Send VDU command to move sprite
    ld  a,23
    rst $10                ; MOS: output byte to VDP
    ld  a,27
    rst $10
    ld  a,4                ; command: select sprite
    rst $10
    ld  a,(sprite_num)
    rst $10

    ld  a,23
    rst $10
    ld  a,27
    rst $10
    ld  a,13                ; command: move sprite to
    rst $10

    ld  a,(sprite_x)        ; X low byte
    rst $10
    ld  a,(sprite_x+1)      ; X high byte
    rst $10
    ld  a,(sprite_y)        ; Y low byte
    rst $10
    ld  a,(sprite_y+1)      ; Y high byte
    rst $10

    ; VDU 23, 27, 15        ; show sprite (update display)
    ld  a,23
    rst $10
    ld  a,27
    rst $10
    ld  a,15
    rst $10
    ret

```

Each RST \$10 sends one byte to the VDP through MOS. The total CPU cost of moving a sprite is approximately 13 bytes sent x ~30 T-states per RST call = ~390 T-states. Compare that to the Spectrum's ~2,400 T-states for a full masked sprite draw. And the Agon version does not need background save/restore — the VDP composites sprites over the background automatically.

The trade-off: latency. The VDP processes commands asynchronously. Between sending the “move sprite” command and the sprite actually appearing at the new position, there is a serial transfer delay and a VDP processing delay. For smooth animation, you need to send all sprite updates early in the frame and trust that the VDP will process them before the next screen refresh.

Tilemap Scrolling

On the Spectrum, horizontal scrolling means shifting every byte of video memory left or right — a chain of RLC or RRC instructions across hundreds of bytes, consuming a substantial fraction of the frame budget (we calculated the cost in Chapter 17). Vertical scrolling requires copying scanlines with awareness of the interleaved memory layout.

On the Agon, the VDP supports hardware tilemaps:

```
; Agon: Set up a tilemap (done once)
; VDU 23, 27, 20, tileWidth, tileHeight
; VDU 23, 27, 21, mapWidth.lo, mapWidth.hi, mapHeight.lo, mapHeight.hi
```

```
; Scroll the tilemap (every frame)
; VDU 23, 27, 24, offsetX.lo, offsetX.hi, offsetY.lo, offsetY.hi
```

```
scroll_tilemap:
    ld    a,23
    rst   $10
    ld    a,27
    rst   $10
    ld    a,24                ; command: set scroll offset
    rst   $10

    ld    hl,(scroll_x)
    ld    a,l
    rst   $10                ; offsetX low
    ld    a,h
    rst   $10                ; offsetX high
    ld    hl,(scroll_y)
    ld    a,l
    rst   $10                ; offsetY low
    ld    a,h
    rst   $10                ; offsetY high
    ret
```

Eight bytes sent. Perhaps 240 T-states of CPU time. On the Spectrum, a full-screen horizontal pixel scroll costs tens of thousands of T-states. The Agon does it in hardware for nearly free.

But you must set up the tilemap first: upload tile definitions, define the map dimensions, populate the map with tile indices. This is a one-time cost at level load, not a per-frame cost. On the Spectrum, your tile data lives in banked RAM and is rendered into the framebuffer by your own code. On the Agon, tile data lives in VDP memory and is rendered by the ESP32. Your role changes from “graphics engine programmer” to “VDP command sequencer.”

Screen Layout

The entire nightmare of the Spectrum’s interleaved screen layout — the split addressing, the DOWN_HL routines, the careful calculations to convert (x, y) coordinates to memory addresses — vanishes. The Agon’s VDP works in screen coordinates. You say “draw at (100, 50)” and the VDP handles the rest.

This means the DOWN_HL routine from Chapter 2, the screen address lookup tables, the attribute address calculations — none of it ports. It is simply deleted. The equivalent operation on the Agon is “send a coordinate pair to the VDP.”

What Needs Rethinking

Some Spectrum patterns are so deeply embedded in the game architecture that you cannot just rewrite the rendering layer. The *design* needs to change.

Memory Architecture

On the Spectrum, you carefully planned which data goes in which bank:

- Banks 0-3: level data, tilesets, sprite graphics
- Banks 4-6: music patterns, sound effects, lookup tables
- Bank 7: shadow screen for double buffering

Every bank switch costs a port write and limits what code can see what data. The game architecture is shaped by the 16 KB window into a 128 KB space.

On the Agon, all 512 KB is visible simultaneously. There is no banking. There is no shadow screen trick (the VDP handles double buffering internally). You can have your entire game — all five levels, all tilesets, all sprites, all music — resident in memory at once. Level transitions do not require loading from tape or disk; you just point to a different region of RAM.

This simplifies development, but it also removes a constraint that forced good architecture. On the Spectrum, you were forced to think about data locality, about what needed to be co-resident, about load sequences. On the Agon, you can be sloppy. Do not be sloppy. The Agon has 512 KB, not infinity. A well-organized memory map is still a virtue.

Typical Agon memory layout for the ported game:

```
$040000 - $04FFFF   Game code (~64 KB)
$050000 - $06FFFF   Level data, all 5 levels (~128 KB)
$070000 - $07FFFF   Music and SFX data (~64 KB)
$080000 - $0FFFFFF   Free / working buffers
```

Everything is addressable with a single LD HL,\$070000 — no bank switching, no port writes.

Loading

On the Spectrum, loading from tape is a minutes-long process with a distinctive audio signature. Even with DivMMC and esxDOS, file access is a sequence of RST \$08 calls:

```
; Spectrum + esxDOS: load a file
ld    a,'*'           ; current drive
ld    ix,filename
ld    b,$01           ; read-only
rst   $08             ; esxDOS call
```

```

DB    $9A                ; F_OPEN
; A = file handle

ld    ix,buffer
ld    bc,size
rst   $08
DB    $9D                ; F_READ
; Data loaded

rst   $08
DB    $9B                ; F_CLOSE

```

On the Agon, MOS provides a file API that reads directly from the SD card:

```

; Agon: load a file using MOS API
ld    hl,filename        ; 24-bit pointer to filename string
ld    de,buffer          ; 24-bit pointer to destination
ld    bc,size            ; 24-bit max size
ld    a,mos_fopen        ; MOS file open function
rst   $08                ; MOS API call
; A = file handle

ld    a,mos_fread        ; MOS file read function
rst   $08
; Data loaded

ld    a,mos_fclose
rst   $08

```

The pattern is similar — open, read, close — but the Agon reads from FAT32 on an SD card, which is fast enough that you can load level data between scenes without a visible delay. No loading screens needed. No tape-loading routines. No block-loading optimization.

Sound

The Spectrum's AY-3-8910 is programmed by writing directly to hardware registers through I/O ports. Every note, every envelope change, every noise burst is a specific register value written at a specific time.

The Agon's audio is handled by the VDP. You send sound commands over the same serial link used for graphics:

```

; Agon: play a note
; VDU 23, 0, 197, channel, volume, freq.lo, freq.hi, duration.lo, duration.hi

```

```

play_note:
ld    a,23
rst   $10
xor   a                ; 0
rst   $10
ld    a,197            ; sound command
rst   $10
ld    a,(channel)

```

```

rst $10
ld  a,(volume)
rst $10
ld  hl,(frequency)
ld  a,l
rst $10
ld  a,h
rst $10
ld  hl,(duration)
ld  a,l
rst $10
ld  a,h
rst $10
ret

```

The Agon's sound system supports multiple waveforms (sine, square, triangle, sawtooth, noise) and per-channel ADSR envelopes — features the AY does not have. But the programming model is completely different. You cannot write raw register values; you send high-level note commands. Your AY music player — the IM2 interrupt handler that reads pattern data and updates 14 AY registers every frame — does not port at all. You need a new music driver that translates your pattern data into VDP sound commands.

One approach: write a thin abstraction layer that both platforms share.

```

; Abstract sound interface
; Spectrum implementation:
sound_play_note:
    ; A = channel, B = note, C = instrument
    ; ... look up AY register values, write to ports $FFFD/$BFFD
    ret

; Agon implementation:
sound_play_note:
    ; A = channel, B = note, C = instrument
    ; ... convert to VDP sound command, send via RST $10
    ret

```

Same call signature. Different internals. The game code calls `sound_play_note` without knowing which platform it is running on.

Input

The Spectrum reads keyboard state by probing port `$FE` with half-row addresses in the accumulator. Kempston joystick is read from port `$1F`. These are raw I/O port reads with specific bit patterns.

The Agon reads keyboard state through MOS:

```

; Spectrum: read keyboard
ld  a,$FD          ; half-row: Q W E R T
in  a,($FE)
bit 0,a            ; bit 0 = Q
; ...

```

```
; Agon: read keyboard via MOS
    ld    a,mos_getkey      ; MOS: get key state
    rst   $08
    ; A = key code, or 0 if no key pressed
```

The Spectrum gives you a bitmask of simultaneously pressed keys, ideal for game input (you can detect multiple keys at once). The Agon’s MOS keyboard API is event-based: it gives you the most recent key pressed. For game input with simultaneous key detection, you typically use the MOS keyboard map — a region of memory updated by MOS that reflects the state of all keys:

```
; Agon: read simultaneous keys from keyboard map
    ld    a,(mos_keymap+KEY_LEFT)  ; 1 if left arrow held, 0 if not
    or    a
    jr    z,.no_left
    ; move player left
.no_left:
    ld    a,(mos_keymap+KEY_RIGHT)
    or    a
    jr    z,.no_right
    ; move player right
.no_right:
```

This is functionally equivalent to the Spectrum’s bitmask approach, just organized differently. The port is straightforward: replace port reads with memory reads from the keyboard map.

eZ80 at 18 MHz: What Still Matters, What Does Not

The eZ80 is roughly five times faster than the Z80A in raw clock speed. But many eZ80 instructions also execute in fewer clock cycles than their Z80 equivalents — single-byte register instructions often complete in 1 cycle instead of 4. The effective speedup for typical code is somewhere between 5x and 20x depending on the instruction mix.

This changes the optimization calculus dramatically.

What Still Matters: Inner Loop Efficiency

Even at 18 MHz with 368,000 T-states per frame, the inner loop still matters for CPU-intensive operations. If you are doing collision checks against a tilemap, iterating over 200 entities, or processing AI state machines for dozens of enemies, the per-iteration cost of your hot loop adds up.

The core Z80 optimization techniques — keeping values in registers instead of memory, using `INC L` instead of `INC HL` where possible, avoiding `IX/IY` indexed instructions for hot paths (they carry a 2-cycle prefix penalty on eZ80, just as they carry a 4-T-state penalty on Z80) — still produce measurable improvements.

```
; Tight entity scan: same optimization principles on both platforms
; Prefer: register-to-register ops, direct addressing, DJNZ
```

; Avoid: IX-indexed loads in hot inner loops when possible

```
scan_entities_fast:
    ld    hl,entity_flags    ; pointer to flags array
    ld    b,MAX_ENTITIES
.loop:
    ld    a,(hl)
    bit   0,a
    call  nz,process_entity
    inc   hl                  ; next entity flag (assume 1-byte stride)
    djnz  .loop
    ret
```

This pattern — minimal memory access, tight register usage, DJNZ loop control — is optimal on both the Z80 at 3.5 MHz and the eZ80 at 18 MHz. Good code is good code.

What Becomes Irrelevant: Memory Conservation Tricks

On the Spectrum, memory pressure drives much of the engineering. Pre-shifted sprites store 4 or 8 copies of each sprite, consuming 4x-8x the memory, as a speed/memory trade-off. Compression is mandatory for fitting demo data into 128 KB. Lookup tables are carefully sized to balance precision against memory cost. Bank switching adds architectural complexity for the sole purpose of addressing more than 64 KB at a time.

On the Agon, with 512 KB of flat RAM and SD card storage for anything that does not need to be resident, these conservation techniques are unnecessary. You can store 8 pre-shifted copies of every sprite without worrying about memory. You can keep all lookup tables at full precision. You can keep all five levels in RAM simultaneously.

This does not mean you should waste memory. But it means that optimization decisions driven by memory scarcity — “should I use a 256-byte sine table or a 128-byte one?” — become irrelevant. Use 256. Use 1,024 if the precision helps.

What Becomes Irrelevant: Self-Modifying Code

On the Spectrum, self-modifying code (SMC) is a standard optimization technique. You write instructions that patch their own immediate operands to avoid memory lookups:

```
; Spectrum: self-modifying code for speed
    ld    a,0                ; operand patched at runtime
    ; ... (the $00 after LD A is overwritten with the real value)
```

On the eZ80, SMC still works (the eZ80 does not have an instruction cache that would invalidate), but the motivation is weaker. The extra cost of a memory load is smaller relative to the total frame budget. More importantly, MOS maps some memory regions as read-only, and certain Agon firmware versions may restrict execution of modified code depending on memory region. SMC is not prohibited on the Agon, but it is rarely necessary and can cause subtle problems.

What Becomes Irrelevant: Stack Tricks for Rendering

On the Spectrum, abusing the stack pointer for fast screen fills (DI, set SP to screen address, PUSH data repeatedly) is a classic trick because PUSH writes 2 bytes in 11 T-states — faster than any other write mechanism. On the Agon, you are not writing to a framebuffer at all. The VDP handles rendering. Stack tricks for display are meaningless.

The Comparison Table

Here is the side-by-side comparison of the same game on both platforms. These numbers are representative of the platformer built in Chapters 21–22.

Metric	ZX Spectrum 128K	Agon Light 2
Game code size	~12 KB	~14 KB
Rendering code	~5 KB (sprite engine, scroll, screen mgmt)	~2 KB (VDP command sequences)
Total code	~17 KB	~16 KB
Level data (all 5)	~40 KB (compressed, loaded per-level)	~60 KB (uncompressed, all resident)
Sprite/tile graphics	~20 KB (packed, 1bpp + masks)	~80 KB (8bpp RGBA, uploaded to VDP)
Music + SFX	~16 KB (PT3 + SFX tables)	~20 KB (converted format + waveform data)
Total data	~76 KB (fits in 128 KB with banking)	~160 KB (fits in 512 KB easily)
Compression needed?	Yes, mandatory	No, optional
Sprite draw cost	~2,400 T/sprite (software)	~400 T/sprite (VDP commands)
Scroll cost per frame	~15,000–30,000 T (software shift)	~240 T (VDP offset command)
Frame budget	~71,680 T	~368,640 T
Achievable fps	25–50 (depends on entity count)	60 (VDP-limited, not CPU-limited)
Dev complexity	High (memory banking, screen layout, rendering engine)	Medium (VDP protocol, MOS API, ADL mode)
Visual result	Monochrome or 2-colour-per-cell sprites, attribute colour, 256x192	Full colour sprites, smooth scrolling, 320x240 or higher

The code size difference is instructive. On the Spectrum, the *rendering engine* is a substantial fraction of the codebase. On the Agon, VDP commands replace most of that code. But the Spectrum game is smaller in total data because everything is compressed and packed tightly. The Agon version uses more memory for richer assets (8-bit-per-pixel sprites instead of 1-bit-per-pixel with masks).

The frame budget comparison is the most striking number. The Agon game is *CPU-idle most of the frame*. After processing game logic, sending sprite updates, and

handling input, the eZ80 has nothing to do until the next frame. On the Spectrum, you are fighting for every cycle to fit everything into the budget.

The Porting Process: Step by Step

Here is the practical sequence for porting the Chapter 21 game to the Agon.

Step 1: Set Up the Agon Project

Create a new project with the Agon toolchain. You will need:

- An eZ80 assembler (ez80asm or the Zilog ZDS tools)
- The MOS API header file (mos_api.inc) defining function numbers and constants
- A way to transfer the binary to the Agon (SD card or serial upload)

Your entry point is different from the Spectrum. Instead of `ORG $8000` with a raw `JP` to your start address, the Agon loads executables at `$040000` and MOS passes control to that address:

```
; Agon: application entry point
    .ASSUME ADL=1          ; we are in ADL mode
    ORG  $040000

    JP  main              ; standard entry

main:
    ; Your game starts here
    ; ... initialize VDP, load assets, enter game loop
```

Step 2: Replace the Rendering Layer

This is the bulk of the work. Delete the Spectrum rendering code and replace it with VDP command sequences:

1. **Upload sprite bitmaps to VDP.** Convert your 1bpp sprite data to the VDP's bitmap format (typically 8bpp RGBA). Send the pixel data using VDU 23,27 sprite definition commands.
2. **Upload tileset to VDP.** Convert your 8x8 tiles from Spectrum attribute format to the VDP's tile format. Define the tilemap dimensions and populate it with tile indices for the current level.
3. **Replace the sprite draw routine** with VDU 23,27 sprite position commands (as shown earlier).
4. **Replace the scroll routine** with VDU 23,27 tilemap scroll offset commands.
5. **Delete the screen address calculation code**, the `DOWN_HL` routine, the attribute address routines, and all direct framebuffer writes.

Step 3: Translate the Game Logic

Walk through the game logic code (entity update, collision detection, physics, AI, state machine) and adjust for ADL mode:

- Change all DW (define word) to DL (define long) for address tables.
- Change pointer arithmetic from 16-bit to 24-bit where addresses are involved.
- Verify that PUSH/POP pairs are balanced — each push is now 3 bytes, not 2.
- Check that LDIR and LDDR block copy counts are correct (BC is 24-bit in ADL mode; if your count fits in 16 bits, the upper byte must be zero).

Step 4: Rewrite Sound

Write a new music driver that reads your pattern data and emits VDP sound commands instead of AY register writes. The pattern data format can remain the same; only the output stage changes.

Step 5: Rewrite Input

Replace port reads with MOS keyboard map reads. The keymap approach is simple and provides simultaneous key detection.

Step 6: Rewrite Loading

Replace esxDOS file operations with MOS file API calls. The pattern is similar; only the function numbers and calling convention differ.

Step 7: Test and Tune

Run the game. Verify sprite positions, collision boxes, scrolling speed, sound timing. The VDP's asynchronous processing means that visual updates may arrive one frame later than expected — adjust your game timing if needed.

What Each Platform Forces You To Do Better

The Spectrum teaches cycle-level efficiency and creative constraint-solving: you learn to count T-states, exploit memory layout, and invent techniques like multi-colour and attribute-only effects that exist only because of the hardware's limitations. The Agon teaches system architecture: managing an asynchronous coprocessor, structuring command pipelines, and building asset conversion tools for larger data volumes. The Spectrum makes you a better optimiser; the Agon makes you a better system designer. Both skills transfer.

A Note on eZ80 Instructions

The eZ80 adds several instructions that Z80 programmers will appreciate. The most useful for game development:

LEA (Load Effective Address). Compute an address from a base register plus an 8-bit signed displacement, without modifying the base register:

```
; eZ80: LEA IX, IY + offset
; Compute IX = IY + displacement without changing IY
    LEA IX,IY+ENT_SIZE    ; IX points to next entity
```

On the Z80, this requires PUSH IY / POP IX / LD DE,ENT_SIZE / ADD IX,DE — four instructions and 40+ T-states. LEA does it in one instruction.

TST (Test Immediate). AND the accumulator with an immediate value and set flags, without modifying A:

```
; eZ80: TST A, mask
; Test bits without destroying A
    TST A,$80            ; test sign bit
    jr  nz,.negative     ; branch if bit 7 set, A unchanged
```

On the Z80, you would need BIT 7,A (which does not work with arbitrary masks) or PUSH AF / AND mask / POP AF (expensive).

MLT (Multiply). 8x8 unsigned multiply, result in a 16-bit register pair:

```
; eZ80: MLT BC
; B * C -> BC (16-bit result)
    ld  b,sprite_width
    ld  c,frame_number
    mlt bc                ; BC = B * C
```

On the Z80, multiplication requires a loop or a lookup table. MLT is a single instruction. For game logic — computing sprite offsets, tile map indices, animation frame positions — this is a substantial simplification.

Summary

- The Agon Light 2 runs the same Z80 instruction set as the Spectrum, but with 24-bit addressing (ADL mode), 512 KB flat RAM, hardware sprites and tilemaps via the VDP coprocessor, and a ~5x larger frame budget.
- **ADL mode** is the native mode. Run your game in ADL mode. Avoid Z80-compatible mode for anything other than running legacy code that cannot be converted. Mode switching via .LIL/.SIS suffixes is available but adds complexity and risk.
- **Game logic ports directly:** entity systems, collision detection, fixed-point physics, state machines, and AI all transfer with minimal changes (mainly widening pointers from 16-bit to 24-bit).
- **Rendering must be rewritten:** the Spectrum's direct framebuffer access is replaced by VDP command sequences for sprites, tilemaps, and scrolling. CPU rendering cost drops dramatically, but you now manage an asynchronous command pipeline.
- **Sound must be rewritten:** AY register writes are replaced by VDP sound commands. The pattern data can remain the same; only the output driver changes.

- **Memory architecture simplifies:** no banking, no shadow screen tricks, no compression mandated by scarcity. All assets can be resident simultaneously.
- **Spectrum tricks that become irrelevant on Agon:** self-modifying code for speed, stack-pointer rendering, pre-shifted sprite copies for memory/speed trade-offs, interleaved screen address calculations, attribute-based visual effects.
- **Spectrum tricks that still matter on Agon:** tight inner loops, register-efficient code, data-oriented struct layout, precomputed lookup tables.
- **Each platform teaches different skills:** the Spectrum teaches cycle-level efficiency and creative constraint-solving; the Agon teaches system architecture, coprocessor communication, and data pipeline management.
- **The eZ80 adds useful instructions:** LEA for address computation, MLT for hardware multiply, TST for non-destructive bit testing — all simplifications that eliminate multi-instruction Z80 patterns.

Sources: Zilog eZ80 CPU User Manual (UM0077); Agon Light 2 Official Documentation, The Byte Attic; Dean Belfield, “Agon Light — Programming Guide” (breakintoprogram.co.uk); FabGL Library Documentation (fabgl.com); Agon MOS API Documentation (github.com/AgonConsole8/agon-docs); Chapters 15–21 of this book.

Chapter 23: AI-Assisted Z80 Development

“Z80 they still don’t know.” – Introspec (spke), Life on Mars, 2024

This book was partially written with AI assistance. The chapter you are reading was drafted by Claude Code. The assembler used to build the examples – MinZ’s `mza` – was built with AI assistance. The “Antique Toy” companion demo that this book documents was coded in a feedback loop between a human and an AI agent. If that makes you uncomfortable, good. That discomfort is worth examining.

This is the most self-aware chapter in the book. We are going to look honestly at what AI assistance means for Z80 development in 2026 – where it genuinely helps, where it confidently fails, and where the answer is a frustrating “it depends.” We will do this with real examples, real code, and real failure cases, because the demoscene has never had patience for hype.

23.1 The Historical Parallel: HiSoft C on ZX Spectrum

Before we talk about AI, let us talk about another attempt to bring higher-level tools to the ZX Spectrum.

In 1998, *Spectrum Expert* #02 – the same issue where Dark and STS published their midpoint 3D method (Chapter 5) – reviewed the HiSoft C compiler for ZX Spectrum. The verdict was mixed. The compiler produced code that ran “10–15x faster than BASIC.” It supported 33 reserved keywords, offered a `stdio.lib` that provided graphics at BASIC capability levels, and included `gam128.h` for 128K memory bank access.

But it had no floating-point support.

Think about that for a moment. A C compiler. On a machine where floating-point is already handled by the ROM’s RST \$28 calculator, sitting there in 16K of free code. And the compiler could not use it.

The conclusion from the *Spectrum Expert* reviewer was precise: “useful for speed-critical work where float isn’t needed.” A tool with clear strengths and hard limits, evaluated honestly.

HiSoft Pascal HP4D told a similar story. The compiler occupied 12K, leaving roughly 21K for programs. It supported real types and trig functions – SIN, COS, SQRT – and was “suitable for data processing and computational mathematics.” But 21K

for your program, on a machine where a single uncompressed screen takes 6,912 bytes, means you are writing small programs or nothing.


Higher-level languages on constrained hardware have always been a compromise. They accelerate certain tasks enormously. They make other tasks impossible. The question was never “is HiSoft C good or bad?” but “what is it good *for*, and what should you still write in assembly?”

AI-assisted Z80 development is the same kind of compromise. Different shape, same question.

23.2 The Claude Code Feedback Loop

Here is how AI-assisted Z80 development actually works in practice. It is not magic. It is a loop.

The Loop



You describe what you want. The AI generates Z80 assembly. You assemble it. It fails – wrong syntax, wrong label format, wrong assembler dialect. You paste the error back. The AI fixes it. You assemble again. This time it compiles. You run it in the emulator. The output is wrong – the screen fills with garbage, the border is the wrong colour, nothing appears at all. You describe what you see. The AI adjusts. You run again.

This loop is not fundamentally different from how a human programmer works, except for one critical detail: the AI generates its first attempt much faster and gets some classes of errors much more predictably.

A Concrete Example

Suppose you ask Claude Code to write a routine that fills the ZX Spectrum attribute area with a diagonal stripe pattern. Here is a plausible first attempt:

```
; Fill attributes with diagonal stripes
; Pattern: alternating INK colours along diagonals
    ORG $8000

fill_diagonal:
    ld    hl, $5800          ; attribute memory start
    ld    b, 24              ; 24 rows

.row_loop:
    ld    c, 32              ; 32 columns per row
    ld    a, 24
    sub   b                  ; A = row index (0-23)

.col_loop:
    push  af
    add   a, c               ; A = row + column
    and   7                 ; clamp to 0-7 (INK colours)
    or    $40               ; set BRIGHT bit
    ld    (hl), a           ; write attribute
    inc   hl
    pop   af
    dec   c
    jr    nz, .col_loop

    djnz  .row_loop
    ret
```

This is reasonable code. It assembles. It runs. But there are problems. The column counter C counts down from 32, not up from 0, so the diagonal goes the wrong direction. The `or $40` sets bit 6 (BRIGHT), which is correct, but the visual effect has diagonals running top-right to bottom-left instead of the more natural top-left to bottom-right. And the PUSH/POP AF in the inner loop costs 22 T-states per attribute byte – unnecessary overhead if you restructure the calculation.

None of these are *wrong* in the sense that the code crashes. They are *wrong* in the sense that a Z80 programmer would not write the code this way. A human who has filled attributes hundreds of times would compute the diagonal index differently, avoid the PUSH/POP, and get the direction right on the first attempt because the pattern of `row + column` is second nature.

Here is the version you arrive at after two iterations:

```
fill_diagonal:
    ld    hl, $5800
    ld    d, 0              ; row index
```

```

.row_loop:
    ld    e, 0                ; column index
    ld    b, 32

.col_loop:
    ld    a, d
    add   a, e                ; diagonal = row + col
    and   7
    or    $40                 ; BRIGHT + INK colour
    ld    (hl), a
    inc   hl
    inc   e
    djnz  .col_loop

    inc   d
    ld    a, d
    cp    24
    jr    nz, .row_loop
    ret

```

Cleaner. No PUSH/POP. Diagonals run the right direction. The inner loop costs $4 + 4 + 7 + 4 + 7 + 7 + 6 + 4 + 13 = 56$ T-states per byte – not brilliant, but functional for a fill routine that runs once.

AI-generated diagonal fill pattern showing bright colour stripes across the ZX Spectrum attribute grid

Figure 47: AI-generated diagonal fill pattern showing bright colour stripes across the ZX Spectrum attribute grid

The point is not that the AI wrote bad code. The point is that the *loop* – prompt, generate, assemble, test, fix, test again – is the actual workflow. AI assistance does not eliminate the need to understand Z80. It shifts the bottleneck from writing code to evaluating code.

What Makes the Loop Fast

The loop is faster with AI than without for specific categories of work:

Boilerplate. The ORG directive, the HALT loop, the border colour timing harness from Chapter 1, the attribute fill skeleton, the AY register write subroutine, the LDIR setup, the interrupt mode setup. Every Z80 project starts with the same 30-50 lines. The AI generates these correctly and instantly. A human types them from memory. The AI is faster.

Iteration on a known pattern. “Now make the diagonal go the other direction.” “Add a frame counter so it animates.” “Make the colours cycle through BRIGHT and non-BRIGHT.” Each iteration is a small delta on existing code. The AI applies the delta faster than manual editing, and the changes are usually correct.

Test harness generation. “Write a test that fills memory at \$C000 with values 0-255, calls the multiply routine at \$8000, and checks the results against a table.”

The AI generates this kind of scaffolding code quickly and reliably. The structure of a test – set up inputs, call routine, compare outputs – is well within AI competence.

Documentation and comments. “Add cycle counts to every instruction in this inner loop.” The AI knows the Z80 timing tables and applies them correctly in straightforward cases. This is tedious human work that machines handle well.

What Makes the Loop Slow

Novel algorithms. When you ask for something the AI has not seen – a new unrolling strategy, a trick that exploits Z80 flag behaviour in a specific way, a code generation scheme tailored to your exact data layout – the AI generates plausible-looking code that is often subtly wrong. Worse, it is wrong in ways that compile and run but produce incorrect results. You spend more time debugging AI-generated code than you would have spent writing it yourself.

Cycle counting under pressure. The AI can count cycles for isolated instructions. But when you need to know the exact cost of a routine that spans contended and uncontended memory, involves conditional branches with different taken/not-taken costs, and must fit within a budget of 2,340 T-states (one scanline minus a few instructions), the AI’s estimates are unreliable. It will tell you “approximately 2,200 T-states” when the actual cost depends on branch probabilities and memory alignment. This is where DeZog becomes essential.

Creative effect design. “Design a visual effect that looks good and fits in 8,000 T-states” is a question the AI cannot answer. It can implement an effect you describe. It cannot invent one. The creative core of demoscene work – finding a computation scheme that produces compelling visuals within a tight budget – remains entirely human.

23.3 DeZog Integration: The Other Half of the Loop

If the AI generates the code, DeZog tells you whether it works.

DeZog is a VS Code extension that provides a Z80 debugger interface. It connects to emulators (ZEsarUX, CSpect, MAME) or its own internal Z80 simulator and gives you breakpoints, memory inspection, register watches, call stacks, and disassembly views – the standard debugging experience that modern developers expect, applied to Z80 code.

The AI + DeZog Workflow

The most productive workflow for AI-assisted Z80 development combines Claude Code with DeZog in a tight loop:

1. **Claude Code generates a routine** – say, an 8x8 multiply.
2. **You assemble it** with `mza` and load it into the DeZog-connected emulator.
3. **You set a breakpoint** at the entry point and step through.
4. **You watch the registers** at each step. Does A contain the right intermediate value after the first `ADD A,B`? Does the carry flag get set when it should?
5. **You spot a divergence** – the high byte of the result is wrong. You take a screenshot of the register state or copy the values.

6. **You paste the divergence back to Claude Code** – “After 6 iterations of the shift loop, A = \$3C but should be \$78. Here are the register values at the breakpoint.”
7. **Claude Code identifies the bug** – usually a missing shift, a wrong register choice, or an off-by-one in the loop count.
8. **You fix, reassemble, retest.**

This workflow is powerful because it gives the AI what it lacks: ground truth. The AI is good at reasoning about code structure but poor at mentally simulating Z80 execution over many iterations. DeZog provides the actual execution state. The AI reasons about the gap between expected and actual state. Together, they converge on correct code faster than either alone.

Memory Inspection for Data-Heavy Code

For routines that manipulate memory – screen fills, table generation, buffer operations – DeZog’s memory view is indispensable. You can set a breakpoint after your sine table generation routine and inspect the 256 bytes at the table address. Are they symmetric? Do they peak at the right value? Do they cross zero at the right position?

This is especially valuable for AI-generated lookup tables. Claude Code can generate a routine that computes a 256-byte sine table using the parabolic approximation from Chapter 4. The routine will usually *almost* work – the shape is right, the range is right, but there might be an off-by-one in the index that shifts the entire table by one position, or a sign error that inverts one quadrant. DeZog lets you see the table directly and compare it to known-good values.

What DeZog Cannot Do (Yet)

DeZog does not currently integrate with AI agents programmatically. You, the human, are the bridge – reading register values, pasting them into the prompt, applying fixes. An AI agent that could set breakpoints and iterate autonomously would close the loop for well-defined problems. For creative and architectural work, the human remains in the loop.

23.4 When AI Helps, When It Does Not

Let us be specific. Not “AI is good at some things” – specific categories with specific assessments.

AI Helps: High Confidence

Instruction encoding and cycle counts. The AI has the Z80 instruction set memorised: opcodes, byte counts, T-state costs. DJNZ taken = 13T, not taken = 8T. LDIR per byte = 21T except last = 16T. It gets these right consistently, with the caveat that it sometimes confuses Pentagon and 48K contended timing.

Boilerplate and scaffolding. ORG directives, HALT loops, AY register writes, screen clear routines, interrupt setup. Patterns seen thousands of times. Generated correctly, saves typing.

Dialect translation and code explanation. Converting between sjasmplus, mza, and z80asm syntax. Explaining what a block of Z80 assembly does – tracing logic, identifying patterns. Reading Z80 is easier than writing it, and the AI reads well.

AI Helps: Medium Confidence

Standard algorithms. Shift-and-add multiply, restoring division, Bresenham line drawing, LDIR-based scrolling. The AI generates working implementations of these, but they are usually textbook versions – correct but not optimised. A human would squeeze out 5-15% more speed through register allocation tricks, flag exploitation, and unrolling that the AI does not think to apply.

Memory layout and addressing. “Set up a 256-byte aligned table at \$xx00” or “calculate the attribute address for screen position (row, col).” The AI understands the Spectrum’s screen layout and generates correct address calculations, though it occasionally gets the third-boundary crossing wrong in the pixel memory interleave.

Simple self-modifying code. Patching an immediate operand, changing a jump target, swapping an instruction. The AI understands the concept and generates correct examples for simple cases. Complex self-modification – where the modified code’s behaviour depends on multiple patches interacting – is unreliable.

AI Does Not Help: Low Confidence

Novel inner loop optimisation. This is the big one. When you need to shave 3 T-states off an inner loop that runs 6,144 times per frame – when 3 T-states is the difference between 50 fps and 48 fps – the AI cannot reliably find the optimisation. It will suggest standard approaches (unrolling, table lookup, register substitution) but will not discover the *specific* trick that this *specific* data layout and register allocation permit.

Introspec’s `ld a,(hl) : inc l : dec h : add a : add a : add (hl)` rotozoomer inner loop from his Illusion analysis (Chapter 7) is 95 T-states for 4 chunky pixels. The genius is in the choice to use `inc l` instead of `inc hl` (saving 2 T-states, 6 for the pair) and to exploit the fact that `add a` (a doubling) is 4T while `sll a` (a shift, which does the same thing) is 8T. These are the kinds of micro-decisions that accumulate into the difference between a demo that runs and a demo that does not. AI does not make these decisions well, because they require understanding the *global* context of register pressure, memory alignment, and frame budget simultaneously.

Contended memory timing. The delay pattern on original Spectrums (6, 5, 4, 3, 2, 1, 0, 0 extra T-states per 8-T-state period) interacts with instruction timing in ways the AI cannot reliably predict. Introspec documented this in “GO WEST” (Hype, 2015). The AI knows the facts but cannot apply them to calculate actual runtime of mixed contended/uncontended routines.

Flag-based tricks and aesthetic judgement. The AI knows that `ADD A,A` sets carry from bit 7 – usable as both branch condition and multiplication – but does not spontaneously combine such facts into novel optimisations. And it cannot make creative decisions: which colours work for a plasma, how a tunnel should feel, whether a scroller should bounce or sine-wave.

23.5 Case Study: Building MinZ

MinZ is a programming language for Z80 and eZ80 systems, built by Alice with substantial AI assistance over the course of 2024-2026. It compiles modern, readable code to efficient Z80 assembly. The project is real, open-source, and at version 0.18.0 as of this writing.

MinZ is relevant to this chapter for two reasons. First, it is a case study in AI-assisted development of Z80-targeting tools. Second, it is itself an example of the HiSoft C pattern – a higher-level language on constrained hardware, with familiar strengths and limits.

What MinZ Is

MinZ provides typed variables (`u8`, `u16`, `i8`, `i16`, `bool`), functions with multiple returns, control flow (`if/else`, `while`, `for i in 0..n`), structs, arrays, and a standard library covering maths, graphics, input, sound, and memory operations. It compiles to Z80 assembly via its own assembler (`mza`), runs on its own emulator (`mzx`), and targets ZX Spectrum, CP/M, MSX, and Agon Light 2.

The toolchain includes four standalone tools:

- **mza** — Z80 assembler with macros, multiple output formats (`.sna`, `.tap`, `.com`, `.rom`, `.bin`), and multi-platform targets
- **mzx** — ZX Spectrum emulator with headless CLI mode, automated screenshots, keystroke injection, and frame-precise capture
- **mzd** — Z80 disassembler with IDA-like recursive descent analysis, cross-references, T-state counting, and reassemblable output
- **MinZ compiler** — compiles MinZ source to Z80 assembly via `mza`

A MinZ program looks like this:

```
import stdlib.input.keyboard;
import stdlib.time.delay;

fun main() -> void {
    clear_screen();
    draw_circle(128, 96, 50);

    loop {
        wait_frame();
        let dx = get_key_dx();
        // Move sprite based on input...
    }
}
```

This compiles to Z80 assembly, assembles to a binary, and runs on real or emulated hardware. The self-contained toolchain – compiler, assembler, emulator, disassembler – means no external dependencies.

Where AI Helped Build MinZ

The compiler itself. MinZ’s compiler is written in Go (~90,000 lines). The bulk of the code generation – translating MinZ’s intermediate representation to Z80

assembly – was written in an AI-assisted loop. The pattern: describe the semantics of a language feature, generate the code generator, test against the emulator, fix discrepancies. For standard features like arithmetic expressions, function calls, and control flow, this loop converged quickly. Claude Code generated correct code generators for if/else and while loops on the first or second attempt.

The assembler. `mza`, the MinZ Z80 assembler, was built with AI assistance. It supports the full Z80 instruction set, macros, multiple output formats, and two-pass assembly. The instruction encoding table – which maps mnemonics to opcodes, handling all the Z80’s irregular prefix byte patterns (CB, DD, ED, FD) – was generated by the AI and verified against the Z80 data sheet. This is exactly the kind of systematic, table-driven code that AI handles well.

The emulator. `mzx` achieves 100% Z80 instruction coverage, including all undocumented opcodes (ED prefix NOPs, DDCB/FDCB indexed bit operations). The AI generated the initial implementation for each instruction from the Z80 manual; the test suite (also AI-generated) caught edge cases – flag behaviour on overflow, the half-carry flag on DAA, interrupt timing. But `mzx`’s most useful feature – built entirely through the AI feedback loop – is its headless CLI mode:

```
mzx --load code.bin@8000,data.bin@C000 --set PC=8000,SP=FFFF,EI,IM=1
mzx --model 128k --tap demo.tap --exec 'LOAD ""' --frames 500
mzx --run effect.bin@8000 --frames DI:HALT --dump-keyframes ./frames/
mzx --model pentagon --trd disk.trd --type "RUN\n" --screenshot grab.png
```

The `--run` flag loads a binary at a given address and starts execution – no ROM, no BASIC, no loading screen. The `--frames DI:HALT` trigger captures the screenshot at the exact moment the code signals “frame complete” by disabling interrupts before a HALT. The `--dump-keyframes` flag saves only frames where the screen changed – an automated visual regression test. The `--exec` and `--type` flags inject BASIC commands and keystrokes, allowing fully automated testing of programs that expect user interaction.

This book’s screenshot pipeline uses `mzx` directly. Every code example screenshot in these pages was generated by:

```
mzx --run build/example.bin@8000 --frames 50 --screenshot build/ch09_plasma.png
```

Twenty-one examples, zero manual intervention, reproducible with make screenshots.

The disassembler. `mzd` performs recursive descent analysis – the same technique used by IDA Pro. Given a binary, it traces all execution paths from entry points, separates code from data, detects strings, generates cross-references, and auto-labels jump targets:

The `--cycles` flag adds T-state counts to every instruction – automating the exact work that Introspec did by hand in his 2017 teardown of X-Trade’s Illusion. The `--target spectrum` flag annotates system calls (RST \$10 for character output, port \$FE for border/keyboard). The `-R` flag produces reassemblable output, closing the disassemble-modify-reassemble loop.

The AI built both `mzd`’s instruction decoder (systematic table-driven work) and its analysis engine (recursive descent, control flow graph construction). The platform-

specific ABI knowledge (which ZX Spectrum ROM calls do what) was partly AI-generated, partly pulled from existing documentation.

The standard library and peephole optimiser. Ten stdlib modules (maths, graphics, input, sound, etc.) and 35+ peephole patterns (“replace LD A,0 with XOR A”). Both were AI-generated and human-refined. The AI knows the instruction set well enough to suggest valid simplifications; the human verifies semantic correctness.

Where AI Did Not Help Build MinZ

True Self-Modifying Code (TSMC). MinZ’s most distinctive feature is TSMC – the compiler can emit code that rewrites its own instructions at runtime for performance. A single-byte opcode patch (7-20 T-states) replaces a conditional branch sequence (44+ T-states). The *concept* of TSMC was Alice’s invention, not the AI’s. The AI could not have proposed “what if the compiled code patched its own opcodes to change behaviour at runtime?” because the idea requires understanding both the compilation model and the Z80’s instruction encoding at a level the AI does not reach unprompted.

The parser. MinZ originally used tree-sitter for parsing but hit out-of-memory issues on large files. The replacement – a hand-written recursive descent parser in Go – was designed by Alice, informed by AI consultation (GPT-4, o4-mini, and Claude were all asked for architectural advice). The AI colleagues agreed that a hand-written parser was the right approach and suggested keeping tree-sitter’s test corpus. But the parser’s actual grammar design – how MinZ’s syntax maps to AST nodes – was human work. The AI could generate parser code for individual grammar rules but could not design the grammar itself.

Register allocation for the code generator. Deciding which variables live in which Z80 registers, when to spill to memory, and how to handle the Z80’s irregular register file (only certain registers can be used with certain instructions) is a constraint satisfaction problem that the AI handles poorly. It generates code that works but wastes registers, uses unnecessary memory stores, and misses opportunities to keep hot values in registers across basic blocks.

The MinZ Verdict

MinZ could not exist without AI assistance. The sheer volume of systematic code – the instruction encoder, the emulator, the disassembler’s analysis engine, the standard library, the peephole patterns – would have taken one developer years to write manually. With AI assistance, MinZ went from concept to a four-tool ecosystem in roughly 18 months.

But MinZ’s *interesting* features – TSMC, the zero-cost lambda-to-function transform, the UFCS method dispatch, mzx’s DI:HALT trigger, mzd’s platform-aware ABI annotations – are human inventions. The AI implemented them, but did not conceive them.

This maps precisely to the HiSoft C pattern. The tool accelerates the routine work enormously. The creative work remains human. The compromise is real and worth making.

23.5b Sidebar: The Other AI — Brute-Force Superoptimisation

MinZ’s peephole optimiser knows 35+ patterns like “replace LD A, 0 with XOR A.” But how do you *find* such patterns? And how do you know which ones are actually safe?

Consider LD A, 0 $\rightarrow$ XOR A. Both set A to zero. Both take fewer bytes in the XOR form (1 byte vs 2). But XOR A clears the carry flag and sets the zero flag; LD A, 0 preserves all flags. If the code after this instruction tests carry, the “optimisation” is a bug. A human expert knows this. A neural-network-based AI *usually* knows this but sometimes forgets. A brute-force superoptimiser *proves* it by testing every possible input state.

z80-optimizer (by oisee, 2025) takes the brute-force approach to its logical conclusion. It enumerates every pair of Z80 instructions — all 406 opcodes $\times$ 406 opcodes = 164,836 pairs — and for each pair, tests whether a shorter replacement produces identical output across all possible register and flag states. No heuristics. No training data. No neural networks. Just exhaustive enumeration with full state equivalence verification.

The results: **602,008 provably correct optimisation rules** from a single run on an Apple M2 (34.7 billion comparisons in 3 hours 16 minutes). Some highlights:

Original sequence	Replacement	Savings
SLA A : RR A	OR A	3 bytes, 12T
LD A, 0 : NEG	SUB A	2 bytes
LD A, B : ADD A, 0	LD A, B : OR A	0 bytes, 4T
SCF : RR A	SCF : RRA	1 byte, 4T

The rules cluster into **83 unique transformation patterns** — families of replacements that share the same structural logic. For instance, the “load-then-test” family: LD A, r : ADD A, 0 $\rightarrow$ LD A, r : OR A applies to all register sources because the optimisation exploits flag behaviour, not register identity.

What makes z80-optimizer interesting for this chapter is not the specific rules — any experienced Z80 coder knows most of the common ones. It is the *methodology*. This is AI in the original sense: a machine that finds knowledge through search, not through learned patterns. The 602,008 rules include thousands that no human has catalogued, because they involve obscure opcode pairs that nobody writes deliberately but that compilers and code generators *do* produce.

The obvious next step — length-3 sequences — requires GPU brute force ($406^3 = 67$ million triples $\times$ all input states). Beyond that, stochastic search (STOKE-style) can explore the space of longer replacements without exhaustive enumeration.

For practical Z80 development, z80-optimizer complements the AI feedback loop from this chapter: Claude Code generates correct-but-unoptimised code, then z80-optimizer can mechanically verify whether any instruction pairs have shorter equivalents. One AI writes the code; the other AI proves how to shrink it.

Source: github.com/oisee/z80-optimizer (MIT license)

23.6 Honest Take: “Z80 They Still Don’t Know”

Introspec’s scepticism about AI’s Z80 capabilities is not generic technophobia. It comes from decades of experience pushing the Z80 to its absolute limits. When he says “Z80 they still don’t know,” he means something specific.

Consider the rotozoomer inner loop from his Illusion analysis. The effect walks through a texture at an angle, producing rotated and zoomed 2x2 chunky pixels. The inner loop is:

```
ld    a, (hl)    ; 7T    read texture byte
inc   l          ; 4T    next column (no carry needed: 256-aligned!)
dec   h          ; 4T    previous row
add   a,a        ; 4T    double (same as SLA A but 4T not 8T)
add   a,a        ; 4T    quadruple
add   a,(hl)     ; 7T    combine with second texture sample
      ; --- 30T per pixel pair
```

The key insight is `inc l` instead of `inc hl`. This saves 2 T-states but only works because the texture is aligned to a 256-byte boundary, so incrementing just L never needs to carry into H. The AI would use `inc hl` – the safe, general choice – and lose 2 T-states per iteration. Over 3,072 pixel pairs per frame, that is 6,144 T-states – nearly 10% of the frame budget on a 48K Spectrum.

This is what “Z80 they still don’t know” means. Not that the AI cannot write Z80 assembly – it can. Not that it gets the instructions wrong – it usually does not. But that it does not understand the *context* in which each instruction choice matters. It does not think in budgets. It does not see that `inc l` is safe here because of the alignment constraint three abstraction layers up. It does not know that 2 T-states per iteration, multiplied by thousands of iterations, is the difference between fitting in one frame and needing two.

What the AI Gets Wrong, Specifically

We tested Claude Code on several Z80 tasks from the demoscene domain. Here are representative results.

Task: Write a DOWN_HL routine. DOWN_HL moves the screen pointer in HL one pixel row down in the Spectrum’s interleaved video memory. It is the most commonly used graphics primitive on the platform. Introspec wrote a definitive article about it (Hype, 2020).

The AI’s first attempt:

```
down_hl:
    inc h          ; 4T    move down one pixel row
    ld  a, h       ; 4T
    and 7          ; 7T
    ret nz         ; 11T/5T still within the character row
    ld  a, l       ; 4T
    add a, 32      ; 7T    next character row
    ld  l, a       ; 4T
    ret c         ; 11T/5T if carry, crossed into next third
    ld  a, h       ; 4T
    sub 8          ; 7T
```

```
ld    h, a        ; 4T
ret                    ; 10T
```

This is actually correct – it handles all three screen thirds properly, including the boundary transitions. The AI can get standard routines like `DOWN_HL` right because the pattern is well-documented and appears in many Z80 references. The routine works: `inc h` advances the pixel row, `add a, 32` advances the character row when needed, the carry from the L addition correctly detects third boundaries, and `sub 8` undoes the spurious TT increment for the common case.

But “correct” is not the same as “good.” Introspec’s article presents a version by RST7 using a dual-counter approach that handles all boundaries in 2,343 T-states for a full-screen traverse. The naive approach above – the standard textbook version – costs 5,922 T-states. The gap between “works” and “works well” is a factor of 2.5x, and the AI does not bridge that gap. It produces the first version any competent programmer would write, not the version an expert would optimise toward.

Task: Generate an unrolled screen fill. Asked to generate an unrolled PUSH-based screen fill (the technique from Chapter 3), the AI produced correct code – PUSH pairs writing two bytes at a time, DI/EI to protect the stack pointer manipulation. But it did not think to arrange the data in reverse order (PUSH writes high byte first, to lower addresses), which means the fill pattern was backwards. A human who has written PUSH fills before accounts for this automatically.

Task: Optimise a given inner loop. Given a working inner loop and asked to make it faster, the AI suggested standard optimisations: unrolling, lookup tables, register substitution. These are valid. But it did not find the non-obvious optimisation – the one where you rearrange memory layout to enable `inc l` instead of `inc hl`, or use the carry flag from an addition as a branch condition instead of a separate comparison. The non-obvious optimisation requires understanding the full context of the routine, and the AI’s context window, while large, does not capture the *spatial* and *temporal* structure of a Z80 program the way a human expert’s mental model does.

Where Introspec Is Right

The deepest Z80 optimisations are not about knowing instructions. They are about understanding the interplay between memory layout, register allocation, instruction encoding, timing constraints, and visual output – simultaneously. This interplay is what Introspec means by “evolving a computation scheme” (Chapter 1). A computation scheme is a holistic design where every decision affects every other decision. The AI operates on code locally. The expert operates on the scheme globally.

The AI does not know Z80 in the sense that Introspec knows Z80. It has memorised the instruction set but not internalised the machine.

Where Introspec Is Not Quite Right

But “Z80 they still don’t know” implies the AI is useless for Z80 work, and that is not true either.

The AI is not trying to replace Introspec. It is trying to help Alice – a programmer who understands Z80 well enough to evaluate AI output but does not have decades

of inner-loop optimisation experience. For Alice, the AI's output is a starting point that is better than a blank screen. She does not need the AI to find the `inc l` trick. She needs it to generate the first 80% of the routine so she can spend her time on the last 20%.

The demoscene has always been about the last 20%. The AI does not change that. It changes how fast you get through the first 80%.

23.7 The “Antique Toy” Demo: AI in Practice

The companion demo for this book – “Antique Toy” – is a deliberate experiment: build a ZX Spectrum demo with AI assistance and document what happens.

The name is a nod to Introspec's *Eager* (2015, 1st place at 3BM openair). We are implementing effects inspired by *Eager* – the attribute tunnel with 4-fold symmetry, the chaos zoomer, 4-phase colour animation – plus Dark's midpoint 3D engine from *Spectrum Expert* #02.

What has worked: Effect prototyping – Claude Code generates working first drafts fast enough to try ideas that would otherwise not be worth the typing time. “What if the tunnel used 8-fold symmetry instead of 4-fold?” takes 15 minutes with AI-generated code instead of 2 hours manually. Tooling – the build system, asset pipeline, Makefile rules, and test harnesses were all AI-generated and work reliably. Code review – feeding the AI a routine and asking “what is wrong?” catches obvious mistakes (off-by-one errors, forgotten DI/EI, wrong port numbers) before they cost debugging hours.

What has not worked: Dark's midpoint 3D engine. The virtual processor with packed 2-bit opcodes and 6-bit point numbers was incorrectly decoded. The averaging instruction computed $(A+B)/2$ using `ADD A,B : SRA A`, which overflows for signed coordinates. Three debugging sessions, longer than writing it from scratch. Music integration failed similarly – the AI generated a player that conflicted with the effect code's use of shadow registers (`EXX, EX AF,AF'`). Both the player and the effect used shadow BC for different purposes, and the `EXX` in the interrupt handler swapped in stale values. This class of bug – system-level register conflicts across interrupt boundaries – requires understanding full system architecture, not just individual routines.

The honest assessment: “Antique Toy” is not finished. The effects work individually. Integration is ongoing. But AI assistance made the project *feasible* for a solo developer working evenings and weekends. The right question is not “does AI match a dedicated human team?” but “does AI assistance let more people make demos?” The answer, provisionally, is yes.

23.8 The Feedback Loop in Practice

A concrete example from the “Antique Toy” project: implementing 4-fold symmetry for the tunnel effect by copying the top-left 16x12 attribute quadrant to the other three quadrants with mirroring.

The prompt was specific: “Write a Z80 routine that copies the top-left 16x12 quadrant of the ZX Spectrum attribute area (\$5800) to the other three quadrants with appropriate mirroring.” Claude Code generated 47 lines that assembled on the first attempt.

Testing revealed the top-right quadrant was offset by one column. DeZog showed the problem: after the mirror loop decremented DE 16 times, the row-advance calculation forgot that DE had already been moved backwards. The code advanced DE by 32 (one row width) instead of the needed 48 (32 for the row + 16 to compensate for the mirror traversal). Pasting the register values into Claude Code – “After row 1, DE = \$581F (should be \$582F)” – produced the fix immediately. The bottom-right quadrant had the same error compounded. One more iteration fixed it.

Total: three iterations, roughly 25 minutes. Manual estimate for an experienced Z80 programmer: 40-60 minutes. For a newcomer: 2-3 hours. The AI saved time on initial generation. The debugging took the same time regardless of who wrote the code.

23.9 Building Your Own AI-Assisted Workflow

The practical setup: VS Code with Z80 Macro Assembler extension and Z80 Assembly Meter. Claude Code (or any code-capable LLM). An assembler (mza or sjasmpplus). DeZog connected to an emulator. A Makefile.

The workflow: **Start with the AI** – describe what you want with specifics (target machine, memory addresses, assembler syntax). **Assemble immediately** – do not read the AI’s code carefully; assemble it, paste errors back. **Test with border colours** – wrap AI-generated routines in the timing harness from Chapter 1. **Debug with DeZog** – set breakpoints, find the first register divergence, report it to the AI. **Iterate** – usually 2-5 rounds for moderate complexity; more than 5 means the AI is failing and you should write it yourself. **Optimise yourself** – once correct, profile and apply the techniques from Chapters 1-14.

Prompt Engineering for Z80

Good prompt: “Write a Z80 routine for ZX Spectrum 128K (Pentagon timing) that copies 16 bytes from the address in HL to screen memory at (DE), with the screen address following the Spectrum interleave pattern. After each byte, advance DE to the next pixel row using the standard down_hl method. Use mza syntax. Include cycle counts.”

Bad prompt: “Write a sprite routine for the Spectrum.”

The good prompt specifies machine, assembler, addresses, behaviour, and output format. The bad prompt leaves everything ambiguous, and the AI will fill the gaps with wrong assumptions.

For optimisation prompts, give a concrete target: “This routine takes ~3,200 T-states. I need it under 2,400. Do not change the interface (HL = source, DE = destination, B = height). Pentagon timing.” A performance target and interface

constraint force the AI to look for real optimisations instead of restructuring the calling convention.

23.10 The Broader Picture

AI assistance does not change the abstraction level of the output – the Z80 still executes the same instructions at the same speeds. What it changes is the speed of the input: how fast you go from idea to working (if unoptimised) code. The demoscene’s experts will still write better inner loops than any AI, but AI-assisted tooling lowers the entry barrier enough that more people can start making demos and learn the deep tricks for themselves.

Summary

- **AI-assisted Z80 development follows a feedback loop:** prompt, generate, assemble, test, debug, iterate. The AI generates the first draft fast; the human evaluates and refines. The loop typically takes 2-5 iterations for a routine of moderate complexity.
- **AI is reliable for** instruction encoding, cycle counts, boilerplate, dialect translation, and code explanation. It is moderately reliable for standard algorithms and simple self-modifying code. It is unreliable for novel optimisation, contended memory timing, creative effect design, and deep flag-based tricks.
- **DeZog integration** closes the gap between AI output and correct code. The human reads register states from the debugger and feeds divergences back to the AI, which reasons about the mismatch. Programmatic AI-debugger integration does not yet exist but is the obvious next step.
- **The MinZ case study** shows the pattern clearly: AI assistance made it possible for one developer to build a complete language toolchain (compiler, assembler, emulator, standard library) in 18 months. The routine work – instruction encoding, test generation, standard library functions – was AI-generated. The creative work – TSMC, zero-cost abstractions, grammar design – was human.
- **Introspec’s scepticism is valid:** AI does not understand Z80 the way an expert does. It does not think in budgets, does not see cross-cutting constraints, does not find non-obvious optimisations. The deepest demoscene work remains beyond AI’s reach.
- **The historical parallel holds:** HiSoft C was “10-15x faster than BASIC” but had no floats. AI-assisted Z80 development is dramatically faster for scaffolding and iteration but cannot match human experts for inner loop optimisation. Higher-level tools on constrained hardware have always been a compromise. The question is not “good or bad?” but “good *for what?*”
- **The practical workflow** combines Claude Code for code generation, DeZog for debugging, `mza` or `sjasmpplus` for assembly, and a Makefile for build automation. Start with AI, assemble immediately, test with border colours, debug with DeZog, optimise yourself.

- **The broader effect** is positive: AI assistance lowers the entry barrier to Z80 development without lowering the ceiling. More people can start; the experts are still needed for the deep work. This is good for the demoscene.
-

Try It Yourself

1. **The boilerplate test.** Ask your AI assistant to generate a ZX Spectrum 128K boot template: ORG at \$8000, disable interrupts, set up IM1, HALT loop with border colour timing harness. Assemble and run it. How many iterations did it take?
 2. **The optimisation test.** Write (or AI-generate) a working attribute fill loop. Measure its cost with border colour timing. Then ask the AI to make it faster. Measure again. Now optimise it yourself using techniques from Chapters 1-3. Compare all three versions: original, AI-optimised, human-optimised.
 3. **The DOWN_HL challenge.** Ask the AI to write a DOWN_HL routine. Test it on all 192 pixel rows. Does it handle the third-boundary transitions correctly? Compare to Introspec's analysis (Hype, 2020). This is a litmus test for AI Z80 competence.
 4. **The MinZ experiment.** Install the MinZ toolchain (mza, mzx, mzd). Assemble a screen fill with mza, run it headless with mzx --run fill.bin@8000 --frames 5 --screenshot fill.png, then disassemble a demo binary with mzd demo.bin --analyze --cycles --target spectrum. Compare the AI-built disassembler's T-state counts to your own hand-counted totals from Chapter 1.
 5. **The automated pipeline.** Write an effect, assemble it, and add it to a Makefile that runs mzx --screenshot for every binary. Run mzx --dump-keyframes to see exactly which frames produce visible changes. This is the same pipeline that generated every screenshot in this book.
 6. **Build something.** Pick an effect from an earlier chapter. Use AI assistance to write the first draft. Iterate until it works. Profile it. Optimise the inner loop by hand. Document each step. You have just experienced the workflow this entire chapter describes.
-

This is the final technical chapter. What follows are the appendices – reference tables, setup guides, and the instruction reference you will reach for every time you write Z80 assembly.

Sources: HiSoft C review (Spectrum Expert #02, 1998); Introspec “Technical Analysis of Illusion” (Hype, 2017); Introspec “DOWN_HL” (Hype, 2020); Introspec “GO WEST Parts 1-2” (Hype, 2015); z80-optimizer (oisee, 2025, github.com/oisee/z80-optimizer)

Glossary

Technical vocabulary used throughout *Coding the Impossible*. Terms are grouped by category; “First” indicates the chapter where the term is introduced or defined, “Also” lists chapters with significant usage.

A. Timing & Performance

Term	Definition	Canonical form	First	Also
T-state	One clock cycle of the Z80 CPU at 3.5 MHz (~286 ns). The fundamental unit of execution time on the Spectrum.	“T-states” in prose; “T” in code comments (e.g., ; 11T)	Ch01	Ch02–Ch23
Frame	One complete display refresh at ~50 Hz (PAL). Duration varies by machine model.	“frame”	Ch01	all

Term	Definition	Canonical form	First	Also
Frame budget	Total T-states between interrupts: Pentagon 71,680, ZX 128K 70,908, ZX 48K 69,888. Practical budget after HALT + ISR + music player: ~66,000-68,000 (Pentagon), ~55,000-60,000 (128K with screen writes during active display). See Ch15 tact-maps for per-region breakdown (top border, active display, bottom border).	“frame budget”	Ch01	Ch04, Ch05, Ch08, Ch10, Ch12, Ch14, Ch16, Ch17, Ch18, Ch21
Scanline	One horizontal line of the display. Width varies: 224 T-states on 48K/Pentagon, 228 T-states on 128K. The frame consists of 320 (Pentagon), 311 (128K), or 312 (48K) scanlines including borders.	“scanline” (one word)	Ch01	Ch02, Ch08, Ch15, Ch17

Term	Definition	Canonical form	First	Also
Con- tended memory	RAM pages (\$4000-\$7FFF) where ULA and CPU share the memory bus on Sinclair hardware. CPU accesses are delayed by 0-6 extra T-states per access in a repeating 8-T-state pattern. Pentagon clones have no contention.	“contended memory”	Ch01	Ch04, Ch05, Ch15, Ch17, Ch18, Ch21, Ch22, Ch23
Machine cycle (M-cycle)	A group of 3-6 T-states within an instruction. The first M-cycle (M1) is always the opcode fetch at 4 T-states.	“machine cycle” or “M-cycle”	Ch01	-
Border time	Scanlines outside the 192-line active display (top/bottom borders). No contention; ~14,000 T-states available on 128K.	“border time”	Ch01	Ch08, Ch15, Ch17
Timing harness	Debugging technique: set border colour to red before code, black after; stripe height on screen shows T-state cost.	“border-colour timing harness”	Ch01	Ch02, Ch03, Ch08, Ch18

B. Hardware — Sinclair & Clones

Term	Definition	Canonical form	First	Also
Z80	Zilog Z80A CPU at 3.5 MHz. 8-bit data bus, 16-bit address bus. The processor in all ZX Spectrum models.	“Z80”	Ch01	all
ULA	Uncommitted Logic Array. Custom chip generating the video signal and handling I/O (keyboard, tape, speaker). Shares the memory bus with the CPU on Sinclair hardware.	“ULA”	Ch01	Ch02, Ch08, Ch15
ZX Spectrum 48K	Original Sinclair model. 69,888 T-states/frame, 312 scanlines, contended lower 32K.	“48K” or “ZX Spectrum 48K”	Ch01	Ch11, Ch15
ZX Spectrum 128K	Extended model with 128K banked RAM, AY sound chip, two display pages. 70,908 T-states/frame, 311 scanlines.	“128K” or “ZX Spectrum 128K”	Ch01	Ch11, Ch15, Ch20, Ch21
Screen memory	\$4000-\$5AFF. Pixel data (\$4000-\$57FF, 6,144 bytes) + attribute area (\$5800-\$5AFF, 768 bytes). Interleaved layout for pixel rows.	“screen memory”	Ch01	Ch02, Ch08, Ch16, Ch17

Term	Definition	Canonical form	First	Also
Attribute area	768 bytes at \$5800-\$5AFF. One byte per 8x8 character cell: FBPPPIII (Flash, Bright, Paper×3, Ink×3). Controls colour.	“attribute area”	Ch02	Ch08, Ch09, Ch10
Attribute clash	Hardware limitation: only 2 colours (ink + paper) per 8x8 cell. Overlapping sprites force colour compromises.	“attribute clash”	Ch02	Ch08, Ch16
Inter-leaved screen layout	Pixel rows in video memory are not sequential. Address encodes Y as 010TTSSS.LLLC-CCCC across two bytes. Three 2,048-byte thirds.	“interleaved screen layout”	Ch02	Ch07, Ch16, Ch17, Ch22
Shadow screen	Second display page at \$C000 (bank 7) on 128K. Switched via port \$7FFD bit 3.	“shadow screen”	Ch08	Ch10, Ch15, Ch21
AY-3-8910	General Instrument Programmable Sound Generator. Three square-wave tone channels, one noise generator, one envelope generator. 14 registers. Standard on 128K models.	“AY-3-8910” (first use), then “AY”	Ch11	Ch12, Ch15, Ch21
Port \$FFFD / BFFD	Ch11	Ch12		
$ AYregisterselect/datawriteportson128K. "FFFD"$ $/$ $"$BFFD"$				

Term	Definition	Canonical form	First	Also
Port \$7FFD	128K memory paging and screen select port.	“\$7FFD”	Ch08	Ch12, Ch15, Ch21
Port <i>FE</i> <i>I/Oport</i> : <i>bits</i> 0 — —2 = <i>bordercolour</i> , <i>bit</i> 3 = <i>MIC</i> , <i>bit</i> 4 = <i>EAR</i> /speaker. Also keyboard input (active— <i>lowhalf</i> — <i>rows</i>). “FE”	Ch01	Ch02, Ch11, Ch18		
IM1 / IM2	Interrupt modes. IM1: handler at \$0038. IM2: vectored via I register + data bus byte; used for custom interrupt handlers (music players, threading).	“IM1” / “IM2”	Ch01	Ch03, Ch05, Ch11, Ch12
DivMMC	SD card interface for modern Spectrum use. Supports esxDOS.	“DivMMC”	Ch15	Ch20
ZX Spec- trum Next	FPGA-based enhanced Spectrum. Z80N CPU (extra instructions), Triple AY (3×AY), hardware sprites, copper co-processor, tilemap, 28 MHz turbo.	“ZX Spectrum Next” or “Next”	Ch11	Ch15

C. Hardware — Soviet/Post-Soviet Ecosystem

Term	Definition	Canonical form	First	Also
Pentagon 128	Most popular Soviet ZX Spectrum clone. No contended memory, 71,680 T-states/frame, 320 scanlines. The reference platform for the ZX Spectrum demoscene.	“Pentagon 128” (first use), then “Pentagon”	Ch01	Ch04, Ch05, Ch08, Ch11, Ch15, Ch16, Ch17, Ch18, Ch19, Ch20, Ch21, Ch22
Scorpion ZS-256	Soviet clone with TurboSound and extended memory. Pentagon-compatible timing.	“Scorpion ZS-256” (first use), then “Scorpion”	Ch04	Ch11, Ch15, Ch20
TurboSound	Two AY chips (2×AY) in one machine, providing 6 sound channels. Standard on Scorpion; available as add-on for Pentagon.	“TurboSound”	Ch11	Ch15, Ch20
TR-DOS	Disk operating system on Soviet clones via Beta Disk 128 interface. File format: .trd. The standard distribution format for demoscene compos and disk magazines.	“TR-DOS”	Ch15	Ch20
Beta Disk 128	Floppy controller interface standard on Pentagon and Scorpion clones.	“Beta Disk 128”	Ch15	Ch20

D. Hardware — Agon Light 2

Term	Definition	Canonical form	First	Also
Agon Light 2	Single-board computer with Zilog eZ80 CPU at 18.432 MHz and ESP32-based VDP. 512KB flat RAM, no banking.	“Agon Light 2” or “Agon”	Ch01	Ch11, Ch15, Ch18, Ch22
eZ80	Zilog eZ80 CPU. Z80-compatible instruction set; most instructions execute in fewer cycles. 24-bit flat addressing (ADL mode). ~368,640 T-states/frame at 50 Hz.	“eZ80”	Ch01	Ch15, Ch22
VDP	Video Display Processor. ESP32 microcontroller running FabGL library. Handles display, sprites, tilemaps, sound synthesis. Communicates with eZ80 over serial at 1,152,000 baud.	“VDP”	Ch02	Ch11, Ch15, Ch18, Ch22
ADL mode	24-bit addressing mode on the eZ80. Flat 512KB address space, no banking.	“ADL mode”	Ch15	Ch22
MOS	Operating system on the Agon. Provides API calls for VDP commands, keyboard, filesystem. waitvblank replaces Spectrum’s HALT for frame sync.	“MOS”	Ch18	Ch22

E. Techniques

Term	Definition	Canonical form	First	Also
Self-modifying code (SMC)	Writing to instruction bytes at runtime. Safe on Z80 (no instruction cache). Common patterns: patching immediate operands, changing jump targets, swapping opcodes.	“self-modifying code (SMC)” on first use; “SMC” subsequently	Ch03	Ch06, Ch07, Ch09, Ch10, Ch13, Ch16, Ch17, Ch22, Ch23
Unrolled loop	Trade code size for speed by repeating the loop body N times, eliminating loop counter overhead. Partial unrolling keeps DJNZ for the outer count.	“unrolled loop”	Ch02	Ch03, Ch07, Ch10, Ch16, Ch17
PUSH trick	Hijack SP to use PUSH for fast memory writes (5.5 T-states/byte vs 21 for LDIR). Must DI first to protect against interrupts using the stack.	“PUSH trick” or “stack-based output”	Ch03	Ch07, Ch08, Ch10, Ch16
LDI chain	Sequence of individual LDI instructions; 24% faster than LDIR for known-size copies. Combined with entry-point arithmetic for variable-length copies.	“LDI chain”	Ch03	Ch07, Ch09

Term	Definition	Canonical form	First	Also
LDPUSH	Fusing display data into LD DE,nn : PUSH DE executable code (21T per 2 bytes). Used in multicolor engines.	"LDPUSH"	Ch08	-
DOWN_HL	Classic routine to advance HL one pixel row down in the Spectrum's interleaved screen memory. 20T common case, 77T worst case (third boundary).	"DOWN_HL"	Ch02	Ch16, Ch17, Ch23
RET-chaining	Set SP to a table of addresses; each routine ends with RET, dispatching to the next. 10T per dispatch vs 17T for CALL.	"RET-chaining"	Ch03	-
Code generation	Writing machine code into a buffer at runtime, then executing it. Eliminates branches and loop counters from inner loops.	"code generation"	Ch03	Ch06, Ch07, Ch09
Compiled sprites	Sprites compiled into a sequence of PUSH/LD instructions with pre-loaded register pairs. Fixed image, maximum speed.	"compiled sprites"	Ch03	Ch16
Double buffering	Maintaining two display pages to avoid tearing. On 128K: Screen 0 (4000) and Screen1 (C000), switched via port \$7FFD.	"double buffering"	Ch05	Ch08, Ch10, Ch12

Term	Definition	Canonical form	First	Also
Dirty rectangles	Save/restore background under sprites before/after drawing. Avoid full-screen clears.	“dirty rectangles”	Ch08	Ch16, Ch21
Multi-color	Changing attribute values between ULA scanline reads to display more than 2 colours per 8x8 cell. Consumes 80–90% of CPU.	“multicolor”	Ch08	–
Page-aligned table	256-byte lookup table placed at an \$xx00 address so H holds the base and L is the index. Single-register indexing with zero overhead.	“page-aligned table”	Ch04	Ch06, Ch07, Ch09, Ch10
Lookup table	Pre-computed table of values for fast runtime access. Avoids expensive calculations in inner loops.	“lookup table”	Ch03	Ch04, Ch07, Ch09, Ch17, Ch19, Ch20, Ch22
Split counters	Restructure screen iteration to match the three-level hierarchy (third/char row/scanline), eliminating branching. 60% faster than naive DOWN_HL traversal.	“split counters”	Ch02	–

Term	Definition	Canonical form	First	Also
4-phase colour	4-frame cycle (2 normal + 2 inverted attributes) at 50 Hz. Persistence of vision averages the colours, creating additional perceived colours per cell.	"4-phase colour"	Ch10	-
Digital drums	Digital PCM sample played through AY volume register as 4-bit DAC (attack phase), transitioning to AY envelope (decay phase). Consumes ~2 frames of CPU per hit.	"digital drums" or "hybrid drums"	Ch12	-
Asynchronous frame generation	Decoupling visual frame production from display via a ring buffer. Generator writes frames ahead; display reads at steady 50 Hz. Absorbs CPU bursts from drum playback.	"asynchronous frame generation"	Ch12	-

F. Assembly Notation & Directives

Term	Definition	Canonical form	First	Also
ORG	Assembler directive setting the code origin address. ORG \$8000 for Spectrum examples.	ORG	Ch01	all code examples

Term	Definition	Canonical form	First	Also
EQU	Define a named constant. SCREEN EQU \$4000.	EQU	Ch02	all code examples
DB / DW / DS	Define Byte / Define Word / Define Space. DB -3 (negative values allowed in sjasmpus).	DB, DW, DS	Ch04	all code examples
ALIGN	Align to a power-of-two boundary. sjasmpus directive.	ALIGN	Ch07	-
IN-CLUE / INCBIN	Include source file / include binary data. sjasmpus directives.	INCLUDE / INCBIN	Ch14	Ch20, Ch21
DEVICE / SLOT / PAGE	sjasmpus directives for 128K memory banking emulation at assembly time.	DEVICE, SLOT, PAGE	Ch15	Ch20, Ch21
DISPLAY	sjasmpus directive printing a value at assembly time. Build-time diagnostics.	DISPLAY	Ch20	Ch21
\$FF	Hex notation. \$ prefix is preferred. #FF also accepted by sjasmpus.	\$FF	Ch01	all
%10101010 .label	Binary notation. Local label, scoped to nearest enclosing global label.	%10101010 .label	Ch01 Ch01	all all

Term	Definition	Canonical form	First	Also
sjasm-plus	Primary assembler (v1.21.1). Full Z80/Z80N instruction set, macros, Lua scripting, DE-VICE/SLOT/PAGE for banking, INCBIN, multiple output formats.	“sjasmpplus”	Ch01	Ch14, Ch20, Ch21, Ch23

G. Demoscene & Culture

Term	Definition	Canonical form	First	Also
Compo	Competition at a demoparty. Categories include demo, intro (size-limited), music, graphics.	“compo”	Ch13	Ch20
De-moparty	Event where demoscene productions are shown and judged. Major ZX parties: Chaos Constructions, DiHalt, CAFe, Multimatograf (Russia); Revision, Forever (Western Europe).	“demoparty” or “party”	Ch20	–
NFO / file_id.diz	Text files bundled with demo releases containing credits, requirements, and sometimes technical notes.	“NFO” / “file_id.diz”	Ch20	–

Term	Definition	Canonical form	First	Also
Making-of	Post-release article documenting the development process and technical decisions of a demo. Published on Hype or in disk magazines.	“making-of”	Ch12	Ch20
Part / effect	A visual section of a demo (tunnel, scroller, sphere, etc.). Multiple parts are sequenced by a demo engine.	“part” or “effect”	Ch09	Ch12, Ch20
Scripting engine	System that sequences demo parts and synchronises them to music. Two-level: outer script (effect sequence) + inner script (parameter variation within an effect).	“scripting engine”	Ch09	Ch12, Ch20
kWORK	Introspec’s command: “generate N frames, then show them independently of generation.” The bridge between scripting and asynchronous frame generation.	“kWORK”	Ch09	Ch12

Term	Definition	Canonical form	First	Also
Zapilator	Russian scene slang for a “precalculator” demo – one that pre-computes all frames before playback. Carries mild disapproval (implies no real-time rendering).	“zapilator”	Ch09	–

Key People

Name	Role	Chapters
Dark	Coder, X-Trade. Author of Spectrum Expert programming articles. Coder of Illusion.	Ch01, Ch04, Ch05, Ch06, Ch07, Ch10
Introspec (spke)	Coder, Life on Mars. Reverse-engineered Illusion. Authored Hype articles (technical analyses, GO WEST series, DOWN_HL). Coded Eager demo.	Ch01–Ch12, Ch15, Ch23
n1k-o	Musician, Skrju. Composed Eager soundtrack. Developed hybrid drum technique with Introspec.	Ch09, Ch12
diver4d	Coder, 4D+TBK (Triebkraft). GABBA demo (CAFe 2019). Pioneered video-editor sync workflow (Luma Fusion).	Ch09, Ch12
DenisGrachev	Coder. Old Tower, GLUF engine, Ringo engine, Dice Legends. RET-chaining technique.	Ch03, Ch08

Name	Role	Chapters
Robus	Coder. Z80 threading system, WAYHACK demo.	Ch12
psndcj (cyberjack)	Coder, 4D+TBK (Triebkraft). AY/TurboSound expertise. Break Space demo (Magen Fractal effect).	Ch01, Ch07, Ch14
Screamer (sq)	Coder, Skrju. Chunky pixel optimisation research (Born Dead #05, Hype). Development environment guide.	Ch01, Ch07
Ped7g	Peter Helcmanovsky. sjasmplus maintainer, ZX Spectrum Next contributor. Signed arithmetic and RLE feedback.	Ch04, Ch14, App F
RST7	Coder. Dual-counter DOWN_HL optimisation.	Ch02

Key Demos & Productions

Name	Author/Group	Year	Chapters
Illusion	X-Trade (Dark)	1996	Ch01, Ch04, Ch05, Ch06, Ch07, Ch10
Eager (to live)	Introspec / Life on Mars	2015	Ch02, Ch03, Ch09, Ch10, Ch12
GABBA	diver4d / 4D+TBK	2019	Ch12
WAYHACK	Robus	–	Ch12
Old Tower	DenisGrachev	–	Ch08
Lo-Fi Motion	–	–	Ch20

Key Publications

Name	Description	Chapters
Spectrum Expert (#01–#02)	Russian ZX disk magazine (1997–98). Dark's programming tutorials.	Ch01, Ch04, Ch05, Ch06, Ch10, Ch11

Name	Description	Chapters
Hype	Russian online demoscene platform (hype.retroscene.org). Technical articles, making-ofs, debates.	Ch01-Ch12, Ch23
Born Dead	Russian demoscene newspaper. sq's chunky pixel research (#05, ~2001).	Ch07
Black Crow	Russian ZX scene magazine. Early multicolor documentation (#05, ~2001).	Ch08

H. Algorithms & Compression

Term	Definition	Canonical form	First	Also
Shift-and-add multiply	Classic 8×8 unsigned multiply. Scan multiplier bits, shift-and-add. 196-204 T-states.	"shift-and-add multiply"	Ch04	Ch05
Square-table multiply	$A \times B = ((A+B)^2 - (A-B)^2)/4$ via lookup. ~61 T-states. Trades 512 bytes of tables for speed.	"square-table multiply"	Ch04	Ch05, Ch07
Logarithmic division	$\log(A/B) = \log(A) - \log(B)$. Two table lookups + subtraction. ~50-70 T-states. Low precision.	"logarithmic division"	Ch04	Ch05
Parabolic sine approximation	Approximate sine with parabola $y = 1 - 2(x/\pi)^2$. Max error ~5.6%. 256-byte table, signed.	"parabolic sine approximation"	Ch04	Ch05, Ch06, Ch09
Bresenham line drawing	Step along major axis with error accumulator for minor-axis steps. ~80 T-states/pixel naive; ~48 with Dark's 8×8 matrix method.	"Bresenham"	Ch04	-

Term	Definition	Canonical form	First	Also
Midpoint method	Rotate only basis vertices fully; derive remaining vertices by averaging. ~36 T-states per derived vertex vs ~2,400 for full rotation.	“midpoint method”	Ch05	-
Fixed-point arithmetic	Represent fractional values in integer registers. Common format: 8.8 (8-bit integer + 8-bit fraction).	“fixed-point”	Ch04	Ch05, Ch18, Ch19
AABB collision	Axis-Aligned Bounding Box overlap test. Four comparisons: left-right and top-bottom for two rectangles. ~70-156 T-states.	“AABB”	Ch19	Ch21
Backface culling	Skip back-facing polygons using cross-product normal Z-component test. ~500 T-states per face.	“backface culling”	Ch05	-
ZX0	Compression format by Einar Saukas. Excellent ratio, moderate decompression speed.	“ZX0”	Ch14	Ch20
LZ4	Fast decompression (~34 T-states/byte). The choice for real-time streaming.	“LZ4”	Ch14	-

Term	Definition	Canonical form	First	Also
Exomizer	High-ratio compressor for 8-bit platforms. Slow decompression.	“Exomizer”	Ch14	–
MegaLZ	Compression format. Good ratio/speed balance.	“MegaLZ”	Ch14	–
hrust1	Compression format common in Russian demoscene.	“hrust1”	Ch14	Ch20

This glossary is extracted from all 23 chapters of “Coding the Impossible.” For detailed treatment of any term, see the chapter listed under “First.”

Appendix A: Z80 Instruction Quick Reference

“An instruction that should take 7 T-states might take 13 if it lands on the worst phase of the contention cycle.” – Chapter 1

This is not a complete Z80 manual. It is a reference card for demoscene and game programmers on ZX Spectrum and Agon Light 2 – the instructions you actually use, the timings you need to know by heart, and the patterns that save T-states in inner loops.

All T-state counts assume **Pentagon timing** (no contention). Byte counts are the instruction encoding length. Flag columns show: **S** (sign), **Z** (zero), **H** (half-carry), **P/V** (parity/overflow), **N** (subtract), **C** (carry). A dash means unchanged; a dot means undefined.

8-Bit Load Instructions

Instruction	Bytes	T-states	Flags	Notes
LD r, r'	1	4	---	Fastest instruction. r, r' = A, B, C, D, E, H, L
LD r, n	2	7	---	Immediate load
LD r, (HL)	1	7	---	Memory read via HL
LD (HL), r	1	7	---	Memory write via HL
LD (HL), n	2	10	---	Immediate to memory
LD A, (BC)	1	7	---	
LD A, (DE)	1	7	---	
LD (BC), A	1	7	---	
LD (DE), A	1	7	---	
LD A, (nn)	3	13	---	Absolute address

Instruction	Bytes	T-states	Flags	Notes
LD (nn),A	3	13	---	Absolute address
LD r,(IX+d)	3	19	---	Indexed. Expensive - avoid in inner loops
LD (IX+d),r	3	19	---	Indexed
LD (IX+d),n	4	23	---	Indexed immediate
LD A,I	2	9	SZ0P0-	P/V = IFF2
LD A,R	2	9	SZ0P0-	P/V = IFF2; R = refresh counter

16-Bit Load Instructions

Instruction	Bytes	T-states	Flags	Notes
LD rr,nn	3	10	---	rr = BC, DE, HL, SP
LD HL,(nn)	3	16	---	
LD (nn),HL	3	16	---	
LD rr,(nn)	4	20	---	rr = BC, DE, SP (ED prefix)
LD (nn),rr	4	20	---	rr = BC, DE, SP (ED prefix)
LD SP,HL	1	6	---	Set up stack pointer
LD SP,IX	2	10	---	
PUSH rr	1	11	---	rr = AF, BC, DE, HL. 5.5T per byte
POP rr	1	10	---	5T per byte - fastest 2-byte read
PUSH IX	2	15	---	
POP IX	2	14	---	

8-Bit Arithmetic and Logic

Instruction	Bytes	T-states	Flags	Notes
ADD A,r	1	4	SZ.V0C	
ADD A,n	2	7	SZ.V0C	
ADD A,(HL)	1	7	SZ.V0C	
ADC A,r	1	4	SZ.V0C	Add with carry
ADC A,n	2	7	SZ.V0C	
SUB r	1	4	SZ.V1C	
SUB n	2	7	SZ.V1C	
SUB (HL)	1	7	SZ.V1C	

Instruction	Bytes	T-states	Flags	Notes
SBC A, r	1	4	SZ.V1C	Subtract with carry
CP r	1	4	SZ.V1C	Compare (SUB without storing result)
CP n	2	7	SZ.V1C	
CP (HL)	1	7	SZ.V1C	
AND r	1	4	SZ1P00	H always set, C always cleared
AND n	2	7	SZ1P00	
OR r	1	4	SZ0P00	Clears H and C
OR n	2	7	SZ0P00	
XOR r	1	4	SZ0P00	XOR A = zero A in 4T/1B (vs LD A, 0 = 7T/2B)
XOR n	2	7	SZ0P00	
INC r	1	4	SZ.V0-	Does not affect carry
DEC r	1	4	SZ.V1-	Does not affect carry
INC (HL)	1	11	SZ.V0-	Read-modify-write
DEC (HL)	1	11	SZ.V1-	Read-modify-write
NEG	2	8	SZ.V1C	A = 0 - A (two's complement negate)
DAA	1	4	SZ.P-C	BCD adjust - rarely used in demos
CPL	1	4	-1-1-	A = NOT A (one's complement)
SCF	1	4	-0-00	Set carry flag. N,H cleared. New behaviour on CMOS

Instruction	Bytes	T-states	Flags	Notes
CCF	1	4	-.0.	Complement carry. H = old C

16-Bit Arithmetic

Instruction	Bytes	T-states	Flags	Notes
ADD HL,rr	1	11	-.?0C	rr = BC, DE, HL, SP. Only affects H, N, C
ADC HL,rr	2	15	SZ.V0C	Full flag set
SBC HL,rr	2	15	SZ.V1C	Full flag set
INC rr	1	6	---	No flags affected
DEC rr	1	6	---	No flags affected
ADD IX,rr	2	15	-.?0C	rr = BC, DE, IX, SP

Key point: INC rr and DEC rr do **not** set the zero flag. You cannot use DEC BC / JR NZ as a 16-bit loop counter. Use DEC B / JR NZ for 8-bit loops with DJNZ, or test BC explicitly.

Rotate and Shift

Instruction	Bytes	T-states	Flags	Notes
RLCA	1	4	-0-0C	Rotate A left, bit 7 to carry and bit 0
RRCA	1	4	-0-0C	Rotate A right, bit 0 to carry and bit 7
RLA	1	4	-0-0C	Rotate A left through carry

Instruction	Bytes	T-states	Flags	Notes
RRA	1	4	-0-0C	Rotate A right through carry. Key for multiply loops
RLC r	2	8	SZ0P0C	CB-prefix. Full flag set
RRC r	2	8	SZ0P0C	
RL r	2	8	SZ0P0C	Rotate left through carry
RR r	2	8	SZ0P0C	Rotate right through carry
SLA r	2	8	SZ0P0C	Shift left arithmetic. Bit 0 = 0
SRA r	2	8	SZ0P0C	Shift right arithmetic. Bit 7 preserved (sign extend)
SRL r	2	8	SZ0P0C	Shift right logical. Bit 7 = 0
RLC (HL)	2	15	SZ0P0C	Read-modify-write
RL (HL)	2	15	SZ0P0C	Scroll pixel data left
RR (HL)	2	15	SZ0P0C	Scroll pixel data right
SLA (HL)	2	15	SZ0P0C	
SRL (HL)	2	15	SZ0P0C	
RLD	2	18	SZ0P0-	Rotate (HL) nibbles left through A. Useful for nybble graphics
RRD	2	18	SZ0P0-	Rotate (HL) nibbles right through A

Demoscene note: RLA/RRA (4T, 1 byte) only affect carry and bits 3,5 of F. The CB-prefix versions RL r/RR r (8T, 2 bytes) set all flags. In multiply loops, the accumulator versions save half the cost.

Bit Manipulation

Instruction	Bytes	T-states	Flags	Notes
BIT b,r	2	8	.Z1.0-	Test bit b of register
BIT b,(HL)	2	12	.Z1.0-	Test bit b of memory
SET b,r	2	8	---	Set bit b of register
SET b,(HL)	2	15	---	Set bit b of memory. Used in line drawing
RES b,r	2	8	---	Reset bit b of register
RES b,(HL)	2	15	---	Reset bit b of memory

Jump, Call, Return

Instruction	Bytes	T-states	Flags	Notes
JP nn	3	10	---	Absolute jump
JP cc,nn	3	10	---	Conditional: NZ, Z, NC, C, PO, PE, P, M. Same speed taken or not
JR e	2	12	---	Relative jump (-128 to +127)
JR cc,e	2	12/7	---	cc = NZ, Z, NC, C only. 7T if not taken

Instruction	Bytes	T-states	Flags	Notes
JP (HL)	1	4	---	Jump to address in HL. Fastest indirect jump
JP (IX)	2	8	---	Jump to address in IX
DJNZ e	2	13/8	---	Dec B, jump if NZ. 13T taken, 8T not taken
CALL nn	3	17	---	Push PC, jump to nn
CALL cc,nn	3	17/10	---	10T if not taken
RET	1	10	---	Pop PC. Used for RET-chaining dispatch
RET cc	1	11/5	---	5T if not taken
RST p	1	11	---	Call to \$00,\$08,\$10,\$18,\$20,\$28

Key comparisons for dispatch:

Method	T-states	Bytes
CALL nn	17	3
RET (as dispatch in RET-chain)	10	1
JP (HL)	4	1
JP nn	10	3

I/O Instructions

Instruction	Bytes	T-states	Flags	Notes
OUT (n),A	2	11	---	Port address = (A << 8) n. Border: OUT (\$FE),A

Instruction	Bytes	T-states	Flags	Notes
IN A, (n)	2	11	---	Port address = (A << 8) n.
OUT (C), r	2	12	---	Keyboard: IN A, (\$FE) Port address = BC. AY register write
IN r, (C)	2	12	SZ0P0-	Port address = BC. Sets flags
OUTI	2	16	.Z..1.	Out (HL) to port (C), inc HL, dec B
OTIR	2	21/16	01..1.	Repeat OUTI until B=0. 16T on last
OUTD	2	16	.Z..1.	Out (HL) to port (C), inc HL, dec B

AY-3-8910 port addresses on ZX Spectrum 128K:

Port	Address	Purpose
Register select	\$FFFD	LD BC, \$FFFD : OUT (C), A
Data write	\$BFFD	LD B, \$BF : OUT (C), r
Data read	\$FFFD	IN A, (C)

Typical AY register write sequence (24T + overhead):

```

out  (c), a      ; 12T  select register number (in A)
ld   b, $BF      ; 7T   switch to data port $BFFD
out  (c), e      ; 12T  write value (in E)
                        ; --- 41T total

```

Block Instructions

Instruction	Bytes	T-states	Flags	Notes
LDI	2	16	-0.0-	(DE) = (HL), inc HL, inc DE, dec BC. P/V = (BC != 0)
LDIR	2	21/16	-000-	Repeat LDI. 21T per byte, 16T last byte
LDD	2	16	-0.0-	(DE) = (HL), dec HL, dec DE, dec BC
DDIR	2	21/16	-000-	Repeat LDD. 21T per byte, 16T last byte
CPI	2	16	SZ.?1-	Compare A with (HL), inc HL, dec BC
CPIR	2	21/16	SZ.?1-	Repeat CPI. Stops when match or BC=0
CPD	2	16	SZ.?1-	Compare A with (HL), dec HL, dec BC
CPDR	2	21/16	SZ.?1-	Repeat CPD

LDI vs LDIR per-byte cost:

Method	Per byte	256 bytes	32 bytes	Savings
LDIR	21T (16T last)	5,371T	672T	-
LDI chain	16T	4,096T	512T	24% faster

An unrolled LDI chain costs 2 bytes per LDI (\$ED \$A0), but saves 5T per byte - 24% faster than LDIR. See Chapter 3 for entry-point arithmetic with LDI chains.

Exchange and Misc

Instruction	Bytes	T-states	Flags	Notes
EX DE,HL	1	4	---	Swap DE and HL. Free pointer swap
EX AF,AF'	1	4	---	Swap AF with shadow AF'
EXX	1	4	---	Swap BC,DE,HL with BC',DE',HL'. 6 registers in 4T
EX (SP),HL	1	19	---	Swap HL with top of stack. Useful for parameter passing
EX (SP),IX	2	23	---	Swap IX with top of stack
DI	1	4	---	Disable interrupts. Required before stack tricks
EI	1	4	---	Enable interrupts. Delayed one instruction
HALT	1	4+	---	Wait for interrupt. The frame sync instruction
NOP	1	4	---	Padding, timing
IM 1	2	8	---	Interrupt mode 1 (RST \$38). Standard Spectrum mode

Instruction	Bytes	T-states	Flags	Notes
IM 2	2	8	—	Interrupt mode 2. Uses I register as vector table high byte

The Demoscene “Fast” Instructions

These are the cheapest instructions per category – the building blocks of every optimised inner loop.

Fastest Register-to-Register Move

LD *r, r'* – **4T, 1 byte**. The minimum cost of any Z80 instruction. Includes LD A,A (effectively a NOP that does not affect flags).

Fastest Way to Zero a Register

XOR A – **4T, 1 byte**. Sets A to zero, sets Z flag, clears carry. Compare with LD A,0 at 7T/2 bytes. Always use XOR A unless you need flags preserved.

Fastest Memory Read

LD A, (HL) – **7T, 1 byte**. The minimum cost for any memory read. Other register sources (LD A, (BC), LD A, (DE)) are also 7T/1 byte, but HL is the only pointer that supports LD *r, (HL)* for all registers.

Fastest Memory Write

LD (HL), *r* – **7T, 1 byte**. Tied with read. Writing to a BC or DE pointer (LD (BC), A, LD (DE), A) is also 7T/1B but only works with A.

Fastest 2-Byte Write

PUSH *rr* – **11T, 1 byte** for 2 bytes = **5.5T per byte**. The fastest way to write data to memory, but only to where SP points (downward). Requires DI and hijacking the stack pointer. See Chapter 3.

Fastest 2-Byte Read

POP *rr* – **10T, 1 byte** for 2 bytes = **5T per byte**. Even faster than PUSH for reading. Use with SP pointing at a data table for ultra-fast lookups.

Fastest Block Copy

Method	Per byte	Notes
PUSH/POP pair	5.25T	Write: 5.5T, Read: 5T. But needs SP hijack
LDI (unrolled chain)	16T	No setup per byte. 24% faster than LDIR
LDIR	21T	Single instruction, but slow per byte
LD (HL), r + INC HL	13T	Manual loop body (no loop counter)
LD (HL), r + INC L	11T	Only works within 256-byte page

Fastest I/O

OUT (n), A – **11T, 2 bytes**. For fixed port addresses (border, etc.). For variable ports (AY), OUT (C), r at 12T/2 bytes is the only option.

Fastest Pointer Swap

EX DE, HL – **4T, 1 byte**. Exchange DE and HL contents instantly. No other register swap is this cheap. EXX also 4T/1 byte but swaps all three pairs at once.

Fastest Conditional Loop

DJNZ e – **13T taken, 8T not taken, 2 bytes**. Decrements B and jumps. Compare with DEC B / JR NZ, e at 4+12 = 16T/3 bytes. DJNZ saves 3T and 1 byte per iteration.

Fastest Indirect Jump

JP (HL) – **4T, 1 byte**. Jumps to the address in HL. Despite the misleading mnemonic, this does NOT read from memory at (HL) – it loads PC with the value of HL. Indispensable for jump tables and computed gotos.

Undocumented Instructions

These instructions are not in the official Zilog documentation but work reliably on all Z80 silicon (NMOS and CMOS), all ZX Spectrum clones, and the eZ80. They are widely used in demoscene code and supported by sjasmpus.

IXH, IXL, IYH, IYL (Half-Index Registers)

The IX and IY registers can be split into 8-bit halves, accessed by prefixing normal H/L instructions with DD/FD:

Instruction	Bytes	T-states	Notes
LD A,IXH	2	8	Read high byte of IX
LD A,IXL	2	8	Read low byte of IX
LD IXH,A	2	8	Write high byte of IX
LD IXL,n	3	11	Immediate into IX low
ADD A,IXL	2	8	Arithmetic with IX halves
INC IXH	2	8	Increment IX high byte
DEC IXL	2	8	Decrement IX low byte

Demoscene use: Two extra 8-bit registers for counters, accumulators, or small values without touching the main register file. Particularly useful when BC/DE/HL are all occupied as pointers. Cost: 4T more than the equivalent main-register operation.

sjasmpplus syntax: IXH, IXL, IYH, IYL (also accepts HX, LX, HY, LY).

SLL r (Shift Left Logical)

Instruction	Bytes	T-states	Notes
SLL r	2	8	Shift left, bit 0 set to 1 (not 0)
SLL (HL)	2	15	Same for memory

SLL shifts left and sets bit 0 to 1 (unlike SLA which sets bit 0 to 0). Opcode: CB 30+r. Occasionally useful for constructing bit patterns.

sjasmpplus syntax: SLL or SLI or SL1.

OUT (C),0

Instruction	Bytes	T-states	Notes
OUT (C),0	2	12	Output zero to port BC

Opcode ED 71. Outputs zero to the port addressed by BC. On CMOS Z80s (including eZ80), this outputs \$FF instead. **Not portable to Agon Light 2.** On NMOS Z80s (all real Spectrums), it outputs \$00 reliably.

sjasmpplus syntax: OUT (C),0.

CB-Prefix Undocumented Bit Operations on (IX+d)

Instructions like SET b, (IX+d), r simultaneously perform a bit operation on memory at (IX+d) and copy the result into register r. These are 4-byte instructions (DD CB dd op) taking 23T. Occasionally useful but rarely critical.

Flag Effects Cheat Sheet

Knowing which instructions set which flags lets you avoid redundant CP or AND A instructions – a common source of wasted T-states.

Instructions That Set All Arithmetic Flags (S, Z, H, P/V, N, C)

- ADD A, r/n/(HL) – P/V = overflow
- ADC A, r/n/(HL) – P/V = overflow
- SUB r/n/(HL) – P/V = overflow
- SBC A, r/n/(HL) – P/V = overflow
- CP r/n/(HL) – Same flags as SUB but A unchanged
- NEG – P/V = overflow
- ADC HL, rr – P/V = overflow
- SBC HL, rr – P/V = overflow

Instructions That Set Z and S (but NOT Carry)

- INC r / DEC r – C unchanged. **Cannot test carry after INC/DEC.**
- INC (HL) / DEC (HL) – Same
- AND r/n/(HL) – C always 0, H always 1
- OR r/n/(HL) – C always 0, H always 0
- XOR r/n/(HL) – C always 0, H always 0
- IN r, (C) – C unchanged
- BIT b, r/(HL) – Z = complement of tested bit, C unchanged
- All CB-prefix rotates/shifts – Full flag set including C

Instructions That Set ONLY Carry-Related Flags

- ADD HL, rr – H and C only (S, Z, P/V unchanged)
- RLCA / RRCA / RLA / RRA – C, H=0, N=0 only (S, Z, P/V unchanged)
- SCF – C=1, H=0, N=0
- CCF – C inverted, H = old C, N=0

Instructions That Set NO Flags

- LD (all variants)
- INC rr / DEC rr (16-bit inc/dec)
- PUSH / POP (except POP AF restores flags)
- EX / EXX
- DI / EI / HALT / NOP
- JP / JR / DJNZ / CALL / RET / RST
- OUT (n), A / IN A, (n) (the non-CB versions)

Practical Tricks

Test A for zero without CP 0:

and a ; 4T same effect, but also sets H

Test carry after 16-bit INC/DEC: You cannot. INC rr/DEC rr set no flags. To test if a 16-bit register reached zero:

or c ; 4T Z set if BC = 0

Skip CP after SUB: If you already performed SUB r, the flags are set – do not follow it with CP or OR A.

INC/DEC preserve carry: Use INC r/DEC r between multi-precision arithmetic without destroying the carry chain.

Register Architecture

Main Register Set

A	F	Accumulator + Flags
B	C	Counter (B for DJNZ) + general
D	E	General purpose pair
H	L	Primary memory pointer (HL is the "accumulator pair")

Special Registers

SP	Stack pointer (16-bit)
PC	Program counter (16-bit)
IX	Index register (16-bit, DD prefix, +4T penalty)
IY	Index register (16-bit, FD prefix, +4T penalty)
	NOTE: IY is used by the Spectrum ROM interrupt handler.
	Do not use IY unless you have DI or IM2 set up.
I	Interrupt vector page (used in IM2)
R	Refresh counter (7-bit, increments every M1 cycle)

Shadow Registers

A'	F'	Swapped with EX AF,AF'
B'	C'	\
D'	E'	Swapped all three with EXX
H'	L'	/

EXX swaps BC/DE/HL with BC'/DE'/HL' in **4T**. This gives you six extra 8-bit registers (or three extra 16-bit pairs) at virtually zero cost. Common use: keep pointers in the shadow set and swap in/out as needed.

Warning: The Spectrum ROM interrupt handler (IM1) uses IY (it must point to the system variables at \$5C3A or to safe memory). Shadow registers BC'/DE'/HL' and AF' are preserved by the ROM ISR and safe to use with interrupts enabled. If your code repurposes IY, disable interrupts first (DI) or switch to IM2 with your own handler.

Register Pairing for Instructions

Pair	Used by	Notes
BC	DJNZ (B only), OUT (C), r, IN r, (C), block instructions (counter)	B = loop counter, C = port low byte
DE	EX DE,HL, LDI/LDIR (destination), LD (DE),A	Destination pointer for block ops
HL	Nearly everything: LD r, (HL), ADD HL,rr, JP (HL), PUSH/POP, LDI (source)	The universal pointer
AF	PUSH AF/POP AF, EX AF,AF'	A = accumulator, F = flags
SP	PUSH/POP, LD SP,HL, EX (SP),HL	Hijack for data tricks

Common Instruction Sequences

Pixel Address Calculation (Screen Address from Y,X)

Convert screen coordinates to ZX Spectrum video memory address. Input: B = Y (0-191), C = X (0-255). Output: HL = screen byte address, A = bit mask.

```
; Input:  B = Y (0-191), C = X (0-255)
; Output: HL = byte address, A = pixel bit position
; Cost:   ~107 T-states
;
pixel_addr:
    ld    a, b          ; 4T   A = Y
    and   $07           ; 7T   scanline within char cell (Y:2-0)
    or    $40           ; 7T   add screen base ($4000 high byte)
    ld    h, a          ; 4T   H = 010 00 SSS (partial)
    ld    a, b          ; 4T   A = Y again
    rra              ; 4T   \
    rra              ; 4T   | shift right 3
    rra              ; 4T   /
    and   $18           ; 7T   mask Y:4-3 (third bits)
    or    h             ; 4T   H = 010 TT SSS
    ld    h, a          ; 4T
    ld    a, b          ; 4T   A = Y again
    and   $38           ; 7T   mask Y:5-3 (character row within third)
    rlca              ; 4T   \ rotate left 2 to get
    rlca              ; 4T   / Y:5-3 in bits 7-5
    ld    l, a          ; 4T   L = RRR 00000 (partial)
    ld    a, c          ; 4T   A = X
    rra              ; 4T   \
    rra              ; 4T   | X / 8
    rra              ; 4T   /
    and   $1F           ; 7T   mask to 5-bit column
    or    l             ; 4T   combine row and column
    ld    l, a          ; 4T   L = RRR CCCCC
```

DOWN_HL: Move One Pixel Row Down

The most-used graphics primitive on the Spectrum. Common case (within character cell) costs only 20T.

```
; Input:  HL = screen address
; Output: HL = address one row below
; Cost:    20T (common), 46T (third boundary), 77T (char boundary)
;
down_hl:
    inc  h                ; 4T   try next scanline
    ld   a, h             ; 4T
    and  7                ; 7T   crossed character boundary?
    ret  nz               ; 5T   no: done (20T total)

    ld   a, l             ; 4T   yes: advance character row
    add  a, 32            ; 7T   L += 32
    ld   l, a             ; 4T
    ret  c                ; 5T   carry = crossed third (46T total)

    ld   a, h             ; 4T   same third: undo extra H increment
    sub  8                ; 7T
    ld   h, a             ; 4T
    ret                     ; 10T (77T total)
```

8x8 Unsigned Multiply (Shift-and-Add)

From Dark / X-Trade, Spectrum Expert #01 (1997). Used in rotation matrices and coordinate transforms.

```
; Input:  B = multiplicand, C = multiplier
; Output: A:C = 16-bit result (A = high, C = low)
; Cost:    196-204 T-states
;
mulu112:
    ld   a, 0             ; 7T   clear accumulator
    ld   d, 8             ; 7T   8 bits

.loop:
    rr   c                ; 8T   shift multiplier bit into carry
    jr   nc, .noadd       ; 7/12T
    add  a, b             ; 4T   add multiplicand
.noadd:
    rra                   ; 4T   shift result right
    dec  d                ; 4T
    jr   nz, .loop        ; 12T
    ret                     ; 10T
```

AY Register Write

Standard sequence for writing to the AY-3-8910 sound chip on ZX Spectrum 128K.


```

; Input:  A = register number (0-15), E = value
; Cost:   41 T-states (plus CALL/RET overhead)
;
ay_write:
    ld    bc, $FFFD      ; 10T  register select port
    out   (c), a          ; 12T  select register
    ld    b, $BF         ; 7T   data port ($BFFD)
    out   (c), e          ; 12T  write value
    ret                  ; 10T

```

16-Bit Compare (HL vs DE)

The Z80 has no direct 16-bit compare. Use SBC and restore.

```

; Destroys: A (if using the OR method for equality)
;
; For equality only:
    or    a                ; 4T   clear carry
    sbc   hl, de           ; 15T  HL = HL - DE, flags set
    add   hl, de           ; 11T  restore HL
                        ; --- 30T total, Z flag valid

```

Stack-Based Screen Fill

The fastest way to fill the screen with a pattern. See Chapter 3.

```

; Input:  HL = 16-bit fill pattern
; Cost:   ~36,000 T-states (vs ~129,000 with LDIR)
;
fill_screen:
    di                          ; 4T
    ld    (restore_sp + 1), sp  ; 20T  save SP (self-modifying)
    ld    sp, $5800            ; 10T  end of pixel area

    ld    b, 192                ; 7T   192 iterations x 16 pushes x 2 bytes =
    ⇨ 6144

.loop:
    REPT 16
        push hl                  ; 11T  x 16 = 176T
    ENDR
    djnz .loop                  ; 13T/8T

restore_sp:
    ld    sp, $0000            ; 10T  self-modified
    ei                          ; 4T
    ret                        ; 10T

```

Fast Pixel-Row Iteration (Split Counters)

From Introspec's DOWN_HL analysis (Hype, 2020). Eliminates all conditional branching from the inner loop. Total cost for 192 rows: 2,343T vs 5,922T for naive DOWN_HL calls.

```

; HL starts at $4000
;
iterate_screen:
    ld    hl, $4000        ; 10T
    ld    c, 3             ; 7T   three thirds

.third:
    ld    b, 8             ; 7T   eight character rows per third

.char_row:
    push hl                ; 11T   save char row start

    REPT 7
        ; ... process row using HL ...
        inc h              ; 4T   next scanline (NO branching)
    ENDR
    ; ... process 8th row ...

    pop  hl                ; 10T   restore char row start
    ld    a, l             ; 4T
    add   a, 32            ; 7T   next character row
    ld    l, a             ; 4T
    djnz .char_row         ; 13T/8T

    ld    a, h             ; 4T
    add   a, 8             ; 7T   next third
    ld    h, a             ; 4T
    dec   c               ; 4T
    jr    nz, .third       ; 12T/7T

```

Quick Cost Comparisons

For inner loop decisions, these comparisons matter most:

Operation	Slow way	Fast way	Savings
Zero A	LD A,0 (7T, 2B)	XOR A (4T, 1B)	3T, 1B
Test A=0	CP 0 (7T, 2B)	OR A (4T, 1B)	3T, 1B
Copy 1 byte (HL)→(DE)	LD A, (HL)+LD (DE), A+INC HL+INC DE (26T, 4B)	LDI (16T, 2B)	10T, 2B per byte
Copy N bytes	LDIR (21T/byte)	N x LDI (16T/byte)	24% faster, costs 2N bytes of code
Fill 2 bytes	LD (HL), A+INC HL x2 (26T)	PUSH rr (11T)	58% faster, needs SP hijack
8-bit loop	DEC B+JR NZ (16T, 3B)	DJNZ (13T, 2B)	3T, 1B per iteration
Indirect call	CALL nn (17T, 3B)	RET via render list (10T, 1B)	7T, 2B per dispatch

Operation	Slow way	Fast way	Savings
Register swap	LD A,H+LD H,D+LD D,A (12T, 3B)	EX DE,HL (4T, 1B)	8T, 2B
Save 6 registers	3 x PUSH (33T, 3B)	EXX (4T, 1B)	29T, 2B

Instruction Encoding Size Reference

For size-coding and estimating code density:

Prefix	Instructions	Extra bytes	Extra T-states
None	Most 8-bit ops, LD, ADD, INC, PUSH, POP, JP, JR	0	0
CB	Bit ops, shifts, rotates on registers	+1	+4 typically
ED	Block ops, 16-bit ADC/SBC, IN/OUT (C), LD rr,(nn)	+1	varies
DD	IX-indexed operations	+1	+4 to +8
FD	IY-indexed operations	+1	+4 to +8
DD CB	Bit/shift/rotate on (IX+d)	+2	+8 to +12

Size-coding tip: Avoid IX/IY-indexed instructions when possible. LD A,(IX+5) is 3 bytes/19T. LD L,5 / LD A,(HL) is 3 bytes/11T if H already holds the page. The index registers are convenient but expensive.

Sources: Zilog Z80 CPU User Manual (UM0080); Sean Young, “The Undocumented Z80 Documented” (2005); Dark / X-Trade, “Programming Algorithms” (Spectrum Expert #01, 1997); Introspec, “Once more about DOWN_HL” (Hype, 2020); Chapter 1 (timing harness); Chapter 3 (toolbox patterns); Chapter 4 (multiply, division)

Appendix B: Sine Table Generation and Trigonometric Tables

“Cosine is just sine offset by a quarter period.” – Raider’s Commandments

Every demo effect that curves – rotation, plasma, scrolling, tunnels – needs a sine table. On the Z80, you pre-compute the values into a lookup table and index by angle. The question is how to store and access that table as efficiently as possible.

This appendix compares eight approaches to sine table storage, ranging from the obvious (256-byte table) to the exotic (2-bit second-order delta encoding). The data comes from `verify/sine_compare.py`, which you can run to reproduce every number here.

The Standard Format

A demoscene sine table has **256 entries**, indexed by angle:

Index	Angle
0	0°
64	90°
128	180°
192	270°
256 (wraps to 0)	360°

Each entry is a **signed byte** (-128 to +127), representing -1.0 to approximately +1.0. The power-of-two period means the angle index wraps naturally with 8-bit overflow, and cosine becomes sine by adding 64:

```
ld    h, high(sin_table) ; 7T    table must be 256-byte aligned
ld    l, a                ; 4T    A = angle (0-255)
ld    a, (hl)             ; 7T    A = sin(angle)
                        ; --- 18 T-states total

; cos(angle) -- offset by quarter period
add   a, 64               ; 7T    cos = sin + 90°
```

```
ld    l, a           ; 4T
ld    a, (hl)        ; 7T
```

This is Raider's key rule: store H once with the table's high byte, then L is the angle and wraps freely.

Comparison of Approaches

#	Approach	Data	Code	Total	RAM	Max Error	RMS
1	Full table (256 bytes)	256	0	256	0	0	0.00
2	Quarter-wave table	65	21	86	0	0	0.00
3	Parabolic approximation	0	38	38	0	8	4.51
4	Quarter-wave + 2nd-order deltas	18	45	63	64	0	0.00
5	Bhaskara I approximation	0	~60	~60	0	1	0.49
5b	Bhaskara I + correction bitmap	1	~80	~81	0	0	0.00
6	Quarter-wave + packed 4-bit deltas	33	43	76	64	0	0.00
7	Full 2nd-order deltas, 2-bit packed	66	30	96	256	0	0.00

The approaches fall into three categories:

- **Lookup-based** (no RAM needed): full table, quarter-wave table
- **Generation-based** (needs RAM buffer at startup): delta and second-order delta approaches
- **Approximate** (no table at all): parabolic, Bhaskara I
- **Approximate + exact correction**: Bhaskara I with correction bitmap

Approach 1: Full 256-Byte Table

The simplest and fastest. Pre-compute all 256 values and embed them as data.

```
ld    h, high(sin_table)
ld    l, a
ld    a, (hl)
```

Cost: 256 bytes of ROM. **Speed:** 18 T-states per lookup. **When to use:** Always, unless you are size-coding. On a 48K Spectrum with ~40K free, 256 bytes is nothing. This is the default choice.

Approach 2: Quarter-Wave Table

A sine wave has four-fold symmetry. The first quadrant (0° to 90°, indices 0 to 64) contains all the information:

- **Second quadrant** (65-128): mirror of first quadrant. $\sin(128 - i) = \sin(i)$.
- **Third quadrant** (129-192): negative of first quadrant. $\sin(128 + i) = -\sin(i)$.
- **Fourth quadrant** (193-255): negative mirror. $\sin(256 - i) = -\sin(i)$.

Store only 65 bytes (indices 0 through 64 inclusive), then reconstruct:

```
; Input:  A = angle (0-255)
; Output: A = sin(angle), signed byte
; Uses:   HL, BC
; Table:  sin_quarter (65 bytes, 256-byte aligned)
;
qsin:
    ld    c, a                ; 4T   save original angle
    and   $7F                 ; 7T   fold to 0-127 (first half)
    cp    65                  ; 7T   past the peak?
    jr    c, .no_mirror       ; 12/7T
    ; Mirror: index = 128 - index
    neg                     ; 8T
    add   a, 128              ; 7T
.no_mirror:
    ld    h, high(sin_quarter) ; 7T
    ld    l, a                ; 4T
    ld    a, (hl)             ; 7T   A = |sin(angle)|
    bit   7, c                ; 8T   was original angle >= 128?
    ret   z                   ; 11/5T no: positive half, done
    neg                     ; 8T   yes: negate for third/fourth quadrant
    ret                     ; 10T
```

Cost: 65 bytes data + ~21 bytes code = **86 bytes total**. **Speed:** ~50-70 T-states per lookup (varies by quadrant). **Error:** Zero. **When to use:** Size-limited demos (256-byte, 512-byte intros) where you need exact values but cannot spare 256 bytes.

Approach 3: Parabolic Approximation (Dark's Method)

From Dark / X-Trade, *Spectrum Expert* #01 (1997). The idea: half a period of cosine looks like a parabola. The approximation $y \approx 1 - 2(x/\pi)^2$ matches closely. In integer terms, each half-period is generated as a piecewise quadratic.

Pure code, zero data. The generation loop needs an 8×8 multiply and some accumulator logic – approximately 38 bytes.

The error is bounded: **maximum absolute error = 8** (out of a 256-step range), or about **6.3%** of full scale. The RMS error is 4.51.

Here is where the parabola diverges from true sine (first quadrant):

Index	True	Para	Diff
0	0	0	+0
4	12	15	-3
8	25	30	-5
12	37	43	-6
16	49	56	-7
20	60	67	-7
24	71	77	-6
28	81	87	-6
32	88	93	-5

The parabola is consistently “ahead” – it rises faster near zero and is flatter near the peak. Maximum divergence: **8 units at index 17** (about 24°).

When to use: Extreme size-coding (64-byte intros, tight loaders). Plasmas, simple scrollers, and wobble effects where the eye does not distinguish exact curvature from approximate. Not suitable for smooth rotation or precise wireframe 3D.

Approach 4: Second-Order Delta Encoding (The Deep Trick)

This is the most mathematically interesting approach, and by far the most compact exact representation.

The Key Insight

The second derivative of $\sin(x)$ is $-\sin(x)$. At 8-bit integer precision, quantising to signed bytes, the second finite difference of the sine table has a remarkable property: **every value is exactly -1, 0, or +1**.

True sine:	[0, 3, 6, 9, 12, 16, 19, 22, 25, ...]
First diff:	[3, 3, 3, 3, 4, 3, 3, 3, 3, ...]
Second diff:	[0, 0, 0, 1, -1, 0, 0, 0, 0, ...]

Three values. Two bits per entry. This is not an approximation – it is exact. The mathematical reason: $d^2(\sin)/dx^2$ is a smooth function with small magnitude, and at 256 entries per period with 8-bit amplitude, the discrete second derivative never exceeds ± 1 .

Full Table via 2-Bit Second-Order Deltas

Store: initial value (1 byte), initial delta (1 byte), then 254 second-order deltas packed at 2 bits each (64 bytes). **Total: 66 bytes data + ~30 bytes decode code = 96 bytes.** Needs 256 bytes of RAM to decode into.

Quarter-Wave via 2-Bit Second-Order Deltas

Combine with quarter-wave symmetry: store only the first 64 second-order deltas. **Total: 18 bytes data + ~45 bytes decode code = 63 bytes.** Needs 64 bytes of RAM.

This is the **smallest exact representation**: 63 bytes total for a perfect 256-entry sine table.

```
; sin_d2_data: 16 bytes of packed 2-bit deltas (64 entries)
; sin_buffer: 64 bytes RAM for decoded quarter-wave
;
decode_quarter_d2:
    ld    hl, sin_buffer      ; destination
    ld    de, sin_d2_data    ; source (packed d2 values)
    xor    a
    ld    (hl), a            ; sin[0] = 0
    inc    hl
    ld    b, a                ; b = current delta (starts at 0)
    ld    c, 63               ; 63 more entries to decode

.loop:
    ; Unpack 2-bit d2 value
    ; 00 = 0, 01 = +1, 11 = -1 (10 unused)
    rr    (de)                ; shift out 2 bits
    rr    (de)
    ; ... (bit extraction logic)

    ; Apply: delta += d2, value += delta
    add    a, b                ; new delta
    ld    b, a
    ld    a, (hl-1)           ; previous value (pseudocode)
    add    a, b
    ld    (hl), a
    inc    hl
    dec    c
    jr    nz, .loop

    ; Now use qsin() lookup on sin_buffer
```

The decode runs once at startup. After that, use the quarter-wave lookup routine from Approach 2 on the decoded buffer.

When to use: Size-coded demos (128-byte, 256-byte intros) where you need exact values, can afford 64 bytes of RAM, and have a brief startup phase. The decode loop runs in under 2,000 T-states – invisible.

Sidebar: Why Not 1 Bit Per Delta?

An intuitive objection: the quarter-wave sine (0° to 90°) is monotonically increasing. First differences $d1$ are always non-negative. In continuous maths, the second derivative of sine in the first quadrant is always negative (the curve is concave). So $d2$ should be ≤ 0 , meaning we only need $\{-1, 0\}$ – a single bit per entry.

The intuition is right for continuous sine, but wrong for quantised integer sine. At 8-bit precision, rounding creates occasional upward bumps in $d1$:

$d1$: 3, 3, 3, 3, 4, 3, 3, ... (that 4 is a rounding correction)
 $d2$: 0, 0, 0, +1, -1, 0, ... (the +1 is load-bearing)

There are 12 such +1 entries out of 63. If you suppress them (cap $d1$ to be monotonically non-increasing), the errors *accumulate*: by index 64 the peak reaches only 108 instead of 127 – a max error of 19, worse than the parabolic approximation. Those +1 corrections carry precisely the information needed to hit the right integer values. You cannot drop them.

A variable-length prefix code ($0 \rightarrow 1$ bit, $\pm 1 \rightarrow 2$ bits) saves 4 bytes of data over fixed 2-bit encoding but costs ~ 15 extra bytes of Z80 decode logic. Net loss. The 2-bit fixed encoding is the practical optimum.

Sidebar: Why Parabolic + Correction Doesn't Help

Another intuitive idea: generate a parabolic approximation (38 bytes of code, max error 8), then store a small correction table to fix it to exact values. The corrections range from -8 to +8, so they should compress well.

The corrections *do* compress well – their first differences are exactly $\{-1, 0, +1\}$, packing into 2 bits per entry. But this is no coincidence. A parabola is a quadratic with constant second derivative. So:

- $d2(\sin) \in \{-1, 0, +1\}$ – the sine's second derivative at integer precision
- $d2(\text{para}) \in \{-1, 0, +1\}$ – the parabola's second derivative (nearly constant)
- $d1(\text{correction}) = d1(\sin) - d1(\text{para}) \in \{-1, 0, +1\}$ – **same entropy**

The correction deltas have *exactly the same structure* as the direct sine $d2$. But the parabolic route adds 38 bytes of generation code, plus ~ 20 bytes to apply corrections. Total: ~ 96 bytes vs 63 bytes for direct $d2$ encoding.

The parabola removes the smooth (low-frequency) component of sine – but the 2-bit $d2$ encoding already handles smooth data perfectly. There is nothing left for the parabola to contribute that $d2$ doesn't already capture. The generation code is pure overhead.

Approach 5: Bhaskara I Approximation (7th Century)

The most surprising entry in our comparison comes from 7th-century Indian mathematician Bhaskara I. His rational approximation to sine, published around 629 CE,

achieves **max error of only 1 unit** at 8-bit precision – dramatically better than the parabolic approximation (max error 8) and nearly exact.

The Formula

For angle x in radians (0 to π):

$$\sin(x) \approx 16x(\pi - x) / (5\pi^2 - 4x(\pi - x))$$

In our integer domain (angle 0-64 for the first quadrant, amplitude 0-127):

$$\begin{aligned}\sin(i) &\approx 127 \times 16i(64 - i) / (5 \times 64^2 - 4 \times i(64 - i)) \\ &= 127 \times 16i(64 - i) / (20480 - 4i(64 - i))\end{aligned}$$

The formula is a ratio of two quadratics. On the Z80 this needs an 8×8 multiply and a 16-bit division – routines that many demos already include for 3D projection or texture mapping.

Accuracy

Across the 65 entries of the first quadrant, Bhaskara I matches the exact integer sine everywhere except **8 positions** (out of 65), where it is off by exactly ± 1 :

Index	True	Bhaskara	Diff
4	12	13	-1
17	51	52	-1
28	81	80	+1
31	88	87	+1
40	106	105	+1
43	111	110	+1
50	120	119	+1
52	122	121	+1

Only 8 positions differ, all by exactly ± 1 . The errors are split: 2 entries where Bhaskara overshoots (near the start), 6 where it undershoots (near the peak). Eight corrections total, which encode as a single byte bitmap.

Z80 Implementation

The implementation requires: - An 8×8→16 multiply routine (~20 bytes, likely already available) - A 16÷16→16 divide routine (~30 bytes, likely already available) - The Bhaskara wrapper itself (~25 bytes) - Quarter-wave folding logic (~15 bytes, shared with Approach 2)

If your demo already has multiply and divide routines, the marginal cost is roughly **25 bytes** for a sine function with max error 1.

If you need the routines from scratch, total is approximately **60 bytes** of code with zero data bytes. This is competitive with the d2 delta approach (63 bytes) but requires no RAM buffer and no startup decode phase. The tradeoff: 1 unit of error vs perfect accuracy.

Bhaskara I + Correction Bitmap (Exact)

To eliminate that last unit of error, store the 8 correction positions as a bitmap. Since the corrections are symmetric (first 4 need +1, last 4 need -1), one byte suffices:

```

push af
ld  a, c
; Look up correction from bitmap (8 specific indices)
; ... (~20 bytes of correction logic)
pop af
add a, correction      ; ±1 or 0

```

Total: ~80 bytes code + 1 byte data = **~81 bytes**, zero RAM, zero startup, exact values. More expensive than d2 deltas (63B) but avoids the RAM buffer and startup decode.

When to Use Bhaskara I

- **You already have multiply/divide routines:** ~25 bytes extra, max error 1. Hard to beat.
- **No RAM available for decode buffer:** Unlike d2 deltas, Bhaskara computes on the fly.
- **Real-time generation needed:** Each value is computed independently - no sequential dependency, so you can compute sin(any angle) without decoding a table first.
- **±1 error is acceptable:** For scrollers, plasmas, and most visual effects, the difference between max error 1 and max error 0 is literally invisible.

Historical Note: Bhaskara I's formula predates European trigonometric tables by nearly a millennium. That a 7th-century rational approximation achieves max error 1 on a 1980s 8-bit processor is a beautiful collision of mathematical elegance and engineering constraints. The formula was published in *Mahabhaskariya* (629 CE), a commentary on Aryabhata's astronomical methods.

Practical Recommendations

Every generation-based approach produces a lookup table at startup. After that, the runtime cost is identical: LD H, high(table) / LD L, A / LD A, (HL) = **18 T-states** for a 256-byte table, or the quarter-wave folding routine at **50-70 T-states** for a 64-byte buffer. The "ROM cost" column below is what matters for size-coding - it is the total bytes your approach occupies in the binary.

Use Case	Approach	ROM Cost	RAM	Init	Lookup	Error
Normal demo / game	Full 256-byte table	256B	0	none	18 T	exact
512-byte intro	Quarter-wave table	86B	0	none	50-70 T	exact

Use Case	Approach	ROM Cost	RAM	Init	Lookup	Error
256-byte intro	Quarter + d2 deltas	63B	64B	~2K T	50-70 T	exact
Has multiply/divide	Bhaskara I (generate to LUT)	~25B extra	256B	~80K T	18 T	± 1 max
128-byte intro	Parabolic (generate to LUT)	38B	256B	~10K T	18 T	± 8 max

The Decision Tree

1. **Do you have 256 bytes to spare?** Use the full table. Do not overthink it. LD L,A / LD A,(HL) at 18 T-states cannot be beaten.
2. **Size-limited but need accuracy?** Quarter-wave table at 86 bytes. No RAM needed, no startup phase. Lookup is 50-70 T-states (the folding logic).
3. **Extreme size limit, exact values needed?** Quarter-wave + second-order delta decoding at 63 bytes. Decode once at startup into a 64-byte quarter buffer, then use the same folding lookup.
4. **Already have multiply/divide?** Bhaskara I at ~25 bytes extra. Generate a full 256-byte LUT at startup, then enjoy 18 T-state lookups with max error 1.
5. **Extreme size limit, approximate OK?** Parabolic at 38 bytes, zero data. Generate to a 256-byte LUT at startup. Max error 8, good for plasmas and wobbles.

What Does Not Work

- **Parabolic + correction table** (123 bytes exact): worse than just using a quarter-wave table (86 bytes). The overhead of computing the parabola *and* looking up a correction defeats the purpose.
- **Delta + RLE** (100-219 bytes): sine deltas vary smoothly rather than repeating in runs. RLE is designed for data with long constant runs – sine is the wrong shape for it.
- **Full delta-encoded table** (152-271 bytes): uses *more* total bytes than the raw 256-byte table. Delta encoding only helps when deltas are significantly smaller than the original values; sine deltas are already bounded to ± 4 , but you still need 256 of them.

Raider's Commandments

In the Hype comments on Introspec's analysis of *Illusion*, veteran coder Raider distilled decades of collective wisdom into informal "commandments" for sine table design:

1. **256 entries per full period.** The angle index wraps with 8-bit overflow. No modular arithmetic needed.

2. **Signed bytes: -128 to +127.** Matches Z80 signed arithmetic.
3. **Page-align the table.** Place it at a 256-byte boundary so H is constant. LD H,high(table) once, then LD L,angle / LD A,(HL) forever.
4. **Cosine is sine + 64.** One ADD A,64 instruction.
5. **Sine of (angle + 128) = -sine(angle).** NEG flips the sign. Use this for phase shifts.
6. **Do not compute sine at runtime** unless you are size-coding. A table lookup is always faster.
7. **Keep the amplitude as a power of two** (64, 127, 128) so multiplication is a shift.
8. **Quarter-wave symmetry** saves 75% storage when every byte matters.
9. **Test at the boundaries.** Index 0 should be exactly 0. Index 64 should be the maximum positive value (+127). Index 128 should be exactly 0. Index 192 should be the maximum negative value (-128 or -127, depending on your convention).

These rules reflect decades of experience. Follow them and your sine tables will be fast, small, and correct.

Reference: The Full 256-Byte Table

For convenience, here is the standard sine table (256 entries, signed, period = 256, amplitude ± 127):

```
; sin(0) = 0, sin(64) = +127, sin(128) = 0, sin(192) = -128
;
```

```
ALIGN 256
sin_table:
    DB  0,  3,  6,  9, 12, 16, 19, 22
    DB 25, 28, 31, 34, 37, 40, 43, 46
    DB 49, 51, 54, 57, 60, 63, 65, 68
    DB 71, 73, 76, 78, 81, 83, 85, 88
    DB 90, 92, 94, 96, 98, 100, 102, 104
    DB 106, 108, 109, 111, 112, 114, 115, 117
    DB 118, 119, 120, 121, 122, 123, 124, 124
    DB 125, 126, 126, 127, 127, 127, 127, 127
    DB 127, 127, 127, 127, 127, 127, 126, 126
    DB 125, 124, 124, 123, 122, 121, 120, 119
    DB 118, 117, 115, 114, 112, 111, 109, 108
    DB 106, 104, 102, 100, 98, 96, 94, 92
    DB 90, 88, 85, 83, 81, 78, 76, 73
    DB 71, 68, 65, 63, 60, 57, 54, 51
    DB 49, 46, 43, 40, 37, 34, 31, 28
    DB 25, 22, 19, 16, 12, 9, 6, 3
    DB 0, -3, -6, -9, -12, -16, -19, -22
    DB -25, -28, -31, -34, -37, -40, -43, -46
    DB -49, -51, -54, -57, -60, -63, -65, -68
    DB -71, -73, -76, -78, -81, -83, -85, -88
    DB -90, -92, -94, -96, -98, -100, -102, -104
    DB -106, -108, -109, -111, -112, -114, -115, -117
```

```

DB -118, -119, -120, -121, -122, -123, -124, -124
DB -125, -126, -126, -127, -127, -127, -127, -127
DB -128, -127, -127, -127, -127, -127, -126, -126
DB -125, -124, -124, -123, -122, -121, -120, -119
DB -118, -117, -115, -114, -112, -111, -109, -108
DB -106, -104, -102, -100, -98, -96, -94, -92
DB -90, -88, -85, -83, -81, -78, -76, -73
DB -71, -68, -65, -63, -60, -57, -54, -51
DB -49, -46, -43, -40, -37, -34, -31, -28
DB -25, -22, -19, -16, -12, -9, -6, -3

```

Copy, paste, assemble, use.

Sources: Dark / X-Trade “Programming Algorithms” (Spectrum Expert #01, 1997) for the parabolic approximation; Bhaskara I, *Mahabhaskariya* (629 CE) for the rational approximation; Raider (Hype comments, 2017) for sine table design principles; `verify/sine_compare.py` for comparative analysis

Appendix C: Compression Quick Reference

“The question is not whether to compress — it is which compressor to use, and when.” – Chapter 14

This appendix is a tear-out reference card for data compression on the ZX Spectrum. Chapter 14 covers the theory, the benchmark data, and the reasoning behind each recommendation. This appendix distils it into lookup tables and decision rules you can pin above your monitor.

All numbers come from Introspec’s 2017 benchmark (“Data Compression for Modern Z80 Coding,” Hype) unless otherwise noted. The test corpus was 1,233,995 bytes of mixed data: Calgary/Canterbury academic benchmarks, 30 ZX Spectrum graphics, 24 music files, and miscellaneous demo data.

Compressor Comparison Table

Com-pressor	Au-thor	Compressed (bytes)	Ra-tio	Decompres-sor Size	Speed (T/byte)	Back-wards	Notes
Ex-omizer 2	Mag-nus Lind	596,161	48.3%	~170 bytes	~250	Yes	Best ra-tio. Slow to de-com-press.
ApLib	Joer-gen Ib-sen	606,833	49.2%	~199 bytes	~105	No	Good all-rounder.

Com-pressor	Au-thor	Compressed (bytes)	Ra-tio	Decompres-sor Size	Speed (T/byte)	Back-wards	Notes
Hrust 1	Alone Coder	613,602	49.7%	~150 bytes	~120	Yes	Re-locatable stack de-packer. Pop-ular in Rus-sian scene.
Pu-Crunch	Pasi Ojala	616,855	50.0%	~200 bytes	~140	No	Orig-inally for C64.
Pletter 5	XL2S	635,797	51.5%	~120 bytes	~69	No	Fast + de-cent ra-tio.
MegaLZ	LVD / Intro-spec	636,910	51.6%	92 bytes (compact)	~98 (compact)	No	Opti-mal parser. Re-vived 2019 with new de-com-pres-sors.
MegaLZ fast	LVD / Intro-spec	636,910	51.6%	234 bytes	~63	No	Fastest MegaLZ vari-ant. Faster than 3x LDIR.

Com- pressor	Au- thor	Compressed (bytes)	Ra- tio	Decompres- sor Size	Speed (T/byte)	Back- wards	Notes
ZX0	Einar Saukas	~642,000*	~52%	~70 bytes	~100	Yes	Suc- ces- sor to ZX7. Opti- mal parser. Mod- ern de- fault.
ZX7	Einar Saukas	653,879	53.0%	69 bytes	~107	Yes	Tiny de- com- pres- sor. The clas- sic size- coder tool.
Bit- buster	Team Bomba	~660,000*	~53.5%	~90 bytes	~80	No	Sim- ple. Good for first projects.
LZ4	Yann Col- let (Z80 port)	722,522	58.6%	~100 bytes	~34	No	Fastest de- com- pres- sion. Byte- aligned to- kens.

Com- pressor	Au- thor	Compressed (bytes)	Ra- tio	Decompres- sor Size	Speed (T/byte)	Back- wards	Notes
Hrum	Hrumer	~642,000*	~52%	~130 bytes	~110	No	Pop- ular in Rus- sian scene. De- clared ob- so- lete by In- tro- spec. ZX0 vari- ant. Slightly bet- ter ra- tio, slightly larger de- com- pres- sor.
ZX1	Einar Saukas	—	~51%	~80 bytes	~90	Yes	Used in RED RE- DUX 256b in- tro (2025). Best ZXn ra- tio.
ZX2	Einar Saukas	—	~50%	~100 bytes	~85	Yes	

* Approximate. ZX0, Bitbuster, and Hrum were not in the original 2017 benchmark; values are estimated from independent tests on similar corpora.

Reading the table:

- **Ratio** = compressed size / original size. Lower is better.

- **Speed** = T-states per output byte during decompression. Lower is faster.
 - **Decompressor Size** = bytes of Z80 code needed for the decompression routine. Lower is better for size-coded intros.
 - **Backwards** = supports decompressing from end to start, allowing in-place decompression when source and destination overlap.
-

Decision Tree: Which Compressor?

Follow top-to-bottom. Take the first branch that matches your situation.

START

```

|
+-- Is this a 256-byte or 512-byte intro?
|   YES --> ZX0 (70-byte decompressor) or custom RLE (<30 bytes)
|
+-- Is this a 1K or 4K intro?
|   YES --> ZX0 (best ratio-to-decompressor-size)
|
+-- Do you need real-time streaming (decompress during playback)?
|   YES --> LZ4 (~34 T/byte = 2+ KB per frame at 50fps)
|
+-- Do you need fast decompression between scenes?
|   YES --> MegaLZ fast (~63 T/byte) or Pletter 5 (~69 T/byte)
|
+-- Is decompression speed irrelevant (one-time load at startup)?
|   YES --> Exomizer (48.3% ratio, nothing beats it)
|
+-- Need a good balance of ratio and speed?
|   YES --> ApLib (~105 T/byte, 49.2% ratio)
|
+-- Is the data mostly runs of identical bytes?
|   YES --> Custom RLE (decompressor < 30 bytes, trivial)
|
+-- Is the data sequential animation frames?
|   YES --> Delta-encode first, then compress with ZX0 or LZ4
|
+-- First project, want something simple?
|   YES --> Bitbuster or ZX0 (both well-documented, easy to integrate)

```

Compressibility of Common ZX Spectrum Data Types

How well different data types compress, and tricks to improve the ratio.

Data Type	Raw Size	Typical ZX0 Ratio	Typical Exomizer Ratio	Notes
Screen pixels (\$4000-\$57FF)	6,144 bytes	40-60%	35-55%	De- pends on image com- plexity. Black back- grounds com- press well.
Attributes (\$5800-\$5AFF)	768 bytes	30-50%	25-45%	Often highly repeti- tive. Solid- colour areas com- press to almost noth- ing.
Full screen (pixels + attrs)	6,912 bytes	40-58%	35-52%	Com- press pixels and at- tributes sepa- rately for 5-10% better ratio.
Sine/co-sine tables	256 bytes	60-75%	55-70%	Smooth curves com- press well. Con- sider genera- tion instead (Ap- pendix B).

Data Type	Raw Size	Typical ZX0 Ratio	Typical Exomizer Ratio	Notes
Tile data (8x8 tiles)	varies	35-55%	30-50%	Re-order tiles by similarity for better ratio. Mask bytes hurt ratio. Store masks separately. Pattern data is repetitive. Empty rows compress well. Highly repetitive between frames. Delta-encode first. Random-looking data compresses poorly. Pre-sort if possible.
Sprite data	varies	45-65%	40-60%	
PT3 music data	varies	40-55%	35-50%	
AY register dumps	varies	30-50%	25-45%	
Lookup tables (arbitrary)	varies	50-80%	45-75%	

Data Type	Raw Size	Typical ZX0 Ratio	Typical Exomizer Ratio	Notes
Font data (96 chars x 8 bytes)	768 bytes	55-70%	50-65%	Lots of zero bytes (de- scen- ders, thin strokes).

Pre-compression tricks

These techniques improve compression ratio by restructuring data before feeding it to the compressor.

Separate pixels from attributes. A full 6,912-byte screen stored as one block forces the compressor to handle a transition from pixel data to attribute data at byte 6,144. Compress the 6,144-byte pixel block and the 768-byte attribute block separately. The attribute block, being highly repetitive, often compresses to under 200 bytes.

Delta-encode animation frames. Store the first frame in full. For each subsequent frame, store only the bytes that differ from the previous frame as (offset, value) pairs. Apply LZ compression to the delta stream. psndcj compressed 122 frames (843,264 bytes raw) into 10,512 bytes using this technique in Break Space.

Reorder data for locality. Tile maps stored in row-major order may compress better if reordered so that similar tiles are adjacent. Sort sprite frames by visual similarity. Group repeated sub-patterns together.

Store constants separately. If a data block contains a repeated header or footer (e.g., tile metadata), factor it out and store it once. Only compress the variable portion.

Interleave planes. For multicolour or masked sprites, storing all mask bytes together and all pixel bytes together often compresses better than interleaving mask-pixel-mask-pixel per row.

Minimal RLE Decompressor

The simplest useful compressor. Only 12 bytes of code. Suitable for 256-byte intros or data with long runs of identical bytes. See Chapter 14 for a full discussion.

```
; Format: [count][value] pairs, terminated by count = 0
; HL = source (compressed data)
; DE = destination (output buffer)
; Destroys: AF, BC
rle_decompress:
    ld     a, (hl)      ; read count      7T
    inc    hl           ;                  6T
    or     a            ; count = 0?      4T
```

```

        ret      z           ; yes: done           5T/11T
        ld      b, a         ; B = count           4T
        ld      a, (hl)      ; read value          7T
        inc     hl           ;                     6T
.fill:  ld      (de), a      ; write value          7T
        inc     de           ;                     6T
        djnz    .fill        ; loop B times        13T/8T
        jr      rle_decompress ; next pair         12T
; Total: 12 bytes of code
; Speed: ~26 T-states per output byte (within long runs)
;       + 46T overhead per [count, value] pair

```

Encoding tool (Python one-liner for simple RLE):

```

out = bytearray()
i = 0
while i < len(data):
    val = data[i]
    count = 1
    while i + count < len(data) and data[i + count] == val and count < 255:
        count += 1
    out.extend([count, val])
    i += count
out.extend([0]) # terminator
return out

```

This naive RLE expands data with no runs (worst case: 2 bytes per 1 byte of input). For mixed data, use escape-byte RLE: a special byte signals a run, and all other bytes are literals. Or just use ZX0.

Transposition trick. RLE benefits dramatically from column-major data layout. If you have a 32×24 attribute block where each row varies but columns are often constant, transposing the data (storing all column 0 values, then column 1, etc.) creates long runs that RLE compresses well. The trade-off: the Z80 must untranspose the data after decompression, which costs an extra pass (~13 T-states per byte for a simple nested-loop copy). Count the total cost (decompressor code + un-transpose code + compressed data) against ZX0 (decompressor + compressed data, no transform needed) to see which wins for your specific data.

ZX0 Standard Decompressor (Z80)

The full standard forward decompressor by Einar Saukas. Approximately 70 bytes. This is the version you will use in most projects.

```

; (c) Einar Saukas, based on Wikipedia description of LZ format
; HL = source (compressed data)
; DE = destination (output buffer)
; Destroys: AF, BC, DE, HL
dzx0_standard:
    ld      bc, $ffff      ; initial offset = -1
    push   bc              ; store offset on stack
    inc     bc              ; BC = 0 (literal length will be read)

```

```

        ld      a, $80          ; init bit buffer with end marker
dzx0s_literals:
        call    dzx0s_elias     ; read number of literals
        ldir    ; copy literals from source to dest
        add     a, a            ; read next bit: 0 = last offset, 1 = new offset
        jr      c, dzx0s_new_offset
        ; reuse last offset
        call    dzx0s_elias     ; read match length
dzx0s_copy:
        ex      (sp), hl       ; swap: HL = offset, stack = source
        push    hl             ; put offset back on stack
        add     hl, de          ; HL = dest + offset = match source address
        ldir    ; copy match
        add     a, a            ; read next bit: 0 = literal, 1 = match/offset
        jr      nc, dzx0s_literals
        ; new offset
dzx0s_new_offset:
        call    dzx0s_elias     ; read offset MSB (high bits)
        ex      af, af'        ; save bit buffer
        dec     b               ; B = $FF (offset is negative)
        rl      c               ; C = offset MSB * 2 + carry
        inc     c               ; adjust
        jr      z, dzx0s_done    ; offset = 256 means end of stream
        ld      a, (hl)         ; read offset LSB
        inc     hl
        rra                     ; LSB bit 0 -> carry = length bit
        push    bc              ; save offset MSB
        ld      b, 0
        ld      c, a            ; C = offset LSB >> 1
        pop     af              ; A = offset MSB (from push bc)
        ld      b, a            ; BC = full offset (negative)
        ex      (sp), hl       ; store offset, retrieve source
        push    bc              ; store offset again
        ld      bc, 1           ; minimum match length = 1
        jr      nc, dzx0s_copy  ; if carry clear: length = 1
        call    dzx0s_elias     ; otherwise read match length
        inc     bc              ; +1
        jr      dzx0s_copy
dzx0s_done:
        pop     hl              ; clean stack
        ex      af, af'        ; restore flags
        ret
; Elias interlaced code reader
dzx0s_elias:
        inc     c               ; C starts at 1
dzx0s_elias_loop:
        add     a, a            ; read bit
        jr      nz, dzx0s_elias_nz
        ld      a, (hl)         ; refill bit buffer
        inc     hl
        rla                     ; shift in carry

```



```

dzx0s_elias_nz:
    ret      nc          ; stop bit (0) = done
    add     a, a          ; read data bit
    jr      nz, dzx0s_elias_nz2
    ld      a, (hl)       ; refill
    inc     hl
    rla
dzx0s_elias_nz2:
    rl      c             ; shift bit into C
    rl      b             ; and into B
    jr      dzx0s_elias_loop

```

Usage:

```

ld      de, $4000          ; destination (e.g., screen)
call    dzx0_standard      ; decompress

```

Backwards variant. ZX0 also provides a backwards decompressor (`dzx0_standard_back`) that reads compressed data from end to start and writes output from end to start. This enables in-place decompression: place the compressed data at the end of the destination buffer, and decompress backwards so the output overwrites the compressed data only after it has been read. Essential when RAM is tight.

Integration Patterns

Pattern 1: Decompress to screen at startup

The most common use case. Load a compressed loading screen and display it.

```

start:
    ld      hl, compressed_screen
    ld      de, $4000          ; screen memory
    call    dzx0_standard
    ; screen is now visible
    ; ... continue with demo/game ...

    include "dzx0_standard.asm"

compressed_screen:
    incbin "screen.zx0"

```

Pattern 2: Decompress to buffer between effects

Decompress the next effect's data into a scratch buffer while the current effect is still running, or during a fade-out.

```

ld      hl, scene2_data_zx0
ld      de, scratch_buffer    ; e.g., $C000 in bank 1
call    dzx0_standard
; scratch_buffer now holds the uncompressed data
; switch to scene 2, which reads from scratch_buffer

```

Pattern 3: Stream decompression during playback

For real-time effects that need a continuous data feed. LZ4 is the only practical choice here.

```
frame_loop:
    ld    hl, (lz4_read_ptr)    ; current position in compressed stream
    ld    de, frame_buffer
    ld    bc, 2048              ; bytes to decompress this frame
    call  lz4_decompress_partial
    ld    (lz4_read_ptr), hl    ; save position for next frame
    ; render from frame_buffer
    ; ...
    jr    frame_loop
```

At ~34 T/byte, LZ4 decompresses 2,048 bytes in 69,632 T-states — fitting within one frame (69,888 T-states on 48K). This is tight. Use border-time decompression or double-buffering for safety.

Pattern 4: Bank-switched compressed data (128K)

Store compressed data across multiple 16KB banks. Decompress from the currently paged bank, then switch banks when you run out.

```
    ld    a, (current_bank)
    or    $10                      ; bit 4 = ROM select
    ld    bc, $7ffd
    out   (c), a                  ; page bank into $C000-$FFFF

    ld    hl, $C000                ; compressed data starts at bank base
    ld    de, dest_buffer
    call  dzx0_standard

    ; Page next bank for next asset
    ld    a, (current_bank)
    inc   a
    ld    (current_bank), a
```

For large demos with many compressed assets, maintain a table of (bank, offset, destination) tuples and loop through them during loading.

Build Pipeline: From Asset to Binary

The compression step belongs in your Makefile, not in your head.

Source asset	Converter	Compressor	Assembler
(PNG)	--> (png2scr)	--> (zx0)	--> (sjasplus) --> .tap
(WAV)	--> (pt3tools)	--> (zx0)	--> (incbin)
(TMX)	--> (tmx2bin)	--> (exomizer)	

Makefile rules:

```

%.zx0: %.scr
    zx0 $< $@

# Compress large assets with Exomizer (one-time load)
%.exo: %.bin
    exomizer raw -c $< -o $@

# Build final binary
demo.bin: main.asm assets/title.zx0 assets/font.zx0
    sjasmpplus main.asm --raw=$@

```

Tool installation:

Tool	Source	Install
ZX0	github.com/einar-saukas/ZX0	gcc -O2 -o zx0 src/zx0.c src/compress.c src/optimize.c src/memory.c
Exomizer	github.com/bitmanipulators/exomizer	make in src/ directory
LZ4	github.com/lz4/lz4	make or brew install lz4
MegaLZ	github.com/Antonio-Cerra/megalzR	Older; check Introspec's Hype article for links

Quick Formulas

Bytes per frame at 50fps with decompressor X:

bytes_per_frame = 69,888 / speed_t_per_byte

Compressor	T/byte	Bytes/frame (48K)	Bytes/frame (128K Pentagon)
LZ4	34	2,055	2,108
MegaLZ fast	63	1,109	1,138
Pletter 5	69	1,012	1,038
ZX0	100	698	716
ApLib	105	665	682
Hrust 1	120	582	597
Exomizer	250	279	286

(128K Pentagon frame = 71,680 T-states)

Memory saved by compression over N screens:

saved = N * 6912 * (1 - ratio)

Example: 8 loading screens with Exomizer at 48.3% ratio save 8 * 6912 * 0.517 = 28,575 bytes — nearly two full 16KB banks.

See Also

- **Chapter 14:** Full discussion of compression theory, Introspec’s benchmark, ZX0 internals, and the delta + LZ pipeline.
- **Appendix B:** Sine table generation — when tables are small enough, consider generating instead of compressing.
- **Appendix A:** Z80 instruction reference — LDIR, PUSH/POP, and other instructions used in decompressors.

Sources: Introspec “Data Compression for Modern Z80 Coding” (Hype, 2017); Introspec “Compression on the Spectrum: MegaLZ” (Hype, 2019); Einar Saukas, ZX0/ZX7/ZX1/ZX2 (github.com/einar-saukas); Break Space NFO (Thesuper, 2016)

Appendix D: Development Environment Setup

“You need five things: an editor, an assembler, an emulator, a debugger, and a Makefile. Everything else is optional.” – Chapter 1

This appendix walks you through setting up a complete Z80 development environment from scratch. By the end, you will be able to compile every assembly example in this book, run them in an emulator, and debug them with breakpoints and register inspection. The instructions cover macOS, Linux, and Windows.

If you have already followed the setup in Chapter 1, you have most of this in place. This appendix adds detail, covers alternative configurations, and serves as a single reference you can return to when setting up a new machine.

1. The Assembler: sjasmplus

Every code example in this book is written for **sjasmplus**, an open-source Z80/Z80N macro assembler by z00m128. It supports the full Z80 instruction set including all IX/IY indexed modes, macros, Lua scripting, multiple output formats, and expressions that make demoscene code practical to write.

Installing from Source

The most reliable way to get sjasmplus is to build it from source. This guarantees you have a known version and avoids platform-specific packaging issues.

```
cd sjasmplus
make
```

On macOS, you need Xcode command-line tools (`xcode-select --install`). On Linux, you need `g++` and `make` (install via your package manager). On Windows, use MinGW or WSL.

After building, copy the sjasmplus binary to somewhere in your PATH:

```
sudo cp sjasmplus /usr/local/bin/
```

```
# Verify
sjasmplus --version
```

You should see version 1.20.x or later. This book was developed and tested with v1.21.1.

Version Pinning

The book's repository pins sjasmplus as a git submodule in tools/sjasmplus/. If you clone the repository with `--recursive`, you get the exact version used to compile every example:

```
cd antique-toy/tools/sjasmplus
make
```

This is the safest approach. Assembler behaviour can change between versions – an expression that works in 1.21 might parse differently in 1.22.

Key Flags

Flag	Purpose	Example
<code>--nologo</code>	Suppress the startup banner	<code>sjasmplus --nologo main.a80</code>
<code>--raw=FILE</code>	Output a raw binary (no header)	<code>sjasmplus --raw=output.bin main.a80</code>
<code>--sym=FILE</code>	Write a symbol file (for debuggers)	<code>sjasmplus --sym=output.sym main.a80</code>
<code>--fullpath</code>	Show full file paths in error messages	Useful with VS Code problem matcher
<code>--msg=war</code>	Suppress info messages, show only warnings and errors	Cleaner build output
<code>--syntax=abf</code>	Enable all syntax features (A as accumulator alias, brackets for indirection, full fake instructions)	Recommended for beginners; allows <code>add a, b</code> alongside <code>add b</code>

A typical build command for a chapter example:

File Extension

All Z80 assembly source files in this book use the `.a80` extension. This is a convention, not a requirement – sjasmplus does not care about extensions. We use `.a80` because it is recognised by the Z80 Macro Assembler VS Code extension and distinguishes our source from other assembly dialects.

Hex Notation

The book uses `$FF` for hexadecimal values. sjasmplus also accepts `#FF` and `0FFh`, but `$FF` is the convention throughout. The standalone `$` symbol represents the current program counter address, used in constructs like `jr $` (infinite loop) or `dw $ + 4`.

2. The Editor: VS Code

Any text editor works. We recommend **Visual Studio Code** because of its Z80-specific extensions and integrated terminal. The entire workflow – edit, build, debug – happens in one window.

Essential Extensions

Install these from the VS Code Extensions marketplace (Ctrl+Shift+X):

Extension	Author	What It Does
Z80 Macro Assembler	mborik (mborik.z80-macroasm)	Syntax highlighting, code completion, symbol resolution for Z80 assembly. Understands sjasmpplus syntax including macros and local labels.
Z80 Assembly Meter	Nestor Sancho	Displays byte count and T-state cost of selected instructions in the status bar. Select a block, see its total cost instantly. Indispensable for cycle counting.
DeZog	Maziac	Z80 debugger. Connects to emulators or its internal simulator. Breakpoints, register watches, memory inspection. See Section 4.

Build Task

Set up a build task so Ctrl+Shift+B compiles your current file. Create `.vscode/tasks.json` in your project root:

```
"version": "2.0.0",
"tasks": [
  {
    "label": "Assemble Z80",
    "type": "shell",
    "command": "sjasmpplus",
    "args": [
      "--fullpath",
      "--nologo",
      "--msg=war",
      "${file}"
    ],
    "group": {
      "kind": "build",
      "isDefault": true
    },
    "problemMatcher": {
      "owner": "z80",
```

```

    "fileLocation": "absolute",
    "pattern": {
      "regexp": "^(.*)\\((\\d+\\):\\s+(error|warning)[^:]*:\\s+(.*)$",
      "file": 1,
      "line": 2,
      "severity": 3,
      "message": 4
    }
  }
}
]
}

```

The `problemMatcher` parses `sjasmplus` error output so that clicking an error in the terminal jumps to the offending line. The `--fullpath` flag ensures file paths are absolute, which VS Code needs to resolve them correctly.

Recommended Settings

Add these to your workspace `.vscode/settings.json` for a clean Z80 editing experience:

```

"editor.tabSize": 8,
"editor.insertSpaces": false,
"editor.rulers": [80],
"files.associations": {
  "*.a80": "z80-macroasm"
}
}

```

Tab size 8 with real tabs matches the traditional assembly convention where mnemonics and operands align at fixed columns.

3. Emulators

You need an emulator to run your assembled code. Different emulators serve different purposes.

ZEsaUX - Feature-Rich Debugging

ZEsaUX by cerebellio is the most feature-rich open-source ZX Spectrum emulator. It supports the full range of Spectrum models (48K, 128K, +2, +3, Pentagon, Scorpion, ZX-Uno, ZX Spectrum Next), has a built-in debugger, and integrates with DeZog for VS Code debugging.

Install:

- macOS: `brew install zesarux`
- Linux: Build from source or use the AppImage from <https://github.com/cherandezba/zesarux>
- Windows: Download the installer from the ZEsaUX website

Why ZEsarUX for this book: Most chapter examples target the ZX Spectrum 128K. ZEsarUX emulates 128K memory banking, AY sound, TurboSound (two AY chips), and the contention patterns of different models. Its built-in debugger shows registers, memory, and disassembly without needing VS Code. And its DeZog integration provides the full VS Code debugging experience when you need it.

Quick launch:

```
zesarux --machine 128k --nosplash output.sna
```

```
# Run a .tap file
```

```
zesarux --machine 128k --nosplash output.tap
```

Fuse - Lightweight and Accurate

Fuse (the Free Unix Spectrum Emulator) by Philip Kendall is lightweight, cycle-accurate, and available on every platform. It is the best choice for quick testing when you do not need the full debugger.

Install:

- macOS: `brew install fuse-emulator`
- Linux: `apt install fuse-emulator-sdl` (Debian/Ubuntu) or `dnf install fuse-emulator` (Fedora)
- Windows: Download from the Fuse website

Quick launch:

```
fuse --machine pentagon output.sna
```

```
# Run as 128K Spectrum
```

```
fuse --machine 128 output.tap
```

Fuse is particularly good for testing timing-sensitive code because its cycle accuracy is well-verified. If your border stripe timing harness (Chapter 1) shows different results in Fuse versus another emulator, trust Fuse.

Unreal Speccy - Windows, Pentagon-Focused

If you develop primarily on Windows and target Pentagon timing, **Unreal Speccy** is a strong choice. It has a built-in debugger with memory map, breakpoints, and AY register monitoring. It emulates Pentagon and Scorpion hardware accurately.

Download from <http://dlcorp.nedopc.com/viewforum.php?f=27> or search for “Unreal Speccy Portable.”

For Agon Light 2

The Agon Light 2 uses an eZ80 CPU and a different hardware architecture. Chapter 22 covers Agon development in detail. For emulation, **Fab Agon Emulator** provides a software simulation of the Agon hardware (eZ80 + ESP32 VDP). It is available at <https://github.com/tomm/fab-agon-emulator> and runs on macOS, Linux, and Windows.

Which Emulator Should I Use?

Situation	Recommended Emulator
Day-to-day development, running examples	Fuse (fast startup, accurate)
Debugging with breakpoints and register watches	ZEsarUX + DeZog
AY/TurboSound music development	ZEsarUX (best AY emulation)
Pentagon timing verification	Fuse or Unreal Speccy
Agon Light 2 development	Fab Agon Emulator
Quick sanity check on Windows	Unreal Speccy

4. The Debugger: DeZog

DeZog by Maziac is a VS Code extension that turns your editor into a Z80 debugger. It connects to ZEsaUX, CSpect, or its own internal Z80 simulator and provides the debugging experience modern developers expect: breakpoints, stepping, register watches, memory inspection, disassembly view, and call stack.

Chapter 23 discusses DeZog in the context of AI-assisted development. This section covers the practical setup.

Installation

1. Open VS Code.
2. Go to Extensions (Ctrl+Shift+X).
3. Search for “DeZog” by Maziac.
4. Click Install.

Connecting to ZEsaUX

DeZog communicates with ZEsaUX over a socket connection. First, launch ZEsaUX with its ZRCP (ZEsaUX Remote Control Protocol) server enabled:

Then create a `.vscode/launch.json` in your project:

```
"version": "0.2.0",
"configurations": [
  {
    "type": "dezog",
    "request": "launch",
    "name": "DeZog + ZEsaUX",
    "remoteType": "zesarux",
    "zesarux": {
```

```

        "port": 10000
    },
    "sjasmpplus": [
        {
            "path": "build/output.sld"
        }
    ],
    "topOfStack": "$FF00",
    "commandsAfterLaunch": [
        "-sprites disable",
        "-patterns disable"
    ]
}
]
}

```

The `sjasmpplus` section points to the `.sld` (Source Level Debug) file that `sjasmpplus` generates with the `--sld=FILE` flag. This gives DeZog source-level debugging – breakpoints on source lines, not just addresses.

To generate the `.sld` file, add the flag to your build command:

Using the Internal Simulator

For quick debugging without launching an external emulator, DeZog includes a built-in Z80 simulator. Change your `launch.json` to:

```

"type": "dezog",
"request": "launch",
"name": "DeZog Internal Simulator",
"remoteType": "zsim",
"zsim": {
    "machine": "spectrum",
    "memoryModel": "ZX128K"
},
"sjasmpplus": [
    {
        "path": "build/output.sld"
    }
],
"topOfStack": "$FF00"
}

```

The internal simulator is faster to start and does not require ZEsarUX to be installed. It lacks accurate contention emulation, so do not use it for timing-critical debugging – but for logic debugging (does my multiply routine produce the right result?), it is perfect.

Key DeZog Features

Breakpoints. Click the gutter next to a source line to set a breakpoint. Execution pauses when the Z80 program counter reaches that address. You can also set

conditional breakpoints (e.g., break when A == \$FF).

Register watches. The Variables panel shows all Z80 registers: AF, BC, DE, HL, IX, IY, SP, PC, and the alternate set (AF', BC', DE', HL'). Individual flags (C, Z, S, P/V, H, N) are broken out for easy reading.

Memory inspection. The Memory panel shows a hex dump of any address range. You can type an address and see what is there. Essential for verifying lookup tables, screen memory contents, and stack state.

Disassembly view. Even without source-level debugging, DeZog disassembles the code around the current PC. Useful for understanding what self-modifying code actually looks like at runtime.

Call stack. DeZog tracks CALL/RET pairs and shows a call stack. This works for conventional code. Self-modifying code and RET-chaining (Chapter 3) will confuse the stack tracker – this is expected.

5. Building the Book's Examples

Clone the Repository

```
cd antique-toy
```

The `--recursive` flag pulls the `sjsmplus` submodule in `tools/sjsmplus/`. If you already cloned without it:

Prerequisites

You need `make` and a C++ compiler (for building `sjsmplus` from the submodule). On most systems these are already present:

- macOS: `xcode-select --install`
- Debian/Ubuntu: `sudo apt install build-essential`
- Fedora: `sudo dnf install make gcc-c++`
- Windows: Use WSL, or install MinGW and GNU Make

Build Commands

The project Makefile handles everything. All compiled output goes to `build/`, which is gitignored.

Command	What It Does
<code>make</code>	Compile all chapter examples with <code>sjsmplus</code>
<code>make test</code>	Assemble all examples, report pass/fail for each
<code>make ch01</code>	Compile only Chapter 1 examples
<code>make ch04</code>	Compile only Chapter 4 examples
<code>make demo</code>	Build the “Antique Toy” companion demo
<code>make clean</code>	Remove all build artifacts

A successful make test run looks like this:

```
OK  chapters/ch01-thinking-in-cycles/examples/timing.a80
OK  chapters/ch04-maths/examples/multiply.a80
OK  chapters/ch05-wireframe-3d/examples/cube.a80
...
---
```

12 passed, 0 failed

If any example fails, the output shows FAIL with the filename. Run the failing file manually with sjasmplus to see the detailed error:

Running an Example

After compiling, load the output binary into your emulator. The exact method depends on the output format:

```
fuse --machine pentagon build/ch01-thinking-in-cycles/examples/timing.sna
```

```
# If you built a raw binary, you need to create a .tap or .sna first,
# or load it at the correct address in the emulator's debugger
```

Most chapter examples use ORG \$8000 and produce raw binaries. To run them, either:

1. Use a .tap wrapper (the Makefile generates these when the source includes the appropriate directives), or
2. Load the binary at address \$8000 in your emulator's debugger and set PC to \$8000.

6. Project Structure for Your Own Code

When you start writing your own Z80 code beyond the book's examples, here is a recommended directory layout:

```
my-demo/
src/
  main.a80          -- entry point, includes other files
  effects/
    plasma.a80      -- individual effect routines
    scroller.a80
  data/
    font.a80        -- DB/DW data tables
    sintable.a80
  lib/
    multiply.a80    -- reusable utility routines
    draw_line.a80
assets/
  title.scr         -- raw screen files
  music.pt3         -- tracker music
build/              -- compiled output (gitignored)
```

```

Makefile
.vscode/
  tasks.json      -- build task
  launch.json     -- DeZog configuration

```

Minimal Makefile

```

BUILD_DIR := build
SJASM_FLAGS := --nologo

.PHONY: all clean run

all: $(BUILD_DIR)/demo.bin

$(BUILD_DIR)/demo.bin: src/main.a80 src/effects/*.a80 src/data/*.a80
↪ src/lib/*.a80
  @mkdir -p $(BUILD_DIR)
  $(SJASMPPLUS) $(SJASM_FLAGS) --raw=$@ --sym=$(BUILD_DIR)/demo.sym
  ↪ --sld=$(BUILD_DIR)/demo.sld $<

run: $(BUILD_DIR)/demo.bin
  fuse --machine 128 $(BUILD_DIR)/demo.sna

clean:
  rm -rf $(BUILD_DIR)

```

Key points:

- The main source file uses `INCLUDE` directives to pull in other files. `sjasmpplus` resolves includes relative to the source file's directory.
- The `--sym` flag generates a symbol file for reference. The `--sld` flag generates source-level debug data for DeZog.
- List all included source files as dependencies so make rebuilds when any file changes.

Include Convention

In your `main.a80`:

```

; --- Entry point ---
start:
  di
  ; set up stack, interrupts, etc.
  ld  sp, $FF00
  ; ...

  include "lib/multiply.a80"
  include "effects/plasma.a80"
  include "data/sintable.a80"

```

`sjasmpplus` processes `INCLUDE` files inline, as if the text were pasted at that point. Keep your includes organised: library routines first, then effects, then data (since

data is typically placed at the end of the binary).

7. Alternative Tools

This book uses sjasmplus because it is the most capable Z80 assembler for demoscene and game development work. But you may encounter other tools in the community.

Other Assemblers

Assembler	Notes
z80asm	Part of the z88dk C cross-compiler suite. Good if you mix C and assembly.
RASM	Different syntax conventions. By Roudoudou. Fast, supports CPC and Spectrum. Popular in the Amstrad CPC scene.
pasmo	Simple, portable, limited features. Suitable for small standalone programs but lacks macros and advanced features needed for larger projects.

The book's examples use sjasmplus-specific features (local labels with `.`, negative DB values, `$` hex prefix) that may not work unmodified in other assemblers. If you want to port an example to a different assembler, the changes are usually minor: label syntax, hex notation, and directive names.

Other VS Code Extensions

Extension	Notes
SjASMPlus Syntax	Alternative syntax highlighting tuned specifically for sjasmplus. Try it if z80-macroasm does not highlight a sjasmplus-specific feature correctly.
Z80 Debugger (Spectron)	An older VS Code Z80 debugger. DeZog has largely superseded it.

CSpect - Next-Focused Emulator

If you have a ZX Spectrum Next or are targeting Next-specific features (copper, layer 2, DMA, 28MHz turbo mode), **CSpect** by Mike Dailly is the reference emulator. DeZog connects to CSpect the same way it connects to ZEsarUX. CSpect is Windows-only but runs under Wine on macOS and Linux.

SpectrumAnalyzer

A browser-based tool that visualises ZX Spectrum screen memory layout, attribute conflicts, and timing. Useful for understanding Chapter 2's discussion of the screen interleave. Available at <https://shiru.undergrund.net/spectrumanalyzer.html> (or search for "ZX Spectrum screen analyzer").

8. Troubleshooting

"sjasmpplus: command not found"

The binary is not in your PATH. Either copy it to `/usr/local/bin/` (macOS/Linux), add its directory to your PATH, or set the `SJASMPPLUS` variable when running `make`:

Compilation errors in the book's examples

First, make sure you are using the `sjasmpplus` version from the submodule (`tools/sjasmpplus/`). Newer or older versions may have different behaviour. Second, check that you are assembling the file from the correct directory – `sjasmpplus` resolves `INCLUDE` paths relative to the source file, not the working directory.

DeZog cannot connect to ZEsarUX

1. Make sure ZEsarUX is running with `--enable-remoteprotocol`.
2. Check that the port number in `launch.json` matches the `--remoteprotocol-port` argument.
3. On macOS, you may need to allow ZEsarUX through the firewall (System Settings > Privacy & Security).
4. Try restarting ZEsarUX before launching the DeZog session.

Emulator shows garbled screen

If you load a raw binary and see garbage, the most likely cause is a wrong load address. The book's examples use `ORG $8000`. Make sure your emulator loads the binary at that address, not at `$0000` or some other default. Using `.sna` or `.tap` output (which includes address information) avoids this problem.

Build output is empty or zero bytes

Check that your source file has an `ORG` directive and at least one instruction. `sjasmpplus` with `--raw=` produces a binary from the first byte emitted to the last. If nothing is emitted (e.g., the file contains only `ORG $8000` with no code), the output is empty.

Tool Reference

Tool	Purpose	Official URL	Book Reference
sjasm-plus	Z80/Z80N macro assembler	https://github.com/z00m128/sjasm-plus	Chapter 1, all examples
VS Code	Editor and IDE	https://code.visualstudio.com	Chapter 1
Z80 Macro Assembler	VS Code syntax extension	Marketplace: mborik.z80-macroasm	Chapter 1
Z80 Assembly Meter	Cycle count display	Marketplace: Nestor Sancho	Chapter 1
DeZog	VS Code Z80 debugger	Marketplace: Maziac / https://github.com/maziac/De-Zog	Chapter 23
ZEsaUX	Feature-rich Spectrum emulator	https://github.com/cher-nandezba/zesarux	Chapter 1, Chapter 23
Fuse	Lightweight Spectrum emulator	https://fuse-emulator.sourceforge.net	Chapter 1
Unreal Speccy	Pentagon-focused emulator	http://dlcorp.net/dopc.com	Chapter 1
CSpect	ZX Spectrum Next emulator	https://daily.blogspot.com	Chapter 22 (Next features)
Fab Agon Emulator	Agon Light 2 emulator	https://github.com/tommfab/agon-emulator	Chapter 22
ZX0	Optimal LZ compressor	https://github.com/einar-saukas/ZX0	Chapter 14, Appendix C
Exomizer	Best-ratio compressor	https://github.com/bit-manipulators/exomizer	Chapter 14, Appendix C
Vortex Tracker II	AY music tracker	https://github.com/ivan-pirog/vortextracker	Chapter 11

See Also

- **Chapter 1:** First practical – setting up VS Code, sjasmplus, and the timing harness.
- **Chapter 22:** Agon Light 2 platform setup, eZ80 toolchain, Fab Agon Emulator.
- **Chapter 23:** DeZog integration with AI-assisted workflows.
- **Appendix A:** Z80 instruction reference – the instructions you will be debugging.
- **Appendix G:** AY register reference – useful when debugging sound code in ZEsaUX.

Appendix E: eZ80 Quick Reference

“The same instruction set, a completely different machine.” – Chapter 22

This appendix is a reference card for Z80 programmers approaching the eZ80 for the first time. It covers the architectural differences, the mode system, the new instructions, and the Agon Light 2 specifics you need for porting. It is not an exhaustive eZ80 manual — it is the subset that matters for the work in Chapter 22.

If you already know the Z80 (and if you have read this far, you do), the eZ80 will feel familiar. The registers have the same names, the instructions have the same mnemonics, the flags work the same way. But three things are different: addresses are 24 bits wide, registers can be 24 bits wide, and there is a mode system that controls which width is active. Everything else follows from that.

1. Architecture Overview

The eZ80 is Zilog’s extended Z80, designed for embedded systems that need more than 64 KB of address space. It is a strict superset of the Z80 — every Z80 opcode is valid on the eZ80, with identical behaviour. The extensions add 24-bit addressing, 24-bit register width, and a handful of new instructions.

Feature	Z80	eZ80
Address width	16-bit (64 KB)	24-bit (16 MB)
Register width	16-bit (HL, BC, DE, SP, IX, IY)	16-bit or 24-bit (mode-dependent)
Stack frame per PUSH/CALL	2 bytes	2 bytes (Z80 mode) or 3 bytes (ADL mode)
Hardware multiply	None	MLT rr (8x8 unsigned)
Instruction prefixes	CB, DD, ED, FD	Same, plus mode suffix prefixes
MBASE register	N/A	Provides upper 8 bits of addresses in Z80 mode
New instructions	N/A	LEA, PEA, MLT, TST, TSTIO, SLP, IN0/OUT0, STMIX/RSMIX

The key mental model: in **ADL mode** (Address Data Long), the eZ80 behaves like a Z80 with 24-bit registers and 24-bit addresses. In **Z80-compatible mode**, it behaves like a standard Z80, with 16-bit registers and the MBASE register providing the missing upper 8 address bits.

2. Mode System

The eZ80 has two operating modes that control register width and address generation. Understanding these modes is the single most important concept for any Z80 programmer approaching the eZ80.

The Two Modes

Z80-compatible mode. Registers are 16 bits wide. Addresses are 16 bits, with MBASE providing the upper 8 bits. `LD HL,$4000` loads a 16-bit value. `PUSH HL` pushes 2 bytes. Code behaves exactly like a standard Z80.

ADL mode (Address Data Long). Registers are 24 bits wide. Addresses are 24 bits. `LD HL,$040000` loads a 24-bit value. `PUSH HL` pushes 3 bytes. This is the native mode of the eZ80 and the default on the Agon Light 2.

Mode Suffixes

Individual instructions can override the current mode using suffix prefixes:

Suffix	Meaning	Register width	Address width
.SIS	Short Immediate, Short	16-bit	16-bit
.LIS	Long Immediate, Short	24-bit	16-bit
.SIL	Short Immediate, Long	16-bit	24-bit
.LIL	Long Immediate, Long	24-bit	24-bit

The first letter (S/L) controls register width for that instruction. The third letter (S/L) controls address width. In ADL mode, .SIS forces an instruction to behave as standard Z80. In Z80 mode, .LIL forces an instruction to behave as full 24-bit.

Mode-Switching Calls and Jumps

Calls and jumps can switch the processor mode at the target:

Instruction	Current mode	Target mode	Return address size
CALL.IS nn	ADL	Z80	3 bytes (ADL convention)
CALL.IL nn	Z80	ADL	3 bytes (long)
JP.SIS nn	any	Z80	N/A (no return address)
JP.LIL nn	any	ADL	N/A (no return address)
RST.LIL \$08	Z80	ADL	3 bytes (long)

The .IS suffix means “Instruction Short” — the target runs in Z80 mode. .IL means “Instruction Long” — the target runs in ADL mode.

The Practical Rule

Stay in ADL mode. MOS boots the Agon in ADL mode. MOS API calls expect ADL mode. VDP commands are sent through MOS routines that assume 24-bit stack frames. If you switch to Z80 mode and call MOS, the stack frame width mismatch will corrupt the stack and crash.

If you need tight 16-bit loops (e.g., porting a Z80 inner loop without rewriting it), use the .SIS suffix on individual instructions rather than switching the entire processor mode.

The MBASE Trap

In Z80-compatible mode, the MBASE register provides the upper 8 bits of every memory address — including instruction fetches. If you change MBASE while executing in Z80 mode, the next instruction fetch uses the new MBASE value. Unless your code exists at the corresponding physical address, execution jumps into garbage.

Rule: if you must use Z80 mode, set MBASE once at startup and leave it alone. Better yet, stay in ADL mode.

3. New Instructions

These instructions exist on the eZ80 but not on the standard Z80. They are the reason the eZ80 is more than just a Z80 with wider registers.

Arithmetic and Test

Instruction	Bytes	Cycles	Description
MLT BC	2	6	8x8 unsigned multiply: B * C -> BC
MLT DE	2	6	8x8 unsigned multiply: D * E -> DE
MLT HL	2	6	8x8 unsigned multiply: H * L -> HL
MLT SP	2	6	8x8 unsigned multiply: SPH * SPL -> SP (rarely useful)
TST A, n	3	7	Test: A AND n, set flags, A unchanged
TST A, r	2	4	Test: A AND r, set flags, A unchanged
TSTIO n	3	12	Test I/O: (C) AND n, set flags

Address Computation

Instruction	Bytes	Cycles	Description
LEA BC, IX+d	3	4	BC = IX + signed displacement d
LEA DE, IX+d	3	4	DE = IX + d
LEA HL, IX+d	3	4	HL = IX + d
LEA IX, IX+d	3	4	IX = IX + d (add displacement to IX)
LEA BC, IY+d	3	4	BC = IY + d
LEA DE, IY+d	3	4	DE = IY + d
LEA HL, IY+d	3	4	HL = IY + d
LEA IY, IY+d	3	4	IY = IY + d (add displacement to IY)
PEA IX+d	3	7	Push (IX + d) onto stack
PEA IY+d	3	7	Push (IY + d) onto stack

LEA computes an effective address without performing a memory access — it is pure register arithmetic. On the standard Z80, computing $HL = IX + 5$ requires a multi-instruction sequence (PUSH IX / POP HL / LD DE,5 / ADD HL,DE). LEA does it in a single instruction at 4 cycles.

I/O and System

Instruction	Bytes	Cycles	Description
IN0 r, (n)	3	7	Read from internal I/O port (8-bit address)
OUT0 (n), r	3	7	Write to internal I/O port (8-bit address)
SLP	2	-	Sleep: halt CPU until interrupt (lower power than HALT)
STMIX	2	4	Set mixed-mode flag (enable ADL/Z80 interleaving)
RSMIX	2	4	Reset mixed-mode flag

IN0/OUT0 use an 8-bit port address (unlike the standard IN/OUT which can use 16-bit port addresses via BC). They are designed for the eZ80's internal peripherals and are rarely used in Agon game code.

4. MLT — The Game Changer

Of all the eZ80's new instructions, MLT is the most impactful for game and demo code. It performs an 8x8 unsigned multiply in a single instruction.

On the standard Z80, 8x8 multiplication requires a shift-and-add loop:

```
; Cost: 196-204 T-states, 14 bytes
mulu_z80:
    ld  a, 0          ; 7T  clear accumulator
    ld  d, 8          ; 7T  8 bits
```

```

.loop:
    rr    c                ; 8T    shift multiplier bit into carry
    jr    nc, .noadd       ; 7/12T
    add   a, b             ; 4T    add multiplicand
.noadd:
    rra                ; 4T    shift result right
    dec   d                ; 4T
    jr    nz, .loop        ; 12T
    ret                   ; 10T ~200 T-states total

```

On the eZ80:

```

; Cost: 6 cycles, 2 bytes
    mlt   bc                ; BC = B * C. Done.

```

Six cycles. Two bytes. No loop, no carry management, no temporary registers. The result lands in the full 16-bit register pair: the high byte of the product is in B, the low byte in C.

What MLT Enables

Sine table indexing. Computing $\text{base} + \text{angle} * \text{stride}$ for variable-stride sine lookups reduces from a subroutine call to two instructions (MLT + ADD).

Sprite offset calculation. Finding the address of sprite frame N in a sprite sheet: $\text{base} + \text{frame} * \text{frame_size}$. With MLT, this is trivial when `frame_size` fits in 8 bits.

Fixed-point arithmetic. Multiplying two 8-bit fixed-point values (e.g., $\text{velocity} * \text{friction}$) becomes a single instruction instead of a 200-T-state loop.

Tile map addressing. Computing $\text{tile_base} + (\text{row} * \text{width} + \text{col})$ where `width` fits in 8 bits: one MLT for the row offset, one ADD for the column.

MLT Limitations

- **Unsigned only.** For signed multiplication, adjust the sign manually after the MLT.
- **8x8 only.** For 16x16 multiply, you still need a multi-step algorithm (though you can build it from MLT components).
- **Result overwrites both operands.** MLT BC destroys both B and C, replacing them with the 16-bit product. Save the inputs first if you need them.

5. Key Differences Summary

Feature	Z80 (ZX Spectrum 128K)	eZ80 (Agon Light 2)
Clock	3.5 MHz	18.432 MHz
Address width	16-bit (64 KB visible)	24-bit (16 MB, 512 KB populated)
Register width	16-bit	16-bit (Z80 mode) or 24-bit (ADL mode)

Feature	Z80 (ZX Spectrum 128K)	eZ80 (Agon Light 2)
Stack per PUSH/CALL	2 bytes	2 bytes (Z80 mode) or 3 bytes (ADL mode)
Multiply instruction	None (shift-and-add loop, ~200T)	MLT rr (6 cycles)
RAM	128 KB banked (8 x 16 KB pages)	512 KB flat
RAM access model	Bank switching via port \$7FFD	Flat 24-bit addressing
Video	ULA: direct memory-mapped at \$4000	VDP (ESP32): command-based via UART
Sound	AY-3-8910 via I/O ports	VDP audio via serial commands
Interrupts	IM1 (RST \$38) or IM2 (vector table)	IM2 with MOS-managed vectors
Frame budget (50 Hz)	~71,680 T-states (Pentagon)	~368,640 T-states
Contended memory	Yes (ULA steals cycles from \$4000-\$7FFF)	No contention
OS	None (bare metal)	MOS (Machine Operating System)
Storage	Tape / DivMMC	SD card (FAT32, MOS file API)

6. Agon Light 2 Specifics

Hardware

- **CPU:** Zilog eZ80F92, 18.432 MHz
- **RAM:** 512 KB external SRAM, flat address space
- **VDP:** ESP32-PICO-D4 running FabGL, communicates with eZ80 via UART at 1,152,000 baud
- **Video modes:** Multiple, up to 640x480x64 colours. Most games use 320x240 or 320x200.
- **Hardware sprites:** Up to 256, managed entirely by VDP
- **Hardware tilemaps:** Scrollable tile layers, managed by VDP
- **Audio:** VDP synthesis — sine, square, triangle, sawtooth, noise. Per-channel ADSR envelopes.
- **Storage:** MicroSD card, FAT32, accessed through MOS file API

Frame Budget

At 18.432 MHz and 50 Hz refresh:

$18,432,000 \text{ cycles/sec} / 50 \text{ frames/sec} = 368,640 \text{ cycles/frame}$

Compare with Pentagon (3.5 MHz, 50 Hz): 71,680 T-states/frame. The Agon has roughly **5.1x** the per-frame budget. But since many eZ80 instructions also execute in fewer cycles than their Z80 equivalents, effective throughput for typical code is 5x-20x greater.

MOS API

MOS (Machine Operating System) provides the system interface. The standard entry point is RST \$08:

```
ld  a, mos_function    ; function number in A
rst $08                ; call MOS
; return value in A (and sometimes HL)
```

Key MOS API functions:

Function	Number	Description
mos_getkey	\$00	Wait for keypress, return ASCII code
mos_load	\$01	Load file from SD card to memory
mos_save	\$02	Save memory region to file on SD card
mos_cd	\$03	Change current directory
mos_dir	\$04	List directory contents
mos_del	\$05	Delete a file
mos_ren	\$06	Rename a file
mos_sysvars	\$08	Get pointer to system variables (keyboard map, VDP state, clock)
mos_fopen	\$0A	Open file, return handle
mos_fclose	\$0B	Close file handle
mos_fread	\$0C	Read bytes from file
mos_fwrite	\$0D	Write bytes to file
mos_fseek	\$0E	Seek within file

VDP Command Protocol

VDP commands are sent as byte streams using RST \$10 (MOS: output byte to VDP). Most commands begin with VDU 23, followed by command-specific parameters:

```
ld  a, byte_value
rst $10                ; MOS: send byte to VDP

; Example: set screen mode
ld  a, 22               ; VDU 22 = set mode
rst $10
ld  a, mode_number     ; 0-3 for standard modes
rst $10

; Example: move hardware sprite
; VDU 23, 27, 4, spriteNum    -- select sprite
; VDU 23, 27, 13, x.lo, x.hi, y.lo, y.hi -- set position
```

The VDP processes commands asynchronously. There is a serial transfer delay between sending a command and the VDP acting on it. For smooth animation, send all updates early in the frame.

7. Porting Checklist

When porting Z80 code from the Spectrum to eZ80 ADL mode on the Agon, walk through this checklist:

Addresses and pointers: - All addresses become 24-bit (3 bytes instead of 2)
 - Change DW (define word) to DL (define long) for address tables - Pointer table indexing changes from * 2 to * 3 - Ensure upper byte of 24-bit addresses is correct (typically \$00 or \$04)

Stack frames: - Every PUSH is 3 bytes, every CALL pushes a 3-byte return address
 - Verify PUSH/POP pairs are balanced — a mismatched pair corrupts 3 bytes, not 2
 - Stack-relative offsets (e.g., accessing parameters pushed by the caller) change

Block operations: - LDIR / LDDR use 24-bit BC in ADL mode — ensure the upper byte of BC is zero if your count fits in 16 bits - PUSH/POP-based block copy tricks push 3 bytes per PUSH, not 2

Multiply: - Replace shift-and-add multiply loops with MLT where applicable - MLT BC = B * C -> BC, MLT DE = D * E -> DE, MLT HL = H * L -> HL

I/O and peripherals: - Replace port I/O (OUT (C), A, IN A, (\$FE)) with MOS/VDP API calls - Replace direct framebuffer writes (\$4000-5AFF) with VDP commands — Replace AY register writes (FFFD/\$BFFD) with VDP sound commands

Memory architecture: - Remove all bank switching logic (port \$7FFD writes) — flat address space - Remove contended memory workarounds — no contention on the Agon - Remove shadow screen tricks — VDP handles double buffering

Code patterns that become unnecessary: - Self-modifying code for speed (still works, rarely worth the complexity) - Stack pointer tricks for fast screen fills (no framebuffer to fill) - Pre-shifted sprite copies (hardware sprites handle sub-pixel positioning) - Interleaved screen address calculations (DOWN_HL, pixel_addr — delete them)

Code patterns that transfer directly: - Entity system loops (just widen pointers) - AABB collision detection (8-bit comparisons, unchanged) - Fixed-point 8.8 arithmetic (byte-level, unchanged) - State machines and jump tables (widen table entries to 24-bit) - DJNZ loops, CPIR searches, flag-based branching (all identical)

See Also

- **Appendix A: Z80 Instruction Quick Reference** — complete Z80 instruction table with T-states, byte counts, and flag effects. Everything in Appendix A also applies to the eZ80 in Z80-compatible mode.
- **Chapter 22: Porting — Agon Light 2** — the full porting walkthrough, with before/after code examples for rendering, sound, input, and game logic.

Sources: Zilog eZ80 CPU User Manual (UM0077); Zilog eZ80F92 Product Specification (PS0153); Agon Light 2 Official Documentation, The Byte Attic; Dean Belfield, “Agon Light — Programming Guide” (breakin-

toprogram.co.uk); Agon MOS API Documentation (github.com/AgonConsole8/agon-docs); Chapter 22 of this book

Appendix F: Z80 Variants — Extended Instruction Sets

“The same instruction set, a completely different machine.” – Chapter 22

The Zilog Z80 did not stay frozen in 1976. Over five decades, the original design has been cloned, extended, reimagined, and — in the case of the ZX Spectrum Next — rebuilt by the very community that spent thirty years cursing its limitations. This appendix surveys the major Z80 variants and their instruction set extensions, with a focus on what matters to demoscene and game programmers. The standard Z80 instruction set is covered in Appendix A; the eZ80 gets a deep dive in Appendix E. This appendix is the big picture — how the variants relate, what each one adds, and why.

1. The Z80 Family Tree

The Z80 was designed by Federico Faggin and Masatoshi Shima at Zilog in 1976 as a software-compatible successor to the Intel 8080. It became the most widely used 8-bit CPU in history, powering everything from CP/M business machines to arcade cabinets to an entire generation of home computers. The instruction set was frozen at launch and never officially extended by Zilog — until the eZ80, two decades later.

But others did extend it. Here is the family at a glance:

Variant	Year	Notable Machine	Key Addition
Z80 (Zilog)	1976	ZX Spectrum, MSX, Amstrad CPC	The original. 158 documented instructions.
KR1858VM1 (Soviet)	~1986	Pentagon, Scorpion	Exact clone. No instruction changes.
NSC800 (National Semi)	1980	Various embedded	CMOS Z80 with 8085-style bus. No new instructions.
R800 (ASCII Corp)	1990	MSX turboR	Z80-compatible, radically different pipeline. MULUB, MULUW.

Variant	Year	Notable Machine	Key Addition
eZ80 (Zilog)	2001	Agon Light 2	24-bit addressing, MLT, LEA, PEA, ADL mode.
Z80N (Next team)	2017	ZX Spectrum Next	Demoscene wish list: MUL, MIRROR, LDIRX, PIXELDN, barrel shifts.

The Z80 and its clones share an identical instruction set. The R800, eZ80, and Z80N each added instructions to solve specific problems — but very different problems, reflecting very different design goals.

2. Z80N — The Demoscener’s Wish List

The Z80N is the CPU in the ZX Spectrum Next. It was designed by Victor Trucco, Fabio Belavenuto, and the Next team — people who had spent decades writing Z80 code and knew exactly where it hurt. Every new instruction addresses a specific, documented pain point from thirty years of Spectrum programming. The Z80N runs at 28 MHz (8x the original Spectrum clock) and adds roughly 40 new instructions, all encoded in previously-unused \$ED xx opcode space.

The best way to understand the Z80N extensions is by the problem each instruction solves.

Screen Navigation (the DOWN_HL problem)

On a standard Spectrum, calculating a screen address from pixel coordinates takes 50-60 T-states and a page of code (see `pixel_addr` in Appendix A). Moving one pixel row down requires the infamous `DOWN_HL` routine — a conditional maze of `INC`, `AND`, `ADD`, and `SUB` that handles character boundaries and third boundaries. The Z80N replaces all of this with single instructions.

Instruction	Bytes	T	What it replaces
PIXELDN	2	8	The 10+ instruction <code>DOWN_HL</code> sequence (check third boundary, handle wrap, adjust H and L). Moves HL one pixel row down in screen memory.
PIXELAD	2	8	Full screen address calculation from (D,E) coordinates. Replaces the <code>pixel_addr</code> routine (~55T, 15+ instructions).
SETAE	2	8	Sets the appropriate pixel bit in A based on the low 3 bits of E (the x-coordinate). Replaces a lookup table or shift sequence.

With these three instructions, the entire pixel-plotting sequence that consumed 70+ T-states and 20+ bytes on the original Z80 becomes:

```

pixelad          ; 8T  HL = screen address from (D,E)
setae            ; 8T  A = pixel bit mask from E
or   (hl)        ; 7T  set pixel (non-destructive)
; ... ld (hl), a to write

```

Sprite Rendering (the masked blit problem)

The single most CPU-intensive operation in any Spectrum game or demo is the masked sprite blit: copying a rectangular block of sprite data to screen memory while skipping transparent pixels. On the standard Z80, this requires an inner loop of LD/AND/OR/LD per byte, typically 30-40 T-states per pixel byte. The Z80N adds block-copy instructions with built-in transparency.

Instruction	Bytes	T	What it replaces
LDIX	2	16	LDI but skips the copy if (HL) == A. One-instruction transparent copy: load A with the transparent colour, point HL at source, DE at destination, and each byte is copied only if it is not the transparent value.
LDDX	2	16	Same as LDIX but decrementing (like LDD).
LDIRX	2	21/16	Repeating LDIX. Hardware masked sprite blit in a single instruction. Copies BC bytes from (HL) to (DE), skipping any byte equal to A. 21T per byte while BC>0, 16T on the final iteration.
LDDRX	2	21/16	Repeating LDDX.
LDPIRX	2	21/16	Pattern fill with transparency from an 8-byte aligned source. Reads from (HL & \$FFF8) + (E & 7), copies to (DE) if not equal to A, increments DE, decrements BC. Hardware tiled background renderer. 21T per byte while BC>0, 16T on the final iteration.

LDIRX alone replaces the most heavily optimised inner loop in decades of Spectrum game code. LDPIRX is even more exotic — it treats the source as a repeating 8-byte pattern, effectively giving you a hardware tile renderer with transparency. Combined with the Next's Layer 2 and tilemap hardware, these instructions make the Spectrum Next a qualitatively different platform for sprite-heavy games.

Arithmetic (the multiply problem)

The Z80 has no multiply instruction. Every multiply in Spectrum code is a shift-and-add loop costing 150-250 T-states (see `mulu112` in Appendix A). The Z80N fixes this with a single instruction.

Instruction	Bytes	T	What it replaces
MUL D,E	2	8	8x8 unsigned multiply, result in DE. Replaces the ~200T shift-and-add loop.

Eight T-states. At 28 MHz, that is 286 nanoseconds. The same operation on a standard 3.5 MHz Spectrum takes roughly 57 microseconds — a 200:1 improvement when you factor in both the faster clock and the faster instruction. Rotation matrices, coordinate transforms, texture mapping, perspective projection — everything that needs multiply is fundamentally cheaper on the Z80N.

Bit Manipulation

Instruction	Bytes	T	What it replaces
MIRROR	2	8	Reverses all 8 bits of A. Horizontal sprite flip without a 256-byte lookup table. On the standard Z80, bit-reversing A requires either an unrolled 18-instruction sequence (<code>LD B,A : XOR A then 8× RR B : RLA</code> , ~104T) or a 256-byte lookup table (11T but costs 256 bytes of RAM).
SWAPNIB	2	8	Swaps the high and low nibbles of A. Replaces <code>RLCA : RLCA : RLCA : RLCA</code> (16T, 4 bytes).
TEST nn	3	11	AND A, nn without storing the result — sets flags but preserves A. Like a CP for bitwise AND. Similar to the eZ80's TST instruction.

MIRROR is particularly valuable for games. Without it, every horizontally flipped sprite either needs a pre-flipped copy in memory (doubling sprite data) or a 256-byte bit-reversal lookup table plus per-byte table lookups. With it, you can flip sprites on the fly at 8T per byte.

Barrel Shifts (the multi-bit shift problem)

On the standard Z80, shifting a 16-bit value by more than one bit requires a loop: `SLA E : RL D` per bit, costing 16T per bit position. Shifting DE left by 5 bits costs 80T. The Z80N adds barrel shift instructions that shift DE by an arbitrary number of positions (specified in B) in constant time.

Instruction	Bytes	T	What it replaces
BSLA DE,B	2	8	Shift DE left by B bits. Replaces B * (SLA E : RL D).
BSRA DE,B	2	8	Arithmetic shift DE right by B bits (sign-extending).
BSRL DE,B	2	8	Logical shift DE right by B bits (zero-filling).
BSRF DE,B	2	8	Shift DE right by B bits, filling with bit 15.
BRLC DE,B	2	8	Rotate DE left by B bits (circular).

These are enormously useful in fixed-point arithmetic, pixel sub-positioning, and any code that converts between integer scales. A common operation like “multiply by 5 and shift right 3” that would take ~50T on a standard Z80 becomes trivial.

Convenience Instructions

Instruction	Bytes	T	What it replaces
PUSH nn	4	23	Push a 16-bit immediate value onto the stack. No register needed. Saves the LD rr, nn : PUSH rr pattern (21T, 4 bytes on the original Z80 — same byte count, 2T <i>slower</i> , but does not clobber a register pair). When the register pair is already in use, PUSH nn saves a PUSH/POP pair around the LD+PUSH, which more than compensates.
ADD HL,A	2	8	Add A to HL. Replaces the 5-instruction, 23T sequence: ADD A,L : LD L,A : ADC A,H : SUB L : LD H,A (or the equivalent using a spare register).
ADD DE,A	2	8	Same as ADD HL,A but for DE.
ADD BC,A	2	8	Same for BC.
NEXTREG reg,val	4	20	Direct write to a Next hardware register. No I/O port setup needed. Replaces LD BC,\$243B : OUT (C),reg : LD BC,\$253B : OUT (C),val — four instructions, 8 bytes, ~48T.
NEXTREG reg,A	3	17	Write A to a Next hardware register.
OUTINB	2	16	OUT (C),(HL) : INC HL combined. Useful for streaming data to I/O ports.

The Big Picture

The Z80N is, in a real sense, thirty years of demoscene frustration cast in silicon. Each instruction is a scar from a specific, well-documented pain:

- **PIXELDN** exists because every Spectrum programmer has written `DOWN_HL` at least once, debugged the third-boundary case at least twice, and wished they never had to again.
- **MIRROR** exists because every game programmer has wasted 256 bytes on a bit-reversal table for horizontal sprite flips.
- **LDIRX** exists because the masked blit inner loop is where most Spectrum games spend most of their CPU time.
- **MUL D,E** exists because the shift-and-add multiply loop is the single most re-implemented subroutine in Z80 history.

Unlike the eZ80 (designed by Zilog for embedded markets) or the R800 (designed by ASCII Corporation for the MSX platform), the Z80N was designed by the community, for the community. The instruction set reads like a demoscene wish list because it is a demoscene wish list — the Next team solicited input from active Spectrum coders and prioritised the instructions that would remove the most pain from the most common operations.

3. eZ80 — The Enterprise Extension

The eZ80 is Zilog's official successor to the Z80, designed for embedded systems that need more than 64 KB of address space. It is a strict superset of the Z80 — every Z80 opcode is valid and behaves identically. The extensions are architectural rather than computational:

- **24-bit addressing and ADL mode.** Registers can be 16-bit or 24-bit depending on the operating mode. ADL mode (Address Data Long) gives you 24-bit registers and a flat 16 MB address space. Z80-compatible mode behaves exactly like a standard Z80 with MBASE providing the missing upper 8 address bits.
- **MLT rr — 8x8 unsigned multiply.** `MLT BC` multiplies B by C and stores the 16-bit result in BC. Similarly for `MLT DE` ($D * E \rightarrow DE$) and `MLT HL` ($H * L \rightarrow HL$). This is more flexible than the Z80N's `MUL D,E`, which only operates on DE. The eZ80 gives you three independent multiply units. At 6 T-states on the eZ80 (running at 18.432 MHz on the Agon), this is blazing fast.
- **LEA and PEA.** Load Effective Address and Push Effective Address — indexed address computation instructions. `LEA rr, IX+d` loads the computed address into a register pair without accessing memory. `PEA IX+d` pushes the computed address onto the stack. Useful for parameter passing and pointer arithmetic.
- **TST (test) and TSTIO.** Non-destructive AND test, similar to the Z80N's `TEST nn`. `TSTIO` tests an I/O port value against a mask.
- **IN0/OUT0.** I/O access to the internal peripheral space (addresses 00 — 0F).

The eZ80 was designed for industrial control, networking equipment, and printers — not retrocomputing. But it landed in the Agon Light 2, and suddenly a CPU

designed for embedded markets became a retro gaming platform. The full eZ80 reference is in Appendix E; the porting story is in Chapter 22.

4. R800 — The MSX turboR Speedster

The R800 is the oddest member of the family. Developed by ASCII Corporation for the MSX turboR (1990, Panasonic FS-A1GT/FS-A1ST), it is Z80-compatible in the sense that it executes the full Z80 instruction set — but its internal architecture is radically different.

Pipeline, not microcode. The original Z80 is microcoded: each instruction is broken into machine cycles (M-cycles) of 3-6 clock ticks each, and complex instructions take many M-cycles. The R800 uses a pipelined design where most instructions complete in 1-2 clock cycles. At 7.16 MHz, this gives it an effective throughput roughly 5-8x faster than a 3.5 MHz Z80 for typical code.

Hardware multiply. The R800 adds two multiply instructions:

Instruction	Operands	Result	Cycles	Notes
MULUB A, r	A * r (8-bit unsigned)	HL = 16-bit product	14	r = B, C, D, E
MULUW HL, rr	HL * rr (16-bit unsigned)	DE:HL = 32-bit product	36	rr = BC, SP

A 16x16 multiply with a 32-bit result in 36 cycles is remarkable for an 8-bit-era CPU. On a standard Z80, a 16x16 multiply takes 600-1000 T-states depending on implementation. The R800 makes 3D transformations, DSP-style filtering, and other multiply-heavy algorithms genuinely practical.

The pipeline trap. Code optimised for Z80 timing can behave unexpectedly on the R800. The Z80 optimisation trick of unrolling LDIR into individual LDI instructions (saving 5T per byte) actually runs *slower* on the R800, because the R800's pipeline handles the LDIR repeat prefix efficiently. Similarly, self-modifying code — a staple of Z80 demoscene technique — can stall the R800's pipeline when a write hits a prefetched instruction. Code that is fast on the Z80 is not necessarily fast on the R800, and vice versa.

Demoscene presence. The MSX turboR has a tiny but dedicated demoscene. The hardware multiply makes real-time 3D feasible, and the raw speed allows effects that would be impossible at Z80 clock rates. But the platform's rarity (only sold in Japan, small production run) means the R800 remains a footnote in the broader Z80 story.

5. Soviet Clones — Behind the Iron Curtain

The Soviet Union produced several Z80 clones to circumvent Western export restrictions. These chips enabled an entire ecosystem of ZX Spectrum-compatible

computers that flourished from the late 1980s through the 1990s — and whose demoscene community remains active today.

KR1858VM1 (KP1858BM1). The primary Soviet Z80 clone. Pin-compatible, instruction-compatible, bug-compatible. Manufactured at the Angstrom factory in Zelenograd using reverse-engineered masks. The KR1858VM1 powered the Pentagon 128 and Scorpion ZS-256 — the two most important Soviet Spectrum clones, and the platforms where much of the modern Russian/CIS ZX demoscene still targets its work.

T34VM1 (T34BM1). A later CMOS version of the same design, with lower power consumption and slightly different electrical characteristics. Functionally identical to the KR1858VM1.

Neither chip adds any new instructions. They are exact functional replicas of the Zilog Z80. The differences are electrical: different fabrication processes, different timing margins on setup and hold, slightly different behaviour on undocumented opcodes and flag bits. For software purposes, code that runs on a Zilog Z80 runs identically on a KR1858VM1.

The historical significance is immense. Without these clones, the ZX Spectrum ecosystem would not have spread across the Soviet Union and its successor states. The Pentagon 128, built around the KR1858VM1, became the *de facto* standard Spectrum platform in Russia, and its uncontended timing (no ULA contention delays) is the reference timing used throughout this book.

6. Comparison Table

Feature	Z80	Z80N	eZ80	R800
Clock (typical)	3.5 MHz	28 MHz	18.4 MHz	7.16 MHz
Address space	64 KB	64 KB + Next regs	16 MB	64 KB
Hardware multiply	No	MUL D,E (8T)	MLT rr (6T)	MULUB (14T), MULUW (36T)
16x16 multiply	No	No	No	MULUW (36T, 32-bit result)
Barrel shift	No	BSLA/BSRA/BSRL/BSRR/ DE,B (8T)	BSR/BRLC	No
Block copy with mask	No	LDIRX, LDPIRX	No	No
Screen address helpers	No	PIXELDN, PIXELAD, SETAE	No	No
Bit reverse	No	MIRROR (8T)	No	No
Nibble swap	No	SWAPNIB (8T)	No	No
24-bit mode	No	No	ADL mode	No
Push immediate	No	PUSH nn (23T)	No	No
Add 8-bit to 16-bit	No	ADD HL/DE/BC, A (8T)	No	No

Feature	Z80	Z80N	eZ80	R800
Test (non-destructive AND)	No	TEST nn (11T)	TST (7T)	No
Hardware register I/O	No	NEXTREG (17-20T)	IN0/OUT0	No
Designed for	General computing	ZX Spectrum Next	Embedded systems	MSX turboR

Effective Multiply Performance

Since all four variants run at different clock speeds, the raw T-state counts do not tell the full story. Here is the wall-clock time for an 8x8 unsigned multiply on each platform:

Variant	Clock	Instruction	Cycles	Wall-clock time
Z80	3.5 MHz	Shift-and-add loop	~200	~57 us
Z80N	28 MHz	MUL D,E	8	~0.29 us
eZ80	18.4 MHz	MLT DE	6	~0.33 us
R800	7.16 MHz	MULUB A,r	14	~1.96 us

The Z80N and eZ80 are effectively tied for multiply performance. The R800 is 30x faster than a standard Z80 but 6-7x slower than the Z80N/eZ80. All three are fast enough to make real-time 3D practical.

7. What This Means for the Book

All assembly code in this book targets the **standard Z80 instruction set**. Every example in every chapter will assemble and run on a stock ZX Spectrum 48K, a Pentagon 128, a Scorpion, an MSX, or any other machine with a Zilog Z80 or compatible clone. No extensions required.

This is a deliberate choice. The optimisation principles — T-state budgets, register allocation, loop structure, self-modifying code, unrolled loops, stack tricks — are universal. Code that is fast on a 3.5 MHz Z80 is fast on a 28 MHz Z80N. Code that fits in 48 KB fits in 16 MB. The constraints of the original Z80 teach you to think in a way that transfers to every variant in the family.

That said, if you have a ZX Spectrum Next, the Z80N extensions are too good to ignore. Chapter 22 covers porting strategies including Z80N-specific optimisations. If you have an Agon Light 2, Appendix E is your eZ80 reference and Chapter 22 walks through a complete Spectrum-to-Agon port. The barrel shifts, hardware multiply, and masked blit instructions do not change *how* you think about optimisation — they change *where the bottleneck moves* once the classic pain points are eliminated.

The fundamentals do not change. The instructions get better.

See Also

- **Appendix A: Z80 Instruction Quick Reference** — complete standard Z80 instruction table with T-states, byte counts, and flag effects. The baseline that all variants share.
- **Appendix E: eZ80 Quick Reference** — full eZ80 reference including mode system, MLT, LEA/PEA, and Agon Light 2 specifics.
- **Chapter 22: Porting — Agon Light 2** — practical porting walkthrough covering both eZ80 and Z80N extensions in context.

T-state audit note: Z80N T-state values in this appendix were corrected using the official documentation at wiki.specnext.dev. The initial values (approximately halved across the board) were identified as incorrect by Ped7g (Peter Helcmanovsky).

Sources: Zilog Z80 CPU User Manual (UM0080); Zilog eZ80 CPU User Manual (UM0077); ZX Spectrum Next User Manual, Issue 2; ZX Spectrum Next Extended Instruction Set Documentation (wiki.specnext.dev); Victor Trucco, Fabio Belavenuto et al., Z80N instruction set design notes; ASCII Corporation R800 Technical Reference (1990); Sean Young, “The Undocumented Z80 Documented” (2005); Introspec, “Once more about DOWN_HL” (Hype, 2020); Dark / X-Trade, “Programming Algorithms” (Spectrum Expert #01, 1997)

Appendix G: AY-3-8910 / TurboSound / Triple AY Register Reference

“Fourteen registers. Three square-wave tone channels. One noise generator. One envelope generator. That is all you get.” – Chapter 11

This appendix is a complete register-level reference for the AY-3-8910 Programmable Sound Generator as used in the ZX Spectrum 128K, TurboSound-equipped clones (Pentagon, Scorpion), and the ZX Spectrum Next’s triple-AY configuration. Chapter 11 covers the musical techniques; this appendix is the datasheet you keep open while coding.

All hex values use the \$FF notation. Binary values use %10101010. Clock frequency assumes the PAL Spectrum 128K standard of 1.7734 MHz unless otherwise noted.

I/O Ports on ZX Spectrum 128K

Port	Direction	Function
<i>FFFD</i> *	Read	Read value of currently selected register
* <i>Write</i> <i>Selectactive</i> register(0–15) * * FFFD		
\$BFFD	Write	Write data to selected register

The Write Sequence

Every AY register write is a two-step operation: select the register, then write the value.

```
; Clobbers: BC
ay_write:
    ld    bc, $FFFD
    out   (c), a        ; 12T  select register
    ld    b, $BF        ; 7T   BC = $BFFD (C stays $FD)
    out   (c), e        ; 12T  write data
    ret                               ; total: 31T + call overhead
```

The trick: only the high byte of BC changes between the two OUT instructions. The low byte \$FD stays in C throughout. This saves 3 bytes and 10 T-states compared to loading a full 16-bit value twice.

Reading a Register

```
; Clobbers: BC
ay_read:
    ld    bc, $FFFD
    out   (c), a        ; select register
    in    a, (c)        ; read value
    ret
```

Reading is useful for preserving the mixer state (R7) when modifying individual channel enables.

Bulk Register Write

For music engines that update all 14 registers per frame, an unrolled loop is fastest:

```
; HL = pointer to 14-byte register buffer (R0-R13)
ay_flush:
    ld    de, $BFFD      ; D = $BF, E = $FD
    ld    c, e           ; C = $FD (shared low byte)
    ld    b, $FF         ; BC = $FFFD (register select port)
    xor   a              ; start at R0
.loop:
    out   (c), a         ; select register A
    ld    b, d           ; BC = $BFFD
    outi                ; write (HL) to port, inc HL, dec B
    ; B is now $BE after OUTI dec, but we reload it anyway
    ld    b, $FF         ; BC = $FFFD
    inc   a
    cp    14
    jr    nz, .loop
    ret
```

Complete Register Map

Overview

Reg	Name	Bits Used	R/W	Reset	Description
R0	Tone A period, low	7-0	R/W	\$00	Channel A tone period, bits 7-0
R1	Tone A period, high	3-0	R/W	\$00	Channel A tone period, bits 11-8

Reg	Name	Bits Used	R/W	Reset	Description
R2	Tone B period, low	7-0	R/W	\$00	Channel B tone period, bits 7-0
R3	Tone B period, high	3-0	R/W	\$00	Channel B tone period, bits 11-8
R4	Tone C period, low	7-0	R/W	\$00	Channel C tone period, bits 7-0
R5	Tone C period, high	3-0	R/W	\$00	Channel C tone period, bits 11-8
R6	Noise period	4-0	R/W	\$00	Noise generator period (0-31)
R7	Mixer / I/O	7-0	R/W	\$FF	Tone/noise enable + I/O port direction
R8	Volume A	4-0	R/W	\$00	Channel A volume or envelope mode
R9	Volume B	4-0	R/W	\$00	Channel B volume or envelope mode
R10	Volume C	4-0	R/W	\$00	Channel C volume or envelope mode
R11	Envelope period, low	7-0	R/W	\$00	Envelope period, bits 7-0
R12	Envelope period, high	7-0	R/W	\$00	Envelope period, bits 15-8
R13	Envelope shape	3-0	W	-	Envelope waveform (writing restarts envelope)
R14	I/O Port A	7-0	R/W	-	General-purpose I/O (directly mapped on AY-3-8910)
R15	I/O Port B	7-0	R/W	-	General-purpose I/O (directly mapped on AY-3-8910)

Note on R14-R15: The AY-3-8910 has two 8-bit I/O ports. On the ZX Spectrum 128K, I/O port A (R14) is active - it reads the keyboard matrix and other hardware signals. R15 (I/O port B) is typically not connected. The AY-3-8912 (used in some clones) has only port A; the AY-3-8913 has no I/O ports at all. For sound programming, R14 and R15 are irrelevant.

R0-R1: Channel A Tone Period (12-bit)

R0: [D7] [D6] [D5] [D4] [D3] [D2] [D1] [D0] bits 7-0 of period

R1: [0] [0] [0] [0] [D11][D10][D9][D8] bits 11-8 of period

- **Range:** 1 to 4095 (12-bit unsigned). Period of 0 behaves as 1 on most implementations.
- **Effect:** Sets the frequency of the square wave on channel A.
- **Formula:** frequency = 1773400 / (16 * period) (PAL Spectrum 128K)
- **Practical range:** Period 1 = 110,837 Hz (ultrasonic), period 4095 = 27 Hz (deep bass).

```
ld  a, 0           ; R0 = Tone A low
ld  e, $A8         ; 424 & $FF = $A8
call ay_write
ld  a, 1           ; R1 = Tone A high
ld  e, $01         ; 424 >> 8 = $01
call ay_write
```

R2-R3: Channel B Tone Period (12-bit)

Identical layout to R0-R1, for channel B.

R4-R5: Channel C Tone Period (12-bit)

Identical layout to R0-R1, for channel C.

R6: Noise Period (5-bit)

R6: [0] [0] [0] [D4] [D3] [D2] [D1] [D0]

- **Range:** 0 to 31. A single noise generator shared by all three channels.
- **Effect:** Lower values = higher-pitched noise (hiss). Higher values = lower-pitched, rougher noise.
- **Formula:** noise_frequency = 1773400 / (16 * period) (same as tone formula)

R6 Value	Character	Typical Use
0-5	High hiss	Hi-hat, cymbal, metallic ring
6-12	Mid noise	Snare drum body, white noise
13-20	Low rumble	Explosion, engine, wind
21-31	Very low	Thunder, distant rumble

The noise generator uses a 17-bit linear feedback shift register (LFSR) to produce pseudo-random output. The output is the same for both the AY-3-8910 and YM2149, but the LFSR tap positions differ, producing subtly different noise textures on each chip.

R7: Mixer Control (The Most Important Register)

R7: [IOB] [IOA] [NC] [NB] [NA] [TC] [TB] [TA]
 bit7 bit6 bit5 bit4 bit3 bit2 bit1 bit0

Bit	Name	0 =	1 =
7	I/O port B direction	Output	Input
6	I/O port A direction	Output	Input
5	Noise C enable	ON	Off
4	Noise B enable	ON	Off
3	Noise A enable	ON	Off
2	Tone C enable	ON	Off
1	Tone B enable	ON	Off
0	Tone A enable	ON	Off

Critical: 0 means ON. The mixer uses active-low logic. A cleared bit enables the corresponding source. This is the single most confusing aspect of AY programming.

Bits 7-6: Always set to 1 (input mode) on the Spectrum. Do not change these unless you know what you are doing with the I/O ports.

Common Mixer Values

Value	Binary (NC NB NA TC TB TA)	Effect
\$38	%00 111 000	All three tones, no noise
\$3E	%00 111 110	Tone A only
\$3D	%00 111 101	Tone B only
\$3B	%00 111 011	Tone C only
\$36	%00 110 110	Tone A + Noise A
\$2D	%00 101 101	Tone B + Noise B
\$1B	%00 011 011	Tone C + Noise C
\$28	%00 101 000	All tones + Noise C (music + drums on C)
\$07	%00 000 111	All noise, no tones
\$00	%00 000 000	Everything on
\$3F	%00 111 111	Everything off (silence)

Note: The binary values above show bits 5-0 only. Bits 7-6 should be %11 for normal operation (I/O ports as inputs), making e.g. \$38 actually %11 111 000 = \$F8. However, on the Spectrum 128K the unused upper bits are ignored by convention, and most code uses the short form. Some engines write the full \$F8 form; both work identically.

```
; Binary: I/O=11, NC=0(on), NB=1, NA=1, TC=0(on), TB=0(on), TA=0(on)
; = %11 011 000 = $D8 (full form) or $28 (short form, bits 7-6 treated as 0)
ld  a, 7                ; R7 = Mixer
ld  e, $28
call ay_write
```

R8-R10: Volume Registers (5-bit)

R8: [0] [0] [0] [M] [D3] [D2] [D1] [D0] Channel A
 R9: [0] [0] [0] [M] [D3] [D2] [D1] [D0] Channel B
 R10: [0] [0] [0] [M] [D3] [D2] [D1] [D0] Channel C

Bit	Function
4 (M)	Mode: 0 = use fixed volume (bits 3-0), 1 = use envelope generator
3-0	Fixed volume level: 0 (silent) to 15 (maximum)

- **Fixed volume mode** (bit 4 = 0): Bits 3-0 set the channel volume directly. 15 = loudest, 0 = silent. The volume curve is logarithmic on a genuine AY-3-8910 (approximately 1.5 dB per step) but linear on the YM2149.
- **Envelope mode** (bit 4 = 1): The channel's volume is controlled by the envelope generator (R11-R13). Bits 3-0 are ignored. Only one envelope generator exists, shared by all channels in envelope mode.

Value	Effect
\$00	Silent
\$08	Half volume (approximately)
\$0F	Maximum fixed volume
\$10	Envelope mode (any value \$10-\$1F activates envelope)

4-bit DAC trick: By rapidly writing successive sample values to a volume register (without envelope mode), you can play digitised audio through the AY. At 8 kHz sample rate, this consumes approximately 437 T-states per sample write, leaving almost no CPU for other work. See Chapter 12 for the hybrid drum technique.

```
ld  a, 8          ; R8 = Volume A
ld  e, $0F        ; volume 15
call ay_write

; Switch channel C to envelope mode
ld  a, 10         ; R10 = Volume C
ld  e, $10        ; bit 4 set = envelope mode
call ay_write
```

R11-R12: Envelope Period (16-bit)

R11: [D7] [D6] [D5] [D4] [D3] [D2] [D1] [D0] bits 7-0
 R12: [D7] [D6] [D5] [D4] [D3] [D2] [D1] [D0] bits 15-8

- **Range:** 1 to 65535 (16-bit unsigned). Period of 0 behaves as 1.
- **Formula:** $\text{envelope_frequency} = 1773400 / (256 * \text{period})$
- **At period = 1:** ~6,927 Hz (one complete envelope cycle in ~144 microseconds)
- **At period = 65535:** ~0.106 Hz (one cycle in ~9.4 seconds)

The envelope period controls how fast the volume ramps up or down according to the shape selected in R13. For buzz-bass, the envelope period replaces the tone period as the pitch control:

$\text{bass_envelope_period} = 1773400 / (256 * \text{desired_frequency})$

For bass C2 (65.4 Hz): $\text{period} = 1773400 / (256 * 65.4) = 106$.

```
ld  a, 11          ; R11 = Envelope period low
ld  e, 106         ; $6A
call ay_write
ld  a, 12          ; R12 = Envelope period high
ld  e, 0
call ay_write
```

R13: Envelope Shape (4-bit, Write-Only)

R13: [0] [0] [0] [0] [CONT] [ATT] [ALT] [HOLD]
 bit3 bit2 bit1 bit0

Bit	Name	Function
3	CONT	Continue: 0 = single-shot (reset to 0 or hold), 1 = repeating
2	ATT	Attack: 0 = start at max, ramp down, 1 = start at 0, ramp up
1	ALT	Alternate: 0 = same direction each cycle, 1 = reverse direction
0	HOLD	Hold: 0 = cycle normally, 1 = hold at final level after first cycle

Critical behaviour: Writing ANY value to R13 immediately restarts the envelope from the beginning of the cycle. This is true even if you write the same value already stored. This restart behaviour is essential for triggering buzz-bass notes and drum sounds.

All 16 Shape Values

Although 4 bits allow 16 values, many produce identical output. There are only 10 unique shapes:

Value	Bits (CONT ATT ALT HOLD)	Shape	Behaviour
\$00	0 0 0 0	__	Decay to 0, hold at 0
\$01	0 0 0 1	__	Same as \$00
\$02	0 0 1 0	__	Same as \$00
\$03	0 0 1 1	__	Same as \$00
\$04	0 1 0 0	/__	Attack to 15, drop to 0, hold at 0

Value	Bits (CONT ATT ALT HOLD)	Shape	Behaviour
\$05	0 1 0 1	/___	Same as \$04
\$06	0 1 1 0	/___	Same as \$04
\$07	0 1 1 1	/___	Same as \$04
\$08	1 0 0 0	\\ \\	Repeating sawtooth down
\$09	1 0 0 1	___	Single decay, hold at 0 (same as \$00)
\$0A	1 0 1 0	\\ \ /	Repeating triangle (down-up)
\$0B	1 0 1 1	\ ^ ^	Single decay, then hold at max (15)
\$0C	1 1 0 0	/// /	Repeating sawtooth up
\$0D	1 1 0 1	/ ^ ^	Single attack, then hold at max (15)
\$0E	1 1 1 0	/ \ \	Repeating triangle (up-down)
\$0F	1 1 1 1	/___	Single attack, drop to 0, hold at 0 (same as \$04)

Envelope Shape Waveforms


\$00-\$03, \$09:




\$04-\$07, \$0F:




\$08 (repeating \\ \\):




\$0C (repeating /// /):



\$0A (repeating \\ \ /):



\$0E (repeating / \ \):



\$0B (decay, hold max): \$0D (attack, hold max):

```

15 | \      15  /|
   |  \      / |
   |   \    /  |
   |    \  /   |
0  |      |    0  |

```

Practical Envelope Uses

Shape	Value	Use Case
\$00	__	Drum decay: sharp hit that fades to silence
\$08	\\\\	Buzz-bass: thick repeating bass tone
\$0C	////	Buzz-bass (inverted phase): same pitch, different timbre
\$0A	\\\/	Metallic/bell tone (triangle modulation)
\$0E	/\\	Metallic/bell tone (inverted triangle)
\$0D	/^^	Fade-in: volume rises to max and holds

Tone Period to Frequency Conversion

Formula

frequency = AY_clock / (16 * period)
period = AY_clock / (16 * frequency)

AY Clock by Platform

Platform	AY Clock	CPU Clock	Relationship
ZX Spectrum 128K (PAL)	1,773,400 Hz	3,546,900 Hz	AY = CPU / 2
ZX Spectrum 128K (NTSC)	1,789,772 Hz	3,579,545 Hz	AY = CPU / 2
Pentagon 128	1,750,000 Hz	3,500,000 Hz	AY = CPU / 2
Amstrad CPC	1,000,000 Hz	4,000,000 Hz	AY = CPU / 4
MSX	1,789,772 Hz	3,579,545 Hz	AY = CPU / 2
Atari ST (YM2149)	2,000,000 Hz	8,000,000 Hz	YM = CPU / 4
ZX Spectrum Next	1,773,400 Hz	varies	Same as 128K

Practical consequence: The same period value produces slightly different pitches on different platforms. A tune composed on Pentagon (1.75 MHz) will sound fractionally sharp on a Spectrum 128K (1.7734 MHz). For most musical purposes, the difference is inaudible.

Complete Note Frequency Table

All period values calculated for AY clock = 1,773,400 Hz (PAL Spectrum 128K).

Octaves 1-2 cover the bass range (buzz-bass territory). Octaves 3-6 are the main melodic range. Octaves 7-8 are high-pitched and increasingly inaccurate due to integer rounding.

Octave 1 (Deep Bass)

Note	Freq (Hz)	Period	R_LO	R_HI
C-1	32.70	3390	\$3E	\$0D
C#1	34.65	3199	\$7F	\$0C
D-1	36.71	3019	\$CB	\$0B
D#1	38.89	2850	\$22	\$0B
E-1	41.20	2690	\$82	\$0A
F-1	43.65	2539	\$EB	\$09
F#1	46.25	2396	\$5C	\$09
G-1	49.00	2262	\$D6	\$08
G#1	51.91	2135	\$57	\$08
A-1	55.00	2015	\$DF	\$07
A#1	58.27	1902	\$6E	\$07
B-1	61.74	1795	\$03	\$07

Octave 2 (Bass)

Note	Freq (Hz)	Period	R_LO	R_HI
C-2	65.41	1695	\$9F	\$06
C#2	69.30	1599	\$3F	\$06
D-2	73.42	1510	\$E6	\$05
D#2	77.78	1425	\$91	\$05
E-2	82.41	1345	\$41	\$05
F-2	87.31	1270	\$F6	\$04
F#2	92.50	1198	\$AE	\$04
G-2	98.00	1131	\$6B	\$04
G#2	103.83	1067	\$2B	\$04
A-2	110.00	1008	\$F0	\$03
A#2	116.54	951	\$B7	\$03
B-2	123.47	897	\$81	\$03

Octave 3

Note	Freq (Hz)	Period	R_LO	R_HI
C-3	130.81	847	\$4F	\$03
C#3	138.59	800	\$20	\$03
D-3	146.83	755	\$F3	\$02
D#3	155.56	713	\$C9	\$02
E-3	164.81	673	\$A1	\$02

Note	Freq (Hz)	Period	R_LO	R_HI
F-3	174.61	635	\$7B	\$02
F#3	185.00	599	\$57	\$02
G-3	196.00	566	\$36	\$02
G#3	207.65	534	\$16	\$02
A-3	220.00	504	\$F8	\$01
A#3	233.08	475	\$DB	\$01
B-3	246.94	449	\$C1	\$01

Octave 4 (Middle Octave)

Note	Freq (Hz)	Period	R_LO	R_HI
C-4	261.63	424	\$A8	\$01
C#4	277.18	400	\$90	\$01
D-4	293.66	378	\$7A	\$01
D#4	311.13	357	\$65	\$01
E-4	329.63	337	\$51	\$01
F-4	349.23	318	\$3E	\$01
F#4	369.99	300	\$2C	\$01
G-4	392.00	283	\$1B	\$01
G#4	415.30	267	\$0B	\$01
A-4	440.00	252	\$FC	\$00
A#4	466.16	238	\$EE	\$00
B-4	493.88	225	\$E1	\$00

Octave 5

Note	Freq (Hz)	Period	R_LO	R_HI
C-5	523.25	212	\$D4	\$00
C#5	554.37	200	\$C8	\$00
D-5	587.33	189	\$BD	\$00
D#5	622.25	178	\$B2	\$00
E-5	659.26	168	\$A8	\$00
F-5	698.46	159	\$9F	\$00
F#5	739.99	150	\$96	\$00
G-5	783.99	142	\$8E	\$00
G#5	830.61	134	\$86	\$00
A-5	880.00	126	\$7E	\$00
A#5	932.33	119	\$77	\$00
B-5	987.77	112	\$70	\$00

Octave 6

Note	Freq (Hz)	Period	R_LO	R_HI
C-6	1046.50	106	\$6A	\$00
C#6	1108.73	100	\$64	\$00

Note	Freq (Hz)	Period	R_LO	R_HI
D-6	1174.66	94	\$5E	\$00
D#6	1244.51	89	\$59	\$00
E-6	1318.51	84	\$54	\$00
F-6	1396.91	79	\$4F	\$00
F#6	1479.98	75	\$4B	\$00
G-6	1567.98	71	\$47	\$00
G#6	1661.22	67	\$43	\$00
A-6	1760.00	63	\$3F	\$00
A#6	1864.66	59	\$3B	\$00
B-6	1975.53	56	\$38	\$00

Octave 7 (High)

Note	Freq (Hz)	Period	R_LO	R_HI
C-7	2093.00	53	\$35	\$00
C#7	2217.46	50	\$32	\$00
D-7	2349.32	47	\$2F	\$00
D#7	2489.02	45	\$2D	\$00
E-7	2637.02	42	\$2A	\$00
F-7	2793.83	40	\$28	\$00
F#7	2959.96	37	\$25	\$00
G-7	3135.96	35	\$23	\$00
G#7	3322.44	33	\$21	\$00
A-7	3520.00	31	\$1F	\$00
A#7	3729.31	30	\$1E	\$00
B-7	3951.07	28	\$1C	\$00

Octave 8 (Very High - Limited Accuracy)

Note	Freq (Hz)	Period	R_LO	R_HI	Actual Freq	Error (cents)
C-8	4186.01	26	\$1A	\$00	4264.23	+32
C#8	4434.92	25	\$19	\$00	4433.50	-1
D-8	4698.63	24	\$18	\$00	4618.23	-30
D#8	4978.03	22	\$16	\$00	5038.07	+21
E-8	5274.04	21	\$15	\$00	5278.27	+1
F-8	5587.65	20	\$14	\$00	5541.88	-14
F#8	5919.91	19	\$13	\$00	5833.55	-26
G-8	6271.93	18	\$12	\$00	6158.75	-32

Note: Above octave 6, integer rounding of the period value produces increasingly audible pitch errors. The “Actual Freq” and “Error” columns for octave 8 show the real output frequency and deviation in cents. For high notes, fine-tuning is not possible – the resolution is simply too coarse.

Compact Note Table for Z80 Code

Most tracker engines store the period table as 16-bit words. Here is a practical 96-note table (octaves 1-8) for inclusion in assembly source:

```
; Each entry is a 16-bit period value (low byte first)
; Index = (octave * 12) + semitone, where C=0, C#=1, ..., B=11
note_table:
    ; Octave 1
    DW 3390, 3199, 3019, 2850, 2690, 2539
    DW 2396, 2262, 2135, 2015, 1902, 1795
    ; Octave 2
    DW 1695, 1599, 1510, 1425, 1345, 1270
    DW 1198, 1131, 1067, 1008, 951, 897
    ; Octave 3
    DW 847, 800, 755, 713, 673, 635
    DW 599, 566, 534, 504, 475, 449
    ; Octave 4
    DW 424, 400, 378, 357, 337, 318
    DW 300, 283, 267, 252, 238, 225
    ; Octave 5
    DW 212, 200, 189, 178, 168, 159
    DW 150, 142, 134, 126, 119, 112
    ; Octave 6
    DW 106, 100, 94, 89, 84, 79
    DW 75, 71, 67, 63, 59, 56
    ; Octave 7
    DW 53, 50, 47, 45, 42, 40
    DW 37, 35, 33, 31, 30, 28
    ; Octave 8
    DW 26, 25, 24, 22, 21, 20
    DW 19, 18, 17, 16, 15, 14

; Returns DE = period value
; Clobbers: HL
get_note_period:
    ld  h, 0
    ld  l, a
    add hl, hl          ; HL = note * 2 (word index)
    ld  de, note_table
    add hl, de
    ld  e, (hl)
    inc hl
    ld  d, (hl)          ; DE = period
    ret
```

Table #5: Natural Tuning (Just Intonation)

The standard note table above uses equal temperament (12-TET) – every semitone is the 12th root of 2 apart. This works well for tone channels, but creates a problem for **buzz-bass (T+E)**: since $\text{envelope_period} = \text{tone_period} / 16$, any tone period not divisible by 16 introduces a rounding error. The envelope drifts against the tone, producing audible beating on sustained bass notes.

Ivan Roshin's "Frequency Table with Zero Error" (2001) and oisee's VTi implementation (2009) solve this by using **just intonation** – integer-ratio intervals for C major / A minor:

C [9/8] D [10/9] E [16/15] F [9/8] G [10/9] A [9/8] B [16/15] C

This produces pure fifths (C-G, E-B, A-E at exact 3:2) and, critically, periods where most main-scale notes divide evenly by 16.

AY clock: 1,520,640 Hz (non-standard; select per-key frequency below):

; 96 notes, C-1 to B-8, for AY clock = 1,520,640 Hz
; C major / A minor only; other keys via chip frequency change

natural_note_table:

```
; Octave 1
DW 2880, 2700, 2560, 2400, 2304, 2160
DW 2025, 1920, 1800, 1728, 1620, 1536
; Octave 2
DW 1440, 1350, 1280, 1200, 1152, 1080
DW 1013, 960, 900, 864, 810, 768
; Octave 3
DW 720, 675, 640, 600, 576, 540
DW 506, 480, 450, 432, 405, 384
; Octave 4
DW 360, 338, 320, 300, 288, 270
DW 253, 240, 225, 216, 203, 192
; Octave 5
DW 180, 169, 160, 150, 144, 135
DW 127, 120, 113, 108, 101, 96
; Octave 6
DW 90, 84, 80, 75, 72, 68
DW 63, 60, 56, 54, 51, 48
; Octave 7
DW 45, 42, 40, 38, 36, 34
DW 32, 30, 28, 27, 25, 24
; Octave 8
DW 23, 21, 20, 19, 18, 17
DW 16, 15, 14, 14, 13, 12
```

Period divisibility check (Octave 2, bass range):

Note	Period	mod 16	Env Period	Clean T+E?
C2	1440	0	90	Yes
C#2	1350	6	84	No
D2	1280	0	80	Yes
D#2	1200	0	75	Yes
E2	1152	0	72	Yes
F2	1080	0	67	~
F#2	1013	5	63	No
G2	960	0	60	Yes
G#2	900	4	56	No
A2	864	0	54	Yes
A#2	810	2	50	No

Note	Period	mod 16	Env Period	Clean T+E?
B2	768	0	48	Yes

Seven of twelve notes (all natural notes of C major) divide cleanly - compared to *zero* in the equal-tempered table.

Transposition via chip frequency: since the table is fixed to C/Am, other keys require a different AY clock. Each step multiplies by $2^{(1/12)}$:

Key	Chip Frequency (Hz)
C/Am	1,520,640
C#/A#m	1,611,062
D/Bm	1,706,861
D#/Cm	1,808,356
E/C#m	1,915,886
F/Dm	2,029,811
F#/D#m	2,150,510
G/Em	2,278,386
G#/Fm	2,413,866
A/F#m	2,557,401
A#/Gm	2,709,472
B/G#m	2,870,586

On hardware the AY clock is fixed; in trackers (Vortex Tracker II/Improved) and emulators, this is a per-module setting. Table #5 is the default for the autosiril MIDI-to-PT3 converter because most converted tracks use buzz-bass heavily. See Chapter 11 for a detailed explanation of the T+E alignment problem.

TurboSound: 2 x AY

TurboSound is a hardware modification adding a second AY chip, giving 6 tone channels, 2 noise generators, and 2 envelope generators. The most common modern implementation is the NedoPC TurboSound card.

Chip Selection

Both AY chips share the same I/O ports (*FFFD*/*BFFD*). A chip select value written to *\$FFFD* switches all subsequent register operations to the selected chip.

Chip Select Value	Target
\$FF	Chip 0 (primary / original AY)
\$FE	Chip 1 (secondary AY)

```
ld    bc, $FFFD
ld    a, $FF
out   (c), a           ; all subsequent R/W goes to chip 0
```

```

; Select chip 1 (secondary)
    ld    bc, $FFFD
    ld    a, $FE
    out   (c), a           ; all subsequent R/W goes to chip 1

```

Important: The chip select persists until changed. After selecting chip 1, all register operations – including reads – target chip 1. Always explicitly select the chip before any register access in your engine.

TurboSound Engine Architecture

A typical TurboSound music engine maintains two 14-byte register buffers and writes them sequentially:

```

ts_update:
    ; --- Chip 0 ---
    ld    a, $FF
    ld    bc, $FFFD
    out   (c), a           ; select chip 0
    ld    hl, ay0_regs      ; 14-byte buffer for chip 0
    call  ay_flush         ; write all 14 registers

    ; --- Chip 1 ---
    ld    a, $FE
    ld    bc, $FFFD
    out   (c), a           ; select chip 1
    ld    hl, ay1_regs      ; 14-byte buffer for chip 1
    call  ay_flush         ; write all 14 registers
    ret

```

```

ay0_regs: DS 14           ; chip 0 register shadow
ay1_regs: DS 14           ; chip 1 register shadow

```

Total cost: approximately 28 OUT instructions, roughly 700-800 T-states including overhead. This is about 1% of a frame budget – negligible.

Stereo Configuration

TurboSound enables true stereo. The two chips can be hard-panned or mixed:

Configuration	Chip 0 (Left)	Chip 1 (Right)	Character
Wide stereo	Lead, bass (A0+B0), drums (C0)	Counter-melody (A1+B1), pads (C1)	Spacious, concert-like
Centered bass	Bass on C0 + C1 (same data)	Lead on A0, harmony on A1	Solid low end
Split drums	Kick (C0)	Snare + hi-hat (C1)	Punchy, wide percussion

Each chip has its own independent noise generator and envelope generator. This

means you can run buzz-bass on chip 0 and a completely independent envelope percussion on chip 1 without interference – impossible on a single AY.

ZX Spectrum Next: Triple AY (3 x AY)

The ZX Spectrum Next includes three AY-compatible sound chips, providing 9 tone channels, 3 noise generators, and 3 envelope generators.

Chip Selection on Next

The Next extends the TurboSound protocol. The third chip is selected via the same mechanism:

Chip Select Value	Target
\$FF	Chip 0
\$FE	Chip 1
\$FD	Chip 2

The Next also provides access through Next register \$06 (Peripheral 2), which configures the sound chip mode:

Next Reg \$06, bits 1-0	Mode
%00	Single AY (Spectrum 128K compatible)
%01	TurboSound (2 x AY)
%10	Triple AY (3 x AY)

Per-Channel Stereo Panning

The Next adds a feature the original AY never had: per-channel stereo panning. Each of the 9 channels can be individually panned left, right, or centre.

Panning is controlled through AY register 14 (R14), repurposed on the Next as a stereo control register when accessed in the AY chip context:

Bit	Channel	0 =	1 =
0	A, right	off	on
1	A, left	off	on
2	B, right	off	on
3	B, left	off	on
4	C, right	off	on
5	C, left	off	on

Set both bits (left + right) for centre panning. Set only one bit for hard left or hard right.

Common Patterns in Z80 Code

Silence the AY

```
ay_silence:
    xor    a                      ; volume = 0
    ld     e, a
    ld     a, 8                   ; R8 = Volume A
    call  ay_write
    ld     a, 9                   ; R9 = Volume B
    call  ay_write
    ld     a, 10                  ; R10 = Volume C
    call  ay_write
    ; Disable all tone and noise
    ld     a, 7                   ; R7 = Mixer
    ld     e, $3F                 ; all off
    call  ay_write
    ret
```

Play a Single Note on Channel A

; Assumes mixer is already configured for tone A

```
play_note_a:
    ld     a, 0                   ; R0 = Tone A low
    ld     e, d                   ; wait -- let's do this properly
    ; Actually: DE has period, E = low byte, D = high byte
    push  de
    ld     a, 0                   ; R0
    call  ay_write                ; E already has low byte
    pop   de
    ld     e, d
    ld     a, 1                   ; R1
    call  ay_write                ; E = high byte
    ld     a, 8                   ; R8 = Volume A
    ld     e, 12                  ; volume 12
    call  ay_write
    ret
```

Trigger a Buzz-Bass Note

; DE = envelope period for desired pitch

```
buzz_bass:
    ld     a, 11                  ; R11 = Envelope period low
    call  ay_write                ; E = low byte of period
    ld     e, d
    ld     a, 12                  ; R12 = Envelope period high
    call  ay_write
    ld     a, 13                  ; R13 = Envelope shape
    ld     e, $08                 ; repeating sawtooth down
    call  ay_write                ; this also restarts the envelope
    ld     a, 10                  ; R10 = Volume C
    ld     e, $10                 ; envelope mode
```

```

call ay_write
ret

```

Trigger Envelope (Restart)

```

; Useful for re-triggering buzz-bass on a new note
; The shape value must be written even if unchanged
env_restart:
    ld    a, 13                ; R13 = Envelope shape
    ld    e, (hl)              ; shape from current instrument data
    call ay_write              ; writing R13 = instant restart
    ret

```

Digital Drum Hit (4-bit DAC)

```

; HL = pointer to 4-bit sample data (values 0-15)
; B = number of samples to play
; Uses channel A (R8) as DAC
;
; WARNING: This consumes all CPU time during playback.
; Disable interrupts before calling.
play_sample:
    di
    ld    a, 8                  ; select R8 (Volume A)
    ld    bc, $FFFD
    out   (c), a
    ld    c, $FD                ; prepare for $BFFD writes
.loop:
    ld    a, (hl)               ; 7T  load sample
    inc   hl                    ; 6T  advance pointer
    ld    b, $BF                ; 7T  BC = $BFFD
    out   (c), a                ; 12T write volume = sample
    ; Timing padding: adjust NOPs to hit target sample rate
    nop                      ; 4T
    nop                      ; 4T
    ld    b, a                  ; 4T  (dummy, preserves timing)
    ld    b, 0                  ; 7T  reset B for next iteration
    ; Total per sample: ~51T = ~14.6 us = ~68.6 kHz
    ; For 8 kHz, add ~386T of padding (96 NOPs or a delay loop)
    dec   b
    jr    nz, .loop
    ei
    ret

```

See Chapter 12 for the hybrid drum technique (digital attack + AY envelope decay) that makes this practical in a demo engine.

Tracker Format Notes

ProTracker 3 (.pt3)

ProTracker 3 (PT3) is the most widely used tracker format on the ZX Spectrum 128K. The cross-platform editor **Vortex Tracker II** natively produces .pt3 files.

Property	Value
Channels	3 (single AY) or 6 (TurboSound)
Pattern length	1-64 rows
Song speed	1-255 (frames per row)
Instruments	Ornament tables + amplitude envelopes
Effects	Arpeggio, portamento, vibrato, envelope control
Player size	~1.2-1.8 KB (Z80 assembly)
Data format	Packed patterns + instrument/ornament tables

How PT3 Data Maps to AY Registers

On each frame (1/50th second), the PT3 player:

1. Decrements the speed counter. If it reaches zero, advance to the next row.
2. For each channel (A, B, C):
 - Read the note, instrument, and effect from the current row.
 - Apply the instrument's ornament (per-frame pitch offset table).
 - Apply the instrument's amplitude envelope (per-frame volume table).
 - Calculate the final tone period and volume.
3. Update the mixer register (R7) based on which channels use tone, noise, or both.
4. Write all calculated values to R0-R13.

PT3 row data:

Note C-4, Instrument 3, Effect: Arpeggio +4+7

```

      |           |           |
      v           v           v
Period=424  Ornament table  +4 semitones, +7 semitones
      |           |           |
      |           |           |
      v           |           |
R0=$A8, R1=$01 |           |
      <-----+           |
Final period adjusted by ornament offset
      <-----+
Period cycles: 424 -> 337 -> 283 -> 424 -> ...
(C4  -> E4  -> G4  -> C4  -> ...)
```

Integration in Your Project

```

; 1. Include the player source
    INCLUDE "pt3player.asm"

; 2. Include the song data
song_data:
    INCBIN "mysong.pt3"
```



```

; 3. Initialise at startup
    ld    hl, song_data
    call music_init          ; PT3 player init routine

; 4. Call from IM2 handler every frame
isr_handler:
    ; ... push registers ...
    call music_play          ; PT3 player frame routine
    ; ... pop registers ...
    ei
    reti

```

Other Tracker Formats

Format	Tracker	Channels	Player Size	Key Feature
.pt3	Pro Tracker 3 / Vortex Tracker II	3-6	~1.2-1.8 KB	Industry standard. Ornaments + samples.
.pt2	Pro Tracker 2	3	~0.8-1.0 KB	Older, simpler. Still used for size reasons.
.stc	Sound Tracker / Sound Tracker Pro	3	~0.5-0.7 KB	Oldest Spectrum format. No ornaments.
.asc	ASC Sound Master	3	~0.8-1.0 KB	Compact player. Russian scene favourite.
.sqt	SQ-Tracker	3	~0.6-0.8 KB	Excellent data compression.
.ay	AY emulation container	varies	N/A	Captures raw register dumps. For emulators, not playback engines.

Vortex Tracker II

Vortex Tracker II is the modern standard for composing AY music. It runs on Windows (and via Wine on Linux/macOS), and directly exports .pt3 files compatible with all standard Z80 players.

Key features for demoscene use: - **TurboSound mode:** Edit 6 channels (2 x AY) in a single module. - **Ornament editor:** Visual per-frame pitch offset tables. - **Sample editor:** Per-frame amplitude + tone/noise control. - **Loop points:** Set loop start for each pattern and for the entire song. - **Export:** .pt3 (native), .txt (plain text dump for analysis), .wav (audio render).

The typical demoscene workflow: 1. Compose in Vortex Tracker II. 2. Export as .pt3. 3. Include the standard pt3player.asm (widely available, multiple versions optimised

for size or speed). 4. INCBIN the .pt3 data. 5. Call `music_init` and `music_play` as shown above.

AY-3-8910 vs YM2149: Differences That Matter

The Yamaha YM2149 (used in the Atari ST and some Spectrum clones) is pin-compatible with the AY-3-8910 but not bit-identical:

Feature	AY-3-8910	YM2149
Volume curve	Logarithmic (1.5 dB/step)	Linear
Noise LFSR	17-bit, specific taps	17-bit, different taps
Envelope precision	16 volume steps (4-bit)	32 volume steps (5-bit internally)
Pin 26 (SEL)	Clock divider	Same, but often hard-wired
Output DAC	4-bit resistor ladder	5-bit resistor ladder

Practical impact: The same tune sounds warmer/bassier on a real AY-3-8910 and brighter/thinner on a YM2149 due to the different volume curves. Emulators typically let you select which chip to emulate. When testing your music, try both – your audience may have either chip in their machine.

Quick Reference Card

Register Summary (Tear-Out)

R0	= Tone A low	(8 bits)	R/W	
R1	= Tone A high	(4 bits)	R/W	
R2	= Tone B low	(8 bits)	R/W	
R3	= Tone B high	(4 bits)	R/W	
R4	= Tone C low	(8 bits)	R/W	
R5	= Tone C high	(4 bits)	R/W	
R6	= Noise period	(5 bits)	R/W	0=highest, 31=lowest
R7	= Mixer	(8 bits)	R/W	0=0N (active low!)
R8	= Volume A	(5 bits)	R/W	bit4=envelope mode
R9	= Volume B	(5 bits)	R/W	bit4=envelope mode
R10	= Volume C	(5 bits)	R/W	bit4=envelope mode
R11	= Envelope low	(8 bits)	R/W	
R12	= Envelope high	(8 bits)	R/W	
R13	= Envelope shape	(4 bits)	W	write = restart!

Ports: \$FFFD = select register \$BFFD = write data

Write: OUT (\$FFFD),reg then OUT (\$BFFD),value

Read: OUT (\$FFFD),reg then IN A,(\$FFFD)

Tone: freq = 1773400 / (16 * period) period = 12-bit (1-4095)

Env: freq = 1773400 / (256 * period) period = 16-bit (1-65535)
 Noise: freq = 1773400 / (16 * period) period = 5-bit (0-31)

Mixer R7 bits: [IOB IOA | NC NB NA | TC TB TA] 0=enable, 1=disable

Envelope R13: [CONT ATT ALT HOLD]

\$00 = ___ \$08 = \\\ \$0A = \\/ \$0B = \^^

\$04 = /___ \$0C = /// \$0E = /\ \$0D = /^^

TurboSound: \$FF->\$FFFD = chip 0 \$FE->\$FFFD = chip 1

Next 3xAY: \$FF = chip 0 \$FE = chip 1 \$FD = chip 2

Sources: General Instrument AY-3-8910 / AY-3-8912 datasheet (1979); Yamaha YM2149 Application Manual; Dark "GS Sound System" (Spectrum Expert #01, 1997); Introspec "Eager" making-of (Hype 2015); Vortex Tracker II documentation (S.V. Bulba); NedoPC TurboSound FM documentation; ZX Spectrum Next User Manual, Issue 2

Appendix H: Storage APIs — TR-DOS and esxDOS

“A technically impressive demo that ships as a .tZX when the rules require .trd will be disqualified.” – Chapter 20

Two storage APIs dominate the ZX Spectrum world: **TR-DOS** (the disk operating system of the Soviet Beta Disk 128 interface, standard on Pentagon and Scorpion clones) and **esxDOS** (the modern SD card operating system running on DivMMC and DivIDE hardware). Most Russian and Ukrainian demoscene releases ship as .trd disk images. Most modern Western releases use .tap tape images or esxDOS-compatible file layouts. If you are releasing a demo or game today, the practical choice is to provide a .trd image for compatibility with the enormous Russian/Ukrainian install base, and a .tap file for everyone else. If your loader supports esxDOS detection (as described in Chapter 21), users with DivMMC hardware get fast SD card loading for free.

This appendix is the API reference you keep open while writing your loader. Chapter 21 covers integration in a full game project. Chapter 15 covers the hardware details of memory banking and port mapping that underpin both APIs.

1. TR-DOS (Beta Disk 128)

Hardware

The Beta Disk 128 interface is the standard floppy disk controller for Pentagon, Scorpion, and most Soviet ZX Spectrum clones. It is based on the Western Digital WD1793 floppy disk controller chip, which communicates with the Z80 through five I/O ports.

The TR-DOS ROM (8 KB) occupies \$0000–\$3FFF when the Beta Disk interface is active. It is paged in automatically when the Z80 executes code at address \$3D13 (the magic entry point), and paged out when execution returns to the main ROM area.

Disk Format

Property	Value
Tracks	80
Sides	2

Property	Value
Sectors per track	16
Bytes per sector	256
Total capacity	640 KB (655,360 bytes)
Image format	.trd (raw disk image, 640 KB)
System track	Track 0, side 0

Track 0 contains the disk directory (sectors 1-8) and the disk information sector (sector 9). The directory holds up to 128 file entries. Each entry is 16 bytes:

Bytes 0-7: Filename (8 characters, space-padded)
 Byte 8: File type: 'C' = code, 'B' = BASIC, 'D' = data, '#' = sequential
 Bytes 9-10: Start address (or BASIC line number)
 Bytes 11-12: Length in bytes
 Byte 13: Length in sectors
 Byte 14: Starting sector
 Byte 15: Starting track

WD1793 Port Map

Port	Read	Write
\$1F	Status register	Command register
\$3F	Track register	Track register
\$5F	Sector register	Sector register
\$7F	Data register	Data register
\$FF	TR-DOS system register	TR-DOS system register

Port \$FF is the Beta Disk system port. It controls drive selection, side selection, head load, and density. The upper bits also reflect the DRQ (Data Request) and INTRQ (Interrupt Request) signals from the WD1793.

WD1793 Commands

Command	Code	Description
Restore	\$08	Move head to track 0. Verify track.
Seek	\$18	Move head to track in data register.
Step In	\$48	Step head one track toward centre.
Step Out	\$68	Step head one track toward edge.
Read Sector	\$88	Read one 256-byte sector.
Write Sector	\$A8	Write one 256-byte sector.
Read Address	\$C0	Read the next sector ID field.
Force Interrupt	\$D0	Abort current command.

The low nibble of each command byte carries modifier flags (step rate, verify, side select, delay). The values above use common defaults. Consult the WD1793 datasheet for the full bit layout.

ROM API: Loading a File

The standard approach to file I/O under TR-DOS is to call the ROM routines at \$3D13. The TR-DOS ROM provides high-level file operations through a command system: you place parameters in registers and in the system area at \$5D00-\$5FFF, then call into the ROM.

```
; Loads a code file ('C' type) to its stored start address
;
; The filename must be placed at $5D02 (8 bytes, space-padded).
; The file type goes to $5D0A.
;
; Call $3D13 with C = $08 (load file command)
```

```
load_trdos_file:
    ; Set up filename at TR-DOS system area
    ld    hl, my_filename
    ld    de, $5D02
    ld    bc, 8
    ldir                      ; copy 8-char filename

    ld    a, 'C'              ; file type: code
    ld    ($5D0A), a

    ld    c, $08              ; TR-DOS command: load file
    call  $3D13              ; enter TR-DOS ROM
    ret
```

```
my_filename:
    db    "SCREEN  "          ; 8 characters, space-padded
```

To load to a specific address (overriding the stored start address):

```
; HL = destination address
; DE = length to load
; Filename already at $5D02, type at $5D0A
```

```
load_trdos_to_addr:
    ld    hl, $4000           ; load to screen memory
    ld    de, 6912           ; 6912 bytes (one screen)
    ld    ($5D03), hl        ; override start address
    ld    ($5D05), de        ; override length
    ld    c, $08             ; load file
    call  $3D13
    ret
```

Direct Sector Access

For demos that stream data from disk — fullscreen animations, music data that exceeds available RAM, or multipart demos that load effects on the fly — direct sector access bypasses the file system entirely. You control the head position, read sectors one at a time, and process the data as it arrives.

```
; B = track number (0-159, with side encoded in bit 0 of $FF)
; C = sector number (1-16)
; HL = destination buffer (256 bytes)
```

```
read_sector:
```

```
    ld    a, b
    out   ($3F), a        ; set track register
    ld    a, c
    out   ($5F), a        ; set sector register

    ld    a, $88          ; Read Sector command
    out   ($1F), a        ; issue command
```

```
    ; Wait for DRQ and read 256 bytes
    ld    b, 0            ; 256 bytes to read
```

```
.wait_drq:
```

```
    in    a, ($FF)        ; read system register
    bit   6, a            ; test DRQ bit
    jr    z, .wait_drq    ; wait until data ready
    in    a, ($7F)        ; read data byte
    ld    (hl), a
    inc   hl
    djnz  .wait_drq
```

```
    ; Wait for command completion
```

```
.wait_done:
```

```
    in    a, ($1F)        ; read status register
    bit   0, a            ; test BUSY bit
    jr    nz, .wait_done
    ret
```

Warning: Direct sector access is timing-sensitive. Interrupts must be disabled during the data transfer loop, or bytes will be lost. The WD1793 asserts DRQ for a limited time window; if the Z80 does not read the data register before the next byte arrives, data is overwritten. At 250 kbit/s (double density), you have approximately 32 microseconds per byte — about 112 T-states on a Pentagon. The tight loop above runs in roughly 50-60 T-states per byte, leaving adequate margin.

Disk Detection

To detect whether a Beta Disk interface is present:

```
; Returns: carry clear if present, carry set if absent
```

```
detect_beta_disk:
```

```
    ; The TR-DOS ROM signature is at $0069 when paged in.
    ; We can check port $FF for a sane response:
    ; If no Beta Disk is present, port $FF reads as floating bus.
    in    a, ($1F)        ; read WD1793 status
    cp    $FF            ; floating bus returns $FF
    scf
    ret   z              ; probably no controller
    or    a              ; clear carry = present
```

```
ret
```

A more robust method is to attempt to call \$3D13 and check if TR-DOS ROM signature bytes are present. Production code typically checks for a known byte sequence at the TR-DOS ROM entry points.

2. esxDOS (DivMMC / DivIDE)

Hardware

DivMMC (and its older sibling DivIDE) is a mass storage interface that connects an SD card to the ZX Spectrum. The esxDOS firmware provides a POSIX-like file API accessible from Z80 code through RST \$08. esxDOS supports FAT16 and FAT32 file systems, long filenames, subdirectories, and multiple open file handles.

DivMMC uses auto-mapping: when the Z80 fetches an instruction from certain “trap” addresses (notably \$0000, \$0008, \$0038, \$0066, \$04C6, \$0562), the DivMMC hardware automatically pages in its own ROM at \$0000-\$1FFF. The RST \$08 trap is the primary API entry point.

API Pattern

Every esxDOS call follows the same pattern:

```
db  function_id      ; function number (byte after RST)
; Returns:
;   Carry clear = success
;   Carry set   = error, A = error code
```

The function number is the byte immediately following the RST \$08 instruction in memory. The Z80 executes RST \$08, which jumps to address \$0008. DivMMC auto-maps its ROM at that address, reads the next byte (the function number), dispatches the call, then un-maps its ROM and returns to the instruction after the DB.

Function Reference

Function	ID	Description	Input	Output
M_GETSETDRV	\$89	Get/set default drive	A = '*' for default	A = drive letter
F_OPEN	\$9A	Open file	IX = filename (zero-terminated), B = mode, A = drive	A = file handle
F_CLOSE	\$9B	Close file	A = file handle	-

Function	ID	Description	Input	Output
F_READ	\$9D	Read bytes	A = handle, IX = buffer, BC = count	BC = bytes read
F_WRITE	\$9E	Write bytes	A = handle, IX = buffer, BC = count	BC = bytes written
F_SEEK	\$9F	Seek in file	A = handle, L = whence, BCDE = offset	BCDE = new position
F_FSTAT	\$A1	File status (by handle)	A = handle, IX = buffer	11-byte stat block
F_OPENDIR	\$A3	Open directory	IX = path (zero- terminated)	A = dir handle
F_READDIR	\$A4	Read directory entry	A = dir handle, IX = buffer	entry at (IX)
F_CLOSEDIR	\$A5	Close directory	A = dir handle	-
F_GETCWD	\$A8	Get current directory	IX = buffer	path at (IX)
F_CHDIR	\$A9	Change directory	IX = path	-
F_STAT	\$AC	File status (by name)	IX = filename	11-byte stat block

File Open Modes

Mode	Value	Description
Read only	\$01	Open existing file for reading
Create/truncate	\$06	Create new or truncate existing for writing
Create new only	\$04	Create new file; fail if exists
Append	\$0E	Open for writing at end of file

Seek Whence Values

Whence	Value	Description
SEEK_SET	\$00	Offset from beginning of file
SEEK_CUR	\$01	Offset from current position
SEEK_END	\$02	Offset from end of file

Code Example: Load a File

```
;
; Uses register conventions from esxDOS API documentation.
```

; Note: F_READ uses IX for the destination buffer, not HL.

```

ld    a, '*'           ; use default drive
ld    ix, filename      ; pointer to zero-terminated filename
ld    b, $01           ; FA_READ: open for reading
rst   $08
db    $9A              ; F_OPEN
jr    c, .error        ; carry set = error

ld    (.file_handle), a ; save file handle

ld    ix, $4000         ; destination buffer (screen memory)
ld    bc, 6912          ; bytes to read (one full screen)
ld    a, (.file_handle)
rst   $08
db    $9D              ; F_READ
jr    c, .error

ld    a, (.file_handle)
rst   $08
db    $9B              ; F_CLOSE
ret

```

```

.error:
; A contains the esxDOS error code
; Common errors:
;   5 = file not found
;   7 = file already exists (on create)
;   9 = invalid file handle
ret

```

```

filename:
db    "screen.scr", 0

```

```

.file_handle:
db    0

```

Code Example: Streaming Data from File

For demos that load data incrementally — decompressing level chunks between frames, streaming a pre-rendered animation, or loading music patterns on demand — the pattern is: open the file once, read a chunk per frame, close when done.

; Call `stream_init` once, then `stream_chunk` from your main loop.

```
CHUNK_SIZE equ 256           ; bytes per frame (tune to budget)
```

```
stream_handle: db 0
stream_done:   db 0

```

```

; Initialise: open the file
stream_init:

```

```

    ld    a, '*'
    ld    ix, stream_file
    ld    b, $01                ; FA_READ
    rst   $08
    db    $9A                  ; F_OPEN
    ret   c                    ; error
    ld    (stream_handle), a
    xor   a
    ld    (stream_done), a     ; not done yet
    ret

; Per-frame: read one chunk into buffer
; Returns: BC = bytes actually read (may be < CHUNK_SIZE at EOF)
stream_chunk:
    ld    a, (stream_done)
    or    a
    ret   nz                    ; already finished

    ld    a, (stream_handle)
    ld    ix, stream_buffer
    ld    bc, CHUNK_SIZE
    rst   $08
    db    $9D                  ; F_READ
    jr    c, .eof

    ; BC = bytes actually read
    ld    a, b
    or    c
    jr    z, .eof              ; zero bytes read = end of file
    ret

.eof:
    ld    a, (stream_handle)
    rst   $08
    db    $9B                  ; F_CLOSE
    ld    a, 1
    ld    (stream_done), a
    ret

stream_file:
    db    "anim.bin", 0

stream_buffer:
    ds    CHUNK_SIZE

```

Detecting esxDOS

```

; Returns: carry clear = esxDOS available, carry set = not available
;
; Strategy: attempt M_GETSETDRV. If esxDOS is present, it returns
; the current drive letter. If not present, RST $08 goes to the

```

; Spectrum ROM's error handler at \$0008 (a benign instruction on
; the 128K ROM) and does not crash.

```
detect_esxdos:
    ld    a, '*'                ; request default drive
    rst   $08
    db    $89                   ; M_GETSETDRV
    ret                                ; carry flag set by esxDOS on error
```

A more conservative approach checks for the DivMMC's trap signature before calling any API functions. In practice, the method above works on all 128K models because the 128K ROM's \$0008 handler does not crash — it executes a benign sequence and returns. On a 48K machine without esxDOS, RST \$08 goes to the error restart, which may need special handling. Chapter 21 discusses this in the context of a production game loader.

3. +3DOS (Amstrad +3)

The Amstrad Spectrum +3, with its built-in 3-inch floppy drive, has its own DOS: +3DOS. The API uses a different mechanism — calls to entry points in the +3DOS ROM at page \$01, accessed through RST \$08 with a different set of function codes.

+3DOS is rarely used in the demoscene for two reasons. First, the +3 was primarily sold in Western Europe and was never the dominant Spectrum model in any scene community. Second, the +3's non-standard memory layout and ROM paging scheme make it incompatible with most demoscene code written for the 128K/Pentagon architecture. If you need +3 compatibility, the +3DOS API is documented in the Spectrum +3 technical manual (Amstrad, 1987). For most demo and game projects, providing a .tap file is sufficient — the +3 loads .tap files natively through its tape compatibility mode.

4. Practical Patterns

Loading Screen from Disk (TR-DOS)

The loading screen is the user's first impression. On TR-DOS, the screen file (SCREEN C, 6912 bytes) loads directly to \$4000 and appears immediately:

```
; The screen appears as it loads, line by line.
load_screen_trdos:
    ld    hl, scr_filename
    ld    de, $5D02
    ld    bc, 8
    ldir
    ld    a, 'C'
    ld    ($5D0A), a
    ld    hl, $4000            ; destination: screen memory
    ld    ($5D03), hl
    ld    de, 6912            ; length: full screen
```

```

ld    ($5D05), de
ld    c, $08          ; load file
call  $3D13
ret

```

```

scr_filename:
db    "SCREEN  "      ; 8 chars, padded

```

Loading Screen from SD (esxDOS)

Same visual result, different API:

```

load_screen_esxdos:
ld    a, '*'
ld    ix, scr_filename_esx
ld    b, $01          ; FA_READ
rst   $08
db    $9A             ; F_OPEN
ret   c

push  af              ; save handle
ld    ix, $4000        ; destination: screen memory
ld    bc, 6912
pop   af
push  af
rst   $08
db    $9D             ; F_READ
pop   af
rst   $08
db    $9B             ; F_CLOSE
ret

scr_filename_esx:
db    "screen.scr", 0

```

Dual-Mode Loader

A production loader should detect the available storage and use it:

```

load_data:
call  detect_esxdos
jr    nc, .use_esxdos  ; carry clear = esxDOS present

call  detect_beta_disk
jr    nc, .use_trdos   ; carry clear = Beta Disk present

; Fall back to tape loading
jp    load_from_tape

.use_esxdos:
jp    load_from_esxdos

```

```
.use_trdos:
    jp    load_from_trdos
```

Streaming Compressed Data

The most powerful pattern combines storage API with compression (Appendix C). Open a file containing compressed data, read chunks into a buffer each frame, decompress into the destination, and advance:

```
Frame 1:  F_READ 256 bytes -> buffer    |  decompress buffer -> screen
Frame 2:  F_READ 256 bytes -> buffer    |  decompress buffer -> screen
Frame 3:  F_READ 256 bytes -> buffer    |  decompress buffer -> screen
...
Frame N:  F_READ < 256 bytes (EOF)      |  decompress, close file
```

At 256 bytes per frame and 50 fps, you stream 12.5 KB/sec from the SD card — enough for a compressed fullscreen animation. On TR-DOS, direct sector reads at one sector per frame give 12.8 KB/sec (256 bytes * 50 fps). The bottleneck is decompression speed, not I/O.

5. File Format Reference

Format	Extension	Usage	Notes
TR-DOS disk image	.trd	Standard for Pentagon/Scorpion releases	Raw 640 KB image. Every emulator supports it.
TR-DOS file container	.scl	Simpler than .trd	Contains files without full disk structure. Good for distribution.
Tape image	.tap	Universal tape format	Works on every Spectrum model and emulator. No file system.
Extended tape image	.tzx	Tape with copy protection / turbo loaders	Preserves exact tape timing. Rarely needed for new releases.
Snapshot (48K/128K)	.sna	Quick load, no file system	Captures full machine state. No loading code needed.
Snapshot (compressed)	.z80	Like .sna but compressed	Multiple versions; .z80 v3 supports 128K.

Format	Extension	Usage	Notes
Next distribution	.nex	ZX Spectrum Next executable	Self-contained binary with header specifying bank layout.

Choosing a release format: For a demoscene release, provide at least two formats:

1. **.trd** for TR-DOS users (the Russian/Ukrainian community, Pentagon/Scorpion owners, and emulator users who prefer disk images). This is the default for parties like Chaos Constructions and DiHalt.
2. **.tap** for everyone else (real 128K hardware with tape input, DivMMC users via .tap loader, and all emulators). sjasmpplus can generate .tap output directly with its SAVETAP directive.

If your demo is small enough (under 48 KB), a .sna snapshot also works well — it loads instantly with no loader code.

6. See Also

- **Chapter 15** — Hardware Anatomy: memory banking, port \$7FFD, the full port map that TR-DOS and esxDOS sit alongside.
- **Chapter 20** — Demo Workflow: release formats, party submission rules, .trd vs .tap requirements.
- **Chapter 21** — Full Game: production-quality tape and esxDOS loading code, dual-mode detection, bank-by-bank loading.
- **Appendix C** — Compression: which compressors to pair with streaming I/O.
- **Appendix E** — eZ80 / Agon Light 2: the MOS file API on Agon, which provides similar file operations (mos_fopen, mos_fread, mos_fclose) through a different mechanism (RST \$08 with MOS function codes in ADL mode).

Sources: WD1793 datasheet (Western Digital, 1983); TR-DOS v5.03 disassembly (various, public domain); esxDOS API documentation (Wikipedia, zxe.io); DivMMC hardware specification (Mario Prato / ByteDelight); Spectrum +3 Technical Manual (Amstrad, 1987); Introspec, “Loading and saving on the Spectrum” (Hype, 2016)

Appendix I: Bytebeat and AY-Beat - Generative Sound on Z80

“At 256 bytes, bytebeat is your only realistic option - there is no room for a pattern player.” - Chapter 13

This appendix covers formula-driven sound generation on the ZX Spectrum - from the original PCM bytebeat concept to the AY-adapted technique that produces structured, evolving music from a handful of Z80 instructions. Chapter 13 introduces AY-beat as a size-coding tool. This appendix is the full reference: the theory, the formulas, the register mappings, and a complete working engine you can drop into a 256-byte intro.

You will need the AY register reference in Appendix G open alongside this appendix. Every register number mentioned here (R0, R7, R8, R11, R13, etc.) is documented there with full bit layouts and port addresses.

1. Classic Bytebeat: The PCM Tradition

In 2011, Ville-Matias Heikkila (Viznut) published a discovery that had been circulating in underground programming circles: a single C expression, evaluated once per sample with an incrementing counter t , can produce complex rhythmic music when the output is interpreted as 8-bit unsigned PCM at 8 kHz.

The core idea:

```
putchar( f(t) );    // pipe to /dev/dsp at 8000 Hz
```

The function $f(t)$ is typically a one-line expression built from bitwise operations, multiplication, and bit shifts. No oscillators, no envelopes, no note tables - just integer arithmetic on a counter.

Famous Formulas

$t*((t \gg 12 | t \gg 8) \& 63 \& t \gg 4)$ - Viznut's original. Cascading rhythmic tones that cycle through pitch relationships, producing an effect somewhere between a music box and a broken telephone. The $t \gg 12$ and $t \gg 8$ create two frequency-divided versions

of the counter; $\&63$ limits the range; $t \gg 4$ gates the output rhythmically. The multiplication by t creates the fundamental pitch sweep.

$t * (t \gg 5 | t \gg 8) \gg (t \gg 16)$ – Evolving rhythmic patterns. The right-shift by $t \gg 16$ means the entire character of the sound changes every ~8 seconds (65536 samples at 8 kHz). Each 8-second section has a different dynamic range and feel.

$(t * 5 \& t \gg 7) | (t * 3 \& t \gg 10)$ – Two interleaved melodic lines. The $t * 5$ and $t * 3$ create two pitch streams at different intervals; the AND with shifted counters gates them independently; the OR merges them. The result sounds like two interlocking melodies playing simultaneously.

Why It Works

Bitwise operations on an incrementing counter create periodic structures at multiple time scales simultaneously. Consider the bit pattern of t as it counts:

- Bit 0 toggles every sample (4000 Hz – inaudible as pitch, but shapes the waveform)
- Bit 7 toggles every 128 samples (~62.5 Hz – bass territory)
- Bit 12 toggles every 4096 samples (~1.95 Hz – rhythmic pulse)
- Bit 15 toggles every 32768 samples (~0.24 Hz – structural change)

A right-shift $t \gg n$ selects which time scale dominates. AND operations create coincidence patterns – moments when two time scales align. OR operations merge independent patterns. Multiplication by small constants creates harmonic relationships (frequency ratios). The 8-bit truncation of the output acts as a natural waveform shaper, folding values back into range and creating additional harmonics.

The result is self-similar: the sound has rhythmic structure at every scale, from individual oscillation cycles up to multi-second phrase structures. This self-similarity is what makes bytebeat sound like music rather than noise – even though no musical knowledge went into the formula.

On the Spectrum: The Beeper Dead End

The ZX Spectrum's beeper output is a 1-bit speaker controlled by bit 4 of port \$FE. You can, in principle, run a bytebeat formula and output the result:

```
; DE = t (16-bit counter)
    ld    de, 0
.loop:
    ; Compute f(t) -- simplified: A = t AND (t >> 8)
    ld    a, e            ; 4T    A = low byte of t
    and    d              ; 4T    A = t_lo AND t_hi
    ; Output bit 4 to speaker
    and    $10            ; 7T    isolate bit 4
    out    ($FE), a       ; 11T   toggle speaker
    inc    de             ; 6T    t++
    jr     .loop          ; 12T
                                ; --- 44T per sample = ~79.5 kHz
```

This runs, and it produces sound. But it consumes 100% of the CPU – the Z80 is doing nothing but computing samples and toggling the speaker. No screen updates,

no visual effects, no input handling. The sample rate is also wrong (too fast), and controlling it precisely requires careful cycle counting with padding NOPs.

For a demo, this is a dead end. The beeper is a 1-bit output that demands constant CPU attention. The real adaptation of bytebeat to the Spectrum requires a different approach entirely.

2. AY-Beat: Bytebeat Reimagined for a Tone Generator

The AY-3-8910 is not a DAC. It does not accept amplitude samples. It is a programmable tone generator: you give it a frequency (as a period value), a volume (0-15), and optional noise and envelope parameters, and its internal oscillators produce the sound autonomously. The CPU is free to do other work.

The key insight of AY-beat: **replace the sample counter with a frame counter, and replace PCM output with AY register values.**

Classic bytebeat computes one amplitude sample at ~8000 Hz. AY-beat computes tone periods, volumes, and noise parameters at 50 Hz – once per video frame, triggered by the HALT instruction. The AY's oscillators handle the actual sound generation between frames.

The frame counter *t* replaces the sample counter. Formulas operate on *t* but produce register values, not waveform samples. Where PCM bytebeat has one degree of freedom (amplitude), AY-beat has many: three independent tone periods (12-bit each), three volumes (4-bit each), one noise period (5-bit), and a 16-bit envelope period with shape selection.

Basic AY-Beat Architecture

```
; Assumes frame_counter is a byte in memory
ay_beat_update:
    ld    a, (frame_counter)
    ld    e, a                ; E = t (keep a copy)

    ; === Channel A: tone period from formula ===
    ; tone_lo = (t * 3) AND $3F -- pentatonic-ish cycling
    add  a, a                 ; 4T  A = t * 2
    add  a, e                 ; 4T  A = t * 3
    and  $3F                 ; 7T  mask to 6 bits (periods 0-63)
    ld    d, a                ; 4T  save tone period
    ; Write R0 (tone A low) = D
    xor  a                    ; 4T  A = 0 (register number)
    call ay_write_d           ; writes D to AY register A

    ; Write R1 (tone A high) = 0
    ld    a, 1                ; 7T
    ld    d, 0                ; 7T
    call ay_write_d

    ; === Channel A: volume from formula ===
```

```

; volume = bits 6-3 of t, giving 0-15 cycling
ld  a, e          ; 4T  reload t
rrca              ; 4T
rrca              ; 4T
rrca              ; 4T
and  $0F          ; 7T  volume 0-15
ld  d, a          ; 4T
ld  a, 8          ; 7T  R8 = Volume A
call ay_write_d

; Advance frame counter
ld  hl, frame_counter
inc (hl)          ; 11T
ret

```

This is the simplest possible AY-beat: one channel, one tone formula, one volume formula, ~30 bytes. It produces a cycling sweep that rises in pitch and fades in and out – not music, but recognisably structured sound.

What Changes from PCM Bytebeat

Aspect	Classic (PCM)	AY-Beat
Update rate	~8000 Hz	50 Hz (frame rate)
Output	8-bit amplitude	Tone period (12-bit), volume (4-bit), noise (5-bit)
Channels	1 (mono speaker)	3 tone + 1 noise + envelope
CPU cost	100% (all cycles)	~200-500 T per frame (~0.3%)
Formula scaling	Fine-grained, fast evolution	Coarse-grained, need wider bit shifts
Sound generation	CPU computes every sample	AY hardware oscillators run autonomously

The 50 Hz frame rate means formulas evolve 160x slower than at 8 kHz. To get equivalent rhythmic density, use larger multipliers and fewer right-shifts. A formula that produces a pleasant rhythm at 8 kHz with $t \gg 12$ (period ~0.5 sec at 8 kHz) needs approximately $t \gg 4$ at 50 Hz for similar timing (~0.3 sec between repetitions). The general rule: divide the PC bytebeat shift amounts by ~7 (log2 of the 160x ratio) and adjust by ear.

3. Drone: Envelope + Tone (E+T Mode)

This is where AY-beat gets genuinely interesting. The AY envelope generator automatically cycles the volume of a channel without any CPU intervention. Set a channel's volume register to envelope mode (bit 4 = 1, i.e. write \$10 to R8/R9/R10),

and the hardware handles volume modulation at the envelope frequency defined by R11-R12.

The result is a drone: a continuously evolving timbre produced by the interaction of the tone oscillator and the envelope oscillator. The CPU cost for maintaining this drone is almost zero – you only need to update the tone period and envelope period once per frame, and the hardware does the rest.

The Drone Recipe

1. Set a tone period from a formula – this defines the base pitch.
2. Set an envelope period from a different formula – this defines the modulation rate.
3. Set the envelope shape to a repeating waveform (shapes \$08, \$0A, \$0C, or \$0E).
4. Set the channel volume to envelope mode (\$10).
5. The hardware produces a continuously evolving drone with zero per-sample CPU cost.

; Tone period evolves per frame, envelope period evolves slower

drone_update:

```
ld  a, (frame_counter)
ld  e, a          ; E = t
```

```
; --- Tone period: slowly sweeping ---
```

```
and $7F          ; 128-frame cycle (2.56 seconds)
ld  d, a
xor  a           ; R0 = Tone A low
call ay_write_d
ld  a, 1         ; R1 = Tone A high
ld  d, 0
call ay_write_d
```

```
; --- Envelope period: evolves on different time scale ---
```

```
ld  a, e         ; reload t
rrca
rrca             ; divide by 4 -- envelope evolves 4x slower
and $3F
add  a, $10      ; offset to avoid very fast envelopes
ld  d, a
ld  a, 11        ; R11 = Envelope period low
call ay_write_d
ld  a, 12        ; R12 = Envelope period high
ld  d, 0
call ay_write_d
```

```
; --- Envelope shape: repeating triangle ---
```

```
; CAUTION: writing R13 restarts the envelope cycle.
```

```
; Only write on the first frame, or when you want a restart.
```

```
ld  a, e
or  a
jr  nz, .skip_shape
```

```

    ld    a, 13                ; R13 = Envelope shape
    ld    d, $0A              ; shape $0A = repeating triangle /\!/
    call ay_write_d
.skip_shape:

    ; --- Volume A: envelope mode ---
    ld    a, 8                ; R8 = Volume A
    ld    d, $10              ; bit 4 set = use envelope
    call ay_write_d

    ret

```

The beauty of E+T mode is the interference between the two frequencies. When the envelope period is close to the tone period, you get amplitude modulation effects – the volume beats at the difference frequency, producing a wavering, organ-like timbre. When the envelope frequency is much lower than the tone frequency, it acts as a slow tremolo. When it is much higher, you get buzz-bass territory (see Appendix G and Chapter 11).

Sweeping the envelope period while the tone period also moves produces continuously evolving textures. The two formulas create a two-dimensional parameter space that the sound explores over time. With the right formula pair, the drone never quite repeats – it wanders through timbral variations, creating an ambient soundscape from fewer than 30 bytes of code.

Byte Cost

Setting up E+T mode with a formula takes approximately 15-25 bytes. For a 256-byte intro, this gives you rich, evolving drone sound essentially for free – no per-frame volume computation needed, no pattern data, just two register values derived from simple formulas. The AY hardware is doing all the oscillation work.

4. Noise Percussion

The AY noise generator (R6) produces pseudo-random noise at a programmable frequency (0-31). The mixer register (R7) controls which channels receive noise. Toggling noise on and off rhythmically, driven by the frame counter, creates percussion patterns.

Basic Kick Pattern

```

; Frame counter in A (already loaded)
    ld    e, a                ; E = t
    and   $07                ; every 8th frame = 6.25 Hz pulse
    jr    nz, .decay

    ; --- Hit: enable noise on C, max volume ---
    ld    a, 7                ; R7 = Mixer
    ld    d, %00100100        ; tone C off, noise C on, others unchanged
    call ay_write_d

```

```

    ld  a, 6           ; R6 = Noise period
    ld  d, 2           ; low period = harsh, punchy
    call ay_write_d
    ld  a, 10          ; R10 = Volume C
    ld  d, $0F         ; maximum volume
    call ay_write_d
    jr  .done

.decay:
    ; --- Decay: reduce volume each frame ---
    ; Simple approach: volume = 15 - (t AND 7)
    ld  a, e
    and $07            ; frames since last hit
    ld  d, a
    ld  a, $0F
    sub d               ; volume = 15 - elapsed
    jr  nc, .vol_ok
    xor a               ; clamp to 0
.vol_ok:
    ld  d, a
    ld  a, 10          ; R10 = Volume C
    call ay_write_d

.done:

```

Percussion Character by Noise Period

R6 Value	Character	Use
0-3	Harsh click, punchy	Kick drum, rimshot
4-8	Crisp hiss	Snare body
10-15	Broad noise	Open hi-hat
20-31	Low rumble	Distant thunder, ambient

Rhythmic Variety from Bit Masks

Different AND masks on the frame counter produce different rhythmic densities:

Mask	Period	Frequency	Character
AND \$03	Every 4 frames	12.5 Hz	Rapid fire, hi-hat
AND \$07	Every 8 frames	6.25 Hz	Standard kick
AND \$0F	Every 16 frames	3.125 Hz	Half-time, sparse
AND \$1F	Every 32 frames	1.5625 Hz	Slow pulse, intro

Combine two masks for polyrhythm: test `t AND $07` for the kick, `t AND $03` for the hi-hat. This costs about 10 extra bytes but adds significant rhythmic complexity.

Using the Envelope for Drum Decay

Instead of manually decaying the volume each frame, use the AY envelope generator in single-shot mode. Set R13 to shape \$00 (decay to zero, hold), and the hardware handles the volume fade automatically:

```

ld    a, e
and   $07                ; every 8th frame
jr    nz, .no_hit
; Set envelope period (controls decay speed)
ld    a, 11
ld    d, $80              ; period = $0080 -- medium decay
call  ay_write_d
ld    a, 12
ld    d, 0
call  ay_write_d
; Trigger: write shape $00 (single decay)
ld    a, 13
ld    d, $00
call  ay_write_d
; Volume C = envelope mode
ld    a, 10
ld    d, $10
call  ay_write_d
.no_hit:

```

This saves several bytes by eliminating the manual decay code. The trade-off: the envelope generator is shared across all channels in envelope mode. If channel A is using E+T drone, channel C cannot independently use the envelope for drum decay. Plan your channel allocation accordingly.

5. Multi-Channel Harmony

The AY has three independent tone channels. AY-beat can derive all three from a single formula using bit rotation, creating the impression of counterpoint from almost no code.

Three Voices from One Formula

```

ld    a, (frame_counter)
ld    e, a

; === Channel A: base formula ===
and   $3F                ; period 0-63
ld    d, a
xor   a                   ; R0
call  ay_write_d

; === Channel B: same formula, rotated 2 bits ===
ld    a, e

```

```

rrca
rrca                ; rotate right by 2
and $3F
ld d, a
ld a, 2              ; R2
call ay_write_d

; === Channel C: same formula, rotated 4 bits ===
ld a, e
rrca
rrca
rrca
rrca                ; rotate right by 4
and $3F
ld d, a
ld a, 4              ; R4
call ay_write_d

```

The bit rotations create phase-shifted versions of the same pattern. The channels play related but offset melodies – they follow the same contour but arrive at each pitch at different times. This creates an impression of counterpoint: multiple independent voices that share an underlying logic.

Why Rotation Creates Harmony

RRCA is a *rotation*, not a shift – bits that fall off the bottom wrap to the top. This means the three channels cycle through the same set of period values in the same order, but offset in time. The offset depends on the rotation amount:

- **RRCA x 2:** The channel is “ahead” by approximately a quarter of the pattern cycle. This often creates intervals that sound like fourths or fifths – not precisely tuned, but harmonically related enough to be pleasant.
- **RRCA x 4:** Half a byte’s worth of offset. This tends to produce octave-like relationships, since bit 4 being rotated to bit 0 effectively halves the period in certain phase alignments.

These are not real musical intervals. They are pseudo-harmonic relationships created by the structure of binary numbers. But the ear is forgiving – if two tones share most of their bit pattern, they sound “related,” and that is enough for a 256-byte intro.

Volume Formulas for Multi-Channel

Give each channel a different volume formula to avoid all three voices being at the same level simultaneously:

```

ld a, e
rrca
rrca
rrca
and $0F
ld d, a
ld a, 8

```



```

call ay_write_d

; Volume B: bits 4-1 of t (different phase)
ld  a, e
rrca
and $0F
ld  d, a
ld  a, 9
call ay_write_d

; Volume C: inverted bits 5-2 of t
ld  a, e
rrca
rrca
and $0F
xor $0F          ; invert -- when A is loud, C is quiet
ld  d, a
ld  a, 10
call ay_write_d

```

The inverted volume on channel C creates a call-and-response dynamic: as one voice fades in, another fades out. This costs 2 extra bytes (the XOR \$0F) but significantly improves the musical texture.

6. Formula Cookbook

The following formulas have been tested on the AY at 50 Hz frame rate. “Bytes” refers to the Z80 implementation cost for computing the formula from a value already in register A (the frame counter). The period mask determines the pitch range.

Tone Period Formulas

#	Formula	Z80 Implementation	Bytes	Sound	Best For
1	t AND \$3F	and \$3F	2	Rising saw-tooth, 1.28-sec cycle	Simple sweep
2	t*3 AND \$3F	ld e,a : add a,a : add a,e : and \$3F	5	Faster sweep, wider intervals	Energetic bass
3	t XOR (t>>3)	ld e,a : rrca : rrca : rrca : xor e	5	Chaotic with periodic structure	Noise texture

#	Formula	Z80 Implementation	Bytes	Sound	Best For
4	(t AND \$0F) XOR \$0F	and \$0F : xor \$0F	4	Triangle wave, ping-pong sweep	Melodic lead
5	t*5 AND t>>2	ld e,a : add a,a : add a,a : add a,e : ld d,a : ld a,e : rrca : rrca : and d	10	Rhythmic gating	Percussion-like
6	(t+t>>4) AND \$1F	ld e,a : rrca : rrca : rrca : rrca : add a,e : and \$1F	6	Slowly modulated sweep	Evolving drone
7	t AND (t>>3) AND \$1F	ld e,a : rrca : rrca : rrca : and e : and \$1F	6	Self-similar, fractal rhythm	Complex patterns
8	(t>>1) XOR (t>>3)	ld e,a : rrca : ld d,a : rrca : rrca : xor d	6	Dual-speed interference	Metallic texture
9	t*7 AND \$7F	ld e,a : add a,a : add a,a : add a,a : sub e : and \$7F	6	Wide sweep, 7x speed	Fast arpeggio feel
10	(t XOR t>>1) AND \$3F	ld e,a : rrca : xor e : and \$3F	5	Gray code sequence	Staircase melody
11	t AND \$07 OR t>>4	ld e,a : and \$07 : ld d,a : ld a,e : rrca : rrca : rrca : rrca : or d	8	Nested loops, two rhythmic layers	Layered rhythm
12	(t+t+t>>2) AND \$3F	ld e,a : add a,a : ld d,a : ld a,e : rrca : rrca : add a,d : and \$3F	8	Accelerating sweep with sub-pattern	Textured lead

Volume Formulas

#	Formula	Z80	Bytes	Effect
V1	t>>3 AND \$0F	rrca : rrca : : rrca : and \$0F	5	Slow fade cycle, 5.12 sec
V2	(t AND \$0F) XOR \$0F	and \$0F : xor \$0F	4	Triangle volume, ping-pong

#	Formula	Z80	Bytes	Effect
V3	t*3>>4 AND \$0F	ld e,a : add a,a : add a,e : rrca : rrca : rrca : rrca : and \$0F	8	Irregular fade pattern
V4	\$0F (constant)	ld d,\$0F	2	Maximum volume, use with envelope mode

How to Read the Table

Pick a tone formula and a volume formula. Combine them. The total byte cost is the sum of both implementations plus the AY register write overhead (~8 bytes per channel for two ay_write calls: register select + data for tone low and volume). A single channel with formula #1 and volume V1 costs approximately 2 + 5 + 16 = 23 bytes including register writes.

Formula #10 (Gray code) deserves special mention. The Gray code sequence only changes one bit per step, so the tone period changes by exactly one unit per frame – a smooth, staircase-like melody. Combined with the AND mask, it cycles through a limited pitch range with pleasant regularity. This is one of the most musical-sounding single formulas.

7. Putting It Together: A Complete AY-Beat Engine

Here is a complete, minimal AY-beat engine that produces 3-channel generative sound with envelope drone. This is the engine you drop into a 256-byte intro alongside your visual effect.

```
; Complete AY-beat engine -- 47 bytes
; Produces 3-channel generative music with envelope drone
; Call once per frame (after HALT)
; Clobbers: AF, BC, DE, HL
; =====

ay_beat:
    ld    hl, .frame
    ld    a, (hl)          ; A = frame counter
    inc   (hl)             ; advance for next frame
    ld    e, a              ; E = t (preserved copy)

    ; --- Mixer: all three tones on, no noise ---
    ; Only needed on first frame, but costs 5 bytes either way
    push af
    ld    a, 7              ; R7
    ld    d, $38            ; tones A+B+C on, noise off
```

```

call .wr
pop  af

; --- Channel A: tone = t AND $3F ---
and  $3F
ld   d, a
xor  a          ; R0
call .wr

; --- Channel B: tone = t*3 AND $3F ---
ld   a, e
add  a, a
add  a, e          ; A = t * 3
and  $3F
ld   d, a
ld   a, 2          ; R2
call .wr

; --- Channel C: tone = (t XOR t>>1) AND $3F ---
ld   a, e
rrca
xor  e
and  $3F
ld   d, a
ld   a, 4          ; R4
call .wr

; --- Volumes: A and B fixed at 12, C = envelope mode ---
ld   a, 8          ; R8 = Volume A
ld   d, 12
call .wr
ld   a, 9          ; R9 = Volume B
ld   d, 10
call .wr
ld   a, 10         ; R10 = Volume C
ld   d, $10        ; envelope mode
call .wr

; --- Envelope: period sweeps with t, triangle shape ---
ld   a, e
rrca
rrca
and  $3F
add  a, $20        ; keep envelope period above $20
ld   d, a
ld   a, 11         ; R11
call .wr
; R12 = 0 (envelope period high)
inc  a             ; A = 12
ld   d, 0
call .wr

```

```

; Shape: only write on frame 0 to avoid constant restarts
ld  a, e
or  a
ret  nz          ; skip shape write on all frames except 0
ld  a, 13        ; R13
ld  d, $0E       ; shape $0E = repeating triangle /\ /\
; fall through to .wr, then ret

.wr:
; Write D to AY register A
ld  bc, $FFFD
out (c), a      ; select register
ld  b, $BF
out (c), d      ; write value
ret

.frame:
DB  0           ; frame counter (self-modifying data)

; =====
; Total: 47 bytes (code) + AY write routine shared
; The .wr routine is 9 bytes. If your intro already has an
; AY write routine, the engine body alone is 38 bytes.
; =====

```

What This Produces

- **Channel A:** A simple ascending sweep, cycling through periods 0-63 every 64 frames (1.28 seconds). The fundamental pattern.
- **Channel B:** The same sweep at 3x speed, creating faster-moving intervals. When it aligns with channel A, you hear consonance; when it diverges, you hear dissonance. The alternation creates rhythmic interest.
- **Channel C:** A Gray code sweep in envelope mode. The triangle envelope creates automatic volume modulation, producing a drone that phases against the tone period. This is the harmonic bed underlying the other two voices.
- **Overall:** An evolving, self-similar texture that cycles through tonal relationships. It sounds alien and mechanical – exactly right for a 256-byte intro.

Customisation Points

Change the tone formulas. Swap any of the AND/RRCA sequences for a different formula from the cookbook (section 6). Each substitution changes the character entirely.

Add noise percussion. Insert a `ld a,e : and $07 : jr nz,.no_hit` block (section 4) to add rhythmic kicks. Cost: ~12 bytes. Steal a channel (typically B) or overlay noise on channel C.

Use pentatonic masking. Instead of `AND $3F` as the final mask, index into a 5-byte pentatonic lookup table. This constrains the tone periods to harmonically related values, making the output sound more deliberately musical. Cost: ~8 bytes (5 for

the table, 3 for the lookup). Chapter 13 discusses this technique.

Vary fixed volumes. Replace the constant volume writes with volume formulas from section 6. Even `ld a,e : rrca : rrca : rrca :` and `$0F` (5 bytes per channel) adds significant dynamic interest.

8. Advanced: Combining Techniques

The preceding sections cover individual building blocks. A well-crafted AY-beat engine combines several:

Architecture for a 256-Byte Intro

```

Frame 0:  Set mixer, envelope shape (one-time setup)
Frame N:  Update tone A (melody formula)
          Update tone B (harmony formula, rotated)
          Update volume A (fade formula)
          Update volume B (inverted fade)
          Channel C in E+T drone mode (auto-evolving)
          Every 8th frame: noise hit on C (toggle mixer)

```

The total CPU cost per frame is approximately 300-500 T-states – well under 1% of the ~70,000 T-states available per frame. The remaining 99% is available for your visual effect.

Register Budget

The AY has 14 writable registers. In a minimal AY-beat engine, you typically write 8-10 per frame:

Register	Written	Source
R0 (Tone A low)	Every frame	Formula
R2 (Tone B low)	Every frame	Formula
R4 (Tone C low)	Every frame or once	Formula or fixed
R1, R3, R5 (Tone high)	Once (set to 0)	Constant
R7 (Mixer)	Every frame or once	Constant or toggled for noise
R8, R9 (Volume A, B)	Every frame	Formula or constant
R10 (Volume C)	Once	\$10 (envelope mode)
R11 (Envelope low)	Every frame	Formula
R13 (Envelope shape)	Once (frame 0)	Constant

Registers you can skip entirely: R6 (noise period – only needed if using noise), R12 (envelope high – set once to 0 for short periods), R14-R15 (I/O ports – irrelevant for sound).

Size Breakdown

For a 256-byte intro, every byte matters. Here is how a typical AY-beat budget looks:

Component	Bytes
AY write routine	9
Frame counter management	5
3 tone formulas (simple)	12-18
3 volume settings	6-15
Mixer setup	5
Envelope setup	8-12
Total	45-64

This leaves 192-211 bytes for the visual effect, the main loop, and any other infrastructure. At 45 bytes, the engine in section 7 is close to optimal for the amount of sound it produces.

9. AY-as-DAC: Classic Bytebeat Through the Volume Register

There is a middle path between the beeper dead end and the AY-beat reimagining. The AY-3-8910's volume registers (registers 8, 9, 10) accept 4-bit values (0-15). If you update a volume register at a high rate – say, during a tight loop – the AY output becomes a 4-bit DAC. This is how digitised speech and sample playback work in Spectrum demos.

Applied to bytebeat: compute $f(t)$, shift right to 4 bits, write to volume register:

```
; Still costly (~80% CPU), but sounds better than beeper
    ld    a, 7                ; mixer: all channels off (tone+noise)
    ld    d, %00111111
    call  ay_write
    ld    de, 0                ; t = 0
.loop:
    ; Compute f(t): t AND (t >> 5) -- classic bytebeat formula
    ld    a, e
    ld    b, d
    srl   b
    rr    a
    srl   b
    rr    a
    srl   b
    rr    a
    srl   b
    rr    a
    srl   b
    rr    a
    and   e                    ; A = t >> 5 (using DE as 16-bit t)
                                ; A = t AND (t >> 5)
    rrca
    rrca
    rrca
    rrca
```

```

and  $0F                ; scale to 4-bit (0-15)
ld   bc, AY_REG
ld   b, $FF             ; select register
push af
ld   a, 8               ; register 8: volume A
out  (c), a
ld   b, $BF
pop  af
out  (c), a             ; write volume
inc  de                 ; t++
jr   .loop

```

This produces recognisable bytebeat – the actual waveform formulas from section 1, audible through the AY. The sound quality is better than the beeper (4-bit resolution vs 1-bit), and the AY’s output stage provides proper audio levels.

The cost is still brutal: ~80% CPU. You get a thin sliver of time for visuals – enough for a slowly-updating attribute effect, not enough for anything ambitious. This technique is useful when you want the *specific sound* of classic bytebeat formulas and are willing to pay the CPU price.

Three Output Paths Compared

Path	Resolution	CPU cost	Sound character	Practical for demos?
Beeper (port \$FE)	1-bit	~100%	Harsh, buzzy	No
AY vol- ume DAC	4-bit	~80%	Classic bytebeat	Barely (attribute effects only)
AY- beat (regis- ters)	Tone/noise	~0.5%	Chip music, generative	Yes – the right choice

For size-coded intros and demos, AY-beat is almost always the correct choice. Reserve AY-as-DAC for art projects where the specific bytebeat sound aesthetic is the point.

10. Music Theory for Algorithms

AY-beat formulas that ignore music theory produce interesting noise. Formulas that *encode* music theory produce actual music. The following techniques add musicality for minimal bytes.

Scale Tables: Constraining Output to Pleasant Notes

A raw formula like `tone = t AND $3F` produces all 64 possible period values – most of which are not musically useful. A **scale table** maps formula output to actual

note periods, ensuring every value sounds good.

Scale	Notes	Table size	Character
Pentatonic	5 (C D E G A)	10 bytes (5 × 2-byte periods)	Always consonant, folk/world feel
Diatonic major	7 (C D E F G A B)	14 bytes	Bright, Western, familiar
Diatonic minor	7 (C D Eb F G Ab Bb)	14 bytes	Dark, melancholic
Blues	6 (C Eb F F# G Bb)	12 bytes	Gritty, expressive
Chromatic	12	24 bytes	Atonal, dissonant – usually wrong for sizecoding

The pentatonic scale is the size-coder’s best friend: 5 notes, 10 bytes, and *any* combination of notes sounds acceptable. You cannot play a wrong note on a pentatonic scale. This is why so many 256-byte intros sound vaguely “Asian” or “folk” – the pentatonic constraint makes random sequences musical.

```
; Input: A = formula output (any value)
; Output: DE = AY tone period
    ; Map to scale index: A mod scale_length
    and $07                ; keep low 3 bits
    cp    5                ; pentatonic has 5 notes
    jr    c, .in_range
    sub   5                ; wrap: 5→0, 6→1, 7→2
.in_range:
    add   a, a              ; ×2 for word entries
    ld    hl, pentatonic
    add   a, l
    ld    l, a
    ld    e, (hl)
    inc   hl
    ld    d, (hl)          ; DE = tone period
```

Octave Derivation: Free Pitch Range

Store one octave of periods. Derive all others by bit-shifting:

- SRL D : RR E = one octave up (period halved, pitch doubled)
- SLA E : RL D = one octave down (period doubled, pitch halved)

Five pentatonic notes × one stored octave × bit-shifting = 5 notes × 5+ octaves = 25+ distinct pitches from 10 bytes of data. The formula selects the note, a separate bit mask selects the octave:

```
; octave = note_index / 5 (0-2)
; note = note_index % 5
; Look up base period, then SRL 'octave' times
```

Arpeggio: Chord Tones in Sequence

An arpeggio cycles through the tones of a chord. In scale-degree terms:

Chord	Scale offsets	Sound
Major triad	0, 2, 4 (root, third, fifth)	Bright, resolved
Minor triad	0, 2, 3 (root, m.third, fifth)	Dark, tense
Power chord	0, 4 (root, fifth)	Open, strong
Suspended	0, 3, 4 (root, fourth, fifth)	Ambiguous, floating

Implementation: `arp_step = (t / speed) % chord_size`, then add the offset to the current root note:

```

ld  a, (frame)
rrca
rrca                      ; A = frame / 4 (arp speed)
; mod 3 for three chord tones
ld  b, a
.mod3:
sub  3
jr   nc, .mod3
add  a, 3                  ; A = 0, 1, or 2
ld   hl, arp_major
add  a, l
ld   l, a
ld   a, (hl)              ; A = scale offset
; add to chord root, look up in scale table
; ...

```

```

arp_major: DB  0, 2, 4      ; root, third, fifth (3 bytes)
arp_minor: DB  0, 2, 3      ; root, min.third, fifth (3 bytes)

```

Three bytes per chord shape. The arpeggio speed is derived from the frame counter — no separate timer needed.

Step Ornaments: Trills, Mordents, and Slides

An ornament is a tiny cyclic pattern of relative pitch offsets applied to a note. In tracker music, ornaments make flat tones come alive:

Ornament	Pattern	Effect	Bytes
Trill	0, +1, 0, -1	Rapid alternation with neighbor	4
Mordent	0, +1, 0, 0	Brief upper neighbor, then settle	4
Slide up	0, 0, +1, +1	Gradual rise	4
Vibrato	0, +1, +1, 0, -1, -1	Smooth wobble	6

Apply by adding the ornament value to the note index before the scale table lookup:

```

ld  a, (frame)
and  $03                  ; mod 4 for 4-step ornament
ld  hl, trill

```

```

    add  a, l
    ld   l, a
    ld   a, (hl)      ; A = pitch offset (-1, 0, or +1)
    add  a, c          ; C = current note index
    ; ... look up modified note in scale table

trill:   DB  0, 1, 0, -1  ; 4 bytes
mordent: DB  0, 1, 0, 0   ; 4 bytes

```

Four bytes transform a static tone into a living voice. Stack multiple ornaments on different channels for rich texture.

Chord Progressions: Harmonic Movement

The chord root can change over time, following a progression. Classical harmony in 4 bytes:

```

progression: DB  0, 3, 4, 0      ; scale degrees

; Select chord: (frame / 64) AND 3
    ld  a, (frame)
    rrca
    rrca
    rrca
    rrca
    rrca
    rrca
    rrca
    rrca
    and  $03
    ld  hl, progression
    add  a, l
    ld  l, a
    ld  a, (hl)      ; A = chord root (scale degree)

```

Four bytes of progression data, cycled by the frame counter, give your AY-beat piece harmonic movement – the sense that it is “going somewhere” rather than looping on one chord. Other progressions:

Progression	Degrees	Bytes	Feel
I-IV-V-I	0, 3, 4, 0	4	Classic resolution
I-V-vi-IV	0, 4, 5, 3	4	Pop/rock standard
i-VI-III-VII	0, 5, 2, 6	4	Epic minor
I-I-I-I	0, 0, 0, 0	1 (or skip)	Drone/meditative

Total Data Budget for Rich Music

Combining all techniques:

Component	Bytes
Pentatonic table (5 notes)	10
Arpeggio pattern (1 chord)	3
Ornament (trill)	4

Component	Bytes
Progression (4 chords)	4
Total	21

21 bytes of musical data — plus ~45 bytes of engine code — produces three-channel music with melody, harmony, chord changes, and ornamentation. The `aybeat.a80` example in this book's companion code demonstrates this approach in 320 bytes, with room left over for visuals.

11. L-System Grammars: Fractal Melodies

Lindenmayer systems (L-systems) are rewriting grammars originally invented to model plant growth. Applied to music, they generate self-similar sequences with long-range structure from tiny rule sets.

The Concept

An L-system has an **axiom** (starting string) and **production rules** (expansion rules). Each iteration replaces every symbol according to its rule:

Axiom: A

Rules: A → A B, B → A

Step 0: A

Step 1: A B

Step 2: A B A

Step 3: A B A A B

Step 4: A B A A B A B A

This is the **Fibonacci L-system**. The sequence grows by the Fibonacci ratio (~1.618x per step). Map the symbols to musical events:

Symbol	Musical meaning
A	Play root note (scale degree 0)
B	Play fifth (scale degree 4)

The resulting melody: root, fifth, root, root, fifth, root, fifth, root... — a sequence that is neither periodic nor random, but *quasi-periodic*. It has structure at every scale, like a fractal. It sounds intentional without being repetitive.

Why L-Systems Work for Music

1. **Self-similarity.** The melody at large scales echoes the melody at small scales. This is what makes composed music feel coherent – themes recur at different levels.
2. **Non-repetition.** Unlike a looped pattern, an L-system sequence never exactly repeats (for irrational growth ratios). It stays interesting.

3. **Tiny encoding.** The rules are a few bytes. The sequence they generate is arbitrarily long.

Useful L-System Rules

Name	Axiom	Rules	Growth	Character
Fibonacci	A	A→AB, B→A	~1.618x	Quasi-periodic, organic
Thue-Morse	A	A→AB, B→BA	2x	Balanced, fair — no long runs
Period-doubling	A	A→AB, B→AA	2x	Increasingly syncopated
Cantor	A	A→ABA, B→BBB	3x	Sparse, with silences (B=rest)

Z80 Implementation

The trick for Z80 is to **not expand the string in memory** (that would require unbounded buffer space). Instead, compute the symbol at position *n* recursively: trace back through the rule applications to determine which original symbol position *n* came from.

For the Fibonacci L-system, there is an elegant shortcut. The symbol at position *n* depends on the Zeckendorf representation (Fibonacci coding) of *n*. But for practical sizecoding, a simpler approach works:

```
; Returns next note in sequence
; Uses position counter in memory
;
; The sequence of symbols can be generated iteratively:
; keep two "previous" bytes and generate the next

lsys_next:
    ld    hl, lsys_state
    ld    a, (hl)          ; prev1
    inc   hl
    ld    b, (hl)          ; prev2
    inc   hl
    ld    c, (hl)          ; position in current generation

    ; Fibonacci rule: output prev1, then swap
    ; When position reaches length, expand to next generation
    ld    d, a              ; D = current symbol to output

    ; Advance: shift the pair
    inc   c
    ld    (hl), c
    dec   hl
    ld    (hl), a           ; prev2 = prev1
    dec   hl
    ; New prev1 from rule: A→A (first output), then A→B (second)
    ; Simplified: alternate symbols based on parity
    ld    a, c
    and   $01
    jr    z, .sym_a
```

```

        ld  a, 1                ; B
        jr  .store
.sym_a:
        xor  a                  ; A (=0)
.store:
        ld  (hl), a

        ; Map symbol to scale degree
        ld  a, d
        or  a
        jr  z, .root
        ; B = fifth
        ld  a, 4                ; scale degree 4 = fifth in pentatonic
        ret
.root:
        xor  a                  ; scale degree 0 = root
        ret

lsys_state:
        DB  0                  ; prev1 (A=0, B=1)
        DB  0                  ; prev2
        DB  0                  ; position

```

A more practical approach for sizecoding: precompute several iterations of the L-system into a short buffer at init time (one iteration of Fibonacci from a 5-symbol axiom yields 8 symbols, two iterations yield 13, three yield 21 — all fitting in a small buffer), then loop through the buffer as a melody sequence:

```

; Axiom: "AABAB" (5 symbols) → 8 → 13 → 21 symbols
; 21 notes of fractal melody from 5 bytes of axiom + expansion code

```

```

lsys_expand:
        ld  hl, lsys_axiom
        ld  de, lsys_buf
        ld  b, 5                ; axiom length
.expand_iter:
        ; One iteration: for each symbol, apply rule
        push bc
        push hl
        ld  hl, lsys_buf
        ld  de, lsys_work      ; expand into work buffer
        ; ...expand according to rules...
        pop hl
        pop bc
        ; Copy work back to buf for next iteration
        ; Repeat for desired number of iterations
        ret

lsys_axiom:
        DB  0, 0, 1, 0, 1      ; A A B A B

```

```

; During playback:

```

```
; melody_index = frame / note_duration
; note = lsys_buf[melody_index % buf_length]
; look up in scale table → AY period
```

Melody as Motion, Not Absolute Notes

The most musical use of L-systems is not mapping symbols to fixed notes, but mapping them to **scale step directions**. A melody is fundamentally about *motion* — up, down, repeat, skip — on a scale. The starting note is arbitrary; the contour is what matters.

Define symbols as movements:

Symbol	Meaning	Scale step
U	Step up	+1
D	Step down	-1
R	Repeat	0
S	Skip up (leap)	+2

Now an L-system generates melodic *contour*, not fixed pitch sequences:

Axiom: U

Rules: U → U R D, D → U, R → U D

```
Step 0: U                      (+1)
Step 1: U R D                  (+1, 0, -1)
Step 2: U R D U D U            (+1, 0, -1, +1, -1, +1)
Step 3: U R D U D U U R D U U R D U D ...
```

The melody walks up and down the current scale, always staying within the scale table. It naturally tends toward the starting pitch (the returns balance the departures), creating the tension-and-resolution arc that makes music feel intentional.

```
; current_note = scale index, modified by each symbol
ld  a, (current_note)
ld  hl, lsys_buf
ld  b, (melody_pos)
add a, l
; ... get motion symbol at current position ...
; D = motion offset from symbol table
add a, d          ; current_note += motion
and $0F           ; wrap to scale range
ld  (current_note), a
; look up in pentatonic table → AY period
```

This is more musical than mapping A=root, B=fifth. The same L-system rules produce different melodies depending on the starting note and the underlying scale — change the scale from pentatonic to blues and the same contour produces a completely different mood.

Tribonacci: Three Symbols for Richer Patterns

The Fibonacci L-system uses two symbols. **Tribonacci** uses three: A→ABC, B→A, C→B. The growth ratio is ~1.839x (the tribonacci constant). Three symbols mean more varied melodic content:

Symbol	As motion	As note
A	Step up (+1)	Root
B	Repeat (0)	Third
C	Step down (-1)	Fifth

```
Axiom: A
Step 1: A B C
Step 2: A B C A B
Step 3: A B C A B A B C A B C
```

The tribonacci sequence has longer non-repeating runs than Fibonacci and a more complex internal structure. Musically, the three-symbol vocabulary gives melodies more variety — they don't just ping-pong between two states.

PRNG Melodies with Curated Seeds

A linear feedback shift register (LFSR) or similar PRNG generates a deterministic pseudo-random sequence from a seed value. The sequence *sounds* random but repeats exactly if you reset the seed. This gives you reproducible melody fragments.

The technique: **pre-test many seeds, keep the ones that sound good**. Store 2-4 seed values (2 bytes each) for different sections of your piece. At runtime, load the seed and let the PRNG generate the melody. The PRNG itself is ~6-8 bytes; each seed is 2 bytes.

```
; HL = seed (determines the melody)
prng_note:
    ld  a, h
    xor l          ; mix bits
    rrca
    rrca
    xor h
    ld  h, a
    ld  a, l
    add a, h
    ld  l, a        ; advance LFSR state (~6 bytes)
    and $07         ; constrain to scale range
    ret             ; A = note index for scale table

; Different seeds → different melodies
seed_verse:  DW  $A73B    ; tested: produces ascending contour
seed_chorus: DW  $1F4D    ; tested: produces energetic pattern
seed_bridge: DW  $8E21    ; tested: produces descending, calm
```

The workflow: write a test harness that plays the PRNG melody for each seed value 0-65535, listen (or analyse), mark the good ones. In practice, a few hours of testing

yields dozens of usable seeds. Store 3-4 of them and switch between sections of your piece.

Combining with scale tables: the PRNG output feeds through the pentatonic table, so even “bad” seeds produce consonant notes. You’re curating for *melodic contour*, not avoiding wrong notes — the scale table already handles that.

Combining with L-systems: use the PRNG to *select which L-system rule to apply* at each step, creating stochastic L-systems. The seed controls the “personality” of the piece; the grammar rules control the structure. This hybrid produces the richest output from the fewest bytes.

Combining L-Systems with Other Techniques

L-systems generate note *sequences*. Combine with the other techniques from this appendix:

- **Scale table** maps L-system symbols to actual AY periods
- **Ornaments** add expression to each note
- **Arpeggio** turns each L-system note into a chord
- **Envelope drone** provides a sustained harmonic bed under the fractal melody
- **Chord progression** changes the root — the L-system melody is transposed to each chord

The result: a tiny program (~60-80 bytes of music code + 20 bytes of data) generating minutes of structurally coherent, non-repeating, harmonically grounded music. This is algorithmic composition, not random noise — and it fits in a size-coded intro.

Other Grammars for Music

Beyond L-systems, other formal grammars produce interesting musical sequences:

Cellular automata. Rule 30 or Rule 110, applied to a row of bits, produce complex patterns. Map bit positions to note on/off events. Cost: ~15 bytes for the CA rule, ~20 bytes for the stepper.

Euclidean rhythms. Distribute k beats evenly across n steps. This algorithm (related to the Euclidean GCD) generates rhythmic patterns found in music worldwide: 3-in-8 is tresillo, 5-in-8 is cinquillo, 7-in-12 is a common West African bell pattern. Implementation is ~20 bytes and produces perfect rhythmic foundations for any AY-beat engine.

See Also

- **Chapter 11** – AY-3-8910 architecture, tone/noise/envelope theory, buzz-bass technique
- **Chapter 12** – Music engine integration, sync to effects, hybrid digital drums
- **Chapter 13** – Sizecoding techniques, where AY-beat fits in the 256b/512b/1K/4K size tiers
- **Appendix G** – Complete AY register reference with bit layouts, port addresses, and note tables

Sources: Viznut (Ville-Matias Heikkila), “Algorithmic symphonies from one line of code – how and why?” (2011); countercomplex.blogspot.com; Chapter 13 of this book; various 256-byte ZX Spectrum intros from Pouet.net

Appendix J: Modern Tools for Retro Demo Production

“The best tool is the one that gets the data out of your head and into the Z80’s memory in the fewest steps.”

J.1 Two Worlds, One Philosophy

The PC demoscene and the Z80 demoscene look like different planets. On one planet, coders write fragment shaders in GLSL, generate meshes procedurally through node graphs, and compress everything into 64K executables that render real-time 3D at 60fps on consumer GPUs. On the other planet, coders hand-craft Z80 assembly routines that fit fourteen effects into 128K of banked RAM, counting T-states to ensure each frame renders in exactly 71,680 cycles.

The gap in hardware capability is enormous. The gap in philosophy is almost zero.

Both worlds worship constraint. A 64K PC intro is constrained by file size just as ruthlessly as a ZX Spectrum demo is constrained by CPU speed. A Shader Showdown competitor writing a raymarcher in 25 minutes faces the same creative pressure as a 256-byte coder optimising register reuse. Procedural generation — building content from algorithms rather than storing it as data — is the central technique on both platforms, because both platforms punish data and reward computation (the PC punishes stored data through file size limits; the Spectrum punishes it through memory limits).

This appendix surveys the modern demoscene toolchain: sync editors, data generators, music synthesisers, executable packers, and shader tools. Most of these tools target x86/GPU platforms and cannot run on a Z80 directly. But they serve three roles in ZX Spectrum production:

1. **Data generators** — compute trajectories, particle positions, procedural textures on a PC, then export the results as compressed binary tables that the Z80 plays back.
2. **Sync planners** — design the timing relationship between music and visuals in an interactive editor, then export frame numbers and parameter curves to Z80 dw tables.
3. **Prototyping environments** — test visual algorithms at full speed on a modern GPU, then translate the working algorithm to Z80 assembly knowing exactly what the target output should look like.

The philosophy is consistent: **use any tool for preparation, but the Z80 runtime is handwritten assembly**. The demo's quality is judged by what runs on the Spectrum, not by what runs on the development PC. The tools are scaffolding; the building is Z80.

J.2 Sync Tools

The Sync Problem

Synchronisation — making the right visual event happen at the right musical moment — is the hardest part of demo production (Chapter 20). At the Z80 level, sync is always a data table: frame numbers paired with actions. The question is how to determine those frame numbers efficiently and iterate on them quickly.

Five tools address this problem, each with a different workflow.

GNU Rocket

What it is. GNU Rocket is a sync-tracker — a tracker-like editor where columns represent named parameters (`camera:x`, `fade:alpha`, `effect:id`) and rows represent time steps (typically frames or musical beats). You set keyframes at specific rows and choose an interpolation mode between them: **step** (instant change), **linear** (constant rate), **smooth** (cubic ease in/out), or **ramp** (exponential). The demo connects to the Rocket editor via TCP during development. You scrub the timeline, edit values, and see the demo update in real-time.

Who uses it. Rocket is the de facto standard sync tool across the PC and Amiga demoscenes. Logicoma, Noice, Loonies, Adapt, and dozens of other groups use it. It has been ported to C, C++, C#, Python, Rust, and JavaScript.

Z80 workflow. A Z80 Spectrum cannot run a TCP client, but the concept transfers directly:

1. Design sync tracks in Rocket on the PC, scrubbing with the music
2. Export keyframe data as binary (Rocket's native export)
3. Run a Python converter: quantise floats to 8-bit or 16-bit fixed-point, emit db/dw tables
4. INCBIN the tables into the demo

The Z80 code just reads a table — no TCP, no floats, no complexity. You get Rocket's interactive editing experience during development, and the Spectrum's minimal runtime overhead in the final binary.

Interpolation. Rocket's four interpolation modes map cleanly to Z80 playback: - **Step** → just use the value directly (0 cycles of interpolation overhead) - **Linear** → precompute the per-frame delta, add it each frame (~20 T-states) - **Smooth/Ramp** → bake the interpolated curve into the exported table (the Z80 reads precalculated values, no interpolation at runtime)

For most ZX Spectrum demos, baking all curves to per-frame values is the simplest approach. A 3,000-frame demo (one minute at 50fps) with 4 sync parameters consumes 12KB of uncompressed data — significant, but compressible to 2-4KB with ZX0.

Source: github.com/rocket/rocket (MIT-like license)

FIGURE: GNU Rocket sync editor
Tracker-like grid with named columns: [camera:x] [camera:y] [fade:alpha] [effect:id]
Rows = frames/time steps. Keyframes shown as bright cells. Between keyframes: interpolation curves (step/linear/smooth/ramp) visualised as lines connecting values.
Bottom: transport controls (play, pause, scrub). Connected to running demo via TCP – edit live.
Screenshot needed: build GNU Rocket from source, create example project with 4 tracks, capture at a point showing all 4 interpolation modes.

Vortex Tracker II

What it is. Vortex Tracker II is the standard ProTracker 3 editor for the ZX Spectrum scene. The critical feature for sync work is the **frame counter** in the bottom-right corner of the window — it shows the absolute interrupt count (frame number) at the current playback position.

Sync workflow. Play the .pt3 file. Watch the frame counter. When you hear a beat, accent, or phrase transition you want to sync to, note the frame number. Write it into your sync table. Rebuild the demo, test, adjust.

Kolnogorov (Vein) describes this as his primary method: “Vortex + video editor. In Vortex the frame is shown in the bottom-right corner — I looked at which frames to hook onto, created a table with dw frame, action entries, and synced from that.”

The VTI fork. The community maintains VTI, a fork of Vortex Tracker II with additional features including improved playback accuracy and expanded format support. For sync work, the original VT2 and VTI are equivalent — the frame counter works the same way.

Limitation. Iterating is slow. Every timing change requires rebuilding the demo and watching it from the beginning. For simple demos with a dozen sync points, this is fine. For complex demos with hundreds of events, a more interactive tool (Rocket, Blender) is worth the setup cost.

FIGURE: Vortex Tracker II – frame counter
VT2 main window with pattern editor visible. Bottom-right: position display showing pattern:row and absolute frame number. Highlight/circle the frame counter.
Caption: "The frame number in VT2's status bar maps directly to the PT3 player's interrupt counter on the Spectrum. What you see here is what your sync table references."

Screenshot needed: open any .pt3 in VTI fork, play to a mid-song position, capture with frame number visible.
--

Blender VSE (Video Sequence Editor)

What it is. Blender's built-in non-linear video editor. For demo sync, it provides a timeline where you can lay colour-coded strips (one per effect), import the music track as an audio strip with a visible waveform, and place markers at sync points.

Sync workflow: 1. Capture each effect running in the emulator as a short video clip (or use solid-colour placeholder strips) 2. Import the clips and the music (.wav) into the VSE 3. Arrange clips on the timeline, scrub to find the perfect cut points 4. Place markers (green diamonds) at each sync event 5. Read off the frame numbers from the marker positions

The visual workflow is powerful: you *see* the audio waveform, you *see* where the beat hits, and you drag the marker to the right spot. No calculations, no BPM-to-frame conversions.

Export. Markers can be exported via Blender's Python console:

```
print(f"dw {m.frame}, 0 ; {m.name}")
```

This outputs a Z80-ready sync table directly. Rename markers to match your effect IDs (plasma, flash, scroll) and the export becomes a complete scene table.

DaVinci Resolve provides similar capabilities (timeline + markers + EDL/CSV export) and its free version is sufficient. Choose whichever you already know.

FIGURE: Blender VSE – demo storyboard

Timeline with 4-5 colour-coded strips (one per effect): [TORUS: blue] [PLASMA: green] [DOTSCROLL: yellow] [ROT0ZOOM: red] Below: audio waveform strip (music.wav) Vertical markers (green diamonds) at sync points. Playhead at a transition point between effects. Screenshot needed: create Blender project with dummy strips + real AY music exported as WAV. Place ~8 markers at beat hits.
--

Blender Graph Editor

What it is. Blender's curve editor for keyframed properties. For demo work, you create custom properties on an object (e.g., scroll_speed, fade_alpha, camera_z) and keyframe them to match the music's energy and phrasing.

Why it matters. The Graph Editor gives you visual, interactive control over how parameters change over time — the same thing GNU Rocket provides, but integrated into Blender's ecosystem. You can see multiple curves simultaneously, adjust keyframe timing by dragging, and switch interpolation modes (constant, linear, Bezier) per keyframe.


Export via Python API:

```

name = fcurve.data_path.split(' ')[1]
print(f"; {name}")
for kf in fcurve.keyframe_points:
    print(f"    dw {int(kf.co.x)}, {int(kf.co.y)}")

```

This outputs Z80-ready sync data directly. The Blender project becomes your storyboard, your sync reference, and your data pipeline in one file.

**Motion Canvas**

A brief mention: **Motion Canvas** is an emerging MIT-licensed, TypeScript-based tool for creating parametric animations programmatically. It is designed for explanatory videos but its timeline-as-code approach could serve as a sync planning tool for coders who prefer writing code to dragging keyframes. Still early in development; watch the project at motioncanvas.io.

Comparison Table

Tool	License	Sync method	Z80 export path	Interactive?
GNU Rocket	MIT-like	Keyframed tracks + interpolation	Binary → Python → dw tables	Yes (TCP live editing)
Vortex Tracker II	Freeware	Frame counter readout	Manual (note frame numbers)	Partial (listen + note)
Blender VSE	GPL	Timeline markers + waveform	Python → dw tables	Yes (visual scrubbing)
Blender Graph Editor	GPL	Keyframed curves	Python API → dw tables	Yes (visual curve editing)

Tool	License	Sync method	Z80 export path	Interactive?
DaVinci Re-solve	Free/Commercial	Timeline markers	EDL/CSV → Python → dw tables	Yes (visual scrubbing)
Motion Canvas	MIT	Code-defined timeline	TypeScript export (custom)	Programmatic

J.3 Pre-visualisation and Storyboarding

Before writing a single line of Z80 assembly, you need to know what the demo looks like. Not in detail — not every palette choice or scroll speed — but the overall structure: which effects, in what order, for how long, transitioning how. This is pre-visualisation, and it saves more time than any code optimisation.

Blender as Storyboard Tool

Blender’s **Grease Pencil** (now fully integrated as a 2D drawing system) lets you sketch rough representations of each effect — coloured rectangles, simple shapes, hand-drawn approximations. These do not need to look like the final effects. They need to communicate: “here, for this many seconds, something blue and swirly happens.”

A practical approach: 1. Create one Grease Pencil object per effect (coloured rectangle with the effect’s name) 2. Arrange them on the VSE timeline with the music track 3. Play back and watch the structure

You are not making an animatic. You are making a *schedule* — a visual representation of which effect runs when, and how long it runs for. The colours help you spot structural problems: three slow blue effects in a row means the pacing is off. A 45-second stretch with no transition means the audience gets bored.

Video Editors for Rough-cut Assembly

Once you have effects running in an emulator, capture them as video clips (OBS Studio, screen recorder, or emulator recording features). Import the clips into a video editor — DaVinci Resolve, Blender VSE, or even iMovie — and arrange them on a timeline with the music.

This rough cut reveals problems that are invisible in code: - An effect that looks impressive in isolation but runs too long in context - A transition that should happen on the beat but lands between beats - Two consecutive effects with similar colour palettes that blur together - A section where the visual energy drops while the music builds

The rough cut costs an hour of work. It saves days of assembly-level timing adjustments.

The GABBA Workflow

diver4d's approach for GABBA (2019) took this further: he used **Luma Fusion**, an iOS video editor, as his primary sync planning tool. The workflow was:

1. Code each effect individually, test in emulator
2. Screen-record each effect running
3. Import recordings + music into Luma Fusion on an iPad
4. Arrange on the timeline, scrub frame-by-frame to find sync points
5. Note the frame numbers, write sync table

The key insight: frame-level synchronisation is a *video editing* problem, not a *programming* problem. By solving it in a tool designed for video editing, diver4d could iterate on timing in seconds instead of minutes. The Z80 code was the implementation layer; the creative decisions happened in the video editor.

Kolnogorov on Sync Planning

Kolnogorov (Vein) articulates the combined approach: "I exported effect clips to video, assembled them in a video editor, attached the music track, and looked at what order the effects work best in, noting the frames where events should happen."

The important word is *looked*. This is a visual, intuitive process. You *see* where the beat hits the waveform. You *see* where the effect transition feels right. And you read the frame number. No calculations, no BPM-to-frame conversions.

J.4 Data Generation Engines

A ZX Spectrum demo plays back precomputed data at 50fps. The question is: where do you compute that data? For simple sine tables and Bresenham coordinates, a Python script suffices. For complex 3D trajectories, organic particle motion, or VR-captured gestures, you need a more powerful environment.

Unity as a Data Generator

Unity is not overkill for ZX Spectrum demos — it is overkill as a *demo engine*, but perfect as a *data generator*. The distinction matters.

VR motion capture. Unity's XR Toolkit captures VR controller position at 90Hz. Draw a trajectory in the air with a VR controller — the organic movement you get from a hand gesture is impossible to replicate with mathematical formulas. Downsample to 50fps, quantise to 8-bit signed values, delta-encode, compress. A 5-second hand-drawn trajectory becomes 250 bytes of packed data that *feels alive* on the Spectrum.

GPU particle systems. Unity's VFX Graph runs millions of particles on the GPU. Prototype a particle fountain, a vortex, or a flocking simulation, then export the particle positions per frame as CSV. On the Spectrum, you plot those positions as dots or attribute cells. The physics simulation that would take months to implement in Z80 runs in milliseconds on a GPU.

Shader prototyping. Write a plasma, tunnel, or rotozoomer as a fragment shader in Unity’s ShaderGraph. Iterate in real-time at full resolution until it looks right. Then translate the algorithm to Z80, knowing exactly what the target visual should be.

Unreal Engine as a Data Generator

Unreal offers equivalent capabilities through different tools: - **OpenXR** for VR motion capture - **Niagara** for GPU particle systems - **Material Editor** for shader prototyping

The choice between Unity and Unreal for data generation is a matter of familiarity. Both export to the same formats (CSV, binary arrays). Both provide more computational power than you will ever need for generating Z80 demo data.

Blender as a Data Generator

For most data generation tasks, Blender is sufficient and fully open-source:

- **Geometry Nodes** for procedural trajectories — define a path as a spline, scatter points along it, animate with noise, export vertex positions per frame
- **Grease Pencil** for hand-drawn animation — draw frames, export point coordinates
- **Python API** for direct export — access any animated property from bpy.data
- **Physics simulation** for particles, cloth, fluid — simulate, bake, export per-frame data

Blender cannot do VR motion capture (without third-party addons) and its particle system is CPU-bound (slower than GPU-based systems for millions of particles). For everything else, it matches or exceeds Unity/Unreal for data generation purposes.

The Export Pipeline

Regardless of which tool generates the data, the pipeline to Z80 follows the same stages:

- Export: float arrays (CSV, JSON, or binary)
- Python: float → 8-bit fixed-point (or 16-bit where needed)
- Delta-encode: store differences between frames (smaller values)
- Transpose: column-major layout (all X values, then all Y values)
- packbench analyze: verify entropy, check suggested transforms
- zx0/pletter: compress
- sjasmpus: INCBIN into demo

Each stage reduces the data size:

Stage	Example size (250 frames × 4 params)
Float CSV	~10 KB
8-bit fixed-point	1,000 bytes
Delta-encoded	1,000 bytes (values shrink, better compression ratio)

Stage	Example size (250 frames × 4 params)
Transposed + compressed	300-500 bytes

The delta-encoding and transposition stages do not reduce raw size — they reshape the data to compress better. Column-major layout groups similar values together (all X deltas, then all Y deltas), which compresses dramatically better than row-major layout where X and Y deltas alternate.

When to Use Which

Need	Tool	Why
Sine tables, lookup tables	Python script	Simplest, no dependencies
Procedural 3D trajectories	Blender (Geometry Nodes)	Free, visual, Python export
Hand-drawn animation paths	Blender (Grease Pencil)	Draw directly, frame-by-frame
VR gesture capture	Unity (XR Toolkit) or Unreal (OpenXR)	Need VR hardware + runtime
GPU particle positions	Unity (VFX Graph) or Unreal (Niagara)	Millions of particles, fast
Shader algorithm prototype	Any (ShaderToy, Unity, Unreal, Blender)	Whichever you know

Kolnogorov’s precalculated vector animations are a case in point: the 3D geometry was computed offline (minutes of calculation are acceptable for a 4K intro), the resulting vertex trajectories were stored as compressed tables, and the Spectrum played them back at 50fps. The tool that generated the trajectories is irrelevant to the audience. What matters is that the data exists and the Z80 plays it.

J.5 The PC Demoscene Toolchain: A Brief History

The PC demoscene has spent twenty-five years building tools for procedural content generation and extreme compression. These tools cannot run on a Z80, but their design philosophy — procedural everything, small is beautiful, constraint as creativity — mirrors exactly what ZX Spectrum coders do by hand. Understanding the PC toolchain helps you recognise which problems have been solved (in different contexts) and which ideas you can adapt.

Farbrausch (1999-2012)

Farbrausch was a German demogroup that redefined what was possible in 64K and 4K executables. Their approach: build tools that generate everything procedurally, then pack the generators into tiny executables.

Werkkzeug (versions 1 through 4) was their flagship tool — a node-based procedural system where textures, meshes, animations, and compositions were defined as operator graphs. No bitmaps were stored; every pixel was computed at run-time from a recipe of mathematical operations. The tool was used internally by Farbrausch; the audience saw only the resulting executables.

Notable productions built with Werkkzeug: - **fr-08: .the .product** (2000) — a real-time 3D music video in 64KB that redefined the 64K intro category. Won the demo compo at The Party 2000. - **.kkrieger** (2004) — a first-person shooter in 97,280 bytes. Textures, meshes, animations, AI, and sound, all procedurally generated. The executable is smaller than this appendix's source file. - **fr-041: debris.** (2007) — a cinematic 177KB demo that pushed real-time rendering quality beyond what many full-size demos achieved. Won at Breakpoint 2007.

kkrunchy was their executable packer for 64K intros — a context-mixing compressor that achieved compression ratios far beyond standard packers like UPX. kkrunchy is still actively used by 64K intro coders today, more than a decade after Farbrausch disbanded.

V2 Synthesizer was a software synthesizer designed for intros — a VSTi plugin for composing, with a tiny runtime player that fit inside the packed executable. Music was stored as note data and synthesis parameters, not as audio.

In 2012, Farbrausch open-sourced their entire toolchain under a BSD-style license: github.com/farbrausch/fr_public. The repository includes Werkkzeug (all versions), kkrunchy, V2, and several other tools. The code is a time capsule of early-2000s demoscene engineering: dense C++, Win32 APIs, Direct3D 9, and creative solutions to problems that modern GPUs solve with brute force.

Relevance to Z80. The ZX Spectrum equivalent of Werkkzeug's procedural textures is your Python build script that generates lookup tables from mathematical functions. The equivalent of kkrunchy is ZX0 or Pletter. The equivalent of V2 is AY-beat or Shiru's AY player. Different scale, same principle: generate content from compact descriptions, compress the result, decompress at runtime.

TiXL (2024-present)

TiXL (formerly Tooll3) is the spiritual successor to Werkkzeug, developed by **Still** (pictur/Thomas Mann) — a demoscener who has worked with Farbrausch members and carried the node-based procedural approach forward.

TiXL is a real-time motion graphics environment built on modern GPU APIs. Like Werkkzeug, it uses a node graph to define procedural content, but with 20 years of GPU evolution behind it: compute shaders, physically-based rendering, GPU particles, and real-time raymarching.

Licensed under MIT (github.com/tixl3d/tixl), TiXL shows where the Werkkzeug philosophy has evolved. The node graph concept — defining content as a recipe of operations rather than storing it as data — is directly applicable to Z80 demo development, even though the specific operations are entirely different.

	FIGURE: TiXL node graph	
	Visual programming canvas with connected nodes:	
	[Time] → [Sine] → [Multiply] → [SetFloat]	

```

A procedural texture pipeline:
[Noise3D] → [Remap] → [ColorGrade] → [RenderTarget]

3D scene graph:
[Mesh:Torus] → [Transform] → [Material] → [DrawMesh]
[Camera] → [Render] → [PostFX:Bloom] → [Output]

Right panel: live preview of the rendered output.
Bottom: timeline with playback controls.

Caption: "TiXL (MIT, 2024) carries forward the Werkzeug
philosophy of node-based procedural content generation. The
Z80 equivalent is a Python build script that generates lookup
tables and compressed data from mathematical functions."

Screenshot needed: install TiXL, open an example project,
capture the node graph + preview in a split view.

```

Bonzomatic

Bonzomatic (by Gargaj / Conspiracy) is the standard tool for **Shader Showdown** competitions — live shader coding battles at demoscene parties. Two coders share a stage, each writing a fragment shader from scratch in 25 minutes, with the code and output projected side-by-side for the audience.

Bonzomatic provides a minimal editor with a live preview window. Every keystroke recompiles the shader; the output updates in real-time. There is no save, no undo history beyond the editor buffer, and no libraries. It is pure constraint-driven creativity — the same spirit as 256-byte coding on the Spectrum.

Source: github.com/Gargaj/Bonzomatic

Crinkler

Crinkler (by Rune Stubbe and Aske Simon Christensen / Loonies) is a **compressing linker** for Windows executables. Where a normal linker produces an .exe and you compress it separately, Crinkler combines linking and compression into one step, achieving compression ratios that no separate packer can match.

Crinkler is the standard for 1K, 4K, and 8K PC intros. Its compression is so effective that coders routinely measure their work in “Crinkler bytes” — the final packed size, which may be 60–70% of the uncompressed code + data.

The Z80 parallel: ZX0 and Pletter serve the same role for Spectrum demos, though compression ratios are more modest (Z80 code has less redundancy than x86).

Source: github.com/runestubbe/Crinkler (zlib license)

Squishy

Squishy (by Logicoma) is a 64K executable packer — the equivalent of kkrunchy for modern 64K intros. Unlike Crinkler (which handles tiny intros), Squishy targets

the 64KB size class where the executable contains substantial shader code, texture generation routines, and music.

Binary-only distribution; no source available. Mentioned here because Logicoma's 64K intros (Happy Coding, Elysian, H - Immersion) represent the current state of the art in the 64K category, and Squishy is part of that pipeline.

Shader Minifier

Shader Minifier (by Ctrl-Alt-Test) is a GLSL/HLSL minifier that renames variables, removes whitespace, and optimises shader code for minimum size. Used in size-coded intros where every byte of shader source counts (shaders are often stored as strings in the executable and compiled at runtime).

Source: github.com/laurentlb/shader-minifier

z80-optimizer

z80-optimizer (by oisee, 2025) is a brute-force Z80 superoptimiser written in Go. It enumerates every pair of Z80 instructions (406×406 opcodes), tests each pair against all possible register and flag states, and reports when a shorter or faster replacement produces identical output. No heuristics, no machine learning — pure exhaustive search with full state equivalence verification.

A single run on an Apple M2 (3h 16m, 34.7 billion comparisons) produces **602,008 provably correct optimisation rules** across **83 unique transformation patterns**. Examples: `SLA A : RR A → OR A` (saves 3 bytes, 12T); `LD A, 0 : NEG → SUB A` (saves 2 bytes); `SCF : RR A → SCF : RRA` (saves 1 byte, 4T). Crucially, it correctly *rejects* `LD A, 0 → XOR A` because the flag behaviour differs — the kind of subtle distinction that hand-maintained peephole tables sometimes get wrong.

Useful as a post-processing pass for compiler-generated Z80 code, or as a reference for hand-optimisation. The output is a machine-readable rule database that can be integrated into assembler toolchains. See Chapter 23 for discussion of brute-force vs neural-network approaches to Z80 optimisation.

Source: github.com/oisee/z80-optimizer (MIT license, v0.1.0)

The Common Philosophy

All these tools share a principle that translates directly to Z80 work: **procedural generation beats stored data**. On the PC, this means generating textures from noise functions instead of storing bitmaps. On the Spectrum, it means computing a sine table from a parabolic approximation instead of storing 256 bytes of precomputed values. The hardware differs by a factor of a million; the approach is the same.

J.6 Music and Sound Tools

The PC demoscene needs tiny synthesisers — software instruments that fit inside a 4K or 64K executable while producing music that sounds professional. The Z80 scene needs AY-3-8910 trackers. There is some overlap.

Sointu

Sointu is a 4K intro synthesiser — a fork of 4klang rewritten in Go for cross-platform support. It provides a VST-like interface for designing patches (oscillators, filters, envelopes, effects) and a compiler that generates a minimal native player.

The player code is designed for extreme size: the entire synth runtime, including all instrument definitions and note data, fits in under 4KB of x86 code + data. This makes it the PC equivalent of what AY-beat (Chapter 13) achieves on the Spectrum: maximum sound from minimum bytes.

Source: github.com/vsariola/sointu (MIT license)

4klang

4klang (by Alcatraz) is the original 4K intro synth that Sointu forked from. It provides a VSTi plugin for composing and an assembler output mode that generates NASM source code for the player — the synthesiser is literally written in x86 assembly, optimised for size.

4klang defined the standard for 4K intro music and remains in active use alongside its successor Sointu.

Source: github.com/gopher-atz/4klang

WaveSabre

WaveSabre (by Logicoma) targets 64K intros — a larger size class where the synth can be more sophisticated. It provides VST-compatible instruments with reverb, delay, chorus, distortion, and other effects that 4K synths must omit. The runtime player is compact enough for 64K but too large for 4K.

WaveSabre powers the music in Logicoma's award-winning 64K intros (Happy Coding, Elysian).

Source: github.com/logicomacorp/WaveSabre (MIT license)

Oidos

Oidos (by Blueberry / Loonies) takes a fundamentally different approach: additive synthesis. Where most synths build sounds from oscillators and filters, Oidos builds them from sums of sine waves with individually controlled frequencies and amplitudes. The result is a distinctive, rich sound that occupies a unique sonic space in 4K intros.

Source: github.com/askeksa/Oidos

Furnace

Furnace is a multi-system chiptune tracker that supports over 80 sound chips — including the **AY-3-8910**. This makes it directly relevant to ZX Spectrum demo production: you can compose AY music in Furnace using a modern interface with features that Vortex Tracker II lacks (undo, multiple tabs, visual envelope editor, per-channel oscilloscope).

Additional chips supported include the SN76489 (Sega Master System, BBC Micro), YM2612 (Mega Drive), SID (C64), Pokey (Atari), and many others. If you are working on multi-platform retro projects, Furnace is the one tracker that covers all targets.

Export formats. Furnace can export to VGM (Video Game Music) format, which captures raw register writes per frame. For the AY-3-8910, this means you get a stream of 14-byte register dumps at 50/60Hz — directly usable on the Spectrum with a minimal player that writes registers each interrupt. Custom export scripts can convert VGM to PT3 or to raw register tables for INCBIN.

Limitation. Furnace’s AY emulation is good but not identical to Vortex Tracker’s PT3 playback. If your demo uses a PT3 player, compose in Vortex Tracker for guaranteed 1:1 playback. Use Furnace when you want a modern editing experience and are willing to handle format conversion, or when targeting the AY directly with raw register writes.

Source: github.com/tildearrow/furnace (GPL-2.0)

Comparison Table

Tool	Size class	License	Z80 relevance
Sointu	4K intros	MIT	Conceptual (same constraint as AY-beat)
4klang	4K intros	OSS	Conceptual
WaveSabre	64K intros	MIT	Conceptual
Oidos	4K intros	OSS	Conceptual
Furnace	Any	GPL-2.0	Direct — AY-3-8910 support, VGM export
Vortex Tracker II	Any	Freeware	Direct — PT3 native, ZX standard

J.7 Practical Recipes

Five step-by-step workflows that bridge modern tools to Z80 assembly. Each recipe starts from a blank project and ends with data you can INCBIN into a demo.

Recipe 1: GNU Rocket → Z80 Sync Table

Goal: Design interpolated sync curves in Rocket, export as Z80 dw tables.

Steps:

1. **Install Rocket.** Clone github.com/rocket/rocket, build from source (CMake). Run the editor.
2. **Create tracks.** Add four tracks: `effect:id`, `fade:alpha`, `scroll:speed`, `flash:border`. Set the BPM to match your music (e.g., 125 BPM at 50fps = 24 frames per beat).

3. **Set keyframes.** At row 0: effect:id = 0 (logo), fade:alpha = 255. At row 150: effect:id = 1 (plasma), fade:alpha ramp from 255 to 0 (smooth interpolation). Continue for the full demo length.
4. **Export.** Use Rocket's sync exporter (or the sync_export tool) to write binary files per track.
5. **Convert with Python:**

```
def rocket_to_z80(track_file, output_file):
    with open(track_file, 'rb') as f:
        data = f.read()
        # Rocket binary: array of (row: u32, value: float32) pairs
        pairs = []
        for i in range(0, len(data), 8):
            row, val = struct.unpack('<If', data[i:i+8])
            pairs.append((row, max(0, min(255, int(val)))))

        with open(output_file, 'w') as f:
            for row, val in pairs:
                f.write(f"    dw {row}, {val}\n")
            f.write("    dw 0 ; end marker\n")
```

```
rocket_to_z80('effect_id.track', 'sync_effect.inc')
```

6. Include in demo:

```
INCLUDE "sync_effect.inc"
```

Baking option. For smooth interpolation without runtime cost, expand the keyframes to per-frame values:

```
for frame in range(total_frames):
    value = interpolate(keyframes, frame) # linear, smooth, or ramp
    print(f"    db {max(0, min(255, int(value)))}")
```

A 3,000-frame track baked to per-frame db values = 3,000 bytes uncompressed, ~800 bytes after ZX0.

Recipe 2: Blender VSE → Frame Number Table

Goal: Plan demo structure visually with music, export sync points.

Steps:

1. **Create Blender project.** Set frame rate to 50fps (Properties → Output → Frame Rate → 50). Set the timeline length to match your music.
2. **Import music.** Add → Sound in the VSE. Import your .pt3 exported as .wav (use Vortex Tracker's "File → Export to WAV"). The waveform appears on the timeline.
3. **Add effect strips.** Add → Color strips for each effect. Name them (plasma, scroll, torus). Colour-code them. Arrange on the timeline so they cover the full demo duration.

4. **Place markers.** Scrub the timeline. When you hear a beat or transition point, press M to add a marker. Rename each marker: plasma_start, flash_1, scroll_begin, etc.
5. **Export via Python console:**

```
# Print sync table
print("; Auto-generated sync table from Blender VSE")
print("sync_table:")
for m in sorted(bpy.context.scene.timeline_markers,
                key=lambda x: x.frame):
    print(f"    dw {m.frame} ; {m.name}")
print("    dw 0 ; end marker")
```

6. **Copy output to your project's .inc file.** Rebuild the demo with make.
7. **Iterate.** Watch the demo, note where timing feels wrong, return to Blender, adjust markers, re-export. Each iteration takes seconds.

Recipe 3: Unity VR → Trajectory Data

Goal: Capture a hand-drawn 3D trajectory using a VR controller, export as compressed Z80 data.

Steps:

1. **Set up Unity project.** Create new 3D project, install XR Toolkit via Package Manager. Configure for your headset (Quest, Index, etc.).
2. **Record controller path.** Create a script that captures transform.position of the right controller every frame:

```
positions.Add(new Vector3(
    controller.transform.position.x,
    controller.transform.position.y,
    controller.transform.position.z
));
}
```

3. **Export CSV.** On recording end, write positions to CSV:

```
positions.Select(p => $"{p.x},{p.y},{p.z}");
```

4. **Convert with Python:**

```
with open('trajectory.csv') as f:
    rows = list(csv.reader(f))

# Downsample from 90Hz to 50Hz
factor = 90 / 50
resampled = [rows[int(i * factor)] for i in range(int(len(rows) / factor))]

# Quantise to signed 8-bit (-128..127)
def q8(val, scale=64.0):
    return max(-128, min(127, int(float(val) * scale)))
```

```

# Delta-encode
prev = [0, 0, 0]
deltas = []
for row in resampled:
    cur = [q8(row[0]), q8(row[1]), q8(row[2])]
    deltas.append([c - p for c, p in zip(cur, prev)])
    prev = cur

# Transpose (column-major) and output
print("; X deltas")
print("traj_dx:")
print("    db " + ", ".join(str(d[0] & 0xFF) for d in deltas))
print("; Y deltas")
print("traj_dy:")
print("    db " + ", ".join(str(d[1] & 0xFF) for d in deltas))
print("; Z deltas")
print("traj_dz:")
print("    db " + ", ".join(str(d[2] & 0xFF) for d in deltas))

```

5. **Compress.** Run each column through ZX0 separately (column-major layout compresses much better than interleaved). INCBIN the compressed blobs.
6. **Playback on Z80.** Decompress each column to a buffer. Each frame, read the next delta and add to the current position. Plot the resulting (X, Y) on screen (project Z for depth if needed).

Recipe 4: Furnace → AY Music

Goal: Compose AY-3-8910 music in Furnace and export for ZX Spectrum playback.

Steps:

1. **Configure Furnace.** Create new project. Add an AY-3-8910 system (Settings → select “AY-3-8910” from the chip list). Set clock to 1.7734 MHz (ZX Spectrum standard). Set refresh rate to 50Hz (PAL).
2. **Compose.** Use Furnace’s pattern editor — similar to Vortex Tracker but with undo, per-channel oscilloscope, and visual envelope editing. Furnace supports AY hardware envelopes, noise mixing, and all standard AY features.
3. **Export as VGM.** File → Export → VGM. This produces a .vgm file containing raw AY register writes per frame — a stream of (register, value) pairs at 50Hz.
4. **Convert VGM to register dump:**

```

def vgm_to_ay_regs(vgm_file):
    with open(vgm_file, 'rb') as f:
        data = f.read()

    # Skip VGM header (find data offset at 0x34)
    data_offset = struct.unpack_from('<I', data, 0x34)[0] + 0x34

```

```

frames = []
current_regs = [0] * 14
pos = data_offset

while pos < len(data):
    cmd = data[pos]
    if cmd == 0xA0: # AY-3-8910 register write
        reg = data[pos + 1]
        val = data[pos + 2]
        if reg < 14:
            current_regs[reg] = val
        pos += 3
    elif cmd == 0x62: # Wait 1/50s
        frames.append(list(current_regs))
        pos += 1
    elif cmd == 0x66: # End of data
        break
    else:
        pos += 1

return frames

frames = vgm_to_ay_regs('music.vgm')

# Output as Z80 include
print(f"music_frames: equ {len(frames)}")
print("music_data:")
for regs in frames:
    print("    db " + ", ".join(f"${r:02X}" for r in regs))

```

5. **Z80 player.** The simplest possible player — 14 OUT instructions per interrupt:

```

ld hl, (music_ptr)
ld b, 14
xor a
.loop:
    out ($FD), a        ; select register
    ld c, (hl)
    inc hl
    push af
    ld a, c
    out ($BF), a        ; write value
    pop af
    inc a
    djnz .loop
    ld (music_ptr), hl
    ret

```

6. **Alternative: PT3 conversion.** If you prefer using a standard PT3 player (smaller, better compression), use the vgm2pt3 tool or compose directly in Vortex Tracker. Furnace's advantage is the modern interface; Vortex Tracker's advantage is guaranteed PT3 compatibility.

Recipe 5: packbench → Pre-compression Analysis

Goal: Analyse raw demo data before compression to identify optimal transforms.

Steps:

1. **Assemble without compression.** Build your demo with raw (uncompressed) INCBIN data.
2. **Run packbench analyse** on each data file:
3. **Read the report.** packbench reports:
 - **Entropy** (bits per byte) — theoretical minimum compressed size
 - **Byte distribution** — shows if values cluster (good) or spread uniformly (bad for compression)
 - **Run lengths** — shows repetition patterns
 - **Suggested transforms** — delta encoding, bit-plane separation, transpose, etc.
4. **Apply suggested transforms.** If packbench suggests delta encoding:

```
deltas = bytes([((data[i] - data[i-1]) & 0xFF) for i in range(1, len(data))])
open('sprites_delta.bin', 'wb').write(bytes([data[0]]) + deltas)
```

5. **Re-analyse.** Run packbench on the transformed data. Entropy should be lower.
6. **Compress.** Run ZX0 or Pletter on the transformed data. Compare compressed size against the original — the transform should yield a smaller result.
7. **Update demo.** The Z80 decompressor runs first (ZX0 decode), then the inverse transform (un-delta) reconstructs the original data. The inverse transform is cheap: ~10 T-states per byte for delta decoding.

Further Reading

- **GNU Rocket:** github.com/rocket/rocket — sync editor + client libraries
- **TiXL:** github.com/tixl3d/tixl — node-based motion graphics (MIT)
- **Farbrausch archive:** github.com/farbrausch/fr_public — Werkzeug, kkrunchy, V2 (BSD)
- **Furnace:** github.com/tildearrow/furnace — multi-system chiptune tracker (GPL-2.0)
- **Sointu:** github.com/vsariola/sointu — 4K intro synth (MIT)
- **WaveSabre:** github.com/logicomacorp/WaveSabre — 64K intro synth (MIT)
- **Crinkler:** github.com/runestubbe/Crinkler — compressing linker (zlib)
- **Bonzomatic:** github.com/Gargaj/Bonzomatic — live shader coding
- **Shader Minifier:** github.com/laurentlb/shader-minifier — GLSL/HLSL optimizer
- **z80-optimizer:** github.com/oisee/z80-optimizer — brute-force Z80 superoptimizer (MIT)
- **Motion Canvas:** motioncanvas.io — parametric animation (MIT)
- **Blender:** blender.org — 3D, VSE, Graph Editor, Geometry Nodes, Grease Pencil (GPL)

Appendix K: GPU Superoptimization — When Brute Force Beats Human Intuition

Dark wrote the fastest general multiply in 1997. In 2026, GPU brute-force proved what's optimal for each specific constant — and it's 10× faster.

K.1 The Idea

A superoptimizer tries every possible instruction sequence to find the provably shortest/fastest one. On a CPU this is impractical beyond 3–4 instructions. On a GPU, testing billions of candidates per second, it becomes feasible up to 15 instructions — enough to cover virtually all useful Z80 arithmetic.

The z80-optimizer project ran exhaustive searches on commodity GPUs (NVIDIA RTX 4060 Ti: ~1 billion candidates/second, AMD Radeon RX 580: ~400M/s, Apple M2: ~200M/s) and cross-verified every result across 5 platforms (CUDA, OpenCL, Vulkan, Metal, CPU reference).

K.2 Constant Multiplication

Dark's Method 1 (shift-and-add loop, Chapter 4): **196-204T** for any 8-bit constant. The GPU found specific-constant sequences that are dramatically faster.

254/254 constants solved — ALL DIRECT. Every u8 multiply has a proven optimal sequence found by direct GPU brute-force. No composition fallbacks needed — last 5 constants (×170, ×171, ×173, ×179, ×181) found at length-11.

Constant	GPU-Optimal Sequence	Insts	T-states	vs Dark's loop
×2	RLA	1	4T	50× faster
×3	LD B,A : ADC A,A : ADD A,B	3	12T	16×
×10	RLA : LD B,A : ADD A,B : ADD A,A : ADD A,B	5	20T	10×

Constant	GPU-Optimal Sequence	Insts	T-states	vs Dark's loop
×42	RLA : LD B,A : ... (8 insts)	8	32T	6×
×128	RRCA : SBC A,A : ADC A,B : RRA	4	16T	12×
×254	RLA : NEG	2	12T	16×
×255	NEG	1	8T	25×

Average: 8× faster than the general loop for constant K.

Key insight: only 14 of 21 candidate instructions appear in any optimal solution. Seven are strictly dominated: SLA A (= ADD A,A but 8T not 4T), SRA A, RLC A, RRC A (CB prefix = slower), OR A, SCF, EX AF,AF'. Removing them accelerated the search 38×

16-bit Multiply (u8 × K → HL)

254/254 complete using only **3 instructions**: ADD HL,HL + ADD HL,BC + LD C,A. This is a complete basis for 16-bit constant multiplication on Z80.

With SWAP_HL (= LD H,L / LD L,0) and SUB HL,BC: 86% code compression via prefix sharing.

×255 (16-bit): LD H,L : LD L,0 : LD C,A : OR A : SBC HL,BC = 30T (byte swap trick: ×256 − ×1).

K.3 Constant Division

Division by constant K on Z80 has no hardware support. The GPU found that the reciprocal multiplication trick produces optimal sequences automatically:

$$n / K = (n \times M) \gg S \text{ where } M \approx 2^S / K$$

247/247 divisors solved (3-255). COMPLETE — zero gaps. Every u8 division has a proven optimal sequence.

Divisor	Insts	T-states	vs general loop (280T)
/3	14	130T	2.2×
/5	14	127T	2.2×
/7	14	123T	2.3×
/9	11	97T	2.9×
/10	14	124T	2.3×
/15	12	102T	2.7×
/19	10	86T	3.3×
/25	10	83T	3.4×
/50	10	80T	3.5×
/57	6	49T	5.7×
/100	9	105T	2.7×
/114	6	46T	6.1×

Divisor	Insts	T-states	vs general loop (280T)
/171	4	27T	10.4×
/205	5	35T	8.0×
/255	6	50T	5.6×

Total: 247/247 divisors — COMPLETE. Average 107T. Fastest: div171 = 4 instructions, 27T (10.4× vs loop). div10 = 124T matches the famous Hacker’s Delight hand-optimized sequence — found automatically by GPU in 11 seconds. Last solved: div129 = 16 instructions, 160T.

K.4 Branchless Idioms

All found via GPU exhaustive search with a 37-op instruction pool:

Idiom	Sequence	Insts	T-states
bool(A)	LD B,A : NEG : ADC A,B	3	16T
NOT(A)	NEG : SBC A,A : INC A	3	16T
is_neg(A)	RLCA : SBC A,A : NEG	3	16T
lsb(A)	LD B,A : NEG : AND B	3	16T
complement(A)	CPL	1	4T
half(A)	RRA	1	4T
nibble_swap(A)	RLCA : RLCA : RLCA : RLCA	4	16T
double_sat(A)	RLCA : LD B,A : SBC A,A : OR B	4	16T
max_0(A)	LD B,A : RLCA : SBC A,A : XOR B : AND B	5	20T
sign(A)	(5 insts)	5	20T
ABS(A)	LD B,A : RLCA : SBC A,A : XOR B : SBC A,B : ADC A,B	6	24T

The SBC A,A carry-to-mask trick appears in many sequences: it sets A to \$00 or \$FF depending on the carry flag. Combined with RLCA (sign bit → carry), this enables instant sign detection without any conditional jumps.

```
LD    B, A
NEG
ADC   A, B    ; carry from NEG + original = 0 or 1
```

; ABS(A): absolute value of signed byte – 6 insts, 24T, branchless

```
LD    B, A
RLCA      ; sign bit → carry
SBC   A, A ; carry → mask (0x00 or 0xFF)
XOR   B    ; conditional complement
SBC   A, B ; adjust
```



```

    ADC  A, B      ; correct

; Sign-extend A → HL: 3 insts, 12T
    ADC  A, L      ; double – overflow if ≥128 sets carry
    SBC  A, A      ; carry → 0xFF mask
    LD   H, A      ; H = sign extension byte

```

K.5 16-bit Idioms

Idiom	Sequence	Insts	T-states
NEG HL (DE=0)	EX DE,HL : OR A : SBC HL,DE	3	23T
NEG HL (universal)	XOR A : SUB L : LD L,A : SBC A,A : SUB H : LD H,A	6	24T
Sign-extend A→HL	ADC A,L : SBC A,A : LD H,A	3	12T
NOT HL	DEC H : XOR H : LD L,A	3	12T
HL » 1	SRL H : RR L	2	16T
HL × 3	LD C,A : ADD HL,BC : ADD HL,BC	3	26T
HL × 10	(5 insts via shift-add)	5	48T
HL × 256	LD H,L : LD L,0	2	11T

NEG HL has 4 variants with different register prerequisites. Alf's universal method works without any prerequisites. The GPU found shorter versions requiring DE=0 or B=0.

K.6 The Packed Arithmetic Cassette

All optimal sequences share instruction prefixes. A packed library with multiple entry points via labels eliminates all redundancy:

```

mul52:  ADD A, A      ; ×52 entry point (×104 = ×52 × 2)
mul26:  ADD A, A      ; ×26 (×52 = ×26 × 2)
mul24:  ADD A, B      ; ×24
mul12:  ADD A, A      ; ×12
mul6:   LD  B, A      ; ×6
        ADD A, B
        ADD A, B
mul2:   RLA           ; ×2
        RET           ; shared – 7 constants, 9 instructions!

```

Rotation and shift sleds cover all amounts with one fall-through chain:

rot6:	RLCA	; 6	shr6:	SRL A	; 6
rot5:	RLCA	; 5	shr5:	SRL A	; 5
rot4:	RLCA	; nibble swap	shr4:	SRL A	; /16
rot3:	RLCA	; 3	shr3:	SRL A	; 3
rot2:	RLCA	; 2	shr2:	SRL A	; 2
rot1:	RLCA	; 1	shr1:	SRL A	; /2
	RET	; 9 bytes		RET	; 16 bytes

Packed library sizes: - 594 bytes for 164 mul8 constants (51% compression) - ~500 bytes for 254 mul16 constants (86% compression) - **Total: ~2KB packed blob covers ALL optimal arithmetic for Z80.** 254 multiplies + 247 divisions + rotation/shift sleds. For ZX Spectrum: just 4% of 48KB RAM.

K.7 How It Works

1. **Define the target:** for each input A (0-255), what should the output be?
2. **Generate** all instruction sequences up to length N
3. **QuickCheck:** test 4 carefully chosen inputs — rejects 99.99% instantly
4. **Full verify:** test all 256 inputs for survivors
5. **Pool reduction:** analyze which ops appear in solutions, remove the rest
6. **Guided search:** abstract chains predict structure, GPU searches focused space

The search runs on commodity GPUs:

- NVIDIA RTX 4060 Ti: ~1 billion candidates/second
- AMD Radeon RX 580: ~400 million/second (via OpenCL)
- Apple M2: ~200 million/second (via Metal)

Cross-verified across 5 platforms (CUDA + OpenCL + Vulkan + Metal + CPU). All results identical across 4 GPU APIs.

K.8 Source & Data

All sequences available at: <https://github.com/oisee/z80-optimizer>

- data/mulopt8_clobber.json — 164 multiply sequences with register annotations
- data/div8_optimal.json — 244 division sequences
- data/mul8_library.asm — packed Z80 assembly with multi-entry points

501 total provably optimal arithmetic sequences for Z80.

Generated by z80-optimizer v1.0.0

What's New

v20 (2026-02-28)

Russian translation updated to v20 + community feedback fixes.

Community feedback (Aki, Rombor via SinDiKat Slack)

Ch11 (Sound Architecture): AY history corrected - Investronica first integrated the AY into the Spanish Spectrum 128K, Sinclair adopted it, Amstrad inherited it for +2/+3. Chip variant clarified: the Spectrum 128K uses the AY-3-8912 (fewer I/O ports), not the 8910. "Without a dedicated sound CPU" reworded to "without dedicating the CPU to sound generation."

Ch12 (Music & Sync): "synthesiser" corrected to "PSG (Programmable Sound Generator)" for consistency with Ch11.

Ch02 illustration: Pixel row 109 → 107 in ch02_screen_layout.png (reported by Rombor). Address bits were correct, only the final row number was wrong: $64 + 5 \times 8 + 3 = 107$.

Full feedback log: `_in/tasks/feedback-2026-02-28.md`

Translation Memory tool + Russian edition

Built `translations/tm.py` - a paragraph-level Translation Memory system that aligns v0.6 EN with existing translations, diffs current EN at block level, and exports only the delta for LLM translation. The tool found that **87% of blocks were unchanged** since v0.6, so only 13% needed retranslation (~89% cost savings).

How it worked: 1. `tm.py build` segmented all 28 items at git tag v0.6, aligned EN blocks with RU/ES/UK translations positionally, stored in `tm.json` (3,976 segments) 2. `tm.py diff ru` classified each block as EQUAL (87%), MODIFIED (4%), or NEW (9%) 3. `tm.py export ru <chapter>` generated structured translation jobs containing only delta blocks, with previous RU as reference for MODIFIED blocks 4. Parallel Claude agents translated all delta blocks (~482 across 26 stale chapters + glossary + appendices) 5. `tm.py apply ru <chapter> <file>` reassembled complete RU chapters from TM (unchanged blocks) + agent output (new/modified blocks) 6. 6 new appendices (D-J) translated from scratch in parallel

Result: All 28 existing RU items updated from v0.6 to v20. 6 new appendices added. 34/34 OK, 0 stale, 0 missing.

New files: - `translations/tm.py` - Translation Memory tool (build/diff/export/apply/stats) - `translations/tm.json` - TM database (~3.5 MB, 3,976 segments) -

translations/README-tm.md - TM tool documentation

v19 (2026-02-28)

z80-optimizer coverage + mermaid rendering fix.

Chapter 23 (AI-Assisted): Superoptimiser sidebar — new section 23.5b covering z80-optimizer (oisee, 2025): - Brute-force Z80 superoptimisation: 602,008 provably correct rules from exhaustive length-2 enumeration - Explains methodology (full state equivalence, 406 opcodes × 406 opcodes, 34.7B comparisons) - Bridges peephole optimisation (MinZ, 23.5) with the “Z80 they still don’t know” discussion (23.6) - Positions exhaustive search as “the other AI” alongside neural-network-based code generation

Appendix J (Modern Tools): z80-optimizer entry in J.5 (PC Demoscene Toolchain): - Tool listing with usage context, example rules, and cross-reference to Ch.23

Build fix: Mermaid fenced code blocks now render correctly in PDF/EPUB. `build_book.py` converts ````mermaid id:xxx` to ````{.mermaid}` for pandoc compatibility.

v18 (2026-02-27)

Ped7g feedback patches + scene identity fixes.

Chapter 4 (Maths): Signed multiply section — fills a gap identified by Ped7g where Ch05 calls `mul_signed` 8 times but it was never implemented: - Two’s complement primer (bit 7 = sign, `$FF` = -1, NEG for abs) - Sign extension idiom: `r1a / sbc a,a` (2 bytes, 8T, branchless) - `mul_signed` (B,C signed → HL, ~240-260T) and `mul_signed_c` wrapper (A,C → HL) matching Ch05 calling conventions - Cost comparison table: unsigned vs signed vs square-table

Chapter 14 (Compression): Ped7g’s self-modifying RLE sidebar — complete 120-byte working mini-intro contributed with permission: - 9-byte core depacker using SP as data pointer and `db $18,0 → jr` self-modifying exit trick - Byte count analysis, interrupt safety note, advanced variants (`jp $C3C3` trick)

Appendix F: Z80N T-state audit — nearly all values were wrong (halved). Corrected against `wiki.specnext.dev`: - `PIXELDN/PIXELAD/SETAE`: 4→8T; `LDIX/LDDX`: 5→16T; `LDIRX/LDDR/LDPIRX`: 5→21/16T - `MIRROR/SWAPNIB`: 4→8T; `TEST`: 7→11T; barrel shifts: 4→8T - `PUSH nn`: 11→23T; `ADD HL/DE/BC,A`: 4→8T; `NEXTREG`: 12→20T / 8→17T; `OUTINB`: 5→16T - `MUL D,E` (8T) was the only correct one

Scene identity corrections across glossary, 5 chapters, and all 3 translations: - `sq ≠ psndcj`: `sq` = Screamer (Skrju), `psndcj` = cyberjack (TBK/Triebkraft + 4D) - “4th Dimension” → “4D+TBK (Triebkraft)” throughout - Added Ped7g and Screamer (`sq`) to glossary Key People tables

v17 (2026-02-27)

Appendix J (modern tools), Ch.20 refactor, Ped7g feedback fixes.

- New Appendix J: Modern Development Tools
- Ch.20 refactored with Unity/Unreal as data generators, Farbrausch sidebar
- Ped7g feedback: shadow register warning fix, pixel_addr fix, Quick Cost table fix, -syntax=abf, VSCode regex, RLE byte count correction

v16 (2026-02-26)

Pre-compression data preparation + packer benchmark tool.

Chapter 14 (Compression) expanded with new section “Pre-Compression Data Preparation”: - Shannon entropy as theoretical floor for compression — with Python code to compute it - Second derivative encoding: sine vs quadratic decay insight (visually identical curves, but quadratic has constant 2nd derivative → compresses to near-zero) - Transposition for tabular data: column-major reordering of interleaved XYZ vertex tables, AY register dumps, animation keyframes — with concrete entropy measurements (7.52 → 2.58 bits/byte) - Mask/pixel plane separation for sprites (10-20% ratio improvement) - Practical transforms table for 8 common demo data types - Key insight: transpose + delta + fast packer often beats raw + slow packer (better ratio AND faster decompression) - Two new exercises: transpose-and-measure, quadratic substitution

New tool: tools/packbench.py — packer benchmark and streaming decompression estimator: - bench mode: run real packer binaries on data files, measure compressed sizes (falls back to Introspec’s benchmark ratios when binaries unavailable) - budget mode: 128K memory map from TOML config — reserved areas, per-effect compressed storage, bank utilization - timeline mode: streaming decompression scheduler — models overlap/pause/streaming phases for seamless demo effect transitions; - -what-if compares all 9 packers per effect - analyze mode: pre-compression data analysis — entropy comparison across transforms (raw/delta/delta²/XOR), stride-based transposition testing, curve fitting (linear/quadratic R²), periodicity detection, actionable suggestions - All modes support --json for Clockwork integration - 9 packer profiles from Introspec’s benchmark + Ped7g’s upkr feedback - 4 platform presets (spectrum48, spectrum128, pentagon128, next)

New config: demo/packbench.toml — Antique Toy demo memory budget and timeline preset (5 effects: torus, plasma, dotscroll, rotozoomer, credits).

New Makefile targets: packbench, packbench-budget, packbench-timeline, packbench-analyze.

v12 (2026-02-24)

Comprehensive review pass. Full technical audit of all 23 chapters with systematic fixes:

12 HIGH severity fixes: - Ch.02: RET cc T-state values corrected (5/11T → 11/5T taken/not-taken) across 6 comments + timing table - Ch.04: Patrik Rak and Raxoft identified as same person throughout - Ch.07: SET n,(DE) / RES n,(DE) rewritten — instructions don’t exist on Z80, replaced with (HL) approach - Ch.10: False claim that LD (addr),SP doesn’t exist removed (ED 73, 20T) - Ch.11: AY mixer value corrected (\$28 → \$18 for noise C) - Ch.15: Contended memory table corrected — only \$4000-\$7FFF contended on 48K, not all RAM - Ch.16: XOR sprite

inner loop timing corrected (98T → 134T), JR NZ annotations fixed throughout - Ch.18: spawn_bullet bug fixed — D register zeroed before ADD IX,DE - Ch.19: AABB worst-case recounted (156T → ~270T), tile_at recounted (40T → ~182T) - Ch.23: DOWN_HL argument rewritten — was claiming correct code was buggy

19 MEDIUM severity fixes: LDIR formula standardised to $(N-1) \times 21 + 16$, PUSH operation order corrected, byte counts fixed, Agon 60Hz corrected, IM2 IY save added, invalid `ld hl,a` replaced, and more.

AI smell cleanup: Removed chain-of-thought leaks (“Wait – that is wrong”), cut philosophical padding sections, eliminated word-level tics (remarkably, deceptively simple, paradigm shift, borders on magic, etc.), trimmed defensive balancing and motivational pep talks across 9 chapters.

Two new tools: - `tools/autotag.py` — semi-automatic code block classifier and tagger (`-preview`, `-apply`, `-stats`) - `tools/audit_tstates.py` — T-state audit comparing inline annotations with computed values (`-scan-chapters`, `-asm-check`)

Code block pipeline: - 79 bare code blocks classified and tagged with language (z80/mermaid/text) - 279 code blocks tagged with `id:chMN_slug` identifiers - 271 .z80 + 8 .mmd canonical listings extracted to `listings/` - Final audit: 0 WRONG, 0 PARTIAL T-state annotations - All 29 assembly units compile with `sjasmpus`

v11 (2026-02-24)

- Art-top/Gogin attribution fix across affected chapters.
- MCC sidebar content added.
- ZXDN research integrated.

v10 (2026-02-24)

- External listings system (`tools/manage_listings.py`): `extract`, `inject`, `verify`, `stats` commands.
- Unified versioning via `version.json` + `build_book.py`.
- UNDER CONSTRUCTION banner for work-in-progress chapters.

v0.8 (2026-02-24)

Chapters expanded. Three thin chapters significantly deepened:

- **Chapter 7 (Rotozoomer):** Three rotozoomer variants compared – monochrome bitmap (full pixel), chunky (2x2/4x4), and attribute-based (32x24 colour grid). New sections on fixed-point stepping mechanics, self-modifying code at the byte level, texture design, and boundary handling. Credits corrected: GOA4K and Refresh are by Exploder^XTM.
- **Chapter 10 (Dotfield Scroller):** POP-trick deep dive with interrupt risks and SP-save patterns. Lissajous, helix, and multi-wave trajectory formulas. 4-phase colour cycling code with attribute byte manipulation. Demoscene lineage from Born Dead to Introspec's Eager.
- **Chapter 13 (Size-Coding):** Now covers the full range from 256 bytes to 4K intros. New size-coder's toolkit (register assumptions, DJNZ tricks, RST

as 1-byte CALL, overlapping instructions). 4K intro section with GOA4K (Exploder^XTM), inal and Megademica (SerzhSoft). AY bytebeat and procedural graphics compo sections added.

Five new appendices (all 9 now complete, A-I):

- **E: eZ80 Quick Reference** – ADL mode, MLT/LEA/PEA/TST, Agon Light 2 specifics.
- **F: Z80 Variants** – Z80N (Next) instructions organised by the pain they solve, R800 (MSX turboR) pipeline and MULUB/MULUW, Soviet clones, comparison table.
- **H: Storage APIs** – TR-DOS (Beta Disk 128) and esxDOS (DivMMC) with complete port maps, ROM API, and working code examples.
- **I: Bytebeat & AY-Beat** – Classic PCM bytebeat adapted for the AY-3-8910. Formula cookbook (12 tone + 4 volume formulas), envelope drone, music theory for algorithms (scale tables, arpeggios, ornaments, chord progressions, L-system grammars, PRNG with curated seeds).

New assembly example: aybeat.a80 – 320-byte AY-beat engine demonstrating pentatonic scales, I-IV-V-I chord progression, arpeggio, step ornaments, envelope drone, and noise percussion. 19 bytes of musical data, three channels of generative sound.

9 new diagrams added across chapters 3, 8, 9, 14, 19, 20, 21, 22, 23. All 23 figures tagged with `<!-- figure: chNN_name -->` for greppability and translation tracking.

29 compilable examples, all passing. ~180K words (English).

v0.7 (2026-02-24)

- Appendix D: Development environment setup (sjasmpplus, VS Code, DeZog, emulators).
- Three new chapter examples: ch21 game skeleton (128K, 438 lines), ch22 Agon entity system (240 lines), ch23 AI-assisted diagonal fill (149 lines).
- README rewritten with honest platform positioning (“this book lives on the ZX Spectrum”), per-chapter platform tags, TL;DR with direct download link.
- 28 compilable examples, all passing.

v0.6 (2026-02-23)

- All 23 chapters drafted. All 4 language editions (EN, ES, RU, UK) built and released.
- Translation manifest system for staleness tracking.
- ~128K words (English), ~575K words total across all languages.